
Gradient-Free Approaches is a Key to an Efficient Interaction with Markovian Stochasticity¹

Anonymous Author(s)²

Affiliation
Address
email

Abstract³

¹ This paper deals with stochastic optimization problems involving Markovian ⁴
² noise with a zero-order oracle. We present and analyze a novel derivative-
³ free method for solving such problems in strongly convex smooth and
⁴ non-smooth settings with both one-point and two-point feedback oracles.
⁵ Using a randomized batching scheme, we show that when mixing time τ of
⁶ the underlying noise sequence is less than the dimension of the problem d , the
⁷ convergence estimates of our method do not depend on τ . This observation
⁸ provides an efficient way to interact with Markovian stochasticity: instead
⁹ of invoking the expensive first-order oracle, one should use the zero-order
¹⁰ oracle. Finally, we complement our upper bounds with the corresponding
¹¹ lower bounds. This confirms the optimality of our results.

¹² 1 Introduction⁵

¹³ Stochasticity is a fundamental aspect of many optimization problems, naturally arising ⁶
¹⁴ in the field of machine learning [48, 28]. Stochastic gradient descent (SGD) [45] and its
¹⁵ accelerated variants [39, 25] have become a de facto optimizers for modern large models
¹⁶ training. Theoretical properties of SGD have been extensively studied under various statistical
¹⁷ frameworks [36, 24, 10, 56], often relying on the assumption that noise is independent
¹⁸ and identically distributed (i.i.d.). However, in many real-world applications — including
¹⁹ reinforcement learning (RL) [6, 16], distributed optimization [35, 31], and bandit problems
²⁰ [3] — noise is not i.i.d., instead exhibiting correlations or *Markovian structure*.

²¹ For instance, in the mentioned growing field of RL, sequential interactions with the environ- ⁷
²² ment induce state-dependent structure of the noise, creating a need for non-i.i.d. noise aware
²³ algorithms. Although several gradient-based methods for Markovian stochastic oracles have
²⁴ been studied in the past decade [14, 18], policy optimization in RL is based solely on reward
²⁵ feedback, making traditional methods inapplicable, since there is no access to first-order
²⁶ information [46, 9, 19]. *Zero-order optimization* (ZOO) methods are specifically developed
²⁷ to address such problems, and are used in scenarios where gradients are unavailable or
²⁸ prohibitively expensive to compute. Apart from RL, ZOO techniques are widely employed in
²⁹ adversarial attack generation [8], hyperparameter tuning [47, 57], continuous bandits [7, 49]
³⁰ and other applications [54, 33]. While the literature on ZOO is extensive, this work is, to
³¹ our knowledge, *the first study of optimization problem with both zero-order information and*
³² *Markovian noise*, aimed at developing an optimal algorithm for a large family of problems
³³ from the intersection of these two areas.

34 1.1 Related works

35 $\diamond$ **Zero-order** methods is one of the key and oldest areas of optimization. There are various 1
 36 zero-order approaches, here we can briefly highlight, e.g., one-dimensional methods [32, 42]
 37 or their high-dimensional analogues [41], ellipsoid algorithms [58] and searches along random
 38 directions [4]. Currently, the most popular and most studied mechanism behind ZOO
 39 methods is the finite-difference approximation of the gradient described in [43, 20, 40]. The
 40 idea is simple: querying two sufficiently close points is essentially equivalent to finding a
 41 value of the directional derivative of the function:

$$\langle \nabla f(x), e \rangle \approx \frac{f(x + te) - f(x)}{t} \approx \frac{f(x + te) - f(x - te)}{2t}, \quad 2 \quad (1)$$

42 where e is a random direction. It can be a random coordinate, a vector from the Euclidean 3
 43 sphere or a sample of the Gaussian distribution. The approximation (1) in turn leads back to
 44 the gradient methods or coordinate algorithms of Nesterov [38]. There are, however, several
 45 differences:

- 46 • First, to get full gradient information, the algorithm would need d queries instead of one 4
 47 gradient oracle call (here d is the dimension of x).
- 48 • Second, if the ZO oracle is inexact, i.e. only noisy values of function are available, then
 49 finite difference schemes can fail if noise components do not cancel out.

50 The setting of the second point, when function evaluations experience zero-mean additive 5
 51 perturbations, is called *Stochastic ZOO*. The stochasticity, as noted before, is abundant in
 52 the modern optimization world. To tackle this issue, additional assumptions about the noise
 53 structure are required. Here we briefly discuss two main ideas adopted in the literature, and
 54 refer the reader to Section 2 for precise definitions.

55 In the case of *two-point feedback*, we assume that for a fixed value of the noise variable one 6
 56 can call the stochastic zero-order oracle at least twice. It means that we can compute the
 57 finite difference approximation of the following form:

$$p(x, \xi, e) = \frac{f(x + te, \xi) - f(x - te, \xi)}{2t} \approx \langle \nabla_x f(x, \xi), e \rangle \quad 7 \quad (2)$$

58 Such approximation produces an estimate for the directional derivative of a noisy realization 8
 59 $f(\cdot, \xi)$ of the function f . As mentioned before, the approximation (2) can be used instead of
 60 the (stochastic) gradient in first-order methods. In the case of independent randomness, a
 61 large number of works are based on this idea. There are results for both non-smooth and
 62 smooth convex problems built on classical and accelerated gradient methods of Nesterov and
 63 Spokoiny [40]. In the scope of our paper, we are interested in the results for smooth strongly
 64 convex problems from [17], namely estimates on zero-order oracle calls to achieve ε -solution
 65 in terms of $\|x - x^*\|$: $\mathcal{O}(\frac{d\sigma_2^2}{\mu^2\varepsilon})$. Here σ_2 is introduced as the variance of the gradient, i.e. it is
 66 assumed that $\mathbb{E}_\xi \nabla f(x, \xi) = \nabla f(x)$ and $\mathbb{E}_\xi \|\nabla f(x, \xi) - \nabla f(x)\|^2 \leq \sigma_2^2$. The main limitation
 67 of two-point approach is that several evaluations with the same noise variable are required,
 68 which is well suited for problems like empirical risk optimization [34], but can be a major
 69 barrier for RL or online optimization.

70 In the *one-point feedback* setting, a more general stochasticity is assumed. In this case, each 9
 71 call to the zero-order oracle generates a new randomness. Now the approximation (1) looks
 72 as follows

$$p(x, \xi^\pm, e) = \frac{f(x + te, \xi^+) - f(x - te, \xi^-)}{2t} \quad 10 \quad (3)$$

73 Using different ξ^+ and ξ^- in (3) renders any conditions on the properties of $\nabla f(\cdot, \xi)$ useless. 11
 74 Instead, it is assumed that $\mathbb{E}_\xi f(x, \xi) = f(x)$ and $\mathbb{E}_\xi |f(x, \xi) - f(x)|^2 \leq \sigma_1^2$. With one-point
 75 feedback, the major problem is choosing the right shift t for the finite difference scheme.
 76 Picking it too small results in an amplification of the additive noise, and taking t too big
 77 leads to a poor gradient estimate. Because of this variance trade-off, the optimal rate for
 78 methods with one-point approximation is worse than for two-point feedback. In particular,
 79 for smooth strongly convex problems we have the following estimate on zero-order oracle
 80 calls [23]: $\mathcal{O}(\frac{d^2\sigma_1^2}{\mu^3\varepsilon^{\frac{1}{2}}})$.

81 Although zero-order gradient approximation schemes suffer from high variance, there is a 1
 82 surprising property that makes them superior in *non-smooth* optimization [22, 44, 49]. The
 83 idea goes back to the 70s and utilizes the fact that

$$\mathbb{E}[e \cdot p(x, \xi^{(\pm)}, e)] = \frac{1}{d} \nabla f_t(x), \text{ where } f_t \text{ is a smoothed function, defined as } 2$$

$$f_t(x) = \mathbb{E}_r[f(x + tr)] \text{ with } r \sim RB_2^d$$

84 In fact, it can be shown that f_t is $\frac{\sqrt{d}G}{t}$ -smooth if f is G -Lipschitz. This makes zero-order 4
 85 approximation a suitable candidate for a stochastic gradient of f_t . Optimizing this function
 86 with a first-order method produces some solution, but it may not be the optima of f [22].
 87 From this point, there is a game – for small t the functions f and f_t are closer and for big t
 88 the function f_t is easier to optimize as it gets smoother.


89 In more recent works, there have been many improvements in theoretical understanding of 5
 90 ZO methods. The authors consider higher-order smoothness of the underlying function [2],
 91 tackle non-convex non-smooth problems [44], take arbitrary Bregman geometry to benefit
 92 in terms of oracle complexity [49, 29], and come up with sharp information-theoretic lower
 93 bounds to understand computational limits [15, 1]. But none of them consider Markovian
 94 stochasticity.

95 $\diamond$ **Markovian first-order methods.** While the literature on stochastic optimization with 6
 96 i.i.d. noise is extensive, research addressing the Markovian setting remains relatively sparse.
 97 In our paper, we focus on the most "friendly" type of uniformly geometrically ergodic Markov
 98 chains (see Section 2 for precise definitions).

99 Duchi et al. [14] conducted pioneering work on non-i.i.d. noise, investigating the Ergodic 7
 100 Mirror Descent algorithm and establishing optimal convergence rates for non-smooth convex
 101 problems. For smooth problems there were different attempts to get record-breaking estimates
 102 on the first-order oracle [12, 11, 59, 18]. Finally, the optimal results were obtained for both
 103 convex and non-convex problems in the works of Beznosikov et al. [5], Solodkin et al. [52]. In
 104 particular, for smooth strongly convex objectives under Markovian noise the authors give the
 105 complexity of the form: $\mathcal{O}(\frac{\tau \sigma_2^2}{\mu^2 \epsilon})$, where τ is defined as the mixing time of the corresponding
 106 Markov chain (see Section 2). Note that these works utilize Multilevel Monte Carlo (MLMC)
 107 batching technique, which helps to effectively interact with Markovian noise. We will need
 108 this approach as well. Note that it was first considered in Markovian gradient optimization
 109 by Dorfman and Levy [13] for automatic adaptation to unknown τ .

110 $\diamond$ **Hypothesis.** The complexity estimate for strongly convex first-order stochastic meth- 8
 111 ods is $\mathcal{O}(\frac{\sigma_2^2}{\mu^2 \epsilon})$ [36, 37]. Lower bounds for the same class of problems and methods
 112 show that the result is unimprovable [58]. As mentioned before, the transition from
 113 i.i.d. stochasticity to Markovian stochasticity increases the estimate by τ times. This
 114 result is also optimal as shown by Beznosikov et al. [5]. At the same time, going
 115 from gradient oracle to zero-order methods adds a multiplier d in the two-point feed-
 116 back and d^2/ϵ in the one-point case. And this estimate is unimprovable as well [1, 15].

117 The hypothesis arises that the transition to zero- 9
 118 order Markov optimization adds two multipliers
 119 at once: $d\tau$ and $d^2\tau/\epsilon$ for two- and one-point. It is
 120 illustrated in the following diagram for two-point
 121 feedback:



122 1.2 Our contribution 11

123 Our main contribution is the answer to the hy-
 124 pothesis above: *surprisingly, it is not true.* In more detail:

125 $\diamond$ **Accelerated SGD.** We present the first analysis of Zero-Order Accelerated SGD under 13
 126 Markovian noise, considering both two-point and one-point feedback. Contrary to the
 127 expected multiplicative scaling of convergence rates with both dimensionality and mixing
 128 time, our analysis reveals a significant acceleration, as presented in Figure 1. It turns out
 129 that if τ is smaller than d , our results do not differ at all from the gradient-free methods

with independent stochasticity. The key technique behind this acceleration is described in Section 2.1. The theory is also numerically validated in Section 3.

◊ **Non-smooth problems.** We also consider non-smooth problems with Markovian noise. Using the smoothing technique we come up with a corresponding upper bounds in this case, as shown in Figure 1. The details of these bounds are presented in Appendix A.

Figure 1: Summary of upper bounds. For notation, see Table 1

	Smooth		Non-smooth	
	IID	Markov.	IID	Markov.
FO	$\frac{\sigma_2^2}{\mu^2 \varepsilon}$ [45]	$\tau \frac{\sigma_2^2}{\mu^2 \varepsilon}$ [5]	$\frac{G^2}{\mu^2 \varepsilon}$ [50]	$\tau \frac{G^2}{\mu^2 \varepsilon}$ [14] ¹
ZO 2P	$d \frac{\sigma_2^2}{\mu^2 \varepsilon}$ [30]	$(d + \tau) \frac{\sigma_2^2}{\mu^2 \varepsilon}$	$d \frac{G^2}{\mu^2 \varepsilon}$ [22]	$(d + \tau) \frac{G^2}{\mu^2 \varepsilon}$
ZO 1P	$d^2 \frac{\sigma_1^2}{\mu^3 \varepsilon^2}$ [2] ²	$d(d + \tau) \frac{L \sigma_1^2}{\mu^3 \varepsilon^2}$	$d^2 \frac{\sigma_1^2 G^2}{\mu^4 \varepsilon^3}$ [23]	$d(d + \tau) \frac{\sigma_1^2 G^2}{\mu^4 \varepsilon^3}$

◊ **Computational efficiency.** First, as noted above, our method gives the same oracle complexity for any $\tau \leq d$. Moreover, if we assume that calling a zero-order oracle is d times cheaper than computing the corresponding gradient, then the gradient method with Markov noise will require resources proportionally to $d \cdot \tau$ — the cost of one oracle call is d and the complexity scales as τ for the first-order method from Figure 1. At the same time, the resource complexity of our zero-order method is proportional to $d + \tau$.

◊ **Lower bounds.** In Section 2.3 we establish the first information-theoretic lower bounds for solving Markovian optimization problems with one-point and two-point feedback. Our results match the convergence guarantee of our algorithm up to logarithmic factors, showing that the analysis is accurate and no further improvement is possible.

Table 1: Notations & Definitions

Sym.	Definition	Sym.	Definition
$\ \cdot\ , \langle \cdot, \cdot \rangle$	Norm, dot product, assumed Euclidean by default	ε	$\ x - x^*\ ^2$
$\mathcal{Z}, \mathcal{Z}$	Complete separable metric space, its Borel σ -algebra	d	Problem dimension
$\mathcal{Q}$	Markov kernel on $\mathcal{Z} \times \mathcal{Z}$	L	Gradient's Lipschitz constant
$\mathbb{P}_\xi, \mathbb{E}_\xi$	Probability, Expectation under initial distribution ξ^3	μ	Strong convexity constant
$\{Z_k\}$	Canonical process with kernel $\mathcal{Q}$	G	Function's Lipschitz constant
RB_2^d, RS_2^d	Uniform distribution on unit ℓ_2 -ball, -sphere	σ_1^2	$ F(x, Z) - f(x) ^2 \leq \sigma_1^2$
e	Random direction, $e \sim RS_2^d$	σ_2^2	$\ \nabla F(x, Z) - \nabla f(x)\ ^2 \leq \sigma_2^2$
$a_n \lesssim b_n$	$\exists c \in \mathbb{R}$ (problem-independent): $a_n \leq cb_n$ for all n	τ	Mixing time of Z
$a_n \simeq b_n$	$a_n \lesssim b_n$ and $b_n \lesssim a_n$	$g, \hat{g}$	Gradient estimators
$T = \tilde{\mathcal{O}}(S)$	$T \leq \text{poly}(\log S) \cdot S$ as $\varepsilon \rightarrow 0$	$f_t(x)$	$\mathbb{E}_r[f(x + tr)], r \sim RB_2^d$

2 Main results

We are now ready for a more formal presentation. In this paper, we study the minimization problem

$$\min_{x \in \mathbb{R}^d} f(x) := \mathbb{E}_{Z \sim \pi} [F(x, Z)], \quad (4)$$

where π is an unknown distribution and access to the function f (not to its gradient ∇f) is available through a stochastic one-point or two-point oracle $F(x, Z)$.

In our analysis, we will use a set of assumptions on the underlying function f and its oracle, starting with smoothness and convexity:

Assumption 1. *The function f is L -smooth on $\mathbb{R}^d$ with $L > 0$, i.e., it is differentiable and there is a constant $L > 0$ such that the following inequality holds for all $x, y \in \mathbb{R}^d$:*

$$\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|. \quad (5)$$

In the two-point feedback setting, we require the following generalization:

Assumption 1'. *For all $Z \in \mathcal{Z}$ the function $F(\cdot, Z)$ is L -smooth on $\mathbb{R}^d$.*

Note that the uniform 1' implies

¹The authors consider general convex case. Using standard restart technique, we get the corresponding bound in the strongly convex case.

²The noise is assumed to be point-independent.

³By construction, for any $A \in \mathcal{Z}$, we have $\mathbb{P}_\xi(Z_k \in A \mid Z_{k-1}) = \mathcal{Q}(Z_{k-1}, A)$, $\mathbb{P}_\xi$ -a.s.

162 **Assumption 2.** The function f is μ -strongly convex on $\mathbb{R}^d$, i.e., it is continuously dif- 1
 163 ferentiable and there is a constant $\mu > 0$ such that the following inequality holds for all
 164 $x, y \in \mathbb{R}^d$:

$$\frac{\mu}{2} \|x - y\|^2 \leq f(x) - f(y) - \langle \nabla f(y), x - y \rangle. \quad 2 \quad (5)$$

165 We now turn to assumptions on the sequence of noise states $\{Z_i\}_{i=0}^\infty$. Specifically, we 3
 166 consider the case where $\{Z_i\}_{i=0}^\infty$ forms a time-homogeneous Markov chain. Let Q denote the
 167 corresponding Markov kernel. We impose the following assumption on Q to characterize its
 168 mixing properties:

169 **Assumption 3.** $\{Z_i\}_{i=0}^\infty$ is a stationary Markov chain on $(Z, \mathcal{Z})$ with Markov kernel Q and 4
 170 unique invariant distribution π . Moreover, Q is uniformly geometrically ergodic with mixing
 171 time $\tau \in \mathbb{N}$, i.e., for every $k \in \mathbb{N}$,

$$\Delta(Q^k) = \sup_{z, z' \in Z} (1/2) \|Q^k(z, \cdot) - Q^k(z', \cdot)\|_{\text{TV}} \leq (1/4)^{\lfloor k/\tau \rfloor}. \quad 5 \quad (6)$$

172 Assumption 3 is common in the literature on Markovian stochasticity [14, 12, 13, 5, 52]. It 6
 173 includes, for instance, irreducible aperiodic finite Markov chains [18]. The mixing time τ
 174 reflects how quickly the distribution of the chain approaches stationarity, providing a natural
 175 measure of the temporal dependence in the data.

176 Next, we specify our assumptions on the oracle. As discussed in Section 1.1, these assumptions 7
 177 differ based on the type of feedback.

178 **Assumption 4** (for one-point). For all $x \in \mathbb{R}^d$ it holds that $\mathbb{E}_\pi[F(x, Z)] = f(x)$. Moreover, 8
 179 for all $Z \in \mathcal{Z}$ and $x \in \mathbb{R}^d$ it holds that

$$|F(x, Z) - f(x)|^2 \leq \sigma_1^2, \quad 9$$

180 **Assumption 4'** (for two-point). For all $x \in \mathbb{R}^d$ it holds that $\mathbb{E}_\pi[\nabla F(x, Z)] = \nabla f(x)$. 10
 181 Moreover, for all $Z \in \mathcal{Z}$ and $x \in \mathbb{R}^d$ it holds that

$$\|\nabla F(x, Z) - \nabla f(x)\|^2 \leq \sigma_2^2. \quad 11$$

182 Recent works on stochastic ZOO methods have considered milder assumptions, such as 12
 183 bounded variance (see Section 1.1). However, the uniform boundedness assumed in Assump-
 184 tions 4 and 4', is standard in analyses under Markovian noise [14, 12, 13, 5, 52]. These
 185 assumptions can be relaxed under stronger conditions, e.g., uniform convexity and smoothness
 186 of $F(\cdot, Z)$ [18].

187 Assumptions 3 and 4 allow us to reduce the variance of the noise via batching, similarly the 13
 188 to i.i.d. setting. This is captured in the following technical lemma:

189 **Lemma 1.** Let Assumptions 3 and 4(4') hold. Then for any $n \geq 1$ and $x \in \mathbb{R}^d$ and any 14
 190 initial distribution ξ on $(Z, \mathcal{Z})$, we have

$$\mathbb{E}_\xi \left[\frac{1}{n} \sum_{i=1}^n F(x, Z_i) - f(x) \right]^2 \lesssim \frac{\tau}{n} \sigma_1^2, \quad \mathbb{E}_\xi \left\| \frac{1}{n} \sum_{i=1}^n \nabla F(x, Z_i) - \nabla f(x) \right\|^2 \lesssim \frac{\tau}{n} \sigma_2^2. \quad 15$$

191 2.1 Batching technique 16

192 In this section, we describe the main tools used to establish the $(d + \tau)$ -type scaling of the 17
 193 error rate. We will focus on reducing the variance and bias of gradient estimators using a
 194 specialized batching approach.

195 We begin by fixing a common building block of our gradient estimators at a point x for both 18
 196 one-point and two-point feedback, as introduced in Section 1.1:

$$\hat{g}(x, Z^{(\pm)}, e) = d \cdot p(x, Z^{(\pm)}, e) \cdot e = e \cdot \begin{cases} d \frac{F(x + te, Z^+) - F(x - te, Z^-)}{2t} & \text{(one-point),} \\ d \frac{F(x + te, Z) - F(x - te, Z)}{2t} & \text{(two-point).} \end{cases} \quad 19$$

197 These estimators exhibit a twofold randomness that affects how rapidly they concentrate 20
 198 around the true gradient, as we will discuss below.

For clarity, we focus our discussion on the one-point case, although our conclusions extend to the two-point case as well.

A widely used variance reduction technique is *mini-batching*, where one computes $F(x, Z_i)$ over a batch of noise variables $\{Z_i\}_{i=1}^n$. The mini-batch gradient estimator is given by:

$$\hat{g}_{mb}(x) = \frac{1}{n} \sum_{i=1}^n \hat{g}(x, Z_i^\pm, e) = e \cdot d \left(\frac{1}{n} \sum_{i=1}^n p(x, Z_i^\pm, e) \right). \quad 3$$

Let us estimate the scaling of its variance $\mathbb{E}_e \mathbb{E}_Z \|\hat{g}_{mb} - \nabla f\|^2$ with the noise level σ_1^2 . As $E_Z \hat{g}_{mb} \approx d \frac{f(x+te) - f(x-te)}{2t} \approx d \langle \nabla f, e \rangle$ we would like to estimate the following for any fixed direction e :

$$\mathbb{E}_Z [d \cdot p_{mb}(x) - d \langle \nabla f, e \rangle]^2 \approx \frac{d^2}{t^2} \mathbb{E}_Z \left[\frac{1}{n} \sum_{i=1}^n F(x+te, Z_i^+) - f(x+te) \right]^2 \stackrel{(1)}{\approx} \frac{d^2 \tau \sigma_1^2}{n t^2}. \quad 5 \quad (7)$$

With that, we bound the variance:

$$\mathbb{E}_e \mathbb{E}_Z \|\hat{g}_{mb} - \nabla f\|^2 \gtrsim \mathbb{E}_e \mathbb{E}_Z \|\hat{g}_{mb} - \mathbb{E}_Z \hat{g}_{mb}\|^2 \approx \mathbb{E}_e \mathbb{E}_Z \|\hat{g}_{mb} - d \langle \nabla f, e \rangle\|^2 \stackrel{(7)}{\approx} \frac{d^2 \tau \sigma_1^2}{n t^2}. \quad 7 \quad (8)$$

Can the mini-batching scheme be improved?

This subsection explores an unexpected source of improvement that contradicts our initial hypothesis. Specifically, we identify an inefficiency in the current use of samples Z_i , which becomes evident from two perspectives. Equation (8) shows the variance scales as $\frac{\tau}{n}$. If we could reduce τ by a factor of k , we would need k -times fewer samples to maintain the same variance. This leads us to the idea of sparsified sampling. We partition the Markov noise chain $\{Z_i\}$ into k subchains $\{Z_{k \cdot i + r}\}$ for $r = 0 \dots k-1$. This corresponds to a mixing time of $\lceil \frac{\tau}{k} \rceil$ for each subchain (see (3)), effectively reducing temporal correlation - a natural consequence of sampling every k -th element of the original chain. Thus, sampling from any single subchain could yield a $\min(k, \tau)$ -fold reduction in the number of samples needed (although such procedure would still require all intermediate oracle calls, yielding no computational speedup).

For a concrete illustration of that inefficiency, consider a lazy Markov chain that remains in the same state for (an average of) τ steps before transitioning uniformly at random. In such a case, all oracle queries $F(x, Z)$ for a fixed x return the same value for τ consecutive steps. Therefore, retaining only every τ -th estimate $\hat{g}$ would yield a mini-batch of equivalent quality.

In summary, we observe that the mini-batching scheme could, in principle, operate just as effectively by retaining only every k -th sample and discarding the rest. This might suggest that better utilization of the samples is possible. First order methods, nevertheless, are unable to exploit this redundancy (as shown by [5]’s lower bound) and are effectively forced to wait out the τ -step mixing window. In contrast, we can exploit this structure by querying finite differences along different directions to estimate the gradient better. Specifically, we construct d subchains, and use the sample from the r -th subchain $Z_{d \cdot i + r}$ to estimate r -th partial derivative $\frac{F(x+te_r, Z) - F(x-te_r, Z)}{2t}$, effectively restoring the full gradient coordinate-wise.

Let us estimate the resulting variance reduction. First, we achieve a d -fold reduction by reconstructing all d gradient coordinates. Second, each coordinate now operates on a chain with mixing time $\lceil \frac{\tau}{d} \rceil$, yielding an additional factor of $\min(d, \tau)$. However, because batches are now split across d coordinates, each batch is d times smaller than before, introducing a factor of d loss. The net variance reduction is therefore $\min(d, \tau)$, and the final scaling becomes $d \cdot \frac{d\tau}{\min(d, \tau)} = d \cdot \max(d, \tau) \simeq d(d + \tau)$.

Random directions

This insight can be extended to a simpler yet equally effective method. Instead of assigning directions deterministically, we associate each sample with a random direction $e \in RS_2^d$, forming the estimator:

$$\hat{g}_{rd}[n](x, Z, e) = \frac{1}{n} \sum_{i=1}^n \hat{g}(x, Z_i, e_i). \quad 15$$

While the above discussion was intuitive, we now outline a more formal approach (see Lemma 5 for details). As lazy Markov chain is effectively equivalent to stochastic i.i.d. τ -point feedback setting, we follow Corollary 2 of [15], who decompose the total variance into two terms:

$$\mathbb{E}\|\hat{g}_{rd} - \nabla f(x)\|^2 \leq 2\mathbb{E}\|\hat{g}_{rd} - \mathbb{E}_e \hat{g}_{rd}\|^2 + 2\mathbb{E}\|\mathbb{E}_e \hat{g}_{rd} - \nabla f(x)\|^2.$$

Each of the two terms individually eliminates one factor from the $d^2\tau$ dependence.

The first term:

$$\begin{aligned} \mathbb{E}\|\hat{g}_{rd} - \mathbb{E}_e \hat{g}_{rd}\|^2 &= \mathbb{E}_Z \mathbb{E}_e \left\| \frac{1}{n} \sum_{i=1}^n \underbrace{[\hat{g}(x, Z_i, e_i) - \mathbb{E}_{e_i} \hat{g}(x, Z_i, e_i)]}_{\mathbb{E}_e[\cdot]=0, \text{ independent w.r.t. } e} \right\|^2 \\ &= \frac{1}{n^2} \sum_{i=1}^n \mathbb{E}\|\hat{g}(x, Z_i, e_i) - \mathbb{E}_{e_i} \hat{g}(x, Z_i, e_i)\|^2 \end{aligned}$$

is independent of τ since Assumption 4 bounds each term directly.

For the second term, we observe that $\mathbb{E}_e \hat{g}_{rd} = \mathbb{E}_e \hat{g}_{mb}$, and thus the bound involves $\mathbb{E}\|\mathbb{E}_e \hat{g}_{mb} - \nabla f(x)\|^2$. This is crucially different from the $d^2\tau$ dependence that appeared in the mini-batch case, when we considered $\mathbb{E}\|\hat{g}_{mb} - \nabla f(x)\|^2$. Intuitively, the expectation over directions helps recover the full gradient rather than a directional component, thereby reducing variance with respect to d .

Multilevel Monte Carlo

The estimator $\hat{g}_{rd}$ is not our final construction. While it controls variance, the temporal correlation in noise may introduce significant bias. A well-established approach to mitigating this is MLMC, widely used in the statistical literature [27, 26], and more recently in gradient optimization [13, 5]. Here is our interpretation.

With parameters J, l, M, B from Table 2, $\{Z_i\}$ - $2^J l$ samples from Z and $\{e_i\}$ - random directions we introduce MLMC estimator:

$$\hat{g}_{ml}(x) = \hat{g}_{rd}[l](x) + \begin{cases} 2^J [\hat{g}_{rd}[2^J l](x) - \hat{g}_{rd}[2^{J-1} l](x)], & \text{if } 2^J \leq M \\ 0, & \text{otherwise} \end{cases}$$

$\hat{g}_{ml}$ is our final gradient estimator, with the following guarantees:

Lemma 2 (for one-point). *Let Assumptions 1, 3 and 4 hold. For any initial distribution¹ ξ on $(Z, \mathcal{Z})$ the gradient estimates $\hat{g}_{ml}$ satisfy $\mathbb{E}[\hat{g}_{ml}] = \mathbb{E}[\hat{g}_{rd}[2^{\lceil \log_2 M \rceil} l]]$. Moreover,*

$$\begin{aligned} \mathbb{E}\|\nabla f_t(x) - \hat{g}_{ml}(x)\|^2 &\lesssim \frac{d\|\nabla f(x)\|^2}{B} + \frac{d^2 L^2 t^2}{B} + \frac{d(d+\tau)\sigma_1^2}{Bt^2}, \\ \|\nabla f_t(x) - \mathbb{E}[\hat{g}_{ml}(x)]\|^2 &\lesssim \frac{d\tau\sigma_1^2}{t^2 BM}. \end{aligned}$$

One can note that although $\hat{g}_{ml}$ requires, on average, $\mathbb{E}[2^J l B] = \log_2^2 M \cdot B$ oracle calls, the variance is only reduced by a factor of B . In contrast, the bias is reduced significantly - by a factor of BM .

2.2 Algorithm

We now present the full version of Algorithm 1, which incorporates the gradient estimators discussed in the previous section and uses a slightly modified variant of Nesterov's Accelerated Gradient Descent at its core.

While technically we prove four separate upper bounds covering both one- and two-point feedback under smooth and non-smooth assumptions, they follow the same scheme which we will illustrate in the one-point smooth case.

¹Note that $\hat{g}_{ml}$ (specifically Z_1) indirectly depends on the chain's initial distribution. As our algorithm is going to repeatedly call $\hat{g}_{ml}$, next iteration's initial distribution is current iteration's final distribution. This fact makes the estimates correlated. We sidestep this problem by assuming any initial distribution.

Table 2: Parameters of Algorithm 1

Hyperparameters		Momentums		Batch hidden parameters	
γ	Stepsize, $\in (0; \frac{3}{4L}]$	β	$\sqrt{\frac{4p^2\mu\gamma}{3}}$	$2^J l$	Batch size. If $2^J l > M$, then 0
t	Approximation step	η	$\frac{3\beta}{2p\mu\gamma} = \sqrt{\frac{3}{\mu\gamma}}$	J	Random, $J \sim \text{Geom}(1/2)$
B	Batch size multiplier	θ	$\frac{p\eta^{-1}-1}{\beta p\eta^{-1}-1}$	M	Batch size limit, $M = \frac{1}{p} + \frac{2}{\beta}$
N	Number of iterations	p	See Appendix	l	$(\lfloor \log_2 M \rfloor + 1) \cdot B$

Lemma 4 establishes key properties of the smoothed objective function. Lemma 5 provides bounds on the bias and variance of the baseline estimator $\hat{g}_{rd}$. Lemma 2 then quantifies how the MLMC scheme amplifies or reduces these statistics. Finally, in Section C.4, we combine the results of these lemmas to prove the first part of Theorem 1, bounding Algorithm 1's error. By tuning the parameters appropriately, we obtain the following iteration complexity bound:

Algorithm 1 Randomized Accelerated ZO GD

- 1: **Initialization:** $x_f^0 = x^0$; see Table 2.
- 2: **for** $k = 0, 1, 2, \dots, N-1$ **do**
- 3: $x_g^k = \theta x_f^k + (1-\theta)x^k$
- 4: Sample $J_k, \{e_i\}, \{F(x_g^k \pm te_i, Z_i^{(\pm)})\}$
- 5: Calculate $\hat{g}^k = \hat{g}_{ml}(x)$
- 6: $x_f^{k+1} = x_g^k - p\gamma\hat{g}^k$
- 7: $x^{k+1} = \eta x_f^{k+1} + (p-\eta)x_f^k + (1-p)(1-\beta)x^k + (1-p)\beta x_g^k$
- 8: **end for**

Theorem 1. Let Assumptions 1 to 4 hold, and consider problem (4) solved by Algorithm 1. Then, for any target accuracy ε and batch size multiplier B (see Tables 1 and 2 for notation), and for a suitable choice of γ, t, p , the number of oracle calls required to ensure $\mathbb{E}\|x^N - x^*\|^2 \leq \varepsilon$ is bounded by

$$B \cdot \tilde{\mathcal{O}} \left[\max\left(1, \frac{d}{B}\right) \sqrt{\frac{L}{\mu}} \log \frac{1}{\varepsilon} + \frac{Ld(d+\tau)\sigma_1^2}{B\mu^3\varepsilon^2} \right] \quad \text{one-point oracle calls.}$$

Theorem 1'. Let Assumptions 1' to 4' hold, and consider problem (4) solved by Algorithm 1. Then, for any target accuracy ε and batch size multiplier B (see Tables 1 and 2 for notation), and for a suitable choice of γ, t, p , the number of oracle calls required to ensure $\mathbb{E}\|x^N - x^*\|^2 \leq \varepsilon$ is bounded by

$$B \cdot \tilde{\mathcal{O}} \left[\max\left(1, \frac{d}{B}\right) \sqrt{\frac{L}{\mu}} \log \frac{1}{\varepsilon} + \frac{(d+\tau)\sigma_2^2}{B\mu^2\varepsilon} \right] \quad \text{two-point oracle calls.}$$

Remark. The iteration complexity of the algorithm, i.e., the number of iterates x^k generated (equal to the oracle complexity divided by B), is bound by $\tilde{\mathcal{O}}\left(\sqrt{\frac{L}{\mu}} \log \frac{1}{\varepsilon}\right)$ as the batch size multiplier B goes to infinity. This matches the optimal convergence rates for optimization with exact gradients [39].

2.3 Lower bounds

Here we present theorems demonstrating that no algorithm can asymptotically outperform Algorithm 1 in the smooth, strongly convex setting with either one- or two-point feedback.

Theorem 2. (Lower bounds) For any (possibly randomized) algorithm that solves the problem (4), there exists a function f that satisfies Assumptions 1 to 4 (1' to 4'), s.t. in order to achieve ε -approximate solution in expectation $\mathbb{E}\|x^N - x^*\|^2 \leq \varepsilon$, the algorithm needs at least

$$\Omega\left(\frac{d(d+\tau)\sigma_1^2}{\mu^2\varepsilon^2}\right) \quad \text{one-point or} \quad \Omega\left(\frac{(d+\tau)\sigma_2^2}{\mu^2\varepsilon}\right) \quad \text{two-point oracle calls.}$$

Remark. These results assume bounded second moments rather than uniform noise bounds. We explain how to adapt them to our setting, incurring only logarithmic overheads, in Section E.2.

Discussion. We now compare our results to existing work. Akhavan et al. [2] analyze a special case of the one-point setting where the noise is independent of the query points. This

aligns with our one-point oracle model and allows i.i.d. sampling as a Markov chain with fixed mixing time $\tau = 1$. The only factor they do not consider is σ_1^2 , which, however, appears in their proof with additional μ^2 factor if used with scaled Gaussian noise. We discuss this further in Appendix E.

In the work of Beznosikov et al. [5], a first-order Markovian oracle is considered, but the hard instance problem is a one-dimensional quadratic function, which makes first-order and zero-order information equivalent. Their result therefore corresponds to the $d = 1$ case in the two-point regime. Duchi et al. [15] provide tight lower bounds for general convex functions under two-point feedback. Their techniques can be extended to the strongly convex case by incorporating a shared quadratic component across the hard instances, as detailed in Appendix E, Theorem 10, yielding the bound we state for the two-point oracle with $\tau = 1$.

Our novel contribution lies in establishing a lower bound that scales as $d\tau$ in the one-point regime for large τ ; see Theorem 9. While our analysis relies on classical tools such as multidimensional hypothesis testing, the Markovian structure requires new bound on distances between joint distributions and the use of clipping. Detailed proofs, discussions, and further remarks on clipping appear in Appendix E.

3 Experiments

This section empirically supports our theoretical convergence rates and lower bounds, with particular focus on the stochastic component where we claim linear scaling in $d + \tau$ instead of $d\tau$.

Setup. Our setup repeats the problem we used to prove the lower bounds (see Appendix E and [51]). We consider a quadratic objective $f(x) = \frac{1}{2}\|x\|^2$ and a two-point Markovian oracle $F(x, Z) = f(x) + \langle x, Z \rangle$. The noise sequence $\{Z_i\}$ is a lazily updated standard Gaussian vector with variance σ_2^2 . Figure 2 illustrates how the optimization error of Algorithm 1 scales with mixing time, problem dimension, and different values of σ_2^2 .




Figure 2: Optimization error $\varepsilon = \|x^N - x^*\|^2$ after $N = 10^3$ iterations. Starting point error $\|x_0 - x^*\|^2 = 10^{-2}$. Stepsize $\gamma = 10^{-3}$, $t = 10^{-5}$. The results are averaged over 10^4 runs.

Discussion. The results confirm the linear dependence of the error on both the problem dimension d and the mixing time τ . The noise parameter σ^2 controls the influence of the stochastic part. In Fig. (a), where $\sigma_2^2 = 10^{-3}$, the stochastic component dominates, while in Fig. (c), with $\sigma_2^2 = 10^{-5}$, it is negligible. Fig. (b) shows an intermediate regime that smoothly interpolates between the two, yet maintains the linear scaling. The deterministic part (c) shows no dependence on mixing time, but grows linearly with d , which aligns with our theory (Theorem 1'). The stochastic part (a) scales as $(d + \tau)$, also matching the bound from the Theorem 1'.

References

- [1] Arya Akhavan, Massimiliano Pontil, and Alexandre Tsybakov. Exploiting higher order smoothness in derivative-free optimization and continuous bandits. *Advances in Neural Information Processing Systems*, 33:9017–9027, 2020.

- 347 [2] Arya Akhavan, Evgenii Chzhen, Massimiliano Pontil, and Alexandre B Tsybakov. 1
348 Gradient-free optimization of highly smooth functions: improved analysis and a new
349 algorithm. *Journal of Machine Learning Research*, 25(370):1–50, 2024.
- 350 [3] Peter Auer. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*,
351 47:235–256, 2002.
- 352 [4] El Houcine Bergou, Eduard Gorbunov, and Peter Richtárik. Stochastic three points
353 method for unconstrained smooth minimization. *SIAM Journal on Optimization*, 30(4):
354 2726–2749, 2020.
- 355 [5] Aleksandr Beznosikov, Sergey Samsonov, Marina Sheshukova, Alexander Gasnikov,
356 Alexey Naumov, and Eric Moulines. First order methods with markovian noise: from
357 acceleration to variational inequalities. *Advances in Neural Information Processing*
358 *Systems*, 36, 2024.
- 359 [6] Jalaj Bhandari, Daniel Russo, and Raghav Singal. A finite time analysis of temporal
360 difference learning with linear function approximation. In *Conference on learning theory*,
361 pages 1691–1692. PMLR, 2018.
- 362 [7] Sébastien Bubeck, Nicolo Cesa-Bianchi, et al. Regret analysis of stochastic and non-
363 stochastic multi-armed bandit problems. *Foundations and Trends® in Machine Learning*,
364 5(1):1–122, 2012.
- 365 [8] Pin-Yu Chen, Huan Zhang, Yash Sharma, Jinfeng Yi, and Cho-Jui Hsieh. Zoo: Zeroth
366 order optimization based black-box attacks to deep neural networks without training
367 substitute models. In *Proceedings of the 10th ACM workshop on artificial intelligence*
368 *and security*, pages 15–26, 2017.
- 369 [9] Krzysztof Choromanski, Mark Rowland, Vikas Sindhwani, Richard Turner, and Adrian
370 Weller. Structured evolution with compact architectures for scalable policy optimization.
371 In *International Conference on Machine Learning*, pages 970–978. PMLR, 2018.
- 372 [10] Aymeric Dieuleveut, Nicolas Flammarion, and Francis Bach. Harder, better, faster,
373 stronger convergence rates for least-squares regression. *Journal of Machine Learning*
374 *Research*, 18(101):1–51, 2017.
- 375 [11] Thinh T Doan. Finite-time analysis of markov gradient descent. *IEEE Transactions on*
376 *Automatic Control*, 68(4):2140–2153, 2022.
- 377 [12] Thinh T Doan, Lam M Nguyen, Nhan H Pham, and Justin Romberg. Convergence
378 rates of accelerated markov gradient descent with applications in reinforcement learning.
379 *arXiv preprint arXiv:2002.02873*, 2020.
- 380 [13] Ron Dorfman and Kfir Yehuda Levy. Adapting to mixing time in stochastic optimization
381 with markovian data. In *International Conference on Machine Learning*, pages 5429–
382 5446. PMLR, 2022.
- 383 [14] John C Duchi, Alekh Agarwal, Mikael Johansson, and Michael I Jordan. Ergodic mirror
384 descent. *SIAM Journal on Optimization*, 22(4):1549–1578, 2012.
- 385 [15] John C Duchi, Michael I Jordan, Martin J Wainwright, and Andre Wibisono. Optimal
386 rates for zero-order convex optimization: The power of two function evaluations. *IEEE*
387 *Transactions on Information Theory*, 61(5):2788–2806, 2015.
- 388 [16] Alain Durmus, Eric Moulines, Alexey Naumov, Sergey Samsonov, and Hoi-To Wai. On
389 the stability of random matrix product with markovian noise: Application to linear
390 stochastic approximation and td learning. In *Conference on Learning Theory*, pages
391 1711–1752. PMLR, 2021.
- 392 [17] Pavel Dvurechensky, Eduard Gorbunov, and Alexander Gasnikov. An accelerated
393 directional derivative method for smooth stochastic convex optimization. *European*
394 *Journal of Operational Research*, 290(2):601–621, 2021.

- [18] Mathieu Even. Stochastic gradient descent under markovian sampling schemes. In *International Conference on Machine Learning*, pages 9412–9439. PMLR, 2023.
- [19] Maryam Fazel, Rong Ge, Sham Kakade, and Mehran Mesbahi. Global convergence of policy gradient methods for the linear quadratic regulator. In *International conference on machine learning*, pages 1467–1476. PMLR, 2018.
- [20] Abraham D. Flaxman, Adam Tauman Kalai, and H. Brendan McMahan. Online convex optimization in the bandit setting: gradient descent without a gradient. In *Proceedings of the Sixteenth Annual ACM-SIAM Symposium on Discrete Algorithms, SODA '05*, page 385–394, USA, 2005. Society for Industrial and Applied Mathematics. ISBN 0898715857.
- [21] Alexander Gasnikov, Darina Dvinskikh, Pavel Dvurechensky, Eduard Gorbunov, Aleksandr Beznosikov, and Alexander Lobanov. *Randomized Gradient-Free Methods in Convex Optimization*, pages 1–15. Springer International Publishing, Cham, 2020. ISBN 978-3-030-54621-2. doi: 10.1007/978-3-030-54621-2_859-1. URL https://doi.org/10.1007/978-3-030-54621-2_859-1.
- [22] Alexander Gasnikov, Anton Novitskii, Vasilii Novitskii, Farshed Abdukhakimov, Dmitry Kamzolov, Aleksandr Beznosikov, Martin Takac, Pavel Dvurechensky, and Bin Gu. The power of first-order smooth optimization for black-box non-smooth problems. In Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvari, Gang Niu, and Sivan Sabato, editors, *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 7241–7265. PMLR, 17–23 Jul 2022. URL <https://proceedings.mlr.press/v162/gasnikov22a.html>.
- [23] Alexander V Gasnikov, Ekaterina A Krymova, Anastasia A Lagunovskaya, Ilnura N Usmanova, and Fedor A Fedorenko. Stochastic online optimization. single-point and multi-point non-linear multi-armed bandits. convex and strongly-convex case. *Automation and remote control*, 78:224–234, 2017.
- [24] Saeed Ghadimi and Guanghui Lan. Stochastic first-and zeroth-order methods for nonconvex stochastic programming. *SIAM journal on optimization*, 23(4):2341–2368, 2013.
- [25] Saeed Ghadimi and Guanghui Lan. Accelerated gradient methods for nonconvex nonlinear and stochastic programming. *Mathematical Programming*, 156(1):59–99, 2016.
- [26] Michael B. Giles. Multilevel monte carlo path simulation. *Operations Research*, 56(3): 607–617, 2008. doi: 10.1287/opre.1070.0496. URL <https://doi.org/10.1287/opre.1070.0496>.
- [27] Peter W. Glynn and Chang-Han Rhee. Exact estimation for markov chain equilibrium expectations. *Journal of Applied Probability*, 51A:377–389, 2014. ISSN 00219002. URL <http://www.jstor.org/stable/43284129>.
- [28] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [29] Eduard Gorbunov, Pavel Dvurechensky, and Alexander Gasnikov. An accelerated method for derivative-free smooth stochastic convex optimization. *SIAM Journal on Optimization*, 32(2):1210–1238, 2022. doi: 10.1137/19M1259225. URL <https://doi.org/10.1137/19M1259225>.
- [30] Elad Hazan and Satyen Kale. Beyond the regret minimization barrier: Optimal algorithms for stochastic strongly-convex optimization. *Journal of Machine Learning Research*, 15(71):2489–2512, 2014. URL <http://jmlr.org/papers/v15/hazan14a.html>.
- [31] Bjorn Johansson, Maben Rabi, and Mikael Johansson. A simple peer-to-peer algorithm for distributed optimization in sensor networks. In *2007 46th IEEE Conference on Decision and Control*, pages 4705–4710, 2007. doi: 10.1109/CDC.2007.4434888.

- [32] J. Kiefer. Sequential minimax search for a maximum. *Proceedings of the American Mathematical Society*, 4(3):502–506, 1953. ISSN 00029939, 10886826. URL <http://www.jstor.org/stable/2032161>.
- [33] Xiangru Lian, Yijun Huang, Yuncheng Li, and Ji Liu. Asynchronous parallel stochastic gradient for nonconvex optimization. *Advances in neural information processing systems*, 28, 2015.
- [34] Sijia Liu, Bhavya Kailkhura, Pin-Yu Chen, Paishun Ting, Shiyu Chang, and Lisa Amini. Zeroth-order stochastic variance reduction for nonconvex optimization. *Advances in Neural Information Processing Systems*, 31, 2018.
- [35] Cassio G. Lopes and Ali H. Sayed. Incremental adaptive strategies over distributed networks. *IEEE Transactions on Signal Processing*, 55(8):4064–4077, 2007. doi: 10.1109/TSP.2007.896034.
- [36] Eric Moulines and Francis Bach. Non-asymptotic analysis of stochastic approximation algorithms for machine learning. In J. Shawe-Taylor, R. Zemel, P. Bartlett, F. Pereira, and K.Q. Weinberger, editors, *Advances in Neural Information Processing Systems*, volume 24. Curran Associates, Inc., 2011. URL https://proceedings.neurips.cc/paper_files/paper/2011/file/40008b9a5380fcacce3976bf7c08af5b-Paper.pdf.
- [37] Deanna Needell, Rachel Ward, and Nati Srebro. Stochastic gradient descent, weighted sampling, and the randomized kaczmarz algorithm. *Advances in neural information processing systems*, 27, 2014.
- [38] Yu Nesterov. Efficiency of coordinate descent methods on huge-scale optimization problems. *SIAM Journal on Optimization*, 22(2):341–362, 2012.
- [39] Yurii Nesterov. A method for solving the convex programming problem with convergence rate $o(1/k^2)$. In *Doklad nauk Sssr*, volume 269, page 543, 1983.
- [40] Yurii Nesterov and Vladimir Spokoiny. Random gradient-free minimization of convex functions. *Foundations of Computational Mathematics*, 17(2):527–566, 2017.
- [41] Donald J Newman. Location of the maximum on unimodal surfaces. *Journal of the ACM (JACM)*, 12(3):395–398, 1965.
- [42] J. Nocedal and S. Wright. *Numerical Optimization*. Springer Series in Operations Research and Financial Engineering. Springer New York, 2006. ISBN 9780387227429. URL <https://books.google.ru/books?id=7wDpBwAAQBAJ>.
- [43] Boris Polyak. *Introduction to Optimization*. Optimization Software - Inc., Publications Division, 1987.
- [44] Yuyang Qiu, Uday Shanbhag, and Farzad Yousefian. Zeroth-order methods for non-differentiable, nonconvex, and hierarchical federated optimization. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 3425–3438. Curran Associates, Inc., 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/file/0a70c9cd8179fe6f8f6135fafa2a8798-Paper-Conference.pdf.
- [45] Herbert Robbins and Sutton Monro. A stochastic approximation method. *The annals of mathematical statistics*, pages 400–407, 1951.
- [46] Tim Salimans, Jonathan Ho, Xi Chen, Szymon Sidor, and Ilya Sutskever. Evolution strategies as a scalable alternative to reinforcement learning. *arXiv preprint arXiv:1703.03864*, 2017.
- [47] Bobak Shahriari, Kevin Swersky, Ziyu Wang, Ryan P. Adams, and Nando de Freitas. Taking the human out of the loop: A review of bayesian optimization. *Proceedings of the IEEE*, 104(1):148–175, 2016. doi: 10.1109/JPROC.2015.2494218.

- 491 [48] Shai Shalev-Shwartz and Shai Ben-David. *Understanding machine learning: From* 1
492 *theory to algorithms*. Cambridge university press, 2014.
- 493 [49] Ohad Shamir. An optimal algorithm for bandit and zero-order convex optimization
494 with two-point feedback. *Journal of Machine Learning Research*, 18(52):1–11, 2017.
- 495 [50] Ohad Shamir and Tong Zhang. Stochastic gradient descent for non-smooth optimization:
496 Convergence results and optimal averaging schemes. In Sanjoy Dasgupta and David
497 McAllester, editors, *Proceedings of the 30th International Conference on Machine*
498 *Learning*, volume 28 of *Proceedings of Machine Learning Research*, pages 71–79, Atlanta,
499 Georgia, USA, 17–19 Jun 2013. PMLR. URL [https://proceedings.mlr.press/v28/](https://proceedings.mlr.press/v28/shamir13.html)
500 [shamir13.html](https://proceedings.mlr.press/v28/shamir13.html).
- 501 [51] Alexander Shapiro, Darinka Dentcheva, and Andrzej Ruszczyński. *Lectures on*
502 *Stochastic Programming*. Society for Industrial and Applied Mathematics, 2009.
503 doi: 10.1137/1.9780898718751. URL [https://epubs.siam.org/doi/abs/10.1137/](https://epubs.siam.org/doi/abs/10.1137/1.9780898718751)
504 [1.9780898718751](https://epubs.siam.org/doi/abs/10.1137/1.9780898718751).
- 505 [52] Vladimir Solodkin, Andrew Veprikov, and Aleksandr Beznosikov. Methods for optimiza-
506 tion problems with markovian stochasticity and non-euclidean geometry. *arXiv preprint*
507 *arXiv:2408.01848*, 2024.
- 508 [53] Sebastian U. Stich. Unified optimal analysis of the (stochastic) gradient method, 2019.
509 URL <https://arxiv.org/abs/1907.04232>.
- 510 [54] Ben Taskar, Vassil Chatalbashev, Daphne Koller, and Carlos Guestrin. Learning
511 structured prediction models: A large margin approach. In *Proceedings of the 22nd*
512 *international conference on Machine learning*, pages 896–903, 2005.
- 513 [55] Alexandre B. Tsybakov. *Lower bounds on the minimax risk*, pages 77–135. Springer New
514 York, New York, NY, 2009. ISBN 978-0-387-79052-7. doi: 10.1007/978-0-387-79052-7_2.
515 URL https://doi.org/10.1007/978-0-387-79052-7_2.
- 516 [56] Sharan Vaswani, Francis Bach, and Mark Schmidt. Fast and faster convergence of sgd
517 for over-parameterized models and an accelerated perceptron. In *The 22nd international*
518 *conference on artificial intelligence and statistics*, pages 1195–1204. PMLR, 2019.
- 519 [57] Jian Wu, Saul Toscano-Palmerin, Peter I Frazier, and Andrew Gordon Wilson. Practical
520 multi-fidelity bayesian optimization for hyperparameter tuning. In *Uncertainty in*
521 *Artificial Intelligence*, pages 788–798. PMLR, 2020.
- 522 [58] David B Yudin and Arkadi S Nemirovskii. Informational complexity and efficient
523 methods for the solution of convex extremal problems. *Matekon*, 13(2):22–45, 1976.
- 524 [59] Yawei Zhao. Markov chain mirror descent on data federation. *arXiv preprint*
525 *arXiv:2309.14775*, 2023.

NeurIPS Paper Checklist

1. Claims 1

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? 2

Answer: [Yes] 3

Justification: main contributions of this paper are described accurately in a dedicated subsection (Section 1.2) of the introduction. 4

Guidelines: 5

- The answer NA means that the abstract and introduction do not include the claims made in the paper. 6
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations 7

Question: Does the paper discuss the limitations of the work performed by the authors? 8

Answer: [Yes] 9

Justification: assumptions we use to prove the main results are presented in Section 2. The motivation for these assumptions as well their limitations are also described there. 10

Guidelines: 11

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper. 12
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs¹

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?²

Answer: [Yes]³

Justification: all assumptions and definitions are carefully stated. The complete proofs appear in the supplemental material and are properly referenced in the main part.⁴

Guidelines:⁵

- The answer NA means that the paper does not include theoretical results.⁶
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility⁷

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?⁸

Answer: [Yes]⁹

Justification: see Section 3. The setup is fully disclosed.¹⁰

Guidelines:¹¹

- The answer NA means that the paper does not include experiments.¹²
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

(d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: our experiments are rather a practical confirmation of theoretical results, and these experiments can be easily reproduced.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: see Section 3, all parameters are described there.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: we use experiments to verify the theoretical rates and have no statistical effects associated with running the experiments. 1

Guidelines: 2

- The answer NA means that the paper does not include experiments. 3
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources 4

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments? 5

Answer: [No] 6

Justification: the experiments performed are not computationally heavy and can be reproduced on an average machine in a fairly reasonable amount of time. 7

Guidelines: 8

- The answer NA means that the paper does not include experiments. 9
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics 10

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>? 11

Answer: [Yes] 12

Justification: the research follows the NeurIPS Code of Ethics. 13

Guidelines: 14

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics. 15
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts 1

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed? 2

Answer: [NA] 3

Justification: there is no societal impact of the work performed – we only develop the theoretical understanding of Optimization. 4

Guidelines: 5

- The answer NA means that there is no societal impact of the work performed. 6
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards 7

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)? 8

Answer: [NA] 9

Justification: the paper poses no such risks. 10

Guidelines: 11

- The answer NA means that the paper poses no such risks. 12
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets 13

Question: Are the creators or original owners of assets (e.g., code, data, models) used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected? 14

Answer: [NA] 15

Justification: the paper does not use existing assets. 1

Guidelines: 2

- The answer NA means that the paper does not use existing assets. 3
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets 4

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets? 5

Answer: [NA] 6

Justification: the paper does not propose new assets. 7

Guidelines: 8

- The answer NA means that the paper does not release new assets. 9
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects 10

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)? 11

Answer: [NA] 12

Justification: the paper does not involve crowdsourcing nor research with human subjects. 13

Guidelines: 14

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 15
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects 16

1240	Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?	1
1241		
1242		
1243		
1244	Answer: [NA]	2
1245	Justification: the paper does not involve crowdsourcing nor research with human subjects.	3
1246		
1247	Guidelines:	4
1248	• The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.	5
1249		
1250	• Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.	
1251		
1252		
1253	• We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.	
1254		
1255		
1256	• For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.	
1257		
1258	16. Declaration of LLM usage	6
1259	Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.	7
1260		
1261		
1262		
1263		
1264	Answer: [NA]	8
1265	Justification: LLMs were used only for editing.	9
1266	Guidelines:	10
1267	• The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.	11
1268		
1269	• Please refer to our LLM policy (https://neurips.cc/Conferences/2025/LLM) for what should or should not be described.	
1270		