
Kernel Trace Distance: Quantum Statistical Metric between Measures through RKHS Density Operators

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Distances between probability distributions are a key component of many statistical
2 machine learning tasks, from two-sample testing to generative modeling, among
3 others. We introduce a novel distance between measures that compares them
4 through a Schatten norm of their kernel covariance operators. We show that
5 this new distance is an integral probability metric that can be framed between a
6 Maximum Mean Discrepancy (MMD) and a Wasserstein distance. In particular, we
7 show that it avoids some pitfalls of MMD, by being more discriminative and robust
8 to the choice of hyperparameters. Moreover, it benefits from some compelling
9 properties of kernel methods, that can avoid the curse of dimensionality for their
10 sample complexity. We provide an algorithm to compute the distance in practice by
11 introducing an extension of kernel matrix for difference of distributions that could
12 be of independent interest. Those advantages are illustrated by robust approximate
13 Bayesian computation under contamination as well as particle flow simulations.

1 INTRODUCTION

15 Statistical distances are ubiquitous in the fundamental theory of machine learning and serve as the
16 backbone of many of its applications, such as: discriminating between the generative model and real
17 data in Generative Adversarial Networks (GAN) [Goodfellow et al., 2014, Arjovsky et al., 2017, Li
18 et al., 2017, Genevay et al., 2018, Birrell et al., 2022], testing whether a dataset is close to another
19 (two-sample test) [Eric et al., 2007, Gretton et al., 2012, Hagrass et al., 2024] or to a particular
20 distribution (goodness-of-fit test), as well as acting as an objective loss function in particle gradient
21 flows [Arbel et al., 2019, Feydy et al., 2019, Korba et al., 2021, Hertrich et al., 2023, Neumayer et al.,
22 2024, Chen et al., 2024], or in minimum distance estimators [Wolfowitz, 1957, Basu et al., 2011].

23 A class of distances between probability distributions, called Integral Probability Metrics
24 (IPM) [Müller, 1997], is defined by measuring the supremum of difference of integrals over a
25 function space. It comprises many popular metrics such as the Total Variation distance, Wasserstein-1
26 distance and the Maximum Mean Discrepancy (MMD) [Gretton et al., 2012] also known as quadratic
27 distance [Lindsay et al., 2008]. IPMs' theoretical properties were largely investigated in the literature,
28 such as their statistical convergence rate [Sriperumbudur et al., 2010], concentration for inference
29 using ABC [Legramanti et al., 2022], PAC-Bayes bounds [Amit et al., 2022], as well as adversarial
30 interpretations [Husain and Knoblauch, 2022]. For instance, the MMD enjoys a fast statistical
31 convergence rate of $O(n^{-\frac{1}{2}})$ while the Wasserstein distance suffers from the curse of dimensionality
32 with a rate no better than $\Theta(n^{-\frac{1}{d}})$ [Kloeckner, 2012]. One could wonder: *how large such a function*
33 *space could be before the curse of dimensionality kicks in?* In this work, we theoretically investigate
34 how to get closer to such frontier by defining an extended family of kernel distances, that write as
35 novel IPM whose dual function space is larger than the one of MMD.

Kernel methods allow to represent a distribution by a vector by associating to a datapoint x a feature map image $\varphi(x)$ in a Hilbert space, and by doing so, embed in a linear way a distribution μ to what is called a *(kernel) mean embedding* $\mathbb{E}_{X \sim \mu}[\varphi(X)] = \int \varphi(x) d\mu(x)$. However mean embeddings for different distributions may have different “energies”, i.e., squared Hilbert norms, which may lead to several pitfalls of MMD. In quantum information theory [Watrous, 2018], a similar idea to mean embedding is called superposition. The quantum equivalent of a datapoint or deterministic Dirac distribution is called a *pure state* and is a projector of rank and trace one, that could be denoted vv^* (or $|v\rangle\langle v|$) for a unit vector v . Its analog for a general probability distribution is called a *mixed state* and is the superposition $\sum_v p(v)|v\rangle\langle v|$ where $p(v)$ are probabilities. A non-trivial mixed state can hardly be confused with a pure state as a linear combination of different projectors is of higher rank than 1: using projecting operators instead of the vectors themselves makes the *linearity less “trivial”*. As those positive definite operators can be diagonalised, by using always the same orthogonal basis and studying the eigenvalues, we recover classical probabilities, and as such we can see quantum probabilities as their extension. Recently, the work of [Bach, 2022] introduced a novel divergence between probability distributions, by plugging a kernel operator embedding of the distributions (which are also positive definite operators) in the Von Neumann relative entropy from quantum information theory (i.e., a Kullback-Leibler divergence between positive Hermitian operators), and whose statistical and geometrical properties were investigated more in depth in [Chazal et al., 2024]. Instead of considering a divergence on such operators, here we propose to draw inspiration from quantum statistical metrics, which enjoy nice geometrical properties such as the triangle inequality. Two of them are well-known and mutually bounding: the Bures metric, and the trace distance, on which we focus here, and which is derived from a (Schatten) norm.

Related works The kernelised version of Bures metric, i.e., a Bures metric between kernel covariance operators, has been studied for instance in [Oh et al., 2020], [Zhang et al., 2019]. The closest work to ours is the one by [Mroueh et al., 2017]. They consider a similar metric to ours, i.e. the trace distance, that they refer to as Covariance Matching IPM. It shares the same dual writing as the metric we consider, yet, in that work, the dual problem is solved through a numerical program involving neural networks that approach kernel features. Hence, they compute an approximate version of their target metric. In contrast, we use kernel features directly in the dual formulation, and derive a closed-form for the metric leveraging a kernel trick. Moreover, we provide theoretical guarantees regarding this metric and investigate different numerical applications than the one of the GAN considered in [Mroueh et al., 2017].

Contributions Our main contributions can be summarized as follows:

- (i) Inspired by quantum statistics, we introduce a novel distance between probability distributions called *kernel trace distance* (d_{KT}).
- (ii) We show that d_{KT} is an IPM and illustrate several of its theoretical properties, mainly: a direct comparison to MMD, robustness to contamination, and statistical convergence rates that do not depend on the dimension.
- (iii) We showcase how to compute d_{KT} and illustrate its practical performance on particle gradient flows and Approximate Bayesian Computation (ABC).

Organisation of the paper In section 2, we provide some background on quantum statistical distances and introduce d_{KT} . In section 3, we explain further the motivation to introduce d_{KT} , notably by comparing it with the other distances, MMD in particular. We show in section 4, under some eigenvalue decay rate assumptions, convergence rates that do not depend on the dimension, as well as robustness. In section 2.3, we explain how to compute d_{KT} . Finally, we illustrate our findings by experiments in section 5.

2 Kernel Trace Distance

For a positive semi-definite kernel $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$, its RKHS $\mathcal{H}$ is a Hilbert space of real-valued functions with inner product $\langle \cdot, \cdot \rangle_{\mathcal{H}}$ and norm $\| \cdot \|_{\mathcal{H}}$. It is associated with a feature map $\varphi : \mathcal{X} \rightarrow \mathcal{H}$ such that $k(x, y) = \langle \varphi(x), \varphi(y) \rangle_{\mathcal{H}}$. We denote $\mathcal{L}(\mathcal{H})$ the space of bounded linear operators from $\mathcal{H}$ to itself. For a vector $v \in \mathcal{H}$, v^* denotes its dual linear form defined by $v^*(w) = \langle v, w \rangle$ for any

87 $w \in \mathcal{H}$. For an operator $T \in \mathcal{L}(\mathcal{H})$, T^* is its adjoint. $\|\cdot\|_p$ denotes the p -Schatten norm explicited
88 below.

89 **Assumption 0.** In the whole paper, we restrict ourselves to the setting of a completely separable set
90 $\mathcal{X}$, endowed with a Borel σ -algebra, and a separable RKHS $\mathcal{H}$ of real-valued functions on $\mathcal{X}$, with a
91 bounded continuous strictly positive kernel.

92 2.1 Background

93 **RKHS density operators** [Bach, 2022]. Let μ a measure on $\mathcal{X}$. Define Φ the kernel covariance
94 operator embedding as:

$$\Phi : \mu \mapsto \Sigma_\mu = \int_{\mathcal{X}} \varphi(x) \varphi(x)^* d\mu(x). \quad (1)$$

95 We will call Σ_μ the RKHS density operator of μ , in reference to the wording of density operator
96 in quantum information theory: this is to insist that Σ_μ is an embedding in itself (with feature map
97 $\varphi(\cdot) \varphi(\cdot)^*$), rather than just the covariance of a mean embedding with feature map φ . The operator Σ_μ
98 is self-adjoint, and positive semidefinite when μ is a probability measure. To keep the analogy with
99 quantum density operators, similarly to [Bach, 2022], we consider kernels respecting the property:

100 **Assumption 1.** $\forall x \in \mathcal{X}, k(x, x) = 1$.

101 to ensure $\text{Tr} \Sigma_\mu = 1$ (as in the sum of all probabilities equals one). If $\forall x \in \mathcal{X}, k(x, x) = M$ for a
102 non-zero constant $M \neq 1$, it is will be easy to generalize many of our results later by dividing by
103 M , so this assumption is not too restrictive. If the kernel does not verify Assumption 1 but is strictly
104 positive, it is could also be normalised using $\tilde{k}(x, y) = \frac{k(x, y)}{\sqrt{k(x, x)k(y, y)}}$ instead.

105 **Schatten norms.** We now provide some background on Schatten norms [Simon, 2005]. For an
106 operator $T \in \mathcal{L}(\mathcal{H})$ and $p \in [1, \infty)$, the p -Schatten norm is defined as $\|T\|_p = (\text{Tr}(|T|^p))^{1/p}$ where
107 $|T| = \sqrt{T^*T}$. If T is compact, this can be rewritten as the p -vectorial norm of the singular values of
108 T . It also admits a dual definition, denoting q such that $1/p + 1/q = 1$:

$$\|T\|_p = \sup_{U \in \mathcal{L}(\mathcal{H}), \|U\|_q = 1} \langle U, T \rangle \quad (2)$$

109 where the inner product is $\langle U, T \rangle = \text{Tr}(U^*T)$.

110 The Schatten 2-norm is the Hilbert-Schmidt norm with respect to this inner product: $\|T\|_2 =$
111 $\sqrt{\text{Tr}(T^*T)}$. Then, the Schatten ∞ -norm is the operator norm : $\|T\|_\infty = \sup_{x \in \mathcal{H} \setminus 0} \frac{\|Tx\|_{\mathcal{H}}}{\|x\|_{\mathcal{H}}}$ i.e., the
112 maximum of the singular values of the operator in absolute value. We have the following inequalities:

- 113 • For $1 \leq p \leq q \leq \infty$: $\forall T \in \mathcal{L}(\mathcal{H}), \|T\|_1 \geq \|T\|_p \geq \|T\|_q \geq \|T\|_\infty$.
- 114 • $\forall T, S \in \mathcal{L}(\mathcal{H}), \|TS\|_1 \leq \|T\|_2 \|S\|_2$.
- 115 • From this, it can be deduced taking T as the identity operator, for $\mathcal{H}$ of finite dimension:

$$\forall S \in \mathcal{L}(\mathcal{H}), \|S\|_1 \leq \sqrt{\dim(\mathcal{H})} \|S\|_2. \quad (4)$$

116 2.2 Definition

117 In quantum information theory, the trace distance is a mathematical tool that can be used to compare
118 density operators by measuring the Schatten 1-norm of their difference. Inspired by this, we define:

119 **Definition 2.1.** The *kernel trace distance* between two probability measures μ, ν on $\mathcal{X}$ is defined as:

$$d_{KT}(\mu, \nu) = \|\Sigma_\mu - \Sigma_\nu\|_1.$$

120 We will also relate it to other distances such as:

- 121 • Wasserstein distances [Villani, 2009]:

$$W_d(\mu, \nu) = \inf_{\pi \in \Pi(\mu, \nu)} \iint d(x, y) d\pi(x, y)$$

122 where $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}^+$ is a cost and $\Pi(\mu, \nu)$ denotes all the possible couplings between μ
123 and ν . The Wasserstein- p distance is obtained by replacing d by its power d^p in the integral
124 and taking the p -root of the whole expression.

- The Bures distance [Bhatia et al., 2019] on positive definite matrices A :

$$d_{BW}(A, B) = \sqrt{\text{Tr } A + \text{Tr } B - 2F(A, B)}$$

where $F(A, B) = \text{Tr}(A^{1/2}BA^{1/2})^{1/2}$ is called the fidelity. It coincides with the Wasserstein-2 distance between two normal distributions (also called Bures-Wasserstein distance) with identical mean, and different covariances A and B . The formula can be extended to operators with finite traces.

- The Kernel Bures distance [Zhang et al., 2019] is defined as:

$$d_{KBW}(\mu, \nu) = d_{BW}(\Sigma_\mu, \Sigma_\nu).$$

- The Total Variation is a special case of the Wasserstein distance where the cost is $d : (x, y) \mapsto 1_{x=y}$ and can be expressed as:

$$\|\mu - \nu\|_{TV} = \frac{1}{2} \int_{\mathcal{X}} |\mu(x) - \nu(x)| dx$$

- The Maximum Mean Discrepancy [Gretton et al., 2012]:

$$\text{MMD}(\mu, \nu) = \left\| \int_{\mathcal{X}} k(x, \cdot) \mu(x) dx - \int_{\mathcal{X}} k(x, \cdot) \nu(x) dx \right\|_{\mathcal{H}}$$

- Integral Probability Metrics (IPM) [Müller, 1997] defined as:

$$d(\mu, \nu) = \sup_{f \in \mathcal{F}} \{|\mathbb{E}_{X \sim \mu}[f(X)] - \mathbb{E}_{X \sim \nu}[f(X)]|\}$$

where the function space $\mathcal{F}$ is rich enough to make this expression a metric. The Wasserstein-1 distance, the TV and MMD are IPMs (with $\mathcal{F}$ being 1-Lipschitz functions w.r.t. $\|\cdot\|$, functions with values in $[-1, 1]$, and a RKHS unit ball respectively).

Proposition 2.2. *If k^2 is characteristic i.e Φ is injective, d_{KT} and d_{KBW} are metrics.*

PROOF. Symmetry, non-negativity, triangle inequality and $d_{KT}(\mu, \mu) = 0$ (resp. $d_{KBW}(\mu, \mu) = 0$) are naturally inherited from the Schatten norm on operators for d_{KT} and from the standard Bures-Wasserstein distance for d_{KBW} . Then, as $d_{KT}(\mu, \nu) = 0$ (resp. $d_{KBW}(\mu, \nu) = 0$) implies $\Sigma_\mu = \Sigma_\nu$, injectivity of Φ enforces $\mu = \nu$.

Examples of characteristic kernels are the family of Gaussian kernels, whose squared kernel also belong to, modulo a change of parameter. On compact set, a sufficient condition for characteristicity is universality [Steinwart, 2001], see for instance [Bach, 2022].

2.3 Computation for discrete measures

As interesting, i.e. expressive RKHS are often of infinite dimension, computations with kernel methods relies on the so-called “kernel trick”, reducing computation on the empirical kernel matrix (Gram matrix of two sets of samples using the kernel inner product) which is of finite dimension. It is well-known that the spectrum of the covariance operator Σ_{μ_n} are the ones of the kernel Gram matrix $(k(x_i, x_j))_{i,j=1}^n$ divided by the number of samples [Bach, 2022, Proposition 6]. Here, we generalise the concept for differences of distributions.

First, notice that $\Sigma_{\mu_n} - \Sigma_{\nu_m} = \Sigma_{\mu_n - \nu_m}$, which incites us to consider the samples from each distribution altogether. We denote without duplicates $(z_k)_{k=1, \dots, r}$ the samples in the union of the sample sets X, Y (corresponding respectively to distributions μ_n, ν_m), where r is the number of distinct elements in X, Y . We note $Z = [\tilde{\varphi}(z_k)]_{k=1 \dots r}$ the column of vectors in $\mathcal{H}$ where $\tilde{\varphi}(z_k) = \sqrt{(\mu_n - \nu_m)(\{z_k\})} \varphi(z_k)$ if $(\mu_n - \nu_m)(\{z_k\}) \geq 0$, $\tilde{\varphi}(z_k) = i \sqrt{|(\mu_n - \nu_m)(\{z_k\})|} \varphi(z_k)$ else.

We can see Z by a slight abuse of notation as the linear map $Z : \mathcal{H} \rightarrow \mathbb{C}^r, v \mapsto [\langle \tilde{\varphi}(z_1), v \rangle, \dots, \langle \tilde{\varphi}(z_r), v \rangle]$ and by duality Z^* (real not Hermitian adjoint) would be the linear map $Z^* : \mathbb{C}^r \rightarrow \mathcal{H}, u \mapsto \sum_{i=1, \dots, r} u_i \tilde{\varphi}(z_i)$.

Then we define the *difference kernel matrix* as $K = Z^* Z$. Typically, in case where all samples are distinct, $X \cap Y = \emptyset$ and $(\mu_n - \nu_m)(\{z_k\}) = \mu_n(\{z_k\}) = 1/n$ for samples $z_k \in X$ from μ_n and $(\mu_n - \nu_m)(\{z_k\}) = -\nu_m(\{z_k\}) = 1/m$ for samples $z_k \in Y$ from ν_m , then

$$K = \begin{bmatrix} \frac{1}{n} K_{XX} & \frac{i}{\sqrt{mn}} K_{XY} \\ -\frac{i}{\sqrt{mn}} K_{YX} & -\frac{1}{m} K_{YY} \end{bmatrix}$$

where $K_{XX}, K_{YY}, K_{YX}, K_{XY}$ are the usual kernel Gram matrices. Other cases are similar, adjusting the probability weights on rows and columns.

Proposition 2.3. *Assume the kernel is such that for any family (x) of distinct elements of $\mathcal{X}$, $(\varphi(x))$ is linearly independent. The difference kernel matrix K as defined just above and $\Sigma_{\mu_n - \nu_m}$ have the same eigenvalues, whose Schatten 1-norm is $d_{KT}(\mu_n, \nu_m)$.*

The proof of Proposition 2.3 is deferred to Appendix A.4. The condition is verified by the Gaussian kernel and more generally it is equivalent to the kernel being strictly positive. It is sufficient to get the eigenvalues by either Autonne-Takagi factorisation [Autonne 1915, Takagi 1924], Schur or Singular Value decomposition, and compute their 1-norm. This SVD is of complexity $O(r^3)$ in general.

3 Discriminative properties

In this section, we study the discriminative properties of the d_{KT} distance and how it relates to alternative distances between distributions introduced previously.

3.1 Comparison with other distances

We first show that our novel distance d_{KT} belongs to the family of Integral Probability Metrics (IPM).

Proposition 3.1.

(i) d_{KT} is an IPM with respect to the function space $\mathcal{F}_1 = \{f : x \mapsto \varphi(x)^* U \varphi(x) | U \in \mathcal{L}(\mathcal{H}), \|U\|_\infty = 1\}$.

Moreover if Assumption 1 is verified:

(ii) functions in $\mathcal{F}_1$ have values in $[-1, 1]$, and

(iii) verify the following ‘‘Lipschitz’’ property: $\forall x, y \in \mathcal{X}, |f(x) - f(y)| \leq 2\|\varphi(x) - \varphi(y)\|_{\mathcal{H}}$.

The proof of Proposition 3.1 is deferred to Appendix A.2. Since the TV distance is an IPM with respect to functions bounded by 1, we have the following corollary:

Corollary 3.2. $d_{KT}(\mu, \nu) \leq \|\mu - \nu\|_{TV}$.

We also have a direct comparison between d_{KT} and a MMD.

Lemma 3.3. *The Schatten 2-norm of the difference of the RKHS density operators of two probability distributions μ, ν on $\mathcal{X}$ can be identified to their Maximum Mean Discrepancy using the kernel k^2 :*

$$\|\Sigma_\mu - \Sigma_\nu\|_2 = \text{MMD}_{k^2}(\mu, \nu)$$

Consequently, since d_{KT} is a Schatten 1-norm of this difference, $\text{MMD}_{k^2}(\mu, \nu) \leq d_{KT}(\mu, \nu)$.

This follows mainly from the fact that $\langle \Sigma_\mu, \Sigma_\nu \rangle = \int_{\mathcal{X}} \int_{\mathcal{Y}} k(x, y) k(x, y) \mu(x) \nu(y) dx dy$ (see Appendix A.1.1). Finally, we can relate d_{KT} to some Wasserstein distance. Denoting $c_k(x, y) = \|\varphi(x) - \varphi(y)\|_{\mathcal{H}} = \sqrt{2(1 - k(x, y))}$ a cost defined from the kernel k , and applying the Lipschitz property of Theorem 3.1 we get the following:

Corollary 3.4. *If Assumption 1 is verified, $d_{KT}(\mu, \nu) \leq 2W_{c_k}(\mu, \nu)$. Furthermore, using the Gaussian kernel with parameter σ ,*

$$d_{KT}(\mu, \nu) \leq 2W_{c_k}(\mu, \nu) \leq \frac{2}{\sigma} W_{\|\cdot\|}(\mu, \nu).$$

The last remark is due to the fact that the Wasserstein-1 distance is an IPM defined by the functions which are 1-Lipschitz w.r.t. $\|\cdot\|$, and for the Gaussian kernel $k(x, y) = e^{-\frac{\|x-y\|^2}{2\sigma^2}}$, we have $c_k(x, y) \leq \frac{\|x-y\|}{\sigma}$. See Appendix A.1.1 for full proof.

Finally our novel distance can be related to other kernelized quantum divergences. Some well-known inequality in quantum information theory relating the trace distance and the fidelity is the following Fuchs and Van De Graaf [1999] inequality :

$$2(1 - F(A, B)) \leq \|A - B\|_1 \leq 2\sqrt{1 - F(A, B)^2} \quad (5)$$

which translates as upper and lower bounds on d_{KT} with respect to d_{KBW} (see proof in Appendix A.1.1) using Assumption 1:

$$d_{KBW}(\mu, \nu)^2 \leq d_{KT}(\mu, \nu) \leq 2d_{KBW}(\mu, \nu) \quad (6)$$

Let $D_{KL}(A|B) = \text{Tr}(A(\log A - \log B))$ the quantum relative entropy. The Kernel-Kullback-Leibler (KKL) divergence introduced in Bach [2022] is defined as the latter applied to the density operators of two distributions μ, ν on $\mathcal{X}$ (in particular, it is infinite if μ is not absolutely continuous w.r.t. ν). Thanks to the (quantum) Pinsker’s inequality, we have then: $\frac{1}{2}d_{KT}(\mu, \nu)^2 \leq D_{KL}(\Sigma_\mu|\Sigma_\nu) := \text{KKL}(\mu|\nu)$. Hence, our distance can be framed within several well-known alternative discrepancies.

3.2 Normalized energy

From our Assumption 1 on the kernel, we have ensured that for any measure μ , $\|\Sigma_\mu\|_1 = 1$ which means that all measures representations considered are somehow “normalised”. On the contrary, for MMD with k^2 (or the Schatten 2-norm), $\|\Sigma_\mu\|_2$ the “internal energy” depends on the measure (and on the kernel parameters such as bandwidth) and it can be smaller for distributions which are very flat, with high variance, as in general $k(x, y) \leq k(x, x)$ for $x \neq y$. This has consequences as intrinsically $\|\Sigma_\mu - \Sigma_\nu\|_2 \leq \sqrt{\|\Sigma_\mu\|_2^2 + \|\Sigma_\nu\|_2^2}$, the maximum value can be already small independently of the differences between μ and ν . When minimizing an objective such as $\mu \mapsto \|\Sigma_\mu - \Sigma_\nu\|_2$ (e.g., with gradient descent on the atoms in the support of μ if it is a discrete measure, as in Arbel et al. [2019]), this has an impact on the shape of the slope. Moreover, the energy depends on the hyperparameters of the kernel, which are hard to tune for both the distributions’ variances and the distance between their means at the same time.




Figure 1: Kernel distances between $\mu = \mathcal{N}(0, 1)$ and $\nu = \mathcal{N}(5, 1)$, as a function of the Gaussian kernel bandwidth σ .

Figure 1 illustrates this by displaying the two distances between sets of $n = 1000$ samples from $\mathcal{N}(0, 1)$ and $\mathcal{N}(5, 1)$. We would expect sample sets to look closer as the Gaussian kernel bandwidth σ grows, but for MMD that is not always the case. Other such phenomena are displayed by varying the variance or the mean of the distributions in the Appendix B.1.

Now let us consider two measures μ, ν on $\mathcal{X}$ such that $\mathbb{E}_{X \sim \mu, Y \sim \nu}[k(X, Y)] \leq \epsilon$ for some small parameter $\epsilon > 0$. Then, $\langle \Sigma_\mu, \Sigma_\nu \rangle \leq \epsilon$ by Cauchy-Schwartz. Consider the density operator of the mixture $\Sigma_{\frac{1}{2}\mu + \frac{1}{2}\nu} = \frac{1}{2}\Sigma_\mu + \frac{1}{2}\Sigma_\nu$, we have:

$$\|\Sigma_{\frac{1}{2}\mu + \frac{1}{2}\nu}\|_1 = 1 = \frac{1}{2}\|\Sigma_\mu\|_1 + \frac{1}{2}\|\Sigma_\nu\|_1, \quad \|\Sigma_{\frac{1}{2}\mu + \frac{1}{2}\nu}\|_2^2 \leq \frac{1}{2} \left(\frac{1}{2}\|\Sigma_\mu\|_2^2 + \frac{1}{2}\|\Sigma_\nu\|_2^2 + \epsilon \right).$$

We see that in contrast to the 1-Schatten norm, the 2-Schatten norm energy bound is roughly divided by 2 (as $\epsilon \rightarrow 0$, e.g. for almost orthogonal Σ_μ, Σ_ν). Then, we reason with distance rather than norm:

Proposition 3.5. *Let us consider distances between two mixtures $P = \frac{1}{2}\mu_1 + \frac{1}{2}\mu_2$ and $Q = \frac{1}{2}\nu_1 + \frac{1}{2}\nu_2$ such that $\Sigma_{\mu_1}, \Sigma_{\nu_1}$ are orthogonal to $\Sigma_{\mu_2}, \Sigma_{\nu_2}$. Then:*

$$d_{KT}(P, Q) = \frac{1}{2}d_{KT}(\mu_1, \nu_1) + \frac{1}{2}d_{KT}(\mu_2, \nu_2)$$

$$\text{MMD}_{k^2}^2(P, Q) = \frac{1}{4}\text{MMD}_{k^2}^2(\mu_1, \nu_1) + \frac{1}{4}\text{MMD}_{k^2}^2(\mu_2, \nu_2).$$

See proof in the Appendix A.2.1. If the distance between μ_2 and ν_2 are the same as between μ_1 and ν_1 (for instance, if the former are respective translation of the latter and the kernel is translation-invariant), we can see that the squared MMD distance loses a factor 2 while d_{KT} behaves similarly to the Total Variation of the mixtures when μ_1, ν_1 have different supports than μ_2, ν_2 . This is the case when taking for instance in $\mathcal{X} = \mathbb{R}^2$ $\mu_1 = \mathcal{N}([0, 0], I_2)$ and $\nu_1 = \mathcal{N}([0.3, 0.3], I_2)$ while $\mu_2 = \mathcal{N}(\Delta, I_2)$ and $\nu_2 = \mathcal{N}(\Delta + [0.3, 0.3], I_2)$ for $\Delta = [10, 10]$. In practice, the RKHS density operators are not perfectly orthogonal unless $\|\Delta\| \rightarrow +\infty$ (in that case $\langle \Sigma_\mu, \Sigma_\nu \rangle \rightarrow 0$ for a fixed bandwidth), but typically they can look so up to numerical precision, when using exponentially decreasing kernels (e.g., Gaussian). Taking $n = 100$ samples each from each μ_1 and ν_1 , and translating them by Δ , the results above from

251 Proposition 3.5 are confirmed numerically: we find empirically $\widehat{d_{KT}}(P, Q) = \widehat{d_{KT}}(\mu_1, \nu_1) = 0.5992$
 252 while $\widehat{\text{MMD}}_{k^2}^2(\mu_1, \nu_1) = 0.0253$ but $\widehat{\text{MMD}}_{k^2}^2(P, Q) = 0.0127$, half of it (for a Gaussian kernel with
 253 bandwidth $\sigma = 0.5$).

254 3.3 Robustness

255 We now turn to investigating the robustness of the kernel trace distance. In particular, we consider the
 256 ϵ -contamination model, where the training dataset is supposedly contaminated by a fraction $\epsilon \in (0, 1)$
 257 of outliers [Huber, 1964]. The following proposition quantifies the robustness of this distance.

258 **Proposition 3.6.** Denote $P_\epsilon = (1 - \epsilon)P + \epsilon C$ where C is some contamination distribution. We have
 259 when Assumption 1 is verified: $|d_{KT}(P_\epsilon, Q) - d_{KT}(P, Q)| \leq 2\epsilon$.

260 The proof relies on the triangular inequality (see Appendix A.3.2). Hence, we see that d_{KT} is robust
 261 while for the Wasserstein distance, a contamination C arbitrarily “far away from the distribution Q ”
 262 will incur an arbitrarily high distance. The proof of robustness also works for MMD.

263 4 Statistical Properties

264 4.1 Convergence rate

265 In this section, we consider a measure μ and its empirical counterpart μ_n for n independent samples
 266 and study the rate of convergence of $d_{KT}(\mu, \mu_n)$. We note $A \lesssim_{\mu \otimes n} b$ where A is r.v., when for any
 267 $\delta > 0$, there exists $c_\delta < \infty$ such that $\mu^{\otimes n}(A \leq c_\delta b) \geq \delta$. With the Schatten 1-norm, it is not enough
 268 to study only the concentration of one (the maximal) eigenvalue as for the operator norm ($p = \infty$),
 269 we need to handle an infinity of eigenvalues (when the RKHS is of infinite dimension), neither can we
 270 use the Cauchy-Schwarz trick as for the Hilbert norm ($p = 2$). However, since the trace of our kernel
 271 density operators are bounded by 1, only a few of the eigenvalues will have a significant contribution.
 272 Therefore, assuming some decay rate on those eigenvalues, we can focus on the convergence of
 273 operators on a subspace of the top eigenvectors, using results from the Kernel PCA literature. We
 274 introduce the population and empirical square loss associated with some projector P :

$$R(P) = \mathbb{E}_{X \sim \mu} \|\phi(X) - P\phi(X)\|_{\mathcal{H}}^2, \quad R_n(P) = \sum_{i=1}^n \frac{1}{n} \|\phi(x_i) - P\phi(x_i)\|_{\mathcal{H}}^2$$

275 where the $(x_i)_{i=1 \dots n}$ are each drawn independently from μ . We first make the following assumption,
 276 as in [Sterge et al., 2020].

277 **Assumption 2.** The eigenvalues $(\lambda_i)_{i \in I}$ of Σ_μ (resp. $(\hat{\lambda}_j)_{j \in J}$ of Σ_{μ_n}) are positive, simple and
 278 w.l.o.g. arranged in decreasing order ($\lambda_1 \geq \lambda_2 \geq \dots$).

279 This allows us to denote $P^l(\Sigma_\mu)$ the projector on the subspace of the l eigenvectors associated with
 280 the l highest eigenvalues $\lambda_1, \dots, \lambda_l$. Note that $\|P^l(\Sigma_\mu)\Sigma_\mu - \Sigma_\mu\|_1 = \sum_{i>l} \lambda_i = R(P^l(\Sigma_\mu))$ (see
 281 for instance [Blanchard et al., 2007], [Rudi et al., 2013]). Similarly we consider $P^l(\Sigma_{\mu_n})$ for Σ_{μ_n} .

282 We now consider different kinds of assumptions on the decay rate of eigenvalues of Σ_μ to get different
 283 corresponding convergence rates, as in [Sterge et al., 2020], [Sterge and Sriperumbudur, 2022].

284 **Assumption P (Polynomial).** For some $\alpha > 1$ and $0 < \underline{A} < \bar{A} < \infty$,

$$\underline{A}i^{-\alpha} \leq \lambda_i \leq \bar{A}i^{-\alpha}. \quad (\text{P})$$

285 **Assumption E (Exponential).** For $\tau > 0$ and $\underline{B}, \bar{B} \in (0, \infty)$,

$$\underline{B}e^{-\tau i} \leq \lambda_i \leq \bar{B}e^{-\tau i}. \quad (\text{E})$$

286 **Lemma 4.1.** Suppose Assumption 1 and 2 are verified. With a polynomial decay rate of order $\alpha > 1$
 287 (Assumption P), for $l = n^{\frac{\theta}{\alpha}}, 0 < \theta \leq \alpha$:

$$\|P^l(\Sigma_\mu)\Sigma_\mu - \Sigma_\mu\|_1 = R(P^l(\Sigma_\mu)) = \Theta\left(n^{-\theta(1-\frac{1}{\alpha})}\right), \quad \|P^l(\Sigma_\mu)\Sigma_\mu - \Sigma_\mu\|_2 = \Theta\left(n^{-\theta(1-\frac{1}{2\alpha})}\right), \quad (7)$$

288 and there exists $N \in \mathbb{N}$ such that for $n > N$:

$$\|P^l(\Sigma_{\mu_n})\Sigma_\mu - \Sigma_\mu\|_2 \lesssim_{\mu \otimes n} \max(n^{-\frac{1}{2} + \frac{1}{4\alpha}}, n^{-\theta + \frac{1}{4\alpha}}). \quad (8)$$

289 With an exponential decay rate (Assumption [E](#)), for $l = \frac{1}{\tau} \log n^\theta, \theta > 0$:

$$\|P^l(\Sigma_\mu)\Sigma_\mu - \Sigma_\mu\|_1 = R(P^l(\Sigma_\mu)) = \Theta(n^{-\theta}), \quad \|P^l(\Sigma_\mu)\Sigma_\mu - \Sigma_\mu\|_2 = \Theta(n^{-\theta}) \quad (9)$$

290 and there exists $N \in \mathbb{N}$ such that for $n > N$:

$$\|P^l(\Sigma_{\mu_n})\Sigma_\mu - \Sigma_\mu\|_2 \lesssim_{\mu \otimes n} \begin{cases} \sqrt{\frac{\log n}{n^\theta}} & \text{if } \theta < 1 \\ \frac{(\log n)}{\sqrt{n}} & \text{if } \theta \geq 1. \end{cases} \quad (10)$$

291 The previous lemma (see proof in Appendix [A.3.1](#)) is crucial to prove our main theorem below, that
292 provides dimension-independent statistical rates.

293 **Theorem 4.2.** Suppose Assumption [1](#) and [2](#) are verified.

- 294 • If the eigenvalues of Σ_μ follow a polynomial decay rate of order $\alpha > 1$ (Assumption [P](#)),
295 then:

$$d_{KT}(\mu, \mu_n) \lesssim_{\mu \otimes n} n^{-\frac{1}{2} + \frac{1}{2\alpha}}.$$

- 296 • If the eigenvalues of Σ_μ follow an exponential decay rate (Assumption [E](#)), then:

$$d_{KT}(\mu, \mu_n) \lesssim_{\mu \otimes n} \frac{(\log n)^{\frac{3}{2}}}{\sqrt{n}}.$$

297 **SKETCH OF PROOF.** For clarity of notation, we abbreviate Σ_μ and Σ_{μ_n} as Σ and Σ_n . By the
298 triangular inequality:

$$\begin{aligned} \|\Sigma - \Sigma_n\|_1 &\leq \|\Sigma - P^l(\Sigma)\Sigma\|_1 + \|(P^l(\Sigma) - P^l(\Sigma_n))\Sigma\|_1 + \|P^l(\Sigma_n)(\Sigma - \Sigma_n)\|_1 \\ &\quad + \|P^l(\Sigma_n)\Sigma_n - \Sigma_n\|_1 := (A) + (B) + (C) + (D) \end{aligned} \quad (11)$$

299 We bound each term of eq. [11](#). Term (A) is bounded using Lemma [4.1](#). Similarly, (D) relates to (A)
300 by a result due to [Blanchard et al. \[2007\]](#) (eq. (30)), see Lemma [A.3](#) in Appendix [A.3.1](#). For (B) and
301 (C), the projections allow to work in a subspace of dimension at most $2l$ and by eq. (3) (Hölder's
302 inequality) to relate to the Schatten 2-norm which has rates like MMD. Finally, we pick $\theta = \frac{1}{2}$ for
303 polynomial decay and $\theta = 1$ for the exponential decay (see Lemma [4.1](#)) to minimise the maximum
304 of the four terms. See Appendix [A.3.1](#) for the full proof.

305 By the Fuchs-van de Graaf inequality (Eq. [\(5\)](#) and [\(6\)](#)), it directly implies (also dimensionally-
306 independent) convergence rates for the Kernel Bures Wasserstein distance, that are novel to the best
307 of our knowledge.

308 **Corollary 4.3.** Suppose Assumption [1](#) and [2](#) verified. If Assumption [P](#) is verified:
309 $d_{KBW}(\mu, \mu_n) \lesssim_{\mu \otimes n} n^{-\frac{1}{4} + \frac{1}{4\alpha}}$. If Assumption [E](#) is verified: $d_{KBW}(\mu, \mu_n) \lesssim_{\mu \otimes n} (\log n)^{\frac{3}{4}} n^{-\frac{1}{4}}$.

310 5 Experiments

311 In this section, we illustrate the interest of our novel kernel trace distance on different experiments.

312 **Approximate Bayesian Computation (ABC)** The purpose of Approximate Bayesian Compu-
313 tation [\[Tavaré et al., 1997\]](#) is to compute an approximation of the posterior when doing Bayesian
314 inference in a likelihood-free fashion. The idea of using a distance d between distributions to build a
315 synthetic likelihood has recently flourished [\[Frazier, 2020, Bernton et al., 2019, Jiang, 2018\]](#). ABC
316 methods based on IPM enjoy theoretical guarantees [\[Legramanti et al., 2022\]](#). The ABC posterior
317 distribution is defined by $\pi(\theta|X^n) \propto \int \pi(\theta) \mathbb{1}_{\{d(X^n, Y^m) < \epsilon\}} p_\theta(Y^m) dY^m$, where $\pi(\theta)$ is a prior over
318 the parameter space Θ , $\epsilon > 0$ is a tolerance threshold, and Y^m are synthetic data generated according
319 to $p_\theta(Y^m) = \prod_{j=1}^m p_\theta(Y_j)$. It is approximately computed by drawing $\theta_i \sim \pi$ for $i = 1, \dots, T$ and
320 simulating synthetic data $Y^m \sim p_{\theta_i}$ and keeping or rejecting θ_i according to whether the synthetic
321 data is close to the real data. The result is a list L_θ of all accepted θ_i (see Algo. [1](#) in the Appendix [B.2](#)).

322 Here, as we are interested in robustness, we will consider a contamination case using Normal
323 distributions but where nonetheless the usual likelihood fails to recover the correct mean as the data is
324 corrupted. We will take as prior $\pi = \mathcal{N}(0, \sigma_0^2)$ and the real data consist of $n = 100$ samples coming
325 following $\mu^* = \mathcal{N}(\theta^* = 1, 1)$ where 10% of the samples are replaced by contaminations from
326 $\mathcal{N}(20, 1)$. We fit the model $p_\theta = \mathcal{N}(\theta, 1)$ by picking the best θ possible. We carry out $T = 10000$
327 iterations, generating each times $m = n$ synthetic data.

328 We consider ABC with the threshold value $\epsilon = 0.05, 0.25, 0.5, 1$. For the proposed distance d_{KT} ,
329 Bayes' rule gives posterior $p(\theta|x) = \mathcal{N}(\frac{\sum_{i=1}^n x_i}{n+\frac{1}{\sigma_0^2}}, \frac{1}{n+\frac{1}{\sigma_0^2}})$. Since $\mathbb{E}[X_i] = 0.9 \times 1 + 0.1 \times 20 = 2.9$
330 the location is therefore in expectation $\mathbb{E}[\frac{\sum_{i=1}^n x_i}{n+\frac{1}{\sigma_0^2}}] = \frac{n}{n+\frac{1}{\sigma_0^2}} 2.9 \approx 2.9$, the contamination significantly
331 impacted the posterior. Similarly, for any model p_θ , the Wasserstein distance with the contaminated
332 mixture $0.9\mathcal{N}(1, 1) + 0.1\mathcal{N}(20, 1)$ will be high, and empirically all of the T iterations are rejected
333 for all the values of ϵ considered. Thus, we disregard the Wasserstein distance from the experiment
334 and compare the performance of MMD to that of d_{KT} . We also consider concurrent methods out of
335 our scope such as MMD with the unbounded energy kernel: $k(x, y) = \frac{1}{2}(\|x\| + \|y\| - \|x - y\|)$
336 Sejdinovic et al. [2013], and others displayed in Appendix B.2.

337 We measure the average Mean Square Error between the target parameter $\theta^* = 1$ and the accepted
338 $\theta_i \in L_\theta$: $\overline{MSE} = \frac{1}{|L_\theta|} \sum_{\theta_i \in L_\theta} \|\theta_i - \theta^*\|^2$ which also corresponds to the average of squared
339 Wasserstein 2-distance as $W_2^2(\mu^*, p_{\theta_i}) = \|\theta_i - \theta^*\|^2$ since we consider only Gaussians with same
340 variance. We picked $\sigma_0 = 5$ for the prior. We repeat 10 times the experiment with fresh samples,
341 the averaged results are shown in Table 1. As expected – and discussed in subsection 3.2 – MMD
342 (gaussian) is too lenient to accept. For $\epsilon = 0.05$ inferior to the contamination level (10%), it still
343 accept 11% of the times, while d_{KT} reject all the times, which can be understood as d_{KT} detecting
344 the contamination, that prevents to match with the Gaussian model. The energy kernel can not help
345 enough to beat d_{KT} . The densities of the obtained posteriors are shown alongside the target in Fig. (4)
346 and (5) in the Appendix B.2.

Table 1: Average MSE of ABC Results.

The Gaussian kernel is used with $\sigma = 1$ (as the variance of p_θ and μ^*). As expected, MMD is too lenient to accept most sampled θ_i leading to a high average MSE unless ϵ is carefully chosen. Whereas the proposed d_{KT} discriminates between the correct and the wrong θ_i for ϵ larger than the contamination threshold 0.1. MMD is assumed to use the Gaussian kernel while MMD_E denotes the MMD with the energy kernel.

ϵ	0.05			0.25			0.5		
distance	MMD	MMD _E	d_{KT}	MMD	MMD _E	d_{KT}	MMD	MMD _E	d_{KT}
#accept.	1092	0	0	2964	0	58	6168	846	828
MSE	0.19	N/A	N/A	1.29	N/A	0.03	7.47	0.17	0.12

347 **Particle Flow** We consider the performance of gradient descent when optimizing $\mu \mapsto d_{KT}(\mu, \nu)$
348 for discrete measures μ, ν on $\mathbb{R}^2$, given an initial point cloud (in red) and a target cloud of points
349 (in blue) both of $n = 100$ points. We run the scheme with a learning rate of 0.005 for 1000 steps,
350 using d_{KT} (Schatten 1-norm) and MMD (Schatten 2-norm), see Appendix B.3 (Figs. 6 and 7). We
351 use the Laplacian kernel: $k(x, y) = e^{-\frac{\|x-y\|_1}{\sigma}}$ where here $\|\cdot\|_1$ means the l_1 norm for vectors.
352 We choose a bandwidth $\sigma = 1$ (as the image size is a unit square) for d_{KT} and for MMD we use
353 k^2 as kernel to match the Schatten 2-norm (i.e. we use $\sigma = 0.5$ instead of $\sigma = 1$, and it gives a
354 better convergence). The inherent internal energy of MMD incites the point cloud to spread out and
355 therefore some particles are still left out far away from the target, which does not happen with d_{KT} .

356 6 Conclusion

357 We introduced a robust distance between probability measures, based on RKHS density (or covariance)
358 operators and their Schatten-1 norm. It is the greatest in a family of kernel-based IPM including
359 MMD, and so is more discriminative as shown in experiments. We show how to compute it between
360 discrete measures via a new kernel trick. Assuming some decay rate of the eigenvalues of the RKHS
361 density operator leads to a statistical convergence rate that can be close to $O(n^{-\frac{1}{2}})$. This implies the
362 first (dimension-independent) rates for the Kernel Bures Wasserstein distance. Future work includes
363 reducing computational complexity via Nystrom method, improving the dependence on the order of
364 decay α , as well as minimax lower bounds.

References

- Ron Amit, Baruch Epstein, Shay Moran, and Ron Meir. Integral probability metrics PAC-bayes bounds. *Advances in Neural Information Processing Systems*, 35:3123–3136, 2022.
- Michael Arbel, Anna Korba, Adil Salim, and Arthur Gretton. Maximum Mean Discrepancy Gradient Flow. *Advances in Neural Information Processing Systems*, 32, 2019.
- Martin Arjovsky, Soumith Chintala, and Léon Bottou. Wasserstein generative adversarial networks. In *International Conference on Machine Learning*, pages 214–223. PMLR, 2017.
- Léon Autonne. *Sur les matrices hypohermitiennes et sur les matrices unitaires*. A. Rey, 1915.
- Francis Bach. Information theory with kernel methods. *IEEE Transactions on Information Theory*, 69(2):752–775, 2022.
- Ayanendranath Basu, Hiroyuki Shioya, and Chanseok Park. *Statistical Inference: the Minimum Distance Approach*. CRC press, 2011.
- Espen Bernton, Pierre E Jacob, Mathieu Gerber, and Christian P Robert. Approximate Bayesian computation with the Wasserstein distance. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 81(2):235–269, 2019.
- Rajendra Bhatia, Tanvi Jain, and Yongdo Lim. On the Bures-Wasserstein distance between positive definite matrices. *Expositiones Mathematicae*, 37(2):165–191, 2019.
- Jeremiah Birrell, Paul Dupuis, Markos A Katsoulakis, Yannis Pantazis, and Luc Rey-Bellet. (f, Gamma)-Divergences: Interpolating between f-Divergences and Integral Probability Metrics. *Journal of Machine Learning Research*, 23(39):1–70, 2022.
- Gilles Blanchard, Olivier Bousquet, and Laurent Zwald. Statistical properties of kernel principal component analysis. *Machine Learning*, 66:259–294, 2007.
- Clémentine Chazal, Anna Korba, and Francis Bach. Statistical and Geometrical properties of regularized Kernel Kullback-Leibler divergence. *Advances in Neural Information Processing Systems*, 2024.
- Zonghao Chen, Aratrika Mustafi, Pierre Glaser, Anna Korba, Arthur Gretton, and Bharath K Sriperumbudur. (De)-regularized Maximum Mean Discrepancy Gradient Flow. *arXiv preprint arXiv:2409.14980*, 2024.
- Moulines Eric, Francis Bach, and Zaïd Harchaoui. Testing for homogeneity with kernel Fisher discriminant analysis. *Advances in Neural Information Processing Systems*, 20, 2007.
- Jean Feydy, Thibault Séjourné, François-Xavier Vialard, Shun-ichi Amari, Alain Trounev, and Gabriel Peyré. Interpolating between Optimal Transport and MMD using Sinkhorn Divergences. In *The 22nd International Conference on Artificial Intelligence and Statistics*, pages 2681–2690, 2019.
- David T Frazier. Robust and efficient approximate Bayesian computation: A minimum distance approach. *arXiv preprint arXiv:2006.14126*, 2020.
- Christopher A Fuchs and Jeroen Van De Graaf. Cryptographic distinguishability measures for quantum-mechanical states. *IEEE Transactions on Information Theory*, 45(4):1216–1227, 1999.
- Aude Genevay, Gabriel Peyré, and Marco Cuturi. Learning Generative Models with Sinkhorn Divergences. In *International Conference on Artificial Intelligence and Statistics*, pages 1608–1617. PMLR, 2018.
- Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative Adversarial Nets. *Advances in neural information processing systems*, 27, 2014.
- Arthur Gretton, Karsten M Borgwardt, Malte J Rasch, Bernhard Schölkopf, and Alexander Smola. A Kernel Two-Sample Test. *The Journal of Machine Learning Research*, 13(1):723–773, 2012.

410 Omar Hagrass, Bharath K Sriperumbudur, and Bing Li. Spectral Regularized Kernel Goodness-of-Fit
411 Tests. *Journal of Machine Learning Research*, 25(309):1–52, 2024.

412 Johannes Hertrich, Christian Wald, Fabian Altekrüger, and Paul Hagemann. Generative sliced MMD
413 flows with Riesz kernels. *arXiv preprint arXiv:2305.11463*, 2023.

414 Peter J Huber. Robust Estimation of a Location Parameter. *The Annals of Mathematical Statistics*, 35
415 (1):73–101, 1964.

416 Hisham Husain and Jeremias Knoblauch. Adversarial interpretation of Bayesian inference. In
417 *International Conference on Algorithmic Learning Theory*, pages 553–572. PMLR, 2022.

418 Bai Jiang. Approximate Bayesian computation with Kullback-Leibler divergence as data discrepancy.
419 In *International Conference on Artificial Intelligence and Statistics*, pages 1711–1721. PMLR,
420 2018.

421 Benoit Kloeckner. Approximation by finitely supported measures. *ESAIM: Control, Optimisation
422 and Calculus of Variations*, 18(2):343–359, 2012.

423 Anna Korba, Pierre-Cyril Aubin-Frankowski, Szymon Majewski, and Pierre Ablin. Kernel stein
424 discrepancy descent. In *International Conference on Machine Learning*, pages 5719–5730. PMLR,
425 2021.

426 Sirio Legramanti, Daniele Durante, and Pierre Alquier. Concentration of discrepancy-based ABC via
427 Rademacher complexity. *arXiv preprint arXiv:2206.06991*, 2022.

428 Chun-Liang Li, Wei-Cheng Chang, Yu Cheng, Yiming Yang, and Barnabás Póczos. Mmd gan:
429 Towards deeper understanding of moment matching network. *Advances in Neural Information
430 Processing Systems*, 30, 2017.

431 Bruce G. Lindsay, Marianthi Markatou, Surajit Ray, Ke Yang, and Shu-Chuan Chen. Quadratic
432 distances on probabilities: A unified foundation. *The Annals of Statistics*, 36(2):983 – 1006, 2008.

433 Youssef Mroueh, Tom Sercu, and Vaibhava Goel. Mcgan: Mean and covariance feature matching
434 gan. In *International Conference on Machine Learning*, pages 2527–2535. PMLR, 2017.

435 Alfred Müller. Integral Probability Metrics and Their Generating Classes of Functions. *Advances in
436 applied probability*, 29(2):429–443, 1997.

437 Sebastian Neumayer, Viktor Stein, Gabriele Steidl, and Nicolaj Rux. Wasserstein gradient flows
438 for Moreau envelopes of f-divergences in reproducing kernel Hilbert spaces. *arXiv preprint
439 arXiv:2402.04613*, 2024.

440 Jung Hun Oh, Maryam Pouryahya, Aditi Iyer, Aditya P Apte, Joseph O Deasy, and Allen Tannenbaum.
441 A novel kernel Wasserstein distance on Gaussian measures: an application of identifying dental
442 artifacts in head and neck computed tomography. *Computers in biology and medicine*, 120:103731,
443 2020.

444 Alessandro Rudi, Guillermo D Canas, and Lorenzo Rosasco. On the Sample Complexity of Subspace
445 Learning. *Advances in Neural Information Processing Systems*, 26, 2013.

446 Dino Sejdinovic, Bharath Sriperumbudur, Arthur Gretton, and Kenji Fukumizu. Equivalence of
447 distance-based and rkhs-based statistics in hypothesis testing. *The Annals of Statistics*, pages
448 2263–2291, 2013.

449 Barry Simon. *Trace ideals and their applications*. Number 120. American Mathematical Society,
450 2005.

451 Bharath K Sriperumbudur and Nicholas Sterge. Approximate kernel PCA: Computational versus
452 statistical trade-off. *The Annals of Statistics*, 50(5):2713–2736, 2022.

453 Bharath K Sriperumbudur, Kenji Fukumizu, Arthur Gretton, Bernhard Schölkopf, and Gert RG
454 Lanckriet. Non-parametric Estimation of Integral Probability Metrics. In *2010 IEEE International
455 Symposium on Information Theory*, pages 1428–1432. IEEE, 2010.

- 456 Ingo Steinwart. On the Influence of the Kernel on the Consistency of Support Vector Machines.
457 *Journal of Machine Learning Research*, 2(Nov):67–93, 2001.
- 458 Nicholas Sterge and Bharath K Sriperumbudur. Statistical Optimality and Computational Efficiency
459 of Nystrom Kernel PCA. *Journal of Machine Learning Research*, 23(337):1–32, 2022.
- 460 Nicholas Sterge, Bharath Sriperumbudur, Lorenzo Rosasco, and Alessandro Rudi. Gain with no
461 Pain: Efficiency of Kernel-PCA by Nyström Sampling. In *International Conference on Artificial
462 Intelligence and Statistics*, pages 3642–3652. PMLR, 2020.
- 463 Teiji Takagi. On an Algebraic Problem related to an Analytic Theorem of Carathéodory and Fejér
464 and on an Allied Theorem of Landau. In *Japanese Journal of Mathematics: transactions and
465 abstracts*, volume 1, pages 83–93. The Mathematical Society of Japan, 1924.
- 466 Simon Tavaré, David J Balding, Robert C Griffiths, and Peter Donnelly. Inferring coalescence times
467 from DNA sequence data. *Genetics*, 145(2):505–518, 1997.
- 468 Joel A Tropp et al. An Introduction to Matrix Concentration Inequalities. *Foundations and Trends®
469 in Machine Learning*, 8(1-2):1–230, 2015.
- 470 Cédric Villani. *Optimal transport: Old and New*, volume 338. Springer, 2009.
- 471 John Watrous. *The Theory of Quantum Information*. Cambridge university press, 2018.
- 472 Jacob Wolfowitz. The Minimum Distance Method. *The Annals of Mathematical Statistics*, pages
473 75–88, 1957.
- 474 Zhen Zhang, Mianzhi Wang, and Arye Nehorai. Optimal transport in reproducing kernel hilbert
475 spaces: Theory and applications. *IEEE transactions on pattern analysis and machine intelligence*,
476 42(7):1741–1754, 2019.

NeurIPS Paper Checklist

(i) Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: Yet the abstract reflect our claims, supported by theorems and simulations in the paper.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

(ii) Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: We mention the cubic complexity of the algorithm computing our novel distance, and we stated with Assumptions to which scope of kernels we can apply our results.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

(iii) Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: We clearly numbered our assumptions and reference them and we have complete proofs in the appendix.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

(iv) **Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: We explicitly state our parameters in the Experiments section and we also provide supplementary code.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

(v) **Open access to data and code**

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The data are simulated and their simulation process are provided in the code.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

(vi) **Experimental setting/details**

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: We include hyperparameters such as kernel bandwidth in our code.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

(vii) **Experiment statistical significance**

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The standard deviations over the 10 runs of 10000 iterations are included in the appendix.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer “Yes” if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).

- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

648 (viii) **Experiments compute resources**

649 Question: For each experiment, does the paper provide sufficient information on the com-
650 puter resources (type of compute workers, memory, time of execution) needed to reproduce
651 the experiments?

652 Answer: [Yes]

653 Justification: We stated the complexity of our algorithm at the end of the section 2.3 and we
654 specify our computer setting in Appendix B.

655 Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

664 (ix) **Code of ethics**

665 Question: Does the research conducted in the paper conform, in every respect, with the
666 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

667 Answer: [Yes]

668 Justification: This is a theoretical work.

669 Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

675 (x) **Broader impacts**

676 Question: Does the paper discuss both potential positive societal impacts and negative
677 societal impacts of the work performed?

678 Answer: [NA]

679 Justification: This is a theoretical work.

680 Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.

- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

(xi) **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: This is a theoretical work with simulated data.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

(xii) **Licenses for existing assets**

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The code is mostly original, external code is mostly packages used which can be seen as import.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

(xiii) **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: Code are separated between mathematical distances functions and simulation experiments. A README is also provided.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

(xiv) **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: This is a theoretical work.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

(xv) **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: This is a theoretical work.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.

787 • We recognize that the procedures for this may vary significantly between institutions
788 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
789 guidelines for their institution.
790 • For initial submissions, do not include any information that would break anonymity (if
791 applicable), such as the institution conducting the review.

792 (xvi) **Declaration of LLM usage**

793 Question: Does the paper describe the usage of LLMs if it is an important, original, or
794 non-standard component of the core methods in this research? Note that if the LLM is used
795 only for writing, editing, or formatting purposes and does not impact the core methodology,
796 scientific rigorousness, or originality of the research, declaration is not required.

797 Answer: [No]

798 Justification: This work only concerns kernel methods and not NLP.

799 Guidelines:

800 • The answer NA means that the core method development in this research does not
801 involve LLMs as any important, original, or non-standard components.

802 • Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>)
803 for what should or should not be described.