
Fence off Anomaly Interference: Cross-Domain Distillation for Fully Unsupervised Anomaly Detection

Anonymous Author(s)

Affiliation

Address

email

Abstract

Fully Unsupervised Anomaly Detection (FUAD) is a practical extension of Unsupervised Anomaly Detection (UAD), aiming to detect anomalies without any labels even when the training set may contain anomalous samples. To achieve FUAD, we pioneer the introduction of Knowledge Distillation (KD) paradigm based on teacher-student framework into the FUAD setting. However, due to the presence of anomalies in the training data, traditional KD methods risk enabling the student to learn the teacher’s representation of anomalies under FUAD setting, thereby resulting in poor anomaly detection performance. To address this issue, we propose a novel Cross-Domain Distillation (CDD) framework based on the widely studied reverse distillation (RD) paradigm. Specifically, we design a Domain-Specific Training, which divides the training set into multiple domains with lower anomaly ratios and train a domain-specific student for each. Cross-Domain Knowledge Aggregation is then performed, where pseudo-normal features generated by domain-specific students collaboratively guide a global student to learn generalized normal representations across all samples. Experimental results on noisy versions of the MVTec AD and VisA datasets demonstrate that our method achieves significant performance improvements over the baseline, validating its effectiveness under FUAD setting.

1 Introduction

In the field of industrial image anomaly detection, acquiring or predefining anomalous samples is often impractical. Consequently, Unsupervised Anomaly Detection (UAD), which relies only on normal samples for training, has been extensively studied [5, 7]. To tackle the challenges of UAD task, a variety of methods are proposed, such as those based on memory banks [26, 2], anomaly synthesis [19, 35], and image reconstruction [4, 1]. In recent years, UAD methods based on Knowledge Distillation (KD) have gained increasing attention [6]. Compared to other techniques, they show greater potential in pixel-level anomaly localization. The KD-based UAD methods typically employ a teacher-student framework, which allows the student network to imitate the feature representations of the teacher on normal samples. Since the student is never exposed to anomalous samples during training, its ability to generate teacher’s anomaly features is limited. This discrepancy in feature mimicking performance becomes a useful signal for anomaly identification.

In real-world scenarios, however, it is often inevitable that the collected data contain a small proportion of anomalous samples. Relying entirely on manual data cleaning incurs high labor costs. This motivates the need for Fully Unsupervised Anomaly Detection (FUAD), a more practical and challenging setting where the training set may contain unlabeled anomalous samples. Although several studies have begun to explore this task, most existing methods rely heavily on memory




Figure 1: (a) Simplified schematic of Cross-Domain Distillation with 3 domains, where the top represents Domain-Specific Training and the bottom depicts Cross-Domain Knowledge Aggregation. (b) Quantitative comparison against other FUAD methods on MVTec AD-noise-0.1 (with 10% anomaly ratio in the train set) [5], with P-AUC (x-axis), PRO (y-axis), and circle size indicating FPS (larger means faster inference), proves that CDD has the best overall performance. (c) Qualitative comparison with other FUAD methods on MVTec AD-noise-0.1.

banks [17, 24, 16], which introduces additional storage overhead. In contrast, Knowledge Distillation paradigm offers a storage-efficient alternative, yet its potential in FUAD has not been fully explored.

We make two key observations regarding the data under the FUAD setting: (1) From a probabilistic perspective, normal pixels still dominate the training set despite the presence of noise; (2) In terms of feature distribution, teacher representations of normal samples tend to be more compact and stable, making them easier for the student network to learn, whereas anomaly features are more dispersed and less likely to be captured. These insights suggest that, even under the FUAD setting, the student network trained via KD inherently focuses on learning normal representations, leading to poor fitting in anomalous regions. This makes the discrepancy between teacher and student features a reliable signal for anomaly detection, which indicates the feasibility and potential of applying the KD paradigm to the FUAD setting.

However, meanwhile, Knowledge Distillation faces a long-standing over-generalization problem [28, 32, 27]. Even though the student network is trained to learn teacher’s representations only on normal pixels, its learned representation ability may still generalize to anomalous ones, which leads to miss detections when testing on anomalous samples. This issue becomes more pronounced under the FUAD setting. If a certain type of anomaly appears frequently in the training data, the student may learn its common feature patterns and become capable of generating teacher-like representations for similar anomalies during inference.

To address the above challenge, we propose a novel cross-domain distillation framework for FUAD, built upon the widely studied KD-based UAD method Reverse Distillation (RD) [9] with an encoder-decoder architecture. First, our intuition is that *reducing the probability of anomalous samples being learned during training mitigates the student’s tendency to overfit to teacher’s anomaly features*. To achieve this, we design a domain division mechanism that distributes high-confidence normal samples across multiple domains while dispersing potentially anomalous samples, thereby lowering the anomaly ratio within each domain without discarding any data. Considering that data distributions vary across domains and that RD’s student decoder generates anomaly-free features for unseen anomalous samples, we hypothesize that *domain-specific students trained on different domains produce pseudo-normal features when applied to other domains*. Building on this insight, we introduce a **Cross-Domain Distillation (CDD)** framework: for each domain, we utilize domain-specific students from other domains to generate pseudo-normal features for its samples, guiding the training of a global student decoder. This global student learns to produce anomaly-free features across all samples, both normal and anomalous. Finally, the distance between the features generated by the global student decoder and the teacher encoder is used to detect and localize anomalies. Our contributions are summarized as follows:

- We are the first to explore the application of the knowledge distillation paradigm to the Fully Unsupervised Anomaly Detection task.
- We propose Domain-Specific Training (DST) as in Fig. 1 (a), which first performs Confidence-Guided Domain Construction to build data domains with low anomaly proba-

bility. Then, each domain is used to train a domain-specific student via Domain-Specific Distillation with Regularization.

- We introduce Cross-Domain Knowledge Aggregation (CDKA), where domain-specific students provide pseudo-normal features for each sample to train a global student that integrates information across all domains as depicted in Fig. 1 (a).
- Experimental verification shows that our CDD is significantly higher than the baseline RD, and has better performance and faster inference speed than the previous FUAD methods.

2 Related Work

Unsupervised Anomaly Detection. Unsupervised Anomaly Detection (UAD) has been widely studied in recent years due to its ability to operate without requiring anomalous samples during training. Existing methods are broadly categorized into the following types: (1) reconstruction-based generative models [4, 1, 29, 25, 33, 37], which learn to reconstruct only normal samples and identify anomalies based on reconstruction errors during inference; (2) density estimation-based methods [8, 12, 39], which assume that normal samples follow a specific distribution in the feature space and detect deviations from this distribution; (3) synthetic anomaly-based approaches [19, 35, 21, 38], which generate pseudo-anomalies using image transformations, external generators, or diffusion models to enhance the model’s ability to perceive anomalies; and (4) methods that incorporate pre-trained models and memory bank mechanisms [26, 2, 15], comparing the features of test samples with those of normal samples to identify anomalies. In recent years, Knowledge Distillation-based UAD methods [6, 20, 28, 32, 27, 3, 22] using the teacher-student framework have emerged as excellent methods for anomaly localization. These methods learn representations of normal regions and detect anomalies by measuring the discrepancy in features between the teacher and student networks on anomalous regions. To mitigate the student’s over-generalization to anomalies, some studies introduce heterogeneous architectures or reverse information flow, such as Reverse Distillation [9] and its variants [30, 13, 11, 18, 14, 36], which further improve anomaly detection accuracy.

Fully Unsupervised Anomaly Detection. Fully Unsupervised Anomaly Detection (FUAD) has attracted increasing attention, owing to its ability to operate without manual annotations and its suitability for tackling noisy training data in real-world scenarios [31]. Existing methods are categorized as follows: (1) SoftPatch [17], based on PatchCore [26], adopts a memory-based patch-level denoising strategy using noise discriminators to mitigate overconfidence. (2) InReaCh [24] builds detection models by associating high-confidence patch channels across training images. (3) FUN-AD [16] leverages nearest-neighbor distances and class homogeneity, employing an iteratively reconstructed memory bank (IRMB) to handle noisy data. However, these methods often rely on explicit memory banks, which impose storage burdens in practice. Knowledge Distillation has shown strong potential in unsupervised anomaly localization without additional storage, but its application to FUAD remains unexplored. This work aims to explore this promising direction.

3 Motivation and Assumptions

3.1 Rethinking Reverse Distillation for FUAD

What is Reverse Distillation? Early KD-based AD methods typically adopt a homogeneous teacher-student framework, where the student only learns the teacher’s representation ability on normal samples. During inference, anomalies are detected by measuring the discrepancy between teacher and student features. Reverse Distillation (RD) [9] builds upon KD by introducing an encoder-decoder structure. The teacher network is a frozen encoder, while the student consists of a trainable one-class bottleneck embedding (OCBE) module $\mathcal{B}(\cdot; \phi)$ and a trainable decoder $\mathcal{D}_S(\cdot; \psi)$.

Let the training set be $\mathcal{I}_{train}$. Given a training image $I_i^{train} \in \mathcal{I}_{train}$, the teacher extracts multi-layer features $\mathcal{F}_{\mathcal{T},i} = \mathcal{T}(I_i^{train}) = \{f_{\mathcal{T},i}^l\}_{l=1}^L$, which are then reconstructed by the student network as $\mathcal{F}_{\mathcal{S},i} = \mathcal{S}(\mathcal{F}_{\mathcal{T},i}; \theta_S) = \{f_{\mathcal{S},i}^l\}_{l=1}^L$. The student network is denoted as $\mathcal{S}(\cdot; \theta_S)$, with parameters $\theta_S = \{\phi, \psi\}$. The training objective is to minimize the cosine distance between teacher and student features across all $L = 3$ layers on normal samples as:

$$\cos(f_1, f_2) = \frac{f_1 \cdot f_2}{\|f_1\| \|f_2\|} \quad (1)$$

123

$$\ell_{cos}(\mathcal{F}_T, \mathcal{F}_S) = \sum_{l=1}^L \left(1 - \cos(f_T^l, f_S^l)\right) \quad 1 \quad (2)$$

124

$$\arg \min_{\theta_S} \mathbb{E}_{I_i \sim \mathcal{I}_{train}} \ell_{cos}(\mathcal{F}_{T,i}, \mathcal{S}(\mathcal{F}_{T,i}; \theta_S)) \quad (3)$$

125 **Why does RD Work for FUAD?** Although Reverse Distillation (RD) is initially designed for training with only 2
 126 normal samples, it demonstrates strong adaptability in Fully Unsupervised Anomaly Detection. We attribute this
 127 to two key factors:

128 (1) **Probability Perspective** - Dominance of Normal Samples 3

129 In industrial scenarios, normal samples are much more common than anomalies, which results in low proportion
 130 of anomalous images in the training set. Moreover, anomalies typically occupy only a small region within an
 131 image. Consequently, the student network, driven by the dominance of normal samples, primarily learns to
 132 represent normal features, while the sparsity of anomalies limits their impact on the optimization process.

133 (2) **Distribution Perspective** - Concentrated Normal vs. Diverse Anomalous 4

134 Normal samples exhibit compact and consistent feature patterns, while anomalous samples are diverse and 5
 135 scattered. This makes it difficult for the student to generalize learned anomalous features.

136 **Challenges of Applying RD to FUAD.** In FUAD task, the training set naturally includes a certain proportion 6
 137 of anomalous samples. If specific anomaly patterns appear repeatedly during training, the student can easily
 138 learn to reconstruct the teacher features of these common anomalies. This results in poor discrimination against
 139 similar anomalies during testing and further intensifies over-generalization. Therefore, the key challenge in
 140 applying RD to FUAD is *how to prevent the student from modeling common anomaly patterns during training,*
 141 *to ensure that it generates anomaly-free features.*

142 3.2 Assumptions 7

143 To address over-generalization problem in FUAD, we propose two assumptions based on the diversity and 8
 144 sparsity of anomalies, guiding the following design of our method Cross-Domain Distillation.

145 **Assumption 1 (Limited Representation of Rare Anomalies)** *When a particular anomaly type is sufficiently 9*
 146 *rare in training data, the student fails to learn its corresponding teacher anomaly features, and instead tends to*
 147 *produce features that closely resemble normal patterns.*

148 Due to the consistency of normal samples and the diversity of anomalies (i.e., anomalies exhibit multiple distinct 10
 149 patterns), we assume the training set contains one normal type and M_{train} anomaly types, expressed as:

$$\mathcal{I}_{train} = \mathcal{N} \cup \mathcal{A} = \mathcal{N} \cup \bigcup_{m=1}^{M_{train}} \mathcal{A}_m \quad 11 \quad (4)$$

150 where $\mathcal{N}$ denotes the set of normal samples, $\mathcal{A}_m$ denotes the set of the m -th anomaly type, and: 12

$$\mathbb{P}(\mathcal{N}) \gg \mathbb{P}(\mathcal{A}_m) \quad \forall m = 1, \dots, M_{train} \quad 13 \quad (5)$$

151 Following Empirical Risk Minimization (ERM), the training objective is to minimize the distance between 14
 152 student features and teacher features over all samples. The empirical risk can be expressed as

$$\mathcal{L} = \mathbb{P}(\mathcal{N}) \cdot \mathbb{E}_{I_i \sim \mathcal{N}} [\ell_{cos}(\mathcal{F}_{T,i}, \mathcal{F}_{S,i})] + \sum_{m=1}^{M_{train}} \mathbb{P}(\mathcal{A}_m) \cdot \mathbb{E}_{I_j \sim \mathcal{A}_m} [\ell_{cos}(\mathcal{F}_{T,i}, \mathcal{F}_{S,i})] \quad 15 \quad (6)$$

153 The gradient of parameters θ_S is: 16

$$\frac{\partial \mathcal{L}}{\partial \theta_S} = \mathbb{P}(\mathcal{N}) \cdot \mathbb{E}_{I_i \sim \mathcal{N}} \left[\frac{\partial \ell_{cos}}{\partial \theta_S} \right] + \sum_{m=1}^{M_{train}} \mathbb{P}(\mathcal{A}_m) \cdot \mathbb{E}_{I_j \sim \mathcal{A}_m} \left[\frac{\partial \ell_{cos}}{\partial \theta_S} \right] \quad 17 \quad (7)$$

154 If $\mathbb{P}(\mathcal{A}_m)$ is small enough, the contribution of the anomaly type $\mathcal{A}_m$ to the gradient is negligible. Thus, the 18
 155 student receives limited learning signals for this type of anomaly and fails to reconstruct the corresponding
 156 teacher features effectively.

157 **Assumption 2 (Lack of Cross-Anomaly Generalization)** *Even if a student learns to reconstruct some specific 19*
 158 *anomaly patterns during training, this reconstruction ability is not generalized to other unseen anomaly types.*

159 This assumption is based on the diversity of anomalies. Anomalies are often unstructured and come from 20
 160 different sources or physical mechanisms. As a result, they follow multiple, structurally different patterns in the
 161 feature space:

$$\mathcal{F}_{T,i} | I_i \in \mathcal{A}_m \sim \mathcal{P}_m \quad 21 \quad (8)$$

162 where each pattern $\mathcal{P}_m$ represents the teacher's feature distribution for the m -th type of anomaly. The total 22
 163 number of types is M , which may even be infinite in practice.




Figure 2: Overall framework of our proposed CDD.

During training, the student network only sees a subset of these anomaly patterns:

$$\mathcal{P}_{\text{train}} = \{\mathcal{P}_1, \dots, \mathcal{P}_{M_{\text{train}}}\}, \quad M_{\text{train}} \ll M \quad (9)$$

According to the **No Free Lunch** theorem, if an input anomalous sample $I_j \sim \mathcal{P}_{m'}$ with $\mathcal{P}_{m'} \notin \mathcal{P}_{\text{train}}$, its distribution is outside the training support. Then, the student may fail to generate the correct teacher features:

$$\mathcal{F}_{S,j} \not\approx \mathcal{F}_{T,j}, \quad I_j \notin \mathcal{P}_{\text{train}} \quad (10)$$

Since normal samples dominate the training set, the student tends to generate features similar to the normal distribution.

4 Method

Problem Definition. In FUAD, we denote the training set as $\mathcal{I}_{\text{train}} = \{I_i^{\text{train}}\}_{i=1}^N$, where each image $I_i^{\text{train}} \in \{\mathcal{N}, \mathcal{A}\}$ is unlabeled and may be normal or anomalous. The test set $\mathcal{I}_{\text{test}} = \{I_j^{\text{test}}\}_{j=1}^M$ comprises both normal and anomalous images, with normal samples following the same distribution as $\mathcal{I}_{\text{train}}$. The objective is to learn the distribution of normal samples from $\mathcal{I}_{\text{train}}$ to detect anomalies in $\mathcal{I}_{\text{test}}$.

Overview. Fig. 2 illustrates the training process of each epoch (top and lower right) and the inference process (lower left). All teacher and student networks follow the design of Reverse Distillation. The teacher is a WideResNet-50 [34] pre-trained on ImageNet [10]. And each student includes an OCBE module and a decoder.

Each training epoch consists of two stages: Domain-Specific Training and Cross-Domain Knowledge Aggregation. In the first stage, we propose Confidence-Guided Domain Construction to extract high-confidence normal samples from the original training set and use them as the intersection between multiple data domains. In this way, each domain has a reduced anomaly ratio compared to the full dataset. Then, we train a domain-specific student for each domain using Domain-Specific Distillation with Regularization. Based on *Assumption 1*, these students ease off from modeling anomaly features and thus focus on modeling normal features in their local domains. The second stage Cross-Domain Knowledge Aggregation mainly explains how to use the domain-specific students obtained in the first stage to train a global student that reconstructs normal features on all samples. According to *Assumption 2*, for anomalous samples in a specific domain k , domain-specific students that are not trained on domain k still generates normal-like features. We use these features as pseudo-normal supervision signals to perform Cross-Domain Pseudo-Normal Feature Distillation for the global student. After that, we further distill the global student using the teacher on high-confidence normal samples, enabling it to effectively learn the reliable reconstruction of normal patterns.

The lower left part of Fig. 2 depicts the inference process. During inference, for each image $I_j^{\text{test}} \in \mathcal{I}_{\text{test}}$, cosine distances across multi-layer features generated by the teacher $\mathcal{T}$ and the global student trained for E epochs $\mathcal{S}_E^{\text{Glo}}$ are fused to generate a pixel-level anomaly map $\mathcal{M}$, whose maximum value serves as the image-level anomaly score s :

$$\mathcal{M}(h, w) = \sum_{l=1}^L \left(1 - \cos(f_T^l(h, w), f_{S_E^{\text{Glo}}}^l(h, w)) \right), \quad s = \max(\mathcal{M}) \quad (11)$$

195 4.1 Domain-Specific Training

196 **Confidence-Guided Domain Construction** Based on *Assumption 1*, reducing the anomaly probability 1
197 in the training set helps the student better learn normal patterns. A naive way to achieve this is to retain only
198 high-confidence normal samples or discard low-confidence anomalous ones. However, such strategies fail to
199 fully utilize the training data, as potentially useful normal regions are also discarded along with the anomalies.

200 To address this issue, we introduce a confidence-guided strategy on top of naive equal partitioning. Specifically, 2
201 we inject a portion of highly confident normal samples into each domain based on normality confidence scores,
202 which ensures that: (1) The anomaly ratio in each domain becomes lower than that in the original training set,
203 reducing the interference of anomalous samples on the modeling of normal patterns. (2) The normal distribution
204 in each domain remains more similar to the overall normal distribution, mitigating the negative impact of domain
205 partitioning on student normality modeling.

206 We use the features output by the global student of the previous epoch S_{e-1}^{Glo} as the basis for confidence evaluation. 3
207 For each training sample I_i , the average cosine similarity between teacher features f_T and global student features
208 $f_{S_{e-1}^G}$ across L layers is calculated to obtain the corresponding $Conf_i$:

$$Conf_i = \sum_{l=1}^L \left\{ \frac{1}{H_l W_l} \sum_{h=1}^{H_l} \sum_{w=1}^{W_l} \cos(f_{T,i}^l(h, w), f_{S_{e-1}^{Glo},i}^l(h, w)) \right\} \quad 4 \quad (12)$$

209 All samples are sorted by confidence in descending order. The top $r(e)$ samples form the high-confidence set 5
210 $\mathcal{D}^{HC}$. The confidence threshold $r(e)$ increases with training, up to 50%. Let e be the current epoch and E the
211 total epochs, $r(e)$ is calculated as

$$r(e) = \min\left(\frac{e}{E}, 0.5\right) \quad 6 \quad (13)$$

212 The remaining low-confidence samples are randomly and evenly divided into K subsets, denoted 7
213 $\mathcal{D}_k^{LC}$, $k = 1, \dots, K$. By combining $\mathcal{D}^{HC}$ and $\mathcal{D}_k^{LC}$, each domain is expressed as

$$\mathcal{D}_k = \mathcal{D}^{HC} \cup \mathcal{D}_k^{LC}, \quad k = 1, \dots, K. \quad 8 \quad (14)$$

214 **Domain-Specific Distillation with Regularization** After domain construction, we train a corresponding 9
215 domain-specific student S_k^{DS} for the k -th domain, who learns to reconstruct the features of samples within its
216 corresponding domain. The initial parameters of each domain-specific student are inherited from the global
217 student of the previous epoch S_{e-1}^{Glo} . This training process of S_k^{DS} follows the basic framework of Reverse
218 Distillation, which minimizes the cosine distance between the features generated by the student $\mathcal{F}_{S_k^{DS}}$ and the
219 features of the teacher $\mathcal{F}_T$. In this way, the domain-specific students are able to model the teacher's feature
220 representation ability of data in their local domain.

221 However, even with controlled anomaly ratios and dispersed common anomalies, the domain-specific student 10
222 may still learn representations of abnormal samples, especially when a particular type of anomaly is overly
223 represented in the domain. To further tackle this problem, we introduce pseudo-normal features generated by
224 the global student obtained from the previous epoch S_{e-1}^{Glo} as the regularization signal. As the global student
225 becomes more and more capable of modeling normal patterns during training, it provides useful guidance to help
226 domain-specific students avoid over-learning anomaly features. The loss $\mathcal{L}_k^{DS}$ used to train each domain-specific
227 student S_k^{DS} combines two terms: the primary distillation loss (from the teacher) and the regularization loss
228 (from the global student), which is expressed as

$$\mathcal{L}_k^{DS} = \mathbb{E}_{I_i \sim \mathcal{D}_k} \left(\underbrace{\ell_{cos}(\mathcal{F}_{T,i}, \mathcal{F}_{S_k^{DS},i})}_{\text{Teacher Guidance } \mathcal{L}_k^{DS-T}} + \lambda(e) \cdot \underbrace{\ell_{cos}(\mathcal{F}_{S_{e-1}^{Glo},i}, \mathcal{F}_{S_k^{DS},i})}_{\text{Regularization } \mathcal{L}_k^{DS-R}} \right) \quad 11 \quad (15)$$

229 where $\lambda(e)$ is a dynamic increasing coefficient that adjusts the regularization strength over the training epochs. 12
230 It is controlled using an S-shaped scheduling function with $p = 4.0$ as

$$\lambda(e) = \frac{(e/E)^p}{(e/E)^p + (1 - e/E)^p} \quad 13 \quad (16)$$

231 4.2 Cross-Domain Knowledge Aggregation 14

232 **Cross-Domain Pseudo-Normal Feature Distillation** Due to the high consistency of normal samples 15
233 across domains, domain-specific students reconstruct correct normal features on normal samples in all domains.
234 Based on *Assumption 2*, the diversity of anomalies prevents domain-specific students from generalizing to out-of-
235 domain anomaly patterns, even if they learn the reconstruction of anomaly features in local domains. Following
236 this idea, we propose using domain-specific students to generate pseudo-normal features for out-of-domain
237 samples, providing supervision for the training of the global student to generate normal features on all samples.
238 To prevent pseudo-normal feature contamination caused by some domain-specific students learning the ability to

Table 1: Anomaly detection and localization results I-AUC / P-AUC / PRO under *No Overlap* setting on MVTec AD-noise-0.1 with the best in bold, and the second best underlined.

Category	Unsupervised		Fully Unsupervised			
	RD [9]	URD [23]	SoftPatch [17]	InReaCh [24]	FUN-AD [16]	CDD (Ours)
bottle	0.997 / 0.983 / 0.955	0.992 / 0.984 / 0.961	1.000 / 0.986 / 0.956	1.000 / 0.981 / 0.915	1.000 / 0.992 / 0.960	1.000 / 0.987 / 0.959
cable	0.931 / 0.835 / 0.768	0.955 / 0.881 / 0.824	0.996 / 0.984 / 0.919	0.958 / 0.978 / 0.862	0.952 / 0.920 / 0.740	0.981 / 0.969 / 0.891
capsule	0.939 / 0.980 / 0.956	0.951 / 0.981 / 0.958	0.961 / 0.990 / 0.965	0.446 / 0.914 / 0.657	0.922 / 0.987 / 0.855	0.942 / 0.984 / 0.950
carpet	0.985 / 0.988 / 0.957	0.993 / 0.991 / 0.972	0.989 / 0.992 / 0.959	0.980 / 0.992 / 0.958	1.000 / 0.995 / 0.953	0.989 / 0.989 / 0.960
grid	0.956 / 0.994 / 0.979	1.000 / 0.990 / 0.976	0.965 / 0.991 / 0.963	0.917 / 0.983 / 0.929	0.991 / 0.993 / 0.935	1.000 / 0.992 / 0.976
hazelnut	1.000 / 0.992 / 0.936	0.994 / 0.993 / 0.953	1.000 / 0.994 / 0.942	0.997 / 0.988 / 0.907	0.999 / 0.991 / 0.885	1.000 / 0.993 / 0.945
leather	1.000 / 0.995 / 0.988	1.000 / 0.995 / 0.990	1.000 / 0.994 / 0.988	1.000 / 0.992 / 0.985	1.000 / 0.998 / 0.986	1.000 / 0.991 / 0.971
metal_nut	0.988 / 0.833 / 0.859	0.994 / 0.848 / 0.869	0.998 / 0.886 / 0.838	0.970 / 0.958 / 0.887	0.997 / 0.992 / 0.864	1.000 / 0.962 / 0.870
pill	0.960 / 0.966 / 0.956	0.961 / 0.956 / 0.950	0.953 / 0.977 / 0.945	0.889 / 0.956 / 0.883	0.939 / 0.972 / 0.893	0.971 / 0.978 / 0.958
screw	0.980 / 0.995 / 0.983	0.954 / 0.994 / 0.977	0.952 / 0.994 / 0.975	0.779 / 0.982 / 0.936	0.913 / 0.981 / 0.772	0.934 / 0.992 / 0.974
tile	0.988 / 0.961 / 0.858	1.000 / 0.964 / 0.897	1.000 / 0.959 / 0.878	0.999 / 0.965 / 0.878	0.999 / 0.978 / 0.939	0.997 / 0.955 / 0.879
toothbrush	1.000 / 0.991 / 0.939	1.000 / 0.992 / 0.943	1.000 / 0.986 / 0.915	0.990 / 0.989 / 0.904	0.972 / 0.981 / 0.850	0.997 / 0.987 / 0.916
transistor	0.943 / 0.882 / 0.753	0.948 / 0.901 / 0.812	0.996 / 0.952 / 0.819	0.929 / 0.982 / 0.786	0.962 / 0.975 / 0.520	0.998 / 0.980 / 0.831
wood	0.990 / 0.978 / 0.906	0.994 / 0.983 / 0.924	0.997 / 0.979 / 0.912	0.947 / 0.962 / 0.875	1.000 / 0.977 / 0.960	0.993 / 0.979 / 0.916
zipper	0.924 / 0.976 / 0.941	0.861 / 0.973 / 0.926	0.974 / 0.989 / 0.969	0.952 / 0.937 / 0.796	0.984 / 0.970 / 0.925	0.958 / 0.980 / 0.950
Average	0.972 / 0.957 / 0.916	0.973 / 0.962 / 0.929	0.985 / 0.977 / 0.930	0.917 / 0.971 / 0.877	0.975 / <u>0.980</u> / 0.869	<u>0.984</u> / 0.981 / 0.930

reconstruct certain types of teacher anomaly features, we design a Consensus-driven Pseudo-Normal Feature Selection strategy.

Specifically, we select the most "consensual" domain-specific student to generate the normal feature supervision for each sample. The core motivation is that for the same sample, multiple domain-specific students that have not been trained on the sample should generate similar normal features. In the implementation, we achieve pseudo-normal feature selection by eliminating outlier features that are more likely to be abnormal features from the output features of domain-specific students with the help of the global student from the previous epoch $\mathcal{S}_{e-1}^{Glo}$.

For a sample I_i from domain $\mathcal{D}_k$, we first extract features $\mathcal{F}_{\mathcal{S}_h^{DS},i} = \{f_{\mathcal{S}_h^{DS},i}^l\}_{l=1}^L, h = \{1, \dots, K\} \setminus k$ using the domain-specific students from domains $\mathcal{D}_h, h = \{1, \dots, K\} \setminus k$, and obtain the reference features $\mathcal{F}_{\mathcal{S}_{e-1}^{Glo},i} = \{f_{\mathcal{S}_{e-1}^{Glo},i}^l\}_{l=1}^L$ the global student from the previous epoch. We construct an affinity matrix $\text{Aff}_i \in \mathbb{R}^{(K-1) \times 1}$, where each element measures the cosine similarity between the flattened features of each student and the global student:

$$\text{Aff}_i(h) = \sum_{l=1}^L \cos(f_{\mathcal{S}_h^{DS},i}^l, f_{\mathcal{S}_{e-1}^{Glo},i}^l), \quad h = \{1, \dots, K\} \setminus k \quad (17)$$

The pseudo-normal feature for the training sample is then selected as the one with the highest similarity:

$$\mathcal{F}_{pseudo,i} = \mathcal{F}_{\mathcal{S}_{h^*}^{DS},i}, \quad h^* = \arg \max_h \text{Aff}_i(h) \quad (18)$$

However, the selected pseudo-normal features may still be contaminated with anomaly features. To prevent the trainable global student from overfitting these pseudo-normal features, we inject Gaussian noise with $\sigma_{noise} = 0.2$ as feature perturbation into its input:

$$\mathcal{F}_{\mathcal{S}_{e-1}^{Glo},i}^* = \mathcal{S}(\mathcal{F}_{\mathcal{T},i} + \delta; \theta_{\mathcal{S}_{e-1}^{Glo}}), \quad \delta \sim \mathcal{N}(0, \sigma_{noise}^2) \quad (19)$$

The loss of Cross-Domain Pseudo-Normal Feature Distillation $\mathcal{L}_{Glo}^{Cross}$ is then defined as:

$$\mathcal{L}_{Glo}^{Cross} = \sum_{k=1}^K \mathbb{E}_{I_i \sim \mathcal{D}_k} \ell_{cos}(\mathcal{F}_{pseudo,i}, \mathcal{F}_{\mathcal{S}_{e-1}^{Glo},i}^*) \quad (20)$$

Confident Distillation for High-Confidence Domain In addition to pseudo-normal feature guidance, we also leverage the previously defined high-confidence sample set $\mathcal{D}^{HC}$, using teacher features as direct supervision to further enhance the global student's ability to model true normal patterns:

$$\mathcal{L}_{Glo}^{HC} = \mathbb{E}_{I_i \sim \mathcal{D}^{HC}} \ell_{cos}(\mathcal{F}_{\mathcal{T},i}, \mathcal{F}_{\mathcal{S}_{e-1}^{Glo},i}) \quad (21)$$

5 Experiments

5.1 Experimental Setup

Datasets. We conduct experiments on two widely-used datasets: MVTec AD and VisA. Since both datasets are originally designed for unsupervised anomaly detection, we adapt them to the FUAD setting following SoftPatch [17]. Specifically, we keep the normal training images unchanged and randomly inject a portion of anomalous test samples into the training set with a predefined anomaly ratio R_{noise} . We evaluate under two settings: (1) *No overlap* setting, where injected anomalous samples are removed from the test set; and (2) *Overlap* setting, where these anomalies remain in the test set, making the task more challenging.

Table 2: Anomaly detection and localization results I-AUC / P-AUC / PRO under *Overlap* setting on MVTec AD-noise-0.1 with the best in bold. and the second best underlined.

	Unsupervised		Fully Unsupervised			
	RD [9]	URD [23]	SoftPatch [17]	InReaCh [24]	FUN-AD [16]	CDD (Ours)
Average	0.708 / 0.818 / 0.901	0.696 / 0.792 / <u>0.909</u>	0.984 / 0.957 / 0.915	0.879 / 0.943 / 0.861	<u>0.976</u> / 0.977 / 0.870	0.971 / <u>0.973</u> / 0.921

(a) I-AUC
(b) P-AUC
(c) PRO

Figure 3: Comparison of anomaly detection performance with baseline RD under different R_{noise} .

Implementation Details. We train a separate model for each category. The backbone follows RD, using a WideResNet-50 pretrained on ImageNet. Following SoftPatch [17], all images are resized to 256×256 and then center cropped to 224×224 during both training and inference. All domain-specific students and the global student are optimized by its own Adam optimizer with a learning rate of 0.005 and trained for 200 epochs. To smooth the obtained anomaly maps, we apply Gaussian filtering with $\sigma = 4$. All our experiments are performed on a single Nvidia GTX 3090 GPU.

Evaluation Metrics. We use the area under the ROC curve (AUC) at both image and pixel levels, denoted as I-AUC and P-AUC, to evaluate anomaly detection and localization performance. The per-region-overlap (PRO) metric is also reported to better evaluate the localization performance of anomalies with small sizes.

5.2 Anomaly Detection under FUAD setting

Results on MVTec AD. We evaluate our proposed Cross-Domain Distillation (CDD) on the MVTec-AD dataset with $R_{\text{noise}} = 0.1$, denoted as MVTec-AD-noise-0.1. CDD is compared with unsupervised KD-based UAD methods including RD [9], and FUAD methods, such as SoftPatch [17], InReaCh [24], and FUN-AD [16]. Tab. 1 and Tab. 2 present the anomaly detection and localization results under *No Overlap* and *Overlap* settings, respectively, where each method reports I-AUC, P-AUC, and PRO metrics, all reproduced through 200 epochs of model training under a unified dataset split. In the *No Overlap* setting, CDD matches SoftPatch’s I-AUC while achieving a P-AUC of 0.981 and PRO of 0.930, surpassing all methods including RD in pixel-level localization. In the *Overlap* setting, despite some methods’ performance dropping sharply, CDD retains robustness with a P-AUC of 0.973 and PRO of 0.921, significantly outperforming the baseline and demonstrating strong resistance to anomaly noise. Furthermore, we compare RD and CDD on MVTec AD-noise- $\{0.2-0.15\}$ as in Fig. 3. At low R_{noise} , CDD’s advantage over RD is subtle, but as R_{noise} rises, RD becomes unstable, especially in the *Overlap* setting, while CDD shows consistent performance with minimal fluctuations.

Results on VisA. For the VisA dataset, we set $R_{\text{noise}} = 0.05$ (VisA-noise-0.05) based on the ratio of normal to anomalous samples in the original dataset and conduct relevant experiments as in Tab. 3. The compared methods include unsupervised and fully unsupervised AD methods. Our method achieves the best performance in both *No Overlap* and *Overlap* settings. Notably, in the *Overlap* setting, we outperform the baseline RD by 28.0% in I-AUC, 6.8% in P-AUC, and 1.9% in PRO, respectively, demonstrating that our cross-domain training strategy effectively enhances the baseline’s resilience to anomaly interference.

Table 3: Anomaly detection and localization results on VisA-noise-0.05.

Setting	Metrics	RD [9]	SoftPatch [17]	InReaCh [24]	CDD (Ours)
<i>No Overlap</i>	I-AUC	0.945	0.927	0.827	0.954
	P-AUC	0.979	0.985	0.974	0.982
	PRO	0.897	<u>0.904</u>	0.793	0.911
<i>Overlap</i>	I-AUC	0.656	0.924	0.725	0.936
	P-AUC	0.909	0.954	0.914	0.977
	PRO	<u>0.892</u>	0.883	0.721	0.911




Figure 4: Comparison of histograms of anomaly scores obtained by RD and our CDD.

Visualization Comparisons. We perform additional visualization experiments to compare our proposed CDD with the baseline RD. First, we obtain anomaly scores on both the training and test sets of MVTec-AD-noise-0.1 using the trained RD and CDD, generating histograms of anomaly scores for all the samples as depicted in Fig. 4. On one hand, RD proves effective in the FUAD setting, yet it inadvertently learns certain anomaly patterns from the training set, impairing its ability to accurately detect anomalies. Notably, our CDD overcomes this limitation, markedly improving anomaly detection ability on the training set.




Figure 5: Visualization comparison of anomaly maps generated by RD and our CDD.

Fig. 5 further compares anomaly maps generated by RD and CDD. Compared to RD, CDD exhibits greater sensitivity to anomalies, intuitively demonstrating its ability to mitigate overfitting to some extent, preventing the student network from excessively learning the teacher’s anomaly representations.

5.3 Ablation Analysis

Effectiveness of Proposed Designs. We first conduct the stepwise ablation experiments to evaluate the effectiveness of module designs based on $K = 2$ as in Tab. 4. Without any additional designs, the setup reverts to the baseline RD. For Domain-Specific Training (DST), *D.C.* represents a simple even-split domain construction, while *C.G.* integrates Confidence-Guided Domain Construction for improved division. Initially, domain-specific training relies exclusively on the teacher for supervision, with *Reg.* introducing regularization by distilling from the previous global student. For Cross-Domain Knowledge Aggregation (CDKA), *P.N.* denotes the basic cross-domain distillation, generating pseudo-normal features across domains for feature distillation. The inclusion of *F.P.* involves applying feature perturbation to the global student’s input during training. Lastly, *Conf.D.* refers to training the global student directly on High-Confidence Domains to learn teacher representations. The results in Tab. 4 confirm that the addition of each module consistently enhances performance over the baseline.

Table 5: Ablation study of domain number K on MVTec-noise-0.1.

K	I-AUC	P-AUC	PRO	Average
2	0.9836	0.9818	0.9287	0.9647
3	0.9821	0.9793	0.9252	0.9622
4	0.9791	0.9811	0.9271	0.9624
$\{2, 3, 3, 2\}$	0.9840	0.9812	0.9297	0.9650
$\{2, 3, 4, 3, 2\}$	0.9837	0.9806	0.9260	0.9634

Table 6: Ablation study of pseudo-normal feature selection strategies on MVTec-noise-0.1.

$K = 3$					
Select Strategy		I-AUC	P-AUC	PRO	Average
<i>All</i>		0.9753	0.9747	0.9251	0.9584
One	<i>Next</i>	0.9510	0.9692	0.9142	0.9448
	<i>Consensual</i>	0.9821	0.9793	0.9252	0.9622

Number of Domains. To investigate the impact of the number of domains K , we conduct an ablation study on MVTec-AD, with performance results in Tab. 5. Moreover, we observe that as training progresses, the student can gradually generate normal teacher features. In this case, appropriately increasing K better isolate anomalies. In the later stages, as the global student learns to generate normal features even in anomaly regions, finer domain division becomes less critical, allowing K to be reduced. To test this, we experimented with dynamic K strategies. Results show that the $\{2, 3, 3, 2\}$ strategy achieves a PRO of 0.9297, a 1% improvement over the fixed $K = 2$. This indicates that dynamically adjusting K effectively balances anomaly suppression and normal feature modeling. Therefore, our final design adopts K varying as $\{2, 3, 3, 2\}$ across epochs.

Selection of Pseudo-Normal Features. We conduct an ablation study on pseudo-normal feature selection strategies, all performed with $K = 3$, with results presented in Tab. 6. One strategy, labeled *All*, uses pseudo-normal features generated by domain-specific students from all other domains for distillation. Alternatively, we select features from only one domain, either via our Consensus-driven Pseudo-Normal Feature Selection (denoted as *Consensual*) or by choosing the next domain’s feature (denoted as *Next*, akin to random selection). Results show that our *Consensual* strategy markedly achieves the best performance, which demonstrates that the Consensus-driven strategy significantly enhances cross-domain distillation quality.

6 Conclusions

In this paper, we propose a novel Cross-Domain Distillation framework to address the FUAD task. To reduce the impact of anomalies during training, we introduce two key strategies: Domain-Specific Training, which constructs multiple low-anomaly domains and trains corresponding domain-specific students; and Cross-Domain Knowledge Aggregation, which transfers pseudo-normal features in a cross-domain manner to guide a global student. Compared with the original Reverse Distillation (RD) baseline, our approach significantly improves robustness and accuracy under noisy training conditions. Compared with the original RD baseline, CDD is less affected by anomaly interference under the FUAD setting, as supported by our experimental results.

Discussion. Although CDD is implemented based on the RD paradigm, the core design is conceptually general and could be extended to other UAD paradigms. However, our experiments are restricted to RD-based architectures. Future work will focus on adapting and validating CDD under other paradigms, such as feature reconstruction, to further demonstrate its generality.

References

- [1] S. Akcay, A. Atapour-Abarghouei, and T. P. Breckon. Ganomaly: Semi-supervised anomaly detection via adversarial training. In *Computer Vision-ACCV 2018: 14th Asian Conference on Computer Vision, Perth, Australia, December 2–6, 2018, Revised Selected Papers, Part III 14*, pages 622–637. Springer, 2019.
- [2] J. Bae, J.-H. Lee, and S. Kim. Pni: industrial anomaly detection using position and neighborhood information. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 6373–6383, 2023.
- [3] K. Batzner, L. Heckler, and R. König. Efficientad: Accurate visual anomaly detection at millisecond-level latencies. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 128–138, 2024.
- [4] P. Bergmann, S. Löwe, M. Fauser, D. Sattlegger, and C. Steger. Improving unsupervised defect segmentation by applying structural similarity to autoencoders. *arXiv preprint arXiv:1807.02011*, 2018.
- [5] P. Bergmann, M. Fauser, D. Sattlegger, and C. Steger. Mvtec ad—a comprehensive real-world dataset for unsupervised anomaly detection. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 9592–9600, 2019.
- [6] P. Bergmann, M. Fauser, D. Sattlegger, and C. Steger. Uninformed students: Student-teacher anomaly detection with discriminative latent embeddings. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 4183–4192, 2020.
- [7] P. Bergmann, K. Batzner, M. Fauser, D. Sattlegger, and C. Steger. Beyond dents and scratches: Logical constraints in unsupervised anomaly detection and localization. *International Journal of Computer Vision*, 130(4):947–969, 2022.
- [8] T. Defard, A. Setkov, A. Loesch, and R. Audigier. Padim: a patch distribution modeling framework for anomaly detection and localization. In *International Conference on Pattern Recognition*, pages 475–489. Springer, 2021.
- [9] H. Deng and X. Li. Anomaly detection via reverse distillation from one-class embedding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9737–9746, 2022.
- [10] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee, 2009.
- [11] Z. Gu, L. Liu, X. Chen, R. Yi, J. Zhang, Y. Wang, C. Wang, A. Shu, G. Jiang, and L. Ma. Remembering normality: Memory-guided knowledge distillation for unsupervised anomaly detection. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 16401–16409, 2023.
- [12] D. Gudovskiy, S. Ishizaka, and K. Kozuka. Cflow-ad: Real-time unsupervised anomaly detection with localization via conditional normalizing flows. In *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pages 98–107, 2022.
- [13] H. Guo, L. Ren, J. Fu, Y. Wang, Z. Zhang, C. Lan, H. Wang, and X. Hou. Template-guided hierarchical feature restoration for anomaly detection. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 6447–6458, 2023.
- [14] J. Guo, L. Jia, W. Zhang, H. Li, et al. Recontrast: Domain-specific anomaly detection via contrastive reconstruction. *Advances in Neural Information Processing Systems*, 36, 2024.
- [15] J. Hyun, S. Kim, G. Jeon, S. H. Kim, K. Bae, and B. J. Kang. Reconpatch: Contrastive patch representation learning for industrial anomaly detection. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 2052–2061, 2024.
- [16] J. Im, Y. Son, and J. H. Hong. Fun-ad: Fully unsupervised learning for anomaly detection with noisy training data. In *2025 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 9447–9456. IEEE, 2025.
- [17] X. Jiang, J. Liu, J. Wang, Q. Nie, K. Wu, Y. Liu, C. Wang, and F. Zheng. Softpatch: Unsupervised anomaly detection with noisy data. *Advances in Neural Information Processing Systems*, 35:15433–15445, 2022.
- [18] Y. Jiang, Y. Cao, and W. Shen. A masked reverse knowledge distillation method incorporating global and local information for image anomaly detection. *Knowledge-Based Systems*, 280:110982, 2023.

- [19] C.-L. Li, K. Sohn, J. Yoon, and T. Pfister. Cutpaste: Self-supervised learning for anomaly detection and localization. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 9664–9674, 2021.
- [20] H. Li, Z. Chen, Y. Xu, and J. Hu. Hyperbolic anomaly detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 17511–17520, 2024.
- [21] J. Lin and Y. Yan. A comprehensive augmentation framework for anomaly detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 8742–8749, 2024.
- [22] X. Liu, J. Wang, B. Leng, and S. Zhang. Dual-modeling decouple distillation for unsupervised anomaly detection. In *ACM Multimedia 2024*, 2024. URL <https://openreview.net/forum?id=TOMVFf5L6Q>.
- [23] X. Liu, J. Wang, B. Leng, and S. Zhang. Unlocking the potential of reverse distillation for anomaly detection. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 5640–5648, 2025.
- [24] D. McIntosh and A. B. Albu. Inter-realization channels: Unsupervised anomaly detection beyond one-class classification. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 6285–6295, 2023.
- [25] P. Perera, R. Nallapati, and B. Xiang. Ocgan: One-class novelty detection using gans with constrained latent representations. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 2898–2906, 2019.
- [26] K. Roth, L. Pemula, J. Zepeda, B. Schölkopf, T. Brox, and P. Gehler. Towards total recall in industrial anomaly detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14318–14328, 2022.
- [27] M. Rudolph, T. Wehrbein, B. Rosenhahn, and B. Wandt. Asymmetric student-teacher networks for industrial anomaly detection. In *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pages 2592–2602, 2023.
- [28] M. Salehi, N. Sadjadi, S. Baselizadeh, M. H. Rohban, and H. R. Rabiee. Multiresolution knowledge distillation for anomaly detection. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 14902–14912, 2021.
- [29] T. Schlegl, P. Seeböck, S. M. Waldstein, G. Langs, and U. Schmidt-Erfurth. f-anogan: Fast unsupervised anomaly detection with generative adversarial networks. *Medical image analysis*, 54:30–44, 2019.
- [30] T. D. Tien, A. T. Nguyen, N. H. Tran, T. D. Huy, S. Duong, C. D. T. Nguyen, and S. Q. Truong. Revisiting reverse distillation for anomaly detection. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 24511–24520, 2023.
- [31] C. Wang, W. Zhu, B.-B. Gao, Z. Gan, J. Zhang, Z. Gu, S. Qian, M. Chen, and L. Ma. Real-iad: A real-world multi-view dataset for benchmarking versatile industrial anomaly detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 22883–22892, 2024.
- [32] G. Wang, S. Han, E. Ding, and D. Huang. Student-teacher feature pyramid matching for anomaly detection. In *32nd British Machine Vision Conference 2021, BMVC 2021, Online, November 22-25, 2021*, page 306. BMVA Press, 2021.
- [33] Z. You, L. Cui, Y. Shen, K. Yang, X. Lu, Y. Zheng, and X. Le. A unified model for multi-class anomaly detection. *Advances in Neural Information Processing Systems*, 35:4571–4584, 2022.
- [34] S. Zagoruyko and N. Komodakis. Wide residual networks. *arXiv preprint arXiv:1605.07146*, 2016.
- [35] V. Zavrtanik, M. Kristan, and D. Skočaj. Draem-a discriminatively trained reconstruction embedding for surface anomaly detection. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 8330–8339, 2021.
- [36] J. Zhang, M. Suganuma, and T. Okatani. Contextual affinity distillation for image anomaly detection. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 149–158, 2024.
- [37] X. Zhang, N. Li, J. Li, T. Dai, Y. Jiang, and S.-T. Xia. Unsupervised surface anomaly detection with diffusion probabilistic model. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 6782–6791, 2023.

- 450 [38] X. Zhang, M. Xu, and X. Zhou. Realnet: A feature selection network with realistic synthetic anomaly 1
451 for anomaly detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern*
452 *Recognition*, pages 16699–16708, 2024.
- 453 [39] Y. Zhou, X. Xu, J. Song, F. Shen, and H. T. Shen. Msflow: Multiscale flow-based framework for
454 unsupervised anomaly detection. *IEEE Transactions on Neural Networks and Learning Systems*, 2024.

NeurIPS Paper Checklist ¹

1. Claims ²

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? ³

Answer: [Yes] ⁴

Justification: The main claims made in the abstract and introduction accurately reflect the paper's contributions and scope. ⁵

Guidelines: ⁶

- The answer NA means that the abstract and introduction do not include the claims made in the paper. ⁷
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations ⁸

Question: Does the paper discuss the limitations of the work performed by the authors? ⁹

Answer: [Yes] ¹⁰

Justification: We've discussed the limitations of our work in Sec. 6. ¹¹

Guidelines: ¹²

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper. ¹³
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs ¹⁴

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof? ¹⁵

Answer: [Yes] ¹⁶

Justification: Sec. 3.2 provides the full set of assumptions to support our work. ¹⁷

Guidelines: ¹⁸

- The answer NA means that the paper does not include theoretical results. ¹⁹
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.

- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility 2

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)? 3

Answer: [Yes] 4

Justification: We have introduced all the details of the algorithm in Sec. 4. 5

Guidelines: 6

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code 8

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material? 9

Answer: [No] 10

Justification: We will make the code public after publication. 11

Guidelines: 12

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.

- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details 2

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results? 3

Answer: [Yes] 4

Justification: We've provided the training and test details in Sec. 5.1. 5

Guidelines: 6

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance 8

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments? 9

Answer: [No] 10

Justification: We do not report error bars. Instead, we provide comparative results under different dataset with different anomaly noise ratio. 11

Guidelines: 12

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources 14

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments? 15

Answer: [Yes] 16

Justification: Please see Sec. 5.1. 17

Guidelines: 18

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.

- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics 2

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes] 4

Justification: Our work conform with the NeurIPS Code of Ethics. 5

Guidelines: 6

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts 8

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes] 10

Justification: Please see Sec. 1. 11

Guidelines: 12

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards 14

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA] 16

Justification: Our work poses no such risks. 17

Guidelines: 18

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets 2

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes] 4

Justification: We've cited the related work. 5

Guidelines: 6

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets 8

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA] 10

Justification: We does not release new assets. 11

Guidelines: 12

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects 14

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA] 16

Justification: Our work does not involve crowdsourcing nor research with human subjects. 17

Guidelines: 18

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects 20

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA] 22

Justification: Our work does not involve crowdsourcing nor research with human subjects¹

Guidelines: ²

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. ³
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage ⁴

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required. ⁵

Answer: [No] ⁶

Justification: We do not use LLMs as the component of the core methods. ⁷

Guidelines: ⁸

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components. ⁹
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.