
On the Complexity of Verifying Quantized GNNs with Readout¹

Anonymous Author(s)²

Affiliation

Address

email

Abstract³

¹ In this paper, we introduce a logical language for reasoning about quantized graph
² neural networks (GNNs) with Global Readout. We then prove that verifying quan-
³ tized GNNs with Global Readout is NEXPTIME-complete. We also experimentally
⁴ show the relevance of quantization in the context of ACR-GNNs.

⁵ 1 Introduction⁵

⁶ Graph neural networks (GNNs) are models used for classification and regression tasks on graphs or
⁷ graph-node pairs, aka pointed graphs. GNNs are applied for recommendation in social network [30],
⁸ knowledge graphs [40], chemistry [29], drug discovery [39], etc.

⁹ Quantization designates the fact that numbers are represented by a small amount of bits, opposed
¹⁰ to e.g., integers or real numbers whose number of bits can be arbitrary long. Standard IEEE 754
¹¹ 64-bit floats, INT8, or FP8 [22] enter in our setting. Essentially, our setting reflects GNNs as they are
¹² practically implemented (e.g., in PyTorch), rather than idealized GNNs that assume integer or perfect
¹³ mathematical real number weights, as studied in previous research comparing GNNs and logic [4],
¹⁴ [24] or [8].

¹⁵ GNNs, as several other machine learning models are difficult to interpret, understand and verify. This
¹⁶ is a major issue for their adoption, morally and legally, with the enforcement of regulatory policies
¹⁷ like the EU AI Act [13]. In the literature, verifying quantized GNNs has already been addressed [32].
¹⁸ The methodology is to design a logical language to represent both the properties to check and the
¹⁹ computation of a GNN. However, global readout has not been considered whereas it is an essential
²⁰ element of GNNs, especially for graph classification.

²¹ In this paper, we focus on verifying Aggregate-Combine Graph Neural Networks with global Readout⁹
²² (ACR-GNNs) and we design a logical framework called $q\mathcal{L}$.

²³ **Example 1.** Assume a class of knowledge graphs (KGs) representing communities of people and
²⁴ animals, where each node corresponds to an individual. Each individual can be Animal, Human, Leg,
²⁵ Fur, White, Black, etc. These concepts can be encoded with features $x_0, x_1, \dots, x_5, \dots$ respectively,
²⁶ taking values 0 or 1. Edges in a KG represent a generic ‘has’ relationship: a human can have an
²⁷ animal (pet); an animal can have a human (owner), a leg, a fur; a fur can have a color; etc. Suppose
²⁸ that $\mathcal{A}$ is a GNN processing those KGs and is trained to supposedly recognize dogs. We can verify
²⁹ that the nodes recognized by $\mathcal{A}$ are animals—arguably a critical property of the domain—by checking
³⁰ the validity (i.e., the non-satisfiability of the negation) of $\varphi_{\mathcal{A}} \rightarrow x_0 = 1$ where $\varphi_{\mathcal{A}}$ is a $q\mathcal{L}$ -formula
³¹ corresponding to $\mathcal{A}$ ’s computation, true in exactly the pointed graphs accepted by $\mathcal{A}$. Ideally, $\mathcal{A}$
³² should not overfit the concept of dog as a perfect prototypical animal. For instance, three-legged
³³ dogs do exist. We can verify that $\mathcal{A}$ lets it be a possibility by checking the satisfiability of the formula
³⁴ $\varphi_{\mathcal{A}} \wedge \Diamond^{\leq 3}(x_2 = 1)$.

35 *More complex $q\mathcal{L}$ formulas can be written to express graph properties to be evaluated against an* 1
 36 *ACR-GNN, that will be formalized later in Example 2: 1. Has a human owner, whose pets are all*
 37 *two-legged. 2. A human in a community that has more than twice as many animals as humans, and*
 38 *more than five animals without an owner¹. 3. An animal in a community where some animals have*
 39 *white and black fur.*

40 **Contribution.** In Section 3, we define logic $q\mathcal{L}$ extending the one from [32] for capturing global 2
 41 readout. It is expressive enough to capture quantized ACR-GNNs with arbitrary activation functions.
 42 Moreover, $q\mathcal{L}$ can serve as a flexible graph property specification language reminiscent of modal
 43 logics [9], for expressing e.g. properties 1-3 in Example 1.

44 Section 4 shows that the satisfiability problem of $q\mathcal{L}$ is in NEXPTIME, i.e. it can be decided by a 3
 45 non-deterministic algorithm in exponential time. To do that, we reuse the concept of mathematical
 46 logic called Hintikka sets [] which are complete sets of subformulas that can be true at a given vertex
 47 of a graph. We then introduce a quantized variant of Quantifier-Free Boolean algebra Presburger
 48 Arithmetic (QFBAPA) logic, denoted by QFBAPA_K, and prove that it is in NP as the original
 49 QFBAPA on integers. We then reduce the satisfiability problem of $q\mathcal{L}$ to the one of QFBAPA_K.

50 In Section 5, we then prove that $q\mathcal{L}$ is NEXPTIME-complete, while it is PSPACE-complete without 4
 51 global readout [32]. In a similar way, we also add global counting to the logic $K^\sharp$ previously
 52 introduced by [24]. We show that it corresponds to AC-GNNs over $\mathbb{Z}$ with global readout and trReLU
 53 activation functions. We prove that the satisfiability problem is NEXPTIME-complete, partially
 54 addressing a problem left open in the literature—that is, for the case of integer values and trReLU
 55 activation functions [7, 8]. Details are in the appendix for keep the main text concise.

56 As NEXPTIME is highly intractable, in Section 6, we relax the satisfiability problem of $q\mathcal{L}$ and 5
 57 ACR-GNNs, searching graph counterexamples whose number of vertices is bounded. This problem
 58 is NP-complete. We provide an implementation in this line.

59 We experimentally show in Section 7 that quantization of GNNs provide minimal accuracy degra- 6
 60 dation. Our results confirm that the quantized models retain strong predictive performance while
 61 achieving substantial reductions in model size and inference cost. These findings demonstrate the
 62 practical viability of quantized ACR-GNNs for deployment in resource-constrained environments.

63 **Related work.** [4] showed that ACR-GNNs are capable of capturing the expressive power of 7
 64 FOC₂, that is, two-variable first-order logic with counting. Recent work has explored the logical
 65 expressiveness of GNN variants in more detail. Notably, [24] and [7] introduced logics to exactly
 66 characterize the capabilities of different forms of GNNs. Similarly, [11] analyzed Max-Sum-GNNs
 67 through the lens of Datalog. [32] considered the expressivity of GNN with quantized parameters but
 68 without global readout.

69 On the verification side, [17] studied the complexity of verification of quantized feedforward neural 8
 70 networks (FNNs), while [31, 34] investigated reachability and reasoning problems for general FNNs
 71 and GNNs. Approaches to verification are proposed via integer linear programming (ILP) by [18]
 72 and [41], and via model checking by [33].

73 From a logical perspective, reasoning over structures involving arithmetic constraints is closely tied 9
 74 to several well-studied logics. Relevant work includes Kuncak and Rinard’s decision procedures for
 75 QFBAPA ([20]), as well as developments by [12], [2], [6], and [14]. These logics form the basis for
 76 the characterizations established in [24, 7].

77 Quantization techniques have studied in neural networks, with surveys such as [15, 23] providing com- 10
 78 prehensive overviews focused on maintaining model accuracy. Although most practical advancements
 79 target convolutional neural networks (CNNs), many of the underlying principles extend to GNNs as
 80 well ([42]). NVIDIA has demonstrated hardware-ready quantization strategies ([38]), and frameworks
 81 like PyTorch ([1]) support both post-training quantization and quantization-aware training (QAT), the
 82 latter simulating quantization effects during training to improve low-precision performance. QAT has
 83 been particularly effective in closing the gap between quantized and full-precision models, especially
 84 for highly compressed or edge-deployed systems ([19]). In the context of GNNs, [35] proposed
 85 Degree-Quant, incorporating node degree information to mitigate quantization-related issues. Based

¹Interestingly, $q\mathcal{L}$ goes beyond graded modal logic and even first-order logic. The property of item 2 in Example 1 cannot be expressed in FOL.




Figure 1: DAG data structure for the formula $\text{agg}(x_1 + x_2) + (x_1 + x_2) \geq 3$.

on this, [43] introduced A^2Q , a mixed-precision framework that adapts bitwidths on graph topology to achieve high compression with minimal performance loss.

2 Background

Let $\mathbb{K}$ be a set of quantized numbers, and let n denote the *bitwidth* of $\mathbb{K}$, that is, the number of bits required to represent a number in $\mathbb{K}$. The bitwidth n is written in unary; this is motivated by the fact that n is small and that we would in any case need to allocate n -bit consecutive memory for storing a number. Formally, we consider a sequence $\mathbb{K}_1, \mathbb{K}_2, \dots$ corresponding to bitwidths 1, 2, etc., but we retain the notation $\mathbb{K}$ for simplicity. We suppose that $\mathbb{K}$ saturates: e.g., if $x \geq 0, y \geq 0, x + y \geq 0$ (i.e., no modulo behavior like in `int` in C for instance). We suppose that $1 \in \mathbb{K}$.

We consider Aggregate-Combine Graph Neural Networks with global Readout (ACR-GNNs), a standard class of message-passing GNNs [4, 16]. An ACR-GNN layer is defined by a triple $(\text{comb}, \text{agg}, \text{agg}_g)$, where $\text{comb} : \mathbb{K}^{3m} \rightarrow \mathbb{K}^n$ is a combination function, and agg, agg_g are local and global aggregation functions that map multisets of vectors in $\mathbb{K}^m$ to a single vector in $\mathbb{K}^m$.

An ACR-GNN is composed of a sequence of such layers $(\mathcal{L}^{(1)}, \dots, \mathcal{L}^{(L)})$ followed by a final classification function $\text{cls} : \mathbb{K}^m \rightarrow \{0, 1\}$. Given a graph $G = (V, E)$ and an initial node labelling $x_0 : V \rightarrow \{0, 1\}^k$, the state of a node u in layer i is recursively defined as:

$$x_i(u) = \text{comb}(x_{i-1}(u), \text{agg}(\{x_{i-1}(v) \mid uv \in E\}), \text{agg}_g(\{x_{i-1}(v) \mid v \in V\}))$$

The final output of the GNN for a pointed graph (G, u) is $\mathcal{A}(G, u) = \text{cls}(x_L(u))$. A more detailed definition is provided in Appendix C.2.

Our study focuses on a specific subclass where both agg and agg_g perform summation over vectors, and where $\text{comb}(x, y, z) = \vec{\sigma}(xC + yA_1 + zA_2 + b)$, using matrices C, A_1, A_2 with entries from $\mathbb{K}$, and a bias $b \in \mathbb{K}$. The classification function is a linear threshold: $\text{cls}(x) = \sum_i a_i x_i \geq 1$ with weights $a_i \in \mathbb{K}$. Moreover, we assume that all arithmetic operations are executed according to the arithmetic related to $\mathbb{K}$. It is assumed that the context makes clear the $\mathbb{K}$ and arithmetic being used. We note $[[\mathcal{A}]]$ the set of pointed graphs (G, u) such that $\mathcal{A}(G, u) = 1$. An ACR-GNN $\mathcal{A}$ is satisfiable if $[[\mathcal{A}]]$ is non-empty. The *satisfiability problem* for ACR-GNNs is: Given a ACR-GNN $\mathcal{A}$, decide whether $\mathcal{A}$ is satisfiable.

3 Logic $q\mathcal{L}$ for Representing GNN Computations and Properties on Graphs

We set up a logical framework called $q\mathcal{L}$ extending the logic in [32] with global aggregation: it is a *lingua franca* to represent GNN computation and properties on graphs.

Syntax. Let F be a finite set of features and $\mathbb{K}$ be some finite-width arithmetic. We consider a set of *expressions* defined by the following grammar in Backus-Naur form:

$$\vartheta ::= c \mid x_i \mid \alpha(\vartheta) \mid \text{agg}(\vartheta) \mid \text{agg}_g(\vartheta) \mid \vartheta + \vartheta \mid c \times \vartheta$$

where c is a number in $\mathbb{K}$, x_i is a feature in F , α is a symbol for denoting the activation function, and agg and agg_g denote the aggregation function for local and global readout respectively. A *formula* is a construction of the formula $\vartheta \geq k$ where ϑ is an expression and k is an element of $\mathbb{K}$. If $-1 \in \mathbb{K}$, and $-\vartheta$ is not, we can write $-\vartheta$ instead of $(-1) \times \vartheta$. Other standard abbreviations can be used.

Formulas are represented as direct acyclic graphs, aka circuits, meaning that we do not repeat the same expressions several times. For instance, the formula $\text{agg}(x_1 + x_2) + (x_1 + x_2) \geq 3$ can be represented as the DAG given in Figure 1. Formulas can also be represented by a sequence of assignments via new fresh intermediate variables. For instance: $y := x_1 + x_2, z := \text{agg}(y) + y, \text{res} := z \geq 3$.

124 **Semantics.** Consider a graph $G = (V, E)$, where vertices in V are labeled via a labeling function 1
 125 $\ell : V \rightarrow \mathbb{K}^n$ with feature values. The value of an expression ϑ in a vertex $u \in V$ is denoted by
 126 $[[\vartheta]]_{G,u}$ and is defined by induction on ϑ :

$$\begin{aligned} [[c]]_{G,u} &= c, & [[c \times \vartheta]]_{G,u} &= c \times_{\mathbb{K}} [[\vartheta]]_{G,u}, \\ [[x_i]]_{G,u} &= \ell(u)_i, & [[\alpha(\vartheta)]]_{G,u} &= [[\alpha]]([[\vartheta]]_{G,u}), \\ [[\vartheta + \vartheta']]_{G,u} &= [[\vartheta]]_{G,u} +_{\mathbb{K}} [[\vartheta']]_{G,u}, & [[agg(\vartheta)]]_{G,u} &= \sum_{v|uEv} [[\vartheta]]_{G,v}, \\ & & [[agg_{\forall}(\vartheta)]]_{G,u} &= \sum_{v \in V} [[\vartheta]]_{G,v}, \end{aligned}$$

128 We define $[[\vartheta \geq k]] = \{G, u \mid [[\vartheta]]_{G,u} \geq_{\mathbb{K}} [[k]]_{G,u}\}$ (we write $\geq$ for the symbol in the syntax and 4
 129 $\geq_{\mathbb{K}}$ for the comparison in $\mathbb{K}$). A formula φ is satisfiable if $[[\varphi]]$ is non-empty. The *satisfiability*
 130 *problem* for $q\mathcal{L}$ is: Given a $q\mathcal{L}$ -formula φ , decide whether φ is satisfiable.

131 **ACR-GNN verification tasks.** We are interested in the following decision problems. Given a GNN 5
 132 $\mathcal{A}$, and a $q\mathcal{L}$ formula φ : (VT1, sufficiency) Do we have $[[\varphi]] \subseteq [[\mathcal{A}]]$? (VT2, necessity) Do we have
 133 $[[\mathcal{A}]] \subseteq [[\varphi]]$? (VT3, consistency) Do we have $[[\varphi]] \cap [[\mathcal{A}]] \neq \emptyset$?

134 **Representing a GNN computation.** To reason formally about ACR-GNNs, we represent their 6
 135 computations using $q\mathcal{L}$. Logic $q\mathcal{L}$ facilitates the modeling of the acceptance condition of ACR-GNNs.

136 We explain this via example. Consider a two-layer ACR-GNN $\mathcal{A}$ with input and output dimension 2, 7
 137 using summation for aggregation, activation via $\alpha(x) := \max(0, \min(1, x))$ —the truncated ReLU—
 138 and a classification function $2x_1 - x_2 \geq 1$. The combination functions are:

$$\begin{aligned} comb_1((x_1, x_2), (y_1, y_2), (z_1, z_2)) &:= \begin{pmatrix} \sigma(2x_1 + x_2 + 5y_1 - 3y_2 + 1) \\ \sigma(-x_1 + 4x_2 + 2y_1 + 6y_2 - 2) \end{pmatrix}, \\ comb_2((x_1, x_2), (y_1, y_2), (z_1, z_2)) &:= \begin{pmatrix} \sigma(3x_1 - y_1 + 2z_2) \\ \sigma(-2x_1 + 5y_2 + 4z_1) \end{pmatrix}. \end{aligned}$$

139 Note that this assumes that $\mathcal{A}$ operates over $\mathbb{K}$ with at least three bits. Then, the corresponding 9
 140 $q\mathcal{L}$ formula $\varphi_{\mathcal{A}}$ is given by: $\psi_1 = \alpha(2x_1 + x_2 + 5agg(x_1) - 3agg(x_1) + 1)$, $\psi_2 := \alpha(-x_1 +$
 141 $4x_2 + 2agg(x_1) + 6agg(x_2) - 2)$, $\chi_1 := \alpha(3\psi_1 - agg(\psi_1) + 2(agg_{\forall}(psi_2)))$, $\chi_2 := \alpha(-2\psi_1 +$
 142 $5(agg(\psi_2)) + 4agg_{\forall}(psi_1))$, $\varphi_{\mathcal{A}} := 2(\chi_1) - \chi_2 \geq 1$. To sum up, given a GNN $\mathcal{A}$, we compute
 143 $q\mathcal{L}$ -formula in poly-time in the size of $\mathcal{A}$ with $[[\mathcal{A}]] = [[\varphi_{\mathcal{A}}]]$ (as done in [32]).

144 **Simulating a modal logic in the logic $q\mathcal{L}$.** In this section, we show that extending $q\mathcal{L}$ with 10
 145 modal operators [9] does not increase the expressivity. We can even compute an equivalent $q\mathcal{L}$
 146 without Boolean connectives and without modal operators in poly-time. It means that formulas like
 147 $\varphi_{\mathcal{A}_1} \rightarrow x_0 = 1$ or $\varphi_{\mathcal{A}_1} \wedge \Diamond^{\leq 3}(x_2 = 1)$ have equivalent formulas in $q\mathcal{L}$.

148 Assume that α is ReLU. Let Atm_0 be the set of atomic formulas of $q\mathcal{L}$ of the form $\vartheta \geq 0$. We 11
 149 suppose that ϑ takes integer values. In general, $\vartheta \geq k$ is an atomic formula equivalent to $\vartheta - k \geq 0$.
 150 Without loss of generality, we thus assume that formulas of $q\mathcal{L}$ are over Atm_0 . Let modal $q\mathcal{L}$ be the
 151 propositional logic on Atm_0 extended with modalities and a restricted variant of graded modalities
 152 where number k in $\mathbb{K}$.

$$\begin{aligned} [[\Box\varphi]] &= \{G, u \mid G, v \in [[\varphi]] \text{ for every } v \text{ s.t. } uEv\} \\ [[\Box_g\varphi]] &= \{G, u \mid G, v \in [[\varphi]] \text{ for every } v \text{ in } V\} \end{aligned}$$

$$[[\Diamond^{\geq k}\varphi]] = \{G, u \mid |\{G, v \mid uEv \text{ and } G, v \in [[\varphi]]\}| \geq_{\mathbb{K}} k\} \quad [[\Diamond_g^{\geq k}\varphi]] = \{G, u \mid |[[\varphi]]| \geq_{\mathbb{K}} k\}$$

154 and modalities $\Diamond^{\leq k}\varphi$ and $\Diamond_g^{\leq k}\varphi$ defined the same way but with $\leq_{\mathbb{K}}$. We can turn back to the graph 14
 155 properties mentioned in Example 1.

156 **Example 2.** We first define a few simple formulas to characterize the concepts of the domain. Let 15
 157 $\varphi_A := x_0 = 1$ (Animal), $\varphi_H := x_1 = 1$ (Human), $\varphi_L := x_2 = 1$ (Leg), $\varphi_F := x_3 = 1$ (Fur),
 158 $\varphi_W := x_4 = 1$ (White), and $\varphi_B := x_5 = 1$ (Black).

- 159 1. Has a human owner, whose all pets are two-legged: $\Diamond(\varphi_H \wedge \Box(\varphi_A \rightarrow \Diamond^2\varphi_L))$. 16
- 160 2. A human in a community that has more than twice as many animals as humans, and more than 17
 161 five animals without an owner: $\varphi_H \wedge (agg_{\forall}(x_0) - 2 \times agg_{\forall}(x_1) \geq 0) \wedge \Diamond_g^{\geq 5}((\varphi_A \wedge \Box(\neg\varphi_H))$.

162 3. An animal in a community where some animals have white and black fur:
 163 $\varphi_A \wedge \Diamond_g(\Diamond(\varphi_F \wedge \Diamond\varphi_W) \wedge \Diamond(\varphi_F \wedge \Diamond\varphi_B)).$

164 We can see the boolean operator $\neg$, and the various modalities as functions from Atm_0 into Atm_0 , 1
 165 and the boolean operator $\vee$ as a function from $Atm_0 \times Atm_0$ to Atm_0 .

$$\begin{aligned} f_{\neg}(\vartheta \geq 0) &:= -\vartheta - 1 \geq 0 & f_{\vee}(\vartheta_1 \geq 0, \vartheta_2 \geq 0) &:= \vartheta_1 + \text{ReLU}(\vartheta_2 - \vartheta_1) \geq 0 & 2 \\ f_{\Box}(\vartheta \geq 0) &:= \text{agg}(-\text{ReLU}(-\vartheta)) \geq 0 \\ f_{\Diamond \geq k}(\vartheta \geq 0) &:= \text{agg}(\text{ReLU}(\vartheta + 1) - \text{ReLU}(\vartheta)) - k \geq 0 \\ f_{\Diamond \leq k}(\vartheta \geq 0) &:= k - \text{agg}(\text{ReLU}(\vartheta + 1) - \text{ReLU}(\vartheta)) \geq 0 \end{aligned}$$

166 For the corresponding global modalities ($f_{\Box_g}(\vartheta \geq 0)$, $f_{\Diamond \geq k}(\vartheta \geq 0)$, and $f_{\Diamond \leq k}(\vartheta \geq 0)$), it suffices to 3
 167 use $\text{agg}_{\vee}$ in place of agg . The previous transformations can be generalized to arbitrary formulas of
 168 modal $q\mathcal{L}$ as follows.

$$\begin{aligned} \text{mod2expr}(\vartheta \geq 0) &:= \vartheta \geq 0 & \text{mod2expr}(\neg\varphi) &:= f_{\neg}(\text{mod2expr}(\varphi)) & 4 \\ \text{mod2expr}(\varphi_1 \vee \varphi_2) &:= f_{\vee}(\text{mod2expr}(\varphi_1), \text{mod2expr}(\varphi_2)) \\ \text{mod2expr}(\Box\varphi) &:= f_{\Box}(\text{mod2expr}(\varphi)), & \Box &\in \{\Box, \Box_g, \Diamond \geq k, \Diamond \leq k, \Diamond_g \geq k, \Diamond_g \leq k\} \end{aligned}$$

169 We can show that formulas of modal $q\mathcal{L}$ can be captured by a unique expression $\vartheta \geq 0$. This is a 5
 170 consequence of the following lemma ².

171 **Lemma 3.** *Let φ be a formula of modal $q\mathcal{L}$. The formulas φ and $\text{mod2expr}(\varphi)$ are equivalent.* 6

172 Now, ACR-GNN verification tasks can be solved by reduction to the satisfiability problem of $q\mathcal{L}$. 7
 173 VT1 by checking that $\varphi \wedge \neg\varphi_{\mathcal{A}}$ is not satisfiable; VT2 by checking that $\neg\varphi \wedge \varphi_{\mathcal{A}}$ is not satisfiable;
 174 VT3 by checking that $\varphi \wedge \varphi_{\mathcal{A}}$ is satisfiable.

175 4 NEXPTIME Membership of the Satisfiability Problem 8

176 In this section, we prove the NEXPTIME membership of reasoning in modal quantized logic, and 9
 177 also of solving of ACR-GNN verification tasks (by reduction to the former). Remember that the
 178 activation function α can be arbitrary in our setting. Our result holds with the loose restriction that
 179 $[[\alpha]]$ is computable in exponential-time in the bit-width n of $\mathbb{K}$.

180 **Theorem 4.** *The satisfiability problem of $q\mathcal{L}$ is decidable and in NEXPTIME, and so is VT3. VT1 10
 181 and VT2 are in coNEXPTIME.*

182 In order to prove Theorem 4, we adapt the NEXPTIME membership of the description logic 11
 183 $\mathcal{ALCS}^{++}$ from [2] to logic $q\mathcal{L}$. The difference resides in the definition of Hintikka sets and
 184 the treatment of quantization. The idea is to encode the constraints of a $q\mathcal{L}$ -formula φ in a formula of
 185 exponential length of a quantized version of QFBAPA, that we prove to be in NP.

186 4.1 Hintikka Sets 12

187 Consider $q\mathcal{L}$ -formula φ . Let $E(\varphi)$ be the set of subexpressions in φ . For instance, if φ is 13
 188 $3 \times \text{agg}(\alpha(x_2 + \text{agg}_{\vee}(x_1))) \geq 5$ then $E(\varphi) := \{\text{agg}(\alpha(x_2 + \text{agg}_{\vee}(x_1))), \alpha(x_2 + \text{agg}_{\vee}(x_1)), x_2,$
 189 $\text{agg}_{\vee}(x_1), x_1\}$. From now on, we consider equality subformulas that are of the form $\vartheta = k$ where ϑ is
 190 a subexpression of φ and $k \in \mathbb{K}$.

191 **Definition 5.** *A Hintikka set H for φ is a subset of subformulas of φ such that:* 14

- 192 1. For all $\vartheta \in E(\varphi)$, there is a unique value $k \in \mathbb{K}$ such that $\vartheta = k \in H$ 15
- 193 2. $\vartheta_1 = k_1, \vartheta_2 = k_2 \in H$ then $\vartheta_1 + \vartheta_2 = k_1 + k_2 \in H$
- 194 3. If $\vartheta \geq k \in H$ then $c \times \vartheta = k' \in H$ where $k' = c \times_{\mathbb{K}} k$
- 195 4. $\vartheta = k \in H$ and $\alpha(\vartheta) = k'$ implies $k' = [[\alpha]](k)$

²For simplicity, we do not present how to handle $\vartheta \geq 0$ when ϑ is not an integer. We could introduce several activation functions α in $q\mathcal{L}$, one of them could be interpreted as the Heavyside step function. In the sequel Definition 5, Point 4 is just repeated for each α .

Informally, a *Hintikka set* is a set of equality subformulas obtained from a choice of a value for each subexpression of φ (point 1), provided that the set is consistent *at the current vertex* (point 2-4). Note that the notion of Hintikka set does not take any constraints about agg and agg_v into consideration since checking consistency of aggregation would require information about the neighbor or the whole graph.

Example 6. If φ is $3 \times agg(\alpha(x_2 + agg_v(x_1))) \geq 5$ then the following set is an example of Hintikka set: $\{agg(\alpha(x_2 + agg_v(x_1))) = 8, \alpha(x_2 + agg_v(x_1)) = 9, x_2 + agg_v(x_1) = 9, x_2 = 7, agg_v(x_1) = 2, x_1 = 5\}$.

Proposition 7. The number of Hintikka sets is bounded by $2^{n|\varphi|}$ where $|\varphi|$ is the size of φ , and n is the bitwidth of $\mathbb{K}$.

4.2 Quantized Version of QFBABA (Quantifier-free Boolean Algebra and Presburger Arithmetics)

A QFBAPA formula is propositional formula where each atom is either an inclusion of sets or equality of sets or linear constraints [20]. Sets are denoted by Boolean algebra expression e.g., $(S \cup S') \setminus S''$, or $\mathcal{U}$ where $\mathcal{U}$ denotes the set of all points in some domain. Here S, S' , etc. are set variables. Linear constraints are over $|S|$ denoting the cardinality of the set denoted by the set expression S . For instance, the QFBAPA-formula $(pianist \subseteq happy) \wedge (|happy| + |\mathcal{U} \setminus pianist| \geq 6) \wedge (|happy| < 2)$ is read as ‘all pianists are happy and the number of happy persons + the number of persons that are not pianists is greater than 6 and the number of happy persons is smaller than 2’.

We now introduce a *quantized* version QFBAPA $_{\mathbb{K}}$ of QFBAPA. It has the same syntax as QFBAPA except that hard-coded numbers in expressions are in $\mathbb{K}$. Concerning the semantics, every numerical expression is interpreted in $\mathbb{K}$. For each set expression S , the interpretation of $|S|$ is not the cardinality c of the interpretation of S , but the result of the computation $1 + 1 + \dots + 1$ in $\mathbb{K}$ with c occurrences of 1 in the sum.

We consider that $\mathbb{K}$ saturates, meaning that if $x + y$ exceed the upper bound limit of $\mathbb{K}$, there is a special value denoted by $+\infty$ such that $x + y = +\infty$.

Proposition 8. If bitwidth n is in unary, and if $\mathbb{K}$ saturates, then satisfiability in QFBAPA $_{\mathbb{K}}$ is in NP.

4.3 Reduction to QFBAPA $_{\mathbb{K}}$

Let φ be a formula of $q\mathcal{L}$. For each Hintikka set H , we introduce the set variable X_H that intuitively represents the H -vertices, i.e., the vertices in which subformulas of H hold. The following QFBAPA $_{\mathbb{K}}$ -formulas say that the interpretation of X_H form a partition of the universe. For each subformula $\vartheta' = k$, we introduce the set variable $X_{\vartheta'=k}$ that intuitively represents the vertices in which $\vartheta' = k$ holds. Formula (1) expresses that $\{X_H\}_H$ form a partition of the universe. Formula (2) makes the bridge between variables $X_{\vartheta'=k}$ and X_H .

$$(1) \quad \left(\bigwedge_{H \neq H'} X_H \cap X_{H'} = \emptyset \right) \wedge \left(\bigcup_H X_H = \mathcal{U} \right) \quad (2) \quad \bigwedge_{\vartheta' \in E(\varphi)} \bigwedge_{k \in \mathbb{K}} (X_{\vartheta'=k} = \bigcup_{H | \vartheta'=k \in H} X_H)$$

We introduce also a variable S_H that denotes the set of all successors of some H -vertex. If there is no H -vertex then the variable S_H is just irrelevant.

The following QFBAPA $_{\mathbb{K}}$ -formula encodes the semantics of $agg(\vartheta)$. More precisely, it says that for all subexpressions $agg(\vartheta)$, for all values k , for all Hintikka sets H containing subformula $agg(\vartheta)=k$, for all H containing $agg(\vartheta)=k$, it says that, if there is some H -vertex (i.e., vertices in S_H), then the aggregation obtained by summing over the successors of some H -vertex is k .

$$\bigwedge_{agg(\vartheta) \in E(\varphi)} \bigwedge_{k \in \mathbb{K}} \bigwedge_{\substack{\text{Hintikka set } H \\ |agg(\vartheta)=k \in H}} [(X_H \neq \emptyset) \rightarrow \sum_{k' \in \mathbb{K}} |S_H \cap X_{\vartheta=k'}| \times k' = k] \quad (3)$$

In the previous sum, we partition S_H into subsets $S_H \cap X_{\vartheta=k'}$ for all possible values k' . Each contribution for a successor in $S_H \cap X_{\vartheta=k'}$ is k' . We rely here on the fact³ that $(1 + 1 + \dots + 1) \times k' =$

³This is true for some fixed-point arithmetics but not for floating-point arithmetics. See Appendix B.




Figure 2: Encoding a torus of exponential size with (modal) $q\mathcal{L}$ formulas. (x, y) are the vertices of the graph that correspond to locations in the torus while φ_N and φ_E denote intermediate vertices indicating the direction (resp., north and east).

$k' + k' + \dots + k'$. We also fix a specific order over values k' in the summation (it means that $agg(\vartheta)$ is computed as follows: first order the successors according to the taken values of ϑ in that specific order, then perform the summation). Finally, the semantics of agg_{ϑ} is captured by the formula:

$$\bigwedge_{agg_{\vartheta}(\vartheta) \in E(\varphi)} \bigwedge_{k \in \mathbb{K}} X_{agg_{\vartheta}(\vartheta)=k} \neq \emptyset \rightarrow \sum_{k' \in \mathbb{K}} |X_{\vartheta=k'}| \times k' = k \quad (4)$$

Note that intuitively Formula (4) implies that for $X_{agg_{\vartheta}(\vartheta)=k}$ is interpreted as the universe, for the value k which equals the semantics of $\sum_{k' \in \mathbb{K}} |X_{\vartheta=k'}| \times k'$.

Given $\varphi = \vartheta \geq k$, we define $tr(\varphi) := \psi \wedge \bigvee_{k' \geq k} X_{\vartheta=k'} \neq \emptyset$ where ψ the conjunction of Formulas 1–4. The function tr requires to compute all the Hintikka sets. So we need in particular to check Point 4 of Definition 5 and we get the following when $[[\alpha]]$ is computable in exponential time in n .

Proposition 9. $tr(\varphi)$ is computable in exponential-time in $|\varphi|$ and n .

Proposition 10. Let φ be a formula of $q\mathcal{L}$. φ is satisfiable iff $tr(\varphi)$ is $QFBAPA_{\mathbb{K}}$ satisfiable.

Finally, in order to check whether a $q\mathcal{L}$ -formula φ is satisfiable, we construct a $QFBAPA_{\mathbb{K}}$ -formula $tr(\varphi)$ in exponential time. As the satisfiability problem of $QFBAPA_{\mathbb{K}}$ is in NP, we obtain that the satisfiability problem of $q\mathcal{L}$ is in NEXPTIME. We proved Theorem 4,

Remark 11. Our methodology can be generalized to reason in subclasses of graphs. For instance, we may tackle the problem of satisfiability in a graph where vertices are of bounded degree bounded by d . To do so, we add the constraint $\bigwedge_H |S_H| \leq d$.

5 Complexity Lower Bound

The NEXPTIME upper-bound is tight. Having defined modalities in $q\mathcal{L}$ and stated Lemma 3, Theorem 12 is proven by adapting the proof of NEXPTIME-hardness of deciding the consistency of $\mathcal{ALCQ-T_C}$ Boxes presented in [36]. So we already have the hardness result for ReLU.

NEXPTIME-hardness is proven via a reduction from the tiling problem by Wang tiles of a torus of size $2^n \times 2^n$. A Wang tile is a square with colors, e.g., , etc. That problem takes as input a number n in unary, and Wang tile types, and an initial condition – let say the bottom row is already given. The objective is to decide whether the torus of $2^n \times 2^n$ can be tiled while colors of adjacent Wang tiles match. A slight difficulty resides in adequately capturing a two-dimensional grid structure—as in Figure 2—with only a single relation. To do that, we introduce special formulas φ_E and φ_N to indicate the direction (east or north). In the formula computed by the reduction, we also need to bound the number of vertices corresponding to tile locations by $2^n \times 2^n$. Thus $\mathbb{K}$ needs to encode $2^n \times 2^n$. We need a bit-width of at least $2n$.

Theorem 12. The satisfiability problem in $q\mathcal{L}$ is NEXPTIME-hard, and so is VT3. VT1 and VT2 are coNEXPTIME-hard.

Remark 13. It turns out that the verification task only needs the fragment of $q\mathcal{L}$ where agg is applied directly on an expression $\alpha(\cdot)$. Indeed, this is the case when we represent a GNN in $q\mathcal{L}$ or when we translate logical formulas in $q\mathcal{L}$ (Lemma 3). Reasoning about $q\mathcal{L}$ when $\mathbb{K} = \mathbb{Z}$ and the activation function is truncated ReLU is also NEXPTIME-complete (see Appendix E).

6 Bounding the Number of Vertices

The satisfiability problem is NEXPTIME-complete, thus far from tractable. The complexity comes essentially because counterexamples can be arbitrary large graphs. However, usually we are search for small counterexamples. Let $\mathcal{G}^{\leq N}$ be the set of pointed graphs with at most N vertices. We consider the $q\mathcal{L}$ and ACR-GNN *satisfiability problems with a bound on the number of vertices*: given a number N given in unary, 1. given a $q\mathcal{L}$ -formula φ , is it the case that $[[\varphi]] \cap \mathcal{G}^{\leq N} \neq \emptyset$, 2. given an ACR-GNN $\mathcal{A}$, is it the case that $[[\mathcal{A}]] \cap \mathcal{G}^{\leq N} \neq \emptyset$.

Theorem 14. *The satisfiability problems with bounded number of vertices are NP-complete.*

We then can extend the methodology of [33] but for verifying GNNs. Our implementation proposal is a Python program that takes a learnt quantized GNN $\mathcal{A}$ as an input, a precondition, a postcondition and a bound N . It then produces a C program that mimics the execution of $\mathcal{A}$ on an arbitrary graph with at most N vertices, and embeds the pre/postcondition. We then apply ESBMC (efficient SMT-based context-bounded model checker) [21] on the C program.

7 Quantization Effects on Accuracy, Performance and Model Size

To confirm that the GNN models considered in this paper are promising, we now investigate the application of Dynamic Post-Training Quantization (PTQ) to Aggregate-Combined Readout Graph Neural Networks (ACR-GNNs). Our experimental design builds on the framework introduced in [4], using their publicly available implementation [5] as the baseline. ACR-GNNs with specific structural configurations are used as the primary model class for evaluation. Dynamic PTQ, implemented in PyTorch [1, 26], converts a pre-trained floating-point model into a quantized version without retraining. This approach quantizes weights to INT8 statically, while activations remain in floating point until dynamically quantized at compute time. This enables efficient INT8-based computation, reducing memory usage and improving inference speed. PyTorch’s implementation employs per-tensor quantization for weights and stores activations in floating-point format between operations. The evaluation focuses on accuracy, model size, and latency. Experiments are conducted on both synthetic and real-world datasets, with the synthetic benchmark—based on dense Erdős–Rényi graph structures and logical labeling schemes—serving as the primary focus.

The synthetic graphs were generated using the dense Erdős–Rényi model, a classical approach for constructing random graphs. Each graph includes five initial node colours, encoded as one-hot feature vectors. Following [4], labels were assigned using formulas from the logic fragment FOC_2 . Specifically, a hierarchy of classifiers $\alpha_i(x)$ was defined as:

$$\alpha_0(x) := \text{Blue}(x), \quad \alpha_{i+1}(x) := \exists^{[N,M]} y (\alpha_i(y) \wedge \neg E(x, y))$$

where $\exists^{[N,M]}$ denotes the quantifier “there exist between N and M nodes” satisfying a given condition. Each classifier $\alpha_i(x)$ can be expressed within FOC_2 , as the bounded quantifier can be rewritten using $\exists \geq N$ and $\neg \exists \geq M+1$. Each property p_i corresponds to a classifier α_i with $i \in 1, 2, 3$. Summary statistics for the dataset are provided in Appendix G, Table 3.

Table 1: Accuracy difference (%) and model size (MB) of the ACR-GNN model before and after dynamic post-training quantization (PTQ) across FO-properties p_1 , p_2 , and p_3 . Values are reported for three model depths (1, 2, and 3 layers) and three dataset splits (Train, Test 1, Test 2). Accuracy values represent the change after quantization (QINT8 – FP32). p_1, p_2, p_3 are FO-properties described in Appendix G.

#	p_1			p_2			p_3			Size (MB)
	Train	Test 1	Test 2	Train	Test 1	Test 2	Train	Test 1	Test 2	
1	-0.452%	-0.760%	+0.522%	-0.127%	-0.183%	+8.891%	-0.299%	-0.648%	-0.693%	0.034
2	-0.001%	0.000%	-0.043%	+0.083%	-0.125%	+0.144%	-0.178%	-0.226%	+0.018%	0.068
3	-0.036%	+0.062%	-0.494%	-0.161%	-0.143%	-0.342%	-0.015%	+0.280%	-0.346%	0.103

Table 1 presents the difference in accuracy and model size between the quantized (QINT8⁴) and original (FP32) versions of the ACR-GNN model across three configurations (1, 2, and 3 layers). The

⁴The difference between INT8 and QINT8 lies in their implementation and is detailed in Appendix G

evaluation is conducted on three FO-properties (p_1, p_2, p_3) over three data splits: Train, Test1, and Test2. The table highlights how quantization affects accuracy at various depths. In most cases, the impact of quantization on accuracy is minor and bounded, with some configurations even showing positive differences. For instance, in the 2-layer configuration—the overall best performer—the accuracy loss remains within ± 0.1 across all properties and splits, while yielding a model size reduction of 0.068 MB. The 1-layer model shows greater fluctuation: while p_2 on Test2 experiences a significant positive spike (+8.891), p_3 on Test2 drops by -0.693 . This suggests sensitivity to quantization in shallow models, likely due to limited representational capacity. The results confirm that dynamic post-training quantization (PTQ) enables significant compression—up to 60% reduction in size—while maintaining acceptable levels of accuracy. Additional breakdowns, including baseline results and extended configurations, are provided in Appendix G.

Table 2: PPI benchmark. Accuracy (%) and size (MB) of the ACR-GNN with ReLU activation function before and after dynamic PTQ across different layer configurations.

#	Original (FP32)				Quantized (QINT8)				Difference			
	Train	Val	Test	Size (MB)	Train	Val	Test	Size (MB)	Train	Val	Test	Size (MB)
1	54.7%	43.1%	39.5%	0.922	55.0%	50.8%	50.2%	0.242	+0.3%	+7.7%	+10.7%	0.680
2	52.5%	44.6%	45.7%	1.718	52.3%	47.8%	47.2%	0.451	-0.2%	+3.2%	+1.5%	1.267
3	52.3%	42.6%	44.0%	2.515	51.9%	45.7%	42.8%	0.660	-0.4%	+3.1%	-1.2%	1.855

Table 2 shows the results of evaluating the ACR-GNN model on the Protein-Protein Interaction (PPI) benchmark before and after applying dynamic post-training quantization (PTQ). The evaluation covers three model configurations (1 to 3 layers) and reports performance in terms of accuracy (Train, Validation, and Test) and model size (in MB). Quantization results in substantial compression across all configurations. The model size decreases from 0.922 MB to 0.242 MB (a 73% reduction) for the 1-layer network, while the 2- and 3-layer models achieve reductions of 1.267 MB and 1.855 MB, respectively. Accuracy-wise, quantization leads to improvements in the Validation and Test sets for shallower networks. The 1-layer model gains +0.077 on validation and +0.107 on test accuracy, indicating potential for enhanced generalization. The 2-layer model shows minor improvements across all splits, with negligible loss in training accuracy. However, the 3-layer configuration reveals a slight drop in test accuracy (-0.012), suggesting increased sensitivity to quantization at greater depth. See Appendix G, Tables 16, 17, and 18 for additional quantitative breakdowns.

8 Conclusion and Future Work

The central result is the NEXPTIME-completeness of the logic $q\mathcal{L}$ in which both the computations of GNNs and modal properties can be expressed. It helps to understand the inherent complexity of verifying quantized GNNs. We also provide a prototype for verifying GNNs over a set of graphs with a bounded number of vertices. Finally some experiments confirmed that the quantization of ACR-GNNs is promising.

There are many directions to go. First, characterizing the modal flavor of $q\mathcal{L}$ for other activation functions than ReLU. New extensions of $q\mathcal{L}$ could be proposed to tackle other classes GNNs. Verification of neural networks is challenging and is currently tackled by the verification community [10]. So it will be for GNNs as well. Our verification tool with a bound on the number of vertices is still preliminary. One obvious path would be to improve the tool, to compare different approaches (bounded model checking vs. linear programming as in [18]) and apply it to real GNN verification scenarios. Designing a practical verification procedure in the general case (without any bound on the number of vertices) and overcoming the high computational complexity is an exciting challenge for future research towards the verification of GNNs.

Limitations. Section 4 and 5 reflect theoretical results. Some practical implementations of GNNs may not fully align with them. In particular, the order in the (non-associative) summation over values in $\mathbb{K}$ is fixed in formulas (3) and (4). It means that we suppose that the aggregation $agg(\vartheta)$ is computed in that order too (we sort the successors of a vertex according to the values of ϑ and then perform the summation). The verification tool discussed in Section 6 remains a prototype, thus its application warrants careful consideration.

References

- [1] Jason Ansel, Edward Yang, Horace He, Natalia Gimelshein, Animesh Jain, Michael Voznesensky, Bin Bao, Peter Bell, David Berard, Evgeni Burovski, Geeta Chauhan, Anjali Chourdia, Will Constable, Alban Desmaison, Zachary DeVito, Elias Ellison, Will Feng, Jiong Gong, Michael Gschwind, Brian Hirsh, Sherlock Huang, Kshiteej Kalambarkar, Laurent Kirsch, Michael Lazos, Mario Lezcano, Yanbo Liang, Jason Liang, Yinghai Lu, CK Luk, Bert Maher, Yunjie Pan, Christian Puhersch, Matthias Reso, Mark Saroufim, Marcos Yukio Siraichi, Helen Suk, Michael Suo, Phil Tillet, Eikan Wang, Xiaodong Wang, William Wen, Shunting Zhang, Xu Zhao, Keren Zhou, Richard Zou, Ajit Mathews, Gregory Chanan, Peng Wu, and Soumith Chintala. Pytorch 2: Faster machine learning through dynamic python bytecode transformation and graph compilation. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2 (ASPLOS '24)*. ACM, April 2024.
- [2] Franz Baader, Bartosz Bednarczyk, and Sebastian Rudolph. Satisfiability and query answering in description logics with global and local cardinality constraints. In Giuseppe De Giacomo, Alejandro Catalá, Bistra Dilkina, Michela Milano, Senén Barro, Alberto Bugarín, and Jérôme Lang, editors, *ECAI 2020 - 24th European Conference on Artificial Intelligence, 29 August-8 September 2020, Santiago de Compostela, Spain, August 29 - September 8, 2020 - Including 10th Conference on Prestigious Applications of Artificial Intelligence (PAIS 2020)*, volume 325 of *Frontiers in Artificial Intelligence and Applications*, pages 616–623. IOS Press, 2020.
- [3] Franz Baader, Ian Horrocks, Carsten Lutz, and Uli Sattler. *Introduction to Description Logic*. Cambridge University Press, 2017.
- [4] Pablo Barceló, Egor V. Kostylev, Mikaël Monet, Jorge Pérez, Juan L. Reutter, and Juan Pablo Silva. The logical expressiveness of graph neural networks. In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net, 2020.
- [5] Pablo Barceló, Egor V. Kostylev, Mikaël Monet, Jorge Pérez, Juan L. Reutter, and Juan Pablo Silva. Gnn-logic. <https://github.com/juanpablos/GNN-logic.git>, 2021.
- [6] Bartosz Bednarczyk, Maja Orłowska, Anna Pacanowska, and Tony Tan. On classical decidable logics extended with percentage quantifiers and arithmetics. In Mikolaj Bojanczyk and Chandra Chekuri, editors, *41st IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science, FSTTCS 2021, December 15-17, 2021, Virtual Conference*, volume 213 of *LIPIcs*, pages 36:1–36:15. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2021.
- [7] Michael Benedikt, Chia-Hsuan Lu, Boris Motik, and Tony Tan. Decidability of graph neural networks via logical characterizations. In Karl Bringmann, Martin Grohe, Gabriele Puppis, and Ola Svensson, editors, *51st International Colloquium on Automata, Languages, and Programming, ICALP 2024, July 8-12, 2024, Tallinn, Estonia*, volume 297 of *LIPIcs*, pages 127:1–127:20. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2024.
- [8] Michael Benedikt, Chia-Hsuan Lu, and Tony Tan. Decidability of graph neural networks via logical characterizations. *CoRR*, abs/2404.18151v4, 2025.
- [9] Patrick Blackburn, Maarten de Rijke, and Yde Venema. *Modal Logic*, volume 53 of *Cambridge Tracts in Theoretical Computer Science*. Cambridge University Press, 2001.
- [10] Lucas C. Cordeiro, Matthew L. Daggitt, Julien Girard-Satabin, Omri Isac, Taylor T. Johnson, Guy Katz, Ekaterina Komendantskaya, Augustin Lemesle, Edoardo Manino, Artjoms Sinkarovs, and Haoze Wu. Neural network verification is a programming language challenge. *CoRR*, abs/2501.05867, 2025.
- [11] David J. Tena Cucala and Bernardo Cuenca Grau. Bridging max graph neural networks and Datalog with negation. In Pierre Marquis, Magdalena Ortiz, and Maurice Pagnucco, editors, *Proceedings of the 21st International Conference on Principles of Knowledge Representation and Reasoning, KR 2024, Hanoi, Vietnam. November 2-8, 2024*, 2024.

- [12] Stéphane Demri and Denis Lugiez. Complexity of modal logics with presburger constraints. *J. Appl. Log.*, 8(3):233–252, 2010.
- [13] European Parliament. Artificial Intelligence Act, 2024.
- [14] Pietro Galliani, Oliver Kutz, and Nicolas Troquard. Succinctness and complexity of ALC with counting perceptrons. In Pierre Marquis, Tran Cao Son, and Gabriele Kern-Isberner, editors, *Proceedings of the 20th International Conference on Principles of Knowledge Representation and Reasoning, KR 2023, Rhodes, Greece, September 2-8, 2023*, pages 291–300, 2023.
- [15] Amir Gholami, Sehoon Kim, Zhen Dong, Zhewei Yao, Michael W Mahoney, and Kurt Keutzer. A survey of quantization methods for efficient neural network inference. In *Low-power computer vision*, pages 291–326. Chapman and Hall/CRC, 2022.
- [16] Justin Gilmer, Samuel S. Schoenholz, Patrick F. Riley, Oriol Vinyals, and George E. Dahl. Neural message passing for quantum chemistry. In Doina Precup and Yee Whye Teh, editors, *Proceedings of the 34th International Conference on Machine Learning, ICML 2017, Sydney, NSW, Australia, 6-11 August 2017*, volume 70 of *Proceedings of Machine Learning Research*, pages 1263–1272. PMLR, 2017.
- [17] Thomas A. Henzinger, Mathias Lechner, and Dorde Zikelic. Scalable verification of quantized neural networks. In *Thirty-Fifth AAAI Conference on Artificial Intelligence, AAAI 2021, Thirty-Third Conference on Innovative Applications of Artificial Intelligence, IAAI 2021, The Eleventh Symposium on Educational Advances in Artificial Intelligence, EAAI 2021, Virtual Event, February 2-9, 2021*, pages 3787–3795. AAAI Press, 2021.
- [18] Pei Huang, Haoze Wu, Yuting Yang, Ieva Daukantas, Min Wu, Yedi Zhang, and Clark W. Barrett. Towards efficient verification of quantized neural networks. In Michael J. Wooldridge, Jennifer G. Dy, and Sriraam Natarajan, editors, *Thirty-Eighth AAAI Conference on Artificial Intelligence, AAAI 2024, Thirty-Sixth Conference on Innovative Applications of Artificial Intelligence, IAAI 2024, Fourteenth Symposium on Educational Advances in Artificial Intelligence, EAAI 2024, February 20-27, 2024, Vancouver, Canada*, pages 21152–21160. AAAI Press, 2024.
- [19] Benoit Jacob, Skirmantas Kligys, Bo Chen, Menglong Zhu, Matthew Tang, Andrew Howard, Hartwig Adam, and Dmitry Kalenichenko. Quantization and training of neural networks for efficient integer-arithmetic-only inference. In *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2704–2713, 2018.
- [20] Viktor Kuncak and Martin Rinard. Towards efficient satisfiability checking for boolean algebra with presburger arithmetic. In Frank Pfenning, editor, *Automated Deduction – CADE-21*, pages 215–230, Berlin, Heidelberg, 2007. Springer Berlin Heidelberg.
- [21] Rafael Menezes, Mohannad Aldughaim, Bruno Farias, Xianzhiyu Li, Edoardo Manino, Fedor Shmarov, Kunjian Song, Franz Brauße, Mikhail R. Gadelha, Norbert Tihanyi, Konstantin Korovin, and Lucas C. Cordeiro. ESBMC 7.4: Harnessing the Power of Intervals. In *30th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS’24)*, volume 14572 of *Lecture Notes in Computer Science*, page 376–380. Springer, 2024.
- [22] Paulius Micikevicius, Dusan Stosic, Neil Burgess, Marius Cornea, Pradeep Dubey, Richard Grisenthwaite, Sangwon Ha, Alexander Heinecke, Patrick Judd, John Kamalu, Naveen Mellem-pudi, Stuart F. Oberman, Mohammad Shoeybi, Michael Y. Siu, and Hao Wu. FP8 formats for deep learning. *CoRR*, abs/2209.05433, 2022.
- [23] Markus Nagel, Marios Fournarakis, Rana Ali Amjad, Yelysei Bondarenko, Mart van Baalen, and Tijmen Blankevoort. A white paper on neural network quantization. *ArXiv*, abs/2106.08295, 2021.
- [24] Pierre Nunn, Marco Sälzer, François Schwarzentruher, and Nicolas Troquard. A logic for reasoning about aggregate-combine graph neural networks. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence, IJCAI 2024, Jeju, South Korea, August 3-9, 2024*, pages 3532–3540. ijcai.org, 2024.

- [25] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [26] PyTorch Team. Quantization — PyTorch 2.x Documentation. <https://pytorch.org/docs/stable/quantization.html>, 2024. Accessed: 2025-05-16.
- [27] PyTorch Team. torch.quantize_per_tensor — pytorch 2.x documentation. https://pytorch.org/docs/stable/generated/torch.quantize_per_tensor.html#torch-quantize-per-tensor, 2024. Accessed: 2025-05-16.
- [28] PyTorch Team. torch.tensor — pytorch 2.x documentation. <https://pytorch.org/docs/stable/tensors.html#torch.Tensor>, 2024. Accessed: 2025-05-16.
- [29] Patrick Reiser, Marlen Neubert, André Eberhard, Luca Torresi, Chen Zhou, Chen Shao, Housam Metni, Clint van Hoesel, Henrik Schopmans, Timo Sommer, and Pascal Friederich. Graph neural networks for materials science and chemistry. *Communications Materials*, 3(93), 2022.
- [30] Amirreza Salamat, Xiao Luo, and Ali Jafari. Heterographrec: A heterogeneous graph-based neural networks for social recommendations. *Knowl. Based Syst.*, 217:106817, 2021.
- [31] Marco Sälzer and Martin Lange. Reachability is NP-complete even for the simplest neural networks. In Paul C. Bell, Patrick Totzke, and Igor Potapov, editors, *Reachability Problems - 15th International Conference, RP 2021, Liverpool, UK, October 25-27, 2021, Proceedings*, volume 13035 of *Lecture Notes in Computer Science*, pages 149–164. Springer, 2021.
- [32] Marco Sälzer, François Schwarzentruher, and Nicolas Troquard. Verifying quantized graph neural networks is pspace-complete. *CoRR*, abs/2502.16244, 2025.
- [33] Luiz H. Sena, Xidan Song, Erickson H. da S. Alves, Iury Bessa, Edoardo Manino, and Lucas C. Cordeiro. Verifying Quantized Neural Networks using SMT-Based Model Checking. *CoRR*, abs/2106.05997, 2021.
- [34] Marco Sälzer and Martin Lange. Fundamental limits in formal verification of message-passing neural networks. In *ICLR*, 2023.
- [35] Shyam Anil Tailor, Javier Fernandez-Marques, and Nicholas Donald Lane. Degree-quant: Quantization-aware training for graph neural networks. In *International Conference on Learning Representations*, 2021.
- [36] Stephan Tobies. The complexity of reasoning with cardinality restrictions and nominals in expressive description logics. *J. Artif. Intell. Res.*, 12:199–217, 2000.
- [37] G. S. Tseitin. *On the Complexity of Derivation in Propositional Calculus*, pages 466–483. Springer Berlin Heidelberg, Berlin, Heidelberg, 1983.
- [38] Hao Wu, Patrick Judd, Xiaojie Zhang, Mikhail Isaev, and Paulius Micikevicius. Integer quantization for deep learning inference: Principles and empirical evaluation. *CoRR*, abs/2004.09602, 2020.
- [39] Jiacheng Xiong, Zhaoping Xiong, Kaixian Chen, Hualiang Jiang, and Mingyue Zheng. Graph neural networks for automated de novo drug design. *Drug Discovery Today*, 26(6):1382–1393, 2021.
- [40] Zi Ye, Yogan Jaya Kumar, Goh Ong Sing, Fengyan Song, and Junsong Wang. A comprehensive survey of graph neural networks for knowledge graphs. *IEEE Access*, 10:75729–75741, 2022.
- [41] Yedi Zhang, Zhe Zhao, Guangke Chen, Fu Song, Min Zhang, Taolue Chen, and Jun Sun. Qvip: An ILP-based formal verification approach for quantized neural networks. In *Proceedings of the 37th IEEE/ACM International Conference on Automated Software Engineering, ASE ’22*, New York, NY, USA, 2023. Association for Computing Machinery.

- 501 [42] Jie Zhou, Ganqu Cui, Shengding Hu, Zhengyan Zhang, Cheng Yang, Zhiyuan Liu, Lifeng 1
502 Wang, Changcheng Li, and Maosong Sun. Graph neural networks: A review of methods and
503 applications. *AI open*, 1:57–81, 2020.
- 504 [43] Zeyu Zhu, Fanrong Li, Zitao Mo, Qinghao Hu, Gang Li, Zejian Liu, Xiaoyao Liang, and Jian
505 Cheng. A²Q: Aggregation-aware quantization for graph neural networks. In *The Eleventh*
506 *International Conference on Learning Representations*, 2023.

A Proofs of statements in the main text ¹

Lemma 3. *Let φ be a formula of modal $q\mathcal{L}$. The formulas φ and $\text{mod2expr}(\varphi)$ are equivalent.* ²

Proof. We have to prove that for all G, u , we have $G, u \models \varphi$ iff $G, u \models \text{mod2expr}(\varphi)$. We proceed by induction on φ . ³
⁴

- The base case is obvious: $G, u \models \varphi$ iff $G, u \models \text{mod2expr}(\varphi)$ is $G, u \models \varphi$ iff $G, u \models \text{mod2expr}(\varphi)$.
- $G, u \models \neg\varphi$ iff $G, u \not\models \varphi$
iff (by induction) $G, u \not\models \text{mod2expr}(\varphi)$
iff (by writing $\text{mod2expr}(\varphi) = \vartheta \geq 0$) $G, u \models \vartheta \geq 0$
iff $G, u \models \vartheta < 0$
iff $G, u \models \vartheta \leq -1$ (because we suppose that ϑ takes its value in the integers)
iff $G, u \models \vartheta + 1 \leq 0$
iff $G, u \models -\vartheta - 1 \geq 0$.
- $G, u \models (\varphi_1 \vee \varphi_2)$
iff $G, u \models \varphi_1$ or $G, u \models \varphi_2$
iff $G, u \models (\vartheta_1 \geq 0)$ or $G, u \models (\vartheta_2 \geq 0)$
iff $G, u \models \vartheta_1 + \text{ReLU}(\vartheta_2 - \vartheta_1) \geq 0$
Indeed, $(\Rightarrow)$ if $G, u \models (\vartheta_1 \geq 0)$ then $G, u \models \vartheta_1 + \text{ReLU}(\vartheta_2 - \vartheta_1) \geq \vartheta_1 \geq 0$.
If $G, u \models (\vartheta_2 \geq 0)$ and $G, u \models (\vartheta_1 < 0)$ then $G, u \models \vartheta_1 + \text{ReLU}(\vartheta_2 - \vartheta_1) = \vartheta_1 + \vartheta_2 - \vartheta_1 = \vartheta_2 \geq 0$.
 $(\Leftarrow)$ Conversely, by contrapositive, if $G, u \models (\vartheta_2 < 0)$ and $G, u \models (\vartheta_1 < 0)$, then $G, u \models \vartheta_1 + \text{ReLU}(\vartheta_2 - \vartheta_1) = \vartheta_1 + \vartheta_2 - \vartheta_1 = \vartheta_2 < 0$ or $G, u \models \vartheta_1 + \text{ReLU}(\vartheta_2 - \vartheta_1) = \vartheta_1 + 0 = \vartheta_1 < 0$. In the two cases, $G, u \models \vartheta_1 + \text{ReLU}(\vartheta_2 - \vartheta_1) < 0$.
- $G, u \models \Diamond^{\geq k}\varphi$ iff the number of vertices v that are successors of u and with $G, v \models \varphi$ is greater than k
iff the number of vertices v that are successors of u and with $G, v \models \text{mod2expr}(\varphi)$ is greater than k
iff (written $\vartheta \geq 0$) iff the number of vertices v that are successors of u and with $G, v \models \vartheta \geq 0$ is greater than k
iff the number of vertices v that are successors of u and with $G, v \models \text{ReLU}(\vartheta + 1) - \text{ReLU}(\vartheta) = 1$ is greater than k (since we know by defining of modal $q\mathcal{L}$ that ϑ takes its value in integers)
iff $G, u \models \text{agg}(\text{ReLU}(\vartheta + 1) - \text{ReLU}(\vartheta)) \geq k$
iff $G, u \models \text{mod2expr}(\Diamond^{\geq k}\varphi)$
- Other cases are similar.

□

Proposition 7. *The number of Hintikka sets is bounded by $2^{n|\varphi|}$ where $|\varphi|$ is the size of φ , and n is the bitwidth of $\mathbb{K}$.* ⁵

Proof. For each expression ϑ , we choose a number in $\mathbb{K}$. There is 2^n different numbers. There are $|\varphi|$ number of expressions. So we get $(2^n)^{|\varphi|} = 2^{n|\varphi|}$ possible choices for a Hintikka set. ⁶
□

Proposition 8. *If bitwidth n is in unary, and if $\mathbb{K}$ saturates, then satisfiability in $QFBAPA_{\mathbb{K}}$ is in NP.* ⁷

548 *Proof.* Here is a non-deterministic algorithm for the satisfiability problem in $QFBAPA_{\mathbb{K}}$. 1

549 1. Let χ be a $QFBAPA_{\mathbb{K}}$ formula. 2

550 2. For each set expression B appearing in some $|B|$, guess a non-negative integer number k_B 3
551 in $\mathbb{K}$.

552 3. Let χ' be a (grounded) formula in which we replaced $|B|$ by k_B . 4

553 4. Check that χ' is true (can be done in poly-time since χ' is a grounded formula, it is a 5
554 Boolean formula on variable-free equations and inequations in $\mathbb{K}$).

555 5. If not we reject. 6

6. We now build a standard $QFBAPA$ formula $\delta = \bigwedge_B \text{constraint}(B)$ where: 7

$$\text{constraint}(B) = \begin{cases} |B| = k_B & \text{if } k_B < \infty_{\mathbb{K}} \\ |B| \geq \text{limit} & \text{if } k_B = +\infty_{\mathbb{K}} \end{cases} \quad 8$$

556 where *limit* is the maximum number that is considered as infinity in $\mathbb{K}$. 9

557 7. Run a non-deterministic poly-time algorithm for the $QFBAPA$ satisfiability on δ . Accepts if 10
558 it accepts. Otherwise reject.

559 The algorithm runs in poly-time. Guessing a number n_B is in poly-time since it consists in guessing 11
560 n bits (n in unary). Step 4 is just doing the computations in $\mathbb{K}$. In Step 6, δ can be computed in
561 poly-time.

562 If χ is $QFBAPA_{\mathbb{K}}$ satisfiable, then there is a solution σ such that $\sigma \models \chi$. At step 2, we guess 12
563 $n_B = |\sigma(B)|_{\mathbb{K}}$. The algorithm accepts the input.

564 Conversely, if the algorithm accepts its input, χ' is true for the chosen values n_B . δ is satisfiable. So 13
565 there is a solution σ such that $\sigma \models \delta$. By the definition of *constraint*, $\sigma \models \chi$. $\square$

566 **Remark 15.** If the number n of bits to represent $\mathbb{K}$ is given in unary and if $\mathbb{K}$ is "modulo", 14
567 then the satisfiability problem in $QFBAPA_{\mathbb{K}}$ is also in NP. The proof is similar except than now
568 $\text{constraint}(B) = (|B| = k_B + Ld_B)$ where d_B is a new variable.

569 **Proposition 9.** $tr(\varphi)$ is computable in exponential-time in $|\varphi|$ and n . 15

570 *Proof.* In order to create $tr(\varphi)$, we write an algorithm where each big conjunction, big disjunction, 16
571 big union and big sum is replaced by a loop. For instance, $\bigwedge_{H \neq H'}$ is replaced by two inner loops
572 over Hintikka sets. Note that we create check whether a candidate H is a Hintikka set in exponential
573 time in n since Point 4 can be checked in exponential time in n (thanks to our loose assumption on
574 the computability of $[[\alpha]]$ in exponential time in n . There are $2^{n|\varphi|}$ many of them. In the same way,
575 $\bigwedge_{k \in \mathbb{K}}$ is a loop over 2^n values. There is a constant number of nested loops, each of them iterating
576 over an exponential number (in n and $|\varphi|$ of elements). QED. $\square$

577 **Proposition 10.** Let φ be a formula of $q\mathcal{L}$. φ is satisfiable iff $tr(\varphi)$ is $QFBAPA_{\mathbb{K}}$ satisfiable. 17

578 *Proof.* $\Rightarrow$ Let G, u such that $G, u \models \varphi$. We set $\sigma(X_{\vartheta'=k}) := \{v \mid [[\vartheta']]_{G,v} = k\}$ and $\sigma(X_H) = 18$

579 $\{v \mid G, v \models H\}$ where $G, u \models H$ means that for all $\vartheta' = k \in H$, we have $[[\vartheta']]_{G,v} = k$. For all
580 Hintikka sets H such that there is v such that $G, v \models H$, we set: $\sigma(S_H) := \{w \mid vEw\}$.
581 We check that $\sigma \models tr(\varphi)$. First, σ satisfies Formulas 1 and 2 by definition of σ . Now, σ also satisfies 19
582 Formula 3. Indeed, if $agg(\vartheta') = k \in H$, then if there is no H -vertex in G then the implication is
583 true. Otherwise, consider the H -vertex v . But, then by definition of $X_{agg(\vartheta')=k}$, $[[agg(\vartheta')]]_{G,v} = k$.
584 But then the semantics of agg exactly corresponds to $\sum_{k' \in \mathbb{K}} |S_H \cap X_{\vartheta=k'}| \times k' = k$. Indeed, each
585 $S_H \cap X_{\vartheta=k'}$ -successor contributes with k' . Thus, the contribution of successors where ϑ is k' is
586 $|S_H \cap X_{\vartheta=k'}| \times k'$.

587 Formula 4 is also satisfied by σ . Actually, let k such that $\sigma \models X_{agg_{\forall}(\vartheta)=k} = \mathcal{U}$. This means that 20
588 the value of $agg_{\forall}(\vartheta)$ (which does not depend on a specific vertex u but only on G) is k . The sum
589 $\sum_{k' \in \mathbb{K}} |X_{\vartheta=k'}| \times k' = k$ is the semantics of $agg_{\forall}(\vartheta) = k$.

590 Finally, as $G, u \models \varphi$, and φ is of the form $\vartheta \geq k$, there is $k' \geq k$ such that $[[\vartheta]]_{G,u} = k'$. So 1
 591 $X_{\vartheta=k'} \neq \emptyset$.

592 $\Leftarrow$ Conversely, consider a solution σ of $tr(\varphi)$. We construct a graph $G = (V, E)$ as follows. 2

$$\begin{aligned} V &:= \sigma(\mathcal{U}) \\ E &:= \{(u, v) \mid \text{for some } H, u \in \sigma(X_H) \text{ and } v \in \sigma(S_H)\} \\ \ell(v)_i &:= k \text{ where } v \in X_{x_i=k} \end{aligned} \quad 3$$

593 i.e. the set of vertices is the universe, and we add an edge between any H -vertex u and a vertex 4
 594 $v \in \sigma(S_H)$, and the labeling for features is directly given $X_{x_i=k}$. Note that the labeling is well-
 595 defined because of formulas 1 and 2.

596 As $\sigma \models |X_\varphi| \geq 1$, there exists $u \in \sigma(X_\varphi)$. Let us prove that $G, u \models \varphi$. By induction on ϑ' , we 8
 597 prove that $u \in X_{\vartheta'=k}$ implies $[[\vartheta']]_{G,u} = k$. The base case is obtained via the definition of ℓ . Cases
 598 for $+$, $\times$ and α are obtained because each vertices is in some $\sigma(X_H)$ for some H . As the definition of
 599 Hintikka set takes care of the semantics of $+$, $\times$ and α , we have $[[\vartheta_1 + \vartheta_2]]_{G,u} = [[\vartheta_1]]_{G,u} + [[\vartheta_2]]_{G,u}$,
 600 etc.

601 $[[agg(\vartheta)]]_{G,u} = \sum_{v|uEv} [[\vartheta]]_{G,v}$ and $[[agg_{\forall}(\vartheta)]]_{G,u} = \sum_{v \in V} [[\vartheta]]_{G,v}$ hold because of σ satisfies 6
 602 respectively formula 3 and 4. $\square$

603 **Theorem 12.** The satisfiability problem in $q\mathcal{L}$ is NEXPTIME-hard, and so is VT3 . VT1 and VT2 are 7
 604 coNEXPTIME -hard.

605 *Proof.* We reduce the NEXPTIME-hard problem of deciding whether a domino system $\mathcal{D} =$ 8
 606 (D, V, H) , given an initial condition $w_0 \dots w_{n-1} \in D^n$, can tile an exponential torus [36]. In
 607 the domino system, D is the set of tile types, and V and H respectively are the respectively vertical
 608 and horizontal color compatibility relations. We are going to write a set of modal $q\mathcal{L}$ formulas that
 609 characterize the torus $\mathbb{Z}^{2n+1} \times \mathbb{Z}^{2n+1}$ and the domino system. We use $2n + 2$ features. We use
 610 $x_0, \dots, x_{n-1}$, and $x'_0, \dots, x'_{n-1}$, to hold the (binary-encoded) coordinates of vertices in the torus. We
 611 use the feature x_N to denote a vertex ‘on the way north’ (when $x_N = 1$) and x_E to denote a vertex
 612 ‘on the way east’ (when $x_E = 1$), with abbreviations $\varphi_N := x_N = 1$, and $\varphi_E := x_E = 1$. See
 613 Figure 2.

614 For every $n \in \mathbb{N}$, we define the following set of formulas. $T_n =$ 9

$$\left\{ \begin{array}{ll} \Box_g(x_N = 1 \vee x_N = 0) & , \quad \Box_g(x_E = 1 \vee x_E = 0), \\ \Box_g(\bigwedge_{k=0}^{n-1}(x_i = 1 \vee x_i = 0)) & , \quad \Box_g(\bigwedge_{k=0}^{n-1}(x'_i = 1 \vee x'_i = 0)), \\ \Box_g(\neg(x_N = 1 \wedge x_E = 1)) & , \quad \Box_g(\neg(\varphi_N \vee \varphi_E) \rightarrow agg(1) = 2), \\ \Box_g(\neg(\varphi_N \vee \varphi_E) \rightarrow (agg(x_N) = 1)) & , \quad \Box_g(\neg(\varphi_N \vee \varphi_E) \rightarrow (agg(x_E) = 1)), \\ \Box_g(\varphi_N \rightarrow agg(1) = 1) & , \quad \Box_g(\varphi_E = 1 \rightarrow agg(1) = 1), \\ \Diamond_g^1 \varphi_{(0,0)} & , \quad \Diamond_g^1 \varphi_{(2^n-1, 2^n-1)}, \\ \Box_g(\neg(\varphi_N \vee \varphi_E) \rightarrow \varphi_{east}) & , \quad \Box_g(\neg(\varphi_N \vee \varphi_E) \rightarrow \varphi_{north}), \\ \Diamond_g^{\leq 2^n \times 2^n} \neg(\varphi_N \vee \varphi_E), & \Diamond_g^{\leq 2^n \times 2^n} \varphi_N, \quad \Diamond_g^{\leq 2^n \times 2^n} \varphi_E \end{array} \right\} \quad 10$$

615 where $\varphi_{(0,0)} := \bigwedge_{k=0}^{n-1} x_i = 0 \wedge \bigwedge_{k=0}^{n-1} x'_i = 0$, and $\varphi_{(2^n-1, 2^n-1)} := \bigwedge_{k=0}^{n-1} x_i = 1 \wedge \bigwedge_{k=0}^{n-1} x'_i = 1$ 11
 616 represent two nodes, namely those at coordinates $(0, 0)$ and $(2^n - 1, 2^n - 1)$. The formulas φ_{north} and
 617 φ_{east} enforce constraints on the coordinates of states, such that going north increases the coordinate
 618 encoding using the x_i features by one, leaving the x'_i features unchanged, and going east increases
 619 coordinate encoding using the x'_i features by one, leaving the x_i features unchanged. For every

620 formula φ , $\forall east.\varphi$ stands for $\Box(\varphi_E \rightarrow \Box\varphi)$ and $\forall north.\varphi$ stands for $\Box(\varphi_N \rightarrow \Box\varphi)$. ¹

$$\begin{aligned}
\varphi_{north} &:= \bigwedge_{k=0}^{n-1} \left(\bigwedge_{j=0}^{k-1} (x_j = 1) \right) \rightarrow (((x_k = 1) \rightarrow \forall north.(x_k = 0)) \wedge ((x_k = 0) \rightarrow \forall north.(x_k = 1))) \wedge \\
&\quad \bigwedge_{k=0}^{n-1} \left(\bigvee_{j=0}^{k-1} (x_j = 0) \right) \rightarrow (((x_k = 1) \rightarrow \forall north.(x_k = 1)) \wedge ((x_k = 0) \rightarrow \forall north.(x_k = 0))) \wedge \\
&\quad \bigwedge_{k=0}^{n-1} (((x'_k = 1) \rightarrow \forall north.(x'_k = 1)) \wedge ((x'_k = 0) \rightarrow \forall north.(x'_k = 0))) \\
\varphi_{east} &:= \bigwedge_{k=0}^{n-1} \left(\bigwedge_{j=0}^{k-1} (x'_j = 1) \right) \rightarrow (((x'_k = 1) \rightarrow \forall east.(x'_k = 0)) \wedge ((x'_k = 0) \rightarrow \forall east.(x'_k = 1))) \wedge \\
&\quad \bigwedge_{k=0}^{n-1} \left(\bigvee_{j=0}^{k-1} (x'_j = 0) \right) \rightarrow (((x'_k = 1) \rightarrow \forall east.(x'_k = 1)) \wedge ((x'_k = 0) \rightarrow \forall east.(x'_k = 0))) \wedge \\
&\quad \bigwedge_{k=0}^{n-1} (((x_k = 1) \rightarrow \forall east.(x_k = 1)) \wedge ((x_k = 0) \rightarrow \forall east.(x_k = 0)))
\end{aligned}$$

621 The problem of deciding whether a domino system $\mathcal{D} = (D, V, H)$, given an initial condition ³
622 $w_0 \dots w_{n-1} \in D^n$, can tile a torus of exponential size can be reduced to the problem satisfiability in
623 $q\mathcal{L}$, checking the satisfiability of the set of formulas $T(n, \mathcal{D}, w) = T_n \cup T_{\mathcal{D}} \cup T_w$, where T_n is as
624 above, $T_{\mathcal{D}}$ encodes the domino system, and T_w encodes the initial condition as follows. We define

$$T_{\mathcal{D}} = \{ \quad \Box_g(\bigwedge_{d \in D} (x_d = 1 \vee x_d = 0)), \\
\quad \Box_g(\neg(\varphi_N \vee \varphi_E) \rightarrow (\bigvee_{d \in D} \varphi_d)), \\
\quad \Box_g(\neg(\varphi_N \vee \varphi_E) \rightarrow (\bigwedge_{d \in D} \bigwedge_{d' \in D \setminus \{d\}} \neg(\varphi_d \wedge \varphi_{d'}))), \\
\quad \Box_g(\bigwedge_{d \in D} (\varphi_d \rightarrow (\forall east. \bigvee_{(d, d') \in H} \varphi_{d'}))), \\
\quad \Box_g(\bigwedge_{d \in D} (\varphi_d \rightarrow (\forall north. \bigvee_{(d, d') \in V} \varphi_{d'}))) \quad \}$$

625 where for every $d \in D$, there is a feature x_d and $\varphi_d := x_d = 1$. Finally, we define ⁵

$$T_w = \{ \quad \Box_g(\varphi(0,0) \rightarrow \varphi_{w_0}), \dots, \Box_g(\varphi(n-1,0) \rightarrow \varphi_{w_{n-1}}) \quad \}$$

626 The size of $T(n, \mathcal{D}, w)$ is polynomial in the size of the tiling problem instance, that is in $|D| + |H| +$ ⁷
627 $|V| + n$. The rest of the proof is analogous to the proof of [36, Corollary 3.9]. The NEXPTIME-
628 hardness of $q\mathcal{L}$ follows from Lemma 3 and [36, Corollary 3.3] stating the NEXPTIME-hardness of
629 deciding whether a domino system with initial condition can tile a torus of exponential size.

630 For the complexity of ACR-GNN verification tasks, we observe the following. ⁸

- 631 1. We reduce the satisfiability problem in (modal) $q\mathcal{L}$ (restricted to graded modal logic + graded ⁹
632 universal modality, because it is sufficient to encode the tiling problem) to VT3 in poly-time
633 as follows. Let φ be a $q\mathcal{L}$. We build in poly-time an ACR-GNN $\mathcal{A}$ that recognizes all
634 pointed graphs. We have φ is satisfiable iff $[[\varphi]] \cap [[\mathcal{A}]] \neq \emptyset$ So VT3 is NEXPTIME-hard.
- 635 2. The validity problem of $q\mathcal{L}$ (dual problem of the satisfiability problem, i.e., given a formula
636 φ , is φ true in all pointed graphs G, u ?) is coNEXPTIME-hard. We reduce the validity
637 problem of $q\mathcal{L}$ to VT2. Let φ be a $q\mathcal{L}$ formula. We construct an ACR-GNN $\mathcal{A}$ that accepts
638 all pointed graphs. We have φ is valid iff $[[\mathcal{A}]] \subseteq [[\varphi]]$. So VT2 is coNEXPTIME-hard.
- 639 3. We reduce the validity problem of $q\mathcal{L}$ to VT1. Let ψ be a $q\mathcal{L}$ formula. (again in graded
640 modal logic + graded global modalities). So by [4], We construct in poly-time an ACR-
641 GNN $\mathcal{A}$ that is equivalent to ψ (by [4]). We have ψ is valid iff $[[\top]] \subseteq [[\mathcal{A}]]$. So VT1 is
642 coNEXPTIME-hard.

643

□

644 **Theorem 14.** *The satisfiability problems with bounded number of vertices are NP-complete.*

645 *Proof.* NP upper bound is obtained by guessing a graph with at most N vertices and then check that φ 1
 646 holds. The obtained algorithm is non-deterministic, runs in poly-time and decides the satisfiability
 647 problem with bounded number of vertices. NP-hardness already holds for *agg*-free formulas by
 648 reduction from SAT for propositional logic (the reduction is *mod2expr*, see Lemma 3). $\square$

649 B Checking distributivity 2

650 We provide C source code for checking distributivity. The reader may run the model checker ESBMC 3
 651 on it to see whether distributivity holds or not.

652 C Extension of logic $K^\sharp$ and ACR-GNNs over $\mathbb{Z}$ 4

653 A (labeled directed) graph G is a tuple (V, E, ℓ) such that V is a finite set of vertices, $E \subseteq V \times V$ a 5
 654 set of directed edges and ℓ is a mapping from V to a valuation over a set of atomic propositions. We
 655 write $\ell(u)(p) = 1$ when atomic proposition p is true in u , and $\ell(u)(p) = 0$ otherwise. Given a graph
 656 G and vertex $u \in V$, we call (G, u) a *pointed graph*.

657 C.1 Logic 6

658 Consider a countable set Ap of propositions. We define the language of logic $K^{\sharp, \#}_g$ as the set of 7
 659 formulas generated by the following BNF:

$$\begin{aligned} \varphi &::= p \mid \neg\varphi \mid \varphi \vee \varphi \mid \xi \geq 0 \\ \xi &::= c \mid \mathbb{1}\varphi \mid \sharp\varphi \mid \sharp_g\varphi \mid \xi + \xi \mid c \times \xi \end{aligned} \quad 8$$

660 where p ranges over Ap , and c ranges over $\mathbb{Z}$. We assume that all formulas φ are represented as 9
 661 directed acyclic graph (DAG) and refer by *the size of φ* to the size of its DAG representation.

662 Atomic formulas are propositions p , inequalities and equalities of linear expressions. We consider 10
 663 linear expressions over $\mathbb{1}\varphi$ and $\sharp\varphi$ and $\sharp_g\varphi$. The number $\mathbb{1}\varphi$ is equal to 1 if φ holds in the current
 664 world and equal 0 otherwise. The number $\sharp\varphi$ is the number of successors in which φ hold. The
 665 number $\sharp_g\varphi$ is the number of worlds in the model in which φ hold. The language seems strict but we
 666 write $\xi_1 \leq \xi_2$ for $\xi_2 - \xi_1 \geq 0$, $\xi = 0$ for $(\xi \geq 0) \wedge (\neg\xi \geq 0)$, etc.

667 As in modal logic, a formula φ is evaluated in a pointed graph (G, u) (also known as pointed Kripke 11
 668 model). We define the truth conditions $(G, u) \models \varphi$ (φ is true in u) by

$$\begin{aligned} (G, u) \models p & \quad \text{if } \ell(u)(p) = 1, \\ (G, u) \models \neg\varphi & \quad \text{if it is not the case that } (G, u) \models \varphi, \\ (G, u) \models \varphi \wedge \psi & \quad \text{if } (G, u) \models \varphi \text{ and } (G, u) \models \psi, \\ (G, u) \models \xi \geq 0 & \quad \text{if } [[\xi]]_{G,u} \geq 0, \end{aligned} \quad 13$$

669 12

670 and the semantics $[[\xi]]_{G,u}$ (the value of ξ in u) of an expression ξ by mutual induction on φ and ξ as 14
 671 follows.

$$\begin{aligned} [[c]]_{G,u} &= c, \\ [[\xi_1 + \xi_2]]_{G,u} &= [[\xi_1]]_{G,u} + [[\xi_2]]_{G,u}, \\ [[c \times \xi]]_{G,u} &= c \times [[\xi]]_{G,u}, \\ [[\mathbb{1}\varphi]]_{G,u} &= \begin{cases} 1 & \text{if } (G, u) \models \varphi \\ 0 & \text{otherwise,} \end{cases} \\ [[\sharp\varphi]]_{G,u} &= |\{v \in V \mid (u, v) \in E \text{ and } (G, v) \models \varphi\}| \\ [[\sharp_g\varphi]]_{G,u} &= |\{v \in V \mid (G, v) \models \varphi\}|. \end{aligned} \quad 16$$

672 15

673 A local modality $\Box\varphi$ can be defined as $\Box\varphi := (-1) \times \sharp(\neg\varphi) \geq 0$. That is, to say that φ holds 17
 674 in all successors, we say that the number of successors in which $\neg\varphi$ holds is zero. Similarly, a
 675 global/universal modality can be defined as $\Box_g\varphi := (-1) \times \sharp_g(\neg\varphi) \geq 0$.

676 C.2 Aggregate-Combine Graph Neural Networks 18

677 In this section, we consider a detailed definition of quantized (global) Aggregate-Combine GNNs 19
 678 (ACR-GNN) [4], also called message passing neural networks [16]. We stick to the former term.

679 A (global) ACR-GNN layer $\mathcal{L} = (comb, agg, agg_g)$ is a tuple where $comb : \mathbb{R}^{2m} \rightarrow \mathbb{R}^n$ is a so-called 1
 680 combination function, agg is a so-called local aggregation function, mapping multisets of vectors
 681 from $\mathbb{R}^m$ to a single vector from $\mathbb{R}^n$, agg_g is a so-called global aggregation function, also mapping
 682 multisets of vectors from $\mathbb{R}^m$ to a single vector from $\mathbb{R}^n$. We call m the input dimension of layer $\mathcal{L}$
 683 and n the output dimension of layer $\mathcal{L}$. Then, a (global) ACR-GNN is a tuple $(\mathcal{L}^{(1)}, \dots, \mathcal{L}^{(L)}, cls)$
 684 where $\mathcal{L}^{(1)}, \dots, \mathcal{L}^{(L)}$ are L ACR-GNN layers and $cls : \mathbb{R}^m \rightarrow \{0, 1\}$ is a classification function. We
 685 assume that all GNNs are well-formed in the sense that output dimension of layer $\mathcal{L}^{(i)}$ matches input
 686 dimension of layer $\mathcal{L}^{(i+1)}$ as well as output dimension of $\mathcal{L}^{(L)}$ matches input dimension of cls .

Let $G = (V, E)$ be a graph with atomic propositions $p_1, \dots, p_k$ and $\mathcal{A} = (\mathcal{L}^{(1)}, \dots, \mathcal{L}^{(L)}, cls)$ 2
 an ACR-GNN. We define $x_0 : V \rightarrow \{0, 1\}^k$, called the initial state of G , as $x_0(u) :=$
 $(\ell(u)(p_1), \dots, \ell(u)(p_k))$ for all $u \in V$. Then, the i -th layer of $\mathcal{A}$ computes an updated state of
 G by

$$x_i(u) := comb(x_{i-1}(u), agg(\{\{x_{i-1}(v) \mid uv \in E\}\}), agg_g(\{\{x_{i-1}(v) \mid v \in V\}\}))$$
 3

687 where agg , agg_g , and $comb$ are respectively the local aggregation, global aggregation and combination 4
 688 function of the i -th layer. Let (G, u) be a pointed graph. We write $\mathcal{A}(G, u)$ to denote the application
 689 of $\mathcal{A}$ to (G, u) , which is formally defined as $\mathcal{A}(G, u) = cls(x_L(u))$ where x_L is the state of G
 690 computed by $\mathcal{A}$ after layer L . Informally, this corresponds to a binary classification of node u .

691 In this work, we exclusively consider the following form of ACR-GNN $\mathcal{A}$: all local and global 5
 692 aggregation functions are given by the sum of all vectors in the input multiset, all combination
 693 functions are given by $comb(x, y, z) = \vec{\sigma}(xC + yA_1 + zA_2 + b)$ where $\vec{\sigma}(x)$ is the componentwise
 694 application of the truncated ReLU $\sigma(x) = \max(0, \min(1, x))$, with matrices C , A_1 and A_2 and
 695 vector b of $\mathbb{K}$ parameters, and where the classification function is $cls(x) = \sum_i a_i x_i \geq 1$, where a_i
 696 are from $\mathbb{K}$ as well.

697 We note $[[\mathcal{A}]]$ the set of pointed graphs (G, u) such that $\mathcal{A}(G, u) = 1$. An ACR-GNN $\mathcal{A}$ is satisfiable 6
 698 if $[[\mathcal{A}]]$ is non-empty. The satisfiability problem for ACR-GNNs is: Given a ACR-GNN $\mathcal{A}$, decide
 699 whether $\mathcal{A}$ is satisfiable.

700 D Capturing GNNs with $K^{\sharp, \sharp_g}$ 7

701 In this section, we demonstrate that the expressive power of (global) ACR-GNNs, as defined in 8
 702 Section C.2 and $K^{\sharp, \sharp_g}$, is equivalent. Informally, this means that for every formula φ of $K^{\sharp, \sharp_g}$, there
 703 exists an ACR-GNNs $\mathcal{A}$ that expresses the same query, and vice-versa. To achieve this, we define a
 704 translation of one into the other and substantiate that this translation is efficient. This enables ways to
 705 employ $K^{\sharp, \sharp_g}$ for reasoning about ACR-GNN.

706 We begin by showing that global ACR-GNNs are at least as expressive as $K^{\sharp, \sharp_g}$. We remark that the 9
 707 arguments are similar to the proof of Theorem 1 in [24].

708 **Theorem 16.** Let $\varphi \in K^{\sharp, \sharp_g}$ be a formula. There is $\mathcal{A}_\varphi$ such that for all pointed graphs (G, u) we 10
 709 have $(G, u) \models \varphi$ if and only if $\mathcal{A}_\varphi(G, u) = 1$. Furthermore, $\mathcal{A}_\varphi$ can be built in polynomial time
 710 regarding the size of φ .

711 *Proof sketch.* We construct a GNN $\mathcal{A}_\varphi$ that evaluates the semantics of a given $K^{\sharp, \sharp_g}$ formula φ for 11
 712 some given pointed graph (G, v) . The network consists of n layers, one for each of the n subformulas
 713 φ_i of φ , ordered so that the subformulas are evaluated based on subformula inclusion. The first
 714 layer evaluates atomic propositions, and each subsequent messages passing layer l_i uses a fixed
 715 combination and fixed aggregation function to evaluate the semantics of φ_i .

716 The correctness follows by induction on the layers: the i -th layer correctly evaluates φ_i at each 12
 717 vertex of G , assuming all its subformulas are correctly evaluated in previous layers. Finally, the
 718 classifying function cls checks whether the n -th dimension of the vector after layer l_n , corresponding
 719 to the semantics of φ_n for the respective vertex v , indicates that $\varphi_n = \varphi$ is satisfied by (G, v) . The
 720 network size is polynomial in the size of φ due to the fact that the total number of layers and their
 721 width is polynomially bounded by the number of subformulas of φ . A full formal proof is given in
 722 Appendix F. $\square$

Theorem 17. Let $\mathcal{A}$ be a GNN. We can compute in polynomial time wrt. $|\mathcal{A}|$ a $K^{\sharp, \sharp_g}$ -formula $\varphi_{\mathcal{A}}$,
represented as a DAG, such that $[[\mathcal{A}]] = [[\varphi_{\mathcal{A}}]]$. 1

Proof sketch. We construct a $K^{\sharp, \sharp_g}$ -formula $\varphi_{\mathcal{A}}$ that simulates the computation of a given GNN
 $\mathcal{A}$. For each layer l_i of the GNN, we define a set of formulas $\varphi_{i,j}$, one per output dimension, that
encode the corresponding node features using linear threshold expressions over the formulas from
the previous layer. At the base, the input features are the atomic propositions $p_1, \dots, p_{m_1}$. 2

Each formula $\varphi_{i,j}$ mirrors the computation of the GNN layer, including combination, local aggre-
gation, and global aggregation. The final classification formula $\varphi_{\mathcal{A}}$ encodes the output of the linear
classifier on the top layer features. Correctness follows from the fact that all intermediate node
features remain Boolean under message passing layers with integer parameters and truncated ReLU
activations. This allows expressing each output as a Boolean formula over the input propositions.
The construction is efficient: by reusing shared subformulas via a DAG representation, the total size
remains polynomial in the size of $\mathcal{A}$. 3

E Complexity of the satisfiability of $K^{\sharp, \sharp_g}$ and its implications for ACR-GNN verification 4

In this section, we establish the complexity of reasoning with $K^{\sharp, \sharp_g}$. 5

Instrumentally, we first show that every $K^{\sharp, \sharp_g}$ formula can be translated into a $K^{\sharp, \sharp_g}$ formula
that is equi-satisfiable, and has a tree representation of size at most polynomial in the size of the
original formula. An analogous result was obtained in [24] for $K^{\sharp}$. It can be shown using a
technique reminiscent of [37] and consisting of factorizing subformulas that are reused in the DAG
by introducing a fresh proposition that is made equivalent. Instead of reusing a ‘possibly large’
subformula, a formula then reuses the equivalent ‘small’ atomic proposition. 6

Lemma 18. The satisfiability problem of $K^{\sharp, \sharp_g}$ reduces to the satisfiability of $K^{\sharp, \sharp_g}$ with tree
formulas in polynomial time. 7

Proof. Let φ be a $K^{\sharp, \sharp_g}$ formula represented as a DAG. For every subformula ψ (i.e., for every node
in the DAG representation of φ), we introduce a fresh atomic proposition p_{ψ} . We can capture the
meaning of these new atomic propositions with the formula $\Phi := \bigwedge_{\psi \text{ node in the DAG}} \text{sem}(\psi)$ where:

$$\begin{aligned} \text{sem}(\psi \vee \chi) &:= p_{\psi \vee \chi} \leftrightarrow (p_{\psi} \vee p_{\chi}) \\ \text{sem}(\neg \psi) &:= p_{\neg \psi} \leftrightarrow \neg p_{\psi} \\ \text{sem}(\xi \geq 0) &:= p_{\xi \geq 0} \leftrightarrow \xi' \geq 0 \end{aligned} \quad \text{9}$$

10

$$\begin{aligned} (c)' &:= c & (\xi_1 + \xi_2)' &:= \xi_1' + \xi_2' & (c \times \xi)' &:= c \times \xi' \\ (\mathbb{1}\psi)' &:= \mathbb{1}p_{\psi} & (\sharp\psi)' &:= \sharp p_{\psi} & (\sharp_g\psi)' &:= \sharp_g p_{\psi} \end{aligned} \quad \text{11}$$

Now, define $\varphi_t := p_{\varphi} \wedge \square_g \Phi$, where $\square_g \Phi := (-1) \times \sharp_g(\neg \Phi) \geq 0$, enforcing the truth of Φ in every
vertex. The size of its tree representation is polynomial in the size of φ . Moreover, φ_t is satisfiable iff
 φ is satisfiable. 12

754

Theorem 19. $K^{\sharp, \sharp_g}_{\text{tree}}$ -satisfiability problem is NEXPTIME-complete. 13

Proof. For membership, we translate the problem into the NEXPTIME-complete problem of concept
description satisfiability in the Description Logics with Global and Local Cardinality Constraints [2],
noted $\mathcal{ALCSCC}^{++}$. The Description Logic $\mathcal{ALCSCC}^{++}$ uses the Boolean Algebra with Presburger
Arithmetic [20], noted QFBAPA, to formalize cardinality constraints. See Section H for a presentation
of $\mathcal{ALCSCC}^{++}$ and QFBAPA. 14

Let φ_0 be a $K^{\sharp, \sharp_g}$ formula. 15

For every proposition p occurring in φ_0 , let A_p be an $\mathcal{ALCSCC}^{++}$ concept name. Let R be an
 $\mathcal{ALCSCC}^{++}$ role name. For every occurrence of $\mathbb{1}\varphi$ in φ_0 , let ZOO_{φ} be an $\mathcal{ALCSCC}^{++}$ role name.
 ZOO -roles stand for ‘zero or one’. The rationale for introducing ZOO -roles is to be able to capture 16

765 the value of $\mathbb{1}\varphi$ in $\mathcal{ALCSCC}^{++}$ making it equal to the number of successors of the role ZOO_φ which 1
 766 can then be used in QFBAPA constraints. A similar trick was used, in another context, in [14]. Here,
 767 we enforce this with the QFBAPA constraint

$$\chi_0 = \bigwedge_{\mathbb{1}\varphi \in \varphi_0} ((|ZOO_\varphi| = 0 \vee |ZOO_\varphi| = 1) \wedge \bar{\tau}(\varphi) = \text{sat}(|ZOO_\varphi| = 1)) \quad 2$$

768 which states that ZOO_φ has zero or one successor, and has one successor exactly when (the translation 3
 769 of) φ is true. The concept descriptions $\bar{\tau}(\varphi)$ and arithmetic expressions $\bar{\tau}(\xi)$ are defined inductively
 770 as follows:

$$\begin{array}{ll} \bar{\tau}(p) &= A_p \\ \bar{\tau}(\neg\varphi) &= \neg\bar{\tau}(\varphi) \\ \bar{\tau}(\varphi \vee \psi) &= \bar{\tau}(\varphi) \sqcup \bar{\tau}(\psi) \\ \bar{\tau}(\xi \geq 0) &= \text{sat}(-1 < \bar{\tau}(\xi)) \\ \bar{\tau}(c) &= c \\ \bar{\tau}(\xi_1 + \xi_2) &= \bar{\tau}(\xi_1) + \bar{\tau}(\xi_2) \\ \bar{\tau}(c \times \xi) &= \bar{\tau}(c \cdot \xi) \\ \bar{\tau}(\sharp\varphi) &= |R \cap \bar{\tau}(\varphi)| \\ \bar{\tau}(\mathbb{1}\varphi) &= |ZOO_\varphi| \\ \bar{\tau}(\sharp_g\varphi) &= |\bar{\tau}(\varphi)| \end{array} \quad 4$$

771 Finally, we define the $\mathcal{ALCSCC}^{++}$ concept description $C_{\varphi_0} = \bar{\tau}(\varphi_0) \sqcap \text{sat}(\chi_0)$. 5

772 **Claim 20.** *The concept description C_{φ_0} is $\mathcal{ALCSCC}^{++}$ -satisfiable iff the formula φ_0 is $K^{\sharp, \sharp_g}$ - 6
 773 satisfiable. Moreover, the concept description C_{φ_0} has size polynomial in the size of φ_0 .*

774 *Proof.* From right to left, suppose that φ_0 is $K^{\sharp, \sharp_g}$ -satisfiable. It means that there is a pointed 7
 775 graph (G, u) where $G = (V, E)$ and $u \in V$, such that $(G, u) \models \varphi_0$. Let $I_0 = (\Delta^{I_0}, \cdot^{I_0})$ be
 776 the $\mathcal{ALCSCC}^{++}$ interpretation over N_C and N_R , such that $N_C = \{A_p \mid p \text{ a proposition in } \varphi_0\}$,
 777 $N_R = \{R\} \cup \{ZOO_\varphi \mid \mathbb{1}\varphi \in \varphi_0\}$, $\Delta^{I_0} = V$, $A_p^{I_0} = \{v \mid v \in V, (G, v) \models p\}$ for every p in φ_0 ,
 778 $R^{I_0} = E$, $ZOO_\varphi^{I_0} = \{(v, v) \mid v \in V, (G, v) \models \varphi\}$ for every $\mathbb{1}\varphi$ in φ_0 . We can show that $u \in C_{\varphi_0}^{I_0}$.
 779 Basically I^0 is like G with the addition of adequately looping ZOO -roles. An individual in Δ^{I_0} has
 780 exactly one ZOO_φ -successor (itself), exactly when φ is true, and no successor otherwise; A_p is true
 781 exactly where p is true, and the role R corresponds exactly to E .

782 From left to right, suppose that C_{φ_0} is $\mathcal{ALCSCC}^{++}$ -satisfiable. It means that there is an $\mathcal{ALCSCC}^{++}$ 8
 783 finite interpretation $I_0 = (\Delta^{I_0}, \cdot^{I_0})$ and an individual $d \in \Delta^{I_0}$ such that $d \in C_{\varphi_0}^{I_0}$. Let $G = (V, E)$ be
 784 a graph such that $V = \Delta^{I_0}$, $E = R^{I_0}$, and $\ell(d)(p) = 1$ iff $d \in A_p^{I_0}$. We can show that $(G, d) \models \varphi_0$.

785 Since there are at most $|\varphi_0|$ subformulas in φ_0 , the representation of ZOO_φ for every subformula φ 9
 786 of φ_0 can be done in size $\log_2(|\varphi_0|)$. For every formula φ , the size of the concept description $\bar{\tau}(\varphi)$ is
 787 polynomial (at most $O(n \log(n))$). The overall size of $\bar{\tau}(\varphi_0)$ is polynomial in the size of φ_0 , and so
 788 is the size of $\text{sat}(\xi_0)$ (at most $O(n^2(\log(n))^2)$). $\square$

789 The NEXPTIME-membership follows from Claim 20 and the fact that the concept satisfiability 10
 790 problem in $\mathcal{ALCSCC}^{++}$ is in NEXPTIME (Theorem 25).

791 For the hardness, we reduce the problem of consistency of $\mathcal{ALCQ}$ - T_C Boxes which is NEXPTIME- 11
 792 hard [36, Corollary 3.9]. See Section I and Theorem 27 that slightly adapts Tobies' proof to show
 793 that the problem is hard even with only one role.

794 We define the translation τ from the set of $\mathcal{ALCQ}$ concept expressions and $\mathcal{ALCQ}$ cardinality 12
 795 constraints, with only one role R .

$$\begin{array}{ll} \tau(A) &= p_A \\ \tau(\neg C) &= \neg\tau(C) \\ \tau(C_1 \sqcup C_2) &= \tau(C_1) \vee \tau(C_2) \\ \tau(\geq n R.C) &= \sharp\tau(C) + (-1) \times n \geq 0 \\ \tau(\geq n C) &= \sharp_g\tau(C) + (-1) \times n \geq 0 \\ \tau(\leq n C) &= (-1) \times \sharp_g\tau(C) + n \geq 0 \end{array} \quad 13$$

796 It is routine to check the following claim. 1

797 **Claim 21.** Let TC be an $\mathcal{ALCQ-T}_C$ Box. TC is consistent iff $\bigwedge_{\chi \in TC} \mathcal{I}(\chi)$ is $K^{\sharp, \sharp_g}$ -satisfiable. 2

798 Moreover, the reduction is linear. Hardness thus follows from the NEXPTIME-hardness of consis- 3
799 tency of $\mathcal{ALCQ-T}_C$ Boxes. $\square$

800 Lemma 18 and Theorem 19 yield the following corollary. 4

801 **Corollary 22.** $K^{\sharp, \sharp_g}$ -satisfiability problem is NEXPTIME-complete. 5

802 Furthermore, from Theorem 16 and Corollary 22, we obtain the complexity of reasoning with 6
803 ACR-GNNs with truncated ReLU and integer weights.

804 **Corollary 23.** Satisfiability of ACR-GNN with global readout and truncated ReLU is NEXPTIME- 7
805 complete.

806 The decidability of the problem is left open in [7] and in the recent long version [8] when the weights 8
807 are rational numbers. The theorem answers it positively in the case of integer weights and pinpoints
808 the computational complexity.

809 F Formal proofs 9

810 *Proof of Theorem 16.* Let φ be a $K^{\sharp, \sharp_g}$ formula over the set of atomic propositions $p_1, \dots, p_m$. Let 10
811 $\varphi_1, \dots, \varphi_n$ denote an enumeration of the subformulas of φ such that $\varphi_i = p_i$ for $i \leq m$, $\varphi_n = \varphi$,
812 and whenever φ_i is a subformula of φ_j , it holds that $i \leq j$. Without loss of generality, we assume
813 that all subformulas of the form $\xi \geq 0$ are written as

$$\sum_{j \in J} k_j \cdot \mathbb{1}_{\varphi_j} + \sum_{j' \in J'} k_{j'} \cdot \sharp \varphi_{j'} + \sum_{j'' \in J''} k_{j''} \cdot \sharp_g \varphi_{j''} - c \geq 0, \quad 11$$

814 for some index sets $J, J', J'' \subseteq \{1, \dots, n\}$. 12

815 We construct the GNN $\mathcal{A}_\varphi$ in a layered manner. Note that $\mathcal{A}_\varphi$ is fully specified by defining the 13
816 combination function $comb_i$, including its local and global aggregation, for each layer l_i with $i \in$
817 $\{1, \dots, n\}$ and the final classification function cls . Each $comb_i$ produces output vectors of dimension
818 n . The first layer $comb_1$ has input dimension $2m$ and is defined by $comb_1(x, y, z) = (x, 0, \dots, 0)$,
819 ensuring that the first m dimensions correspond to the truth values of the atomic propositions
820 $p_1, \dots, p_m$, while the remaining entries are initialized to zero. Note that $comb_1$ is easily realized by
821 an FNN with ReLU activations. For $i > 1$, the combination function $comb_i$ is defined as

$$comb_i(x, y, z) = \vec{\sigma}(xC + yA_1 + zA_2 + b), \quad 14$$

822 where C, A_1, A_2 are $n \times n$ matrices corresponding to self, local (neighbor), and global aggregation 15
823 respectively, and $b \in \mathbb{R}^n$ is a bias vector. The parameters are defined sparsely as follows:

824 • $C_{ii} = 1$ for all $i \leq m$ (preserving the atomic propositions), 16

825 • If $\varphi_i = \neg \varphi_j$, then $C_{ji} = -1$ and $b_i = 1$,

826 • If $\varphi_i = \varphi_j \vee \varphi_l$, then $C_{ji} = C_{li} = 1$, and 17

827 • If $\varphi_i = \sum_{j \in J} k_j \cdot \mathbb{1}_{\varphi_j} + \sum_{j' \in J'} k_{j'} \cdot \sharp \varphi_{j'} + \sum_{j'' \in J''} k_{j''} \cdot \sharp_g \varphi_{j''} - c \geq 0$, then 18

$$C_{ji} = k_j, \quad A_{1,j'i} = k_{j'}, \quad A_{2,j''i} = k_{j''}, \quad b_i = -c + 1.$$

828 Note that each $comb_i$ has the same functional form, differing only in the non-zero entries of its 19
829 parameters. The classification function is defined by $cls(x) = x_n \geq 1$.

830 Let l_i denote the i th layer of $\mathcal{A}_\varphi$, and fix a vertex v in some input graph. We show, by induction 20
831 on i , that the following invariant holds: for all $j \leq i$, $(x_i(v))_j = 1$ if and only if $v \models \varphi_j$, and
832 $(x_i(v))_j = 0$ otherwise. Assume that $i = 1$. By construction, $x_1(v)$ contains the truth values of
833 the atomic propositions $p_1, \dots, p_m$ in its first m coordinates. Thus, the statement holds at layer 1.

834 Next, assume the statement holds for layer x_{i-1} . Let $j < i$. By assumption, the semantics of φ_j are 1
 835 already correctly encoded in x_{j-1} and preserved by comb_i due to the fixed structure of C , A_1 , A_2 ,
 836 and b . Now consider $j = i$. The semantics of all subformulas of φ_i are captured in x_{i-1} , either at
 837 the current vertex or its neighbors. By the design of comb_i , which depends only on the values of
 838 relevant subformulas, we conclude that φ_i is correctly evaluated. This holds regardless of whether φ_i
 839 is a negation, disjunction, or numeric threshold formula. Thus, the statement holds for all i , and in
 840 particular for $x_n(v)$ and $\varphi_n = \varphi$. Finally, the classifier cls evaluates whether $x_n(v)_n \geq 1$, which is
 841 equivalent to $G, v \models \varphi$. The size claim is obvious given that n depends polynomial on the size of φ .
 842 We note that this assumes that the enumeration of subformulas of φ does not contain duplicates. $\square$

843 *Proof of Theorem 17.* Let $\mathcal{A}$ be a GNN composed of layers $l_1, \dots, l_k$, where each comb_i has input 2
 844 dimension $2m_i$, output dimension n_i , and parameters C_i , $A_{i,1}$, $A_{i,2}$, and b_i . The final classification
 845 is defined via a linear threshold function $\text{cls}(x) = a_1x_1 + \dots + a_{n_k}x_{n_k} \geq 1$. We assume that the
 846 dimensionalities match across layers, i.e. $m_i = n_{i-1}$ for all $i \geq 2$, so that the GNN is well-formed.

847 We construct a formula $\varphi_{\mathcal{A}}$ over the input propositions $p_1, \dots, p_{m_1}$ inductively, mirroring the structure 3
 848 of the GNN computation.

849 We begin with the first layer l_1 . For each $j \in \{1, \dots, n_1\}$, we define: 4

$$\varphi_{1,j} = \sum_{k=1}^{m_1} (C_1)_{kj} \cdot \mathbb{1}p_k + (A_{1,1})_{kj} \cdot \sharp p_k + (A_{1,2})_{kj} \cdot \sharp_g p_k + (b_1)_j \geq 1. \quad 5$$

850 Now suppose that we have already constructed formulas $\varphi_{i-1,1}, \dots, \varphi_{i-1,n_{i-1}}$ for some layer $i \geq 2$. 6
 851 Then, for each output index $j \in \{1, \dots, n_i\}$, we define:

$$\varphi_{i,j} = \sum_{k=1}^{m_i} (C_i)_{kj} \cdot \mathbb{1}\varphi_{i-1,k} + (A_{i,1})_{kj} \cdot \sharp\varphi_{i-1,k} + (A_{i,2})_{kj} \cdot \sharp_g\varphi_{i-1,k} + (b_i)_j \geq 1. \quad 7$$

852 Once all layers have been encoded in this way, we define the final classification formula as 8

$$\varphi_{\mathcal{A}} = a_1 \mathbb{1}\varphi_{k,1} + \dots + a_{n_k} \mathbb{1}\varphi_{k,n_k} \geq 1. \quad 9$$

853 Let G, v be a pointed graph. The correctness of our translation follows directly from the following 10
 854 observations: all weights and biases in $\mathcal{A}$ are integers, and the input vectors $x_0(u)$ assigned to
 855 nodes u in G are Boolean. Moreover, each layer applies a linear transformation followed by a
 856 pointwise truncated ReLU, which preserves the Boolean nature of the node features. It follows that
 857 the intermediate representations $x_i(v)$ remain in $\{0, 1\}^{n_i}$ for all i . Consequently, each such feature
 858 vector can be expressed via a set of Boolean $K^{\sharp, \sharp_g}$ -formulas as constructed above. Taken together,
 859 this ensures that the overall formula $\varphi_{\mathcal{A}}$ faithfully simulates the GNN's computation.

860 It remains to argue that this construction can be carried out efficiently. Throughout, we represent 11
 861 the (sub)formulas using a shared DAG structure, avoiding duplication of equivalent subterms. This
 862 ensures that subformulas $\varphi_{i-1,k}$ can be reused without recomputation. For each layer, constructing all
 863 $\varphi_{i,j}$ requires at most $n_i \cdot m_i$ steps, plus the same order of additional operations to account for global
 864 aggregation terms. Since the number of layers, dimensions, and parameters are bounded by $|\mathcal{A}|$, and
 865 each operation can be performed in constant or linear time, the total construction is polynomial in the
 866 size of $\mathcal{A}$. $\square$

867 G Experimental data and further analyses 12

868 This study investigates the application of dynamic Post-Training Quantization (PTQ) to Aggregate-13
 869 Combined Readout Graph Neural Networks (ACR-GNNs). Implemented in PyTorch [1, 26], dynamic
 870 PTQ transforms a pre-trained floating-point model into a quantized version without requiring retrain-
 871 ing. In this approach, model weights are statically quantized to INT8, while activations remain in
 872 floating-point format until they are dynamically quantized at compute time. This hybrid representation
 873 enables efficient low-precision computation using INT8-based matrix operations, thereby reducing
 874 memory footprint and improving inference speed. PyTorch's implementation applies per-tensor
 875 quantization to weights and stores activations as floating-point values between operations to balance
 876 precision and performance.

We adopt INT8 and QINT8 representations as the primary quantization format. According to theory, INT8 refers to 8-bit signed integers that can encode values in the range $[-128, 127]$. In contrast, QINT8, as defined in the PyTorch documentation [1, 27, 28], is a quantized tensor format that wraps INT8 values together with quantization metadata: a scale (defining the float value represented by one integer step) and a zero-point (the INT8 value corresponding to a floating-point zero). This additional information allows QINT8 tensors to approximate floating-point representations efficiently while enabling high-throughput inference.

To evaluate the practical impact of quantization, we conducted experiments on both synthetic and real datasets. The synthetic data setup was based on the benchmark introduced by [4]. Graphs were generated using the dense Erdős–Rényi model, a classical method for constructing random graphs, and each graph was initialized with five node colours encoded as one-hot feature vectors. The dataset is structured as follows, as shown in Table 3. The training set consists of 5000 graphs, each with 40 to 50 nodes and between 560 and 700 edges. The test set is divided into two subsets. The first subset comprises 500 graphs with the same structure as the training set, featuring 40 to 50 nodes and 560 to 700 edges. The second subset contains 500 larger graphs, with 51 to 69 nodes and between 714 and 960 edges. This design allows us to evaluate the model’s generalization capability to unseen graph sizes.

Table 3: Dataset statistics summary.

Classifier	Dataset	Node			Edge		
		Min	Max	Avg	Min	Max	Avg
p_1	Train	40	50	45	560	700	630
	Test1	40	50	45	560	700	633
	Test2	51	60	55	714	960	832
p_2	Train	40	50	45	560	700	630
	Test1	40	50	44	560	700	628
	Test2	51	60	55	714	960	832
p_2	Train	40	50	44	560	700	629
	Test1	40	50	45	560	700	630
	Test2	51	60	55	714	960	831

For this experiment, we used simple ACR-GNN models with the following specifications. We applied the *sum* function for both the aggregation and readout operations. The combination function was defined as: $comb(x, y, z) = \vec{\sigma}(xC + yA + zR + b)$, where $\vec{\sigma}$ denotes the activation function. Following the original work, we set the hidden dimension to 64, used a batch size of 128, and trained the model for 20 epochs using the Adam optimizer with default PyTorch parameters. We used two activation functions for the experimental part, ReLU and truncated ReLU. For implementation, we used PyTorch [1]: `nn.ReLU` and `nn.Hardtanh(0, 1)` in accordance.

We trained ACR-GNN on complex formulas FOC_2 for labeling. They are presented as a classifier $\alpha_i(x)$ that constructed as:

$$\alpha_0(x) := \text{Blue}(x), \alpha_{i+1}(x) := \exists^{[N,M]} y (\alpha_i(y) \wedge \neg E(x, y))$$

where $\exists^{[N,M]}$ stands for “there exist between N and M nodes”. satisfying a given property.

Observe that each $\alpha_i(x)$ is in FOC_2 , as $\exists^{[N,M]}$ can be expressed by combining $\exists^{\geq N}$ and $\neg \exists^{\geq M+1}$.

The data set has the following specifications: Erdős–Rényigraphs and is labeled according to $\alpha_1(x)$, $\alpha_2(x)$, and $\alpha_3(x)$:

- $\alpha_0(x) := \text{Blue}(x)$
- $p_1 : \alpha_1(x) := \exists^{[8,10]} y (\alpha_0(y) \wedge \neg E(x, y))$
- $p_2 : \alpha_2(x) := \exists^{[10,30]} y (\alpha_1(y) \wedge \neg E(x, y))$
- $p_3 : \alpha_3(x) := \exists^{[10,30]} y (\alpha_2(y) \wedge \neg E(x, y))$

In this section, we present experiments for two activation functions: ReLU and truncated ReLU (implemented via `nn.Hardtanh(0, 1)`) to study the influence of the activation function on the model. Experiments for the ACR-GNN were conducted with different numbers of hidden layers, ranging from 1 to 10. To measure the precision of the results, we use the strategy as [4]: accuracy is calculated as the total number of correctly classified nodes among all nodes in all graphs in the dataset.

Table 4: Accuracy of the ACR-GNN with ReLU according to the number of layers.

Layer	p_1			p_2			p_3		
	Train	Test 1	Test 2	Train	Test 1	Test 2	Train	Test 1	Test 2
1	96.9%	96.4%	74.8%	69.8%	71.0%	56.7%	69.1%	68.8%	75.4%
2	100.0%	100.0%	99.5%	83.7%	84.5%	75.3%	76.6%	76.8%	77.0%
3	97.6%	97.3%	87.2%	83.6%	84.2%	75.1%	76.7%	76.4%	66.9%
4	68.6%	68.4%	67.3%	83.5%	84.0%	76.1%	77.7%	76.3%	46.6%
5	68.5%	68.3%	67.0%	83.5%	83.9%	77.6%	78.2%	76.8%	34.1%
6	68.5%	68.4%	66.1%	83.6%	84.1%	79.6%	77.6%	75.8%	34.8%
7	68.5%	68.5%	67.3%	83.5%	83.8%	80.5%	77.1%	77.7%	49.4%
8	68.5%	68.4%	65.8%	83.4%	83.8%	73.2%	76.7%	75.7%	75.1%
9	68.5%	68.3%	66.7%	83.0%	83.4%	79.1%	77.3%	76.9%	48.0%
10	68.6%	68.3%	65.5%	83.1%	83.7%	77.3%	76.4%	75.6%	37.4%

Table 4 presents the accuracy of the ACR-GNN model with ReLU activation across three FO-properties (p_1 , p_2 , and p_3), evaluated on Train, Test1, and Test2 splits. For p_1 , the model achieves high accuracy in the first three layers, peaking at 99.5% on Test2 at layer 2. From layer 4 and beyond, the accuracy on Test2 declines and stabilizes around 66–67%, suggesting a decreased performance in deeper models for this property. For p_2 , initial accuracy is modest (e.g., 69.8% on Train and 56.7% on Test2 at layer 1), but improves rapidly with depth, surpassing 83% from layer 2 onward on Train and Test1. In particular, the accuracy of Test2 continues to improve with depth, reaching a peak at 80.5% in layer 7, indicating that p_2 benefits from deeper architectures. In contrast, p_3 exhibits less consistent behavior. Accuracy improves early, reaching 77.0% on Test2 at layer 2, but then drops sharply: Test2 accuracy drops to 46.6% at layer 4 and reaches a minimum of 34.1% at layer 5. Some recovery is observed at layers 7 and 8, yet performance remains unstable, with Test2 accuracy at 37.4% by layer 10. Overall, the results demonstrate that model depth significantly affects performance depending on the target property. While p_2 benefits from deeper configurations, both p_1 and p_3 achieve higher generalization performance in shallower networks, with deeper layers leading to overfitting or reduced representation quality on unseen data.

Table 5: Accuracy of the ACR-GNN with ReLU after dynamic PTQ according to the number of layers.

Layer	p_1			p_2			p_3		
	Train	Test 1	Test 2	Train	Test 1	Test 2	Train	Test 1	Test 2
1	96.5%	95.7%	75.3%	69.7%	70.8%	65.6%	68.8%	68.2%	74.7%
2	100.0%	100.0%	99.4%	83.8%	84.4%	75.5%	76.4%	76.6%	77.0%
3	97.6%	97.4%	86.7%	83.5%	84.1%	74.7%	76.7%	76.7%	66.5%
4	68.6%	68.5%	66.9%	83.3%	84.2%	76.2%	77.6%	76.1%	44.6%
5	68.5%	68.2%	67.2%	83.4%	84.0%	77.8%	78.3%	76.6%	33.4%
6	68.6%	68.4%	66.2%	83.5%	83.9%	80.3%	77.4%	75.6%	35.8%
7	68.5%	68.4%	67.1%	83.3%	83.6%	80.6%	77.1%	77.6%	48.7%
8	68.5%	68.3%	65.8%	83.3%	83.7%	73.2%	76.7%	75.5%	74.6%
9	68.5%	68.3%	66.6%	83.0%	83.6%	78.9%	77.1%	76.2%	44.3%
10	68.5%	68.2%	58.1%	83.0%	83.7%	77.5%	76.3%	75.4%	36.6%

Table 5 presents the node-level accuracy of the ACR-GNN model with ReLU activation after applying dynamic post-training quantization (PTQ). Results are reported for three FO-properties (p_1 , p_2 ,

and p_3), evaluated across the Train, Test1, and Test2 splits. For p_1 , the quantized model achieves near-perfect accuracy at layer 2 (Train: 100.0%, Test1: 100.0%, Test2: 99.4%), indicating optimal performance at this depth. Beyond layer 3, accuracy gradually degrades, with Test2 accuracy falling to 58.1% by layer 10. This suggests that deeper networks may amplify quantization-related degradation, especially in generalization. For p_2 , the quantized model demonstrates stable and robust accuracy across most depths. Starting from moderate performance in layer 1 (Train: 69.7%, Test2: 65.6%), accuracy increases quickly and exceeds 83.0% from layer 2 onward in Train and Test1 splits. In particular, the accuracy of Test2 continues to improve up to layer 7 (80.6%), showing resilience to quantization effects even in deeper architectures. In contrast, p_3 exhibits more irregular behavior. Accuracy improves slightly in the early layers (Test2 peaks at 77.0% at layer 2), but then drops substantially, reaching a low of 33.4% at layer 5. Despite stable Train and Test1 accuracy (76–78%), the significant reduction in Test2 suggests overfitting and reduced generalization performance in deeper networks due to quantization. Dynamic PTQ preserves performance well for p_2 in depths, but negatively impacts p_1 and especially p_3 in deeper configurations. This underscores the need for depth-sensitive or property-sensitive quantization strategies when deploying GNNs under resource constraints.

Table 6: Difference in the percentages of the accuracy of ACR-GNN with ReLU before and after dynamic PTQ, rounded to two decimal places.

Layer	p_1			p_2			p_3		
	Train	Test 1	Test 2	Train	Test 1	Test 2	Train	Test 1	Test 2
1	-0.45%	-0.76%	0.52%	-0.13%	-0.18%	8.89%	-0.30%	-0.65%	-0.69%
2	0.00%	0.00%	-0.04%	0.08%	-0.13%	0.14%	-0.18%	-0.23%	0.02%
3	-0.04%	0.06%	-0.49%	-0.16%	-0.14%	-0.34%	-0.02%	0.28%	-0.35%
4	0.01%	0.02%	-0.40%	-0.19%	0.19%	0.06%	-0.05%	-0.20%	-1.99%
5	-0.06%	-0.13%	0.19%	-0.11%	0.06%	0.26%	0.03%	-0.22%	-0.73%
6	0.02%	0.01%	0.06%	-0.03%	-0.18%	0.70%	-0.23%	-0.25%	0.95%
7	0.00%	-0.11%	-0.16%	-0.19%	-0.26%	0.12%	-0.00%	-0.17%	-0.75%
8	-0.03%	-0.09%	-0.01%	-0.12%	-0.12%	-0.02%	-0.05%	-0.28%	-0.49%
9	-0.03%	-0.01%	-0.04%	0.01%	0.21%	-0.13%	-0.26%	-0.72%	-3.74%
10	-0.00%	-0.10%	-7.38%	-0.14%	0.05%	0.20%	-0.08%	-0.14%	-0.78%

Table 6 reports the accuracy differences in percentage points between the original ACR-GNN model with ReLU activation and its dynamically quantized counterpart, using Post-Training quantization (PTQ). The results cover three FO properties (p_1 , p_2 , p_3), three dataset splits (Train, Test1, Test2). Positive values indicate better accuracy after quantization, while negative values indicate degradation. For p_1 , quantization generally causes negligible or negative changes in accuracy. For example, at layer 2, the differences are minimal (Train: 0.00%, Test1: 0.00%, Test2: -0.04%), showing near-identical behavior between the models. However, deeper networks experience more substantial performance drops, especially at layer 10 in Test2 (-7.38%), indicating increased instability due to depth quantization. These patterns highlight a general sensitivity to depth, particularly when generalizing to larger test graphs. In contrast, p_2 exhibits greater resilience to quantization, with occasional performance gains. A notable improvement appears in layer 1 on Test2 (+8.89%), along with smaller gains in layers 5 (+0.26%), 6 (+0.70%) and 10 (+0.20%). However, inconsistencies are still present, for example, a Test2 drop at layer 3 (-0.34%) – which implies that while p_2 benefits more than p_1 , gains are not uniform across the board. p_3 , on the other hand, exhibits the most erratic behavior and is generally more susceptible to quantization. Although a modest gain appears in layer 6 in Test2 (+0.95%), severe degradation is observed in layer 4 (-1.99%) and layer 9 (-3.74%). Across layers and divisions, accuracy losses dominate, suggesting that p_3 is particularly sensitive to quantization, especially in deeper models. In summary, dynamic PTQ results in non-uniform effects across properties, dataset splits, and depths. Although p_2 shows the most consistent tolerance and even improvement in certain cases, p_1 and p_3 are more susceptible to degradation, especially in the Test2 split in deeper configurations. These results emphasize the importance of property-specific and depth-aware quantization strategies to maintain performance in FO-property learning with GNN.

Table 7 presents the accuracy of the ACR-GNN model with truncated ReLU activation on three FO properties (p_1 , p_2 , and p_3), evaluated on the Train, Test1, and Test2 datasets as the number

Table 7: Accuracy of the ACR-GNN with truncated ReLU according to the number of layers. 1

Layer	p_1			p_2			p_3		
	Train	Test 1	Test 2	Train	Test 1	Test 2	Train	Test 1	Test 2
1	98.7%	98.4%	87.0%	77.2%	78.3%	51.1%	69.9%	69.8%	71.5%
2	100.0%	100.0%	98.3%	69.8%	70.0%	63.7%	75.2%	76.5%	75.3%
3	63.1%	61.7%	57.9%	67.8%	67.6%	62.9%	66.3%	65.7%	70.6%
4	58.4%	58.0%	48.6%	66.4%	66.3%	61.3%	61.2%	59.2%	50.3%
5	55.7%	54.3%	50.4%	63.0%	64.3%	39.6%	64.4%	65.1%	66.5%
6	55.5%	54.6%	50.1%	63.0%	64.3%	39.5%	58.2%	57.3%	34.6%
7	53.8%	54.2%	51.4%	63.4%	64.9%	41.7%	57.1%	56.0%	23.3%
8	52.7%	53.6%	50.8%	63.1%	64.0%	40.0%	61.4%	61.5%	55.3%
9	52.5%	52.5%	51.1%	65.0%	65.0%	49.2%	57.2%	56.0%	24.7%
10	54.7%	54.8%	51.1%	63.0%	64.3%	39.6%	57.2%	55.6%	23.4%

of GNN layers increases from 1 to 10. For p_1 , the model exhibits strong performance in shallow configurations, peaking at layer 2 with 100.0% (Train), 100.0% (Test1), and 98.3% (Test2) accuracy. However, performance deteriorates significantly beyond this point: by layer 3, Test2 accuracy drops to 57.9%, and continues to decline in deeper layers, stabilizing around 51.1% by layer 10. This trend suggests overfitting, as training accuracy remains high while generalization performance on Test2 degrades with depth. The accuracy profile of p_2 is more stable. While initial performance is moderate (Test2: 51.1% at layer 1), the model maintains consistent accuracy from layer 3 onward, with minor fluctuations. The narrower gap between training and testing accuracy indicates that p_2 is less sensitive to overfitting and more robust to increasing depth. For p_3 , the model initially performs well, reaching 75.3% on Test2 at layer 2. However, deeper architectures result in a steep decline in generalization performance: Test2 accuracy falls to 50.3% at layer 4, 34.6% at layer 6, and just 23.3% by layer 7. Despite relatively stable scores on Train and Test1, the Test2 drop—evidenced by a gap of over 38 percentage points at layer 7—reflects significant overfitting. In summary, ACR-GNN model with truncated ReLU benefits most from shallow architectures for p_1 and p_3 , whereas p_2 exhibits more resilient behavior across network depths. These results highlight the need for depth-aware design when targeting different FO properties under quantization constraints.

Table 8: Accuracy of the ACR-GNN with truncated ReLU after dynamic PTQ according to the number of layers. 4

Layer	p_1			p_2			p_3		
	Train	Test 1	Test 2	Train	Test 1	Test 2	Train	Test 1	Test 2
1	98.8%	98.8%	86.4%	76.2%	77.8%	59.5%	69.4%	69.3%	74.8%
2	100.0%	100.0%	94.4%	69.6%	69.7%	42.4%	74.8%	76.3%	59.6%
3	61.5%	59.1%	54.9%	67.8%	68.0%	63.6%	66.1%	65.3%	70.7%
4	58.3%	57.7%	47.9%	66.2%	66.7%	43.1%	61.0%	57.5%	46.0%
5	55.4%	54.0%	50.5%	63.0%	64.3%	39.6%	63.9%	57.4%	65.5%
6	55.5%	55.8%	50.0%	63.0%	64.3%	39.8%	57.5%	56.8%	32.5%
7	53.4%	53.1%	50.9%	62.4%	62.5%	44.8%	56.8%	56.2%	24.5%
8	52.5%	53.6%	51.0%	61.4%	63.0%	40.0%	61.4%	62.7%	50.0%
9	52.6%	52.4%	51.2%	65.0%	65.7%	53.7%	57.2%	55.6%	23.7%
10	54.8%	53.9%	51.3%	63.1%	64.3%	39.6%	56.9%	55.1%	23.6%

Table 8 reports the accuracy of the ACR-GNN model after applying dynamic PTQ across three logical query patterns (p_1 , p_2 , p_3) and a range of GNN layers (l from 1 to 10). A general observation is that dynamic PTQ causes more pronounced performance degradation as the number of layers increases, particularly for p_1 and p_3 . While accuracy remains high for shallow configurations, especially at $l = 1$ and $l = 2$ (e.g., p_1 reaches 98.8% on Test1 at $l = 1$ and 100.0% on Train and Test1 at $l = 2$)—a sharp decline follows beyond $l = 2$. For instance, p_1 training accuracy drops from 100.0% at $l = 2$

to 61.5% at $l = 3$, with continued degradation in deeper layers. In contrast, p_2 starts with slightly lower accuracy but exhibits relatively stable behavior across layers. Its accuracy remains in the 60–78% range across all datasets, showing less sensitivity to depth. However, a gradual decline in the precision of Test2 is noticeable, ranging from 59.5% at $l = 1$ to 39.6% at $l = 10$, suggesting that generalization to more complex test graphs is still affected by quantization. The pattern p_3 is the most affected. Although some recovery is observed at intermediate layers (e.g., 70.7% Test2 accuracy at $l = 3$), performance deteriorates with increasing depth, reaching only 23.6% on Test2 at $l = 10$. In summary, dynamic PTQ enables significant model compression for ACR-GNNs, but at the cost of accuracy, particularly in deeper architectures and complex FO-query patterns such as p_1 and p_3 . Shallow configurations (e.g., $l \leq 2$) maintain good performance after quantization, indicating that careful depth-aware quantization strategies are essential for preserving generalization.

Table 9: Difference in the percentages of the accuracy of ACR-GNN with truncated ReLU before and after dynamic PTQ.

Layer	p_1			p_2			p_3		
	Train	Test 1	Test 2	Train	Test 1	Test 2	Train	Test 1	Test 2
1	0.1%	0.3%	-0.6%	-1.0%	-0.5%	8.4%	-0.5%	-0.5%	3.4%
2	0.0%	0.0%	-3.9%	-0.2%	-0.3%	-21.3%	-0.5%	-0.2%	-15.7%
3	-1.6%	-2.7%	-3.0%	0.0%	0.4%	0.7%	-0.2%	-0.4%	0.1%
4	-0.2%	-0.3%	-0.8%	-0.2%	0.5%	-18.2%	-0.2%	-1.7%	-4.3%
5	-0.3%	-0.3%	0.2%	0.0%	0.0%	0.0%	-0.6%	-7.7%	-1.0%
6	-0.0%	1.2%	-0.1%	-0.0%	0.0%	0.3%	-0.6%	-0.5%	-2.2%
7	-0.4%	-1.2%	-0.5%	-1.0%	-2.3%	3.1%	-0.4%	0.2%	1.2%
8	-0.2%	0.0%	0.2%	-1.7%	-1.0%	-0.0%	0.0%	1.3%	-5.3%
9	0.2%	-0.1%	0.1%	0.0%	0.7%	4.5%	0.1%	-0.5%	-1.0%
10	0.1%	-0.9%	0.3%	0.0%	0.0%	0.0%	-0.3%	-0.5%	0.2%

Table 9 presents the percentage changes in accuracy of the ACR-GNN model with truncated ReLU after applying Dynamic Post-Training quantization (PTQ), across three query patterns (p_1, p_2, p_3) and for different numbers of GNN layers ($l = 1$ to $l = 10$). The difference is calculated as the quantized accuracy minus the original, scaled to a percentage. In the case of this table, we can see changes layer by layer. Here, where $l = 1$, we observe small improvements in accuracy. If we examine this more precisely, for p_1 , the precision improves across all datasets, with the highest gain in Test2 (+11.1%). p_2 shows a mixed pattern with small increases in Train / Test1, but a decrease in Test2 (-6.1%). p_3 remains stable, showing minimal change ($\leq 1.2\%$). When $l = 2$, the results show early degradation, as p_2 suffers significant drops, especially on Test2 (-33.0%), while p_3 sees a drop in Test2 of -17.4%, p_1 remains unchanged on Train / Test1 and slightly lower (-5.0%) on Test2. A major drop occurs when $l = 3$ for p_1 , with -36.1% on Train and -38.3% on Test1. p_2 also shows a negative trend, but Test2 is impacted less than in Layer 2. Interestingly, p_3 has a positive change in Test2 (+4.2%), indicating some robustness in this setting. The continuous trend for layers from 4 to 9. For $l = 10$, p_1 appears to recover slightly in Test2 (-6.8%, compared to -15% previously). However, p_2 and p_3 still show substantial losses (-37.9% and -13.1% respectively), suggesting that deeper architectures struggle consistently after dynamic quantization. In summary, Table 9 highlights the accuracy losses due to dynamic PTQ. This correlates with the literature [15], where the authors noted some loss in accuracy, but the quantized model should provide better results in comparing the size. Although some early layers benefit slightly, deeper layers consistently show reduced accuracy, especially in Test2, the data set with larger, more complex graphs. The pattern confirms that dynamic PTQ, though efficient, can harm generalization, particularly in deeper and more expressive GNN configurations.

After presenting the accuracy results before and after applying dynamic Post-Training Quantization (PTQ), we proceed to analyze the influence of the activation function on the performance of the model. This comparison is provided both graphically and in tabular form. For the graphical representation, we utilized box plots, a statistical tool designed to visualize the distribution of a variable in terms of its quartiles. In these plots, the box itself spans from the first quartile (Q1) to the third quartile (Q3), with the median value (Q2) marked by a line within the box. The whiskers of the box plot extend to the minimum and maximum values that do not qualify as outliers, providing insight into the spread

and concentration of the data. In addition to these visualizations, a detailed table complements the analysis by presenting summary statistics. The table includes the mean, standard deviation, minimum, and maximum values for each configuration. It also presents the three quartiles: Q1, which represents the 25th percentile, Q2, or the median, which is the 50th percentile, and Q3, the 75th percentile. These quartiles divide the data into four equal parts, helping to identify the central tendency and variability. Furthermore, we calculate the interquartile range (IQR), defined as the difference between the third quartile (Q3) and the first quartile (Q1), which serves as a measure of statistical dispersion. Based on the IQR, we also determine the lower and upper bounds using the standard rule, which involves subtracting 1.5 times the IQR from Q1 and adding it to Q3, respectively. These bounds enable the identification of potential outliers and provide a more comprehensive understanding of how the activation function and quantization impact the distribution of model accuracy. All metrics were applied to all datasets: Train, Test1, and Test2. For the visualization part, we used the Python library Plotly.




Figure 3: Detailed summary statistics across configurations for p_1 formula.

Table 10: Detailed summary statistics across configurations for p_1 formula.

Statistic	ReLU	ReLU + PTQ	Truncated ReLU	Truncated ReLU + PTQ
Mean	0.758	0.755	0.628	0.623
Std	0.132	0.134	0.178	0.177
Min	0.655	0.581	0.486	0.479
25% (Q1)	0.683	0.682	0.525	0.524
50% (Median)	0.685	0.685	0.547	0.544
75% (Q3)	0.841	0.839	0.609	0.589
Max	1.000	1.000	1.000	1.000
IQR	0.158	0.157	0.084	0.065
Lower Bound	0.446	0.447	0.399	0.427
Upper Bound	1.078	1.073	0.734	0.686

Table 10 and Figure 3 present summary statistics for the accuracy results obtained from four configurations of the ACR-GNN model: ReLU, ReLU with dynamic Post-Training Quantization (PTQ), Truncated ReLU, and Truncated ReLU with PTQ. The results show that the highest mean accuracy is achieved with the ReLU configuration (0.758), closely followed by ReLU + PTQ (0.755). This indicates that applying dynamic quantization to the ReLU model does not significantly reduce the average accuracy. In contrast, both Truncated ReLU (0.628) and Truncated ReLU + PTQ (0.623) result in noticeably lower mean values, suggesting that this activation function may degrade performance on the p_1 query pattern. The median values align with the mean, further confirming this trend. In terms of variability, the standard deviation is lower for the ReLU-based models (0.13), whereas the truncated ReLU configurations show higher variability (0.18). This pattern is also reflected in the interquartile range (IQR): ReLU configurations exhibit wider IQRs (0.158 and 0.157), while truncated versions have narrower ranges (0.084 and 0.065). Despite the narrower spread, the performance is consistently lower with truncated ReLU. All configurations include samples that achieve a maximum

accuracy of 1.0, indicating that optimal predictions are possible in all cases. However, minimum accuracy drops more sharply in truncated ReLU models (0.486 and 0.479) compared to ReLU (0.655 and 0.581), indicating a higher risk of underperformance. The lower and upper bounds provide insight into potential outliers. The lower bounds are lower in the truncated models, while the upper bounds are higher in ReLU configurations (exceeding 1.0 due to statistical calculation), indicating a wider spread and potentially higher ceiling for performance.




Figure 4: Detailed summary statistics across configurations for p_2 formula.

Table 11: Detailed summary statistics across configurations for p_2 formula.

Statistic	ReLU	ReLU + PTQ	Truncated ReLU	Truncated ReLU + PTQ
Mean	0.7992	0.8020	0.6064	0.5967
Std	0.0615	0.0511	0.1085	0.1122
Min	0.5670	0.6560	0.3950	0.3960
25% (Q1)	0.7738	0.7758	0.6170	0.5515
50% (Median)	0.8340	0.8330	0.6385	0.6305
75% (Q3)	0.8370	0.8368	0.6598	0.6608
Max	0.8450	0.8440	0.7830	0.7780
IQR	0.0632	0.0610	0.0428	0.1093
Lower Bound	0.6789	0.6843	0.5529	0.3876
Upper Bound	0.9319	0.9282	0.7239	0.8246

Table 11 and Figure 4 present a comprehensive overview of the accuracy results in four model configurations: ReLU, ReLU with dynamic post-training quantization (PTQ), Truncated ReLU, and Truncated ReLU with PTQ - for the query formula p_2 . From the mean accuracy values, ReLU and ReLU + PTQ clearly outperform the other configurations, achieving 0.7992 and 0.8020, respectively. This indicates that both setups yield strong overall performance, with dynamic quantization having a slightly positive effect on average accuracy in this case. In contrast, Truncated ReLU (0.6064) and Truncated ReLU + PTQ (0.5967) show substantially lower mean values, highlighting a notable drop in predictive performance when using truncated activation. Looking at the variability, the standard deviation is lower for the ReLU configurations (0.0615 and 0.0511), suggesting a more consistent accuracy. The truncated versions, especially the quantized one (0.1122), are more dispersed, reflecting greater instability. This is further emphasized by the IQR values: 0.0632 and 0.0610 for ReLU and ReLU + PTQ versus 0.0428 for Truncated ReLU and a larger 0.1093 for Truncated ReLU + PTQ. The larger IQR for Truncated ReLU + PTQ implies a larger fluctuation in the middle 50% of the data, despite its lower central values. The median values confirm this trend: both ReLU configurations cluster around 0.833–0.834, while truncated versions fall between 0.6305 and 0.6385. The lower bounds, derived from $Q1 - 1.5 \times IQR$, are also lower in the Truncated ReLU + PTQ case (0.3876), indicating a greater potential for underperformance and a higher risk of poor accuracy. The maximum and minimum values highlight the performance extremes. ReLU configurations reach up to 0.845 and 0.844, significantly higher than the 0.783 and 0.778 of truncated variants. The lower minimum accuracy (0.395–0.396) in truncated settings further reinforces concerns about their reliability.




Figure 5: Detailed summary statistics across configurations for p_3 formula.

Table 12: Detailed summary statistics across configurations for p_3 formula.

Statistic	ReLU	ReLU + PTQ	Truncated ReLU	Truncated ReLU + PTQ
Mean	0.6883	0.6844	0.5821	0.5694
Std	0.1434	0.1466	0.1441	0.1427
Min	0.3410	0.3340	0.2330	0.2360
25% (Q1)	0.6888	0.6835	0.5600	0.5575
50% (Median)	0.7635	0.7615	0.6020	0.5750
75% (Q3)	0.7688	0.7670	0.6645	0.6545
Max	0.7820	0.7830	0.7650	0.7630
IQR	0.0800	0.0835	0.1045	0.0970
Lower Bound	0.5687	0.5582	0.4032	0.4120
Upper Bound	0.8888	0.8922	0.8213	0.8000

Table 12 and Figure 5 provide descriptive statistics for the accuracy of the ACR-GNN model under four configurations—ReLU, ReLU with dynamic Post-Training Quantization (PTQ), Truncated ReLU, and Truncated ReLU with PTQ—for the p_3 query formula. Starting with the mean accuracy, ReLU (0.6883) and ReLU + PTQ (0.6844) again outperform the Truncated ReLU configurations, which register noticeably lower means of 0.5821 and 0.5694, respectively. This indicates that models that use ReLU activations are generally more effective for p_3 . The standard deviation values are relatively similar across all configurations (approximately 0.14), suggesting that while the truncated configurations perform worse on average, they do not fluctuate more widely than the ReLU-based ones. The minimum values further emphasize the performance gap: ReLU models maintain minimum accuracies above 0.33, while truncated variants drop to as low as 0.233. This shows that truncated configurations are more prone to poor performance in the worst-case scenarios. In terms of quartiles, ReLU and ReLU + PTQ have Q1 and Q3 clustered around 0.68–0.77, indicating that the middle 50% of their results are concentrated within a tight and relatively high accuracy range. Truncated ReLU variants have their Q1 around 0.56 and Q3 near 0.65, which not only shows lower performance but also a wider IQR (0.1045 for Truncated ReLU and 0.0970 for Truncated ReLU + PTQ). This reflects more variability across the central portion of the data in the truncated setups. The median accuracy is again higher in ReLU configurations (around 0.76), compared to 0.60 and 0.575 for truncated ones, reinforcing the conclusion that ReLU configurations are more reliable. Examining the bounds, the ReLU models show a lower bound above 0.55 and upper bounds above 0.88, suggesting strong and consistent performance. Truncated models exhibit lower bounds near 0.40 and upper bounds around 0.80, indicating both a lower floor and a lower ceiling in performance.

Across all query patterns (p_1 , p_2 , and p_3), ReLU and ReLU + PTQ consistently demonstrate higher average accuracy and more stable performance, making them the most reliable configurations. In contrast, Truncated ReLU and its quantized variant result in lower accuracy and greater variability, especially in worst-case scenarios. Dynamic PTQ tends to maintain or slightly enhance performance in ReLU models, but its effect on truncated activations is less favorable, often introducing further

1113 inconsistency. Overall, ReLU-based configurations—quantized or not—are better suited for the 1
1114 ACR-GNN model across the evaluated formulas.

1115 Other parameters of interest to us are the time and size of the models. In the event of changes in 2
1116 size, it is easy to compare the data using the bar plots presented in Figure 6. The size changes in
1117 percentages we calculated according to the formula:

$$\text{Difference in percentages} = \frac{\text{Value}_{dPTQ} - \text{Value}_{original}}{\text{Value}_{original}} * 100\% \quad 3$$

1118 In other words, this formula shows how much the dynamic PTQ value deviates from the original 4
1119 value as a percentage of the original value.

1120 In this section, we compare parameters for different activation functions. We observe that the results 5
1121 of size changes in the following models remain unchanged when we modify the training dataset. We
1122 present the results not only graphically but also in a tabular format. In the plots, it is possible to see
1123 the trends and, in the tabular format, the numerical changes.

Table 13: Detailed information about the size of the model. The size values are in megabytes and 6
refer to the file sizes of the GNNs.

Layer	Original Size (MB)	Quantized Size (MB)	Difference (MB)	Reduction (%)
1	0.057	0.023	0.034	59.604%
2	0.112	0.044	0.068	60.993%
3	0.167	0.064	0.103	61.559%
4	0.221	0.085	0.137	61.804%
5	0.276	0.105	0.171	61.975%
6	0.331	0.126	0.206	62.068%
7	0.386	0.146	0.240	62.148%
8	0.441	0.167	0.274	62.194%
9	0.496	0.187	0.309	62.230%
10	0.551	0.208	0.343	62.251%


1124 Table 13 provides a detailed comparison of the model sizes before and after applying dynamic 8
1125 post-training quantization (PTQ). As the number of layers increases, both the original and quantized
1126 model sizes grow; however, the percentage reduction remains remarkably consistent, ranging from
1127 approximately 60.993% at 2 layers to 62.251% at 10 layers. This stable percentage reduction,
1128 approximately 60–62%—indicates that PTQ effectively compresses the model regardless of its depth,
1129 significantly reducing the memory footprint without altering the underlying architecture of the GNN.
1130 Such a reduction is particularly crucial for deployments in resource-constrained environments.

1131 Furthermore, after presenting the tabular data, our graphs (Figure 6) reveal a clear trend: While the 9
1132 absolute sizes of the original and quantized models increase with the number of layers, the relative
1133 reduction achieved through dynamic PTQ remains consistent. The size of the original model increases
1134 approximately linearly from 0.057 MB for $l = 1$ to 0.551 MB at $l = 10$, while the quantized model
1135 grows from 0.023 MB to 0.208 MB, preserving the growth structure, but on a reduced scale. The
1136 absolute size difference increases from 0.034 MB in $l = 1$ to 0.343 MB in $l = 10$, demonstrating that
1137 quantization becomes more beneficial for deeper models. Overall, the consistent percentage reduction
1138 across all tested configurations confirms that PTQ scales effectively, delivering stable compression
1139 rates and making it an attractive option for deeper GNN deployments in real-world edge or mobile
1140 environments.


1141 Moreover, we observed that the query property had no noticeable impact on the model size. This can 10
1142 be clearly seen in the bar plots in Figure 6a, Figure 6c, and Figure 6e.

1143 We also measured the change over time. Specifically, we considered three distinct time metrics: 11
1144 Elapsed time (the time taken during training), Time Original (the time required for inference
1145 on the test datasets using the original trained model), and Time quantized (the inference time on
1146 the test datasets using the quantized model). These results are presented in Figure 7.


1147 The data in Figure 7 reflect the impact of dynamic PTQ on the ACR-GNN model in three query 12
1148 patterns (p_1 , p_2 , and p_3) and for GNN depths ranging from 1 to 10 layers. Across all patterns,




(a) Size changes in MB for the first formula




(b) Size changes in MB for the first formula. Difference present in percentage.




(c) Size changes in MB for the second formula



(d) Size changes in MB for the second formula. Difference present in percentage.



(e) Size changes in MB for the third formula



(f) Size changes in MB for the third formula. Difference present in percentage.


Figure 6: Impact of dynamic Post-Training quantization on model size (MB). Changes of size in percentages

quantized models consistently require more inference time than their original counterparts. This increased time is expected as a result of the real-time quantization of weights and activations during inference. Additionally, both the original and quantized models exhibit a consistent, near-linear increase in inference time with model depth, suggesting that computational complexity grows linearly as layers are added.


Despite this overhead, which ranges between 0.1 and 0.9 s depending on the number of layers, the significant reduction in model size (as demonstrated in Table 13 and the corresponding graphs) makes quantized models especially attractive for resource-constrained environments where minimizing the memory footprint is more critical than achieving the lowest possible latency.

To test the technique not only on synthetic data, we chose the Protein-Protein Interactions (PPI) benchmark. The PPI dataset consists of graph-level mini-batches, with separate splits for Training, Validation, and Testing.


In Table 14, we present a summary of the PPI dataset, which consists of 20 training graphs, 2 validation graphs, and 2 test graphs. Each graph contains nodes with 50-dimensional features and supports multi-label classification with 121 possible labels. On average, each node is associated with



(a) Time changes in seconds for the first formula



(b) Time changes in seconds for the second formula



(c) Time changes in seconds for the third formula

Figure 7: Impact of dynamic Post-Training quantization on Latency (sec)

Table 14: Dataset summary.

Dataset	Num Graphs	Node Feature Dim	Label Dim	Avg Active Labels/Node	Avg Degree
Train	20	50	121	37.20	54.62
Validation	2	50	121	35.64	61.07
Test	2	50	121	36.22	58.64

approximately 36 labels, indicating a densely labelled dataset. The average node degree is also high, ranging from 54.6 in the training set to 61.1 in the validation set, reflecting the dense connectivity of the protein-protein interaction graphs. The dataset presents a complex multi-label classification task with consistently rich structure across all splits.

Table 15: Dataset statistics summary.

Dataset	Node			Edge		
	Min	Max	Avg	Min	Max	Avg
Train	591	3480	2245.30	7708	106754	61318.40
Validation	3230	3284	3257.00	97446	101474	99460.00
Test	2300	3224	2762.00	61328	100648	80988.00

The statistics of the dataset presented in Table 15 contain large graphs with varying sizes between the train, the validation, and the test splits. Training graphs range from 591 to 3,480 nodes, with an average of 2,245 nodes per graph, and between 7,708 and 106,754 edges (average 61,318 edges). Validation graphs are more consistent in size, with 3,230 to 3,284 nodes and 97,446 to 101,474 edges, averaging 3,257 nodes and 99,460 edges. The test graphs have 2,300 to 3,224 nodes, averaging 2,762 nodes, and 61,328 to 100,648 edges, averaging 80,988. These statistics confirm that the dataset

contains large and densely connected graphs and demonstrate a distributional shift in graph size and edge count between training and test data. This information is helpful in evaluating the model’s ability to generalize to unseen and variable graph structures.

One key difference between the synthetic data and the PPI dataset is that the latter involves a multi-label classification task, rather than a binary classification task, because the PPI dataset is a common benchmark where each node (representing proteins) can have multiple labels, such as protein functions or interactions. Also, it is important to mention the key differences between the synthetic data and the real one. Here, the authors used the code function `EarlyStopping`: Utility for stopping training early if no further improvement is observed. The second difference is that the code is structured to run multiple experiments to collect statistics (mean and standard deviation) of the model performance, ensuring that the results are robust across different random initializations. In this case, we performed the experiments 10 times for each model, with a combination layer equal to 1 and a number of layers ranging from 1 to 10. The number of hidden dimensions is equal to 256.

For these experiments, we used two activation functions to compare the results with synthetic data. The presentation of the results follows the same approach as for synthetic data. Moreover, in the case of real data [4] used the F1 Score as an evaluation metric. This metric is commonly used to evaluate classification tasks.

According to the Scikit-learn library [25], the F1 score is defined in the following way. The F1 score can be interpreted as a harmonic mean of precision and recall, where an F1 score reaches its best value at 1 and its worst score at 0. The relative contribution of precision and recall to the F1 score is equal. The formula for the F1 score is as follows:

$$F1 = \frac{2TP}{2TP + FP + FN}$$

where, TP – is the number of true positives, FN – is the number of false negatives, FP – is the number of false positives. F1 is calculated by default as 0.0 when there are no true positives, false negatives, or false positives.

The reference code’s results [5] are structured as follows: a table showing the loss and accuracy for each dataset (train, validation, and test). Here, we present only the accuracy of the model according to the number of layers, as we do for the synthetic data. For better representation, we formed the model’s output in a tabular representation.

Table 16: Accuracy for the original and quantized (dynamic PTQ) models. PPI Benchmark.

(a) Accuracy of the ACR-GNN with ReLU according to the number of layers.

(b) Accuracy of the ACR-GNN with ReLU after dynamic PTQ according to the number of layers.

Layer	Train	Validation	Test
1	54.7%	43.1%	39.5%
2	52.5%	44.6%	45.7%
3	52.3%	42.6%	44.0%
4	52.3%	39.2%	40.6%
5	49.6%	39.7%	39.1%
6	49.3%	43.5%	43.3%
7	51.7%	39.9%	38.5%
8	50.8%	36.3%	35.8%
9	48.0%	43.8%	33.2%
10	47.1%	36.9%	36.8%

Layer	Train	Validation	Test
1	55.0%	50.8%	50.2%
2	52.3%	47.8%	47.2%
3	51.9%	45.7%	42.8%
4	51.9%	37.4%	34.1%
5	48.9%	39.1%	40.8%
6	48.9%	42.9%	43.8%
7	51.4%	43.0%	40.6%
8	50.5%	35.9%	36.8%
9	47.7%	40.8%	40.9%
10	46.5%	36.2%	38.7%

Table 16 reports the precision of the ACR-GNN model with ReLU activation in varying numbers of layers, both in its original form and after applying dynamic post-training quantization (dPTQ). The results are presented for the training, validation, and test sets of the PPI benchmark. For both versions of the model, the performance does not increase consistently with the number of layers. Instead, accuracy typically peaks within the first few layers and tends to degrade or fluctuate as the network’s depth increases. In particular, the highest accuracies for the training, validation, and test

sets are achieved with 1 or 2 layers, indicating that shallower architectures are better suited for this task. Specifically, the original model achieves its best test accuracy (45.7%) at 2 layers, while the quantized model achieves an even higher test accuracy (50.2%) at just 1 layer. Dynamic quantization slightly improves generalization performance in the early layers. At layer 1, the quantized model surpasses the original in both validation (50.8% vs. 43.1%) and test accuracy (50.2% vs. 39.5%), suggesting that quantization can have a regularizing effect in low-depth configurations. However, as the number of layers increases beyond 4, the performance of both models tends to decline, likely due to over-smoothing or optimization difficulties common in deep GNNs.

Table 17: Difference in accuracy of ACR-GNN with ReLU before and after dynamic PTQ. PPI Benchmark.

Layer	Train	Validation	Test
1	0.3%	7.7%	10.7%
2	-0.2%	3.2%	1.5%
3	-0.4%	3.1%	-1.2%
4	-0.4%	-1.8%	-6.5%
5	-0.7%	-0.6%	1.7%
6	-0.4%	-0.6%	0.5%
7	-0.3%	3.1%	2.1%
8	-0.3%	-0.4%	1.0%
9	-0.3%	-3.0%	7.7%
10	-0.6%	-0.7%	1.9%

Table 17 reports the absolute difference in precision between the quantized and original ACR-GNN model with ReLU on the PPI benchmark, between training, validation and test sets for varying numbers of layers. Positive values indicate better performance after quantization, while negative values reflect performance degradation. At layer 1, the quantized model shows the largest gains, with improvements of 7.7% on validation and 10.7% on the test set, suggesting a clear generalization advantage in shallow architectures. Smaller, but consistent improvements are also observed at layers 2 and 7, particularly in the validation and test sets. In contrast, certain layers exhibit minor drops in accuracy. For example, layer 4 shows the largest decrease in the test set (6.5%). Overall, the results indicate that dynamic quantization can lead to modest accuracy improvements, particularly in shallow to mid-depth GNNs, with negligible or slightly negative effects in deeper configurations. This highlights the potential of quantization for lightweight deployment with minimal accuracy trade-offs.

Table 18: Detailed information about the model size before and after quantization. PPI Benchmark. Sizes are in megabytes.

Layer	Original Model (MB)	Quantized Model (MB)	Difference (MB)	Reduction (%)
1	0.922	0.242	0.680	-73.749%
2	1.718	0.451	1.267	-73.765%
3	2.515	0.660	1.855	-73.772%
4	3.311	0.868	2.443	-73.776%
5	4.108	1.077	3.031	-73.778%
6	4.904	1.286	3.618	-73.779%
7	5.701	1.495	4.206	-73.780%
8	6.497	1.704	4.794	-73.780%
9	7.294	1.912	5.382	-73.781%
10	8.090	2.121	5.969	-73.781%

Table 18 presents the memory footprint of the ACR-GNN model at different layer depths, comparing the original model (complete precision) with its dynamically quantized counterpart. The table also includes both absolute and percentage differences in size, highlighting the compression effect introduced by dynamic post-training quantization. Across all layers, the quantized model consistently exhibits a size reduction of approximately 73.78% compared to the original model. For example, at 10

layers, the model size decreases from 8.09MB to 2.12MB, yielding an absolute reduction of 5.97MB. This trend is consistent and proportional across all depths, indicating that the memory savings scale linearly with the model’s complexity (i.e., the number of layers). These results demonstrate the effectiveness of dynamic quantization in significantly reducing model size without the need for retraining.

Table 19: Elapsed times (in seconds) for the original and quantized (dynamic PTQ) models. PPI Benchmark.

(a) Elapsed times for the original model.

Layer	Train	Validation	Test
1	0.913	0.115	0.113
2	1.400	0.158	0.182
3	1.447	0.188	0.172
4	1.982	0.257	0.224
5	2.225	0.295	0.247
6	2.846	0.318	0.236
7	3.420	0.442	0.328
8	3.120	0.437	0.343
9	3.626	0.433	0.390
10	4.011	0.410	0.376

(b) Elapsed times for the quantized model.

Layer	Train	Validation	Test
1	0.921	0.134	0.112
2	1.469	0.178	0.129
3	1.410	0.211	0.173
4	1.694	0.252	0.181
5	2.538	0.322	0.304
6	2.878	0.307	0.313
7	3.538	0.328	0.299
8	3.236	0.360	0.342
9	3.936	0.605	0.481
10	3.783	0.464	0.375

Table 21 reports the inference times of the original and dynamically post-training quantized ACR-GNN models across training, validation, and test datasets, measured at various layer depths. The results reveal that quantization does not significantly reduce inference time in most configurations and, in some cases, results in slightly higher latency. For the training set, the execution time of the quantized model closely follows that of the original, with negligible differences across all layers. In the validation and test sets, while some improvements are observed at shallow depths (e.g., the layer 2 test time reduces from 0.182 to 0.129 s), the overall pattern indicates no consistent speedup from quantization. In fact, certain configurations, such as layers 9 and 10 in the validation set, exhibit increased latency in the quantized version compared to the original.

Table 20: Difference in elapsed time (in seconds) and corresponding percentage difference of ACR-GNN with ReLU before and after dynamic PTQ on the PPI Benchmark.

Layer	Train		Validation		Test	
	Diff (s)	% Diff	Diff (s)	% Diff	Diff (s)	% Diff
1	-0.008	0.915%	-0.019	16.307%	0.001	-1.085%
2	-0.069	4.931%	-0.020	12.308%	0.053	-29.114%
3	0.037	-2.525%	-0.023	12.238%	-0.001	0.309%
4	0.288	-14.531%	0.005	-1.990%	0.043	-19.096%
5	-0.313	14.091%	-0.027	9.291%	-0.057	23.218%
6	-0.032	1.131%	0.011	-3.463%	-0.077	32.455%
7	-0.118	3.465%	0.114	-25.741%	0.029	-8.918%
8	-0.116	3.709%	0.077	-17.556%	0.001	-0.276%
9	-0.310	8.555%	-0.172	39.611%	-0.091	23.218%
10	0.228	-5.678%	-0.054	13.105%	0.001	-0.192%

Table 20 presents the difference in inference time between the original and dynamically quantized (dPTQ) ACR-GNN models, reported in absolute (seconds) and relative (%) terms, across various layer depths. The results show that quantization has an inconsistent effect on inference time, with no clear trend of improvement. In some configurations, dynamic quantization slightly reduces inference time; for example, layer 2 shows a 0.053s reduction on the test set, corresponding to a

29.11% improvement. Similarly, layer 5 achieves an improvement in test time of 23.22%, and layer 6 shows the largest test time speedup of 32.46%. However, in other cases, such as layer 4 in the training set (+0.288s, -14.53%) and layer 10 (+0.228s, -5.68%), quantization increases execution time. The relative differences on the validation set also vary widely, with notable slowdowns at layers 7 (-25.74%) and 9 (-39.61%). These inconsistencies highlight that run-time performance does not always benefit from dynamic quantization, and the effectiveness likely depends on the specific computation pattern and how well the underlying hardware supports quantized operations.

Table 21: Elapsed time (in seconds) for ACR-GNN with and without dynamic post-training quantization (dPTQ). PPI Benchmark

Layer	Train		Validation		Test	
	Original	dPTQ	Original	dPTQ	Original	dPTQ
1	0.780	0.858	0.102	0.112	0.077	0.094
2	0.986	0.966	0.130	0.131	0.109	0.107
3	1.138	1.161	0.157	0.159	0.149	0.140
4	1.371	1.366	0.159	0.204	0.156	0.160
5	1.645	1.682	0.201	0.211	0.173	0.199
6	1.833	1.766	0.242	0.256	0.188	0.205
7	2.166	2.156	0.282	0.261	0.239	0.242
8	2.355	2.534	0.317	0.300	0.241	0.283
9	2.539	2.652	0.337	0.349	0.302	0.292
10	2.842	3.122	0.386	0.461	0.326	0.348

Table 21 reports the elapsed time (in seconds) required to perform inference on the training, validation, and test sets using the ACR-GNN model with ReLU activation, both in its original form and after applying dynamic post-training quantization (dPTQ). The measurements reflect the running time of the trained models only; the time required for model training is not included in these results. The values indicate that inference time generally increases with the number of layers, as expected, and the impact of quantization on runtime varies across depths. In some cases, dPTQ slightly reduces inference time (e.g., Layer 6, Train), while in others it introduces moderate overhead, particularly for deeper models.

The experiments were run on a Samsung Galaxy Book4 laptop with an Intel Core i7-150U processor, 16 GB RAM, and 1 TB SSD storage. Additional experiments were conducted using Kaggle’s cloud platform with an NVIDIA Tesla P100 GPU (16 GB RAM).

H Description logics with global and local cardinality constraints

The Description Logic $\mathcal{ALCSCC}^{++}$ [2] extends the basic Description Logic $\mathcal{ALC}$ [3] with concepts that capture cardinality and set constraints expressed in the quantifier-free fragment of Boolean Algebra with Presburger Arithmetic (QFBAPA) [20].

We assume that we have a set of *set variables* and a set of *integer constants*.

A QFBAPA *formula* is a Boolean combination ($\wedge, \vee, \neg$) of *set constraints* and *cardinality constraints*.

A *set term* is a Boolean combination ($\cup, \cap, \neg$) of *set variables*, and *set constants* $\mathcal{U}$, and $\emptyset$. If S is a set term, then its cardinality $|S|$ is an *arithmetic expressions*. Integer constants are also arithmetic expressions. If T_1 and T_2 are arithmetic expressions, so is $T_1 + T_2$. If T is an arithmetic expression and c is an integer constant, then $c \cdot T$ is an arithmetic expression.

Given two set terms B_1 and B_2 , the expressions $B_1 \subseteq B_2$ and $B_1 = B_2$ are *set constraints*. Given two arithmetic expressions T_1 and T_2 , the expressions $T_1 < T_2$ and $T_1 = T_2$ are *cardinality constraints*. Given an integer constant c and an arithmetic expression T , the expression $c \text{ dvd } T$ is a *cardinality constraint*.

A *substitution* σ assigns $\emptyset$ to the set constant $\emptyset$, a finite set $\sigma(\mathcal{U})$ to the set constant $\mathcal{U}$, and a subset of $\sigma(\mathcal{U})$ to every set variable. A substitution is first extended to set terms by applying the standard

1285 set-theoretic semantics of the Boolean operations. It is further extended to map arithmetic expressions 1
 1286 to integers, in such that way that every integer constant c is mapped to c , for every set term B , the
 1287 arithmetic expression $|B|$ is mapped to the cardinality of the set $\sigma(B)$, and the standard semantics for
 1288 addition and multiplication is applied.

1289 The substitution σ (QFBAPA) satisfies the set constraint $B_1 \subseteq B_2$ if $\sigma(B_1) \subseteq \sigma(B_2)$, the set 2
 1290 constraint $B_1 = B_2$ if $\sigma(B_1) = \sigma(B_2)$, the cardinality constraint $T_1 < T_2$ if $\sigma(T_1) < \sigma(T_2)$, the
 1291 cardinality constraint $T_1 = T_2$ if $\sigma(T_1) = \sigma(T_2)$, and the cardinality constraint $c \text{ dvd } T$ if c divides
 1292 $\sigma(T)$.

1293 We can now define the syntax of $\mathcal{ALCSCC}^{++}$ concept descriptions and their semantics. Let N_C be 3
 1294 a set of concept names, and N_R be a set of role names, such that $N_C \cap N_R = \emptyset$. Every $A \in N_C$
 1295 is a *concept description* of $\mathcal{ALCSCC}^{++}$. Moreover, if $C, C_1, C_2, \dots$ are *concept descriptions* of
 1296 $\mathcal{ALCSCC}^{++}$, then so are: $C_1 \sqcap C_2, C_1 \sqcup C_2, \neg C$, and $\text{sat}(\chi)$, where χ is a set or cardinality QFBAPA
 1297 constraint, with elements of N_R and concept descriptions $C_1, C_2, \dots$ used in place of set variables.

1298 A *finite interpretation* is a pair $I = (\Delta^I, \cdot^I)$, where Δ^I is a finite non-empty set of individuals, and 4
 1299 $\cdot^I$ is a function such that: every $A \in N_C$ is mapped to $A^I \subseteq \Delta^I$, and every $R \in N_R$ is mapped to
 1300 $R^I \subseteq \Delta^I \times \Delta^I$. Given an element of $d \in \Delta^I$, we define $R^I(d) = \{d' \mid (d, d') \in R^I\}$.

1301 The semantics of the language of $\mathcal{ALCSCC}^{++}$ makes use QFBAPA substitutions to interpret QFBAPA 5
 1302 constraints in terms of $\mathcal{ALCSCC}^{++}$ finite interpretations. Given an element $d \in \Delta^I$, we can define
 1303 the substitution σ_d^I in such a way that: $\sigma_d^I(\mathcal{U}) = \Delta^I$, $\sigma_d^I(\emptyset) = \emptyset$, and $A \in N_C$ and $R \in N_R$ are
 1304 considered QFBAPA set variables and substituted as $\sigma_d^I(A) = A^I$, and $\sigma_d^I(R) = R^I(d)$.

1305 The finite interpretation I and the QFBAPA substitutions σ_d^I are mutually extended to complex ex- 6
 1306 pressions such that: $\sigma_d^I(C_1 \sqcap C_2) = (C_1 \sqcap C_2)^I = C_1^I \cap C_2^I$; $\sigma_d^I(C_1 \sqcup C_2) = (C_1 \sqcup C_2)^I =$
 1307 $C_1^I \cup C_2^I$; $\sigma_d^I(\neg C) = (\neg C)^I = \Delta^I \setminus C^I$; and $\sigma_d^I(\text{sat}(\chi)) = (\text{sat}(\chi))^I = \{d' \in \Delta^I \mid$
 1308 $\sigma_{d'}^I(\text{QFBAPA}) \text{ satisfies } \chi\}$.

1309 **Definition 24.** The $\mathcal{ALCSCC}^{++}$ concept description C is satisfiable if there is a finite interpretation 7
 1310 I such that $C^I \neq \emptyset$.

1311 **Theorem 25** ([2]). The problem of deciding whether an $\mathcal{ALCSCC}^{++}$ concept description is satisfiable 8
 1312 is NEXPTIME-complete.

1313 I $\mathcal{ALCQ}$ and T_C Boxes consistency 9

1314 $\mathcal{ALCQ}$ is the Description Logic adding qualified number restrictions to the standard Description 10
 1315 Logic $\mathcal{ALC}$, analogously to how Graded Modal Logic extends standard Modal Logic with graded
 1316 modalities.

1317 Let N_C and N_R be two non-intersecting sets of concept names, and role names respectively. A 11
 1318 concept name $A \in N_C$ is an $\mathcal{ALCQ}$ concept expressions of $\mathcal{ALCQ}$. If C is an $\mathcal{ALCQ}$ concept
 1319 expression, so is $\neg C$. If C_1 and C_2 are $\mathcal{ALCQ}$ concept expressions, then so is $C_1 \sqcap C_2$. If C is an
 1320 $\mathcal{ALCQ}$ concept expression, $R \in N_R$, and $n \in \mathbb{N}$, then $\geq n R.C$ is an $\mathcal{ALCQ}$ concept expression.

1321 A *cardinality restriction* of $\mathcal{ALCQ}$ is an expression of the form $(\geq n C)$ or $(\leq n C)$, where C an 12
 1322 $\mathcal{ALCQ}$ concept expression and $n \in \mathbb{N}$.

1323 An $\mathcal{ALCQ}$ - T_C Box is a finite set of cardinality restrictions. 13

1324 An *interpretation* is a pair $I = (\Delta^I, \cdot^I)$, where Δ^I is a non-empty set of individuals, and $\cdot^I$ is 14
 1325 a function such that: every $A \in N_C$ is mapped to $A^I \subseteq \Delta^I$, and every $R \in N_R$ is mapped
 1326 to $R^I \subseteq \Delta^I \times \Delta^I$. Given an element of $d \in \Delta^I$, we define $R^I(d) = \{d' \mid (d, d') \in R^I\}$.
 1327 An interpretation I is extended to complex concept descriptions as follows: $(\neg C)^I = \Delta^I \setminus C^I$;
 1328 $(C_1 \sqcap C_2)^I = C_1^I \cap C_2^I$; and $(\geq n R.C)^I = \{d \mid |R^I(d) \cap C^I| \geq n\}$.

1329 An interpretation I satisfies the cardinality restriction $(\geq n C)$ iff $|C^I| \geq n$ and it satisfies 15
 1330 the cardinality restriction $(\leq n C)$ iff $|C^I| \leq n$. A T_C Box TC is *consistent* if there exists an
 1331 interpretation that satisfies all the cardinality restrictions in TC .

1332 **Theorem 26** ([36]). Deciding the consistency of $\mathcal{ALCQ}$ - T_C Boxes is NEXPTIME-hard. 16

1333 The proof can be slightly adapted to show that the result holds even when there is only one role. 17




Figure 8: Encoding a torus of exponential size with an $\mathcal{ALCQ}\text{-}T_C$ Box with one role.

Some abbreviations are useful. For every pair of concepts C and D , $C \rightarrow D$ stands for $\neg C \sqcup D$. For every concept C , role R , and non-negative integer n , we define: $(\leq n R.C) := \neg(\geq (n+1) R.C)$, $(\forall R.C) := (\leq 0 R.\neg C)$, $(\forall C) := (\leq 0 \neg C)$, $(= n R.C) := (\geq n R.C) \sqcap (\leq n R.C)$, and $(= n C) := (\geq n C) \sqcap (\leq n C)$.

Theorem 27. Deciding the consistency of $\mathcal{ALCQ}\text{-}T_C$ Boxes is NEXPTIME-hard even if $|N_R| = 1$.

Proof. Let $next$ be the unique role in N_R . We use the atomic concepts N to denote an individual ‘on the way north’ and E to denote an individual ‘on the way east’. See Figure 8.

For every $n \in \mathbb{N}$, we define the following $\mathcal{ALCQ}\text{-}T_C$ Box.

$$T_n = \left\{ \begin{array}{ll} (\forall \neg(N \sqcup E) \rightarrow (= 1 next.N)) & , \quad (\forall \neg(N \sqcup E) \rightarrow (= 1 next.E)) \\ (\forall N \rightarrow (= 1 next.\top)) & , \quad (\forall E \rightarrow (= 1 next.\top)) \\ (= 1 C_{(0,0)}) & , \quad (= 1 C_{(2^n-1, 2^n-1)}) \\ (\forall \neg(N \sqcup E) \rightarrow D_{east}) & , \quad (\forall \neg(N \sqcup E) \rightarrow D_{north}) \\ (\leq (2^n \times 2^n) \neg(N \sqcup E)), & (\leq (2^n \times 2^n) N), \quad (\leq (2^n \times 2^n) E) \end{array} \right\}$$

such that the concepts $C_{(0,0)}$, $C_{(2^n-1, 2^n-1)}$ are defined like in [36, Figure 3], and so are the concepts D_{north} and D_{east} , except that for every concept C , $\forall east.C$ now stands for $\forall next.(E \rightarrow \forall next.C)$ and $\forall north.C$ now stands for $\forall next.(N \rightarrow \forall next.C)$.

The problem of deciding whether a domino system $\mathcal{D} = (D, V, H)$, given an initial condition $w_0 \dots w_{n-1}$, can tile a torus of exponential size can be reduced to the problem of consistency of $\mathcal{ALCQ}\text{-}T_C$ Boxes, checking the consistency of $T(n, \mathcal{D}, w) = T_n \cup T_{\mathcal{D}} \cup T_w$, where T_n is as above, $T_{\mathcal{D}}$ encodes the domino system, and T_w encodes the initial condition as follows.

$$T_{\mathcal{D}} = \left\{ \begin{array}{l} (\forall \neg(N \sqcup E) \rightarrow (\bigsqcup_{d \in D} C_d)), \\ (\forall \neg(N \sqcup E) \rightarrow (\prod_{d \in D} \prod_{d' \in D \setminus \{d\}} \neg(C_d \sqcap C_{d'}))), \\ (\forall \prod_{d \in D} (C_d \rightarrow (\forall east. \bigsqcup_{(d, d') \in H} C_{d'}))), \\ (\forall \prod_{d \in D} (C_d \rightarrow (\forall north. \bigsqcup_{(d, d') \in V} C_{d'}))) \end{array} \right\}$$

$$T_w = \left\{ (\forall C_{(0,0)} \rightarrow C_{w_0}), \dots, (\forall C_{(n-1,0)} \rightarrow C_{w_{n-1}}) \right\}$$

The rest of the proof remains unchanged.

1350	NeurIPS Paper Checklist 1
1351	1. Claims 2
1352	Question: Do the main claims made in the abstract and introduction accurately reflect the 3
1353	paper's contributions and scope?
1354	Answer: [Yes] 4
1355	Justification: We introduce a logical language for reasoning about quantized graph neural 5
1356	networks (GNNs) with Global Readout in Section 3. We then prove that verifying quantized
1357	GNNs with Global Readout is NEXPTIME-complete in Section 4 and Section 5. We also
1358	experimentally show the relevance of quantization in the context of ACR-GNNs in Section 7.
1359	2. Limitations 6
1360	Question: Does the paper discuss the limitations of the work performed by the authors? 7
1361	Answer: [Yes] 8
1362	Limitations are addressed in Section 8. 9
1363	3. Theory assumptions and proofs 10
1364	Question: For each theoretical result, does the paper provide the full set of assumptions and 11
1365	a complete (and correct) proof?
1366	Answer: [Yes] 12
1367	Justification: All the theorems, formulas, and proofs in the paper are numbered and cross- 13
1368	referenced. The assumptions are stated and the full proofs are present in the appendix, with
1369	sketches of proofs in the main text.
1370	4. Experimental result reproducibility 14
1371	Question: Does the paper fully disclose all the information needed to reproduce the main ex- 15
1372	perimental results of the paper to the extent that it affects the main claims and/or conclusions
1373	of the paper (regardless of whether the code and data are provided or not)?
1374	Answer: [Yes] 16
1375	Justification: The authors provide the replication package with code and description of the 17
1376	files.
1377	5. Open access to data and code 18
1378	Question: Does the paper provide open access to the data and code, with sufficient instruc- 19
1379	tions to faithfully reproduce the main experimental results, as described in supplemental
1380	material?
1381	Answer: [Yes] 20
1382	Justification: We provided clear instructions on how to access the data and reproduce the ex- 21
1383	perimental results in the supplemental materials, including required scripts and environment
1384	setup.
1385	6. Experimental setting/details 22
1386	Question: Does the paper specify all the training and test details (e.g., data splits, hyper- 23
1387	parameters, how they were chosen, type of optimizer, etc.) necessary to understand the
1388	results?
1389	Answer: [Yes] 24
1390	Justification: The experimental setting is described in sufficient detail in the main body of the 25
1391	paper, including datasets, tools, parameters, and evaluation metrics, to support understanding
1392	and reproducibility of the results.
1393	7. Experiment statistical significance 26
1394	Question: Does the paper report error bars suitably and correctly defined or other appropriate 27
1395	information about the statistical significance of the experiments?
1396	Answer: [Yes] 28
1397	Justification: The authors provided a code in the supplementary materials that generates the 29
1398	detailed summary statistics across configurations for FOC_2 . The method for computing
1399	these plots is included in the code.

1400 **8. Experiments compute resources** 1

1401 Question: For each experiment, does the paper provide sufficient information on the com-2
1402 puter resources (type of compute workers, memory, time of execution) needed to reproduce
1403 the experiments?

1404 Answer: [Yes] 3

1405 Justification: The experiments were run on a Samsung Galaxy Book4 laptop with an Intel 4
1406 Core i7-150U processor, 16 GB RAM, and 1 TB SSD storage. Additional experiments
1407 were conducted using Kaggle's cloud platform with an NVIDIA Tesla P100 GPU (16 GB
1408 RAM). The runtime for the synthetic dataset experiments is reported in Table 21, and full
1409 instructions for reproducing the results are provided in the supplementary materials.

1410 **9. Code of ethics** 5

1411 Question: Does the research conducted in the paper conform, in every respect, with the 6
1412 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

1413 Answer: [Yes] 7

1414 Justification: The research conducted in the paper conforms, in every respect, with the 8
1415 NeurIPS Code of Ethics.

1416 **10. Broader impacts** 9

1417 Question: Does the paper discuss both potential positive societal impacts and negative 10
1418 societal impacts of the work performed?

1419 Answer: [Yes] 11

1420 Justification: Broader impacts are addressed in the introduction, explaining that the black-box 12
1421 nature of NN is a major issue for their adoption, morally and legally, with the enforcement
1422 of regulatory policies like the EU AI Act. NN that can be formally verified solve this. We
1423 do not think that this work may have negative societal impacts.

1424 **11. Safeguards** 13

1425 Question: Does the paper describe safeguards that have been put in place for responsible 14
1426 release of data or models that have a high risk for misuse (e.g., pretrained language models,
1427 image generators, or scraped datasets)?

1428 Answer: [NA] 15

1429 Justification: The paper poses no such risks. 16

1430 **12. Licenses for existing assets** 17

1431 Question: Are the creators or original owners of assets (e.g., code, data, models), used in 18
1432 the paper, properly credited and are the license and terms of use explicitly mentioned and
1433 properly respected?

1434 Answer: [Yes] 19

1435 Justification: For the reference ACR-GNN, we used the original paper [4] and the official 20
1436 implementation available at [5]. The code is distributed under the MIT License, and we
1437 have properly credited the authors and complied with the license terms.

1438 **13. New assets** 21

1439 Question: Are new assets introduced in the paper well documented and is the documentation 22
1440 provided alongside the assets?

1441 Answer: [Yes] 23

1442 Justification: We are releasing new code introduced in this work under the MIT License. The 24
1443 repository includes a README with setup instructions, usage examples, and description of
1444 each module, enabling other researchers to reproduce our results.

1445 **14. Crowdsourcing and research with human subjects** 25

1446 Question: For crowdsourcing experiments and research with human subjects, does the paper 26
1447 include the full text of instructions given to participants and screenshots, if applicable, as
1448 well as details about compensation (if any)?

1449 Answer: [NA] 27

1450	Justification: The paper does not involve crowdsourcing nor research with human subjects.	1
1451	15. Institutional review board (IRB) approvals or equivalent for research with human	2
1452	subjects	
1453	Question: Does the paper describe potential risks incurred by study participants, whether	3
1454	such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)	
1455	approvals (or an equivalent approval/review based on the requirements of your country or	
1456	institution) were obtained?	
1457	Answer: [NA]	4
1458	Justification: The paper does not involve crowdsourcing nor research with human subjects	5
1459	16. Declaration of LLM usage	6
1460	Question: Does the paper describe the usage of LLMs if it is an important, original, or	7
1461	non-standard component of the core methods in this research? Note that if the LLM is used	
1462	only for writing, editing, or formatting purposes and does not impact the core methodology,	
1463	scientific rigorousness, or originality of the research, declaration is not required.	
1464	Answer: [NA]	8
1465	Justification: The core method development in this research does not involve LLMs as any	9
1466	important, original, or non-standard components.	