
Outlier-Robust Phase Retrieval in Nearly-Linear Time¹

Anonymous Author(s)²

Affiliation

Address

email

Abstract³

1 Phase retrieval is a fundamental problem in signal processing, where the goal is to⁴
2 recover a (complex-valued) signal from phaseless intensity measurements. In this
3 paper, we propose and study the (real-valued) outlier-robust phase retrieval problem.
4 Specifically, we seek to recover a vector $x \in \mathbb{R}^d$ from n intensity measurements
5 $y_i = (a_i^\top x)^2$, where a small fraction of the (a_i, y_i) pairs are adversarially corrupted.
6 Our main result is a near-sample-optimal and nearly-linear-time algorithm that
7 provably recovers the ground-truth vector. Our algorithm first solves a lightweight
8 convex program to find an initial point close to the ground truth, and then runs
9 a robust version of gradient descent to achieve exact recovery. Our approach is
10 conceptually simple and provides a framework for developing robust algorithms
11 for other non-convex optimization problems.

12 1 Introduction⁵

13 Phase retrieval is a fundamental problem in signal processing with applications in various fields, in-⁶
14 cluding electron microscopy [35], crystallography [36, 39], astronomy [13], and optical imaging [40].
15 In these applications, one often has access to only the magnitudes of the Fourier transforms of a
16 complex signal. This is because measuring magnitude (e.g., by aggregating energy over time) is
17 much easier than measuring phase (which requires detecting rapid changes). We refer the reader to
18 the survey articles [40, 29] for more details about the theory and applications of phase retrieval.

19 In this paper, we focus on the (real-valued) generalized phase retrieval problem, where we are given⁷
20 intensity measurements of an arbitrary linear operator.

21 **Definition 1.1** (Phase Retrieval). *Let $x \in \mathbb{R}^d$ be the ground-truth vector. Let $a_1, \dots, a_n \in \mathbb{R}^d$ be⁸
22 n sampling vectors and let $y_i = \langle a_i, x \rangle^2 \in \mathbb{R}$ be the corresponding intensity measurements. Given
23 $(a_i, y_i)_{i=1}^n$ as input, the task is to recover x .*

24 Note that it is impossible to distinguish between x and $-x$, thus recovering either is sufficient. This⁹
25 problem has been extensively studied by the machine learning community. Under certain assumptions
26 on the distribution of the sampling vectors, such as when the a_i 's are independent and Gaussian
27 distributed, $\Theta(d)$ input pairs (a_i, y_i) are necessary and sufficient for exact recovery [8]. Additionally,
28 it has been shown that the problem can be solved in linear time with respect to the size of the input
29 [8]. This was first achieved using semidefinite programming (SDP) relaxations [6]. In practice, the
30 problem is often solved by applying first-order optimization methods (e.g., gradient descent) to a
31 suitable objective function such as

$$\min_{z \in \mathbb{R}^d} f(z) = \sum_{i=1}^n (y_i - \langle a_i, z \rangle^2)^2. \quad 10 \quad (1)$$

32 Even though many natural formulations of the phase retrieval objective are nonconvex, including¹¹
33 the one in (1), prior work has shown that, depending on the input distribution, they may not have
34 spurious local optima and thus can be solved using first-order optimization methods [37, 4, 43].

35 However, these analyses of the objective landscape rely on strong assumptions, such as the sampling 1
 36 vectors a_i being i.i.d. Gaussian. Our work is motivated by the following questions: Can we relax
 37 the assumptions used to prove landscape results for tractable nonconvex problems? In the context of
 38 phase retrieval, can we recover the ground-truth vector x when a small subset of the (a_i, y_i) pairs are
 39 adversarially corrupted?

40 Our focus on this adversarial setting, where an ϵ -fraction of the input data is corrupted, is inspired 2
 41 by recent advances in high-dimensional robust statistics. An example problem in robust statistics
 42 is to estimate the mean of a d dimensional spherical Gaussian when ϵ -fraction of the samples are
 43 arbitrarily corrupted. The goal of high-dimensional statistics is often to design efficient algorithms
 44 that can achieve dimension-independent error guarantees. Early work in robust statistics [45, 26, 28]
 45 provided sample-efficient estimators for various tasks, but with runtimes exponential in the dimension.
 46 A recent line of work initiated by [16] and [31] has developed computationally efficient robust
 47 algorithms for many fundamental statistical and learning tasks. Significant progress has been made in
 48 the algorithmic aspects of robust high-dimensional statistics (see, e.g., [15]).

49 We now formally define the ϵ -corruption model we study in this paper. For clarity, we define it 3
 50 directly in the context of phase retrieval.

51 **Definition 1.2** (ϵ -Corrupted Set of Samples). *Let $x \in \mathbb{R}^d$ be the ground-truth vector. Let $\epsilon > 0$. First, 4
 52 the algorithm specifies the sample complexity n . Then, n sampling vectors $(a_1, \dots, a_n)$ are drawn
 53 from some known distribution D and the corresponding intensity measurements $y_i = \langle a_i, x \rangle^2$ are
 54 calculated. The adversary is allowed to replace ϵn pairs (a_i, y_i) with arbitrary data. We call a set of
 55 (a_i, y_i) pairs ϵ -corrupted if it is generated by this process.*

56 In this paper, we say *samples* to refer to (a_i, y_i) pairs. Note that we allow corruption in both the 5
 57 sampling vectors $a_i \in \mathbb{R}^d$ and the intensity measurements $y_i \in \mathbb{R}$, as long as the fraction of corruption
 58 is at most ϵ . We now formally define outlier-robust phase retrieval, the main problem we study in this
 59 paper.

60 **Problem 1.3** (Outlier-Robust Phase Retrieval). *Let $x \in \mathbb{R}^d$ be the ground-truth with $\|x\|_2 = 1$. Let 6
 61 $\epsilon > 0$. Let $(a_1, \dots, a_n)$ be n sampling vectors drawn i.i.d. from $\mathcal{N}(0, I) \in \mathbb{R}^d$, and let $y_i = \langle a_i, x \rangle^2$
 62 be the corresponding intensity measurements. An adversary arbitrarily corrupts an ϵ -fraction of these
 63 (a_i, y_i) pairs, and gives the ϵ -corrupted set of samples as input to the algorithm. The task is to find a
 64 vector $z \in \mathbb{R}^d$ such that $\min\{\|z - x\|_2, \|z + x\|_2\} \leq \Delta$ for some precision parameter $\Delta > 0$.*

65 A simpler adversary setting where corruption is restricted to the intensity measurements y_i has 7
 66 been studied previously [27, 51, 19]. In our work, we study a more comprehensive setting where
 67 we also allow corruption in the measurement vectors. This models realistic scenarios where the
 68 measurement process is affected by hardware noise, miscalibration, or adversarial tampering, leading
 69 to perturbations in the sampling vectors a_i as well as in y_i . Concurrent work [3] studies this same
 70 problem. However, their algorithm uses a black-box subroutine for robust covariance estimation and
 71 thus requires at least $\Omega(d^2)$ time, making it impractical in high dimensions.

72 We are interested in designing a scalable and provably robust algorithm for Problem 1.3. We would 8
 73 like to resolve the following question:

74 *Can we design a provably robust algorithm for the outlier-robust phase retrieval 9
 75 problem (Problem 1.3) that has near-optimal sample complexity and runs in
 76 nearly-linear time?*

77 1.1 Our Results 10

78 We answer the question outlined in the previous subsection affirmatively. 11

79 **Theorem 1.4** (Main). *Consider the outlier-robust phase retrieval problem as defined in Problem 1.3, 12
 80 where $x \in \mathbb{R}^d$ is the ground-truth vector. Let $0 < \epsilon \leq \epsilon_0$ for sufficiently small universal constant ϵ_0 .
 81 Let $\Delta > 0$ be the desired precision. Given an ϵ -corrupted set of $n = \tilde{\Omega}(d \log^2(1/\Delta))$ samples, we
 82 can compute $z \in \mathbb{R}^d$ in time $\tilde{O}(nd)$ such that $\min(\|z - x\|_2, \|z + x\|_2) \leq \Delta$ with probability at
 83 least 0.95.¹*

¹Throughout the paper, we use $\tilde{O}(\cdot)$ and $\tilde{\Omega}(\cdot)$ to hide logarithmic factors in their parameters.

84 A few remarks are in order regarding Theorem 1.4. First, even without corruption, phase retrieval 1
 85 requires $\Omega(d)$ samples [8]. When $\Delta \geq \frac{1}{\text{poly}(d)}$, Theorem 1.4 requires $\tilde{O}(d)$ samples, runs in nearly
 86 linear time (since the input size is $O(nd)$), and achieves exact recovery. Therefore, our algorithm
 87 simultaneously achieves the best possible error, sample complexity, and runtime (up to logarithmic
 88 factors).

89 Second, the success probability in Theorem 1.4 can be boosted to $1 - \tau$ for any $\tau > 0$ by incurring 2
 90 an additional $O(\log(1/\tau))$ factor in the sample complexity and runtime. This can be achieved, for
 91 example, by partitioning the input, repeating the algorithm, letting candidate solutions vote for those
 92 within distance 2Δ , and finally selecting the solution with the most votes.

93 Lastly, ϵ_0 in Theorem 1.4 is an absolute constant that is independent of n or d , and our algorithm 3
 94 works for any corruption level $0 \leq \epsilon \leq \epsilon_0$. An important observation is that the optimal sample
 95 complexity is $\Theta(d)$, which is independent of ϵ . This follows from the fact that exact recovery is
 96 possible as long as the clean samples provide enough constraints to fully determine x , which has
 97 $d - 1$ degrees of freedom.² This is in contrast to problems in robust high-dimensional statistics, such
 98 as robust mean estimation, where exact recovery is impossible with a finite number of samples (even
 99 without corruption).

100 Since our algorithm guarantees exact recovery (to arbitrary precision Δ) for any corruption level ϵ as 4
 101 long as $\epsilon < \epsilon_0$, any input with corruption level $\epsilon < \epsilon_0$ can be treated as if it were corrupted at a fixed
 102 level ϵ_0 . This explains why neither the sample complexity nor the runtime of our algorithm depends
 103 on ϵ . For simplicity, we refer to the corruption level ϵ as a sufficiently small universal constant
 104 throughout the remainder of the paper unless otherwise noted.

105 1.2 Our Approach and Techniques 5

106 When there are infinitely many samples and no corruption, the objective function $f(z)$ simplifies to 6

$$f(z) = \mathbb{E}_{a \sim \mathcal{N}(0, I_d)} \left[(\langle a_i, z \rangle^2 - y_i)^2 \right] = 3 \|x\|_2^4 + 3 \|z\|_2^4 - 2 \|x\|_2^2 \|z\|_2^2 - 4 \langle x, z \rangle^2. \quad 7$$

107 Despite being nonconvex, it is known that f has no spurious local optima [37, 4, 43]. 8

108 Our approach follows a two-step structure used in many local convergence results for nonconvex 9
 109 problems (e.g., Candès et al. [4]), where the goal is to first initialize into a region free of saddle points
 110 and then perform gradient descent. Both steps are vulnerable to adversarial attacks and we develop
 111 provably robust and nearly-linear time algorithms for both steps. In the first step, we use spectral
 112 techniques to obtain an initial guess that is sufficiently close to the ground truth. In the second step,
 113 we run a robust gradient descent algorithm to refine this guess and converge to the final solution.

114 **Step 1: Robust Spectral Initialization.** Consider the following matrix $Y = \frac{1}{n} \sum_{i=1}^n y_i a_i a_i^\top$ 10
 115 where $y_i = \langle a_i, x \rangle^2$. When there is no corruption and the a_i 's are drawn i.i.d. from $\mathcal{N}(0, I)$,
 116 we have $\mathbb{E}[Y] = I + 2xx^\top$. Hence, when there are enough samples and no corruption, we can
 117 obtain a good estimate of the ground truth x (or $-x$) by computing the largest eigenvector of Y .
 118 However, the corrupted (a_i, y_i) pairs can arbitrarily change the largest eigenvector of Y . One natural
 119 approach, which was explored in concurrent work [3], is to apply known robust covariance estimation
 120 algorithms [11, 1] to estimate Y . While this can recover the top eigenvector, the runtime is at least
 121 $\Omega(d^2)$, which is too slow for our goal of designing a nearly linear time algorithm.

122 One of the main technical insights of our work is that it is not necessary to robustly estimate the entire 11
 123 matrix Y , we only need to recover its few largest eigenspaces. We will assign a weight w_i to each
 124 sample such that the weighted intensity-based matrix $Y_w = \sum_{i=1}^n w_i y_i a_i a_i^\top$ is close to $I + 2xx^\top$.
 125 We show that this can be done in nearly-linear time, which is highly non-trivial: the ground-truth
 126 vector x is unknown, and explicitly computing Y_w via fast matrix multiplication takes $\Omega(d^2)$ time.

127 A key observation is that the corrupted samples can only add directions to Y_w but cannot remove 12
 128 directions, because each individual term $y_i a_i a_i^\top$ is a PSD matrix. (We assume w.l.o.g. that $y_i \geq 0$ for

²To build intuition why exact recovery is possible, consider the univariate case with an arbitrary positive ground-truth $x \in \mathbb{R}$. In this case, one can compute the multiset $\{\sqrt{y_i/a_i^2}\}$, and as long as $\epsilon < 1/2$, its most frequent element will be the correct x with probability 1.

all i , because any input with $y_i < 0$ must be corrupted.) Consequently, if we can compute a weight vector w that minimizes the sum of the top two eigenvalues of Y_w (which is a convex optimization problem), we can recover a matrix that is close to the unknown unbiased expectation $I + 2xx^\top$. We show that this optimization problem can be solved in $\tilde{O}(nd)$ time by leveraging algorithmic techniques developed for list-decodable mean estimation [12].

Step 2: Robust Gradient Descent. Starting with the initial guess z given by the robust spectral initialization, we want to run gradient descent to recover the ground truth x . Without corruption, if the initialization is close enough to x , each iteration will bring z closer to x by a constant factor. Intuitively, approximating the gradient at a specific point amounts to a robust mean estimation problem (for the underlying distribution of the gradients). When the input data is ϵ -corrupted, the gradients of the n samples can be viewed as an ϵ -corrupted set of vectors.

We can approximate the true gradient by running robust mean estimation on this ϵ -corrupted set of n gradients. To show convergence, the estimation of the gradient needs to be more accurate as we get closer to the optimal solution, and we show that this is possible because the variance of the gradient on clean samples decreases as the solution gets closer to the optimal solution.

1.3 Related and Prior Works

Phase Retrieval. Phase retrieval arises in various fields of science and engineering [13, 36]. Early research introduced error-reduction algorithms [25, 20, 21]. Convex and nonconvex optimization with various objective functions were later proposed and achieved exact recovery [47, 4–6, 48, 49, 41].

Outlier-Robust Phase Retrieval. Robust phase retrieval has been explored in the literature [51, 30, 8, 7, 34]. A simpler setting where corruption is restricted to the intensity measurements y_i has been studied previously [27, 51, 19]. In Appendix C, we show that methods developed for this setting do not work in ours. Concurrent work [3] is the only one that studies the general corruption model that we consider, allowing corruption in both the sampling vectors a_i and the intensity measurements y_i . The algorithm in [3] achieves near-optimal sample complexity, but relies on robust covariance estimation in a black-box manner, resulting in a slower runtime compared to ours. We emphasize that algorithm runs in nearly-linear time and achieves near-optimal sample complexity, which demonstrates that allowing corruption in both a_i and y_i incurs almost no penalty (asymptotically) in terms of statistical or computational complexity.

Nonconvex Optimization. Besides phase retrieval, it is known that all local optima are globally optimal for natural nonconvex formulations of various learning problems, such as matrix completion [24], matrix sensing [2], dictionary learning [42], and tensor decomposition [23] (we refer the interested reader to Chapter 7 of the book by [50]). A recent line of work explored the robustness of such landscape results: [33] studied matrix sensing in the ϵ -corruption model, [9] and [22] studied semi-random matrix completion and matrix sensing.

High-Dimensional Robust Statistics. Recent works developed nearly-linear time algorithms for robust mean estimation [10, 18, 32]. The robust gradient descent algorithm we use is closely related to algorithms proposed in previous works for finding *first-order* stationary points in robust stochastic optimization [38, 17].

2 Preliminary and Background

Notation. Let $[n] = \{1, 2, \dots, n\}$. For a vector x , we denote its i^{th} coordinate by x_i . We use $\|x\|_1$, $\|x\|_2$, and $\|x\|_\infty$ to denote the ℓ_1 , ℓ_2 , and ℓ_∞ norm of x , respectively. For two vectors x and y , we use $\langle x, y \rangle = x^\top y$ to denote their inner product.

We write I for the identity matrix. For a matrix A , we use $\|A\|_2$ to denote its spectral norm. A symmetric matrix A is positive semidefinite (PSD) if $x^\top Ax \geq 0$ for all vectors x . For two symmetric matrices A and B , we write $A \preceq B$ if $B - A$ is PSD. We write $\lambda_k(A)$ as the k^{th} largest eigenvalue of A , and $\bar{\lambda}_k(A)$ as the sum of the k largest eigenvalues of A . The Ky Fan k -norm of a matrix A is the sum of its k largest singular values, which is equal to $\bar{\lambda}_k(A)$ when A is PSD.

Ky Fan Norm Packing SDP. In our robust spectral initialization step, we solve a Ky Fan norm packing semidefinite program (SDP) of the following form: 1

$$\max_{w \in \mathbb{R}_{\geq 0}^n} \|w\|_1 \quad \text{subject to} \quad \sum_{i=1}^n w_i A_i \preceq I, \quad \bar{\lambda}_k \left(\sum_{i=1}^n w_i B_i \right) \leq k \quad (**) \quad 2$$

We use the nearly-linear time Ky Fan norm SDP solver from [12]. 3

Lemma 2.1 (Ky Fan Norm SDP Solver [12]). *Given an SDP (**) with positive semidefinite matrices $A_i \in \mathbb{R}^{d_1 \times d_1}$ and $B_i \in \mathbb{R}^{d_2 \times d_2}$ with $A_i = C_i C_i^\top$ and $B_i = D_i D_i^\top$ for all $i \in [n]$, integer $k > 0$, error tolerance $\epsilon_1 \geq 1/n^2$, and failure probability $\tau > 0$, one can in time $\tilde{O}((t_C + t_D + d_1 + d_2) \text{poly}(1/\epsilon_1, \log(1/\tau)))$ output $w' \in \mathbb{R}_{\geq 0}^n$ such that $\|w'\|_1 \geq (1 - \epsilon_1) \text{OPT}$ with probability at least $1 - \tau$. Here OPT is the optimal value of (**), t_{C_i} and t_{D_i} are the time taken to perform a matrix-vector product with C_i and D_i respectively, and $t_C = \sum_{i=1}^n t_{C_i}$ and $t_D = \sum_{i=1}^n t_{D_i}$.* 4

Top Eigenvector Computation. We use the power method to compute an approximate top eigenvector. We refer to the analysis of the power method for PSD matrices by Trevisan [44]. 5

Lemma 2.2 (Top Eigenvector via Power Method [44]). *Let $A \in \mathbb{R}^{d \times d}$ be a PSD matrix. Let λ_1 be the largest eigenvalue of A . For any $\epsilon_2 > 0$, one can compute a unit vector $x \in \mathbb{R}^d$ in time $O((t_A + d) \log(d)/\epsilon_2)$ such that $x^\top A x \geq (1 - \epsilon_2) \lambda_1$ with probability at least 0.99, where t_A is the time taken to perform a matrix-vector multiplication with A .* 6

Robust Mean Estimation. In the robust gradient descent step, we use nearly-linear time robust mean estimation algorithms for bounded-covariance distributions [10, 14, 18] to approximate the true gradient. 7

Lemma 2.3 (Robust Mean Estimation [18]). *Let D be a distribution on $\mathbb{R}^d$ with unknown mean μ and unknown covariance matrix Σ where $\Sigma \preceq \sigma^2 I$. Let $\epsilon_3 > 0$ be a sufficiently small universal constant. Let $0 < \epsilon \leq \epsilon_3$ and $\tau > 0$. Given an ϵ -corrupted set of n samples drawn from D , one can output a vector $\hat{\mu} \in \mathbb{R}^d$ in time $\tilde{O}(nd \log(1/\tau))$ such that, with probability at least $1 - \tau - \exp(-n\epsilon)$, we have $\|\hat{\mu} - \mu\|_2 = O\left(\sqrt{\epsilon} + \sqrt{\frac{d}{n\tau}} + \sqrt{\frac{d(\log d + \log(1/\tau))}{n}}\right) \sigma$.* 8

3 Outlier-Robust Phase Retrieval 9

In this section, we present key technical lemmas for the two stages of our algorithm: robust spectral initialization (Lemma 3.1) and robust gradient descent (Lemma 3.2). We then use these lemmas to prove our main result (Theorem 1.4). 10

Lemma 3.1 shows that we can compute an initial guess close to the ground truth. 11

Lemma 3.1 (Robust Spectral Initialization). *Consider the setting of Problem 1.3, where $x \in \mathbb{R}^d$ is the ground-truth. Let ϵ be a sufficiently small universal constant. Given an ϵ -corrupted set of $n = \tilde{\Omega}(d)$ samples, we can compute $z_1 \in \mathbb{R}^d$ in time $\tilde{O}(nd)$ such that $\min\{\|z_1 - x\|_2, \|z_1 + x\|_2\} \leq \frac{1}{8}$ with probability at least 0.97.* 12

Lemma 3.2 shows that after the initialization, a robust gradient descent algorithm can recover the ground-truth vector to arbitrary precision. 13

Lemma 3.2 (Robust Gradient Descent). *Consider the setting of Problem 1.3, where $x \in \mathbb{R}^d$ is the ground-truth vector. Let $\Delta > 0$ be the desired precision. Let ϵ be a sufficiently small universal constant. Given an ϵ -corrupted set of $n = \tilde{\Omega}(d \log^2(1/\Delta))$ samples and an initial guess z_1 with $\|z_1 - x\|_2 \leq \frac{1}{8}$, we can compute a vector $z \in \mathbb{R}^d$ in time $\tilde{O}(nd)$ such that $\|z - x\|_2 \leq \Delta$ with probability at least 0.98.* 14

Lemma 3.1 is proved in Section 4. Lemma 3.2 is proved in Section 5. We first use these lemmas to prove Theorem 1.4. 15

Theorem 1.4. Let ϵ_0 be the minimum of the two universal constants of Lemma 3.1 and Lemma 3.2. Let $0 \leq \epsilon \leq \epsilon_0$. We assume that we have access to two separate ϵ -corrupted sets of samples, and use 16

one set for Lemma 3.1 and the other for Lemma 3.2. Formally, one could randomly partition the input samples into two sets and apply Chernoff bounds to show that both sets are (2ϵ) -corrupted with high probability. 1

By Lemma 3.1, we can compute a vector $z_1 \in \mathbb{R}^d$ such that $\min\{\|z_1 - x\|_2, \|z_1 + x\|_2\} \leq \frac{1}{8}$ for the ground-truth vector x . Given z_1 , by Lemma 3.2, we can output a vector $z \in \mathbb{R}^d$ such that $\min\{\|z - x\|_2, \|z + x\|_2\} \leq \Delta$ for the desired precision parameter $\Delta > 0$. 2

Let $n_1 = \tilde{\Omega}(d)$ and $n_2 = \tilde{\Omega}(d \log^2(1/\Delta))$ denote the number of samples used in Lemma 3.1 and Lemma 3.2, respectively. The overall sample complexity is therefore $n = n_1 + n_2 = \tilde{\Omega}(d \log^2(1/\Delta))$. The overall runtime is $\tilde{O}(n_1 d) + \tilde{O}(n_2 d) = \tilde{O}(nd)$. The overall success probability is at least 0.95 by a union bound over Lemma 3.1 and Lemma 3.2. 3 $\square$

4 Robust Spectral Initialization 4

In this section, we prove Lemma 3.1: given an ϵ -corrupted set of samples $\{(a_i, y_i)\}_{i \in [n]}$, we can compute an initial guess $z_1 \in \mathbb{R}^d$ that is close to the ground truth x or $-x$. 5

Consider the matrix $Y = \frac{1}{n} \sum_{i=1}^n y_i a_i a_i^\top$. When there is no corruption, where $a_i \sim \mathcal{N}(0, I)$ and $y_i = \langle a_i, x \rangle^2$, we have $\mathbb{E}[Y] = I + 2xx^\top$. However, the corrupted (a_i, y_i) 's can change Y arbitrarily. To address this, we propose a nearly-linear time initialization step (Algorithm 1) that computes a nonnegative weight vector $w \in \mathbb{R}^n$ such that the weighted sum $Y_w = \sum_{i=1}^n w_i y_i a_i a_i^\top$ is close to $I + 2xx^\top$. Consequently, we can show that the largest eigenvector of Y_w is close to $\pm x$. 6

Let $G \subset [n]$ be the set of indices of the remaining good samples. An ideal approach would be to assign weight $\frac{1}{|G|} = \frac{1}{(1-\epsilon)n}$ to every sample in G , and weight 0 to the corrupted samples. Formally, we consider weight vectors w in the set $\Delta_{n,\epsilon} = \left\{ w \in \mathbb{R}_{\geq 0}^n : \|w\|_1 = 1 \text{ and } \|w\|_\infty \leq \frac{1}{(1-\epsilon)n} \right\}$. Algorithm 1 computes a near-optimal solution $\hat{w}$ to the following optimization problem: 7

$$\min_{w \in \Delta_{n,\epsilon}} \lambda_1(Y_w) + \lambda_2(Y_w), \quad 8$$

and then returns the largest eigenvector of $Y_{\hat{w}}$. 9

Algorithm 1: Robust Spectral Initialization 10

Input: ϵ -corrupted set of n samples $\{(a_i, y_i)\}_{i \in [n]}$. 11

Output: An initial guess $z_1 \in \mathbb{R}^d$ of the ground-truth x s.t. $\min\{\|z_1 - x\|_2, \|z_1 + x\|_2\} \leq \frac{1}{8}$. 12

1 $\hat{w} \leftarrow$ a near-feasible, near-optimal solution to: $\min_{w \in \Delta_{n,\epsilon}} [\lambda_1(Y_w) + \lambda_2(Y_w)]$ using Lemma A.2, 13

where $Y_w = \sum_{i=1}^n w_i y_i a_i a_i^\top$;

2 $z_1 \leftarrow$ an approximate top eigenvector of $Y_{\hat{w}}$ using the power method (Lemma 2.2); 14

3 **return** z_1 ; 15

To prove that the largest eigenvector of $Y_{\hat{w}}$ is close to x , we will show that $x^\top Y_{\hat{w}} x$ is large, and there is a gap between the first and second largest eigenvalues of $Y_{\hat{w}}$. 16

Lemma 4.1. Consider the setting of Problem 1.3, where $x \in \mathbb{R}^d$ is the ground-truth vector. Fix $\delta > 0$. There exists constants $\epsilon(\delta)$ and $c(\delta)$ such that if we are given an $\epsilon(\delta)$ -corrupted set of $n = \tilde{\Omega}(c(\delta)d)$ samples $(a_i, y_i)_{i \in [n]}$, Algorithm 1 outputs $\hat{w} \in \mathbb{R}^n$ in time $\tilde{O}(nd \text{ poly}(1/\epsilon(\delta)))$ such that, with probability at least 0.98, the following conditions hold: 17

$$x^\top Y_{\hat{w}} x \geq 3 - O(\delta), \quad |\lambda_1(Y_{\hat{w}}) - 3| \leq O(\delta), \text{ and } |\lambda_2(Y_{\hat{w}}) - 1| \leq O(\delta). \quad 18$$

where $Y_{\hat{w}} = \sum_{i=1}^n \hat{w}_i y_i a_i a_i^\top$. 19

We defer the proof of Lemma 4.1 to Appendix A.1 and first use it to prove the correctness and runtime of Algorithm 1 (Lemma 3.1). 20

252 *Proof of Lemma 3.1.* Let $Y_{\hat{w}} = \sum_{i=1}^d \lambda_i v_i v_i^\top$ be the eigendecomposition of $Y_{\hat{w}}$, where $\lambda_1 \geq \lambda_2 \geq$ 1
 253 $\dots \geq \lambda_d$. Let $x = \sum_{i=1}^d \alpha_i v_i$. Note that $\sum_{i=1}^d \alpha_i^2 = \|x\|_2^2 = 1$ and $|\alpha_i| \leq 1$. If the event of
 254 Lemma 4.1 is true, with $\epsilon < \epsilon(\delta)$ and $m = \tilde{\Omega}(c(\delta)d)$, then

$$\begin{aligned} 3 - O(\delta) &\leq x^\top Y_{\hat{w}} x = \sum_{i=1}^d \lambda_i \alpha_i^2 \leq \lambda_1 \alpha_1^2 + \lambda_2 (1 - \alpha_1^2) \\ &\leq (3 + O(\delta)) \alpha_1^2 + (1 + O(\delta))(1 - \alpha_1^2) \leq 1 + O(\delta) + 2\alpha_1^2. \end{aligned} \quad 2$$

255 This implies $|\alpha_1| \geq \alpha_1^2 \geq 1 - O(\delta)$. Consequently, 3

$$\begin{aligned} \min\{\|v_1 - x\|_2^2, \|v_1 + x\|_2^2\} &= \min\{(1 - \alpha_1)^2, (1 + \alpha_1)^2\} + \sum_{i=2}^d \alpha_i^2 \\ &= \min\{2 - 2\alpha_1, 2 + 2\alpha_1\} = 2 - 2|\alpha_1| \leq O(\delta). \end{aligned} \quad 4$$

256 We choose δ as a sufficiently small constant so that $\min\{\|v_1 - x\|_2, \|v_1 + x\|_2\} \leq O(\sqrt{\delta}) \leq \frac{1}{16}$. 5
 257 Note that since it is sufficient to choose δ as a small universal constant, $c(\delta)$ and $\epsilon(\delta)$ can also be
 258 treated as universal constants.

259 We use the power method to approximate the largest eigenvector z_1 of $Y_{\hat{w}}$. By Lemma 2.2, $\|z_1\|_2 = 1$ 6
 260 and $z_1^\top Y_{\hat{w}} z_1 \geq (1 - \epsilon_2) \lambda_1$. Choosing $\epsilon_2 = O(\delta)$ and using Lemma 4.1, we have $z_1^\top Y_{\hat{w}} z_1 \geq 3 - O(\delta)$.
 261 By the same arguments, we can show that $\min\{\|v_1 - z_1\|_2, \|v_1 + z_1\|_2\} \leq \frac{1}{16}$, and then by the
 262 triangle inequality, we conclude that $\min\{\|z_1 - x\|_2, \|z_1 + x\|_2\} \leq \frac{1}{8}$.

263 The success probability of Algorithm 1 is at least 0.97, as $\hat{w}$ satisfies Lemma 4.1 with probability at 7
 264 least 0.98, and the power method succeeds with probability at least 0.99 by Lemma 2.2.

265 The required number of samples is $n = \tilde{\Omega}(d)$ by Lemma 4.1. Algorithm 1 runs in time $\tilde{O}(nd)$: It 8
 266 takes time $\tilde{O}(nd \text{poly}(1/\epsilon(\delta))) = \tilde{O}(nd)$ to compute $\hat{w}$ by Lemma 4.1. The power method can
 267 approximate the largest eigenvector of $Y_{\hat{w}}$ in time $O(nd \log(d)/\delta) = \tilde{O}(nd)$ by Lemma 2.2, since
 268 the matrix-vector product $Y_{\hat{w}} v = \sum_{i=1}^n \hat{w}_i y_i \langle a_i, v \rangle a_i$ can be computed in time $O(nd)$ for any
 269 $v \in \mathbb{R}^d$. $\square$

270 5 Robust Gradient Descent 9

271 After the robust initialization in Section 4, we have an initial guess $z_1 \in \mathbb{R}^d$ that is close to the ground 10
 272 truth x or $-x$. We can assume without loss of generality that z_1 is closer to x than to $-x$.

273 In this section, we prove Lemma 3.2: Given an initial guess z_1 with $\|z_1 - x\|_2 \leq \frac{1}{8}$, we can use a 11
 274 robust gradient descent algorithm (Algorithm 2) to recover x to any desired precision $\Delta > 0$. We
 275 will show that Algorithm 2 converges geometrically even when the input is ϵ -corrupted.

276 Consider the following nonconvex optimization problem: 12

$$\min_{z \in \mathbb{R}^d} \sum_{i=1}^n f_i(z) \quad \text{where} \quad f_i(z) = (\langle a_i, z \rangle^2 - y_i)^2. \quad 13$$

277 Let g_i denote the gradient of f_i with respect to z . Let $\mathcal{D}_z$ denote the distribution of $g_i(z)$ when there 14
 278 is no corruption. Formally, $g(z) \sim \mathcal{D}_z$ is distributed as

$$g(z) = \frac{\partial}{\partial z} \left[(\langle a, z \rangle^2 - \langle a, x \rangle^2)^2 \right] = 4 (\langle a, z \rangle^2 - \langle a, x \rangle^2) \langle a, z \rangle a \quad \text{where } a \sim \mathcal{N}(0, I). \quad 15 \quad (2)$$

279 To perform gradient descent, we want to approximate the *expected true gradient* 16

$$\mu_z = \mathbb{E}_{g(z) \sim \mathcal{D}_z} [g(z)] = (12 \|z\|_2^2 - 4 \|x\|_2^2) z - 8 \langle x, z \rangle x. \quad 17 \quad (3)$$

280 However, the input $\{(a_i, y_i)\}_{i \in [n]}$ is ϵ -corrupted, so the corresponding gradients $\{g_i(z)\}_{i \in [n]}$ are an 18
 281 ϵ -corrupted set of vectors drawn from $\mathcal{D}_z$. We will run nearly-linear time robust mean estimation
 282 algorithms (e.g., [18]) on $\{g_i(z)\}_{i \in [n]}$ to approximate the true gradient μ_z .

283 The accuracy of robust mean estimation algorithms depends on the covariance matrix Σ_z of the 19
 284 distribution $\mathcal{D}_z$. The following lemma upper bounds the spectral norm of Σ_z .

Algorithm 2: Robust Gradient Descent 1

Input: An ϵ -corrupted set of n samples $\{(a_i, y_i)\}_{i \in [n]}$, an initial guess z_1 with $\|z_1 - x\| \leq \frac{1}{8}$, 2
and desired precision $\Delta > 0$.

Output: $z \in \mathbb{R}^d$ such that $\|z - x\|_2 \leq \Delta$, where x is the ground-truth vector. 3

1 $T \leftarrow O(\log(1/\Delta)), \eta \leftarrow \frac{1}{300}$; 4
2 $\{N_1, \dots, N_T\} \leftarrow$ a random disjoint partition of $[n]$ such that $|N_t| = \frac{n}{T}$ for all $t \in [T]$;
3 **for** $t = 1, 2, \dots, T$ **do** 5
4 $g_i(z_t) \leftarrow 4 \left(\langle a_i, z_t \rangle^2 - y_i \right) \langle a_i, z_t \rangle a_i$;
5 $\hat{\mu}_{z_t} \leftarrow$ robust mean estimation on input $\{g_i(z_t)\}_{i \in N_t}$ using Lemma 5.2;
6 $z_{t+1} \leftarrow z_t - \eta \hat{\mu}_{z_t}$;
7 **end** 6
8 **return** z_{T+1} ; 7

285 **Lemma 5.1.** Let $x \in \mathbb{R}^d$ be the ground-truth vector. Let $\mathcal{D}_z$ be the distribution of gradients at z as 8
286 defined in Equation (2). For any z with $\|z - x\|_2 \leq 1$, the covariance matrix Σ_z of $\mathcal{D}_z$ satisfies

$$\Sigma_z \preceq O\left(\|z - x\|_2^2\right) I. \quad 9$$

287 We defer the proof of Lemma 5.1 to Appendix B. For technical reasons, we randomly partition the 10
288 input $\{(a_i, y_i)\}_{i \in [n]}$ into T subsets and use one subset in each iteration. With high probability, each
289 partition has at most (2ϵ) -fraction of corrupted samples. The next lemma shows that, given the
290 covariance bound in Lemma 5.1, we can approximate the true gradient μ_z from a (2ϵ) -corrupted set
291 of gradients with a small error.

292 **Lemma 5.2.** Let $x \in \mathbb{R}^d$ be the ground-truth vector. Consider any $z \in \mathbb{R}^d$ with $\|z - x\|_2 \leq 1$. Let 11
293 $\mathcal{D}_z$ be the distribution defined in Equation (2) and let μ_z be the mean of $\mathcal{D}_z$. Let $c > 0$, and let
294 $\tau \in (0, 1/4)$ be a small constant. There exists a constant $\epsilon(c)$ such that given a $2\epsilon(c)$ -corrupted set of
295 $m = \Omega(d \log(d)/(c^2 \tau))$ vectors drawn from $\mathcal{D}_z$, we can compute $\hat{\mu}_z \in \mathbb{R}^d$ in time $\tilde{O}(md \log(1/\tau))$
296 such that $\|\hat{\mu}_z - \mu_z\|_2 \leq c \|z - x\|_2$ with probability at least $1 - 2\tau$.

297 *Proof of Lemma 5.2.* We constraint $\epsilon(c) \leq \epsilon_3/2$, where ϵ_3 is the universal constant in the robust 12
298 mean estimation algorithm stated in Lemma 2.3. By Lemma 5.1, the covariance matrix Σ_z of $\mathcal{D}_z$
299 satisfies $\Sigma_z \preceq O\left(\|z - x\|_2^2\right) I$. Lemma 2.3 guarantees that robust mean estimation returns a vector
300 $\hat{\mu}_z \in \mathbb{R}^d$ such that

$$\|\hat{\mu}_z - \mu_z\|_2 \leq O\left(\sqrt{2\epsilon(c)} + \sqrt{\frac{d}{m\tau}} + \sqrt{\frac{d(\log d + \log(1/\tau))}{m}}\right) \sqrt{\|\Sigma_z\|_2} \leq c \|z - x\|_2. \quad 13$$

301 The last inequality follows by properly choosing the constant $\epsilon(c) = O(c^2)$. By Lemma 2.3, the 14
302 runtime is $\tilde{O}(md \log(1/\tau))$ and the success probability is at least $1 - \tau - \exp(-\epsilon_0 m) \geq 1 - 2\tau$. $\square$

303 The next lemma shows that the approximate gradient from Lemma 5.2 is sufficient for gradient 15
304 descent to converge, reducing the distance to the ground truth x by a constant factor in each iteration.
305 We provide a proof sketch for Lemma 5.3 and defer the full proof to Appendix B.

306 **Lemma 5.3.** Let $x \in \mathbb{R}^d$ be the ground-truth vector. Suppose at iteration t of Algorithm 2, the 16
307 current solution z_t satisfies $\|z_t - x\|_2 \leq \frac{1}{8}$. Let μ_{z_t} denote the expected true gradient at z_t defined
308 in Equation (3). Suppose the estimated gradient $\hat{\mu}_{z_t} \in \mathbb{R}^d$ satisfies $\|\hat{\mu}_{z_t} - \mu_{z_t}\|_2 \leq c \|z_t - x\|_2$ for
309 $c = 4$. Then, we have

$$\|z_{t+1} - x\|_2^2 \leq 0.99 \|z_t - x\|_2^2. \quad 17$$

310 *Proof Sketch of Lemma 5.3.* Even though the objective function is nonconvex, it is known that gradi- 18
311 ent descent is well-behaved when initialized close enough to a global optimum [37]. More specifically,

for any z with $\|z - x\|_2 \leq \frac{1}{8}$, we can show that the expected true gradient at z aligns with the direction $(z - x)$:

$$\langle \mu_z, z - x \rangle \geq 7.5 \|z - x\|_2^2 \quad \text{and} \quad \|\mu_z\|_2 \leq 29 \|z - x\|_2, \quad 2$$

which is sufficient for proving geometric convergence. 3

Note that this analysis is robust to small error in μ_z . When $\|\hat{\mu}_z - \mu_z\|_2 \leq c \|z - x\|_2$, we have 4

$$\langle \hat{\mu}_z, z - x \rangle \geq (7.5 - c) \|z - x\|_2^2 \quad \text{and} \quad \|\hat{\mu}_z\|_2 \leq (29 + c) \|z - x\|_2. \quad 5$$

When $c < 7.5$, we can choose an appropriate step size η such that the distance between z and x decreases by a constant factor in each iteration. 6 $\square$

We are now ready to prove Lemma 3.2, which states the correctness and runtime of Algorithm 2. 7

Lemma 3.2. First, we analyze the success probability of Algorithm 2. Algorithm 2 can fail in two ways: (i) some N_t has more than (2ϵ) -fraction of corrupted samples, or (ii) robust gradient estimation fails in some iteration t . The probability of event (i) is at most 0.01 for our choice of n , which follows from a standard application of Hoeffding's inequality and a union bound. For event (ii), we choose $\tau \leq 0.005/T$ in Lemma 5.2, so each robust gradient estimation fails with probability at most $2\tau = 0.01/T$. By a union bound over T iterations, the probability of event (ii) is at most 0.01. For the rest of the proof, we assume these bad events do not happen. 8

Next, we prove the correctness of Algorithm 2. Since $\|z_1 - x\|_2 \leq \frac{1}{8}$, we can use Lemma 5.2 to obtain an approximation $\hat{\mu}_{z_1}$ of the true gradient μ_{z_1} such that $\|\hat{\mu}_{z_1} - \mu_{z_1}\|_2 \leq c \|z_1 - x\|_2$ with $c = 4$. Then, by Lemma 5.3, we have $\|z_2 - x\|_2 \leq 0.99 \|z_1 - x\|_2$ after one iteration of gradient descent. Applying these two lemmas repeatedly, after $T = O(\log(1/\Delta))$ iterations, we have $\|z_{T+1} - x\|_2 \leq \Delta$. 9

Finally, we analyze the sample complexity and runtime of Algorithm 2. The algorithm requires in total $n = mT = \Omega(d \log d \log^2(1/\Delta))$ samples. A random partition can be computed in $O(n)$ time via random shuffling. In each iteration, the m gradients in N_t can be computed in time $O(md)$ using Equation (2). By Lemma 5.2, the true gradient can be robustly estimated in time $\tilde{O}(md \log(1/\tau)) = \tilde{O}(md \log T) = \tilde{O}(md \log \log(1/\Delta))$, and z_t can be updated in time $O(d)$. The overall runtime of Algorithm 2 is $\tilde{O}(n + Tmd \log \log(1/\Delta)) = \tilde{O}(nd)$. 10 $\square$

Remark. There are two technical details worth noting. First, robust mean estimation algorithms (Lemma 2.3) require a *known* upper bound σ^2 on the spectral norm of the covariance matrix of $\mathcal{D}_z$. By Lemma 5.1, a *known* upper bound on $\|z - x\|_2$ suffices. We can indeed maintain such an upper bound, which starts at $\frac{1}{8}$ and decreases geometrically, as shown in Lemma 5.3. Second, at each iteration t , to apply Lemma 5.2 to robustly estimate the gradient at z_t , we need a (2ϵ) -corrupted set of gradients drawn from $\mathcal{D}_{z_t}$. This is why we use a set of fresh samples N_t . By the principle of deferred decisions, we can view $(a_i, y_i)_{i \in N_t}$ as being generated and corrupted *after* z_t is chosen. 11

6 Conclusions and Future Directions 12

In this paper, we propose and study the problem of outlier-robust phase retrieval, where a small fraction of the input data is corrupted. Importantly, we allow adversarial corruption in both the sampling vectors $a_i \in \mathbb{R}^d$ and the intensity measurements $y_i \in \mathbb{R}$. We present a near-sample-optimal and nearly-linear-time algorithm for this problem with provable guarantees. One conceptual contributions of our work is that phase retrieval can be solved using a robust first-order methods even when the input is slightly misspecified or corrupted. Our algorithmic framework provides a general approach for developing robust algorithms for a wide range of tractable nonconvex problems, by first robustly initializing into a region free of saddle points and then using robust gradient descent to converge to a global optimum. 13

An immediate technical question is whether our sample complexity can be tightened by removing the $\log(1/\Delta)$ factors. One potential approach is to examine the stability conditions required by robust mean estimation algorithms and see if these conditions can be proved without using fresh samples in each iteration. 14

References

- [1] P. Abdalla and N. Zhivotovskiy. Covariance estimation: Optimal dimension-free guarantees for adversarial corruption and heavy tails. *Journal of the European Mathematical Society*, 2024.
- [2] S. Bhojanapalli, B. Neyshabur, and N. Srebro. Global optimality of local search for low rank matrix recovery. *Advances in Neural Information Processing Systems*, 29, 2016.
- [3] A. Buna and P. Rebeschini. Robust gradient descent for phase retrieval. In *The 28th International Conference on Artificial Intelligence and Statistics*.
- [4] E. J. Candès, X. Li, and M. Soltanolkotabi. Phase retrieval via wirtinger flow: Theory and algorithms. *IEEE Trans. Inf. Theory*, 61(4):1985–2007, 2015a.
- [5] E. J. Candès, X. Li, and M. Soltanolkotabi. Phase retrieval from coded diffraction patterns. *Applied and Computational Harmonic Analysis*, 39(2):277–299, Sept. 2015b. ISSN 1063-5203.
- [6] E. J. Candès, Y. C. Eldar, T. Strohmer, and V. Voroninski. Phase retrieval via matrix completion. *SIAM Rev.*, 57(2):225–251, 2015c.
- [7] J. Chen, L. Wang, X. Zhang, and Q. Gu. Robust Wirtinger Flow for Phase Retrieval with Arbitrary Corruption, Jan. 2018.
- [8] Y. Chen and E. Candes. Solving Random Quadratic Systems of Equations Is Nearly as Easy as Solving Linear Systems. In *Advances in Neural Information Processing Systems*, volume 28. Curran Associates, Inc., 2015.
- [9] Y. Cheng and R. Ge. Non-convex matrix completion against a semi-random adversary. In *Proc. 31st Conference on Learning Theory (COLT)*, volume 75, pages 1362–1394, 2018.
- [10] Y. Cheng, I. Diakonikolas, and R. Ge. High-dimensional robust mean estimation in nearly-linear time. In *Proc. 30th ACM-SIAM Symposium on Discrete Algorithms (SODA)*, pages 2755–2771, 2019.
- [11] Y. Cheng, I. Diakonikolas, R. Ge, and D. P. Woodruff. Faster algorithms for high-dimensional robust covariance estimation. In A. Beygelzimer and D. Hsu, editors, *Proceedings of the Thirty-Second Conference on Learning Theory*, volume 99 of *Proceedings of Machine Learning Research*, pages 727–757. PMLR, 25–28 Jun 2019.
- [12] Y. Cherapanamjeri, S. Mohanty, and M. Yau. List decodable mean estimation in nearly linear time. In *Proc. 61st IEEE Symposium on Foundations of Computer Science (FOCS)*, pages 141–148, 2020.
- [13] C. Dainty and J. R. Fienup. Phase retrieval and image reconstruction for astronomy. *Image recovery: theory and application*, 231:275, 1987.
- [14] J. Depersin and G. Lecué. Robust sub-Gaussian estimation of a mean vector in nearly linear time. *The Annals of Statistics*, 50(1):511–536, 2022.
- [15] I. Diakonikolas and D. M. Kane. *Algorithmic High-Dimensional Robust Statistics*. Cambridge University Press, 2023.
- [16] I. Diakonikolas, G. Kamath, D. M. Kane, J. Li, A. Moitra, and A. Stewart. Robust estimators in high dimensions without the computational intractability. In *57th Annual IEEE Symposium on Foundations of Computer Science—FOCS 2016*, pages 655–664. IEEE Computer Soc., Los Alamitos, CA, 2016.
- [17] I. Diakonikolas, G. Kamath, D. Kane, J. Li, J. Steinhardt, and A. Stewart. Sever: A robust meta-algorithm for stochastic optimization. In K. Chaudhuri and R. Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 1596–1606. PMLR, 09–15 Jun 2019.
- [18] Y. Dong, S. B. Hopkins, and J. Li. Quantum entropy scoring for fast robust mean estimation and improved outlier detection. In *Proc. 33rd Advances in Neural Information Processing Systems (NeurIPS)*, 2019.

- [19] J. C. Duchi and F. Ruan. Solving (most) of a set of quadratic equalities: Composite optimization for robust phase retrieval. *Information and Inference: A Journal of the IMA*, 8(3):471–529, 2019.
- [20] J. R. Fienup. Reconstruction of an object from the modulus of its Fourier transform. *Optics Letters*, 3(1):27–29, July 1978. ISSN 1539-4794.
- [21] J. R. Fienup. Phase retrieval algorithms: A comparison. *Applied Optics*, 21(15):2758–2769, Aug. 1982. ISSN 2155-3165.
- [22] X. Gao and Y. Cheng. Robust matrix sensing in the semi-random model. *Proceedings of the 37th Conference on Neural Information Processing Systems (NeurIPS)*, 2023.
- [23] R. Ge, F. Huang, C. Jin, and Y. Yuan. Escaping from saddle points—online stochastic gradient for tensor decomposition. In *Conference on Learning Theory*, pages 797–842, 2015.
- [24] R. Ge, J. D. Lee, and T. Ma. Matrix completion has no spurious local minimum. In *Advances in Neural Information Processing Systems*, pages 2973–2981, 2016.
- [25] R. W. Gerchberg. A practical algorithm for the determination of phase from image and diffraction plane pictures. *Optik*, Jan. 1972.
- [26] F. R. Hampel, E. M. Ronchetti, P. J. Rousseeuw, and W. A. Stahel. *Robust statistics. The approach based on influence functions*. Wiley New York, 1986.
- [27] P. Hand and V. Voroninski. Corruption robust phase retrieval via linear programming. *CoRR*, abs/1612.03547, 2016.
- [28] P. J. Huber and E. M. Ronchetti. *Robust statistics*. Wiley New York, 2009.
- [29] K. Jaganathan, Y. C. Eldar, and B. Hassibi. Phase retrieval: An overview of recent developments. *Optical Compressive Imaging*, pages 279–312, 2016.
- [30] R. Kolte and A. Özgür. Phase Retrieval via Incremental Truncated Wirtinger Flow, June 2016.
- [31] K. A. Lai, A. B. Rao, and S. Vempala. Agnostic estimation of mean and covariance. In *focs2016*, pages 665–674, 2016.
- [32] G. Lecué, M. Lerasle, and T. Mathieu. Robust classification via MOM minimization. *Mach. Learn.*, 109(8):1635–1665, 2020.
- [33] S. Li, Y. Cheng, I. Diakonikolas, J. Diakonikolas, R. Ge, and S. Wright. Robust second-order nonconvex optimization and its application to low rank matrix sensing. In *Proc. 37th Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [34] J.-W. Liu, Z.-J. Cao, J. Liu, X.-L. Luo, W.-M. Li, N. Ito, and L.-C. Guo. Phase Retrieval via Wirtinger Flow Algorithm and Its Variants. In *2019 International Conference on Machine Learning and Cybernetics (ICMLC)*, pages 1–9, July 2019.
- [35] J. Miao, T. Ishikawa, B. Johnson, E. H. Anderson, B. Lai, and K. O. Hodgson. High resolution 3D X-ray diffraction microscopy. *Physical review letters*, 89(8):088303, 2002.
- [36] R. P. Millane. Phase retrieval in crystallography and optics. *JOSA A*, 7(3):394–411, 1990.
- [37] P. Netrapalli, P. Jain, and S. Sanghavi. Phase retrieval using alternating minimization. In *Proc. 27th Advances in Neural Information Processing Systems (NeurIPS)*, pages 2796–2804, 2013.
- [38] A. Prasad, A. S. Suggala, S. Balakrishnan, and P. Ravikumar. Robust estimation via robust gradient estimation. *Journal of the Royal Statistical Society. Series B. Statistical Methodology*, 82(3):601–627, 2020.
- [39] W. H. Robert. Phase problem in crystallography. *JOSA a*, 10(5):1046–1055, 1993.
- [40] Y. Shechtman, Y. C. Eldar, O. Cohen, H. N. Chapman, J. Miao, and M. Segev. Phase retrieval with application to optical imaging: a contemporary overview. *IEEE signal processing magazine*, 32(3):87–109, 2015.

- 450 [41] M. Soltanolkotabi. Structured signal recovery from quadratic measurements: Breaking sample 1
451 complexity barriers via nonconvex optimization. *IEEE Trans. Inf. Theory*, 65(4):2374–2400,
452 2019.
- 453 [42] J. Sun, Q. Qu, and J. Wright. Complete dictionary recovery over the sphere i: Overview and the
454 geometric picture. *IEEE Trans. Inf. Theor.*, 63(2):853–884, 2 2017.
- 455 [43] J. Sun, Q. Qu, and J. Wright. A geometric analysis of phase retrieval. *Found. Comput. Math.*,
456 18(5):1131–1198, 2018.
- 457 [44] L. Trevisan. Lecture notes on graph partitioning, expanders and spectral methods. *University of*
458 *California, Berkeley*, <https://people.eecs.berkeley.edu/luca/books/expanders-2016.pdf>, 2017.
- 459 [45] J. Tukey. Mathematics and picturing of data. In *Proceedings of ICM*, volume 6, pages 523–531,
460 1975.
- 461 [46] R. Vershynin. *High-dimensional probability*, volume 47 of *Cambridge Series in Statistical and*
462 *Probabilistic Mathematics*. Cambridge University Press, Cambridge, 2018.
- 463 [47] V. Voroninski. *PhaseLift: A Novel Methodology for Phase Retrieval*. PhD thesis, UC Berkeley,
464 2013.
- 465 [48] G. Wang, G. Giannakis, Y. Saad, and J. Chen. Solving most systems of random quadratic
466 equations. *Advances in Neural Information Processing Systems*, 30, 2017.
- 467 [49] G. Wang, G. B. Giannakis, Y. Saad, and J. Chen. Phase retrieval via reweighted amplitude flow.
468 *IEEE Transactions on Signal Processing*, 66(11):2818–2833, 2018.
- 469 [50] J. Wright and Y. Ma. *High-dimensional data analysis with low-dimensional models: Principles,*
470 *computation, and applications*. Cambridge University Press, 2022.
- 471 [51] H. Zhang, Y. Chi, and Y. Liang. Provable non-convex phase retrieval with outliers: Median
472 truncated wirtinger flow. In *International conference on machine learning*, pages 1022–1031.
473 PMLR, 2016.

474 A Omitted Proofs in Section 4 ¹

475 A.1 Proof of Lemma 4.1 ²

476 In this section, we prove Lemma 4.1: We can compute $\hat{w} \in \mathbb{R}^n$ in nearly-linear time such that $x^\top Y_{\hat{w}} x$ ³
477 is large and the two largest eigenvalues of $Y_{\hat{w}}$ are approximately 3 and 1.

478 The following lemma shows that, for any $w \in \Delta_{n,2\epsilon}$, the contribution of the remaining good samples ⁴
479 to Y_w is close to $I + 2xx^\top$. This allows us to lower bound $\lambda_1(Y_{\hat{w}})$, $\lambda_2(Y_{\hat{w}})$, and $x^\top Y_{\hat{w}} x$.

480 **Lemma A.1.** *Let $x \in \mathbb{R}^d$ be the ground-truth vector. Fix $\delta > 0$. There exists constants $\bar{\epsilon}(\delta)$ and ⁵
481 $c(\delta)$ such that if we let $0 < \epsilon \leq \bar{\epsilon}(\delta)$, and we are given an ϵ -corrupted set of $n = \tilde{\Omega}(c(\delta)d)$ samples
482 $(a_i, y_i)_{i \in [n]}$, with probability at least 0.99, for all $w \in \Delta_{n,2\epsilon}$,*

$$\|Y_{G,w} - (I + 2xx^\top)\|_2 \leq \delta, \quad 6$$

483 where G is the set of indices of the remaining good samples and $Y_{G,w} = \sum_{i \in G} w_i y_i a_i a_i^\top$. ⁷

484 Intuitively, Lemma A.1 holds because the moments of Gaussian distributions are stable when a small ⁸
485 fraction of the samples are removed. We defer the proof of Lemma A.1 to Appendix A.

486 The next lemma shows that, assuming Lemma A.1 holds, we can compute a near-optimal $\hat{w}$ that ⁹
487 minimizes $\lambda_1(Y_{\hat{w}}) + \lambda_2(Y_{\hat{w}})$ in nearly-linear time.

488 **Lemma A.2.** *Let $\delta \in (0, 1/2)$. There exists a constant $\epsilon(\delta)$ such that if we are given an $\epsilon(\delta)$ -¹⁰
489 corrupted set of n samples $(a_i, y_i)_{i \in [n]}$ that satisfies Lemma A.1, we can compute $\hat{w} \in \Delta_{n,2\epsilon(\delta)}$ in
490 time $\tilde{O}(nd \text{ poly}(1/\epsilon(\delta)))$ such that with probability at least 0.99,*

$$\lambda_1(Y_{\hat{w}}) + \lambda_2(Y_{\hat{w}}) \leq 4 + O(\delta), \quad 11$$

491 where d is the dimension of a_i and $Y_{\hat{w}} = \sum_{i=1}^n \hat{w}_i y_i a_i a_i^\top$. ¹²

492 *Proof.* We reduce the optimization problem $\min_{w \in \Delta_{n,\epsilon(\delta)}} [\lambda_1(Y_w) + \lambda_2(Y_w)]$ to the following Ky ¹³
493 Fan norm packing semidefinite program (SDP):

$$\max_{w \in \mathbb{R}_{\geq 0}^n} \|w\|_1 \quad \text{subject to} \quad \sum_{i=1}^n w_i A_i \preceq I, \quad \bar{\lambda}_2 \left(\sum_{i=1}^n w_i B_i \right) \leq 4 + O(\delta) \quad 14 \quad (*)$$

494 in which $A_i = (1 - \epsilon(\delta))n e_i e_i^\top$ where $e_i \in \mathbb{R}^n$ is the i^{th} basis vector, $B_i = y_i a_i a_i^\top$, and $\bar{\lambda}_2$ is the ¹⁵
495 sum of the two largest eigenvalues.

496 Let G be the set of indices of the remaining good samples. Consider the weight vector $w^* \in \mathbb{R}^n$ that ¹⁶
497 is uniform on G :

$$w_i^* = \begin{cases} \frac{1}{|G|} = \frac{1}{(1-\epsilon)n} & i \in G, \\ 0 & i \notin G. \end{cases} \quad 17$$

498 Let $\epsilon(\delta) = \min\{\delta, \bar{\epsilon}(\delta)\}$, where $\bar{\epsilon}(\delta)$ is the constant of Lemma A.1. By Lemma A.1, $Y_{w^*} \preceq$ ¹⁸
499 $(1 + \delta)I + 2xx^\top$, which implies $\bar{\lambda}_2(Y_{w^*}) \leq 4 + O(\delta)$. Since w^* is feasible, the optimal value OPT
500 of (*) must be at least $\|w^*\|_1 = 1$.

501 We invoke Lemma 2.1 to solve (*) with failure probability $\tau = 0.01$ and error tolerance parameter ¹⁹
502 $\epsilon_1 = \epsilon(\delta)$. The resulting solution $w' \in \mathbb{R}_{\geq 0}^n$ satisfies $\|w'\|_1 \geq (1 - \epsilon_1)\text{OPT} \geq 1 - \epsilon(\delta)$. Define
503 $\hat{w} = \frac{w'}{\|w'\|_1}$ so that $\|\hat{w}\|_1 = 1$. The constraint with A_i guarantees that $\|w'\|_\infty \leq \frac{1}{(1-\epsilon(\delta))n}$, so
504 $\|\hat{w}\|_\infty \leq \frac{\|w'\|_\infty}{\|w'\|_1} \leq \frac{1}{(1-2\epsilon(\delta))n}$ and thus $w \in \Delta_{n,2\epsilon(\delta)}$. The constraint with B_i implies that $\bar{\lambda}_2(Y_{w'}) \leq$
505 $4 + O(\delta)$, and after scaling we have $\bar{\lambda}_2(Y_{\hat{w}}) \leq \frac{4+O(\delta)}{1-\epsilon(\delta)} \leq 4 + O(\delta)$ since $\epsilon(\delta) \leq \delta < \frac{1}{2}$. The success
506 probability is at least $1 - \tau = 0.99$.

507 We can write $A_i = C_i C_i^\top$ with $C_i = \sqrt{(1-\epsilon)n} e_i$, and $B_i = D_i D_i^\top$ with $D_i = \sqrt{y_i} a_i$. It takes ²⁰
508 $O(1)$ time to perform a matrix-vector product with C_i , and $O(d)$ time with D_i . Therefore, $t_C + t_D$
509 in Lemma 2.1 is $O(nd)$, and the runtime of Lemma 2.1 is $\tilde{O}(nd \text{ poly}(1/\epsilon))$ for $\tau = 0.01$. $\square$

510 We now proceed to prove Lemma 4.1. ²¹

511 *Proof of Lemma 4.1.* Without loss of generality, we can assume that $y_i \geq 0$ for all $i \in [n]$. Because y_i 1
 512 should be $\langle a_i, x \rangle^2$, any sample with $y_i < 0$ must be corrupted and can be discarded. By Lemma A.2,
 513 we can compute $\hat{w} \in \Delta_{n, 2\epsilon(\delta)}$ such that $\lambda_1(Y_{\hat{w}}) + \lambda_2(Y_{\hat{w}}) \leq 4 + O(\delta)$.

514 By Lemma A.1 and $y_i \geq 0$, we have 2

$$Y_{\hat{w}} = \sum_{i \in S} \hat{w}_i y_i a_i a_i^\top \succeq \sum_{i \in G} \hat{w}_i y_i a_i a_i^\top = Y_{G, \hat{w}} \succeq (1 - \delta)I + 2xx^\top, \quad 3$$

515 which gives a lower bound $x^\top Y_{\hat{w}} x \geq 3 - O(\delta)$, as well as lower bounds on the eigenvalues of $Y_{\hat{w}}$: 4

$$\lambda_1(Y_{\hat{w}}) \geq 3 - O(\delta) \quad \text{and} \quad \lambda_2(Y_{\hat{w}}) \geq 1 - O(\delta). \quad 5$$

516 Putting the upper and lower bounds together, we obtain 6

$$\begin{aligned} \lambda_1(Y_{\hat{w}}) &= \bar{\lambda}_2(Y_{\hat{w}}) - \lambda_2(Y_{\hat{w}}) \leq 4 + O(\delta) - (1 - O(\delta)) \leq 3 + O(\delta), \quad \text{and} \\ \lambda_2(Y_{\hat{w}}) &= \bar{\lambda}_2(Y_{\hat{w}}) - \lambda_1(Y_{\hat{w}}) \leq 4 + O(\delta) - (3 - O(\delta)) \leq 1 + O(\delta). \end{aligned} \quad 7$$

517 Lemma A.1 holds with probability at least 0.99. Assuming Lemma A.1 holds, Lemma 2.1 succeeds 8
 518 with probability at least 0.99, so the success probability of Lemma 4.1 is at least 0.98.

519 For the initialization step, it suffices to use Lemma A.2 to find a $(1 - \epsilon(\delta))$ -optimal solution $\hat{w} \in$ 9
 520 $\Delta_{n, 2\epsilon(\delta)}$ to the SDP (*), and the runtime to compute such $\hat{w}$ is $\tilde{O}(nd \text{poly}(1/\epsilon(\delta)))$. $\square$

521 A.2 Proof of Lemma A.1 10

522 This section is devoted to the proof of Lemma A.1. We will use the following concentration results. 11

523 **Lemma A.3** ([4], Section A.4.2). *Let $x \in \mathbb{R}^d$. For any $\delta > 0$, there exists a constant $C(\delta) > 0$ such* 12
 524 *that when $n > C(\delta) \cdot d \log d$ and we are given a set of n samples $\{(a_i, y_i)\}_{i=1}^n$ with $a_i \sim \mathcal{N}(0, I)$*
 525 *independently and $y_i = \langle a_i, x \rangle^2$ for all $i \in [n]$, then with probability at least 0.99, it holds*

$$\left\| \frac{1}{n} \sum_{i=1}^n y_i a_i a_i^\top - (I + 2xx^\top) \right\|_2 \leq \delta. \quad 13$$

526 **Proposition A.4.** *Let $\alpha \in (0, 2/e)$. Let $X_1, \dots, X_m$ be m random variables drawn i.i.d. from* 14
 527 $\mathcal{N}(0, 1)$. Define

$$H = \{i \in [m] : |X_i| \geq 4 \ln^2(2/\alpha)\}. \quad 15$$

528 *With probability at least $1 - 10^{-3}$, the following are all true:* 16

$$\begin{aligned} (a) \quad & |H| \leq O(m\alpha), \\ (b) \quad & \sum_{i \in H} X_i^4 = O(m\alpha \log^8(2/\alpha)), \\ (c) \quad & \max_{i \in H} X_i^2 = O(\log m). \end{aligned} \quad 17$$

529 *Proof.* For $X \sim \mathcal{N}(0, 1)$ and $t > 0$, it holds that $\Pr[|X| \geq t] \leq 2 \exp(-t^2/2)$. By setting 18
 530 $t = 4 \ln^2(2/\alpha)$, we have $\Pr[X^4 \geq t] \leq \alpha$. Let $Y_i \in \{0, 1\}$ be the indicator random variable
 531 for the event “ $i \in H$ ”, so that $|H| = \sum_{i \in H} Y_i$.

532 Because $\mathbb{E}[\sum_{i \in H} Y_i] \leq m\alpha$, by Markov’s inequality, we have $|H| = O(m\alpha)$ with probability at 19
 533 least $1 - 10^{-3}/3$. We assume this holds for the rest of the proof.

534 By the principle of deferred decisions, an equivalent way to draw $X_1, \dots, X_m$ from $\mathcal{N}(0, 1)$ is to 20
 535 first draw Y_i , and then draw X_i conditioned on the value of Y_i . Note that H is fixed after drawing
 536 $Y_1, \dots, Y_m$.

$$\mathbb{E} \left[\sum_{i \in H} X_i^4 \right] = \sum_{i \in H} \mathbb{E}[X_i^4 \mid |X_i| \geq 4 \ln^2(2/\alpha)] \leq O(m\alpha) \mathbb{E}[X_i^4 \mid |X_i| \geq 4 \ln^2(2/\alpha)]. \quad 21 \quad (4)$$

537 For any $t \geq 4$, by the definition of conditional expectation, and using the fact that X_i is normally 1
 538 distributed:

$$\mathbb{E}[X_i^4 \mid |X_i| \geq t] = \frac{\frac{2}{\sqrt{2\pi}} \int_t^\infty x^4 e^{-x^2/2} dx}{\frac{2}{\sqrt{2\pi}} \int_t^\infty e^{-x^2/2} dx} = \frac{\int_t^\infty x^4 e^{-x^2/2} dx}{\int_t^\infty e^{-x^2/2} dx} \leq 2t^4. \quad 2 \quad (5)$$

539 We use Inequality (5) to upper bound (4): 3

$$\mathbb{E}\left[\sum_{i \in H} X_i^4\right] \leq O(m\alpha) (2 \cdot (4 \ln^2(2/\alpha))^4) = O(m\alpha \log^8(2/\alpha)). \quad 4$$

540 By Markov's inequality, with probability at least $1 - 10^{-3}/3$, we have $\mathbb{E}[\sum_{i \in H} X_i^4] =$ 5
 541 $O(m\alpha \log^8(2/\alpha))$. Finally, since $X_1, \dots, X_m$ are drawn i.i.d. from $\mathcal{N}(0, 1)$, we have that
 542 $\max_{i \in [m]} |X_i| = O(\sqrt{\log m})$ with probability at least $1 - 10^{-3}/3$. $\square$

543 **Proposition A.5.** Let $K_1 \geq 0$ and $K_2 \geq 0$. Let $a_1, \dots, a_m \geq 0$ such that $\max_{i \in [m]} a_i^2 \leq K_1$ and 6
 544 $\sum_{i \in [m]} a_i^4 \leq K_2$. Let $X_1, \dots, X_m$ be m random variables drawn i.i.d. from $\mathcal{N}(0, 1)$. Then, with
 545 probability at least $1 - 10^{-3}12^{-d}$,

$$\sum_{i \in [m]} a_i^2 X_i^2 = O(\sqrt{dK_2} + dK_1). \quad 7$$

546 *Proof.* Since $X_i \sim \mathcal{N}(0, 1)$, the random variable X_i^2 is sub-exponential. Applying Bernstein's 8
 547 inequality for sub-exponential random variable [46, Theorem 2.8.2], for every $t \geq 0$,

$$\Pr\left[\sum_{i \in [m]} a_i X_i^2 \geq t\right] \leq \exp\left(-c \min\left(\frac{t^2}{\sum_{i \in [m]} a_i^4}, \frac{t}{\max_i a_i^2}\right)\right), \quad 9 \quad (6)$$

548 where $c > 0$ is a universal constant. 10

549 We need to choose a value of t such that the right-hand side of (6) is upper bounded by $10^{-3}12^{-d}$. 11
 550 Given our assumptions on $a_1, \dots, a_m$, it is sufficient to choose t such that $t = \Omega(\sqrt{dK_2})$ and
 551 $t = \Omega(dK_1)$. $\square$

552 **Proposition A.6.** Let $\alpha \in (0, 1)$. Let $X_1, \dots, X_m$ be m random variables drawn i.i.d. from $\mathcal{N}(0, 1)$. 12
 553 With probability at least $1 - 10^{-3}12^{-d}$, it holds that:

$$\max_{L \subseteq [m]: |L| = \alpha m} \sum_{i \in L} X_i^2 = O(m\alpha \log(1/\alpha) + d). \quad 13$$

554 *Proof.* We define the threshold function 14

$$h_r(z) = \begin{cases} 0, & z \leq r \\ z, & z > r \end{cases} \quad 15$$

555 with $r = 8 \ln(1/\alpha)$. Since $z \leq r + h_r(z)$ for all $z > 0$, 16

$$\max_{L \subseteq [m], |L| = \alpha m} \sum_{i \in L} X_i^2 \leq \max_L \sum_{i \in L} r + \max_L \sum_{i \in L} h_r(X_i^2) \leq m\alpha \cdot r + \sum_{i \in [m]} h_r(X_i^2). \quad 17$$

556 We will use Chernoff-bound like arguments to obtain a high-probability upper bound on 18
 557 $\sum_{i \in [m]} h_r(X_i^2)$. For any $c > 0$ and $t > 0$, we have:

$$\begin{aligned} \Pr\left[\sum_{i \in [m]} h_r(X_i^2) \geq t\right] &= \Pr\left[\exp\left(c \sum_{i \in [m]} h_r(X_i^2)\right) \geq \exp(c \cdot t)\right] \quad 19 \\ &\leq e^{-ct} \mathbb{E}\left[\exp\left(c \sum_{i \in [m]} h_r(X_i^2)\right)\right] \\ &= e^{-ct} \prod_{i \in [m]} \mathbb{E}[\exp(c \cdot h_r(X_i^2))] , \end{aligned} \quad (7)$$

558 where the inequality follows from Markov's inequality. 1

559 Thus, it is sufficient to upper bound $\mathbb{E}[\exp(c \cdot h_r(X_i^2))]$. For any $c < 1/2$, we have 2

$$\begin{aligned} \mathbb{E}[\exp(c \cdot h_r(X_i^2))] &= 1 \cdot \Pr[h_r(X_i^2) = 0] + \frac{1}{\sqrt{2\pi}} \int_{\sqrt{r}}^{\infty} e^{cx^2} e^{-x^2/2} dx \quad 3 \\ &\leq 1 + \frac{1}{\sqrt{2\pi}\sqrt{1-2c}} \int_{\sqrt{r(1-2c)}}^{\infty} e^{-y^2/2} dy \\ &= 1 + \frac{1}{\sqrt{1-2c}} \Pr[X_i \geq \sqrt{r(1-2c)}] \\ &\leq 1 + \frac{1}{\sqrt{1-2c}} \exp(-r(\frac{1}{2} - c)) , \end{aligned}$$

560 where the second inequality is obtained by substituting $x\sqrt{1-2c} = y$ in the integral, and the 4
561 last inequality uses the Gaussian tail bound that $\Pr[X_i \geq z] \leq e^{-z^2/2}$ for all $z > 0$. Recall that
562 $r = 8 \ln(1/\alpha)$. We set $c = 1/4$, so that $\exp(-r/4) = \alpha^2$. Thus, we have that:

$$\mathbb{E}\left[\exp\left(\frac{1}{4} \cdot h_r(X_i^2)\right)\right] \leq 1 + \sqrt{2}\alpha^2 \leq e^{\sqrt{2}\alpha^2} \quad 5 \quad (8)$$

563 We substitute the upper bound (8) into (7) and obtain that 6

$$\Pr\left[\sum_{i \in [m]} h_r(X_i^2) \geq t\right] \leq \exp\left(\frac{-t}{4} + \sqrt{2}m\alpha^2\right). \quad 7$$

564 We can choose $t = 4\sqrt{2}m\alpha^2 + \Omega(d)$, and then conclude that with probability at least $1 - 10^{-3}12^{-d}$, 8

$$\max_{L \subseteq [m]: |L|=\alpha m} \leq m\alpha r + \sum_{i \in [m]} h_r(X_i^2) = O(m\alpha r + m\alpha^2 + d) = O(m\alpha \log(1/\alpha) + d). \quad 9$$

565 10 □

566 **Lemma A.1.** Let $x \in \mathbb{R}^d$ be the ground-truth vector. Fix $\delta > 0$. There exists constants $\bar{\epsilon}(\delta)$ and 11
567 $c(\delta)$ such that if we let $0 < \epsilon \leq \bar{\epsilon}(\delta)$, and we are given an ϵ -corrupted set of $n = \tilde{\Omega}(c(\delta)d)$ samples
568 $(a_i, y_i)_{i \in [n]}$, with probability at least 0.99, for all $w \in \Delta_{n,2\epsilon}$,

$$\|Y_{G,w} - (I + 2xx^\top)\|_2 \leq \delta, \quad 12$$

569 where G is the set of indices of the remaining good samples and $Y_{G,w} = \sum_{i \in G} w_i y_i a_i a_i^\top$. 13

570 *Proof of Lemma A.1.* We recall the definition of $Y_{G,w} = \sum_{i \in G} w_i y_i a_i a_i^\top$. Let $\ell \leq \epsilon \cdot n$ and let 14
571 $\{(a_{n+i}, y_{n+i})\}_{i=1}^\ell$ be the set of samples that were removed by the ϵ -corruption adversary. Let
572 $G' = G \cup \{n+1, \dots, n+\ell\}$, $n' = n + \ell$, and $\epsilon' = \epsilon/(1 + \epsilon)$. Note that without loss of generality,
573 we can assume that $|G| = (1 - \epsilon)n$ and $|G'| = (1 - \epsilon')n' = n$.

574 We define a mapping $\sigma : \Delta_{n,2\epsilon} \rightarrow \Delta_{n',3\epsilon'}$ such that 15

$$\sigma(w)_i = \begin{cases} w_i & i \in [n] \\ 0 & \text{otherwise} \end{cases}. \quad 16 \quad (9)$$

575 In other words, all the weights are the same for the samples with index in the set $[n]$, and are equal to 17
576 0 for the samples removed by the adversary. We can verify that $\sigma(w) \in \Delta_{n',3\epsilon'}$ for all $w \in \Delta_{n,2\epsilon}$
577 since $\sigma(w)_i \leq w_i \leq 1/(1 - 2\epsilon)n = 1/(1 - 3\epsilon')n'$ for all $i \in [n']$, and $\|\sigma(w)\|_1 = \|w\|_1 = 1$.
578 Furthermore, we have $Y_{G,w} = Y_{G',\sigma(w)}$ for all $w \in \Delta_{n,2\epsilon}$. We denote with $w^* \in \Delta_{n',3\epsilon'}$ the desired
579 uniform weighting of the samples with index in G' , i.e., $w_i^* = \frac{1}{(1-\epsilon')n'} \mathbb{1}_{i \in G'}$.

580 By triangle inequality, for any $w \in \Delta_{n,2\epsilon}$, it holds 18

$$\|Y_{G',\sigma(w)} - (I + 2xx^\top)\|_2 \leq \|Y_{G',w^*} - (I + 2xx^\top)\|_2 + \|Y_{G',\sigma(w)-w^*}\|_2, \quad 19 \quad (10)$$

Thus, it suffices to show both $\|Y_{G',w^*} - (I + 2xx^\top)\|_2 \leq \delta/2$ and $\|Y_{G',\sigma(w)} - w^*\|_2 \leq \delta/2$. We upper bound the first term. By using the definition of w^* , note that

$$\begin{aligned} \|Y_{G',\sigma(w)} - (I + 2xx^\top)\|_2 &= \left\| \sum_{i \in G'} w_i^* y_i a_i a_i^\top - (I + 2xx^\top) \right\|_2 \\ &= \left\| \sum_{i \in G'} \frac{1}{|G'|} y_i a_i a_i^\top - (I + 2xx^\top) \right\|_2. \end{aligned}$$

Since $\mathbb{E}[y_i a_i a_i^\top] = I + 2xx^\top$ for any $i \in G'$, we can use a concentration inequality to upper bound this term. By Lemma A.3, as long as $n \geq C(\delta/2) \cdot d \log d$, with probability at least 0.995, we have

$$\|Y_{G',w^*} - (I + 2xx^\top)\|_2 \leq \delta/2. \quad (11)$$

It remains to show a high-probability upper bound to the second term $\|Y_{G',w^* - \sigma(w)}\|_2 \leq \delta/2$ that holds for any $w \in \Delta_{n,2\epsilon}$. To achieve this goal, we will provide a high-probability upper bound to the following quantity:

$$J = \sup_{w \in \Delta_{n,2\epsilon}} \|Y_{G',w^* - \sigma(w)}\|_2 = \sup_{w \in \Delta_{n,2\epsilon}} \left\| \sum_{i \in G'} (w_i^* - \sigma(w)_i) y_i a_i a_i^\top \right\|_2.$$

Note that for every i , the matrix $y_i a_i a_i^\top$ is positive semidefinite since $y_i \geq 0$. Thus, it holds that:

$$J \leq \sup_{w \in \Delta_{n,2\epsilon}} \left\| \sum_{i \in G'} |w_i^* - \sigma(w)_i| y_i a_i a_i^\top \right\|_2.$$

For any $w \in \Delta_{n,2\epsilon}$ and any $i \in [n']$, it is easy to see that $0 \leq |w_i^* - \sigma(w)_i| \leq \frac{1}{(1-2\epsilon)n}$. Additionally the weighting w^* and $\sigma(w)$ cannot be too different. In particular, we can show the following upper bound:

$$\sum_{i \in G'} |w_i^* - \sigma(w)_i| \leq \sum_{i=1}^{n'} |w_i^* - \sigma(w)_i| \leq \sup_{w, w' \in \Delta_{n',3\epsilon'}} \sum_{i=1}^{n'} |w_i - w'_i|. \quad (12)$$

We observe that $\Delta_{n',3\epsilon'}$ can be seen as the convex combination of all possible uniform weighting over subsets of $n'(1-3\epsilon')$ samples. Thus, the maximum distance will be between two points of the convex hull, and we can upper bound (12) as:

$$\sum_{i \in G'} |w_i^* - \sigma(w)_i| \leq \sup_{w, w' \in \Delta_{n',3\epsilon'}} \sum_{i=1}^{n'} |w_i - w'_i| \leq \frac{6\epsilon' n}{n'(1-3\epsilon')} \leq 6\epsilon. \quad (13)$$

Consider the family of weights defined as $\Gamma = \left\{ \beta \in \mathbb{R}^{n'} : \sum_i \beta_i \leq 6\epsilon \text{ and } 0 \leq \beta_i \leq \frac{1}{(1-2\epsilon)n} \right\}$. By the discussion above, we have that

$$J \leq \sup_{w \in \Delta_{n,2\epsilon}} \left\| \sum_{i \in G'} |w_i^* - \sigma(w)_i| y_i a_i a_i^\top \right\|_2 \leq \sup_{\beta \in \Gamma} \left\| \sum_{i \in G'} \beta_i y_i a_i a_i^\top \right\|_2. \quad (14)$$

Since the map $\beta \mapsto \left\| \sum_{i \in G'} \beta_i y_i a_i a_i^\top \right\|_2$ is convex with respect to β , and Γ is a convex set, the supremum over $\beta \in \Gamma$ of (14) is achieved at one of the extreme points of Γ . Thus, it holds:

$$J \leq \sup_{\beta \in \Gamma} \left\| \sum_{i \in G'} \beta_i y_i a_i a_i^\top \right\|_2 \leq \frac{1}{(1-2\epsilon)n} \max_{L \subseteq G', |L|=6\epsilon n} \left\| \sum_{i \in L} y_i a_i a_i^\top \right\|_2.$$

For any vector $v \in \mathbb{S}^{d-1}$ in the unit sphere, let

$$J(v) = \max_{L \subseteq G', |L|=6\epsilon n} \sum_{i \in L} y_i (v_i^\top a_i)^2,$$

and note that $J \leq \frac{1}{(1-2\epsilon)n} \sup_{v \in \mathbb{S}^{d-1}} J(v)$.

Without loss of generality, assume that $x = e_1$, where $e_1 \in \mathbb{R}^d$ is the first canonical vector. Given any vector $u \in \mathbb{R}^d$, we will denote with u_1 the first coordinate, and with $\tilde{u} \in \mathbb{R}^{d-1}$ the remaining $d-1$ coordinates, i.e., $u = (u_1, \tilde{u})$. Assuming $x = e_1$, we have that $y_i = (x^\top a_i)^2 = a_{i,1}^2$ for any $i \in G'$. Let $H = \{i \in G' : |a_{i,1}| \geq 4 \ln^2(2/\epsilon)\}$. We consider a set $L \subseteq G'$ that always contains H and then picks $6\epsilon n$ additional elements. In particular, it holds:

$$J(v) = \max_{L \subseteq G', |L|=6\epsilon n} \sum_{i \in L} y_i (a_i^\top v_i)^2 \leq \left[\sum_{i \in H} y_i (a_i^\top v_i)^2 + \max_{L \subseteq G' \setminus H, |L|=6\epsilon n} \sum_{i \in L} y_i (a_i^\top v_i)^2 \right]. \quad (15)$$

The first term of the right-hand side of (15) can be rewritten as follows:

$$\sum_{i \in H} y_i (a_i^\top v)^2 = \sum_{i \in H} a_{i,1}^2 \left(\sum_{j=1}^d a_{i,j} v_j \right)^2 \leq 2 \sum_{i \in H} [a_{i,1}^4 v_{i,1}^2 + a_{i,1}^2 (\tilde{a}_i^\top \tilde{v})^2]. \quad (16)$$

For the second term of the right-hand side of (15), note that for any $i \in G' \setminus H$, it holds that $y_i = a_{i,1}^2 < 16 \ln^4(2/\epsilon)$ due to the definition of H . Thus, we have that:

$$\max_{L \subseteq G' \setminus H, |L|=6\epsilon n} \sum_{i \in L} y_i (a_i^\top v_i)^2 \leq 16 \ln^4(2/\epsilon) \max_{L \subseteq G' \setminus H, |L|=6\epsilon n} \sum_{i \in L} (a_i^\top v_i)^2 \quad (17)$$

$$\leq 32 \ln^4(2/\epsilon) \max_{L \subseteq G' \setminus H, |L|=6\epsilon n} \sum_{i \in L} [a_i^2 v_1^2 + (\tilde{v}^\top \tilde{a}_i)^2]. \quad (18)$$

Also, using the definition of H , note that for any $L \subseteq G' \setminus H$ with $|L| = 6\epsilon n$, we have that:

$$\sum_{i \in L} a_{i,1}^2 v_1^2 \leq \sum_{i \in L} a_{i,1}^2 = O(n \epsilon \ln^4(2/\epsilon)). \quad (19)$$

By combining (16), (17), and (19) with (15), we obtain that:

$$J(v) = O \left(n \epsilon \log^8(1/\epsilon) + \sum_{i \in H} a_{i,1}^4 + \sum_{i \in H} a_{i,1}^2 (\tilde{a}_i^\top \tilde{v})^2 + \max_{L \subseteq G' \setminus H, |L|=6\epsilon n} \sum_{i \in L} (\tilde{v}^\top \tilde{a}_i)^2 \right). \quad (20)$$

Let E_1 be the event of Proposition A.4 with $\alpha = \epsilon$ and $n = m$. That is, with probability at least $1 - 10^{-3}$, we have that $|H| = O(\epsilon n)$, $\sum_{i \in H} a_{i,1}^4 = O(n \epsilon \log^8(1/\epsilon))$ and $\max_i a_{i,1}^2 = O(\log n)$. For the remaining of this proof, assume that E_1 is true, and thus:

$$J(v) = O \left(n \epsilon \log^8(1/\epsilon) + \sum_{i \in H} a_{i,1}^2 (\tilde{a}_i^\top \tilde{v})^2 + \max_{L \subseteq G' \setminus H, |L|=6\epsilon n} \sum_{i \in L} (\tilde{v}^\top \tilde{a}_i)^2 \right). \quad (20)$$

Denote with $\bar{J}(v) = \sum_{i \in H} (\tilde{a}_i^\top \tilde{v})^2 + \max_{L \subseteq G' \setminus H, |L|=6\epsilon n} \sum_{i \in L} (\tilde{v}^\top \tilde{a}_i)^2$ the last two terms of the right-hand side of (20). We will upper bound each term of $\bar{J}$ individually. First, observe that the random variables $Z_i = \tilde{a}_i^\top \tilde{v} / \|\tilde{v}\|_2$ for $i \in G'$ are independent standard normal random variables. Let E_2^v be the event of Proposition A.5 for the random variables $\{Z_i : i \in H\}$ and weights $\{a_{i,1}^2 : i \in H\}$. That is, with probability at least $1 - 10^{-3} 12^{-d}$, it holds that $\sum_{i \in H} a_{i,1}^2 Z_i^2 = O(\epsilon^{1/2} \sqrt{dn} \log^4(1/\epsilon) + d \log n)$. For the second term of $\bar{J}$, we can invoke Proposition A.6 over the random variables $\{Z_i : i \in G' \setminus H\}$ with $\alpha = 12\epsilon$ and $m = |G' \setminus H| \geq n/2$ (if $\epsilon < 1/2$). Let E_3^v be the event of this proposition, that is, with probability at least $1 - 10^{-3} 12^{-d}$, it holds that:

$$\max_{L \subseteq G' \setminus H, |L|=6\epsilon n} \sum_{i \in L} (\tilde{v}^\top \tilde{a}_i)^2 = O(\epsilon n \log(1/\epsilon) + d). \quad (21)$$

By taking a union bound of the events E_2^v and E_3^v , we have that:

$$\bar{J}(v) = O \left(\epsilon n \log(1/\epsilon) + \epsilon^{1/2} \sqrt{dn} \log^4(1/\epsilon) + d \log n \right). \quad (21)$$

623 Consider a $1/4$ -net $\mathcal{N}$ of $\mathbb{S}^{d-1}$, where $|\mathcal{N}| = O(12^d)$. Note that for any $v \in \mathbb{S}^{d-1}$, it holds 1
 624 $\sup_{v \in \mathbb{S}^{d-1}} \bar{J}(v) \leq 2 \sup_{v \in \mathcal{N}} \bar{J}(v)$. Thus, by taking a union bound over the event described in (21)
 625 for all $v \in \mathcal{N}$, we obtain that with probability at least $1 - 2 \cdot 10^{-3}$, we have:

$$\sup_{v \in \mathbb{S}^{d-1}} \bar{J}(v) = O\left(\epsilon n \log(1/\epsilon) + d \log n + \sqrt{dn\epsilon} \log^2(1/\epsilon)\right) \quad 2 \quad (22)$$

626 We finally combine (22) and (20) to conclude that with probability at least $1 - 0.995$, it holds that 3

$$J \leq O\left(\frac{1}{n} \left[d \log n + \epsilon^{1/2} \sqrt{dn} \log^4(1/\epsilon) + n \sqrt{\epsilon} \log^8(1/\epsilon) \right]\right) \quad 4$$

627 We can pick a sufficiently small ϵ (depending only on δ) and $n \geq c(\delta) d \log d$ for a sufficiently large 5
 628 constant $c(\delta)$ so that with probability at least $1 - 0.995$, it holds that $J \leq \delta/2$. Utilizing this result
 629 along with (11) to upper bound (10) yields the desired statement. $\square$

630 B Omitted Proofs in Section 5 6

631 **Lemma 5.1.** Let $x \in \mathbb{R}^d$ be the ground-truth vector. Let $\mathcal{D}_z$ be the distribution of gradients at z as 7
 632 defined in Equation (2). For any z with $\|z - x\|_2 \leq 1$, the covariance matrix Σ_z of $\mathcal{D}_z$ satisfies

$$\Sigma_z \preceq O\left(\|z - x\|_2^2\right) I. \quad 8$$

633 *Proof of Lemma 5.1.* Recall that $g \sim \mathcal{D}_z$ is distributed as 9

$$g = \frac{\partial}{\partial z} \left[\left(\langle a, z \rangle^2 - \langle a, x \rangle^2 \right)^2 \right] = 4 \left(\langle a, z \rangle^2 - \langle a, x \rangle^2 \right) \langle a, z \rangle a \quad \text{where } a \sim \mathcal{N}(0, I). \quad 10$$

634 Let $\mu_z = \mathbb{E}_{g \sim \mathcal{D}_z} [g]$. We have 11

$$0 \preceq \Sigma_z = \mathbb{E}_{g \sim \mathcal{D}_z} [gg^\top] - \mu_z \mu_z^\top \preceq \mathbb{E}_{g \sim \mathcal{D}_z} [gg^\top]. \quad 12$$

635 Consequently, it suffices to upper bound the spectral norm of $\mathbb{E}_{g \sim \mathcal{D}_z} [gg^\top]$. Let $h = z - x$. 13

$$\begin{aligned} \|\Sigma_z\|_2 &\leq \left\| \mathbb{E}_{g \sim \mathcal{D}_z} [gg^\top] \right\|_2 & 14 \\ &= \max_{\|v\|_2=1} v^\top \mathbb{E}_{g \sim \mathcal{D}_z} [gg^\top] v \\ &= \max_{\|v\|_2=1} \mathbb{E}_{g \sim \mathcal{D}_z} [\langle g, v \rangle^2] \\ &= 16 \max_{\|v\|_2=1} \mathbb{E}_{a \sim \mathcal{N}(0, I)} \left[\left((\langle a, z \rangle^2 - \langle a, x \rangle^2) \langle a, z \rangle \langle a, v \rangle \right)^2 \right] \\ &= 16 \max_{\|v\|_2=1} \mathbb{E}_{a \sim \mathcal{N}(0, I)} \left[\left((\langle a, h \rangle^2 + 2 \langle a, x \rangle \langle a, h \rangle) \langle a, x + h \rangle \langle a, v \rangle \right)^2 \right] \\ &= 16 \max_{\|v\|_2=1} \mathbb{E}_{a \sim \mathcal{N}(0, I)} \left[\langle a, h \rangle^2 \langle a, 2x + h \rangle^2 \langle a, x + h \rangle^2 \langle a, v \rangle^2 \right] \\ &\leq 16 \max_{\|v\|_2=1} \left(\mathbb{E}_a [\langle a, h \rangle^8] \mathbb{E}_a [\langle a, 2x + h \rangle^8] \mathbb{E}_a [\langle a, x + h \rangle^8] \mathbb{E}_a [\langle a, v \rangle^8] \right)^{1/4} \\ &= 16 \max_{\|v\|_2=1} \left(105^4 \|h\|_2^8 \|2x + h\|_2^8 \|x + h\|_2^8 \|v\|_2^8 \right)^{1/4} \\ &= (16 \cdot 105) \|h\|_2^2 \|2x + h\|_2^2 \|x + h\|_2^2 \\ &= O(\|h\|_2^2). \end{aligned}$$

636 The last inequality follows from the Cauchy-Schwarz inequality. The last step uses the fact that 15
 637 $\|2x + h\|_2 = O(1)$ and $\|x + h\|_2 = O(1)$, which follows from $\|x\|_2 = 1$ and $\|h\|_2 \leq 1$. $\square$

638 **Lemma 5.3.** Let $x \in \mathbb{R}^d$ be the ground-truth vector. Suppose at iteration t of Algorithm 2, the
 639 current solution z_t satisfies $\|z_t - x\|_2 \leq \frac{1}{8}$. Let μ_{z_t} denote the expected true gradient at z_t defined
 640 in Equation (3). Suppose the estimated gradient $\hat{\mu}_{z_t} \in \mathbb{R}^d$ satisfies $\|\hat{\mu}_{z_t} - \mu_{z_t}\|_2 \leq c \|z_t - x\|_2$ for
 641 $c = 4$. Then, we have

$$\|z_{t+1} - x\|_2^2 \leq 0.99 \|z_t - x\|_2^2. \quad 2$$

642 *Proof of Lemma 5.3.* Recall that $g \sim \mathcal{D}_z$ is distributed as 3

$$g = \frac{\partial}{\partial z} \left[\left(\langle a, z \rangle^2 - \langle a, x \rangle^2 \right)^2 \right] = 4 \left(\langle a, z \rangle^2 - \langle a, x \rangle^2 \right) \langle a, z \rangle a \quad \text{where } a \sim \mathcal{N}(0, I). \quad 4$$

643 We can compute the mean μ_z of $\mathcal{D}_z$ using moments of Gaussian: 5

$$\mu_z = \mathbb{E}_{g \sim \mathcal{D}_z} [g] = \left(12 \|z\|_2^2 - 4 \|x\|_2^2 \right) z - 8 \langle x, z \rangle x. \quad 6$$

644 Consider one step of gradient descent in Algorithm 2: $z_{t+1} = z_t - \eta \hat{\mu}_{z_t}$, where $\hat{\mu}_{z_t}$ is close to μ_z . 7
 645 We have

$$\|z_{t+1} - x\|_2^2 = \|z_t - \eta \hat{\mu}_{z_t} - x\|_2^2 = \|z_t - x\|_2^2 - 2\eta \langle \hat{\mu}_{z_t}, z_t - x \rangle + \eta^2 \langle \hat{\mu}_{z_t}, \hat{\mu}_{z_t} \rangle \quad 8$$

646 To prove convergence, we need to lower bound $\langle \hat{\mu}_{z_t}, z_t - x \rangle$ and upper bound $\langle \hat{\mu}_{z_t}, \hat{\mu}_{z_t} \rangle$. 9

647 Let $z = z_t$ and $h = z - x$. Substituting $z = x + h$ in the expression for μ_z , we get: 10

$$\begin{aligned} \mu_z &= \left(12 \|x + h\|_2^2 - 4 \|x\|_2^2 \right) (x + h) - 8 \langle x, x + h \rangle x \\ &= \left(16 \langle x, h \rangle + 12 \|h\|_2^2 \right) x + \left(8 \|x\|_2^2 + 24 \langle x, h \rangle + 12 \|h\|_2^2 \right) h. \end{aligned} \quad 11$$

648 We will be using the assumptions of this lemma: $\|x\|_2 = 1$, $\|h\|_2 \leq \frac{1}{8}$, and $\|\hat{\mu}_z - \mu_z\|_2 \leq c \|h\|_2$. 12

649 First we lower bound $\langle \hat{\mu}_z, h \rangle$.

$$\begin{aligned} \langle \hat{\mu}_z, h \rangle &= \langle \mu_z, h \rangle + \langle \hat{\mu}_z - \mu_z, h \rangle \\ &= 16 \langle x, h \rangle^2 + 36 \langle x, h \rangle \|h\|_2^2 + 8 \|x\|_2^2 \|h\|_2^2 + 12 \|h\|_2^4 + \langle \hat{\mu}_z - \mu_z, h \rangle \\ &\geq -\frac{81}{4} \|h\|_2^4 + 8 \|x\|_2^2 \|h\|_2^2 + 12 \|h\|_2^4 - c \|h\|_2^2 \\ &\geq \left(-\frac{81}{256} + 8 + \frac{12}{64} - c \right) \|h\|_2^2 \\ &\geq (7.5 - c) \|h\|_2^2. \end{aligned} \quad 13$$

650 The first inequality uses the fact that $16 \langle x, h \rangle^2 + 36 \langle x, h \rangle \|h\|_2^2$ is a quadratic function of $\langle x, h \rangle$, 14
 651 which has minimum value $-\frac{81}{4} \|h\|_2^4$ for all $\langle x, h \rangle \in \mathbb{R}$.

652 Next we upper bound $\|\hat{\mu}_z\|_2$ using the triangle inequality. 15

$$\begin{aligned} \|\hat{\mu}_z\|_2 &\leq \|\mu_z\|_2 + \|\hat{\mu}_z - \mu_z\|_2 \\ &\leq \left(16 \langle x, h \rangle + 12 \|h\|_2^2 \right) \|x\|_2 + \left(8 \|x\|_2^2 + 24 \langle x, h \rangle + 12 \|h\|_2^2 \right) \|h\|_2 + c \|h\|_2 \\ &\leq \left(16 + \frac{12}{8} + 8 + \frac{24}{8} + \frac{12}{64} + c \right) \|h\|_2 \\ &\leq (29 + c) \|h\|_2. \end{aligned} \quad 16$$

653 Putting everything together, we have 17

$$\begin{aligned} \|z_{t+1} - x\|_2^2 &= \|z_t - x\|_2^2 - 2\eta \langle \hat{\mu}_{z_t}, z_t - x \rangle + \eta^2 \langle \hat{\mu}_{z_t}, \hat{\mu}_{z_t} \rangle \\ &\leq \left[1 - 2(7.5 - c)\eta + (29 + c)^2 \eta^2 \right] \|z_t - x\|_2^2. \end{aligned} \quad 18$$

654 Choosing $c = 4$ and $\eta = 1/300$ gives that $\|z_{t+1} - x\|_2^2 \leq 0.99 \|z_t - x\|_2^2$. 19

□

655 C A Counter-Example for Previous Algorithms

656 In this work, we consider a general setting that allows adversarial corruption in both the measurement 1
 657 vectors a_i 's and the intensity measurements y_i 's. Prior work in robust phase retrieval has addressed a
 658 special case where the adversarial corruption is restricted to the y_i 's, still assuming that the measuring
 659 vectors a_i 's are independently sampled from the Gaussian distribution [27, 51]. In this section, we
 660 construct a counter-example demonstrating the failure of algorithms developed for the restricted
 661 corruption setting when applied to the more general setting considered in our paper.

662 The Median Truncated Wirtinger Flow Algorithm [51] is an algorithm to address the robust phase 2
 663 retrieval problem with adversarial corruption limited to the y_i 's. The algorithm first initializes $z^{(0)}$
 664 using the spectral method. Let $\alpha \geq 3$. In particular, $z^{(0)}$ is computed as the top eigenvector of
 665 the empirical matrix $Y := \frac{1}{m} \sum_{i=1}^m y_i a_i a_i^\top \mathbb{1}_{|y_i| \leq \alpha^2 \text{med}(\{y_i\}_{i=1}^m)}$ that only uses a truncated set of
 666 samples, where the threshold is determined by $\text{med}(\{y_i\}_{i=1}^m)$, the median over all y_i 's. The analysis
 667 of the algorithm relies on the fact that as long as the fraction of outliers is not too large and the sample
 668 complexity is large enough, the initialization is guaranteed to be within a small neighborhood of the
 669 ground truth.

670 We show that this initialization can fail to remove the distortion introduced by the adversarial 3
 671 if corruption is allowed for both a_i 's and y_i 's. Let $x \in \mathbb{S}^{d-1}$ be the ground truth unit vector.
 672 We construct an ϵ -corruption adversary that can manipulate the top eigenvector of the empirical
 673 covariance matrix $Y = \sum_{i=1}^n y_i a_i a_i^\top$, even when all y_i 's are accurately calculated as $y_i = (a_i^\top x)^2$.

674 Let $u \in \mathbb{S}^{d-1}$ be a unit vector such that $x^\top u = 0$. Suppose the adversary changes 1% of the 4
 675 a_i 's to $a_i = \sqrt{d-1/25} \cdot u + (1/5) \cdot x$, and suppose that all the y_i 's are accurate. In particular,
 676 the length of the corrupted a_i 's is comparable to the length of a random Gaussian vector, and
 677 the corresponding intensity measurements satisfy $y_i = (a_i^\top x)^2 = 1/25$. Let $z = (a^\top x)^2$ for a
 678 random vector $a \sim \mathcal{N}(0, I)$. By direct computation, note that $\Pr[z \geq 0.2] \geq 0.6$. Thus, with
 679 high-constant probability, the median-truncated initialization in [51] is not able to filter out any of
 680 those samples. However, after the adversarial corruption, the top eigenvector of $\mathbb{E}[\sum_{i=1}^n y_i a_i a_i^\top] \approx$
 681 $O(d)uu^\top + O(\sqrt{d})(ux^\top + xu^\top) + O(1)(I + 2xx^\top)$ will be manipulated to u , which is far from
 682 the ground truth x .

683 NeurIPS Paper Checklist

684 The checklist is designed to encourage best practices for responsible machine learning research, 1
685 addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove
686 the checklist: **The papers not including the checklist will be desk rejected.** The checklist should
687 follow the references and follow the (optional) supplemental material. The checklist does NOT count
688 towards the page limit.

689 Please read the checklist guidelines carefully for information on how to answer these questions. For 2
690 each question in the checklist:

- 691
- You should answer [Yes], [No], or [NA]. 3
 - [NA] means either that the question is Not Applicable for that particular paper or the
692 relevant information is Not Available.
 - Please provide a short (1–2 sentence) justification right after your answer (even for NA). 3
- 694

695 **The checklist answers are an integral part of your paper submission.** They are visible to the 4
696 reviewers, area chairs, senior area chairs, and ethics reviewers. You will be asked to also include it
697 (after eventual revisions) with the final version of your paper, and its final version will be published
698 with the paper.

699 The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. 5
700 While "[Yes]" is generally preferable to "[No]", it is perfectly acceptable to answer "[No]" provided a
701 proper justification is given (e.g., "error bars are not reported because it would be too computationally
702 expensive" or "we were unable to find the license for the dataset we used"). In general, answering
703 "[No]" or "[NA]" is not grounds for rejection. While the questions are phrased in a binary way, we
704 acknowledge that the true answer is often more nuanced, so please just use your best judgment and
705 write a justification to elaborate. All supporting evidence can appear either in the main paper or the
706 supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification
707 please point to the section(s) where related material for the question can be found.

708 **IMPORTANT, please:** 6

- 709
- **Delete this instruction block, but keep the section heading "NeurIPS Paper Checklist",** 7
 - **Keep the checklist subsection headings, questions/answers and guidelines below.**
 - **Do not modify the questions and only use the provided macros for your answers.**
- 711

712 1. Claims 8

713 Question: Do the main claims made in the abstract and introduction accurately reflect the 9
714 paper's contributions and scope?

715 Answer: [Yes] 10

716 Justification: We provide formal statements on the guarantee of our algorithm, and state our 11
717 main theorem. The rest of the paper is devoted in proving this theorem.

718 Guidelines: 12

- 719
- The answer NA means that the abstract and introduction do not include the claims 13
720 made in the paper.
 - The abstract and/or introduction should clearly state the claims made, including the
721 contributions made in the paper and important assumptions and limitations. A No or
722 NA answer to this question will not be perceived well by the reviewers.
 - The claims made should match theoretical and experimental results, and reflect how
723 much the results can be expected to generalize to other settings.
 - It is fine to include aspirational goals as motivation as long as it is clear that these goals
724 are not attained by the paper.
- 726

728 2. Limitations 14

729 Question: Does the paper discuss the limitations of the work performed by the authors? 15

730 Answer: [Yes] 16

Justification: Limitations and questions for future work are discussed in the conclusion. 1

Guidelines: 2

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper. 3
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs 4

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof? 5

Answer: [Yes] 6

Justification: Sections 3-5 and the first two sections of the appendix are devoted to the proof of our main result. 7

Guidelines: 8

- The answer NA means that the paper does not include theoretical results. 9
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility 10

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)? 11

Answer: [NA] 12

Justification: The paper does not contain experimental results. 13

Guidelines: 14

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code 2

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [NA] 4

Justification: The paper does not include experiments requiring code. 5

Guidelines: 6

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [NA]

Justification: The paper does not include experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [NA]

Justification: The paper does not include experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [NA]

Justification: The paper does not include experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.

- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics 2

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes] 4

Justification: The research follows the NeurIPS Code of Ethics. 5

Guidelines: 6

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts 8

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA] 10

Justification: There is no societal impact of the work performed. 11

Guidelines: 12

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards 14

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA] 16

Justification: The paper is theoretical and poses no such risks. 17

Guidelines: 18

- The answer NA means that the paper poses no such risks. 19

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets 2

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected? 3

Answer: [NA] 4

Justification: The paper is theoretical and does not use existing assets. 5

Guidelines: 6

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets 8

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets? 9

Answer: [NA] 10

Justification: The paper does not release new assets. 11

Guidelines: 12

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects 14

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)? 15

Answer: [NA] 16

Justification: The paper does not involve crowdsourcing experiments. 17

991
992
993
994
995
996
997
998
999
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028
1029
1030
1031

Guidelines: 1

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 2
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects 3

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained? 4

Answer: [NA] 5

Justification: The paper does not involve crowdsourcing nor research with human subjects. 6

Guidelines: 7

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 8
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage 9

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required. 10

Answer: [NA] 11

Justification: The core method development in this research does not involve LLMs as any important, original, or non-standard components. 12

Guidelines: 13

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components. 14
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.