
LayerGuard: Poisoning-Resilient Federated Learning via Layer-Wise Similarity Analysis

Anonymous Author(s)

Affiliation

Address

email

Abstract

In recent years, model poisoning attacks have gradually evolved from conventional global parameter manipulations to more stealthy and strategic Targeted Layer Poisoning (TLP) attacks. These attacks achieve high attack success rates by selectively poisoning only a subset of layers. However, most existing defenses rely on evaluation of the entire network and are thus ineffective against TLP attacks, posing new challenges to the security of Federated Learning (FL). In this paper, we propose **LayerGuard**, a comprehensive defense framework featuring dynamic detection and adaptive aggregation to protect FL against advanced model poisoning attacks. Diverging from traditional methods that analyze the entire network collectively, **LayerGuard** performs layer-wise similarity analysis to detect anomalous clients and adaptively identifies layers under attack based on the clustering behavior of malicious updates, facilitating more precise threat detection. Building on this, we introduce a joint weighting mechanism in the aggregation process, which evaluates each client’s credibility at the layer level from two complementary informational dimensions: inter-layer and intra-layer, balancing attack mitigation and benign contribution retention. Extensive experiments across various datasets and model architectures demonstrate that **LayerGuard** successfully reduces the average attack success rate of TLP attacks to around 5%. Moreover, when confronted with other advanced model poisoning attacks, **LayerGuard** consistently maintains global model accuracy—even under high poisoning rates and severe non-IID conditions—comparable to that of FedAvg under no-attack settings, marking a significant improvement over existing defenses.

1 Introduction

Background and Problem. Federated Learning (FL) has gained widespread adoption in privacy-sensitive domains such as healthcare[3] and finance[4], due to its ability to enable collaborative model training without sharing raw data[1, 2]. However, the decentralized nature of FL also makes it vulnerable to model poisoning attacks[11–18], where adversaries manipulate a subset of client updates to degrade global model performance, posing a serious threat to the reliability of FL systems. Recently, poisoning strategies have evolved from coarse-grained global parameter perturbations[12–18] to more stealthy and strategic forms of Targeted Layer Poisoning (TLP)[11]—where only specific layers of the model are maliciously altered, effectively bypassing existing defenses that rely on evaluation of the entire network[2, 5–10]. In parallel, more general and disruptive poisoning variants such as advanced untargeted attacks[12–15] have emerged, by exhibiting high similarity to benign updates in gradient space, thereby disguising malicious behavior and subtly interfering with the training process. Existing defense mechanisms[5–10] struggle to detect fine-grained layer-level anomalies, and often misclassify benign clients, ultimately harming the overall performance of the global model. These limitations highlight the urgent need for a more fine-grained and robust defense

framework capable of resisting TLP and other advanced poisoning strategies while preserving model utility.

Limitations of Previous Works. To defend against model poisoning attacks, researchers have proposed various defenses. These include defenses that leverage cross-client information, such as Krum, Trimmed Mean, Norm Bound, and FLAME[5–8], where the server compares statistical properties across clients—e.g., Euclidean distance, cosine similarity—to identify anomalous updates. Defense methods that utilize global information, such as FLTrust and FLDetector[9, 10], in which the server uses trusted gradients or reference models to determine whether a client behaves abnormally. Although progress has been made, existing defenses still face significant limitations against emerging attack strategies. First, these defenses rely on evaluation of the entire network, which fails to capture TLP attacks where only a subset of model layers is manipulated[11]. As a result, such attacks can bypass detection with minimal effort. Second, under highly non-IID data distributions, current defenses struggle to distinguish malicious updates from genuinely benign ones[26]. In such settings, benign clients may generate statistically deviant updates that are nonetheless critical for model performance. These challenges call for a new defense paradigm that can effectively detect localized anomalies while preserving benign diversity, enabling more robust and fine-grained protection in FL.

Our Work. To address the above limitations, we propose **LayerGuard**, a comprehensive defense framework featuring dynamic detection and adaptive aggregation. Diverging from traditional methods that evaluate client behavior at the whole-model level, **LayerGuard** performs layer-wise similarity analysis to detect anomalous clients and adaptively identifies layers likely to be compromised, based on the clustering behavior of malicious updates. Building on this analysis, we introduce a joint weighting mechanism in the aggregation process that evaluates the credibility of each client across individual layers from two complementary informational dimensions: *user-level weights* analyze inter-layer information, while *layer-specific weights* capture intra-layer behavior. Extensive experiments on diverse datasets and model architectures show that **LayerGuard** reduces the average attack success rate of TLP attacks from approximately 90% to around 5%. When confronted with other advanced model poisoning attacks, **LayerGuard** maintains global model accuracy under high poisoning rates and severe non-IID conditions, comparable to that of FedAvg in no-attack settings.

Contribution. The main contributions are:

- (a) We uncover a limitation of existing defenses: their coarse-grained evaluation based on the entire network fails to detect localized anomalies, rendering them ineffective against Targeted Layer Poisoning (TLP) attacks.
- (b) We propose **LayerGuard**, a novel defense framework that operates at a finer granularity by analyzing each layer individually to identify anomalous clients and adaptively detect layers under attack. This novel approach facilitates more precise threat detection.
- (c) We design a joint weighting mechanism for aggregation, which evaluates each client’s credibility at the layer level based on two complementary informational dimensions: inter-layer and intra-layer. This design enables precise suppression of malicious updates while retaining the contribution of benign ones.
- (d) We conduct extensive experiments on **LayerGuard** against TLP attacks, other advanced model poisoning attacks, and adaptive attacks.

2 Related Works

2.1 Model Poisoning Attacks

In FL, model poisoning attacks directly manipulate client gradients during training and pose greater threats than data poisoning[18, 13], which relies on altering local training data[19, 20]. Depending on their objectives, these attacks are typically categorized into untargeted attacks[12–15], which aim to degrade overall model performance, and targeted attacks[15–18], such as backdoor insertion[16, 17], which manipulate specific outputs. Untargeted poisoning is a more severe threat to model prediction performance than the targeted one in FL[14]. In this work, we primarily focus on untargeted poisoning attacks. For targeted attacks, we focus on a recently proposed backdoor attack based on TLP[11].

88 **Untargeted Attack.** We consider five representative advanced untargeted model poisoning attacks
 89 in our evaluation. LIE[12] perturbs the average of benign gradients with calibrated noise to subtly
 90 degrade model performance while evading detection. Fang[13] is an optimization-based method
 91 that manipulates gradient directions by solving for a global scaling coefficient λ . Min-Sum[14]
 92 constrains the sum of squared distances between malicious and benign gradients so that it remains
 93 within the maximum squared distance among benign updates, while Min-Max[14] instead limits
 94 the maximum distance, effectively camouflaging malicious updates among benign ones. MPAF[15]
 95 leverages momentum from historical gradients to craft stealthy and disruptive updates that are harder
 96 to detect.

97 2.2 Targeted Layer Poisoning Attacks

98 Targeted Layer Poisoning (TLP) attacks achieve high attack success rates by selectively poisoning
 99 only a subset of model layers. We refer to the poisoned layers as targeted layers, and the unaltered
 100 layers as non-targeted layers. The Layer-wise Poisoning (LP) attack[11], proposed by Zhuang et
 101 al., is a backdoor-based TLP attack. It identifies Backdoor Critical (BC) layers using Backdoor
 102 Success Rate (BSR) as the evaluation metric, and poisons only these targeted layers. Malicious clients
 103 optimize their local models through layer-wise analysis, continuously injecting backdoors across
 104 multiple communication rounds. The BC layers are not fixed and may vary from round to round,
 105 which increases the stealthiness of the attack while maintaining high BSR. Currently, no effective
 106 defense strategy exists to counter this type of attack.

107 3 Design of LayerGuard

108 3.1 Design Challenges

109 Our core idea is inspired by previous research[13, 14] that emphasizes the necessity of a certain
 110 degree of similarity among malicious updates in order to significantly compromise FL. However,
 111 under more complex and stealthy threat scenarios—such as Targeted Layer Poisoning (TLP) attacks
 112 and advanced untargeted attacks—designing defense mechanisms based on this principle presents the
 113 following technical challenges:

114 **C1-** In contrast to conventional poisoning attacks that target the entire model update, TLP attacks
 115 selectively poison specific layers. This raises the question: how can the similarity among malicious
 116 updates be effectively quantified when only partial-layer poisoning is involved?

117 **C2-** Given that TLP attacks affect only selected layers while leaving others largely benign[11], how
 118 can a defense mechanism suppress malicious updates in the targeted layers without disrupting the
 119 aggregation of benign updates in the non-targeted layers?

120 **C3-** Do benign updates sometimes exhibit high similarity similar to that of malicious ones? How can
 121 we ensure that such benign updates are not mistakenly classified as malicious?

122 **C4-** How can malicious updates be effectively mitigated while minimizing the impact on benign
 123 contributions?

124 3.2 Overview

125 To address the aforementioned challenges, we propose a novel and advanced defense mechanism for
 126 FL, termed **LayerGuard**. The design motivations and inspirations behind this method are detailed
 127 in Appendix A. The overall architecture of **LayerGuard** is depicted in Figure 1. Specifically, the
 128 process consists of the following core components:

129 **(1) Anomalous User Identification.** Identify highly similar anomalous users per layer based on layer-
 130 wise similarity scores(LCSS). **(2) High-Risk Layer Detection.** By examining the distribution of
 131 anomalous users across different layers, **LayerGuard** identifies the high-risk layers where malicious
 132 activity is primarily concentrated. **(3) User-Level Weight Calculation.** Given the identified high-
 133 risk layers, **LayerGuard** assigns a user-level weight to each client. **(4) Layer Update Boundary**
 134 **Definition.** Adaptively establish benign and malicious boundaries for updates on individual layers.
 135 **(5) Layer-Specific Weight Calculation.** Based on layer update boundaries, LayerGuard dynamically




Figure 1: Design of LayerGuard Framework.

136 assigns layer-specific weights to each client at different layers. **(6) Global Aggregation.** Finally,
 137 **LayerGuard** integrates both user-level and layer-specific weights to perform a weighted aggregation.

138 3.3 Layer-wise Cosine Similarity Score

139 Inspired by the **Layer-Wise MultiKrum** defense method (details in Appendix A), and employing
 140 cosine similarity as the core measurement(Section 3.1), we introduce a novel metric termed the
 141 Layer-wise Cosine Similarity Score (LCSS). Diverging from traditional approaches that analyze
 142 client updates at the network level, LCSS enables independent analysis of updates at each individual
 143 layer.

144 Formally, let Δw_i^l denote the update of client i at the l -th layer, where $l \in [L]$ indexes the L layers of
 145 the global model and $i, j \in 1, \dots, N$ represent client indices among the total of N participants. The
 146 cosine similarity matrix S^l at layer l is defined as:

$$S_{i,j}^l = \frac{\langle \Delta w_i^l, \Delta w_j^l \rangle}{\|\Delta w_i^l\|_2 \cdot \|\Delta w_j^l\|_2} \quad (1)$$

147 where $\langle \cdot, \cdot \rangle$ denotes the inner product and $\|\cdot\|_2$ is the ℓ_2 -norm. This matrix quantifies the pairwise
 148 directional similarity of updates between clients at layer l .

149 For each client i , we identify the m nearest neighbors in similarity space based on S^l , and compute
 150 the average of the corresponding similarity scores to obtain the Layer-wise Cosine Similarity Score
 151 (LCSS) at layer l :

$$\text{LCSS}_i^l = \frac{1}{m} \sum_{j \in \mathcal{N}_i} S_{i,j}^l \quad (2)$$

152 where $\mathcal{N}_i$ denotes the set of the m most similar clients to client i at layer l . By averaging the m most
 153 similar clients at each layer, the metric captures local similarity patterns, enhancing the sensitivity to
 154 collusive behavior among malicious clients that may remain undetected at the global level(addressing
 155 C1). The effect of different m values on defense performance is explored in Section 4.4.

156 To evaluate the effectiveness of LCSS, we analyze its distribution under untargeted attack MPAF and
 157 LP attack on the CIFAR-10[22] and FashionMNIST[21] datasets, respectively. Each setup includes
 158 10 clients, among which 3 are malicious. Detailed dataset, model configurations, and other FL
 159 settings see Section 4.1. The results are shown in Figure 2. Note that in the LP attack, the targeted
 160 layers are L5, L7, and L8. For untargeted attacks, malicious clients exhibit higher LCSS values than
 161 benign ones. Similarly, under LP attack, this distinction also holds within the targeted layers.




Figure 2: LCSS distributions under the untargeted attack MPAF(Left) and the LP attack (Right).

3.4 Anomalous User Identification

The goal of anomalous user identification is to detect a set of potentially anomalous users $\{B_l\}_{l=1}^L$ that exhibit high LCSS across individual layers. This is based on the observed trend that malicious clients show higher LCSS on the poisoned layers.

To detect anomalous users in each layer $l \in \{1, \dots, L\}$, we apply a threshold-based filtering strategy on the LCSS. Starting from an initial threshold τ_0 , we iteratively lower the threshold until at least one user exceeds it. For each layer l , let the set of anomalous users be denoted by B_l . We define:

$$f_l(\tau) = \left\{ i \in \{1, \dots, N\} \mid \text{LCSS}_i^l \geq \tau \right\} \quad (3)$$

We then determine the final anomalous set B_l by selecting the largest threshold τ in the descending sequence $\{\tau_0, \tau_0 - 0.05, \dots, 0\}$ such that $f_l(\tau)$ is non-empty:

$$B_l = f_l(\tau^*) \quad \text{where} \quad \tau^* = \max \{ \tau \in \mathcal{T} \mid f_l(\tau) \neq \emptyset \} \quad (4)$$

where $\mathcal{T} = \{\tau_0, \tau_0 - 0.05, \dots, 0\}$. If the threshold drops to $\tau \leq 0$, it indicates the likely absence of further malicious clients. The resulting anomalous user sets for all layers are denoted by $\{B_l\}_{l=1}^L$.

Based on the patterns observed in Figure 2, we set the initial similarity threshold to $\tau_0 = 0.95$ and introduce an adaptive threshold adjustment mechanism. Specifically, τ is gradually decreased with a step size of 0.05 to accommodate the varying detection requirements across different layers.

3.5 User-Level Weight Assignment

To quantify the behavioral credibility of each client, we assign a user-level weight α_i to each client $i \in \{1, \dots, N\}$, based on the set of potentially anomalous users $\{B_l\}_{l=1}^L$. This process begins by identifying high-risk layers likely to have been targeted by poisoning, followed by calculating user-level weights α_i based on the anomalous user sets associated with those layers.

High-Risk Layer Detection. Anomalous user identification computes potentially high-similarity anomalous users for all layers, including non-targeted layers that are not poisoned by TLP attacks. The anomalous user sets in non-targeted layers are likely to consist mainly of benign clients. Therefore, it is necessary to filter the layers before analyzing $\{B_l\}_{l=1}^L$, in order to identify the truly targeted layers that have been poisoned. To address this challenge, we adopt a subtle idea: malicious updates tend to be more tightly clustered than benign ones, due to the higher similarity among poisoned updates. This clustering behavior is reflected in the distribution of LCSS scores—malicious updates exhibit LCSS values concentrated in a narrow range, while benign updates in non-targeted layers display a more dispersed, random-like distribution. This pattern is illustrated in Figure 2.

Building on it, we propose a method for high-risk layer detection. We identify a set of high-risk layers, denoted by $\mathcal{L}_{\text{HR}}$, based on the number of anomalous users in each layer. A layer is considered high-risk if its anomalous set size $|B_l|$ is greater than or equal to a threshold τ_{HR} . We initialize the threshold as $\tau_{\text{HR}} = \frac{N}{2} - 1$, which corresponds to the maximum possible number of malicious clients.

Formally, the high-risk layer set is defined as:

$$\mathcal{L}_{\text{HR}} = \{l \in \{1, \dots, L\} \mid |B_l| \geq \tau_{\text{HR}}\} \quad (5)$$

Here, τ_{HR} is progressively decreased from $\frac{N}{2} - 1$ to 1 until $\mathcal{L}_{\text{HR}} \neq \emptyset$.

User-Level Weight Calculation. For each user i , we compute the number of high-risk layers in which the user is flagged as anomalous:

$$|H_i| = \sum_{l \in \mathcal{L}_{\text{HR}}} \mathbb{I}(i \in B_l) \quad (6)$$

The user-level weight is then defined as:

$$\alpha_i = 1 - \frac{|H_i|}{|\mathcal{L}_{\text{HR}}|} \quad (7)$$

This step yields a user-level weight vector $\{\alpha_i\}_{i=1}^N$, where lower values indicate a higher frequency of being flagged in high-risk layers, and thus a greater likelihood of malicious behavior. Under

201 this formulation, even if certain non-targeted layers are mistakenly classified as high-risk due to
 202 distributional irregularities, or if a few benign updates in targeted layers exhibit high LCSS values
 203 similar to malicious ones due to special cases, these benign users are still assigned relatively higher
 204 user-level weights, as they appear in only a small fraction of high-risk layers(addressing C3).

205 3.6 Layer-Specific Weight Assignment

206 After computing the user-level weight vector $\{\alpha_i\}_{i=1}^N$, we propose a more fine-grained and flexible
 207 layer-specific weight $\{\beta_i^l\}_{i=1, l=1}^{N, L}$, building upon the user credibility quantified by the user-level
 208 weights. The user-level weight alone is insufficient to precisely suppress the impact of malicious
 209 updates for two key reasons. First, in the case of TLP attacks, malicious clients only poison the
 210 targeted layers, while their updates in non-targeted layers remain benign[11]. Applying a uniform
 211 user-level weight in this context would undesirably suppress not only the malicious updates in the
 212 targeted layers but also the benign updates in the non-targeted layers. Second, the purpose of the
 213 user-level weight is to accurately quantify user credibility based on the anomalous user set for each
 214 layer. This process does not analyze all users within each layer. Although this yields a reasonably
 215 accurate credibility estimate, it indirectly limits the suppression strength by failing to fully capture
 216 per-layer behaviors. The computation consists of two steps.

217 **Layer Update Boundary Definition.** For each layer l , we first define two behavioral boundaries
 218 based on the $LCSS_i^l$:

- 219 • The malicious boundary γ_{mal}^l is defined as the minimum similarity score among all users identified
 220 as anomalous in layer l :

$$\gamma_{\text{mal}}^l = \min_{i \in B_l} LCSS_i^l \quad (8)$$

- 221 • The benign boundary γ_{ben}^l is computed as the average similarity score of the top 50% of users with
 222 the highest user-level weights α_i , based on the presumption that malicious clients comprise less than
 223 50% of the total.:

$$\gamma_{\text{ben}}^l = \frac{1}{|\mathcal{U}_{\text{ben}}|} \sum_{i \in \mathcal{U}_{\text{ben}}} LCSS_i^l \quad (9)$$

224 Here, $\mathcal{U}_{\text{ben}}$ denotes the set of top 50% users sorted in descending order of α_i . The average is used
 225 because the LCSS of benign users tend to exhibit relatively random distribution patterns, and the
 226 mean provides a more robust representation of their typical behavior.

227 **Layer-Specific Weight Calculation.** Each user's weight in layer l is then determined according to
 228 the following piecewise function:

$$\beta_i^l = \begin{cases} 0, & \text{if } LCSS_i^l \geq \gamma_{\text{mal}}^l \\ 1, & \text{if } LCSS_i^l \leq \gamma_{\text{ben}}^l \\ \frac{\gamma_{\text{mal}}^l - LCSS_i^l}{\gamma_{\text{mal}}^l - \gamma_{\text{ben}}^l}, & \text{otherwise} \end{cases} \quad (10)$$

229 This step yields the matrix of layer-specific weights $\{\beta_i^l\}_{i=1, l=1}^{N, L}$. Considering the distributional
 230 variation of LCSS scores across different layers, we apply linear interpolation between the malicious
 231 and benign boundaries to determine each user's weight in each layer. In this way, the layer-specific
 232 weight addresses the near-threshold distribution phenomenon, which the user-level weight fails to
 233 capture. In this scenario, certain malicious updates in certain layers have LCSS values close to the
 234 threshold τ (e.g., 0.95), but do not exceed it (e.g., 0.945, 0.940), and thus are not included in the
 235 anomalous user set. Moreover, the benign boundary is dynamically adjusted per layer: in targeted
 236 layers, it tends to take lower values, whereas in non-targeted layers, it shifts toward higher values.
 237 This design helps mitigate the unintended suppression of benign updates in non-targeted layers by
 238 the layer-specific weighting mechanism.

239 3.7 Aggregation with Weighted Contributions

240 In each communication round t , the aggregation is performed in a layer-wise manner and incorporates
 241 both user-level reliability weights α_i and layer-specific contribution weights β_i^l .

242 To ensure that a user regarded as benign in a particular layer is not penalized by their user-level
 243 weight, we override the user-level weight in that layer: if $\beta_i^l = 1$, then α_i is replaced with 1 when
 244 computing the aggregation for layer l (addressing C2). Formally, the adjusted user-layer weight is
 245 defined as $\tilde{\alpha}_i^l = 1$ if $\beta_i^l = 1$, and $\tilde{\alpha}_i^l = \alpha_i$ otherwise. The global model is then updated as follows:

$$G^{(t)} = \sum_{l=1}^L \frac{\sum_{i=1}^N |D_i| \cdot \tilde{\alpha}_i^l \cdot \beta_i^l \cdot \Delta w_i^l}{\sum_{j=1}^N |D_j| \cdot \tilde{\alpha}_j^l \cdot \beta_j^l} \quad (11)$$

246 Here, $|D_i|$ (or $|D_j|$) denotes the number of local data samples held by client i (or j). This aggregation
 247 rule ensures that each user’s contribution is weighted according to both their overall credibility (based
 248 on inter-layer information) and their behavior in each layer (based on intra-layer information), allowing
 249 malicious updates to be suppressed without unnecessarily affecting benign ones (addressing C4).

250 4 Experiments

251 4.1 Setup

252 **Datasets and Models.** For untargeted attacks, we evaluate our method on three widely used
 253 FL datasets: FashionMNIST[21], CIFAR-10[22], and CINIC[23], using two different CNN
 254 architectures[26] in total. For the LP attack, we adopt the exact dataset–model combinations used
 255 in its original paper to ensure fair comparison: a CNN architecture[9, 10] on FashionMNIST, and
 256 ResNet18[24] and VGG19[25] on CIFAR-10. Detailed dataset and model configurations are provided
 257 in Appendix B.1 and Appendix B.2.

258 **Attacks and Compared Defenses.** We consider six state-of-the-art model poisoning attacks,
 259 including five untargeted attacks—LIE, Fang, Min-Max, Min-Sum, and MPAF[12–15]—as well
 260 as the Layer-wise Poisoning (LP) attack[11]. We compare our method against six advanced de-
 261 fense baselines, including four defenses that leverage cross-client information—Krum, Trimmed
 262 Mean, Norm Bound, and FLAME[5–8]—two defenses that utilize global information—FLTrust and
 263 FLDetector[9, 10].

264 **FL Settings.** We conduct all experiments using a NVIDIA A100 GPU. By default, for untargeted
 265 attacks, 30 out of 100 clients are selected per round, with a poisoning rate of 30%. A Dirichlet
 266 distribution is used to simulate non-IID data across clients[27, 28], with heterogeneity parameter
 267 $\beta = 0.1$. Each client trains locally for 1 epoch, over 300 communication rounds. For the LP attack,
 268 we follow its original FL setup but increase the poisoning rate to 30% for consistency. Note that
 269 in the original setting, non-IID distribution is controlled by parameter q [9, 13, 11], which we set to
 270 0.5. For **LayerGuard**, based on the discussion in Section 4.4, we set $m = 2$ to balance defense
 271 performance and computational efficiency. Detailed FL settings are provided in Appendix B.3.

272 **Metric.** We consider a set of metrics for evaluating detection and defense effectiveness, including
 273 Accuracy, Backdoor Success Rate (BSR), False Positive Rate (FPR), and False Negative Rate (FNR).

274 4.2 Main Results

275 In this section, we evaluate the defensive effectiveness of **LayerGuard** against both advanced
 276 untargeted attacks and the LP attack. In addition, we assess its detection accuracy and compare it
 277 with the advanced defense FLDetector. Due to space limitations, the detection accuracy results are
 278 presented in Appendix C.1.

279 **Defensive Effectiveness.** **LayerGuard** outperforms existing defenses. Table 1 and Table 2 present
 280 the end-to-end performance under untargeted and LP attacks, respectively. As shown, **LayerGuard**
 281 consistently achieves superior results against both types of attacks. Notably, under five advanced
 282 untargeted attacks, its accuracy is comparable to that of FedAvg under no-attack settings. Against
 283 the LP attack, **LayerGuard** achieves an average BSR of around 5% across all model–dataset
 284 configurations, successfully defending against LP attack that existing defenses fail to mitigate. This
 285 robustness stems from **LayerGuard**’s fine-grained evaluation of client credibility across layers. It is
 286 worth noting that **LayerGuard** incurs a moderate accuracy drop under no-attack settings compared

287 to FedAvg and some baselines. This is because the weighting mechanism still assigns lower weights
288 to certain benign updates with relatively high LCSS, when all clients are benign. Additionally, under
289 the LP attack, **LayerGuard** shows somewhat less stable BSR performance on deeper models, as a
290 larger number of layers increases the chance of errors in user-level weight assignment.

Table 1: Comparisons of final test accuracy under untargeted attacks. Results are averaged over three runs. “a ± b” indicates mean and standard deviation for LayerGuard.

Dataset	Attack	FedAVG	Krum	Trimmed Mean	Norm Bound	FLAME	FLTrust	FLDetector	LayerGuard
CINIC	No Attack	53.48	24.16	53.67	53.22	47.96	41.23	53.33	50.18±0.59
	Lie	21.73	11.84	32.22	21.60	48.40	42.96	14.11	53.26±0.18
	Fang	18.14	33.68	25.75	37.90	48.22	50.30	49.89	53.60±0.14
	Min-Max	38.52	9.99	35.22	39.20	50.66	40.78	13.23	53.35±0.20
	Min-Sum	46.40	12.20	32.41	46.65	47.17	41.81	14.33	53.57±0.17
	MPAF	11.71	10.90	13.63	12.53	51.52	34.14	10.91	53.66±0.29
FashionMNIST	No Attack	87.81	60.77	87.10	87.25	83.04	87.76	87.53	86.15±0.54
	Lie	63.60	30.78	77.78	62.77	80.54	82.74	79.27	87.14±0.09
	Fang	29.37	55.83	35.93	74.03	83.11	85.36	86.09	87.12±0.14
	Min-Max	80.55	10.01	78.99	83.15	82.88	84.14	79.62	86.96±0.22
	Min-Sum	85.49	54.85	78.52	85.92	79.45	82.32	84.33	86.88±0.17
	MPAF	20.46	10.22	73.49	37.73	83.88	81.74	10.00	87.09±0.23
CIFAR-10	No Attack	63.26	30.07	63.84	64.17	52.59	58.51	64.29	59.64±0.42
	Lie	24.10	12.05	35.72	24.02	58.95	51.11	33.45	63.82±0.35
	Fang	11.16	39.49	32.57	41.86	56.28	56.10	60.34	63.99±0.11
	Min-Max	44.12	10.33	41.15	49.31	59.88	56.49	9.59	64.43±0.20
	Min-Sum	57.56	12.60	40.58	55.94	57.64	55.29	19.62	64.51±0.33
	MPAF	12.95	10.42	18.22	12.91	57.66	50.33	10.00	63.91±0.29

Table 2: Comparison of final test accuracy and backdoor success rate under LP attack. Results are averaged over three runs. “a ± b” indicates mean and standard deviation for LayerGuard.

Model (Dataset)	VGG19 (CIFAR-10)		CNN (FashionMNIST)		ResNet18 (CIFAR-10)	
	Accuracy	BSR	Accuracy	BSR	Accuracy	BSR
FLAME	55.55	91.57	87.94	97.70	68.42	96.27
FLTrust	72.36	80.17	88.90	95.35	68.12	92.42
LayerGuard	75.41±0.86	5.22±2.32	89.88±0.34	0.41±0.16	73.04±0.51	7.45±3.45

291 4.3 Impact of FL Setting

292 In this part, we study the influence of different FL settings on our defense. By default, we conduct
293 experiments on CIFAR-10 for untargeted attacks and on FashionMNIST for LP attack. All other
294 settings follow the defaults in Section 4.1. Due to limited space, we use MPAF as the representative
295 untargeted attack. The results on other untargeted attacks are provided in Appendix C.2.

296 **Impact of the fraction of malicious clients.** Figure 3 presents the robustness of **LayerGuard** across
297 different poisoning rates. For untargeted attacks, as the number of malicious clients increases, the
298 impact of poisoning becomes more severe. However, **LayerGuard**’s accuracy remains consistently
299 stable and unaffected by the increasing poisoning rate. For LP attack, as the ratio of malicious
300 clients increases, the BSR stays consistently below 3%. Note that a limitation of **LayerGuard** is its
301 requirement for more than one malicious client to function effectively. In the LP attack experiment,
302 the 10% poisoning ratio means there is only one malicious client, thus, we do not show results for
303 this setting.




Figure 3: Impact of malicious client fraction.




Figure 4: Impact of non-IID degree.

Impact of non-IID degree. As shown in Figure 4, we evaluate the impact of different levels of non-IID data, including IID, on defense performance. Specifically, for untargeted attacks, we consider non-IID levels generated via a Dirichlet distribution[27, 28], where smaller β values correspond to stronger heterogeneity. For LP attacks, following the original paper’s setup, we use parameter q to control the non-IID level[9, 13, 11], where larger q indicates higher non-IIDness. Our results show that even under extreme non-IID conditions (e.g., $\beta = 0.05$, $q = 0.9$), **LayerGuard** maintains high accuracy and a low BSR. This robustness stems from the high similarity among malicious updates being unaffected by data heterogeneity. In contrast, existing advanced defenses that rely on the consistency of distribution across clients often degrade under high non-IID settings.

4.4 Ablation Study

The ablation study investigates two main questions: whether the choice of m in LCSS influences the results, and whether both User-Level Weight (ULW) and Layer-Specific Weight (LSW) are necessary for optimal performance, including the impact of removing either component.

Effect of Different m Values. As shown in Figure 5, no significant performance differences are observed with varying values of m for both untargeted and LP attacks. Notably, the accuracy and BSR are relatively worse when $m = 1$, possibly due to insufficient reliability in computing LCSS with only a single nearest neighbor. In our main experiments, we choose $m = 2$, which not only ensures excellent defense performance but also maintains a reasonable computational cost.




Figure 5: Impact of the value m .




Figure 6: Comparison of defensive effectiveness on ULW-only, LSW-only, and their combined usage.

Effect of User-Level and Layer-Specific Weights. Figure 6 shows the defense performance of ULW-only, LSW-only, and their combined usage. For untargeted attacks, both ULW-only and LSW-only achieve the highest accuracy, and using both together does not interfere with each other’s performance. This is largely attributed to the effectiveness of LCSS in detecting collusive behaviors under untargeted attacks. For LP attacks, we experiment with more deeper model VGG19. We observe that neither ULW-only nor LSW-only alone is sufficient to defend against the LP attack, indicating that for fine-grained, layer-specific TLP attack, combining both inter-layer and intra-layer update information is essential for effective defense.

4.5 Adaptive Attack

We propose an adaptive attack strategy that leverages knowledge of **LayerGuard**’s mechanisms by using decoy updates combined with controlled similarity manipulation. The detailed formulation and experimental evaluation of this attack are presented in Appendix D.

5 Conclusion

In this work, we proposed **LayerGuard**, a novel defense for FL against advanced model poisoning attacks, especially TLP. **LayerGuard** innovatively refines the defense perspective from evaluating the entire network to analyzing both intra-layer and inter-layer information, achieving precise anomaly detection and robust model updates. In the comprehensive evaluation, **LayerGuard** significantly outperforms current state-of-the-art defenses and successfully defends against the LP attack that previous methods fail to mitigate. We believe **LayerGuard** offers a promising direction for enhancing FL robustness in real-world applications.

References

- [1] Konečný, J., McMahan, H. B., Yu, F. X., Richtárik, P., Suresh, A. T., and Bacon, D. Federated learning: Strategies for improving communication efficiency. *arXiv preprint arXiv:1610.05492*, 2016.
- [2] McMahan, H. B., Moore, E., Ramage, D., Hampson, S., and y Arcas, B. A. *Communication-efficient learning of deep networks from decentralized data*. In AISTATS, 2017.
- [3] Rieke, N., Hancox, J., Li, W., Milletari, F., Roth, H. R., Albarqouni, S., Bakas, S., Galtier, M. N., Landman, B. A., Maier-Hein, K., et al. *The future of digital health with federated learning*. In NPJ Digital Medicine, 2020.
- [4] Yang, W., Zhang, Y., Ye, K., Li, L., and Xu, C.-Z. *Ffd: A federated learning based method for credit card fraud detection*. In International Conference on Big Data. Springer, 2019.
- [5] Blanchard, P., El Mhamdi, E. M., Guerraoui, R., and Stainer, J. *Machine learning with adversaries: Byzantine tolerant gradient descent*. In NeurIPS, 2017.
- [6] Yin, D., Chen, Y., Kannan, R., and Bartlett, P. *Byzantine-robust distributed learning: Towards optimal statistical rates*. In ICML, 2018.
- [7] Sun, Z., Kairouz, P., Suresh, A. T., and McMahan, H. B. *Can you really backdoor federated learning?* arXiv preprint arXiv:1911.07963, 2019.
- [8] Nguyen, T. D., Rieger, P., De Viti, R., Chen, H., Brandenburg, B. B., Yalame, H., Möllering, H., Fereidooni, H., Marchal, S., Miettinen, M., et al. *Flame: Taming backdoors in federated learning*. In USENIX Security Symposium, 2022.
- [9] Cao, X., Fang, M., Liu, J., and Gong, N. Z. *Fltrust: Byzantine-robust federated learning via trust bootstrapping*. In NDSS, 2021a.
- [10] Zhang, Z., Cao, X., Jia, J., and Gong, N. Z. *Fldetector: Defending federated learning against model poisoning attacks via detecting malicious clients*. In KDD, 2022.
- [11] Zhuang, H., Yu, M., Wang, H., Hua, Y., Li, J., and Yuan, X. *Backdoor federated learning by poisoning backdoor-critical layers*. arXiv preprint arXiv:2308.04466, 2023.
- [12] Baruch, G., Baruch, M., and Goldberg, Y. *A little is enough: Circumventing defenses for distributed learning*. In NeurIPS, 2019.
- [13] Fang, M., Cao, X., Jia, J., and Gong, N. *Local model poisoning attacks to byzantine-robust federated learning*. In USENIX Security Symposium, 2020.
- [14] Shejwalkar, V. and Houmansadr, A. *Manipulating the byzantine: Optimizing model poisoning attacks and defenses for federated learning*. In NDSS, 2021.
- [15] Cao, X. and Gong, N. Z. *Mpaf: Model poisoning attacks to federated learning based on fake clients*. In CVPR Workshops, 2022.
- [16] Xie, C., Huang, K., Chen, P.-Y., and Li, B. *DBA: Distributed backdoor attacks against federated learning*. In ICLR, 2019.
- [17] Bagdasaryan, E., Veit, A., Hua, Y., Estrin, D., and Shmatikov, V. *How to backdoor federated learning*. In AISTATS, 2020.
- [18] Bhagoji, A. N., Chakraborty, S., Mittal, P., and Calo, S. *Analyzing federated learning through an adversarial lens*. In Proceedings of ICML, 2019.
- [19] Tolpegin, V., Truex, S., Gursoy, M. E., and Liu, L. *Data poisoning attacks against federated learning systems*. In ESORICS, 2020.
- [20] Fung, C., Yoon, C. J., and Beschastnikh, I. *Mitigating sybils in federated learning poisoning*. arXiv preprint arXiv:1808.04866, 2018.
- [21] Xiao, H., Rasul, K., and Vollgraf, R. *Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms*. arXiv preprint arXiv:1708.07747, 2017.

- [22] Krizhevsky, A., Hinton, G., et al. *Learning multiple layers of features from tiny images*. 2009.
- [23] Darlow, L. N., Crowley, E. J., Antoniou, A., and Storkey, A. J. *Cinic-10 is not imagenet or cifar-10*. arXiv preprint arXiv:1810.03505, 2018.
- [24] He, K., Zhang, X., Ren, S., and Sun, J. *Deep Residual Learning for Image Recognition*. In Proceedings of IEEE CVPR, 2016.
- [25] Simonyan, K. and Zisserman, A. *Very deep convolutional networks for large-scale image recognition*. arXiv preprint arXiv:1409.1556, 2014.
- [26] Xie, Y., Fang, M., and Gong, N. Z. *FedREDefense: Defending against model poisoning attacks for federated learning using model update reconstruction error*. In Forty-first International Conference on Machine Learning, 2024. <https://openreview.net/forum?id=Wjq2bS7fTK>
- [27] Pi, R., Zhang, W., Xie, Y., Gao, J., Wang, X., Kim, S., and Chen, Q. *Dynafed: Tackling client data heterogeneity with global dynamics*. In CVPR, 2023.
- [28] Gu, H., Guo, B., Wang, J., Sun, W., Liu, J., Liu, S., and Yu, Z. *Fedaux: An efficient framework for hybrid federated learning*. In ICC, 2022.
- [29] Rieger, P., Nguyen, T. D., Miettinen, M., and Sadeghi, A.-R. *DeepSight: Mitigating backdoor attacks in federated learning through deep model inspection*. arXiv preprint arXiv:2201.00763, 2022.
- [30] Rieger, P., Krauß, T., Miettinen, M., Dmitrienko, A., and Sadeghi, A.-R. *CrowdGuard: Federated backdoor detection in federated learning*. arXiv preprint arXiv:2210.07714, 2023.

A Motivation and Inspiration

To validate the robustness of the LP attack, its researchers propose a targeted adaptive defense method named Layer-wise MultiKrum[5, 11], which extends the conventional MultiKrum strategy to a per-layer granularity. Specifically, Layer-wise MultiKrum independently performs the MultiKrum selection operation at each individual layer, aiming to identify anomalous updates within particular layers. Nevertheless, due to the high stealthiness of LP attack, updates in certain BC layers remain indistinguishable from benign clients, thereby successfully evading detection by Layer-wise MultiKrum.

The design philosophy of Layer-wise MultiKrum, along with its limitations in resisting LP attack, provides three critical insights for the development of our proposed method:

Insight 1: Layer-wise MultiKrum refines the defensive strategy from the holistic model level to a finer per-layer granularity, enabling more precise detection of anomalous behavior at individual layers.

Insight 2: Its core metric, the gradient score, is computed through aggregated distances from multiple nearest-neighbor gradients, effectively enhancing the discrimination capability toward abnormal gradients. This approach reduces the adverse impact of isolated anomalous gradients, thereby improving decision robustness.

Insight 3: Layer-wise MultiKrum aggregates final gradient scores by simply averaging several lowest gradient scores. This relatively coarse handling lacks a more sophisticated mechanism, which limits its effectiveness in detecting highly stealthy LP attack.

These three insights significantly influence our proposed method in terms of metric design and metric processing: Insights 1 and 2 inspire us to incorporate information across various layers and integrate multi-dimensional gradient features to enhance abnormal behavior detection. Insight 3 motivates us to refine the processing mechanism of our metrics, facilitating more accurate identification of highly covert attacks.

B Additional Experimental Setups

B.1 Datasets

FashionMNIST[21] is a grayscale image dataset containing 10 categories of clothing items, with 60,000 training images and 10,000 test images. CIFAR-10[22] is a color image dataset comprising 10 classes of everyday objects, including 50,000 training and 10,000 testing samples. CINIC[23] extends CIFAR-10 by incorporating downsampled images from ImageNet, resulting in a total of 270,000 images evenly divided into training, validation, and test sets.

B.2 Models

- The two CNN architectures[26](under the Apache 2.0 license) used in the defense experiments against untargeted attacks are specified in Table 3 and Table 4 for FashionMNIST and CIFAR-10/CINIC, respectively.
- In the defense experiments against the LP attack, we use three models: a CNN architecture[9, 10](under the MIT license), ResNet18[24], and VGG19[25]. The CNN model is trained on FashionMNIST, while ResNet18 and VGG19 are trained on CIFAR-10. The detailed structure of the CNN model is provided in Table 5.

Table 3: Model architecture for FashionMNIST(untargeted attack defense).

Layer Type	Size
Convolution + ReLU	$3 \times 3 \times 30$
Max Pooling	2×2
Convolution + ReLU	$3 \times 3 \times 50$
Max Pooling	2×2
Fully Connected + ReLU	100
Softmax	10

Table 4: Model architecture for CIFAR-10 and CINIC(untargeted attack defense).

Layer Type	Size
Convolution + ReLU	$3 \times 3 \times 32$
Max Pooling	2×2
Convolution + ReLU	$3 \times 3 \times 64$
Max Pooling	2×2
Fully Connected + ReLU	512
Softmax	10

Table 5: Model architecture for FashionMNIST(LP attack defense).

Layer	Size
Input	$28 \times 28 \times 1$
Convolution + ReLU	$3 \times 3 \times 32$
Convolution + ReLU	$3 \times 3 \times 64$
Max Pooling	2×2
Dropout	0.5
Fully Connected + ReLU	128
Dropout	0.5
Fully Connected	10

B.3 FL Settings

- In the experiments defending against untargeted attacks, all clients participate in each training round, with malicious clients launching attacks in every round. Local training uses a batch size

of 32. Stochastic Gradient Descent (SGD) is employed as the optimizer with a learning rate of 1×10^{-3} . Following prior work[27, 28], a Dirichlet distribution is used to simulate non-IID data across clients, with the heterogeneity parameter β set to 0.1. A smaller β indicates a higher degree of data heterogeneity.

- In the experiments against the LP attack, we follow the exact FL settings used in the original paper, except that the poisoning rate is increased to 30% for consistency. Among 100 clients, 10% are selected in each round. Each selected client trains for 2 local epochs using a batch size of 64. The global model is trained over 200 communication rounds. SGD is used as the optimizer, with a learning rate of 1×10^{-2} for FashionMNIST and 1×10^{-1} for CIFAR-10. In the original setup[11, 9, 13], non-IID data distribution is controlled by a parameter q , q is set to 0.5. Clients are divided into X groups corresponding to the X classes in the dataset. The probability of assigning samples with label x to the x -th group is q , and to other groups is $\frac{1-q}{X-1}$. Samples within each group are then uniformly distributed to clients.

C Additional Experimental Results

C.1 Detection Accuracy

Although **LayerGuard** is not specifically designed for detecting malicious clients but rather for suppressing their impact through adaptive weighting, its mechanism inherently provides a certain capability for malicious client detection. During the user-level weighting process, **LayerGuard** identifies high-risk layers and records the frequency with which a client’s updates appear in these layers. Based on this, we define a client as malicious only if it is consistently flagged across all high-risk layers. Conversely, if a client is not marked in even one high-risk layer, it is treated as benign. In other words, a client is considered malicious only when its user-level weight is exactly zero. This criterion is deliberately stringent to avoid false positives. Despite such a conservative definition, **LayerGuard** still demonstrates strong detection performance.

We evaluate the detection performance for identifying malicious clients, compared with FLDetector[10], an advanced defense capable of client-level detection, as shown in Table 6. **LayerGuard** consistently achieves perfect detection, with 0.00% FPR and FNR in all cases. In contrast, FLDetector suffers from high error rates, completely failing under attacks like MPAF. Its poor performance largely stems from its core assumption that clients with consistent updates are benign—an assumption that breaks down when malicious clients deliberately craft highly consistent updates, especially under non-IID conditions. These findings highlight **LayerGuard**’s clear advantage in delivering robust and precise detection, even under challenging attack scenarios and data heterogeneity.

Table 6: Comparison of detection performance (FPR and FNR) on different datasets.

Attack	Detector	FashionMNIST		CIFAR-10		CINIC	
		FPR	FNR	FPR	FNR	FPR	FNR
LIE	FLDetector	0.00	100.00	100.00	100.00	89.21	100.00
	LayerGuard	0.00	0.00	0.00	0.00	0.00	0.00
Fang	FLDetector	9.25	0.00	8.33	0.00	12.64	0.00
	LayerGuard	0.00	0.00	0.00	0.00	0.00	0.00
Min-Max	FLDetector	0.00	100.00	100.00	100.00	0.00	100.00
	LayerGuard	0.00	0.00	0.00	0.00	0.00	0.00
Min-Sum	FLDetector	0.00	100.00	100.00	100.00	84.29	100.00
	LayerGuard	0.00	0.00	0.00	0.00	0.00	0.00
MPAF	FLDetector	100.00	100.00	100.00	100.00	100.00	100.00
	LayerGuard	0.00	0.00	0.00	0.00	0.00	0.00

C.2 Impact of FL Setting

Impact of the fraction of malicious clients. The impact of different poisoning rates on the defense performance of LayerGuard under four other advanced untargeted attacks is shown in Figure 7.




Figure 7: Impact of malicious client fraction under four other untargted attacks.

Impact of non-IID degree. The effect of different non-IID degrees on the defense performance of LayerGuard under four other advanced untargted attacks is shown in Figure 8.




Figure 8: Impact of non-IID degree under four other untargted attacks.

D Detailed Description of the Adaptive Min-Sum Attack

Attack Strategies Against LayerGuard. The core defense mechanism of **LayerGuard** is its use of User-Level Weight (ULW) and Layer-Specific Weight (LSW) to suppress malicious updates. A key step in ULW calculation is the identification of high-risk layers, where the condition for marking a layer as high-risk is that it contains a sufficient number of malicious updates with LCSS above a threshold τ (Details are provided in Section 3). Malicious updates with LCSS below this threshold within a given layer are not considered in ULW but are instead handled by LSW.

A potential strategy to bypass **LayerGuard** is as follows: it assumes t malicious clients and constructs $m + 1$ malicious updates with high mutual similarity as decoys to manipulate the user-level weight (ULW) mechanism. The reason for constructing $m+1$ such updates is to ensure that the LCSS of the malicious updates remain consistently high and stable. These decoy updates exhibit LCSS above the similarity threshold, thereby triggering high-risk layer identification. The remaining $t - m - 1$ malicious updates—intended for the actual attack—are crafted with LCSS below the threshold. As a result, they are excluded from the anomalous user set and remain concealed under the cover of the decoys, effectively circumventing ULW’s suppression. To further evade detection by the layer-specific weight (LSW) mechanism, the $t - m - 1$ attackers reduce their internal similarity, weakening LSW’s ability to assign low weights to them, while maintaining the effectiveness of the attack.

Adaptive Min-Sum. Based on the bypass strategy described above, further derive an adaptive attack specifically targeting **LayerGuard**, **Adaptive Min-Sum**, which leverages the knowledge and settings of **LayerGuard**, aiming to bypass ULW-based suppression while minimizing the impact imposed by LSW. The original Min-Sum attack[14] generates malicious updates by first computing the average of all benign updates as a reference point, and then applying a shared perturbation to minimize the sum of distances between the malicious and benign updates. Since all malicious clients adjust their updates relative to the same global reference, their final updates exhibit highly similar directions—resulting in strong internal similarity.

To evade **LayerGuard**, Adaptive Min-Sum first constructs $m + 1$ standard Min-Sum-style malicious updates as decoys, leveraging their naturally high similarity to trigger high-risk layer detection. For the remaining $t - m - 1$ malicious clients, the attack introduces a new parameter n , where each such client randomly samples n benign updates and computes their mean as a personalized reference. This modification lowers the internal similarity among the remaining malicious updates, thereby reducing the suppressive effect of LSW on them. The parameter n controls the similarity between these malicious updates, smaller n results in lower mutual similarity.

The detailed procedure is as follows:

518 Assume there are t malicious clients in total. Among them, the first $k = m + 1$ clients act as decoys
 519 and follow the standard Min-Sum strategy, while the remaining $t - k$ clients adopt an adaptive strategy
 520 to reduce internal similarity and evade detection.

521 **Step 1.** For each adaptive malicious client $j \in \{k + 1, \dots, t\}$, randomly sample a subset B_j of size
 522 n from the set of benign clients B :

$$B_j \subset B, \quad |B_j| = n \quad (12)$$

523 **Step 2.** Let g_i denote the update from benign client i . The average of the selected subset is used as
 524 the new reference point:

$$\bar{g}_j = \frac{1}{n} \sum_{i \in B_j} g_i \quad (13)$$

525 **Step 3.** Each adaptive malicious client j then constructs its update by shifting from the reference
 526 point $\bar{g}_j$ along a fixed bias direction d :

$$g_{\text{mal}}^{(j)} = \bar{g}_j - \lambda d \quad (14)$$

527 The perturbation term λd follows the original Min-Sum attack strategy.

528 For the decoy clients $i \in \{1, \dots, k\}$, a shared reference $\bar{g}$ is computed as the average of all benign
 529 updates:

$$\bar{g} = \frac{1}{|B|} \sum_{i \in B} g_i, \quad g_{\text{mal}}^{(i)} = \bar{g} - \lambda d \quad (15)$$

530 Summary of malicious updates:

$$g_{\text{mal}}^{(u)} = \begin{cases} \bar{g} - \lambda d, & u = 1, \dots, k \\ \bar{g}_u - \lambda d, & u = k + 1, \dots, t \end{cases} \quad (16)$$

531 **Evaluation.** In the adaptive experiment, the value of m for **LayerGuard** is set to 3, and Adaptive
 532 Min-Sum is assumed to have prior knowledge of this parameter. Figure 9 presents the accuracy of
 533 various defenses under different values of n in the Adaptive Min-Sum attack. The **Trmean+Filter**
 534 approach first removes the four high-similarity decoy updates of Adaptive Min-Sum using a predefined
 535 filter. The remaining five malicious updates are then handled by the Trmean defense, to detect
 536 the remaining five malicious updates' attack potential. When $n = 7$ and $n = 4$, the remaining
 537 five malicious updates in Adaptive Min-Sum are still effective in launching the attack. However,
 538 **LayerGuard** successfully defends against Adaptive Min-Sum. As shown in Figure 10, even when
 539 $n = 4$ and ULW is completely ineffective, LSW still provides some defense. When $n = 1$, both
 540 ULW and LSW almost lose their effectiveness. In this case, **LayerGuard** experiences a slight loss in
 541 accuracy, but it remains within an acceptable range. While the Adaptive Min-Sum attack is more
 542 potent than traditional data poisoning attacks, it does not cause substantial harm to the system. This is
 543 due to a trade-off: in attempting to evade detection by our method, the attacker inadvertently weakens
 544 the potency of their own attack.




Figure 9: ULW and LSW performance with varying parameter n values under the Adaptive Min-Sum attack, based on the remaining 5 malicious updates.




Figure 10: Accuracy of three defense methods—LayerGuard, Trmean, and Trmean+Filter—under the Adaptive Min-Sum attack with varying parameter n .

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [NA].
- [NA] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for NA).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will be asked to also include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While "[Yes]" is generally preferable to "[No]", it is perfectly acceptable to answer "[No]" provided a proper justification is given (e.g., "error bars are not reported because it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). In general, answering "[No]" or "[NA]" is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- Delete this instruction block, but keep the section heading "NeurIPS Paper Checklist",
- Keep the checklist subsection headings, questions/answers and guidelines below.
- Do not modify the questions and only use the provided macros for your answers.

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification:

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Please refer to Section 4.2 and Section 4.3.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[NA\]](#)

Justification: The paper focuses entirely on algorithmic and experimental contributions without theoretical results.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.

- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Please refer to Section 4.1.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: <https://github.com/Immorash/LayerGuard>

Guidelines:

- The answer NA means that paper does not include experiments requiring code.

- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Please refer to Section 4.1.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Please refer to Section 4.2.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.

- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Please refer to Section 4.2.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification:

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Please refer to Section 5.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.

- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: The paper poses no such risks.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Please refer to Appendix B.2.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA]

Justification: The paper does not release new assets.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: LLMs were only used for editing and understanding background concepts.

Guidelines:

- 888 • The answer NA means that the core method development in this research does not
889 involve LLMs as any important, original, or non-standard components.
- 890 • Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>)
891 for what should or should not be described.