
PETIMOT: A Novel Framework for Inferring Protein Motions from Sparse Data Using SE(3)-Equivariant Graph Neural Networks¹

Anonymous Author(s)²

Affiliation

Address

email

Abstract³

1 Proteins move and deform to ensure their biological functions. Despite significant⁴
2 progress in protein structure prediction, approximating conformational ensembles at
3 physiological conditions remains a fundamental open problem. This paper presents
4 a novel perspective on the problem by directly targeting continuous compact rep-
5 resentations of protein motions inferred from sparse experimental observations.
6 We develop a task-specific loss function enforcing data symmetries, including
7 scaling and permutation operations. Our method PETIMOT (Protein sEquence and
8 sTructure-based Inference of MOTions) leverages transfer learning from pre-trained
9 protein language models through an SE(3)-equivariant graph neural network. When
10 trained and evaluated on the Protein Data Bank, PETIMOT shows superior perfor-
11 mance in time and accuracy, capturing protein dynamics, particularly large/slow
12 conformational changes, compared to state-of-the-art flow-matching approaches
13 and traditional physics-based models. *Our code and protocols are available at*
14 *<https://anonymous.4open.science/r/PETIMOT-4ED4/>.*

15 1 Introduction⁵

16 Proteins orchestrate biological processes in living organisms by interacting with their environment⁶
17 and adapting their three-dimensional (3D) structures to engage with cellular partners, including
18 other proteins, nucleic acids, small-molecule ligands, and co-factors. In recent years, spectacular
19 advances in high-throughput deep learning (DL) technologies have provided access to reliable
20 predictions of protein 3D structures at the scale of entire proteomes [1]. These breakthroughs have
21 also highlighted the complexities of protein conformational heterogeneity. State-of-the-art predictors
22 struggle to model alternative conformations, fold switches, large-amplitude conformational changes,
23 and solution ensembles [2].

24 The success of AlphaFold2 [3] has stimulated machine-learning approaches focused on inference-⁷
25 time interventions in the model to generate structural diversity. They include enabling or increasing
26 dropout [4, 5], or manipulating the evolutionary information given as input to the model [6–9].
27 Despite promising results on specific families, several studies have emphasised the difficulties in
28 rationalising the effectiveness of these modifications and interpreting them [10, 11]. Moreover,
29 these cannot be transferred to protein language model-based predictors that do not rely on multiple
30 sequence alignments. Researchers have also actively engaged in the development of deep-learning
31 frameworks based on diffusion, or the more general flow matching, to generate conformational
32 ensembles [12]. While family-specific models proved useful in exploring native-like conformational
33 landscapes, models trained across protein families still fail to approximate solution ensembles [13].

34 This work presents a new glance at the protein conformational diversity problem. Instead of learning 1
35 and sampling from multi-dimensional empirical distributions, we propose to learn eigenspaces (the
36 structure) of the positional covariance matrices in collections of experimental 3D structures and
37 generalize these over different homology levels. Our motivation is that the diversity present within
38 different 3D structures of the same protein or close homologs is a good proxy for the conformational
39 heterogeneity of proteins in solution [14] and can generally be (almost fully) explained by a small set
40 of linear vectors, also referred to as modes [15, 16]. Although linear spaces may not be well-suited
41 for capturing highly complex non-linear motions, such as loop deformations, they offer multiple
42 advantages. These include faster learning due to the reduced complexity of the model, improved
43 explainability as the components directly correspond to interpretable data dimensions, faster inference,
44 and the straightforward combination or integration of multiple data dimensions.

45 To summarize, our main contributions are: 2

- 46 • We provide a novel formulation of the protein conformational diversity problem. 3
- 47 • We present a novel benchmark representative of the Protein Data Bank structural diversity
48 and compiled with a robust pipeline [15], along with data- and task-specific metrics.
- 49 • We develop a SE(3)-equivariant Graph Neural Network architecture equipped with a novel
50 symmetry-aware loss function for comparing linear subspaces, with invariance to permu-
51 tation and scaling. Our model, PETIMOT, leverages embeddings from pre-trained protein
52 language models (pLMs), building on prior proof-of-concept work demonstrating that they
53 encode information about functional protein motions [17].
- 54 • PETIMOT is trained on sparse experimental data without any use of simulation data, in
55 contrast with Timewarp for instance [18]. Moreover, our model does not require physics-
56 based guidance or feedback, unlike [12] for instance.
- 57 • Our results demonstrate the capability of PETIMOT to generalise across protein families
58 (contrary to variational autoencoder-based approaches) and to compare favorably in running
59 time and accuracy to AlphaFlow, ESMFlow, and the Normal Mode Analysis.

60 2 Related Works 4

61 **Protein structure prediction.** AlphaFold2 was the first end-to-end deep neural network to achieve 5
62 near-experimental accuracy in predicting protein 3D structures, even for challenging cases with
63 low sequence similarity to proteins with resolved structures [3]. It extracts information from an
64 input multiple sequence alignment (MSA) and outputs all-atom 3D coordinates. Later works have
65 shown that substituting the input alignment by embeddings from a protein language model can yield
66 comparable performance [19–22].

67 **Generating conformational ensembles.** Beyond the single-structure frontier, several studies have 6
68 underscored the limitations and potential of protein structure predictors (PSP) for generating al-
69 ternative conformations [23, 24, 11, 2]. Approaches focused on re-purposing AlphaFold2 include
70 dropout-based massive sampling [4, 5], guiding the predictions with state-annotated templates
71 [25, 26], and inputting shallow, masked, corrupted, subsampled or clustered alignments [6–9]. De-
72 spite promising results, these approaches remain computationally expensive and their generalisability,
73 interpretability, and controllability remain unclear [11, 2]. More recent works have aimed at overcom-
74 ing these limitations by directly optimising PSP learnt embeddings under low-dimensional ensemble
75 constraints [27].

76 Another line of research has consisted in fine-tuning or re-training AlphaFold2 and other single- 7
77 state PSP under diffusion or flow matching frameworks [28, 13, 29]. For instance, the Al-
78 phaFlow/ESMFlow method progressively denoises samples drawn from a harmonic prior under
79 flow field controlled by AlphaFold or ESMFold [28]. It compares favourably with MSA subsam-
80 pling or clustering baselines, with a substantially superior precision-diversity Pareto frontier. More
81 generally, diffusion- and flow matching-based models allow for efficiently generating diverse confor-
82 mations conditioned on the presence of ligands or cellular partners [30, 31, 12, 32, 33]. Despite their
83 strengths, these techniques are prone to hallucination.

84 Parallel related works have sought to directly learn generative models of equilibrium Boltzmann 8
85 distributions using normalising flows [34, 18], or machine-learning force fields based on equivariant

86 graph neural network (GNN) representations [35], to enhance or replace molecular dynamics (MD) 1
 87 simulations.

88 **Protein conformational heterogeneity manifold learning.** Unsupervised, physics-based Normal 2
 89 Mode Analysis (NMA) has long been effective for inferring functional modes of deformation by
 90 leveraging the topology of a single protein 3D structure [36–38]. While appealing for its compu-
 91 tational efficiency, the accuracy of NMA strongly depends on the initial topology [39], limiting
 92 its ability to model extensive secondary structure rearrangements. Recent efforts have sought to
 93 address these limitations by directly learning continuous, compact representations of protein motions
 94 from sparse experimental 3D structures. These approaches employ dimensionality reduction tech-
 95 niques, from classical manifold learning methods [15] to neural network architectures like variational
 96 auto-encoders [40]. By projecting motions onto a learned low-dimensional manifold, these meth-
 97 ods enable reconstruction of accurate, physico-chemically realistic conformations, both within the
 98 interpolation regime and near the convex hull of the training data [15]. Additionally, they assist in
 99 identifying collective variables from molecular dynamics (MD) simulations, supporting importance-
 100 sampling strategies [41–45]. Despite these advances, such approaches are currently constrained to
 101 family-specific models.

102 **E(3)-equivariant graph neural networks.** Graph Neural Networks (GNN) have been extensively 3
 103 used to represent protein 3D structures. They are robust to transformations of the Euclidean group,
 104 namely rotations, reflections, and translations, as well as to permutations. In their simplest formulation,
 105 each node represents an atom and any pair of atoms are connected by an edge if their distance is
 106 smaller than a cutoff or among the smallest k interatomic distances. Many works have proposed
 107 to enrich this graph representation with SE(3)-equivariant features informing the model about
 108 interatomic directions and orientations [46–49, 31, 50]. For instance, VisNet captures the full local
 109 geometric information, including bonds, angles, as well as dihedral torsion and improper angles with
 110 node-wise high-order geometric tensors [50]. Moreover, to go beyond local 3D neighbourhoods while
 111 maintaining sub-quadratic complexity, Chroma adds in randomly sampled long-range connections
 112 [31].

113 3 Methods 4

114 3.1 Data representation 5

115 To generate training data, we exploit experimental protein single chain structures available in the 6
 116 Protein Data Bank (PDB) [51]. We first clustered these chains based on their sequence similarity.
 117 Then, within each cluster, we aligned the protein sequences and used the resulting mapping for
 118 superimposing the 3D coordinates [15]. It may happen that some residues in the multiple sequence
 119 alignment do not have resolved 3D coordinates in all conformations. To account for this uncertainty,
 120 we assigned a confidence score w_i to each residue i computed as the proportion of conformations
 121 including this residue. The 3D superimposition puts the conformations’ centers of mass to zero
 122 and then aims at determining the optimal least-squares rotation minimizing the Root Mean Square
 123 Deviation (RMSD) between any conformation and a reference conformation, while accounting for
 124 the confidence scores [52, 53],

$$E = \frac{1}{\sum_i w_i} \sum_i w_i (\vec{r}_{ij} - \vec{r}_{i0})^2, \quad 7 \quad (1)$$

125 where $\vec{r}_{ij} \in \mathbb{R}^3$ is the i th centred coordinate of the j th conformation and $\vec{r}_{i0} \in \mathbb{R}^3$ is the i th centred 8
 126 coordinate of the reference conformation. Next, we defined our ground-truth targets as eigenspaces
 127 of the coverage-weighted C α -atom positional covariance matrix,

$$C = \frac{1}{m-1} R^c W (R^c)^T = \frac{1}{m-1} (R - R^0) W (R - R^0)^T, \quad 9 \quad (2)$$

128 where R is the $3N \times m$ positional matrix with N the number of residues and m is the number of 10
 129 conformations, R^0 contains the coordinates of the reference conformation, and W is the $3N \times 3N$
 130 diagonal coverage matrix. The covariance matrix is a $3N \times 3N$ square matrix, symmetric and real.
 131 We decompose C as $C = Y D Y^T$, where Y is a $3N \times 3N$ matrix with each column defining a
 132 coverage-weighted eigenvector or a principal component that we interpret as a *linear motion*. D is a

133 diagonal matrix containing the eigenvalues. The latter highly depend on the sampling biases in the 1
 134 PDB and thus we do not aim at predicting them.

135 3.2 Problem formulation 2

136 For a protein of length N , let Y be $3N \times K$ *orthogonal* ground-truth deformations, 3

$$Y^T Y = I_K. \quad 4 \quad (3)$$

137 Our goal is to find coverage-weighted vectors $X \in \mathbb{R}^{3N \times L}$ whose components l *approximate* some 5
 138 components k of the ground truth Y :

$$W^{\frac{1}{2}} \tilde{\mathbf{x}}_l \approx \mathbf{y}_k. \quad 6 \quad (4)$$

139 Below, we provide three alternative formulations for this problem. 7

140 3.3 Geometric Loss 8

141 **The least-square formulation.** PETIMOT’s loss function serves two key purposes: it enables 9
 142 effective training of the network to predict subspaces representing multiple distinct modes of deformations – *i.e.*, with low overlap between the subspace’s individual linear vectors, while preventing
 143 convergence to a single dominant mode. For each protein of length N with a coverage W , we
 144 compare ground-truth directions Y with predicted motion directions X by computing a weighted
 145 pairwise *least-square difference* $\mathcal{L}_{kl}$ for each pair of a k direction in the ground truth and an l direction
 146 in the prediction,
 147

$$\mathcal{L}_{kl} = \frac{1}{N} \sum_{i=1}^N \|\vec{y}_{ik} - w_i^{1/2} c_{kl} \vec{x}_{il}\|^2 = \frac{1}{N} \mathbf{y}_k^T \mathbf{y}_k - \frac{1}{N} \frac{(\mathbf{y}_k^T W^{\frac{1}{2}} \mathbf{x}_l)^2}{\mathbf{x}_l^T W \mathbf{x}_l}, \quad 10 \quad (5)$$

148 where we scaled the ground-truth tensors such that $Y^T Y = N I_K$ and we used the fact that the 11
 149 optimal scaling coefficients c_{kl} between the k -th ground truth vector and the l -th prediction are given
 150 by

$$c_{kl} = \frac{\sum_{i=1}^N w_i^{\frac{1}{2}} \mathbf{y}_{ik}^T \mathbf{x}_{il}}{\sum_{i=1}^N w_i \mathbf{x}_{il}^T \mathbf{x}_{il}} = \frac{\mathbf{y}_k^T W^{\frac{1}{2}} \mathbf{x}_l}{\mathbf{x}_l^T W \mathbf{x}_l}. \quad 12 \quad (6)$$

151 This invariance to global scaling is motivated by the fact that we aim at capturing the relative 13
 152 magnitudes and directions of the motion patterns rather than their sign or absolute amplitudes.

153 **Linear assignment problem.** We then formulate an *optimal linear assignment problem* to find the 14
 154 minimum-cost matching between the ground-truth and the predicted directions. Specifically, we aim
 155 to solve the following assignment problem for the least-square (LS) costs,

$$\text{LS Loss} = \min_{\pi \in S_J} \sum_{k=1}^{\min(K,L)} \mathcal{L}_{k, \pi(k)} \quad 15 \quad (7)$$

subject to: 16

$$\pi : \{1, \dots, \min(K, L)\} \rightarrow \{1, \dots, L\}, \quad \pi(k) \neq \pi(k') \text{ for } k \neq k', \quad 17$$

156 where K and L are the number of ground-truth and predicted directions respectively, and $\pi(k)$ 18
 157 represents the index of the predicted direction assigned to the k -th ground truth direction. This
 158 formulation ensures an optimal one-to-one matching, while accommodating cases where the number
 159 of predicted and ground-truth directions differs. We backpropagate the loss only through the optimally
 160 matched pairs, using *scipy* `linear_sum_assignment`. We have also tested a smooth version of the loss
 161 above with continuous gradients, but it did not improve the performance.

162 **The subspace coverage formulation.** We propose another formulation of the problem in terms of 19
 163 the subspace coverage metrics [54–56]. Specifically, we sum up *squared sinus* (SS) dissimilarities

164 between ground-truth and predicted directions (formally computed as one minus squared cosine 1
 165 similarity),

$$\text{SS Loss} = 1 - \frac{1}{K} \sum_{k=1}^K \sum_{l=1}^K (\mathbf{y}_k^T W^{\frac{1}{2}} \mathbf{x}_l^\perp)^2, \quad (8)$$

166 where the subspace $\{\mathbf{x}_l^\perp\}$ is obtained by orthogonalising the coverage-weighted predicted linear 3
 167 subspace $\{W^{\frac{1}{2}} \mathbf{x}_l\}$, where $\mathbf{x}_l^T W \mathbf{x}_l = 1$, using the Gram–Schmidt process. This operation ensures
 168 that the loss ranges from zero for mutually orthogonalising subspaces to one for identical subspaces
 169 and avoids artificially inflating the SS loss due to redundancy in the predicted motions. The order
 170 in which the predicted vectors are orthogonalised does not influence the loss, guaranteeing stable
 171 training. Appendix A proves this statement.

172 **Independent Subspace (IS) Loss.** We can substitute the orthogonalisation procedure by using an 4
 173 auxiliary loss component for maximising the rank of the predicted subspace. For this purpose, we
 174 chose the squared cosine similarity computed between pairs of predicted vectors. The final expression
 175 for the *independent subspace* (IS) loss is

$$\text{IS Loss} = \frac{1}{K^2} \sum_{k=1}^K \sum_{l=1}^K (\mathbf{x}_k^T W \mathbf{x}_l)^2 - \frac{1}{K^2} \sum_{k=1}^K \sum_{l=1}^K (\mathbf{y}_k^T W^{\frac{1}{2}} \mathbf{x}_l)^2, \quad (9)$$

176 where the predictions $\{\mathbf{x}_l\}$ are normalised prior to the loss computation such that $\mathbf{x}_l^T W \mathbf{x}_l = 1$ and 6
 177 the scaling factor K^2 ensures that the loss ranges between 0 and 1. Appendix A analyses the stability
 178 of this formulation.

179 3.4 Architecture 7




Figure 1: **PETIMOT’s architecture overview.** The model processes both sequence embeddings (s) 9
 and motion vectors ($\vec{x}$) through 15 message-passing blocks. Each block updates both representations
 by aggregating information from neighboring residues. Neighbor features are computed in the
 reference frame of the central residue i , ensuring SE(3) equivariance. The geometric features encoded
 in the edges capture the relative spatial relationships between residue pairs. Three types of losses
 (LS, SS, and IS) are computed, with prior normalization of the predictions for the IS and SS losses,
 and an additional orthogonalisation of the predictions for the SS loss.

180 **Dual-Track Representation.** PETIMOT processes protein sequences through a message-passing 10
 181 neural network that simultaneously handles residue embeddings and motion vectors in local coordinate
 182 frames (Fig. 1). For each residue i , we define and update a node embedding $s_i \in \mathbb{R}^d$ initialized
 183 from protein language model features and a set of K motion vectors $\{\vec{x}_{ik}\}_{k=1}^K \in \mathbb{R}^{3 \times K}$ initialized
 184 randomly. The message passing procedure is detailed in Algorithm B.1 of Appendix B.2.

Graph Construction. The protein is represented as a graph where nodes correspond to the residues, and edges capture spatial relationships. For each residue i , we connect to its k nearest neighbors based on $C\alpha$ distances and l randomly selected residues. This hybrid connectivity scheme ensures both local geometric consistency and global information flow, while maintaining sparsity for computational efficiency. Indeed, our model scales *linearly* with the length N of a protein. In our base model we set $k = 5$ and $l = 10$.

Node features. We chose ProstT5 as our default protein language model for initialising node embeddings [57]. This structure-aware pLM offers an excellent balance between model size – including the number of parameters and embedding dimensionality – and performance [17].

Local Reference Frames. Each residue’s backbone atoms (N, CA, C) define a local reference frame through a rigid transformation $T_i \in SE(3)$. For each residue pair (i, j) , we compute their relative transformation $T_{ij} = T_i^{-1} \circ T_j$ from which we extract the rotation $R_{ij} \in SO(3)$ and translation $\vec{t}_{ij} \in \mathbb{R}^3$. Under global rotations and translations of the protein, these relative transformations remain invariant.

Edge Features. Edge features e_{ij} provide an SE(3)-invariant encoding of the protein structure through relative orientations, translational offsets, protein chain distance, and a complete description of peptide plane positioning captured by pairwise backbone atom distances. See Appendix B.3 for more details. The training procedure is detailed in Appendix B.4.

4 Results

Metrics	PETIMOT	AlphaFlow	ESMFlow	NMA
Success Rate (%) $\uparrow$	43.57	31.80	26.82	24.88
Running time $\downarrow$	15.82s	38h 07min	10h 41min	43.59s
Min. LS Error $\downarrow$	0.61 $\pm$ 0.22	0.68 $\pm$ 0.21	0.70 $\pm$ 0.22	0.72 $\pm$ 0.20
Min. Magnitude Error $\downarrow$	0.21 $\pm$ 0.12	0.24 $\pm$ 0.12	0.26 $\pm$ 0.14	0.27 $\pm$ 0.14
OLA LS Error $\downarrow$	0.83 $\pm$ 0.10	0.86 $\pm$ 0.10	0.87 $\pm$ 0.10	0.88 $\pm$ 0.10
OLA Magnitude Error $\downarrow$	0.41 $\pm$ 0.14	0.43 $\pm$ 0.14	0.47 $\pm$ 0.15	0.48 $\pm$ 0.15
Global SS Error $\downarrow$	0.73 $\pm$ 0.14	0.78 $\pm$ 0.14	0.80 $\pm$ 0.14	0.79 $\pm$ 0.14

Table 1: Success rate and average performance on the test set. We report performance metrics on a subset of our test set, namely the 824 test proteins that we could process with AlphaFlow (smaller than 450 amino acids). The success rate is defined as the proportion of test proteins with at least one motion predicted at a reasonable accuracy, namely an LS error below 0.6. The other values are averages computed over the test proteins. Min. stands for the best matching pair of predicted and ground-truth vectors. OLA refers to the optimal linear assignment between all predicted and ground-truth vectors. Arrows indicate whether higher ($\uparrow$) or lower ($\downarrow$) metrics values are better. Best results are shown in **bold**. Running times are recorded on a Intel(R) Xeon(R) W-2245 CPU @ 3.90GHz equipped with GeForce RTX 3090. PETIMOT (with 4 directions), AlphaFlow (50 models), and ESMFlow (50 models) were executed on a GPU, while NOLB NMA (with 10 lowest modes) only used CPU.

Training and evaluation. We trained PETIMOT against linear motions extracted from all $\sim 750,000$ protein chains from the PDB (as of June 2023) clustered at 80% sequence identity and coverage. Our full training data comprises 7 335 conformational collections split into 70% for training, 15% for validation and 15% for test. We augmented the training and validation data by computing the motions with respect to 5 reference conformations per collection. As a result, the full training set comprises 25,595 samples. We set the numbers of predicted and ground-truth motions, $K = L = 4$. See Appendix B.1 for more details.

At inference, we consider $w_i = 1, \forall i = 1..N$. We rely on four main evaluation metrics aimed at addressing the following questions:




Figure 2: **Cumulative error curves computed on the test proteins.** **a-b.** Comparison between PETIMOT base model and three other methods. **c-d.** Comparison between different losses implemented in PETIMOT. The loss of the base model is LS + SS. **a,c.** Minimum LS error corresponding to the best matching pair of predicted and ground-truth motions. **b,d.** SS error computed between the entire predicted and ground-truth subspaces.

- Is PETIMOT able to approximate at least one of the main linear motions of a given protein? For this, we rely on the minimum LS error over all possible pairs of predicted and ground-truth vectors.
- To what extent does PETIMOT capture the main motion linear subspace of a given protein? For this, we use the global SS error.
- Is PETIMOT able to identify the residues that move the most? Here, we rely on the magnitude error, $\frac{1}{N} \sum_{i=1}^N (\|\vec{y}_{ik}\|^2 - \|c_{kl}\vec{x}_{il}\|^2)$.
- How fast is PETIMOT at inference?

Comparison with other methods. PETIMOT showed a better capacity to approximate individual motions and to globally capture motion subspaces than the flow matching-based frameworks AlphaFlow and ESMFlow, and also the unsupervised physics-based Normal Mode Analysis (Fig. 2a-b). It approximated at least one motion with reasonable accuracy (LS error below 0.6) for 43.57% of the test proteins, while the success rate was only 31.80%, 26.82%, and 24.88% for AlphaFlow, ESMFlow, and the NMA, respectively (Table 1). PETIMOT’s best predicted vector better matched a ground-truth vector than any other methods in 43.57% of the cases (Fig. 3a). PETIMOT was also better at identifying which residues contribute the most to the motions (Table 1). Furthermore, PETIMOT was significantly faster at inference - it took about 16s for the whole test set, followed by NOLB (44s), ESMFlow (11h) and AlphaFlow (38h). See Appendix B.5 for more evaluation details. For all the metrics PETIMOT performs the best, with a particular striking difference in performance for the success rate metrics. However, it maybe not very informative to look at a single value averaged over the whole test set. Thus, we also suggest to analyze more informative plots, *e.g.* those in Fig. 2 and Fig. 3a.

Comparison of problem formulations. Our base model combining the LS and SS losses with equal weights outperforms all three individual losses, LS, SS, and LS (Fig. 2c-d). It strikes an excellent balance between approximating individual motions with high accuracy (Fig. 2c) and globally covering the motion subspaces (Fig. 2d). By comparison, the SS and IS losses tend




Figure 3: Individual predictions. **a.** The per-protein minimum LS errors, computed for the best-matching pairs between predicted and ground-truth vectors, are reported for PETIMOT (black), the NMA (red), AlphaFlow (blue) and ESMFlow (green). The values are in ascending order of the errors computed for PETIMOT, from best to worse. **b-c.** Trajectories generated by deforming a protein structure along PETIMOT best predicted motion. Five trajectory snapshots are shown colored from yellow to orange. **b.** *Bacillus subtilis* xylanase A (PDB id: 3EXU, chain A). **c.** Murine Fab fragment (PDB id: 7SD2, chain A).

to underperform on individual motions while the LS loss tends to provide lower coverage of the ground-truth subspaces. See Appendix C for additional results.

Contribution of sequence and structure features. We performed an ablation study to assess the contribution of sequence and structure information to our architecture. Our results show that ProstT5 slightly outperforms the more recent and larger pLM, ESM-Cambrian 600M [58] (Fig. B.2). Geometrical information about protein structure provides the most significant contribution, as replacing ProstT5 embeddings with random numbers has only a small impact on network performance. Conversely, the network’s performance without structural information strongly depends on the chosen pLM. While the structure-aware embeddings from ProstT5 partially compensate for missing 3D structure information, relying solely on ESM-C embeddings results in poor performance (Fig. B.2). Moreover, connecting each residue to its 15 nearest neighbours (sorted according to $C\alpha$ - $C\alpha$ distances) in the protein graph results in lower performance compared to introducing randomly chosen edges or even fully relying on random connectivity (Fig. B.4).

Generalisation capability. Because our conformational collections are defined based on 80% sequence identity and coverage thresholds, some test proteins may be close homologs of the training proteins. Yet, the sequence identity and structural similarity (TM-score) of the test proteins with respect to the training set do not determine the quality of PETIMOT predictions (Fig. C.3). PETIMOT provides high-quality predictions for a number of test proteins that do not share any detectable similarity and only weak structural similarity (TM-score below 0.5) to the training set. To further assess PETIMOT’s generalisation capability across protein families, we implemented a more stringent train-validation-test split using a 30% sequence similarity threshold. This process involved three key steps: (1) removing from training and validation sets any proteins sharing >30% sequence similarity with test proteins, (2) redefining these sets through random splitting at the 30% sequence identity level, and (3) reducing redundancy within each set to ensure balanced representation of protein families (detailed in Appendix B.6). When trained under this stringent protocol, PETIMOT achieved a success rate of 40.87% on our non-redundant set of 734 test proteins, consistently outperforming all baseline methods (see Appendix Fig. C.4a-b). This demonstrates our architecture’s robustness, as the more challenging similarity restrictions did not dramatically reduce performance. For direct comparison between our two model versions *PETIMOT-default* (trained with 80% similarity threshold) and *PETIMOT-stringent* (trained with 30% non-redundant threshold) we evaluated performance on 474 test proteins that share less than 30% sequence similarity with any protein used in training or

validation of either model. *PETIMOT-stringent* achieved a slightly higher success rate (43.67%) compared to *PETIMOT-default* (41.98%) (see Appendix Fig. C.4c-d). This result emphasizes the benefit of encouraging better generalization to evolutionary distant protein families through enhanced data splitting strategies.

Conformation generation. PETIMOT allows straightforwardly generating conformational ensembles or trajectories by deforming an initial protein 3D structure along one or a combination of predicted motions. We showcase this functionality on two example proteins, the xylanase A from *Bacillus subtilis* and the periplasmic domain of Gliding motility protein GldM from *Capnocytophaga canimorsus* (Fig. 3b-c). We used PETIMOT predictions to generate physically realistic conformations representing either the open-to-closed transition of xylanase A thumb (Fig. 3b) or the flexibility of the heavy chain antibody IgE/Fab anti-profilin Hev b 8 (Fig. 3c). Figures C.5 and C.6 compare predicted motions for these proteins with the ground truth.

Limitations. One current limitation of our approach is the scarcity of functional motions in the training set, raising concerns about its accuracy and completeness. The success rates computed for PETIMOT, although higher than the baseline methods, remain modest, suggesting that some ground-truth motions may lack biological or physical relevance. They may result from experimental artifacts, most often of crystallographic origin. Nevertheless, establishing objective criteria to reliably distinguish artifactual conformations from biologically relevant ones remains infeasible. Our working assumption is that a subset of the conformational manifold in a protein collection always represents functional motions. To address this challenge, we designed our training loss function specifically to evaluate submanifolds by calculating the minimum error between each reference motion and the set of predicted motions, allowing the model to capture conformational diversity while mitigating the impact of potential artefacts.

Another limitation is that we do not account for the different time scales at which different types of protein motions occur. We do not have this temporal information in the PDB training set. The only information we have is the proxy of the amplitudes of the motions, represented by the eigenvalues in the original collections. However, these amplitudes might not be relevant as they strongly depend on the uneven sampling in the PDB and different experimental conditions and temperatures. This motivated us to only reconstruct the shape of the motion manifolds, not their amplitudes or effective timescales.

In addition, our approach is limited to modeling protein motions as linear displacement vectors. While this approximation is sufficient to describe most of the observed conformational heterogeneity, it remains inadequate for modeling highly complex non-linear deformations. Furthermore, deforming proteins structures along linear motion direction may produce unrealistic conformations at large amplitudes. A possible solution yet to be investigated would be the nonlinear extrapolation technique that have been widely used in molecular mechanics [37]. Finally, a direction for future improvement would be to provide per-residue confidence estimates for the predicted motions.

5 Conclusion

In this work, we have proposed a new perspective on the problem of capturing protein continuous conformational heterogeneity. Compared to state-of-the-art methods, our approach goes beyond generating alternative protein conformations by directly inferring compact and continuous representations of protein motions. Our comprehensive analysis of PETIMOT's predictive capabilities demonstrates its performance and utility for understanding how proteins deform to perform their functions. Our work opens ways to future developments in protein motion manifold learning, with exciting potential applications in protein engineering and drug development.

Broader impacts. This paper presents work aimed at advancing the field of computational biology and machine learning, specifically in the context of understanding protein dynamics. While the goal is primarily to enhance scientific understanding and methodological rigor, there are potential societal consequences to consider. These include the facilitation of drug discovery and protein engineering applications, as well as ethical concerns regarding the dual use of protein modeling. We emphasize the importance of using these tools responsibly to maximize benefits while mitigating risks. Beyond these considerations, there are no specific societal implications we feel require urgent attention.

References

- [1] Mihaly Varadi, Damian Bertoni, Paulyna Magana, Urmila Paramval, Ivanna Pidruchna, Malarvizhi Radhakrishnan, Maxim Tsenkov, Sreenath Nair, Milot Mirdita, Jingsi Yeo, et al. Alphafold protein structure database in 2024: providing structure coverage for over 214 million protein sequences. *Nucleic acids research*, 52(D1):D368–D375, 2024.
- [2] Devlina Chakravarty, Myeongsang Lee, and Lauren L Porter. Proteins with alternative folds reveal blind spots in alphafold-based protein structure prediction. *Current Opinion in Structural Biology*, 90:102973, 2025.
- [3] John Jumper, Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, Olaf Ronneberger, Kathryn Tunyasuvunakool, Russ Bates, Augustin Žídek, Anna Potapenko, Alex Bridgland, Clemens Meyer, Simon A. A. Kohl, Andrew J. Ballard, Andrew Cowie, Bernardino Romera-Paredes, Stanislav Nikolov, Rishub Jain, Jonas Adler, Trevor Back, Stig Petersen, David Reiman, Ellen Clancy, Michal Zielinski, Martin Steinegger, Michalina Pacholska, Tamas Berghammer, Sebastian Bodenstein, David Silver, Oriol Vinyals, Andrew W. Senior, Koray Kavukcuoglu, Pushmeet Kohli, and Demis Hassabis. Highly accurate protein structure prediction with alphafold. *Nature*, 596(7873):583–589, 2021. doi: 10.1038/s41586-021-03819-2. URL <https://doi.org/10.1038/s41586-021-03819-2>.
- [4] Nessim Raouraoua, Claudio Mirabello, Thibaut Véry, Christophe Blanchet, Björn Wallner, Marc F. Lensink, and Guillaume Brysbaert. MassiveFold: unveiling AlphaFold’s hidden potential with optimized and parallelized massive sampling. *Nature Computational Science*, 4(11):824–828, 2024. doi: 10.1038/s43588-024-00714-4. URL <https://doi.org/10.1038/s43588-024-00714-4>.
- [5] Björn Wallner. Afsample: improving multimer prediction with alphafold using massive sampling. *Bioinformatics*, 39(9):btad573, 2023.
- [6] Yogesh Kalakoti and Björn Wallner. Afsample2: Predicting multiple conformations and ensembles with alphafold2. *bioRxiv*, pages 2024–05, 2024.
- [7] Hannah K Wayment-Steele, Adedolapo Ojoawo, Renee Otten, Julia M Apitz, Warintra Pitsawong, Marc Hömberger, Sergey Ovchinnikov, Lucy Colwell, and Dorothee Kern. Predicting multiple conformations via sequence clustering and alphafold2. *Nature*, pages 1–3, 2023.
- [8] Diego Del Alamo, Davide Sala, Hassane S Mchaourab, and Jens Meiler. Sampling alternative conformational states of transporters and receptors with alphafold2. *Elife*, 11:e75751, 2022.
- [9] Richard A Stein and Hassane S Mchaourab. Speech_af: Sampling protein ensembles and conformational heterogeneity with alphafold2. *PLOS Computational Biology*, 18(8):e1010483, 2022.
- [10] Lauren L Porter, Irina Artsimovitch, and César A Ramírez-Sarmiento. Metamorphic proteins and how to find them. *Current opinion in structural biology*, 86:102807, 2024.
- [11] Patrick Bryant and Frank Noé. Structure prediction of alternative protein conformations. *Nature Communications*, 15(1):7328, 2024.
- [12] Y Wang, L Wang, Y Shen, Y Wang, H Yuan, Y Wu, and Q Gu. Protein conformation generation via force-guided se (3) diffusion models, 2024. doi: 10.48550. *arXiv preprint ARXIV:2403.14088*, 2025.
- [13] Josh Abramson, Jonas Adler, Jack Dunger, Richard Evans, Tim Green, Alexander Pritzel, Olaf Ronneberger, Lindsay Willmore, Andrew J. Ballard, Joshua Bambrick, Sebastian W. Bodenstein, David A. Evans, Chia-Chun Hung, Michael O’Neill, David Reiman, Kathryn Tunyasuvunakool, Zachary Wu, Akvilė Žemgulytė, Eirini Arvaniti, Charles Beattie, Ottavia Bertolli, Alex Bridgland, Alexey Cherepanov, Miles Congreve, Alexander I. Cowen-Rivers, Andrew Cowie, Michael Figurnov, Fabian B. Fuchs, Hannah Gladman, Rishub Jain, Yousuf A. Khan, Caroline M. R. Low, Kuba Perlin, Anna Potapenko, Pascal Savy, Sukhdeep Singh, Adrian Stecula, Ashok Thillaisundaram, Catherine Tong, Sergei Yakneen, Ellen D. Zhong, Michal Zielinski, Augustin Žídek, Victor Bapst, Pushmeet Kohli, Max Jaderberg, Demis

372 Hassabis, and John M. Jumper. Accurate structure prediction of biomolecular interactions with
373 alphafold 3. *Nature*, 630(8016):493–500, 2024. doi: 10.1038/s41586-024-07487-w. URL
374 <https://doi.org/10.1038/s41586-024-07487-w>.

375 [14] Robert B Best, Kresten Lindorff-Larsen, Mark A DePristo, and Michele Vendruscolo. Relation
376 between native ensembles and experimental structures of proteins. *Proceedings of the National
377 Academy of Sciences*, 103(29):10901–10906, 2006.

378 [15] Valentin Lombard, Sergei Grudinin, and Elodie Laine. Explaining conformational diversity in
379 protein families through molecular motions. *Scientific Data*, 11(1):752, 2024.

380 [16] Lee-Wei Yang, Eran Eyal, Ivet Bahar, and Akio Kitao. Principal component analysis of
381 native ensembles of biomolecular structures (pca_nest): insights into functional dynamics.
382 *Bioinformatics*, 25(5):606–614, 2009.

383 [17] Valentin Lombard, Dan Timsit, Sergei Grudinin, and Elodie Laine. Seamoon: Prediction of
384 molecular motions based on language models. *bioRxiv*, pages 2024–09, 2024.

385 [18] Leon Klein, Andrew Foong, Tor Fjelde, Bruno Mlodozeniec, Marc Brockschmidt, Sebastian
386 Nowozin, Frank Noé, and Ryota Tomioka. Timewarp: Transferable acceleration of molecular
387 dynamics by learning time-coarsened dynamics. *Advances in Neural Information Processing
388 Systems*, 36, 2024.

389 [19] Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Nikita Smetanin,
390 Robert Verkuil, Ori Kabeli, Yaniv Shmueli, Allan Dos Santos Costa, Maryam Fazel-Zarandi,
391 Tom Sercu, Salvatore Candido, and Alexander Rives. Evolutionary-scale prediction of atomic-
392 level protein structure with a language model. *Science*, 379(6637):1123–1130, 2023.

393 [20] Thomas Hayes, Roshan Rao, Halil Akin, Nicholas J. Sofroniew, Deniz Oktay, Zeming Lin,
394 Robert Verkuil, Vincent Q. Tran, Jonathan Deaton, Marius Wiggert, Rohil Badkundri, Irhum
395 Shafkat, Jun Gong, Alexander Derry, Raul S. Molina, Neil Thomas, Yousuf Khan, Chetan
396 Mishra, Carolyn Kim, Liam J. Bartie, Matthew Nemeth, Patrick D. Hsu, Tom Sercu, Salvatore
397 Candido, and Alexander Rives. Simulating 500 million years of evolution with a language
398 model. *bioRxiv*, 2024. doi: 10.1101/2024.07.01.600583. URL [https://www.biorxiv.org/
399 content/early/2024/07/02/2024.07.01.600583](https://www.biorxiv.org/content/early/2024/07/02/2024.07.01.600583).

400 [21] Konstantin Weissenow, Michael Heinzinger, and Burkhard Rost. Protein language-model
401 embeddings for fast, accurate, and alignment-free protein structure prediction. *Structure*, 30(8):
402 1169–1177, 2022.

403 [22] Ruidong Wu, Fan Ding, Rui Wang, Rui Shen, Xiwen Zhang, Shitong Luo, Chenpeng Su, Zuofan
404 Wu, Qi Xie, Bonnie Berger, et al. High-resolution de novo structure prediction from primary
405 sequence. *BioRxiv*, pages 2022–07, 2022.

406 [23] Tadeo Saldaño, Nahuel Escobedo, Julia Marchetti, Diego Javier Zea, Juan Mac Donagh,
407 Ana Julia Velez Rueda, Eduardo Gonik, Agustina García Melani, Julieta Novomisky Nechcoff,
408 Martín N Salas, et al. Impact of protein conformational diversity on alphafold predictions.
409 *Bioinformatics*, 38(10):2742–2748, 2022.

410 [24] Thomas J Lane. Protein structure prediction has reached the single-structure frontier. *Nature
411 Methods*, 20(2):170–173, 2023.

412 [25] Bulat Faizov and Roland L Dunbrack Jr. Alphafold2 models of the active form of all 437
413 catalytically-competent typical human kinase domains. *bioRxiv*, pages 2023–07, 2023.

414 [26] Lim Heo and Michael Feig. Multi-state modeling of g-protein coupled receptors at experimental
415 accuracy. *Proteins: Structure, Function, and Bioinformatics*, 90(11):1873–1885, 2022.

416 [27] Zongxin Yu, Yikai Liu, Guang Lin, Wen Jiang, and Ming Chen. ESMAdam: a plug-and-play
417 all-purpose protein ensemble generator. *bioRxiv*, pages 2025–01, 2025.

418 [28] Bowen Jing, Bonnie Berger, and Tommi Jaakkola. Alphafold meets flow matching for generating
419 protein ensembles. *arXiv preprint arXiv:2402.04845*, 2024.

- [29] Rohith Krishna, Jue Wang, Woody Ahern, Pascal Sturmfels, Preetham Venkatesh, Indrek Kalvet, Gyu Rie Lee, Felix S Morey-Burrows, Ivan Anishchenko, Ian R Humphreys, et al. Generalized biomolecular modeling and design with rosettafold all-atom. *Science*, 384(6693):eadl2528, 2024.
- [30] Bowen Jing, Ezra Erives, Peter Pao-Huang, Gabriele Corso, Bonnie Berger, and Tommi Jaakkola. Eigenfold: Generative protein structure prediction with diffusion models. *arXiv preprint arXiv:2304.02198*, 2023.
- [31] John B Ingraham, Max Baranov, Zak Costello, Karl W Barber, Wujie Wang, Ahmed Ismail, Vincent Frappier, Dana M Lord, Christopher Ng-Thow-Hing, Erik R Van Vlack, et al. Illuminating protein space with a programmable generative model. *Nature*, 623(7989):1070–1078, 2023.
- [32] Junqi Liu, Shaoning Li, Chence Shi, Zhi Yang, and Jian Tang. Design of ligand-binding proteins with atomic flow matching. *arXiv preprint arXiv:2409.12080*, 2024.
- [33] Shuxin Zheng, Jiyan He, Chang Liu, Yu Shi, Ziheng Lu, Weitao Feng, Fusong Ju, Jiaxi Wang, Jianwei Zhu, Yaosen Min, He Zhang, Shidi Tang, Hongxia Hao, Peiran Jin, Chi Chen, Frank Noé, Haiguang Liu, and Tie-Yan Liu. Predicting equilibrium distributions for molecular systems with deep learning. *Nature Machine Intelligence*, 6(5):558–567, 2024. doi: 10.1038/s42256-024-00837-3. URL <https://doi.org/10.1038/s42256-024-00837-3>.
- [34] Frank Noé, Simon Olsson, Jonas Köhler, and Hao Wu. Boltzmann generators: Sampling equilibrium states of many-body systems with deep learning. *Science*, 365(6457):eaaw1147, 2019.
- [35] Tong Wang, Xinheng He, Mingyu Li, Yatao Li, Ran Bi, Yusong Wang, Chaoran Cheng, Xiangzhen Shen, Jiawei Meng, He Zhang, et al. Ab initio characterization of protein molecular dynamics with ai2bmd. *Nature*, pages 1–9, 2024.
- [36] Sergei Grudinin, Elodie Laine, and Alexandre Hoffmann. Predicting protein functional motions: an old recipe with a new twist. *Biophysical journal*, 118(10):2513–2525, 2020.
- [37] Alexandre Hoffmann and Sergei Grudinin. Nolib: Nonlinear rigid block normal-mode analysis method. *Journal of chemical theory and computation*, 13(5):2123–2134, 2017.
- [38] Steven Hayward and Nobuhiro Go. Collective variable description of native protein dynamics. *Annual review of physical chemistry*, 46(1):223–250, 1995.
- [39] Elodie Laine and Sergei Grudinin. Hopma: Boosting protein functional dynamics with colored contact maps. *The Journal of Physical Chemistry B*, 125(10):2577–2588, 2021.
- [40] Venkata K Ramaswamy, Samuel C Musson, Chris G Willcocks, and Matteo T Degiacomi. Deep learning protein conformational space with convolutions and latent interpolations. *Physical Review X*, 11(1):011052, 2021.
- [41] Haochuan Chen, Benoît Roux, and Christophe Chipot. Discovering reaction pathways, slow variables, and committor probabilities with machine learning. *Journal of Chemical Theory and Computation*, 19(14):4414–4426, 2023.
- [42] Zineb Belkacemi, Paraskevi Gkeka, Tony Lelièvre, and Gabriel Stoltz. Chasing collective variables using autoencoders and biased trajectories. *Journal of chemical theory and computation*, 18(1):59–78, 2021.
- [43] Luigi Bonati, Giovanni Maria Piccini, and Michele Parrinello. Deep learning the slow modes for rare events sampling. *Proceedings of the National Academy of Sciences*, 118(44):e2113533118, 2021.
- [44] Yihang Wang, Joao Marcelo Lamim Ribeiro, and Pratyush Tiwary. Machine learning approaches for analyzing and enhancing molecular dynamics simulations. *Current opinion in structural biology*, 61:139–145, 2020.

- [45] João Marcelo Lamim Ribeiro, Pablo Bravo, Yihang Wang, and Pratyush Tiwary. Reweighted autoencoded variational bayes for enhanced sampling (rave). *The Journal of chemical physics*, 149(7), 2018.
- [46] John Ingraham, Vikas Garg, Regina Barzilay, and Tommi Jaakkola. Generative models for graph-based protein design. *Advances in neural information processing systems*, 32, 2019.
- [47] Bowen Jing, Stephan Eismann, Patricia Suriana, Raphael JL Townshend, and Ron Dror. Learning from protein structure with geometric vector perceptrons. *arXiv preprint arXiv:2009.01411*, 2020.
- [48] Justas Dauparas, Ivan Anishchenko, Nathaniel Bennett, Hua Bai, Robert J Ragotte, Lukas F Milles, Basile IM Wicky, Alexis Courbet, Rob J de Haas, Neville Bethel, et al. Robust deep learning-based protein sequence design using proteinmpnn. *Science*, 378(6615):49–56, 2022.
- [49] Lucien F Krapp, Luciano A Abriata, Fabio Cortés Rodriguez, and Matteo Dal Peraro. Pesto: parameter-free geometric deep learning for accurate prediction of protein binding interfaces. *Nature communications*, 14(1):2175, 2023.
- [50] Yusong Wang, Tong Wang, Shaoning Li, Xinheng He, Mingyu Li, Zun Wang, Nanning Zheng, Bin Shao, and Tie-Yan Liu. Enhancing geometric representations for molecules with equivariant vector-scalar interactive message passing. *Nature Communications*, 15(1):313, 2024.
- [51] Helen M Berman, John Westbrook, Zukang Feng, Gary Gilliland, Talapady N Bhat, Helge Weissig, Ilya N Shindyalov, and Philip E Bourne. The protein data bank. *Nucleic acids research*, 28(1):235–242, 2000.
- [52] W. Kabsch. A solution for the best rotation to relate two sets of vectors. *Acta Crystallographica Section A*, 32(5):922–923, Sep 1976. doi: 10.1107/S0567739476001873. URL <https://doi.org/10.1107/S0567739476001873>.
- [53] Simon K Kearsley. On the orthogonal transformation used for structural comparisons. *Acta Crystallographica Section A: Foundations of Crystallography*, 45(2):208–210, 1989.
- [54] Andrea Amadei, Marc A Ceruso, and Alfredo Di Nola. On the convergence of the conformational coordinates basis set obtained by the essential dynamics analysis of proteins’ molecular dynamics simulations. *Proteins: Structure, Function, and Bioinformatics*, 36(4):419–424, 1999.
- [55] Alejandra Leo-Macias, Pedro Lopez-Romero, Dmitry Lupyan, Daniel Zerbino, and Angel R Ortiz. An analysis of core deformations in protein superfamilies. *Biophysical journal*, 88(2):1291–1299, 2005.
- [56] Charles C David and Donald J Jacobs. Characterizing protein motions from structure. *Journal of Molecular Graphics and Modelling*, 31:41–56, 2011.
- [57] Michael Heinzinger, Konstantin Weissenow, Joaquin Gomez Sanchez, Adrian Henkel, Martin Steinegger, and Burkhard Rost. ProStt5: Bilingual language model for protein sequence and structure. *bioRxiv*, pages 2023–07, 2023.
- [58] . ESM Team. Esm cambrian: Revealing the mysteries of proteins with unsupervised learning. Evolutionary Scale Website, 2024. URL <https://evolutionaryscale.ai/blog/esm-cambrian>. Available from: <https://evolutionaryscale.ai/blog/esm-cambrian>.
- [59] Robbie P Joosten, Fei Long, Garib N Murshudov, and Anastassis Perrakis. The pdb_redo server for macromolecular structure model optimization. *IUCrJ*, 1(4):213–220, 2014.
- [60] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=Bkg6RiCqY7>.
- [61] Gustaf Ahndritz, Nazim Bouatta, Christina Floristean, Sachin Kadyan, Qinghui Xia, William Gerecke, Timothy J O’Donnell, Daniel Berenberg, Ian Fisk, Niccolò Zanichelli, et al. Openfold: Retraining alphafold2 yields new insights into its learning mechanisms and capacity for generalization. *Nature Methods*, 21(8):1514–1524, 2024.

- 516 [62] Martin Steinegger and Johannes Söding. Mmseqs2 enables sensitive protein sequence searching 1
517 for the analysis of massive data sets. *Nature Biotechnology*, 35(11):1026–1028, Nov 2017.
518 ISSN 1546-1696. doi: 10.1038/nbt.3988. URL <https://doi.org/10.1038/nbt.3988>.
- 519 [63] Yang Zhang and Jeffrey Skolnick. Tm-align: a protein structure alignment algorithm based on
520 the tm-score. *Nucleic acids research*, 33(7):2302–2309, 2005.

521 Appendices

522 A Invariance of the proposed losses ¹

523 **Theorem A.1.** *SS Loss is invariant under unitary transformations of X and Y subspaces.* ²

524 *Proof.* Without loss of generality, let us assume that we apply a unitary transformation $U \in \mathbb{R}^{K \times K}$ ³
 525 to a subspace $X^\perp \in \mathbb{R}^{3N \times K}$, such that the result $X' = X^\perp U$, with $X' \in \mathbb{R}^{3N \times K}$, spans the same
 526 subspace as $X^\perp$, as it is a linear combination of the original basis vectors from $X^\perp$. Then, let us
 527 rewrite the SS loss as

$$\text{SS Loss} = 1 - \frac{1}{K} \sum_{k=1}^K \sum_{l=1}^K (\mathbf{y}_k^T W^{\frac{1}{2}} \mathbf{x}_l^\perp)^2 = 1 - \frac{1}{K} \|Y^T W^{\frac{1}{2}} X^\perp\|_F^2. \quad (A.1) \quad 4$$

528 As the Frobenius matrix norm is invariant under orthogonal, or more generally, unitary, transforma- ⁵
 529 tions, $\|Y^T W^{\frac{1}{2}} X^\perp U\|_F^2 = \|Y^T W^{\frac{1}{2}} X^\perp\|_F^2$, which completes the proof. $\square$

530 **Corollary A.1.1.** *The SS loss is invariant to the direction permutations in the Gram-Schmidt* ⁶
 531 *orthogonalization process.*

532 *Proof.* Let us consider two linear subspaces $X_1^\perp$ and $X_2^\perp$ resulting from the Gram-Schmidt orthogo- ⁷
 533 nalization of X , where we arbitrarily choose the order of the orthogonalization vectors. Both $X_1^\perp$
 534 and $X_2^\perp$ will span the same subspace as X , and since both $X_1^\perp$ and $X_2^\perp$ are also orthogonal, one is a
 535 unitary transformation of the other, $X_2^\perp = X_1^\perp U$, which completes the proof. $\square$

536 **Theorem A.2.** *IS Loss is invariant under unitary transformations of X and Y subspaces.* ⁸

537 *Proof.* Following the previous proof, without loss of generality, let us assume that we apply an ⁹
 538 orthogonal (unitary) transformation $U \in \mathbb{R}^{K \times K}$ to a subspace $X \in \mathbb{R}^{3N \times K}$, such that the result
 539 $X' = XU$, with $X' \in \mathbb{R}^{3N \times K}$, spans the same subspace as X . Then, let us rewrite the IS loss as

$$\text{IS Loss} = \frac{1}{K^2} \sum_{k=1}^K \sum_{l=1}^K (\mathbf{x}_k^T W \mathbf{x}_l)^2 - \frac{1}{K^2} \sum_{k=1}^K \sum_{l=1}^K (\mathbf{y}_k^T W^{\frac{1}{2}} \mathbf{x}_l)^2 = \frac{1}{K^2} \|X^T W X\|_F^2 - \frac{1}{K^2} \|Y^T W^{\frac{1}{2}} X\|_F^2. \quad (A.2) \quad 10$$

540 As the Frobenius matrix norm is invariant under orthogonal transformations, $\|Y^T W^{\frac{1}{2}} XU\|_F^2 =$ ¹¹
 541 $\|Y^T W^{\frac{1}{2}} X\|_F^2$, and $\|U^T X^T W X U\|_F^2 = \|X^T W X\|_F^2$, which completes the proof. $\square$

542 B Methods details ¹²

543 B.1 Training data ¹³

544 **Conformational collections.** To generate the training data, we utilized DANCE [15] to construct ¹⁴
 545 a non-redundant set of conformational collections representing the entire PDB as of June 2023.
 546 Wherever possible, we enhanced the data quality by replacing raw PDB coordinates with their
 547 updated and optimized counterparts from PDB-REDO [59]. Each conformational collection was
 548 designed to include only closely related homologs, ensuring that any two protein chains within the
 549 same collection shared at least 80% sequence identity and coverage. Collections with insufficient
 550 data points were excluded as we require at least 5 conformations. To simplify the data, we retained
 551 only C α atoms (option -c) and accounted for coordinate uncertainty by applying weights (option
 552 -w).

553 **Handling missing data.** The conformations in a collection may have different lengths reflected ¹⁵
 554 by the introduction of gaps when aligning the amino acid sequences. We fill these gaps with the
 555 coordinates of the conformation used to center the data. In doing so, we avoid introducing biases
 556 through reconstruction of the missing coordinates. Moreover, to explicitly account for data uncertainty,

we assign confidence scores to the residues and include them in the structural alignment step and the eigen decomposition. The confidence score of a position i reflects its coverage in the alignment,

$$w_i = \frac{1}{m} \sum_S 1_{a_i^S \neq "X"}, \quad (B.1)$$

where "X" is the symbol used for gaps and m is the number of conformations. The structural alignment of the j th conformation onto the reference conformation amounts to determining the optimal rotation that minimises the following function [52, 53],

$$E = \frac{1}{\sum_i w_i} \sum_i w_i (r_{ij}^c - r_{i0}^c)^2, \quad (B.2)$$

where r_{ij}^c is the i th centred coordinate of the j th conformation and r_{i0}^c is the i th centred coordinate of the reference conformation. The resulting aligned coordinates are then multiplied by the confidence scores prior to the PCA, as we explain below.

Eigenspaces of positional covariance matrices. The Cartesian coordinates of each conformational ensemble can be stored in a matrix R of dimension $3N \times m$, where N is the number of residues (or positions in the associated multiple sequence alignment) and n is the number of conformations. Each position is represented by a C- α atom. We compute the coverage-weighted (to account for missing data, as explained above) covariance matrix as in Eq. 2. The covariance matrix is a $3N \times 3N$ square matrix, symmetric and real.

We decompose C as $C = VDV^T$, where V is a $3N \times 3N$ matrix with each column defining a sqrt-coverage-weighted eigenvector or a principal component that we interpret as a *linear motion*. D is a diagonal matrix containing the eigenvalues. Specifically, the k th principal component was expressed as a set of 3D (sqrt-coverage-weighted) displacement vectors $\tilde{x}_{ik}^{GT}$, $i = 1, 2, \dots, L$ for the L C α atoms of the protein residues. To enable cross-protein comparisons, the vectors were normalized such that $\sum_i i = 1^L |\tilde{x}_{ik}^{GT}|^2 = L$. The sum of the eigenvalues $\sum_{k=1}^{3m} \lambda_k$ amounts to the total positional variance of the ensemble (measured in $\text{\AA}^2$) and each eigenvalue reflects the amount of variance explained by the associated eigenvector.

Data augmentation. The reference conformation used to align and center the 3D coordinates corresponds to the protein chain with the most representative amino acid sequence. To increase data diversity, four additional reference conformations were defined for each collection. At each iteration, the new reference conformation was selected as the one with the highest RMSD relative to the previous reference. This iterative strategy maximizes the variability of the extracted motions by emphasizing the impact of changing the reference.

B.2 Message passing

The node embeddings and predicted motion vectors are updated iteratively according to the following algorithm.

Algorithm B.1 PETIMOT Message Passing Block

```

1: function MESSAGEPASSING( $\{\mathbf{s}_i\}, \{\vec{x}_i\}, \{\mathcal{N}eigh(i)\}, \{R_{ij}, e_{ij}\}$ ):
2:   #  $\{\mathbf{s}_i\}_{i=1}^N$                                 ▷ Node embeddings
3:   #  $\{\vec{x}_i\}_{i=1}^N$                                 ▷ Motion vectors in local frames
4:   #  $\{\mathcal{N}eigh(i)\}_{i=1}^N$                         ▷ Node neighborhoods
5:   #  $\{R_{ij}, e_{ij}\}$                                 ▷ Relative geometric features
6:   for  $i = 1$  to  $N$  do
7:     for  $j \in \mathcal{N}eigh(i)$  do
8:        $\vec{x}_j^i \leftarrow R_{ij} \vec{x}_j$                                 ▷ Project motion in frame  $i$ 
9:        $m_{ij} \leftarrow \text{MessageMLP}(\mathbf{s}_i, \mathbf{s}_j, \vec{x}_i, \vec{x}_j^i, e_{ij})$ 
10:    end for
11:     $m_i \leftarrow \text{Mean}_j(m_{ij})$                                 ▷ Aggregate messages
12:     $\mathbf{s}_i \leftarrow \mathbf{s}_i + \text{LayerNorm}(m_i)$                                 ▷ Update embedding
13:     $\vec{x}_i \leftarrow \vec{x}_i + \text{Linear}([\mathbf{s}_i, \vec{x}_i])$                                 ▷ Update motion
14:  end for
15:  return  $\{\mathbf{s}_i\}_{i=1}^N, \{\vec{x}_i\}_{i=1}^N$ 
16: end function

```

588 **B.3 SE(3)-equivariant features** 2

589 We represent protein structures as attributed graphs. The node embeddings are computed with the 3
590 pre-trained protein language model ProST5 [57]. It is a fine-tuned version of the sequence-only model
591 T5 that translates amino acid sequences into sequences of discrete structural states and reciprocally.

592 The edge embeddings are computed using SE(3)-invariant features derived from the input backbone, 4
593 similarly to prior works [31, 48, 46]. Specifically, the features associated with the edge e_{ij} from node
594 (atom) i to node (atom) j are:

- 595 • **Quaternion representation:** A 4-dimensional quaternion encoding the relative rotation R_{ij} 5
596 between the local reference frames of residues i and j .
- 597 • **Relative translation:** A 3-dimensional vector representing the translation $\vec{t}_{ij}$ between the
598 local reference frames.
- 599 • **Chain separation:** The sequence separation between residues i and j , encoded as $\log(|i -$
600 $j| + 1)$.
- 601 • **Spatial separation:** The logarithm of the Euclidean distance between residues i and j ,
602 computed as $\log(\|\vec{t}_{ij}\| + \epsilon)$, where $\epsilon = 10^{-8}$.
- 603 • **Backbone atoms distances:** Distances between all backbone atoms (N, C α , C, O) at
604 residues i and j , encoded through a radial basis expansion. For each pairwise distance d_{ab} ,
605 we compute:

$$f_k(d_{ab}) = \exp\left(-\frac{(d_{ab} - \mu_k)^2}{2\sigma^2}\right), \quad 6 \quad (\text{B.3})$$

606 where $\{\mu_k\}_{k=1}^{20}$ are centers spaced linearly in $[0, 20]$ Å and $\sigma = 1$ Å. This creates a $16 \times$ 7
607 $20 = 320$ dimensional feature vector, as we have 16 pairwise distances (4×4 atoms) each
608 expanded in 20 basis functions.

609 **B.4 Training procedure** 8

610 We randomly split the 7,335 conformational collections defined with DANCE into training, validation, 9
611 and test sets with a ratio 70:15:15. The data augmentation procedure resulted in $5,119 \times 5 = 25,595$
612 training samples and $1,099 \times 5 = 5,495$ validation samples.

613 The model was optimized using AdamW [60] with a learning rate of 5e-4 and weight decay of 0.01. 10
614 We employed gradient clipping with a maximum norm of 10.0 and mixed precision training with
615 PyTorch’s Automatic Mixed Precision. The learning rate was adjusted using torch’s ReduceLROn-
616 Plateau scheduler, which monitored the validation loss, reducing the learning rate by a factor of
617 0.2 after 10 epochs without improvement. Training was performed with a batch size of 32 for both

training and validation sets. We implemented early stopping with a patience of 50 epochs, monitoring the validation loss. The model achieving the best validation performance was selected for final evaluation. We trained the model on a single NVIDIA A100-SXM4-80GB GPU. One epoch took about 9 minutes of real time.

B.5 Evaluation procedures

We report all evaluation metrics on 824 test proteins out of a total of 1 117. The 293 protein we excluded were too long (>450 amino acids) to be handled by AlphaFlow and ESMFlow in a reasonable amount of time using our computing resources.

Comparison with AlphaFlow and ESMFlow. We compared our approach with the flow-matching based frameworks AlphaFlow and ESMFlow for generating conformational ensembles. For this, we downloaded the distilled "PDB" models from <https://github.com/bjing2016/alphafold>. We executed AlphaFlow using the following command,

```
python predict.py --noisy_first --no_diffusion --mode alphafold
--input_csv seqs.csv --msa_dir msa_dir/
--weights alphafold_pdb_distilled_202402.pt --samples 50
--outpdb output_pdb/
```

AlphaFlow relies on OpenFold [61] to retrieve the input multiple sequence alignment (MSA). ESMFlow was launched using the same command with an additional `--mode esmfold` flag and its corresponding weights. We used AlphaFlow and ESMFlow to generate 50 conformations for each test protein and then we treated each ensemble as a conformational collection. We then aligned all members of the created collections to the reference conformations of the ground-truth collections. We used the identity coverage weights here. Finally, from the aligned collections, we extracted the principal linear motions. We shall additionally mention that we did not filter or adapt our test set to the AlphaFlow and ESMFlow methods. In other words, there can be certain data leakage between AlphaFlow/ESMFlow train data and our test examples.

Comparison with the Normal Mode Analysis. We also compared our approach with the physics-based unsupervised Normal Mode Analysis (NMA) method [38]. The NMA takes as input a protein 3D structure and builds an elastic network model where the nodes represent the atoms and the edges represent springs linking atoms located close to each other in 3D space. The four lowest normal modes are obtained by diagonalizing the mass-weighted Hessian matrix of the potential energy of this network. We used the highly efficient NOLB method, version 1.9, downloaded from <https://team.inria.fr/nano-d/software/nolb-normal-modes/> [37] to extract the first K normal modes from the test protein 3D conformations. Specifically, we used the following command

```
NOLB INPUT.pdb -c 10 -x -n 4 --linear -s 0 --format 1 --hetatm
```

We retained only the $C\alpha$ atoms and defined the edges in the elastic network using a distance cutoff of 10 Å.

B.6 Generalisability assessment across protein families

Protein folds are not evenly represented in the PDB, with some folds being significantly more abundant than others. For our default training procedure, we chose not to correct for this bias because the same fold may exhibit different motions in different collections. This represents a fundamental difference from structure prediction tasks, where distant homologs often exhibit nearly identical 3D structures. In contrast, for protein motion prediction, even proteins with the same fold can display dramatically different conformational changes depending on their specific biological context, binding partners, or experimental conditions. Furthermore, the diversity of motions within a single fold family provides valuable information for our model to learn the relationship between sequence, structure, and dynamic behavior.

Nevertheless, to rigorously assess PETIMOT's capability to generalise across protein families, we re-trained and re-evaluated the model using a more stringent and non-redundant training-validation-test splitting scheme. Specifically, we considered clusters of protein chains defined at 30% sequence

identity and 80% sequence coverage using MMseqs2 [62]. These clusters define distant protein families and we refer to them as clus-30 in the following. We kept the 824 test proteins used for comparison with other methods as is and we removed all collections belonging to the same clus-30 clusters as these proteins from the training and validation sets. Then, we re-defined a training-validation random split at the level of the clus-30 clusters with a 9:1 ratio. This operation ensures that any pair of training-validation, training-test or validation-test collections do not share more than 30% sequence identity. Finally, for each training or validation clus-30 cluster, we randomly drew 5 samples. This step ensures that each protein family is evenly represented in the training and validation sets. We also redundancy reduced the test set of 824 proteins (test-824), keeping only one protein for each clus-30 cluster, which led to 734 proteins (test-734).

To be able to compare the model’s default version, *PETIMOT-default*, and re-trained version, *PETIMOT-stringent*, we computed their performance metrics on a subset of 474 test proteins from test-734. These proteins do not fall in any of the clus-30 clusters used for training or validation any of the two versions.

B.7 Ablation Studies

To understand the impact of different components on the performance of our model, we carried out ablation studies. We list them blow.

Model architecture variations.

- Network depth: We experimented with different numbers of message-passing layers (5 and 10 layers compared to our default value of 15 layers).
- Layer sharing: We tested a variant where all message-passing layers share the same parameters, as opposed to our default where each layer has unique parameters.
- Reduced internal embedding dimension: We tested a model with a smaller internal embedding dimension of 128 instead of the default 256.

Figure B.1 shows the evaluation of these modifications. A shallow 5-layers network underperforms on all evaluation metrics. The difference between other variants is not very significant.

Structure and sequence information ablation.

- Structure ablation: We removed all structural information from the model to assess the importance of geometric features and the performance with the PLM embeddings only. We did it by removing the edge attributes of the input of the message passing MLP.
- Sequence ablation: We ablated sequence information by replacing protein language model embeddings with random embeddings, testing them both with and without structural information.
- Embedding variants: We evaluated a different protein language model (ESMC-600M), both with and without structural tokens.

The evaluation results are shown in Fig. B.2. The results demonstrate that while both ProstT5 and ESM-Cambrian 600M perform similarly when combined with structural information, removing structural features leads to markedly different outcomes. ProstT5 embeddings partially compensate for the missing structural information, likely due to their structure-aware training, while relying solely on ESM-C embeddings results in poor performance.

Problem formulation ablation. We analyzed different combinations of our loss terms (compared to our default balanced weights of LS + SS):

- Least Square loss (LS): Using only the LS loss (weight 1.0).
- Squared Sinus loss (SS): Using only the SS loss (weight 1.0).
- Independent Subspaces (IS): Using only the IS loss (weight 1.0).

Figure B.3 compares three individual losses with the default option. The IS problem formulation underperforms on all the metrics. The default LS + SS formulation performs slightly better than those with individual loss components.




Figure B.1: Network depth ablation. We report cumulative curves for LS error (a-b), magnitude error (c-d), and SS error (e). For each protein, we computed the error either for the best-matching pair of predicted and ground-truth vectors (a,c) or for the best combination of four pairs of predicted and ground-truth vectors (b,d). We vary the number of layers in the network and the embedding dimension.

Graph connectivity ablation. We investigated different approaches to constructing the protein graph:

- Nearest neighbor-only: Using 15 nearest neighbors (sorted according to the corresponding $C\alpha-C\alpha$ distances) without random edges.
- Random connections-only: Using 15 random edges without nearest neighbors. This set is updated between every layer at each epoch.




Figure B.2: **Structure and sequence information ablation study.** We report cumulative curves for LS error (a-b), magnitude error (c-d), and SS error (e). For each protein, we computed the LS and magnitude errors either for the best-matching pair of predicted and ground-truth vectors (a,c) or for the best combination of four pairs of predicted and ground-truth vectors (b,d).

- Static connectivity: Using a fixed set of random neighbors between the layers. This set is updated at each epoch.

Figure B.4 shows the ablation results. We can see that the nearest neighbor-only setup underperforms on all the metrics. Among other options, the random connectivity-only option gets lower results at higher metrics values. The default option performs on par with the static connectivity, showing slightly better results on the optimal assignment magnitude error metrics.

C Additional results

Figure C.1 evaluates PETIMOT against NMA, ESMFlow and AlphaFlow approaches using additional metrics. These include the minimum magnitude error, the optimal assignment magnitude error, and the optimal assignment LS error. On all the metrics we see that PETIMOT outperforms the three other tested approaches.

We also experimented with a different number of predicted components. For these experiments, we trained additional models with the LS loss only, which are listed below:




Figure B.3: Performance comparison of different problem formulations. We report cumulative curves for magnitude error (a,b) and LS error (c). For each protein, we computed the error either for the best-matching pair of predicted and ground-truth vectors (a) or for the best combination of four pairs of predicted and ground-truth vectors (b,c).

- Single component prediction (1 mode).
- Reduced component prediction (2 modes).
- Extended component prediction (8 modes).

We compare these options with our default setting of 4 components. Figure C.2 shows the results. Increasing the number of predicted components from 1 to 8 improves the minimum LS errors, as having more predicted vectors naturally increases the chance of matching at least one ground-truth motion well. However, when evaluating the optimal linear assignment metrics, which measures overall subspace alignment, models with 1 or 2 components have an artificial advantage since they face fewer matching constraints. The 8-components model similarly benefits from having more candidate vectors to match against the 4 ground-truth components.

Figure C.3 compares the accuracy of the predicted test proteins (minimum LS loss) with the structural (TM-score) and sequence (sequence identity) distances to the training set. We do not see a clear correlation between the prediction accuracy and the similarity to the training examples. Please also see Fig. 3b-c for comparison.

Figure C.4 further assesses PETIMOT's generalisation capability across protein families. To do so, we implemented a more stringent train-validation-test split using a 30% sequence similarity threshold (see main text and Appendix B.6). When trained under this stringent protocol, PETIMOT still substantially outperforms all baseline methods according to minimum L and SS metrics (Fig. C.4a-b). Moreover, we observe a slight generalisation improvement of *PETIMOT-stringent* over

753 *PETIMOT-default* on a test set of 474 proteins evolutionary distant from any protein used for training 1
754 or validation of any of the two model versions (Fig. C.4c-d).

755 Figures C.5 and C.6 show predicted (blue arrows) and ground-truth (red arrows) motion vectors for 2
756 the xylanase A from *Bacillus subtilis* and the periplasmic domain of Gliding motility protein GldM
757 from *Capnocytophaga canimorsus*, respectively.

758 D Licenses for used resources 3

759 In this work, we utilize several existing resources. The protein structures were obtained from 4
760 the Protein Data Bank (PDB, <https://www.rcsb.org/>, version accessed on June 2023) which is dis-
761 tributed under the CC0 1.0 Universal Public Domain Dedication license (CC0 1.0). We com-
762 plemented PDB data with data from PDB-redo (accessed June 2023) developed by Joosten *et*
763 *al.* [59], available at <https://pdb-redo.eu> under the licence specified at <https://pdb-redo.eu/license>. For protein language modeling, we employed ProST5 developed by Heinzinger
764 *et al.* [57], available under the MIT license at <https://huggingface.co/Rostlab/ProST5>,
765 and ESM-Cambrian 600M (version esmc-600m-2024-12) developed by the EvolutionaryScale
766 Team [58], available under the Cambrian Non-Commercial License at [https://huggingface.co/](https://huggingface.co/EvolutionaryScale/esmc-600m-2024-12)
767 [EvolutionaryScale/esmc-600m-2024-12](https://huggingface.co/EvolutionaryScale/esmc-600m-2024-12). Additional resources include the DANCE method
768 (version of Oct 8, 2024) developed by Lombard *et al.*, available under the MIT license at
769 <https://github.com/PhyloSofS-Team/DANCE>. As baselines, we ran the NOLB method (ver-
770 sion 1.9) developed by Hoffmann *et al.* [37] and available at [https://team.inria.fr/nano-d/](https://team.inria.fr/nano-d/software/nolb-normal-modes/)
771 [software/nolb-normal-modes/](https://team.inria.fr/nano-d/software/nolb-normal-modes/), and the AlphaFlow (version AlphaFlow-PDB distilled) and
772 ESMFlow (version ESMFlow-PDB distilled) models developed by Jing *et al.* [28] and available at
773 <https://github.com/bjing2016/alphaflow>. We used TM-align (version 20220412) developed
774 by Zhang and Skolnick [63] and available at <https://zhanggroup.org/TM-align/> to perform
775 all-to-all pairwise structural alignments between train and test protein conformations and compute
776 TM-scores.
777




Figure B.4: Graph connectivity ablation. We report cumulative curves for LS error (a-b), magnitude error (c-d), and SS error (e). For each protein, we computed the error either for the best-matching pair of predicted and ground-truth vectors (a,c) or for the best combination of four pairs of predicted and ground-truth vectors (b,d). Only Random Neighbors: each residue (node) is connected to 15 randomly chosen residues and the connectivity changes after each layer. Only Nearest Neighbors: each residue (node) is connected to its 15 nearest neighbors in the input 3D structure. Fixed Random Connectivity: each residue (node) is connected to 15 residues randomly chosen at the beginning.




Figure C.1: **Performance comparison with other methods on the test proteins.** We report cumulative curves for magnitude error (a,b) and LS error (c). For each protein, we computed the error either for the best-matching pair of predicted and ground-truth vectors (a) or for the best combination of four pairs of predicted and ground-truth vectors (b,c).




Figure C.2: Impact of the number of predicted components. We report cumulative curves for LS error (a-b) and magnitude error (c-d). For each protein, we computed the error either for the best-matching pair of predicted and ground-truth vectors (a,c) or for the best combination of all pairs of predicted and ground-truth vectors using optimal linear assignment (b,d). We compare models trained to predict different numbers of components (modes): 1, 2, 4, or 8, using only the LS loss.




Figure C.3: **Relationship between PETIMOT's prediction accuracy and structural/sequence similarity with the training set.** The minimum LS error is plotted against the maximum TM-score between each test protein and any protein in the training set. Points are colored by the maximum sequence identity to the training samples.




Figure C.4: **Cumulative error curves computed on the test proteins.** **a-b.** Comparison between *PETIMOT-stringent* model and three other methods on a non-redundant set of 734 test proteins. *PETIMOT-stringent* was trained on a stringent and non-redundant training-validation-test split defined using a 30% sequence identity threshold. **c-d.** Comparison between *PETIMOT-default*, *PETIMOT-stringent* and three other methods on 474 test proteins that share less than 30% sequence similarity with any protein used in training or validation of either model. **a,c.** Minimum LS error corresponding to the best matching pair of predicted and ground-truth motions. **b,d.** SS error computed between the entire predicted and ground-truth subspaces.




Figure C.5: **Visualization of predicted (blue arrows) and ground-truth (red arrows) motion vectors for PDB structure 3EXU (chain A), with LS error of 0.20.** The predicted deformation was used to generate the interpolated conformations shown in Fig. 3b.




Figure C.6: **Visualization of predicted (blue arrows) and ground-truth (red arrows) motion vectors for PDB structure 7SD2, with LS error of 0.18.** The predicted deformation was used to generate the interpolated conformations shown in Fig. 3c.

NeurIPS Paper Checklist

1. Claims 1

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? 2

Answer: [Yes] 3

Justification: We provide a novel formulation of the protein conformational diversity problem, sections 3.1 and 3.2. We present a novel benchmark representative of the PDB structural diversity along data- and task-specific metrics in the result section and SI. We develop a SE(3)-equivariant Graph Neural Network architecture equipped with a novel symmetry-aware loss function for comparing linear subspaces, with invariance to permutation and scaling, sections 3.3 and 3.4. Our results demonstrate the performance of PETIMOT with respect to other methods, see Results, and the capability of PETIMOT to generalise across protein families, subsection Generalisation capability in the Results section. 4

Guidelines: 5

- The answer NA means that the abstract and introduction do not include the claims made in the paper. 6
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations 7

Question: Does the paper discuss the limitations of the work performed by the authors? 8

Answer: [Yes] 9

Justification: We have written a dedicated subsection "Limitations" in the Results section. 10

Guidelines: 11

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper. 12
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs 1

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof? 2

Answer: [Yes] 3

Justification: Proof of the invariance of the losses is available in Appendix A. 4

Guidelines: 5

- The answer NA means that the paper does not include theoretical results. 6
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility 7

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)? 8

Answer: [Yes] 9

Justification: The methods are explained in the the Methods sections and supplementary details are given in the Methods Details section in Appendix B. 10

Guidelines: 11

- The answer NA means that the paper does not include experiments. 12
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code 1

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material? 2

Answer: [Yes] 3

Justification: Code and instructions to retrieve data are available at <https://anonymous.4open.science/r/PETIMOT-4ED4/>. 4

Guidelines: 5

- The answer NA means that paper does not include experiments requiring code. 6
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details 7

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results? 8

Answer: [Yes] 9

Justification: Essential details are given in the main text, other information can be found in Appendix B. 10

Guidelines: 11

- The answer NA means that the paper does not include experiments. 12
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance 13

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments? 14

Answer: [Yes] 15

Justification: We provide standard deviations in Table 1. Additional statistical significance tests and error bars are provided in SI. 16

Guidelines: 17

- The answer NA means that the paper does not include experiments. 18
- The authors should answer “Yes” if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources 2

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments? 3

Answer: [Yes] 4

Justification: We provide information on the running time, CPU and GPU resources for each model in Tables and in the main text. 5

Guidelines: 6

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics 8

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>? 9

Answer: [Yes] 10

Justification: We followed the NeurIPS Code of Ethics when conducting our study and presenting the work. 11

Guidelines: 12

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts 14

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed? 15

Answer: [Yes] 16

Justification: We have written a dedicated section "Broader Impacts" that concludes the manuscript. 17

Guidelines: 18

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards 2

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)? 3

Answer: [NA] 4

Justification: We do not see any misuse of our model or theoretical results. 5

Guidelines: 6

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets 8

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected? 9

Answer: [Yes] 10

Justification: We have written a dedicated section in Appendix D "Licenses for used resources". 11

Guidelines: 12

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets 2

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets? 3

Answer: [Yes] 4

Justification: Code, instructions to retrieve data, and data splitting scripts are available at <https://anonymous.4open.science/r/PETIMOT-4ED4/>. 5

Guidelines: 6

- The answer NA means that the paper does not release new assets. 7
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects 8

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)? 9

Answer: [NA] 10

Justification: 11

Guidelines: 12

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 13
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects 14

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained? 15

Answer: [NA] 16

Justification: 17

Guidelines: 18

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 19

- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage 2

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA] 4

Justification: 5

Guidelines: 6

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.