
Dominated Actions in Imperfect-Information Games¹

Anonymous Author(s)²

Affiliation

Address

email

Abstract³

1 Dominance is a fundamental concept in game theory. In strategic-form games dom-⁴
2 inated strategies can be identified in polynomial time. As a consequence, iterative
3 removal of dominated strategies can be performed efficiently as a preprocessing
4 step for reducing the size of a game before computing a Nash equilibrium. For
5 imperfect-information games in extensive form, we could convert the game to
6 strategic form and then iteratively remove dominated strategies in the same way;
7 however, this conversion may cause an exponential blowup in game size. In this pa-
8 per we define and study the concept of dominated actions in imperfect-information
9 games. Our main result is a polynomial-time algorithm for determining whether an
10 action is dominated (strictly or weakly) by any mixed strategy in n -player games,
11 which can be extended to an algorithm for iteratively removing dominated actions.
12 This allows us to efficiently reduce the size of the game tree as a preprocessing
13 step for Nash equilibrium computation. We explore the role of dominated actions
14 empirically in the “All In or Fold” No-Limit Texas Hold’em poker variant.

15 1 Introduction⁵

16 In a strategic-form game, (mixed) strategy σ_i for player i is *strictly dominated* if there exists another⁶
17 (mixed) strategy σ'_i such that σ'_i performs strictly better than σ_i regardless of the strategy used by
18 the opponent(s): formally, if $u_i(\sigma'_i, s_{-i}) > u_i(\sigma_i, s_{-i})$ for all pure strategy profiles (vector of pure
19 strategies) $s_{-i} \in S_{-i}$ for the opposing agents. (Note that the requirement that this holds for all
20 opposing pure strategy profiles is enough to ensure that it also holds for all mixed strategy profiles of
21 the opponents as well). Clearly it would be irrational for an agent to play a strictly dominated strategy,
22 as another strategy will do strictly better regardless of the beliefs of what the opponents would play.
23 Strategy σ'_i *weakly dominates* σ_i if the inequality holds weakly (though strictly for at least one pure
24 strategy profile): formally, $u_i(\sigma'_i, s_{-i}) \geq u_i(\sigma_i, s_{-i})$ for all $s_{-i} \in S_{-i}$ where the inequality is strict
25 for at least one s_{-i} . (The condition that at least one inequality is strict is simply to rule out saying
26 that a strategy is dominated by an identical strategy). Similar to strict domination it seems clearly
27 irrational for an agent to play a strategy that is weakly dominated, as another strategy will perform at
28 least as well (and sometimes strictly better) regardless of the strategies employed by the opponent(s).

29 It seems natural to simplify a game by eliminating strategies that are dominated from the game⁷
30 to reduce its size and focus analysis on a smaller game. It can easily be shown that applying an
31 iterative process of removing one dominated strategy for one player, then removing one for another
32 (or possibly the same) player in the reduced game, etc., will ultimately result in a smaller game
33 that contains a Nash equilibrium from the original game. For this reason, this procedure of *iterated*
34 *removal of dominated strategies* is often performed as a preprocessing step to reduce the size of
35 a game before computing a Nash equilibrium (or other desired solution concept). As it turns out,
36 all processes of iterated removal of strictly dominated strategies produce the same reduced game
37 regardless of the order of elimination (while this is not necessarily the case for iterated removal of
38 weakly dominated strategies) [5, 8]. While iterated removal of weakly dominated strategies can

39 sometimes reduce the number of equilibria, it can never create new equilibria, and therefore even this 1
40 procedure is very useful as a preprocessing step for Nash equilibrium computation.

41 There is a linear-time algorithm for determining whether a (mixed) strategy σ_i is strictly dominated 2
42 by any pure strategy for player i [12]. This algorithm simply iterates over each pure strategy s_i for
43 player i and tests whether it performs strictly better than σ_i against each opposing pure strategy
44 profile s_{-i} . This procedure has complexity $O(|A|)$, where $A = \times_i A_i$ is the set of joint action (i.e.,
45 pure strategy) profiles, and so takes time linear in the size of the game. The procedure can also be
46 easily adapted to produce an algorithm for determining whether a mixed strategy profile is weakly
47 dominated by any pure strategy in linear time. Note that it is possible for a strategy to be dominated
48 by a mixed strategy and not be dominated by any pure strategy [12]. In order to determine whether a
49 (mixed) strategy is strictly dominated by a mixed strategy, while the above procedure does not work,
50 it turns out that there exists a linear programming formulation that runs in polynomial time, and there
51 also exists a linear programming formulation that determines whether a (mixed) strategy is weakly
52 dominated by a mixed strategy [1]. Therefore, regardless of whether the strategies being tested as
53 dominated or dominating are mixed or pure, it can be checked in polynomial time.¹

54 2 Extensive-form games 3

55 While the strategic form can be used to model simultaneous actions, another representation, called 4
56 the *extensive form*, is generally preferred when modeling settings that have sequential moves. The
57 extensive form can also model simultaneous actions, as well as chance events and imperfect informa-
58 tion (i.e., situations where some information is available to only some of the agents and not to others).
59 Extensive-form games consist primarily of a game tree; each non-terminal node has an associated
60 player (possibly *chance*) that makes the decision at that node, and each terminal node has associated
61 utilities for the players. Additionally, game states are partitioned into *information sets*, where the
62 player whose turn it is to move cannot distinguish among the states in the same information set.
63 Therefore, in any given information set, a player must choose actions with the same distribution at
64 each state contained in the information set. If no player forgets information that they previously knew,
65 we say that the game has *perfect recall*. A (behavioral) *strategy* for player i , $\sigma_i \in \Sigma_i$, is a function
66 that assigns a probability distribution over all actions at each information set belonging to i .

67 In theory, every extensive-form game can be converted to an equivalent strategic-form game; however, 5
68 there is an exponential blowup in the size of the game representation, and therefore such a conversion
69 is undesirable. Instead, new algorithms have been developed that operate on the extensive form
70 representation directly. It turns out that the complexity of computing equilibria in extensive-form
71 games is similar to that of strategic-form games; a Nash equilibrium can be computed in polynomial
72 time in two-player zero-sum games (with perfect recall) [7], while the problem is hard for two-
73 player general-sum and multiplayer games. One algorithm for computing an equilibrium in two-
74 player zero-sum extensive-form games with perfect recall is based on solving a linear programming
75 formulation [7]. This formulation works by modeling each sequence of actions for each player as a
76 variable, and is often called the *sequence form LP* algorithm. Note that while the number of pure
77 strategies is exponential in the size of the game tree, the number of action sequences is linear. The
78 method uses several matrices defined as follows. For player 1, the matrix $\mathbf{E}$ is defined where each row
79 corresponds to an information set (including an initial row for the “empty” information set), and each
80 column corresponds to an action sequence (including an initial row for the “empty” action sequence).
81 In the first row of $\mathbf{E}$ the first element is 1 and all other elements are 0; subsequent rows have -1 for
82 the entries corresponding to the action sequence leading to the root of the information set, and 1 for
83 all actions that can be taken at the information set (and 0 otherwise). Thus $\mathbf{E}$ has dimension $c_1 \times d_1$,
84 where c_i is the number of information sets for player i and d_i is the number of action sequences for
85 player i . Matrix $\mathbf{F}$ is defined analogously for player 2. The vector $\mathbf{e}$ is defined to be a column vector
86 of length c_1 with 1 in the first position and 0 in other entries, and vector $\mathbf{f}$ is defined with length c_2
87 analogously. The matrix $\mathbf{A}$ is defined with dimension $d_1 \times d_2$ where entry A_{ij} gives the expected
88 payoff for player 1 when player 1 plays action sequence d_1 and player 2 plays action sequence d_2

¹However, most other computational questions related to dominance can be solved in polynomial time for strict but are NP-hard for weak dominance. For example, determining whether there exists some elimination path under which strategy s_i is eliminated, given action subsets $A'_i \subset A_i$ determining whether there exists a maximally reduced game where each player has the actions A'_i , and given constants k_i determining whether there exists a maximally reduced game where each player has exactly k_i actions [12, 6].

(with the expectation being over possible moves of chance along the paths of play to leaf nodes). The matrix $\mathbf{B}$ of player 2's expected payoffs is defined analogously. In zero-sum games $\mathbf{B} = -\mathbf{A}$.

Given these matrices we can solve one of two linear programming problems to compute a Nash equilibrium in zero-sum extensive-form games [7]. In the first formulation the primal variables $\mathbf{x}$ correspond to player 1's mixed strategy while the dual variables correspond to player 2's strategy. In the second formulation, which is the dual problem of the first formulation, the primal decision variables $\mathbf{y}$ correspond to player 2's strategy while the dual variables correspond to player 1's strategy.

$$\begin{array}{ll} \max_{\mathbf{x}, \mathbf{q}} & -\mathbf{q}^T \mathbf{f} \\ \text{s.t.} & \mathbf{x}^T (-\mathbf{A}) - \mathbf{q}^T \mathbf{F} \leq \mathbf{0} \\ & \mathbf{x}^T \mathbf{E}^T = \mathbf{e}^T \\ & \mathbf{x} \geq \mathbf{0} \end{array}$$

$$\begin{array}{ll} \min_{\mathbf{y}, \mathbf{p}} & \mathbf{e}^T \mathbf{p} \\ \text{s.t.} & -\mathbf{A} \mathbf{y} + \mathbf{E}^T \mathbf{p} \geq \mathbf{0} \\ & -\mathbf{F} \mathbf{y} = -\mathbf{f} \\ & \mathbf{y} \geq \mathbf{0} \end{array}$$

3 Dominated actions

In extensive-form games, we can consider analogous concepts of strict, weak, and iterated dominance of strategies as for strategic-form games. However, unlike in the strategic-form setting, identification of a dominated extensive-form strategy does not necessarily allow us to reduce the size of the game, since it is possible that some of the actions played by the dominated strategy are also played by non-dominated strategies. In order to obtain the computational advantage of game size reduction, we must consider a stronger concept of *dominated actions*. We first present several plausible candidate definitions for dominated actions which we demonstrate to be problematic. Our first candidate definition is given as Flawed Definition 1. This definition states that action a_i for player i at information set I_i is strictly dominated by action b_i at the same information set if every leaf node succeeding a_i produces a strictly smaller payoff for player i than every leaf node succeeding action b_i . An analogous definition for weak dominance is given in Flawed Definition 2.

Flawed Definition 1. If for every leaf node N^{a_i} that follows action a_i for player i at information set I_i and every leaf node N^{b_i} that follows action b_i for player i at the same information set I_i , $u_i(N^{b_i}) > u_i(N^{a_i})$, then b_i strictly dominates a_i .

Flawed Definition 2. If for every leaf node N^{a_i} that follows action a_i for player i at information set I_i and every leaf node N^{b_i} that follows action b_i for player i at the same information set I_i , $u_i(N^{b_i}) \geq u_i(N^{a_i})$ where inequality is strict for at least one node, then b_i weakly dominates a_i .

The problem with these definitions are that they are too strong; it is still possible for player 1 to strictly prefer to take action b_i to a_i regardless of the strategy used by player 2 even if Flawed Definition 1 holds. This is illustrated in the following game depicted in Figure 1. In this game, chance makes an initial move taking each of two actions with probability $\frac{1}{2}$. Then player 1 (red) selects one of two actions at a single information set. Player 2 (blue) then takes one of two actions after observing both the moves of chance and player 1. It is clear that action 2 for player 2 at their top information set and action 1 for player 2 at their bottom information set are both strictly dominated according to Flawed Definition 1. The smaller game obtained after removing these actions is depicted in Figure 2. In the smaller game, the expected utility of playing action 1 is $0.5(-100) + 0.5(100) = 0$ and the expected utility of playing action 2 is $0.5(-50) + 0.5(-50) = -50$. Since player 2 does not take any actions, player 1 always achieves a strictly higher expected utility by playing action 1. However, action 2 is not strictly dominated according to Flawed Definition 1 because in one leaf node succeeding action 1 player 1 obtains payoff -100 which is lower than the payoff of -50 at leaf nodes following action 2.

Flawed Definitions 1 and 2 provide sufficient conditions for an action to be dominated, but we have demonstrated that it is clearly possible for actions that do not satisfy these definitions to also be dominated. Thus these conditions are too strong, and we will refer to strategies that satisfy them as being *strongly dominated* (strictly or weakly, respectively). The concept of strong dominance is not without merit, as it can be verified very efficiently by simply iterating over the leaf nodes




Figure 1: Example two-player imperfect-information extensive-form game. 2




Figure 2: Result of removing two dominated actions from game in Figure 1. 4

succeeding each action. Thus it may be useful to first remove actions that are strongly dominated as a preprocessing step before performing potentially more costly computations to remove other actions. 5

We next consider candidate definition given by Flawed Definition 3, where u_i denotes the expected utility accounting for randomized moves of chance as well as potential randomization in the players' strategies. This definition states that action a_i for player i at information set I_i is strictly dominated (potentially by a probability distribution over other actions at the same information set), if there exists a strategy $\sigma_i^{-a_i}$ that never plays a_i at I_i that always has strictly higher expected utility than every strategy that plays a_i at I_i . Again this definition clearly provides a sufficient condition for a_i to be dominated; however, the issue with this definition is that the strategies may potentially take actions early in the game tree that prevent the game from ever reaching information set I_i . Consider the simple example game in Figure 3. Suppose we want to apply Flawed Definition 3 to determine whether action 2 is strictly dominated for player 2. Then $\sigma_i^{-a_i}$ is the strategy for player 2 that plays action 1 with probability 1, and $\sigma_i^{a_i}$ must be the strategy for player 2 that plays action 2 with probability 1. Now suppose that player 1 plays action 2 with probability 1, which we denote as σ_{-i} . Then clearly $u_i(\sigma_i^{-a_i}, \sigma_{-i}) = u_i(\sigma_i^{a_i}, \sigma_{-i}) = 0$, since both strategy profiles will result in reaching the bottom leaf node yielding payoff 0. 6

Flawed Definition 3. Action a_i for player i is strictly dominated at information set I_i if there exists a mixed strategy $\sigma_i^{-a_i}$ that plays action a_i at I_i with probability 0 such that for every mixed strategy $\sigma_i^{a_i}$ for player i that plays action a_i with probability 1 at I_i , $u_i(\sigma_i^{-a_i}, \sigma_{-i}) > u_i(\sigma_i^{a_i}, \sigma_{-i})$ for all opposing strategies $\sigma_{-i} \in \Sigma_{-i}$. 7




Figure 3: Extensive-form game demonstrating problem with Flawed Definition 3. 2

The problem with Flawed Definition 3 is that it allows the players to deviate from the path of play leading to the relevant information set I_i . We address this limitation in our new definitions. *Strictly-dominated actions* are defined in Definition 1 and *weakly-dominated actions* are defined in Definition 2. In these definitions $\Sigma_{-i}^{I_i}$ denotes the set of mixed strategy profiles for the opponents that always take actions leading to information set I_i when possible. Note that these definitions apply to games with any number of players. They also consider actions that are dominated by any mixed strategy (not necessarily just a pure action at I_i). It is easy to verify that these definitions address the issues that arose in the examples from Figure 1 and Figure 3. We can apply these definitions repeatedly in succession to perform iterative removal of dominated actions, as for the strategic form.

Definition 1. Action a_i for player i is strictly dominated at information set I_i if there exists a mixed strategy $\sigma_i^{-a_i}$ that always plays to get to I_i and plays action a_i at I_i with probability 0 such that for every mixed strategy $\sigma_i^{a_i}$ for player i that always plays to get to I_i and plays action a_i with probability 1 at I_i , $u_i(\sigma_i^{-a_i}, \sigma_{-i}) > u_i(\sigma_i^{a_i}, \sigma_{-i})$ for all opposing strategies $\sigma_{-i} \in \Sigma_{-i}^{I_i}$.

Definition 2. Action a_i for player i is weakly dominated at information set I_i if there exists a mixed strategy $\sigma_i^{-a_i}$ that always plays to get to I_i and plays action a_i at I_i with probability 0 such that for every mixed strategy $\sigma_i^{a_i}$ for player i that always plays to get to I_i and plays action a_i with probability 1 at I_i , $u_i(\sigma_i^{-a_i}, \sigma_{-i}) \geq u_i(\sigma_i^{a_i}, \sigma_{-i})$ for all opposing strategies $\sigma_{-i} \in \Sigma_{-i}^{I_i}$, where the inequality is strict for at least one σ_{-i} .

The example games considered were created using the open-source software package Gambit [11]. Gambit has tools that allow the user to remove “strictly dominated” or “strictly or weakly dominated” actions, and these procedures can be repeated to iteratively remove multiple actions sequentially; however, there is no documentation regarding the algorithms applied or definitions of dominance used. In the example game from Figure 1, Gambit correctly identifies that action 2 for player 2 at their top information set and action 1 for player 2 at their bottom information set are both strictly dominated, and removes these to construct the smaller game in Figure 2. However, as for strong domination, Gambit fails to recognize that action 2 is strictly dominated by action 1 for player 1 in the reduced game. This example demonstrates that Gambit’s procedure does not remove all dominated actions (though it does not necessarily imply that Gambit is only removing strongly dominated actions).

4 Algorithm for identifying dominated actions 7

Suppose we want to determine whether an action c for player 1 is dominated at information set I in a two-player extensive-form game G . Using the sequence form representation, suppose that action c is the final action in the sequence with index i for player 1. Let S_I denote the set of indices corresponding to action sequences leading to I . We would like to solve the following problem.

$$\begin{aligned} \max_{\mathbf{x}_2} \min_{\mathbf{x}_1, \mathbf{y}} \quad & \mathbf{x}_2^T \mathbf{A} \mathbf{y} - \mathbf{x}_1^T \mathbf{A} \mathbf{y} \\ \text{s.t.} \quad & \mathbf{x}_{1,i} = 1 \\ & \mathbf{x}_{2,i} = 0 \\ & \mathbf{x}_{1,k} = \mathbf{x}_{2,k} = 1 \text{ for all } k \in S_I \\ & \mathbf{y}_k = 1 \text{ for all } k \in S_I \\ & \mathbf{x}_1^T \mathbf{E}^T = \mathbf{e}^T \\ & \mathbf{x}_2^T \mathbf{E}^T = \mathbf{e}^T \\ & \mathbf{y}^T \mathbf{F}^T = \mathbf{f}^T \\ & \{\mathbf{x}_1, \mathbf{x}_2, \mathbf{y}\} \geq \mathbf{0} \end{aligned} \tag{1}$$

185 Consider the following problem, where now player 2 controls two action sequences $\mathbf{y}_1, \mathbf{y}_2$: 1

$$\begin{aligned}
 \max_{\mathbf{x}_2} \min_{\mathbf{x}_1, \mathbf{y}_1, \mathbf{y}_2} \quad & \mathbf{x}_2^T \mathbf{A} \mathbf{y}_2 - \mathbf{x}_1^T \mathbf{A} \mathbf{y}_1 \\
 \text{s.t.} \quad & \mathbf{x}_{1,i} = 1 \\
 & \mathbf{x}_{2,i} = 0 \\
 & \mathbf{x}_{1,k} = \mathbf{x}_{2,k} = 1 \text{ for all } k \in S_I \\
 & \mathbf{y}_{1,k} = 1 \text{ for all } k \in S_I \\
 & \mathbf{y}_{2,k} = 1 \text{ for all } k \in S_I \\
 & \mathbf{x}_1^T \mathbf{E}^T = \mathbf{e}^T \\
 & \mathbf{x}_2^T \mathbf{E}^T = \mathbf{e}^T \\
 & \mathbf{y}_1^T \mathbf{F}^T = \mathbf{f}^T \\
 & \mathbf{y}_2^T \mathbf{F}^T = \mathbf{f}^T \\
 & \{\mathbf{x}_1, \mathbf{x}_2, \mathbf{y}_1, \mathbf{y}_2\} \geq \mathbf{0}
 \end{aligned} \tag{2}$$

186 **Proposition 1.** The optimal objective values in Problem 1 and Problem 2 are the same. 3

187 *Proof.* Let f_1 be the optimal objective value in Problem 1 and f_2 be the optimal objective value 4
 188 in Problem 2. Suppose that the optimal variables in Problem 1 are $\mathbf{x}_1^1, \mathbf{x}_2^1, \mathbf{y}^1$. Now set $\mathbf{x}_1^2 = \mathbf{x}_1^1$,
 189 $\mathbf{x}_2^2 = \mathbf{x}_2^1, \mathbf{y}_1^2 = \mathbf{y}^1, \mathbf{y}_2^2 = \mathbf{y}^1$. Then $(\mathbf{x}_1^2, \mathbf{x}_2^2, \mathbf{y}_1^2, \mathbf{y}_2^2)$ gives a feasible solution to Problem 2 with
 190 objective value equal to f_1 . So $f_2 \geq f_1$. Now suppose that the optimal variables in Problem 2 are
 191 $\mathbf{x}_1^2, \mathbf{x}_2^2, \mathbf{y}_1^2, \mathbf{y}_2^2$. Now set $\mathbf{x}_1^1 = \mathbf{x}_1^2, \mathbf{x}_2^1 = \mathbf{x}_2^2$, and set $\mathbf{y}^1$ equal to the strategy that follows $\mathbf{y}_1^2$ at states
 192 following action c for player 1, and follows $\mathbf{y}_2^2$ otherwise. Then $(\mathbf{x}_1^1)^T \mathbf{A} \mathbf{y}^1 = (\mathbf{x}_1^1)^T \mathbf{A} \mathbf{y}_1^2$, since both
 193 players only take actions to get to information set I , $\mathbf{x}_1^1$ takes action c at I , and the strategies $\mathbf{y}^1$ and
 194 $\mathbf{y}_1^2$ are identical after player 1 takes action c at I . Similarly, $(\mathbf{x}_2^1)^T \mathbf{A} \mathbf{y}^1 = (\mathbf{x}_2^1)^T \mathbf{A} \mathbf{y}_2^2$, since both
 195 players only take actions to get to I , $\mathbf{x}_1^1$ does not take action c at I , and the strategies $\mathbf{y}^1$ and $\mathbf{y}_2^2$
 196 are identical after player 1 does not take action c at I . So $f_1 \geq f_2$. So we conclude that $f_1 = f_2$. $\square$

197 Proposition 1 allows us to divide Problem 2 into the following two subproblems. If f is the optimal 5
 198 objective value of Problem 2, f_1 is the optimal objective of Problem 3, and f_2 is the optimal objective
 199 of Problem 4, then we have $f = f_1 - f_2$.

$$\begin{aligned}
 \max_{\mathbf{x}_2} \min_{\mathbf{y}} \quad & \mathbf{x}_2^T \mathbf{A} \mathbf{y} \\
 \text{s.t.} \quad & \mathbf{x}_{2,i} = 0 \\
 & \mathbf{x}_{2,k} = 1 \text{ for all } k \in S_I \\
 & \mathbf{y}_k = 1 \text{ for all } k \in S_I \\
 & \mathbf{x}_2^T \mathbf{E}^T = \mathbf{e}^T \\
 & \mathbf{y}^T \mathbf{F}^T = \mathbf{f}^T \\
 & \{\mathbf{x}_2, \mathbf{y}\} \geq \mathbf{0}
 \end{aligned} \tag{3}$$

$$\begin{aligned}
 \max_{\mathbf{x}_1} \max_{\mathbf{y}} \quad & \mathbf{x}_1^T \mathbf{A} \mathbf{y} \\
 \text{s.t.} \quad & \mathbf{x}_{1,i} = 1 \\
 & \mathbf{x}_{1,k} = 1 \text{ for all } k \in S_I \\
 & \mathbf{x}_1^T \mathbf{E}^T = \mathbf{e}^T \\
 & \mathbf{y}^T \mathbf{F}^T = \mathbf{f}^T \\
 & \{\mathbf{x}_1, \mathbf{y}\} \geq \mathbf{0}
 \end{aligned} \tag{4}$$

200 Let us first look at Problem 3 and consider the inner subproblem for a fixed $\mathbf{x}_2$. 8

$$\begin{aligned}
 \min_{\mathbf{y}} \quad & \mathbf{x}_2^T \mathbf{A} \mathbf{y} \\
 \text{s.t.} \quad & \mathbf{y}_k = 1 \text{ for all } k \in S_I \\
 & \mathbf{y}^T \mathbf{F}^T = \mathbf{f}^T \\
 & \mathbf{y} \geq \mathbf{0}
 \end{aligned} \tag{9}$$

The Lagrangian is 10

$$L(\mathbf{y}, \boldsymbol{\lambda}, \boldsymbol{\gamma}, \mathbf{r}) = \mathbf{x}_2^T \mathbf{A} \mathbf{y} - (\mathbf{f}^T - \mathbf{y}^T \mathbf{F}^T) \boldsymbol{\lambda} - \sum_{k \in S_I} \gamma_k (\mathbf{y}_k - 1) - \mathbf{r}^T \mathbf{y} \tag{11}$$

$$\frac{\partial L}{\partial \mathbf{y}} = \mathbf{x}_2^T \mathbf{A} + \lambda^T \mathbf{F} - \sum_{k \in S_I} \gamma_k \mathbf{e}_k - \mathbf{r}^T \quad 1$$

201 The dual problem is 2

$$\begin{aligned} \max_{\lambda, \gamma} \quad & -\mathbf{f}^T \lambda - \sum_{k \in S_I} \gamma_k \\ \text{s.t.} \quad & \mathbf{x}_2^T \mathbf{A} + \lambda^T \mathbf{F} - \sum_{k \in S_I} \gamma_k \mathbf{e}_k \geq \mathbf{0} \end{aligned} \quad 3$$

202 So Problem 3 is equivalent to: 4

$$\begin{aligned} \max_{\mathbf{x}_2, \lambda, \gamma} \quad & -\mathbf{f}^T \lambda - \sum_{k \in S_I} \gamma_k \\ \text{s.t.} \quad & \mathbf{x}_2^T \mathbf{A} + \lambda^T \mathbf{F} - \sum_{k \in S_I} \gamma_k \mathbf{e}_k \geq \mathbf{0} \\ & \mathbf{x}_{2,i} = 0 \\ & \mathbf{x}_{2,k} = 1 \text{ for all } k \in S_I \\ & \mathbf{x}_2^T \mathbf{E}^T = \mathbf{e}^T \\ & \mathbf{x}_2 \geq \mathbf{0} \end{aligned} \quad 5 \quad (5)$$

203 Next consider Problem 4. Note that both players are aligned in trying to maximize the objective. Let 6
 204 us define a new problem $\bar{G}$ where player 1 selects the actions for both players. We denote player 1's
 205 strategy in this modified game by $\bar{\mathbf{x}}_1$. We modify $\mathbf{E}$ and resize $\mathbf{e}$ accordingly and denote them by $\mathbf{E}'$
 206 and $\mathbf{e}'$. The payoffs can now be represented as a vector $\mathbf{a}$. In the new representation let i' denote the
 207 index of the sequence with concluding action c , and let S'_I denote the set of indices corresponding to
 208 action sequences leading to I .

$$\begin{aligned} \max_{\bar{\mathbf{x}}_1} \quad & \bar{\mathbf{x}}_1^T \mathbf{a} \\ \text{s.t.} \quad & \bar{\mathbf{x}}_{1,i'} = 1 \\ & \bar{\mathbf{x}}_{1,k} = 1 \text{ for all } k \in S'_I \\ & \bar{\mathbf{x}}_1^T \mathbf{E}'^T = \mathbf{e}'^T \\ & \bar{\mathbf{x}}_1 \geq \mathbf{0} \end{aligned} \quad 7 \quad (6)$$

209 Let u_2 denote the optimal objective value for optimization problem 5, and u_1 denote the optimal 8
 210 objective value for problem 6. If $u_2 > u_1$ then we conclude that action c is strictly dominated. If
 211 $u_2 < u_1$ then we conclude that action c is not strictly or weakly dominated. If $u_2 = u_1$, let u_3 denote
 212 the optimal objective value for problem 7, and let u_4 denote the optimal objective value for problem 8.
 213 If $u_3 > u_4$ then we conclude that action c is weakly dominated, and if $u_3 = u_4$ we conclude that
 214 action c is not strictly or weakly dominated (note that we cannot have $u_3 < u_4$).

$$\begin{aligned} \max_{\bar{\mathbf{x}}_2, \lambda, \gamma} \quad & \bar{\mathbf{x}}_2^T \mathbf{a} \\ \text{s.t.} \quad & -\mathbf{f}^T \lambda - \sum_{k \in S_I} \gamma_k = u_2 \\ & \mathbf{x}_2^T \mathbf{A} + \lambda^T \mathbf{F} - \sum_{k \in S_I} \gamma_k \mathbf{e}_k \geq \mathbf{0} \\ & \bar{\mathbf{x}}_{2,i'} = 0 \\ & \bar{\mathbf{x}}_{2,k} = 1 \text{ for all } k \in S'_I \\ & \bar{\mathbf{x}}_2^T \mathbf{E}'^T = \mathbf{e}'^T \\ & \bar{\mathbf{x}}_2 \geq \mathbf{0} \end{aligned} \quad 9 \quad (7)$$

The components of $\bar{\mathbf{x}}_2$ for player 1 in $\bar{G}$ correspond to $\mathbf{x}_2$ in G . 10

$$\begin{aligned} \min_{\bar{\mathbf{x}}_1} \quad & \bar{\mathbf{x}}_1^T \mathbf{a} \\ \text{s.t.} \quad & \bar{\mathbf{x}}_{1,i'} = 1 \\ & \bar{\mathbf{x}}_{1,k} = 1 \text{ for all } k \in S'_I \\ & \bar{\mathbf{x}}_1^T \mathbf{E}'^T = \mathbf{e}'^T \\ & \bar{\mathbf{x}}_1 \geq \mathbf{0} \end{aligned} \quad 11 \quad (8)$$

215 We can perform this procedure for every action at each information set of player 1, and analogously 12
 216 for player 2. Since the number of actions is linear in the size of the game tree, the overall procedure
 217 involves solving a linear number of linear programs and therefore runs in polynomial time. We can
 218 repeat the procedure iteratively until no more actions are removed for any player. Thus, iterative

219 removal of dominated actions can be performed in polynomial time. Note that the procedure applies 1
 220 to all two-player games and does not assume that they are zero sum. The procedure also removes
 221 actions that are dominated by any mixed strategy (which may play a probability distribution over
 222 actions at the same information set I), not just actions that are dominated by a pure action.

223 Now suppose we want to determine if an action is dominated for player 1 in a game with n players 2
 224 for $n > 2$. We can construct a new two-player game where player 2 now controls all actions that
 225 were previously controlled by any player other than player 1. Then we can run the above procedure
 226 in this new game. Thus, we can perform iterative removal of (strictly or weakly) dominated strategies
 227 in polynomial time in n -player extensive-form games of imperfect information.

228 **Theorem 1.** *There exists a polynomial-time algorithm that determines whether an action is strictly 3
 229 dominated in an n -player extensive-form game of imperfect information.*

230 **Theorem 2.** *There exists a polynomial-time algorithm that determines whether an action is weakly 4
 231 dominated in an n -player extensive-form game of imperfect information.*

232 **Theorem 3.** *Iterated removal of strictly and weakly dominated actions can be performed in polyno- 5
 233 mial time in n -player extensive-form games of imperfect information.*

234 5 Experiments 6

235 Now that we have defined dominated actions and showed that they can be computed efficiently, we 7
 236 would like to investigate whether they can be a useful concept in practice. Poker has been widely
 237 studied as a test domain for imperfect-information games. The most popular variant regularly played
 238 by humans is No-Limit Texas Hold'em (NLHE). Two-player NLHE works as follows. Initially two
 239 players each have a *stack* of chips. One player, called the *small blind*, initially puts k worth of chips
 240 in the middle, while the other player, called the *big blind*, puts in $2k$. The chips in the middle are
 241 known as the *pot*, and will go to the winner of the hand. Next, there is an initial round of betting. The
 242 player whose turn it is to act can choose from three available options:

- 243 • *Fold*: Give up on the hand, surrendering the pot to the opponent. 8
- 244 • *Call*: Put in the minimum number of chips needed to match the number of chips put into the
 245 pot by the opponent. For example, if the opponent has put in \$1000 and we have put in \$400,
 246 a call would require putting in \$600 more. A call of zero chips is also known as a *check*.
- 247 • *Bet*: Put in additional chips beyond what is needed to call. A bet can be of any size from
 248 1 chip up to the number of chips a player has left in their stack, provided it exceeds some
 249 minimum value and is a multiple of the smallest chip denomination. A bet of all of one's
 250 remaining chips is called an *all-in* bet (aka a *shove*).

251 The initial round of betting ends if a player has folded, if there has been a bet and a call, or if both 9
 252 players have checked. If the round ends without a player folding, then three public cards are revealed
 253 face-up on the table (called the *flop*) and a second round of betting takes place. Then one more public
 254 card is dealt (called the *turn*) and a third round of betting, followed by a fifth public card (called the
 255 *river*) and a final round of betting. If a player ever folds, the other player wins all the chips in the pot.
 256 If the final betting round is completed without a player folding (or if a player is all-in at an earlier
 257 round), then both players reveal their private cards, and the player with the best five-card hand (out of
 258 their two private cards and the five public cards) wins the pot (it is divided equally if there is a tie).

259 In some situations the blinds are very large relative to stack sizes. This can happen frequently at later 10
 260 stages in poker *tournaments*, where the blinds increase after a certain time duration. A common rule
 261 is that when stack sizes are below around 8 big blinds a *shove-or-fold* strategy should be employed
 262 where each player only goes all-in or folds [9]. Study of optimal shove-or-fold strategies has been
 263 considered for 2-player [10] and 3-player [3, 4] poker tournament endgames. The poker site Americas
 264 Cardroom² has specific “All-in or Fold” tables with up to 4 players where players are only allowed to
 265 play shove-or-fold strategies. The initial stack sizes at these tables are either 8 or 10 times the big
 266 blind; the highest stake available has blinds of \$100 and \$200 with initial stacks of \$2000.

267 In the two-player NLHE shove-or-fold game, each player has 169 strategically distinct hands with 11
 268 which they can choose to shove or fold (13 pocket pairs and $\frac{13 \cdot 12}{2} = 78$ combinations of each

²<https://www.americascardroom.eu/>

non-paired offsuit and suited hand). Let player 1 denote the small blind and player 2 denote the big blind. We assume that the blinds are small blind $k = 100$ and big blind $2k = 200$, and initial stacks are 1600 (8 times the big blind). We first remove all strictly dominated actions for player 1, followed by removing strictly dominated actions for player 2 (note that removing weakly dominated actions does not provide additional benefit in this game). It turns out that 85 actions for player 1 are removed and 99 actions for player 2 are removed. Thus, the initial game with 169 hands per player can be reduced to a game where player 1 must make a decision with 84 hands and player 2 must make a decision with 70 hands; the number of decision points has been reduced by over 50%. It turns out that performing an additional round of removing dominated actions does not remove any further actions.³

Next we consider the setting where the stacks are 5 times the big blind (i.e., 1000). In this game five rounds of iterated removal of dominated actions are needed. The first round removes 108 dominated actions for player 1 and 129 for player 2; the second round removes 20 for player 1 and 16 for player 2; the third round removes 8 for player 1 and 6 for player 2; the fourth round removes 7 for player 1 and 0 for player 2. In the final reduced game player 1 must make a decision with only 25 hands while player 2 must make a decision with only 16 hands. Table 5 shows the non-dominated actions for each player with parentheses indicating the iteration at which the alternative action was removed. ‘S’ indicated shove, ‘F’ indicates fold, and ‘?’ indicates that neither action was dominated. If stacks are 4 times the big blind the game is solved completely after 4 rounds of removing dominated actions, and for stacks of 3 big blinds the game is solved completely after 2 rounds. These results demonstrate that iteratively removing dominated actions can significantly reduce the size of realistic games. While for simplicity we considered two-player zero-sum games, for which the full game can be solved directly by a linear program, the computational benefit for games with more than two players could be much more significant.

	A	K	Q	J	T	9	8	7	6	5	4	3	2		A	K	Q	J	T	9	8	7	6	5	4	3	2
A	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	A	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)
K	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	K	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)
Q	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	Q	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)
J	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	J	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(3)	S(3)	?	?
T	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(3)	?	T	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(3)	?	?	?	?
9	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(3)	?	?	?	9	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(3)	?	F(4)	F(3)	F(2)	F(2)
8	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(1)	S(2)	S(4)	?	?	?	8	S(1)	S(1)	S(1)	S(1)	S(3)	S(1)	S(1)	S(1)	?	F(2)	F(2)	F(1)	F(1)
7	S(1)	S(1)	S(1)	S(1)	S(1)	S(3)	S(5)	S(1)	S(2)	S(3)	?	?	F(3)	7	S(1)	S(1)	S(1)	S(2)	?	?	?	S(1)	S(1)	S(1)	F(2)	F(1)	F(1)
6	S(1)	S(1)	S(1)	S(1)	?	?	?	?	S(1)	S(3)	?	?	F(4)	6	S(1)	S(1)	S(1)	?	?	?	F(3)	F(2)	S(1)	S(1)	F(2)	F(1)	F(1)
5	S(1)	S(1)	S(1)	S(1)	?	?	?	?	?	S(1)	S(4)	?	?	5	S(1)	S(1)	S(1)	?	F(4)	F(2)	F(2)	F(2)	F(1)	S(1)	S(1)	F(1)	F(1)
4	S(1)	S(1)	S(1)	S(2)	F(4)	F(2)	F(2)	F(3)	F(4)	?	S(1)	?	F(4)	4	S(1)	S(1)	S(1)	?	F(2)	F(2)	F(1)	F(1)	F(1)	F(1)	S(1)	F(1)	F(1)
3	S(1)	S(1)	S(1)	?	F(2)	F(2)	F(2)	F(2)	F(2)	F(2)	F(2)	S(1)	F(3)	3	S(1)	S(1)	S(1)	?	F(2)	F(1)	F(1)	F(1)	F(1)	F(1)	F(1)	S(1)	F(1)
2	S(1)	S(1)	S(1)	F(4)	F(2)	F(2)	F(2)	F(2)	F(2)	F(2)	F(2)	F(2)	S(1)	2	S(1)	S(1)	S(1)	F(2)	F(2)	F(1)	F(1)	F(1)	F(1)	F(1)	F(1)	F(1)	S(1)

Figure 4: Dominated actions in 2-player NLHE all-in-fold with 5 big blind stacks (player 1 left, player 2 right). Suited hands are in the upper right and unsuited hands are in the lower left.

6 Conclusion

Dominance is a fundamental concept in game theory. It is well-understood in strategic-form games, but its impact in imperfect-information games has so far been unexplored. We consider several plausible definitions of dominated actions which we demonstrate to be problematic; however, one of them which we denote as *strong domination* can be useful as an efficient preprocessing step. We present a new definition that addresses these limitations. We show that both strictly and weakly dominated actions can be identified in polynomial time, and that iterative removal can be performed in polynomial time in n -player games. Our algorithms identify actions that are dominated by any mixed strategy, not necessarily a pure action. We demonstrate empirically that removing dominated actions can play a significant role in reducing the size of realistic imperfect-information games. This can serve as an efficient preprocessing step before computation of a Nash equilibrium. In practice our algorithms could be sped up by several heuristics such as traversing the information sets in decreasing order of their depth. Recent work has shown that some games contain many “mistake” actions that are played with probability zero in all Nash equilibria but are not dominated [2]. Thus, there is potentially more future ground to explore on efficient game reduction by elimination of poor actions.

³We note that no actions are strongly dominated for either player in this game (i.e., Flawed Definition 1). For example, consider the action of player 1 folding with pocket aces (the best starting hand). The payoff of this action is $-\$100$. If player 1 goes all-in, there are leaf nodes where player 1 will receive a payoff of $-\$1600$.

References 1

- [1] Vincent Conitzer and Tuomas Sandholm. Complexity of (iterated) dominance. In *Proceedings of the ACM Conference on Electronic Commerce (ACM-EC)*, pages 88–97, Vancouver, Canada, 2005. ACM.
- [2] Sam Ganzfried. Mistakes in games. In *Proceedings of the International Conference on Distributed Artificial Intelligence (DAI)*, 2019.
- [3] Sam Ganzfried and Tuomas Sandholm. Computing an approximate jam/fold equilibrium for 3-player no-limit Texas hold 'em tournaments. In *Proceedings of the International Conference on Autonomous Agents and Multi-Agent Systems (AAMAS)*, 2008.
- [4] Sam Ganzfried and Tuomas Sandholm. Computing equilibria in multiplayer stochastic games of imperfect information. In *Proceedings of the 21st International Joint Conference on Artificial Intelligence (IJCAI)*, 2009.
- [5] Itzhak Gilboa, Ehud Kalai, and Eitan Zemel. On the order of eliminating dominated strategies. *Operations Research Letters*, 9:85–89, 1990.
- [6] Itzhak Gilboa, Ehud Kalai, and Eitan Zemel. The complexity of eliminating dominated strategies. *Mathematics of Operation Research*, 18:553–565, 1993.
- [7] Daphne Koller, Nimrod Megiddo, and Bernhard von Stengel. Fast algorithms for finding randomized strategies in game trees. In *Proceedings of the 26th ACM Symposium on Theory of Computing (STOC)*, pages 750–760, 1994.
- [8] Michael Maschler, Eilon Solan, and Shmuel Zamir. *Game Theory*. Cambridge University Press, 2013.
- [9] mersenneary. Correctly applying “shove or fold” small blind endgame strategy (basics article), 2011. <https://husng.com/content/correctly-applying-“shove-or-fold”-small-blind-endgame-strategy-basics-article>.
- [10] Peter Bro Miltersen and Troels Bjerre Sørensen. A near-optimal strategy for a heads-up no-limit Texas Hold'em poker tournament. In *Proceedings of the International Conference on Autonomous Agents and Multi-Agent Systems (AAMAS)*, 2007.
- [11] Rahul Savani and Theodore L. Turocy. Gambit: The package for computation in game theory, Version 16.3.0, 2025.
- [12] Yoav Shoham and Kevin Leyton-Brown. *Multiagent Systems Algorithmic, Game-Theoretic, and Logical Foundations*. Cambridge University Press, 2009.

NeurIPS Paper Checklist ¹

1. Claims ²

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? ³

Answer: [Yes] ⁴

Justification: ⁵

Guidelines: ⁶

- The answer NA means that the abstract and introduction do not include the claims made in the paper. ⁷
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations ⁸

Question: Does the paper discuss the limitations of the work performed by the authors? ⁹

Answer: [Yes] ¹⁰

Justification: ¹¹

Guidelines: ¹²

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper. ¹³
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory Assumptions and Proofs ¹⁴

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof? ¹⁵

Answer: [Yes] ¹⁶

Justification: 1

Guidelines: 2

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental Result Reproducibility 4

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)? 5

Answer: [Yes] 6

Justification: 7

Guidelines: 8

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code 10

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material? 11

Answer: [No] 1

Justification: 2

Guidelines: 3

- The answer NA means that paper does not include experiments requiring code. 4
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental Setting/Details 5

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results? 6

Answer: [Yes] 7

Justification: 8

Guidelines: 9

- The answer NA means that the paper does not include experiments. 10
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment Statistical Significance 11

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments? 12

Answer: [Yes] 13

Justification: 14

Guidelines: 15

- The answer NA means that the paper does not include experiments. 16
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.

- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments Compute Resources 2

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments? 3

Answer: [Yes] 4

Justification: 5

Guidelines: 6

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code Of Ethics 8

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>? 9

Answer: [Yes] 10

Justification: 11

Guidelines: 12

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader Impacts 14

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed? 15

Answer: [Yes] 16

Justification: 17

Guidelines: 18

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to

generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.

- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards 2

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [Yes] 4

Justification: 5

Guidelines: 6

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets 8

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes] 10

Justification: 11

Guidelines: 12

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New Assets 14

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes] 1

Justification: 2

Guidelines: 3

- The answer NA means that the paper does not release new assets. 4
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and Research with Human Subjects 5

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)? 6

Answer: [NA] 7

Justification: 8

Guidelines: 9

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 10
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects 11

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained? 12

Answer: [NA] 13

Justification: 14

Guidelines: 15

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 16
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.