
Inconsistency-Aware Minimization: Improving Generalization with Unlabeled Data¹

Anonymous Author(s)²

Affiliation

Address

email

Abstract³

1 Estimating the generalization gap and devising optimization methods that general-⁴
2 ize better are crucial for deep learning models, both for theoretical understanding
3 and for practical applications. The ability to leverage unlabeled data for these pur-
4 poses offers significant advantages in real-world scenarios. This paper introduces
5 a novel generalization measure, termed *local inconsistency*, developed from an
6 information-geometric perspective of the neural network’s parameter space. A
7 key feature of local inconsistency is its computability from unlabeled data. We
8 establish its theoretical underpinnings by connecting local inconsistency to the
9 Fisher information matrix and the loss Hessian. Empirically, we demonstrate
10 that local inconsistency correlates with the generalization gap also exhibiting
11 characteristics comparable to *sharpness*. Based on these findings, we propose
12 Inconsistency-Aware Minimization (IAM) that incorporates local inconsistency
13 into the objective. We demonstrate that in standard supervised learning settings,
14 IAM enhances generalization, achieving performance comparable to existing meth-
15 ods such as Sharpness-Aware Minimization. Furthermore, IAM exhibits efficacy
16 in semi- and self-supervised learning scenarios, where the local inconsistency is
17 computed from the unlabeled data.

18 1 Introduction⁵

19 Estimating the generalization gap and developing optimization techniques to enhance performance⁶
20 on unseen data are pivotal challenges in both the theory and practice of deep learning. There are
21 numerous reports linking the flatness of the loss landscape to generalization [16, 7, 20, 9]. Conse-
22 quently, numerous minimization methods leveraging sharpness have been proposed, demonstrating
23 improvements in generalization [8, 19, 17, 35]. However, the efficacy of such sharpness-based
24 approaches is not fully understood as recent studies, including Andriushchenko et al. [1], indicate
25 that sharpness alone does not reliably predict the generalization gap.

26 On the other hand, Jiang et al. [14], Johnson and Zhang [15] have investigated alternative gener-⁷
27 alization measures, such as *disagreement* and *inconsistency*. Under certain conditions (e.g., zero
28 training error or low randomness of final states), these metrics have been shown to estimate the
29 generalization gap more accurately than ones based on sharpness. From a practical standpoint, it is
30 highly desirable to leverage unlabeled data to compute these measures. However, existing methods
31 for computing disagreement typically require training two separate models under identical conditions
32 (while still being subject to training randomness), and estimating inconsistency often necessitates
33 training multiple models on each of distinct datasets, thereby incurring substantial computational
34 overhead. Such prerequisites hinder the direct minimization of these measures within a single-model
35 training paradigm.

36 In many real-world applications, unlabeled data is far more abundant and accessible than labeled data. 1
37 Therefore, methods capable of estimating the generalization gap and subsequently minimizing it using
38 only unlabeled data within a single-model framework (i.e., without resorting to auxiliary models or
39 requiring held-out labels) are crucial for resource-efficient deployment in practical scenarios.

40 Addressing this need, this paper introduces *local inconsistency*, a novel generalization measure derived 2
41 from an information-geometric perspective of the parameter space. This measure is theoretically
42 grounded through its connection to the Fisher Information Matrix (FIM) and the loss Hessian.
43 Crucially, local inconsistency can be computed using only **unlabeled** data and from a **single** trained
44 model. We demonstrate that local inconsistency exhibits a predictive capability for the generalization
45 gap comparable to sharpness-based measures in some settings. It also shows distinct characteristics
46 where traditional sharpness measures may falter. Building upon local inconsistency, we propose
47 **Inconsistency-Aware Minimization (IAM)**, a novel optimization strategy that incorporates local
48 inconsistency into the training objective. In standard supervised learning settings on CIFAR-10
49 and CIFAR-100[18], IAM achieves generalization performance comparable or superior to existing
50 methods like Sharpness-Aware Minimization (SAM) [8]. Furthermore, IAM demonstrates its efficacy
51 in semi- and self-supervised learning, achieving higher test accuracy than standard SGD when
52 integrated with SimCLR [6].

53 Our contributions can be summarized as follows: 3

- 54 • We propose *local inconsistency*, a novel generalization measure rooted in information 4
55 geometry, designed to capture the sensitivity of model outputs to perturbations in the
56 parameter space. This measure offers significant practical advantages as it can be computed
57 using (i) only unlabeled data and (ii) a single trained model.
- 58 • We establish the theoretical underpinnings of local inconsistency by linking it to the Fisher
59 Information Matrix (FIM) and the loss Hessian. Furthermore, we discuss its approximate
60 relationship with the inconsistency measure previously proposed by Johnson and Zhang
61 [15].
- 62 • We develop Inconsistency-Aware Minimization (IAM), a novel optimization framework
63 with two variants (IAM-D and IAM-S), that directly incorporates local inconsistency into
64 the training objective to seek flatter minima in terms of output sensitivity. IAM demonstrates
65 improved generalization performance across diverse learning paradigms, including standard
66 supervised learning and self-supervised learning settings, showcasing its broad applicability.

67 2 Related work 5

68 Understanding and improving generalization in deep neural networks, especially given their large 6
69 capacity and tendency to overfit [33], remains a central challenge. While networks can memorize
70 random labels [33] and learn simple patterns before noise [2], phenomena like double descent [25]
71 and the inadequacy of uniform convergence theory [24] highlight the need for novel generalization
72 measures beyond loss-based metrics.

73 Traditional measures like VC-dimension often fall short. While spectrally-normalized margin bounds 7
74 [3] and PAC-Bayes approaches offer insights, no single measure consistently predicts generalization
75 [13]. Recently, disagreement [14] and inconsistency [15] have shown promise, correlating well
76 with the generalization gap, even when computed on unlabeled data. However, their reliance on
77 training multiple models poses practical limitations for direct optimization in a single-model setup,
78 underscoring the need for efficient, label-free, single-model generalization measures.

79 The geometry of the loss landscape, particularly the flatness of minima, has been extensively linked 8
80 to generalization [16, 20]. However, the utility of sharpness as a sole predictor is debated due to
81 issues like scale invariance [7] and its correlation with training hyperparameters rather than true
82 generalization [1]. Indeed, some studies suggest that output inconsistency and instability can be more
83 reliable predictors than sharpness [15]. Information geometry has inspired reparametrization-invariant
84 sharpness measure [12], but these can be computationally expensive. This context motivates our
85 exploration of “local inconsistency”, an alternative geometric measure focusing on output sensitivity
86 within a parameter neighborhood, computable from unlabeled data using a single model.

87 Various regularization techniques, both explicit (e.g., dropout [30], batch normalization [28], Mixup 9
88 [34]) and implicit (e.g., SGD’s bias [10, 29]), aim to improve generalization. Methods like Sharpness-

89 Aware Minimization (SAM, [8]) and ASAM [19] directly optimize for flat minima and have shown 1
 90 significant improvements. Despite their success, the precise role of sharpness in generalization
 91 remains an active area of research [13, 1], further motivating the development of complementary
 92 approaches like our proposed IAM.

93 3 Background and preliminaries 2

94 In this section, we briefly review fundamental concepts and notations essential for understanding 3
 95 our proposed metric and its theoretical connections. We focus on probabilistic classification models,
 96 information geometry, and aspects of the loss landscape.

97 3.1 Notation and problem setup 4

98 We consider probabilistic classification models. Let $x \in \mathcal{X}$ be a data point from the input space 5
 99 $\mathcal{X}$, and $y \in [C] = \{0, 1, \dots, C-1\}$ be the corresponding class label, where C is the total number
 100 of classes. The data pair (x, y) are assumed to be drawn from an underlying distribution $\mathcal{D}$ over
 101 $\mathcal{X} \times [C]$. A model, parameterized by $\theta \in \mathbb{R}^m$, outputs a probability distribution over classes for a
 102 given input x . This is typically achieved by transforming a logit vector $z(x, \theta)$ through a softmax
 103 function: $f(x, \theta) = \text{softmax}(z(x, \theta))$. Thus, $f(x, \theta) = [p(0|x; \theta), p(1|x; \theta), \dots, p(C-1|x; \theta)]^\top$.
 104 Given a training dataset $Z_n = \{(x_i, y_i) : i = 1, \dots, n\}$ drawn i.i.d. from $\mathcal{D}$, the model is typically
 105 trained by minimizing a loss function. For classification, the empirical Cross-Entropy (CE) loss will
 106 be written as $L(\theta) = \frac{1}{n} \sum_{i=1}^n l_i(\theta)$, where per-sample loss is $l_i(\theta) = l(x_i, y_i; \theta) = -\log p(y_i|x_i; \theta)$.

107 3.2 Fisher information matrix (FIM) and KL divergence 6

108 The Fisher information matrix (FIM), $F(\theta)$, for the family of probability density $p(x, y; \theta) =$ 7
 109 $p(x)p(y|x; \theta)$ parameterized by a parameters θ is defined as

$$\begin{aligned} F(\theta) &= \mathbb{E}_{x \sim p(x)} [\mathbb{E}_{y \sim p(y|x; \theta)} [\nabla_\theta l(x, y; \theta) \nabla_\theta l(x, y; \theta)^\top]] \\ &= \mathbb{E}_{x \sim p(x)} [\nabla_\theta z(x, \theta) (\text{diag}(f(x, \theta)) - f(x, \theta)f(x, \theta)^\top) \nabla_\theta z(x, \theta)^\top]. \end{aligned} \quad (1) \quad 8$$

110 In practice, the expectation $\mathbb{E}_{p(x)}$ is often approximated by an empirical average over the available 9
 111 data (e.g., training data $\{x_i\}_{i=1}^n$ or unlabeled data).

112 The Kullback-Leibler (KL) divergence between the output distributions of a model with parameters θ 10
 113 and a slightly perturbed model $\theta + \delta$, $f(x, \theta)$ and $f(x, \theta + \delta)$, respectively, can be locally approximated
 114 using a second-order Taylor expansion with respect to δ as:

$$\mathbb{E}_{x \sim p(x)} [\text{KL}(f(x, \theta) \| f(x, \theta + \delta))] = \frac{1}{2} \delta^\top F(\theta) \delta + O(\|\delta\|_2^3). \quad (2) \quad 11$$

115 3.3 Loss Hessian and Gauss-Newton approximation 12

116 The geometry of the empirical loss surface $L(\theta)$ is described by its Hessian matrix $H(\theta) = \nabla_\theta^2 L(\theta)$. 13
 117 For the Cross-Entropy (CE) loss, the Hessian can be approximated by the Gauss-Newton (GN) matrix,
 118 $G(\theta)$. The second derivative of the per-sample CE loss $\ell_i(\theta)$ with respect to the logits $z_i = z(x_i, \theta)$,
 119 $\nabla_z^2 \ell_i(\theta) = \text{diag}(f(x_i, \theta)) - f(x_i, \theta)f(x_i, \theta)^\top$, depends only on the model's output probabilities
 120 $f(x_i, \theta)$. Consequently, the per-sample GN term, $G_i(\theta) = \nabla_\theta z_i^\top (\nabla_z^2 \ell_i) \nabla_\theta z_i$, is equivalent to the
 121 FIM contribution in Eq. (1). The empirical GN matrix, $G(\theta) = \frac{1}{n} \sum_{i=1}^n G_i(\theta)$, thus often termed the
 122 empirical FIM, provides a positive semi-definite approximation to $H(\theta)$ and is frequently used in
 123 optimization [22, 26].

124 4 Accessing generalization gap via local inconsistency 14

125 This section introduces our proposed measure, local inconsistency, designed to capture the gen- 15
 126 eralization gap. We first define local inconsistency and elucidate its theoretical underpinnings by
 127 connecting it to the FIM and the loss Hessian. We then discuss its relationship with inconsistency
 128 [15]. Finally, we present empirical results demonstrating the correlation between local inconsistency
 129 and the generalization gap, comparing it with other common measures.

130 4.1 Local inconsistency, $S_\rho(\theta)$

131 We introduce local inconsistency, $S_\rho(\theta)$, defined as: 1

$$S_\rho(\theta) = \max_{\|\delta\| \leq \rho} \mathbb{E}_{x \sim p(x)} [\text{KL}(f(x, \theta) \| f(x, \theta + \delta))], \quad 2 \quad (3)$$

132 which represents the sensitivity of the model’s output distribution $f(x, \theta)$ with respect to the worst 3
133 perturbations δ , within an Euclidean ball of radius ρ around the parameter θ . Intuitively, a high
134 value of $S_\rho(\theta)$ indicates that the model’s output distribution is highly sensitive to small perturbations
135 in parameter space. This sensitivity suggests potential instability or uncertainty in the model’s
136 predictions associated with the vicinity of θ .

137 **Practical Advantages of S_ρ** Local inconsistency shares a practical advantage with sharpness- 4
138 based measures [16, 8] in that it can be calculated using a **single** trained model. Furthermore, like
139 disagreement [14] and inconsistency [15], our metric can be estimated using only **unlabeled** data.
140 A notable advantage over inconsistency and disagreement estimation is that evaluating S_ρ does not
141 require training multiple model instances derived from the same training procedure. This potentially
142 makes S_ρ more computationally efficient and practical to compute, especially when model training is
143 resource-intensive.

144 4.2 Connection to FIM and Hessian 5

145 The relationship between our metric S_ρ and the Fisher Information Matrix (FIM) can be established 6
146 by leveraging the local quadratic approximation of the KL divergence, as outlined in Section 3. With
147 this quadratic approximation, we can approximate $S_\rho(w)$ with the maximum eigenvalue of FIM,
148 scaled by $\rho^2/2$:

$$S_\rho(\theta) \approx \max_{\|\delta\| \leq \rho} \frac{1}{2} \delta^\top F(\theta) \delta = \frac{1}{2} (\rho v_{\max})^\top F(\theta) (\rho v_{\max}) = \frac{1}{2} \rho^2 \lambda_{\max}, \quad 7$$

149 where $v_{\max}$ is the eigenvector corresponding to the largest eigenvalue $\lambda_{\max}$ of $F(\theta)$. Remarkably, 8
150 this approximation requires only the model θ and unlabeled data (used to compute the expectation).

151 The Fisher Information Matrix $F(\theta)$, to which $S_\rho(\theta)$ is related via its maximum eigenvalue, also 9
152 connects to the Hessian of the loss function $H(\theta)$. As detailed in Section 3, for Negative Log
153 Likelihood losses such as Cross-Entropy, the Hessian can be approximated by the Gauss-Newton
154 matrix $G(\theta)$, equivalent to empirical FIM computed using training data.

155 Consequently, when calculating $S_\rho(\theta)$ using the training data, it approximates $\frac{1}{2} \rho^2 \lambda_{\max}(G(\theta))$. 10
156 Given that $G(\theta)$ often provides a good approximation to the true loss Hessian near a local minimum,
157 $S_\rho(\theta)$ therefore offers insights into the maximum curvature of the loss landscape in that vicinity.

158 4.3 Local Inconsistency and Generalization Bound 11

159 Beyond its connection to the local geometry of the loss landscape, our proposed local inconsistency 12
160 measure, $S_\rho(\theta)$, can also be linked to the generalization ability of the model. Inspired by PAC-
161 Bayesian analyses that connect the geometry of the loss neighborhood to generalization [8], we can
162 sketch a theoretical argument suggesting that controlling $S_\rho(\theta)$ contributes to a generalization bound.
163 We present an informal theorem that captures this intuition, with the detailed heuristic derivation
164 provided in Appendix C.

165 **Theorem 1 (Informal Generalization Bound with Local Inconsistency)** *Under certain assump-* 13
166 *tions regarding the relationship between the worst-case empirical loss increase in a ρ -neighborhood*
167 *and the local inconsistency $S_\rho(\theta)$ (evaluated on the training set Z_n), the true risk $L_{\mathcal{D}}(\theta) =$*
168 *$\mathbb{E}_{(x,y) \sim \mathcal{D}} [l(x, y; \theta)]$ can be bounded with high probability as:*

$$L_{\mathcal{D}}(\theta) \lesssim L(\theta) + S_\rho(\theta) + \mathcal{R}(\|\theta\|^2 / \rho^2) \quad 14 \quad (4)$$

169 where $\mathcal{R} : \mathbb{R} \rightarrow \mathbb{R}^+$ is a strictly increasing function. 15

170 This bound (Eq. (4)) suggests that minimizing a combination of the empirical loss $L_S(\theta)$ and the 16
171 local inconsistency $S_\rho(\theta)$ can lead to a lower upper bound on the true risk. This provides a theoretical
172 motivation for our Inconsistency-Aware Minimization (IAM) framework, which aims to find solutions
173 that are not only accurate on the training data but also exhibit low output sensitivity in the parameter
174 space, as measured by $S_\rho(\theta)$.

175 4.4 Relation with inconsistency in Johnson and Zhang [15]

176 Local inconsistency exhibits an interesting relationship to the inconsistency in Johnson and Zhang 1
177 [15] defined as:

$$\mathcal{C}_P = \mathbb{E}_{Z_n} \mathbb{E}_{\theta, \theta' \sim \Theta_{P|Z_n}} \mathbb{E}_{x \sim p(x)} [\text{KL}(f(x, \theta) \| f(x, \theta'))]. \quad 2$$

178 We consider the conditional inconsistency for a fixed Z_n , denoted $\mathcal{C}_{P|Z_n}$, without outer expectation. 3
179 Then our proposed metric, $S_\rho(\theta_{Z_n})$, is approximately proportional to the conditional inconsistency
180 $\mathcal{C}_{P|Z_n}$:

$$\frac{m}{2C} \mathcal{C}_{P|Z_n} \lesssim S_\rho(\theta_{Z_n}) \lesssim \frac{m}{2} \mathcal{C}_{P|Z_n}, \quad 4 \quad (5)$$

181 under certain assumptions, such as assuming the parameter posterior $\Theta_{P|Z_n}$ as a distribution with 5
182 isotropic covariance and θ_{Z_n} as mean. This connection arises because both metrics are related to
183 the local geometry captured by the FIM at θ_{Z_n} , with S_ρ being linked to its maximum eigenvalue
184 and $\mathcal{C}_{P|Z_n}$ to its trace. Practically, the eigenspectra of the FIM of a neural network are observed
185 to be dominated by a few large eigenvalues (specifically related to the number of classes, C in
186 classification task) while remaining eigenvalues are near zero. This observation indicates that the ratio
187 $\lambda_{\max}(F(\theta))/\text{Tr}(F(\theta))$ is larger than $\frac{1}{C}$ ($C \ll m$). For detailed derivation, please see Appendix B.

188 4.5 Estimating $S_\rho(w)$ 6

189 Directly computing $S_\rho(w)$ requires solving the maximization problem over the high-dimensional 7
190 parameter perturbation δ . For deep neural networks, finding the exact maximum within the L_2 -ball
191 of radius ρ is generally intractable. Therefore, we resort to numerical approximation methods.

192 For small perturbations δ , the expected KL divergence can be accurately approximated by a second- 8
193 order Taylor expansion involving the Fisher Information Matrix (FIM), $F(\theta)$, as Eq. (2) in Section 3.
194 Under quadratic approximation, as discussed in Section 4.2, the optimal perturbation $\delta^* = \rho v_{\max}$,
195 the maximum value is then $S_\rho(\theta) = \frac{1}{2} \rho^2 \lambda_{\max}$, and the gradient of the approximated KL divergence
196 with respect to δ is $F(\theta)\delta$.


197 This connection motivates not an usual Projected Gradient Ascent, that update $\delta_{k+1} \leftarrow$ 9
198 $\Pi_{\{\delta_k: \|\delta_k\| \leq \rho\}}(\delta_k + \eta F(\theta)\delta_k)$, but an iterative gradient ascent approach that update

$$\delta_{k+1} = \frac{\rho}{\|F(w)\delta_k\|} F(w)\delta_k, \quad \delta_0 = \varepsilon \sim \mathcal{N}\left(0, \frac{\sigma^2}{m} I_m\right), \quad 10$$

199 where σ^2 is initial noise scale. Iterative gradient ascent is precisely one iteration of the Power Iteration 11
200 method used to find the dominant eigenvector of $F(w)$.

Algorithm 1 Estimation of $S_\rho(w)$ 12

1: **Input:** model parameter $w \in \mathbb{R}^m$, noise scale σ^2 ,
2: radius $\rho > 0$, number of steps $K \geq 1$
3: **Initialize** δ_0 randomly with $\mathcal{N}(0, \frac{\sigma^2}{m} I_m)$
4: **for** $k = 0$ to $K - 1$ **do**
5: Compute $g_k = \nabla_\delta \mathbb{E}_{x \sim p(x)} \text{KL}(f(x, \theta) \| f(x, \theta + \delta))|_{\delta=\delta_k}$
6: Update perturbation: $\delta_{k+1} = \rho \frac{g_k}{\|g_k\|_2}$
7: **end for**
8: **return** $\mathbb{E}_{x \sim p(x)} \text{KL}(f(x, \theta) \| f(x, \theta + \delta_K))$



202 4.5.1 Algorithm for estimating $S_\rho(w)$ 16

203 Based on the above, we propose Algorithm 1 to estimate $S_\rho(w)$. This algorithm performs K steps of 17
204 normalized gradient ascent (effectively, Power Iteration under the quadratic approximation) to find an
205 approximate maximizing perturbation δ^* .

4.6 Empirical results

To assess the predictive capability of local inconsistency S_ρ for the generalization gap, we conducted experiments on CIFAR-10. We trained two distinct architectures, a 6-layer CNN (6CNN) and a Wide Residual Network (WRN28-2)[32], under various hyperparameter settings (details in Appendix E). S_ρ was estimated using a disjoint, unlabeled data set. For comparison, we also computed two common sharpness-based measures: the trace, $\text{Tr}(H)$, and the maximum eigenvalue, $\lambda_{\max}(H)$.

Figure 1 presents scatter plots of these metrics against the generalization gap, with Kendall’s Tau (τ) reported for each. For the simpler 6CNN model (top row), S_ρ ($\tau = 0.5141$) exhibited a positive correlation with the generalization gap, comparable to $\text{Tr}(H)$ ($\tau = 0.5444$) and $\lambda_{\max}(H)$ ($\tau = 0.5175$). This suggests that for smaller models, various geometric measures may similarly capture aspects of generalization. However, for the larger WRN28-2 model with data augmentation (bottom row), a more nuanced behavior emerged. As noted by Andriushchenko et al. [1], different training configurations can form distinct solution subgroups. In our WRN28-2 experiments, $\text{Tr}(H)$ and $\lambda_{\max}(H)$ showed positive correlations only within such subgroups, but exhibited negative overall correlations globally ($\tau = -0.0439$ and $\tau = -0.1200$, respectively). In stark contrast, our S_ρ maintained a positive, albeit reduced, correlation across all settings ($\tau = 0.3658$).

This divergence, particularly with larger models and data augmentation, suggests that local inconsistency captures information about the generalization gap that is distinct from, or complementary to, traditional Hessian-based sharpness. While the predictive utility of sharpness metrics can be confounded by these subgroup effects, S_ρ demonstrates more consistent global predictiveness, hinting at its potential as a more robust generalization indicator in complex training scenarios.




Figure 1: Local inconsistency and sharpness measures vs the generalization gap.

5 Inconsistency-Aware Minimization (IAM): incorporating local inconsistency into the objective

Our empirical findings suggest that local inconsistency, $S_\rho(\theta)$ defined in Eq. (3), correlates with the generalization gap. This motivates its use as a regularizer to guide the optimization towards solutions that not only fit the training data, but also exhibit low sensitivity in their output distributions with respect to parameter perturbations. We propose two strategies to incorporate local inconsistency into the training objective.

1. Direct Regularization (IAM-D): This approach directly penalizes local inconsistency by adding it to the standard training loss $L(\theta)$:

$$L_{\text{IAM-D}}(\theta) = L(\theta) + \beta S_\rho(\theta) = L(\theta) + \beta \max_{\|\delta\|_2 \leq \rho} \mathbb{E}_X [\text{KL}(f(X, \theta) \| f(X, \theta + \delta))], \quad (6)$$

where $\beta > 0$ is a hyperparameter balancing the trade-off. This objective seeks parameter values θ for which the model outputs are robust across the neighborhood defined by ρ .

2. SAM-like Approach (IAM-S): Inspired by SAM [8], this method aims to find parameters θ that reside in a neighborhood of uniformly low loss by minimizing the loss at an adversarially perturbed point $\theta + \delta^*$:

$$L_{\text{IAM-S}}(\theta) = L(\theta + \delta^*), \quad \text{where } \delta^* = \arg \max_{\|\delta\|_2 \leq \rho} \mathbb{E}_{\mathcal{X}}[\text{KL}(f(X, \theta) \| f(X, \theta + \delta))]. \quad (7)$$

Here, δ^* is the perturbation that maximizes the local inconsistency term. Note that the objective minimizes the original loss L at the perturbed point $\theta + \delta$:

$$L(\theta + \delta) \approx L(\theta) + \delta^\top \nabla_\theta L(\theta) + \frac{1}{2} \delta^\top G(\theta) \delta.$$

Thus, IAM-S implicitly minimizes the principal eigenvalues of $G(\theta)$, equivalent to empirical FIM.

In the following subsections, we detail the algorithm for IAM-S and provide an analysis of its objective. The algorithm for IAM-D involves a similar inner maximization for $S_\rho(\theta)$ followed by a standard gradient descent step on $L_{\text{IAM-D}}(\theta)$.

5.1 Algorithm for IAM-D and IAM-S

Optimizing $L_{\text{IAM-S}}(\theta)$ and $L_{\text{IAM-D}}(\theta)$ involves a min-max procedure. The inner maximization to find δ^* (i.e., computing $S_\rho(\theta)$ and the corresponding δ^*) is performed using an Algorithm 1, typically for $K = 1$ step for efficiency. IAM-D simply add the $\beta S_\rho(\theta)$ with δ_K to the $L(\theta)$, and then update θ with standard SGD. The outer minimization step of IAM-S updates θ based on the gradient of the loss $L(\theta + \delta_K)$ dropping the second-order terms same with SAM: $\nabla_\theta L_{\text{IAM-S}}(\theta) \approx \nabla_\theta L(\theta)|_{\theta=\theta+\delta_K}$. This two-step process is summarized in Algorithm 2 in Appendix D.

5.2 Empirical evaluation in supervised learning

We evaluated the performance of IAM against SGD and SAM in image classification tasks. WRN16-8[32] served as the baseline model, trained on CIFAR-10, 100 with basic augmentations. Optimal hyperparameters for IAM-D were found to be $\beta = 1.0, \rho = 0.1$ for CIFAR-10, and $\beta = 10.0, \rho = 0.1$ for CIFAR-100, and for IAM-S were $\rho = 0.1, 0.5$ in CIFAR-10 and CIFAR-100 respectively. Table 1 summarizes the test error rates. Both IAM-D (Direct Regularization) and IAM-S (SAM-like Approach) variants not only reduce test error compared to SGD but also achieve performance comparable to SAM. Notably, on CIFAR-100, IAM-S outperforms SAM by a margin of 0.75%, demonstrating its effectiveness in more complex datasets.

Table 1: Test Error (mean $\pm$ stderr) of IAM, SAM, and SGD on WRN-16-8 trained with CIFAR-10, CIFAR-100

	SGD	SAM	IAM-D	IAM-S
CIFAR-10	3.95 $\pm$ 0.048	3.31 $\pm$ 0.010	3.28 $\pm$ 0.060	3.30 $\pm$ 0.042
CIFAR-100	19.17 $\pm$ 0.192	17.63 $\pm$ 0.119	17.16 $\pm$ 0.028	16.88 $\pm$ 0.021

Figure 2 illustrates the evolution of local inconsistency $S_\rho(\theta)$ and test accuracy during training for SGD and IAM-D. IAM-D effectively suppresses the increase in $S_\rho(\theta)$ and mitigates overfitting, particularly evident after learning rate decay points where test accuracy for SGD can degrade. Both on CIFAR-10, 100 (Figure 2), IAM-D maintains $S_\rho(\theta)$ below SGD. Although second LR decay temporarily reduces inconsistency for both, SGD’s inconsistency quickly rebounds, unlike the stable behavior of IAM-D. These observations suggest that minimizing local inconsistency helps confine the model to parameter regions with smoother output distributions, correlating with the generalization improvements shown in Table 1.

5.3 IAM for Learning with Limited or No Explicit Labels

A key advantage of local inconsistency is its computability from unlabeled data, making IAM well-suited for scenarios with limited or no explicit supervision. We demonstrate this in semi-supervised and self-supervised learning settings. Detailed settings are listed in Appendix E.




Figure 2: The evolution of the local inconsistency $S_\rho(\theta)$ and test accuracy with SGD and IAM-D.

Semi-Supervised Learning. In many practical scenarios, labeled data is scarce while unlabeled data is abundant. We simulated semi-supervised learning on CIFAR-10 and CIFAR-100 by masking 80% to 99% of training labels. IAM-D was configured to optimize a joint objective: cross-entropy loss on the labeled subset plus the local inconsistency penalty computed over the entire mini-batch (both labeled and unlabeled examples). In contrast, SGD and SAM utilized only the labeled examples. As summarized in Table 2, IAM-D consistently outperforms both baselines across most missing-label ratios, confirming that leveraging unlabeled data via the local inconsistency term enhances robustness when supervision is sparse.

Table 2: Test Error (mean $\pm$ stderr) of IAM-D, SAM, and SGD on WRN-16-8 trained with CIFAR-10 and CIFAR-100 with varying label rates.

Dataset	Model	Label Rate			
		1%	5%	10%	20%
CIFAR-10	SGD	56.54 $\pm$ 0.159	28.64 $\pm$ 1.648	21.93 $\pm$ 0.234	17.42 $\pm$ 0.430
	SAM	55.83 $\pm$ 0.728	28.45 $\pm$ 0.119	20.68 $\pm$ 0.102	14.00 $\pm$ 0.075
	IAM-D	52.78 $\pm$ 0.497	25.66 $\pm$ 0.723	19.44 $\pm$ 0.354	14.24 $\pm$ 0.142
CIFAR-100	SGD	89.35 $\pm$ 0.098	72.65 $\pm$ 0.519	60.86 $\pm$ 0.204	50.01 $\pm$ 0.299
	SAM	89.31 $\pm$ 0.156	72.02 $\pm$ 0.337	58.04 $\pm$ 0.594	45.64 $\pm$ 0.085
	IAM-D	88.36 $\pm$ 0.292	69.99 $\pm$ 0.649	57.60 $\pm$ 0.251	45.05 $\pm$ 0.956

Self-Supervised Learning (SSL). The label-agnostic nature of IAM makes it directly applicable to SSL objectives. We integrated IAM-D into the SimCLR framework [6], training a ResNet-18[11] encoder on CIFAR-10. Performance was evaluated via linear probing. The local inconsistency term for IAM-D was computed using the model’s projection-head outputs. Figure 3 shows that SimCLR trained with IAM-D (SimCLR-IAM) achieves higher test accuracy on the downstream linear classification task compared to vanilla SimCLR (SimCLR-SGD). Furthermore, SimCLR-IAM tends to converge faster in terms of test error and also minimizes the SimCLR training loss more rapidly, despite the additional local inconsistency regularization. This suggests that controlling local inconsistency is beneficial even when no explicit labels are available during representation learning.




Figure 3: Test accuracy on linear probe and SimCLR training loss for ResNet-18 on CIFAR-10, comparing SimCLR trained with SGD (SimCLR-SGD) versus SimCLR with IAM-D (SimCLR-IAM).

301 5.4 IAM as implicit output entropy regularization

302 Improving generalization in deep neural networks often involves modulating the entropy of model 1
303 output distributions. Techniques such as Label Smoothing (LS) [31, 23, 4] and Entropy Regularization
304 (ER) [27, 5, 34] mitigate the common issue of overconfidence of cross-entropy loss minimization, by
305 specifically targeting this output entropy.

306 Our Inconsistency-Aware Minimization (IAM) framework regularizes local inconsistency, $S_\rho(\theta)$. To 2
307 understand its regularizing effect, we can leverage the identity $\text{KL}(P\|Q) = -H(P) + \text{CE}(P, Q)$,
308 where $H(P)$ is the entropy of P and $\text{CE}(P, Q)$ is cross entropy of P and Q . Since the expectation is
309 linear and $H(f(x; \theta))$ does not depend on δ , we can rewrite S_ρ as:

$$S_\rho(\theta) = -\mathbb{E}_{x \sim p(x)}[H(f(x, \theta))] + \max_{\|\delta\|_2 \leq \rho} \mathbb{E}_{x \sim p(x)}[\text{CE}(f(x, \theta), f(x, \theta + \delta))]. \quad 3 \quad (8)$$

310 Minimizing $S_\rho(\theta)$ as defined in Eq. (8) thus involves two concurrent objectives. First, maximizing 4
311 the expected output entropy $\mathbb{E}_X[H(f(x, \theta))]$. This discourages overconfident (low-entropy) pre-
312 dictions. Second, Minimizing the worst-case expected cross-entropy between the original model’s
313 output $f(x, \theta)$ and the perturbed model’s output $f(x, \theta + \delta)$, which promotes stability of the output
314 distribution under parameter perturbations. This dual objective resonates with the goals of established
315 regularization techniques like LS and ER. LS effectively increases the entropy of the target distribu-
316 tion, thereby preventing the model from becoming overly confident in its predictions for the training
317 labels. ER directly penalizes the low-entropy output distributions of the model. Formally, LS can
318 be interpreted as minimizing a KL divergence to a smoothed target (e.g., $\text{KL}(q_{LS}(y|x) \| p(y|x; \theta))$),
319 while ER often involves minimizing $\text{KL}(p(y|x; \theta) \| u)$, while u is a uniform distribution, both encour-
320 aging the model to avoid excessive output certainty.

321 The first term in Eq. (8) ($-\mathbb{E}_X[H(f(x, \theta))]$) directly aligns with the aim of ER to penalize low- 5
322 entropy outputs. The second term, by minimizing the “distance” (via CE) to the output of a perturbed
323 model, enforces a form of local distributional stability. If highly confident predictions are less stable
324 in parameter perturbation, then minimizing $S_\rho(\theta)$ would implicitly penalize such an overconfidence.

325 Therefore, minimizing local inconsistency $S_\rho(\theta)$ acts as a regularization strategy that not only 6
326 promotes prediction stability against minor parameter variations but also indirectly encourages higher
327 output entropy, akin to mechanisms in LS and ER. This multifaceted regularization is anticipated to
328 yield improved generalization by encouraging the model to learn more robust and less overconfident
329 representations, making them less susceptible to training data idiosyncrasies or parameter instabilities
330 [34, 8].

331 6 Conclusion 7

332 In this work, we introduced “local inconsistency,” a novel information-geometric generalization 8
333 measure computable from a single model using only unlabeled data. We theoretically linked it to the
334 Fisher Information Matrix (FIM) and the loss Hessian. Empirically, local inconsistency correlates
335 with the generalization gap and exhibits distinct characteristics from traditional sharpness-based
336 metrics.

337 Based on this, we proposed Inconsistency-Aware Minimization (IAM), an optimization framework 9
338 that directly incorporates local inconsistency into the training objective. IAM enhances generalization
339 in supervised learning, matching or exceeding that of Sharpness-Aware Minimization (SAM). Cru-
340 cially, IAM proves effective in semi- and self-supervised learning by leveraging unlabeled data for
341 local inconsistency computation, improving performance in label-scarce settings. We also elucidated
342 IAM’s mechanism as an implicit regularizer of model output entropy.

343 These findings offer a practical and theoretically-grounded approach to improving model gener- 10
344 alization, particularly valuable in real-world applications where labeled data is limited. Future
345 research could focus on a more rigorous establishment of the theoretical relationship between local
346 inconsistency and generalization bounds (such as the informal bound presented in Theorem 11) and
347 on exploring the scalability and applicability of IAM to a wider array of model architectures and
348 large-scale datasets.

References

- [1] M. Andriushchenko, F. Croce, M. Müller, M. Hein, and N. Flammarion. A modern look at the relationship between sharpness and generalization, 2023. URL <https://arxiv.org/abs/2302.07011>.
- [2] D. Arpit, S. Jastrzębski, N. Ballas, D. Krueger, E. Bengio, M. S. Kanwal, T. Maharaj, A. Fischer, A. Courville, Y. Bengio, and S. Lacoste-Julien. A closer look at memorization in deep networks, 2017. URL <https://arxiv.org/abs/1706.05394>.
- [3] P. L. Bartlett, D. J. Foster, and M. J. Telgarsky. Spectrally-normalized margin bounds for neural networks. In I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017. URL https://proceedings.neurips.cc/paper_files/paper/2017/file/b22b257ad0519d4500539da3c8bcf4dd-Paper.pdf.
- [4] K. Chandrasegaran, N.-T. Tran, Y. Zhao, and N.-M. Cheung. Revisiting label smoothing and knowledge distillation compatibility: What was missing?, 2022. URL <https://arxiv.org/abs/2206.14532>.
- [5] P. Chaudhari, A. Choromanska, S. Soatto, Y. LeCun, C. Baldassi, C. Borgs, J. Chayes, L. Sagun, and R. Zecchina. Entropy-sgd: Biasing gradient descent into wide valleys, 2017. URL <https://arxiv.org/abs/1611.01838>.
- [6] T. Chen, S. Kornblith, M. Norouzi, and G. Hinton. A simple framework for contrastive learning of visual representations, 2020. URL <https://arxiv.org/abs/2002.05709>.
- [7] L. Dinh, R. Pascanu, S. Bengio, and Y. Bengio. Sharp minima can generalize for deep nets. In D. Precup and Y. W. Teh, editors, *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1019–1028. PMLR, 06–11 Aug 2017. URL <https://proceedings.mlr.press/v70/dinh17b.html>.
- [8] P. Foret, A. Kleiner, H. Mobahi, and B. Neyshabur. Sharpness-aware minimization for efficiently improving generalization, 2021. URL <https://arxiv.org/abs/2010.01412>.
- [9] T. Garipov, P. Izmailov, D. Podoprikin, D. P. Vetrov, and A. G. Wilson. Loss surfaces, mode connectivity, and fast ensembling of dnns. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc., 2018. URL https://proceedings.neurips.cc/paper_files/paper/2018/file/be3087e74e9100d4bc4c6268cdbe8456-Paper.pdf.
- [10] M. Hardt, B. Recht, and Y. Singer. Train faster, generalize better: Stability of stochastic gradient descent. In M. F. Balcan and K. Q. Weinberger, editors, *Proceedings of The 33rd International Conference on Machine Learning*, volume 48 of *Proceedings of Machine Learning Research*, pages 1225–1234, New York, New York, USA, 20–22 Jun 2016. PMLR. URL <https://proceedings.mlr.press/v48/hardt16.html>.
- [11] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition, 2015. URL <https://arxiv.org/abs/1512.03385>.
- [12] C. Jang, S. Lee, F. Park, and Y.-K. Noh. A reparametrization-invariant sharpness measure based on information geometry. *Advances in neural information processing systems*, 35:27893–27905, 2022.
- [13] Y. Jiang, B. Neyshabur, H. Mobahi, D. Krishnan, and S. Bengio. Fantastic generalization measures and where to find them, 2019. URL <https://arxiv.org/abs/1912.02178>.
- [14] Y. Jiang, V. Nagarajan, C. Baek, and J. Z. Kolter. Assessing generalization of sgd via disagreement, 2022. URL <https://arxiv.org/abs/2106.13799>.
- [15] R. Johnson and T. Zhang. Inconsistency, instability, and generalization gap of deep neural network training, 2023. URL <https://arxiv.org/abs/2306.00169>.

- [16] N. S. Keskar, D. Mudigere, J. Nocedal, M. Smelyanskiy, and P. T. P. Tang. On large-batch training for deep learning: Generalization gap and sharp minima, 2017. URL <https://arxiv.org/abs/1609.04836>.
- [17] M. Kim, D. Li, S. X. Hu, and T. M. Hospedales. Fisher sam: Information geometry and sharpness aware minimisation, 2022. URL <https://arxiv.org/abs/2206.04920>.
- [18] A. Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009. URL <https://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf>.
- [19] J. Kwon, J. Kim, H. Park, and I. K. Choi. Asam: Adaptive sharpness-aware minimization for scale-invariant learning of deep neural networks. In M. Meila and T. Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 5905–5914. PMLR, 18–24 Jul 2021. URL <https://proceedings.mlr.press/v139/kwon21b.html>.
- [20] H. Li, Z. Xu, G. Taylor, C. Studer, and T. Goldstein. Visualizing the loss landscape of neural nets. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc., 2018. URL https://proceedings.neurips.cc/paper_files/paper/2018/file/a41b3bb3e6b050b6c9067c67f663b915-Paper.pdf.
- [21] S. Mandt, M. D. Hoffman, and D. M. Blei. Stochastic gradient descent as approximate bayesian inference, 2018. URL <https://arxiv.org/abs/1704.04289>.
- [22] J. Martens. New insights and perspectives on the natural gradient method. *Journal of Machine Learning Research*, 21(146):1–76, 2020. URL <http://jmlr.org/papers/v21/17-678.html>.
- [23] R. Müller, S. Kornblith, and G. Hinton. When does label smoothing help?, 2020. URL <https://arxiv.org/abs/1906.02629>.
- [24] V. Nagarajan and J. Z. Kolter. Uniform convergence may be unable to explain generalization in deep learning. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d’Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/file/05e97c207235d63ceb1db43c60db7bbb-Paper.pdf.
- [25] P. Nakkiran, G. Kaplun, Y. Bansal, T. Yang, B. Barak, and I. Sutskever. Deep double descent: where bigger models and more data hurt*. *Journal of Statistical Mechanics: Theory and Experiment*, 2021(12):124003, dec 2021. doi: 10.1088/1742-5468/ac3a74. URL <https://dx.doi.org/10.1088/1742-5468/ac3a74>.
- [26] R. Pascanu and Y. Bengio. Revisiting natural gradient for deep networks, 2014. URL <https://arxiv.org/abs/1301.3584>.
- [27] G. Pereyra, G. Tucker, J. Chorowski, Łukasz Kaiser, and G. Hinton. Regularizing neural networks by penalizing confident output distributions, 2017. URL <https://arxiv.org/abs/1701.06548>.
- [28] S. Santurkar, D. Tsipras, A. Ilyas, and A. Madry. How does batch normalization help optimization? In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc., 2018. URL https://proceedings.neurips.cc/paper_files/paper/2018/file/905056c1ac1dad141560467e0a99e1cf-Paper.pdf.
- [29] D. Soudry, E. Hoffer, M. S. Nacson, S. Gunasekar, and N. Srebro. The implicit bias of gradient descent on separable data. *Journal of Machine Learning Research*, 19(70):1–57, 2018. URL <http://jmlr.org/papers/v19/18-188.html>.
- [30] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov. Dropout: a simple way to prevent neural networks from overfitting. *J. Mach. Learn. Res.*, 15(1):1929–1958, Jan. 2014. ISSN 1532-4435.

- 446 [31] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens, and Z. Wojna. Rethinking the inception architec- 1
447 ture for computer vision, 2015. URL <https://arxiv.org/abs/1512.00567>.
- 448 [32] S. Zagoruyko and N. Komodakis. Wide residual networks, 2017. URL <https://arxiv.org/abs/1605.07146>.
449
- 450 [33] C. Zhang, S. Bengio, M. Hardt, B. Recht, and O. Vinyals. Understanding deep learning requires
451 rethinking generalization, 2017. URL <https://arxiv.org/abs/1611.03530>.
- 452 [34] H. Zhang, M. Cisse, Y. N. Dauphin, and D. Lopez-Paz. mixup: Beyond empirical risk mini-
453 mization, 2018. URL <https://arxiv.org/abs/1710.09412>.
- 454 [35] J. Zhuang, B. Gong, L. Yuan, Y. Cui, H. Adam, N. Dvornek, S. Tatikonda, J. Duncan, and
455 T. Liu. Surrogate gap minimization improves sharpness-aware training, 2022. URL <https://arxiv.org/abs/2203.08065>.
456

457 A Theoretical Analysis: Generalization Bound with Local Inconsistency

458 In this section, we provide a theoretical sketch to connect our proposed local inconsistency measure, 1
 459 $S_\rho(\theta)$, to the generalization error of a model. Our goal is to show that minimizing $S_\rho(\theta)$, along with
 460 the empirical loss, can lead to a tighter upper bound on the true risk, thereby offering a theoretical
 461 motivation for our Inconsistency-Aware Minimization (IAM) approach. We adapt insights from PAC-
 462 Bayesian theory, particularly drawing from the analysis of Sharpness-Aware Minimization (SAM)
 463 [8].

464 Let $L_{\mathcal{D}}(\theta) = \mathbb{E}_{(X,Y) \sim \mathcal{D}}[l(f(X, \theta), Y)]$ be the true risk and $L_{\mathcal{S}}(\theta) = \frac{1}{n} \sum_{i=1}^n l(f(X_i, \theta), Y_i)$ be the 2
 465 empirical risk on a training set $\mathcal{S}$ of size n . Foret et al. [8] provide a PAC-Bayesian generalization
 466 bound (informally stated, see their Appendix A for the full theorem [Thm 1, 8]) which, with high
 467 probability $1 - \xi$ over the draw of $\mathcal{S}$, states:

$$L_{\mathcal{D}}(\theta) \leq \max_{\|\epsilon\|_2 \leq \rho} L_{\mathcal{S}}(\theta + \epsilon) + \mathcal{R}(\theta, \rho, n, m, \xi) \quad 3 \quad (9)$$

where $\mathcal{R}(\theta, \rho, n, m, \xi)$ is a complexity term that depends on the norm of the parameters $\|\theta\|_2^2$, the 4
 perturbation radius ρ , the number of training samples n , the number of parameters m , and the
 confidence ξ . Specifically,

$$\mathcal{R}(\theta, \rho, n, m, \xi) = \sqrt{\frac{m \log \left(1 + \frac{\|\theta\|_2^2}{\rho^2} \left(1 + \sqrt{\frac{\log n}{m}} \right)^2 \right) + 4 \log \frac{n}{\xi} + \tilde{O}(1)}{n-1}}. \quad 5$$

468 The first term on the right-hand side of Eq. (9) can be rewritten as: 6

$$\max_{\|\epsilon\|_2 \leq \rho} L_{\mathcal{S}}(\theta + \epsilon) = L_{\mathcal{S}}(\theta) + \left[\max_{\|\epsilon\|_2 \leq \rho} (L_{\mathcal{S}}(\theta + \epsilon) - L_{\mathcal{S}}(\theta)) \right]. \quad 7 \quad (10)$$

469 The term in the square brackets, let's call it $\mathcal{SHARP}_\rho(\theta)$, represents the sharpness of the loss 8
 470 landscape at θ within a ρ -neighborhood, as considered by SAM. Our aim is to relate this sharpness
 471 term to our proposed local inconsistency measure $S_\rho(\theta)$.

472 Recall our definition of local inconsistency (calculated with respect to the empirical distribution over 9
 473 the training set $\mathcal{S}$ for this analysis):

$$S_\rho(\theta) = \max_{\|\delta\|_2 \leq \rho} \mathbb{E}_{X \sim \mathcal{S}}[\text{KL}(f(X, \theta) \| f(X, \theta + \delta))]. \quad 10 \quad (11)$$

474 For the cross-entropy loss $l(f(X, \theta), Y) = -\log f(X, \theta)_Y$, under certain conditions, particularly 11
 475 near a good minimizer where the model's predictions $f(X, \theta)$ are close to the true label distribution
 476 (or if we consider $f(X, \theta)$ as a reference distribution), the change in loss due to parameter perturbation
 477 can be related to the KL divergence of the output distributions. Specifically, for a single sample
 478 (X, Y) , a second-order Taylor expansion of the loss $l(f(X, \theta + \delta), Y)$ around $\delta = 0$ gives:

$$L_{\mathcal{S}}(\theta + \delta) - L_{\mathcal{S}}(\theta) = \nabla_{\theta} L_{\mathcal{S}}(\theta)^{\top} \delta + \frac{1}{2} \delta^{\top} \nabla_{\theta}^2 L_{\mathcal{S}}(\theta) \delta + O(\|\delta\|^3). \quad 12 \quad (12)$$

479 If θ is a point where $\nabla_{\theta} L_{\mathcal{S}}(\theta) \approx 0$ (i.e., near a minimum of the empirical risk), the first-order term in 13
 480 the empirical average $L_{\mathcal{S}}(\theta + \delta) - L_{\mathcal{S}}(\theta)$ becomes small. Furthermore, as discussed in Section 4.2
 481 (and your Background section), for cross-entropy loss, the Hessian $\nabla_{\theta}^2 L_{\mathcal{S}}(\theta)$ can be approximated
 482 by the empirical Fisher Information Matrix (or Gauss-Newton matrix) $F_{\mathcal{S}}(\theta)$. Also, we know that
 483 $\mathbb{E}_{X \sim \mathcal{S}}[\text{KL}(f(X, \theta) \| f(X, \theta + \delta))] \approx \frac{1}{2} \delta^{\top} F_{\mathcal{S}}(\theta) \delta$ for small δ .

484 ****Assumption (Approximate Equivalence of Loss Increase and Output KL-Divergence):**** We 14
 485 heuristically assume that, for well-trained models near a local minimum, the worst-case increase in
 486 empirical loss due to parameter perturbation is approximately proportional to the local inconsistency
 487 (our $S_\rho(\theta)$ defined on the empirical data $\mathcal{S}$):

$$\max_{\|\epsilon\|_2 \leq \rho} (L_{\mathcal{S}}(\theta + \epsilon) - L_{\mathcal{S}}(\theta)) \approx c \cdot S_\rho(\theta) \quad 15 \quad (13)$$

488 for some positive constant c , which may depend on factors like the temperature of a Gibbs distribution 16
 489 if one were to formally link the loss to KL divergence from a Bayesian perspective (e.g., $c = 1$ if the

loss essentially measures KL divergence to empirical targets). This assumption relies on the idea that changes in model output distributions (measured by KL divergence) are primary drivers of changes in the CE loss, especially for worst-case perturbations. While this is an approximation, it captures the intuition that models whose outputs are highly sensitive to parameter changes (high $S_\rho(\theta)$) are likely to experience larger increases in loss under such perturbations.

Substituting Eq. (13) into Eq. (10), and then into Eq. (9), we obtain the following generalization bound:

$$L_{\mathcal{D}}(\theta) \lesssim L_S(\theta) + c \cdot S_\rho(\theta) + \mathcal{R}(\theta, \rho, n, m, \xi) \quad (14)$$

where $\lesssim$ indicates that the inequality relies on the approximation in Eq. (13).

Interpretation and Implications. The bound in Eq. (14) suggests that the true risk $L_{\mathcal{D}}(\theta)$ is upper-bounded by the sum of the empirical risk $L_S(\theta)$, our local inconsistency measure $S_\rho(\theta)$ (scaled by a constant c), and the PAC-Bayesian complexity term $\mathcal{R}$. This provides a theoretical rationale for our IAM procedure: by minimizing an objective that includes both the empirical loss and the local inconsistency $S_\rho(\theta)$ (as in IAM-Direct, or implicitly by IAM-S seeking regions where $S_\rho(\theta)$ allows $L_S(\theta + \delta^*)$ to be low), we are effectively attempting to minimize this upper bound on the true risk. A smaller $S_\rho(\theta)$ contributes to a tighter bound, potentially leading to better generalization.

The constant c and the tightness of the approximation in Eq. (13) warrant further investigation. However, this sketch provides a plausible pathway to connect $S_\rho(\theta)$ with generalization guarantees by leveraging existing PAC-Bayesian frameworks that deal with loss landscape geometry. Rigorously establishing the relationship in Eq. (13) or deriving a similar bound with $S_\rho(\theta)$ appearing more directly through the KL term in a PAC-Bayes analysis (perhaps by defining a posterior whose “spread” is related to $S_\rho(\theta)$) are important directions for future theoretical work.

B Relation between our metric and inconsistency

This section outlines an approximate derivation relating the model output inconsistency $\mathcal{C}_P$, as defined by Johnson and Zhang [15], to the local sensitivity metric $S_\rho(w)$ defined previously. We will show simple demonstrations that these two metrics are related primarily through the Fisher Information Matrix (FIM), under specific assumptions like isotropic covariance. Then we will show results with anisotropic covariance.

Definitions

- **Inconsistency ($\mathcal{C}_P$):** Measures the average difference (in terms of KL divergence) between the outputs of models generated by a stochastic training procedure P applied to the same training data Z_n . The average is taken over draws of the training data Z_n and pairs of models (Θ, Θ') drawn from the conditional distribution $\Theta_{P|Z_n}$.

$$\mathcal{C}_P = \mathbb{E}_{Z_n} \mathbb{E}_{\Theta, \Theta' \sim \Theta_{P|Z_n}} \mathbb{E}_X [\text{KL}(f(\Theta, X) \| f(\Theta', X))] \quad (11)$$

Here, $\Theta_{P|Z_n}$ denotes the distribution over parameters resulting from applying procedure P to dataset Z_n .

- **Local Sensitivity ($S_\rho(w)$):** Measures the expected maximum change in the model’s output distribution within a ρ -radius ball around a specific parameter vector w . For consistency with the derivation below, we use the form where the expectation is inside the maximization.

$$S_\rho(w) = \max_{\|\delta\|_2 \leq \rho} \mathbb{E}_X [\text{KL}(f(X, w + \delta) \| f(X, w))] \quad (14)$$

Here, $\delta \in \mathbb{R}^d$ is a perturbation to the parameters w .

Assumptions The following derivation relies on several key assumptions:

1. **Isotropic Covariance Posterior Assumption:** For a given training set Z_n , the conditional parameter distribution $\Theta_{P|Z_n}$ can be approximated by an isotropic distribution centered at a specific parameter vector w_{Z_n} derived from Z_n : $\mathbb{E}[\Theta_{P|Z_n}] = w_{Z_n}$, $\text{Cov}[\Theta_{P|Z_n}] = s^2 \mathbf{I}_d$, where s^2 is a small variance. This approximation is motivated by studies interpreting

533 Stochastic Gradient Descent (SGD) as a form of approximate Bayesian inference, where the 1
 534 distribution of parameters after training can resemble a Gaussian centered near a mode of a
 535 posterior distribution related to the loss function [21].

536 **2. Validity of Second-Order KL Approximation:** The KL divergence between outputs of 2
 537 models with slightly different parameters can be accurately approximated by a quadratic
 538 form involving the Fisher Information Matrix (FIM). This relies on the parameter difference
 539 being small, implying s^2 must be small.

540 **3. Effective FIM Constancy in Expectation:** The variations of the FIM $F(\Theta')$ for $\Theta' \sim$ 3
 541 $\mathcal{N}(w_{Z_n}, s^2 \mathbf{I}_d)$ around $F(w_{Z_n})$ are assumed to average out sufficiently within the expecta-
 542 tion required to calculate $\mathcal{C}_{P|Z_n}$. This allows the approximation $\mathcal{C}_{P|Z_n} \approx s^2 \text{Tr}(F(w_{Z_n}))$.

543 **Approximation of $\mathcal{C}_P$** We first consider the conditional inconsistency for a fixed Z_n , denoted 4
 544 $\mathcal{C}_{P|Z_n}$, by removing the outer expectation $\mathbb{E}_{Z_n}$:

$$\mathcal{C}_{P|Z_n} = \mathbb{E}_{\Theta, \Theta' \sim \Theta_{P|Z_n}} \mathbb{E}_X [\text{KL}(f(\Theta, X) \| f(\Theta', X))] \quad 5$$

545 Applying the isotropic covariance posterior assumption, $\Theta = w_{Z_n} + \delta$ and $\Theta' = w_{Z_n} + \delta'$, where 6
 546 δ, δ' are independent perturbations ($\mathbb{E}[\delta] = \mathbb{E}[\delta'] = 0$, $\text{Cov}[\delta] = \text{Cov}[\delta'] = s^2 \mathbf{I}_d$).

$$\mathcal{C}_{P|Z_n} \approx \mathbb{E}_{\delta, \delta'} \mathbb{E}_X [\text{KL}(f(w_{Z_n} + \delta, X) \| f(w_{Z_n} + \delta', X))] \quad 7$$

547 Using the second-order Taylor expansion for KL divergence taking the expectation over X , valid for 8
 548 small $\|\delta - \delta'\|$ (i.e., small s^2):

$$\mathbb{E}_X [\text{KL}(f(w_{Z_n} + \delta, X) \| f(w_{Z_n} + \delta', X))] = \frac{1}{2} (\delta - \delta')^T F(w_{Z_n} + \delta') (\delta - \delta') + O(\|\delta\|^3) \quad 9$$

549 Let $u = \Theta - \Theta' = \delta - \delta'$. Since δ, δ' are independent, $u \sim \mathcal{N}(0, 2s^2 \mathbf{I}_d)$. Substituting this into the 10
 550 expression for $\mathcal{C}_{P|Z_n}$:

$$\begin{aligned} \mathcal{C}_{P|Z_n} &= \mathbb{E}_u \left[\frac{1}{2} u^T F(\Theta') u \right] + O(\|\delta\|^3) \quad 11 \\ &= \mathbb{E}_u \left[\frac{1}{2} u^T F(w_{Z_n}) u \right] + O(\|\delta\|^3) \quad (\text{FIM Constancy in Expectation Assumption}) \\ &= \frac{1}{2} \text{Tr}(\text{Cov}(u) F(w_{Z_n})) + \frac{1}{2} \mathbb{E}[u]^T F(w_{Z_n}) \mathbb{E}[u] + O(\|\delta\|^3) \\ &= \frac{1}{2} \text{Tr}(2s^2 \mathbf{I}_d F(w_{Z_n})) + 0 + O(\|\delta\|^3) \quad (\mathbb{E}[u] = 0) \\ &\approx s^2 \text{Tr}(F(w_{Z_n})) \end{aligned}$$

551 Thus, the conditional inconsistency for a fixed Z_n is approximately proportional to the trace of the 12
 552 FIM evaluated at w_{Z_n} :

$$\mathcal{C}_{P|Z_n} \approx s^2 \text{Tr}(F(w_{Z_n})) \quad 13 \quad (15)$$

553 The overall inconsistency $\mathcal{C}_P$ is the expectation of this quantity over Z_n : $\mathcal{C}_P \approx \mathbb{E}_{Z_n} [s^2 \text{Tr}(F(w_{Z_n}))]$. 14

554 **Approximation of $S_\rho(w_{Z_n})$** Applying the same second-order KL approximation to the definition 15
 555 of $S_\rho(w_{Z_n})$:

$$S_\rho(w_{Z_n}) = \max_{\|\delta\|_2 \leq \rho} \frac{1}{2} \delta^\top F(w_{Z_n}) \delta + O(\|\delta\|^3) \quad 16$$

556 The maximum value of the quadratic form $\delta^T A \delta$ for a positive semi-definite matrix A subject to 17
 557 $\|\delta\|_2 \leq \rho$ is achieved when δ is aligned with the eigenvector corresponding to the largest eigenvalue
 558 ($\lambda_{\max}(A)$) and has norm ρ . Thus:

$$S_\rho(w_{Z_n}) = \frac{1}{2} \rho^2 \lambda_{\max}(F(w_{Z_n})) \quad 18 \quad (16)$$

559 This shows that the local sensitivity S_ρ is approximately proportional to the largest eigenvalue of the 19
 560 FIM.

561 **Connecting $\mathcal{C}_{P|Z_n}$ and $S_\rho(w_{Z_n})$** For a $d \times d$ positive semi-definite matrix A , the relationship 1
 562 between its trace and largest eigenvalue is given by $\frac{1}{d}\text{Tr}(A) \leq \lambda_{\max}(A) \leq \text{Tr}(A)$. Applying this to
 563 the FIM $F(w_{Z_n})$:

$$\frac{1}{d}\text{Tr}(F(w_{Z_n})) \leq \lambda_{\max}(F(w_{Z_n})) \leq \text{Tr}(F(w_{Z_n})) \quad 2$$

564 Substituting this into the approximation for $S_\rho(w_{Z_n})$ from Eq. (16): 3

$$\frac{\rho^2}{2d}\text{Tr}(F(w_{Z_n})) \leq S_\rho(w_{Z_n}) \leq \frac{\rho^2}{2}\text{Tr}(F(w_{Z_n})) \quad 4$$

565 Let's assume a plausible connection, for instance, $s^2 = \rho^2/d$. Substituting this into the approximation 5
 566 for $\mathcal{C}_{P|Z_n}$ from Eq. (15), we get $\mathcal{C}_{P|Z_n} \approx \frac{\rho^2}{d}\text{Tr}(F(w_{Z_n}))$. Combining this with the bounds for
 567 $S_\rho(w_{Z_n})$:

$$\frac{1}{2}\left(\frac{\rho^2}{d}\text{Tr}(F(w_{Z_n}))\right) \leq S_\rho(w_{Z_n}) \leq \frac{d}{2}\left(\frac{\rho^2}{d}\text{Tr}(F(w_{Z_n}))\right) \quad 6$$

568 This leads to the final approximate relationship between the conditional inconsistency (for a fixed 7
 569 Z_n) and the local sensitivity (at the corresponding w_{Z_n}):

$$\frac{1}{2}\mathcal{C}_{P|Z_n} \leq S_\rho(w_{Z_n}) \leq \frac{d}{2}\mathcal{C}_{P|Z_n} \quad 8 \quad (17)$$

570 This result suggests that, under the stated assumptions, the conditional inconsistency $\mathcal{C}_{P|Z_n}$ and the 9
 571 local sensitivity $S_\rho(w_{Z_n})$ are approximately proportional, with the proportionality factor potentially
 572 depending on the parameter dimension d .

anisotropic covariance Let $\text{Cov}[\Theta_{P|Z_n}] = s^2\Sigma$, where $s^2 = \frac{\rho^2}{d}$. Starting from $\mathcal{C}_{P|Z_n} =$ 10
 $\frac{1}{2}\text{Tr}(\Sigma F(w_{Z_n}))$,

$$\begin{aligned} \lambda_{\min}(\Sigma)\text{Tr}(F) &\leq \text{Tr}(\Sigma F) \leq \lambda_{\max}(\Sigma)\text{Tr}(F) & 11 \\ \lambda_{\min}(\Sigma)\lambda_{\max}(F) &\leq \text{Tr}(\Sigma F) \leq \lambda_{\max}(\Sigma)d\lambda_{\max}(F) \\ \frac{\rho^2}{2d\lambda_{\max}(\Sigma)}\text{Tr}(\Sigma F) &\leq \frac{\rho^2}{2}\lambda_{\max}(F) \leq \frac{\rho^2}{2\lambda_{\min}(\Sigma)}\text{Tr}(\Sigma F) \\ \frac{1}{\lambda_{\max}(\Sigma)}\mathcal{C}_{P|Z_n} &\leq S_\rho(w_{Z_n}) \leq \frac{d}{\lambda_{\min}(\Sigma)}\mathcal{C}_{P|Z_n} \end{aligned}$$

573 **Practical Considerations: Eigenvalue Spectrum of Neural Networks** In practice, for deep 12
 574 learning models, the FIM often exhibits a sparse eigenvalue spectrum: many eigenvalues are close to
 575 zero, and only a few are significantly large. In such cases:

- 576 • The trace $\text{Tr}(F) = \sum \lambda_i$ is dominated by the sum of the few large eigenvalues. 13
- 577 • The ratio $\lambda_{\max}(F)/\text{Tr}(F)$ might be closer to $1/m'$ than $1/d$, where $m' \ll d$ is the “effective
 578 rank” or number of dominant eigenvalues.

579 This implies that the bounds relating $\lambda_{\max}(F)$ and $\text{Tr}(F)$ might be tighter than the general $1/d$ and 14
 580 factors suggest. Consequently, the relationship between $\mathcal{C}_{P|Z_n}$ (related to trace) and S_ρ (related to
 581 max eigenvalue) could be closer to direct proportionality than Eq. (5) indicates, especially if s^2 is
 582 appropriately related to ρ^2 .

583 **Summary and Limitations** This analysis provides a heuristic argument suggesting a connection 15
 584 between conditional inconsistency $\mathcal{C}_{P|Z_n}$ and local sensitivity $S_\rho(w_{Z_n})$. Under assumptions of
 585 a Gaussian posterior, small variance s^2 , validity of second-order KL approximations, local FIM
 586 constancy, and a specific link between s^2 and ρ^2 (e.g., $s^2 = \rho^2/d$), we find that $S_\rho(w_{Z_n})$ is
 587 approximately proportional to $\mathcal{C}_{P|Z_n}$, potentially up to a factor related to dimension d . This connection
 588 is mediated by the trace and the maximum eigenvalue of the Fisher Information Matrix. The practical
 589 observation of sparse FIM eigenvalues might strengthen this relationship.

590 C Decision boundary of neural networks and principal eigenspace of FIM

591 To intuitively analysis the role of δ_1 in training of neural network, we conducted experiments using 1
 592 3-layer fully-connected neural network on two-dimensional synthetic data. the data is generated
 593 from a mixture of three Gaussian distributions, a setup analogous to that employed by [12] in their
 594 investigation of the characteristic of the FIM eigensubspace. Their work demonstrated that perturbing
 595 parameters along the principal eigenvectors of the FIM can lead to significant modifications in the
 decision boundary, such as increasing or decreasing the margins of specific classes.




Figure 4: A synthetic classification example. the black, blue, orange lines correspond to decision 3
 boundaries of the NN with trained parameter values, and parameter values perturbed by δ_1 . Each plot
 use different noise.

596 Our investigation focuses on whether δ_1 , despite being derived from only a single gradient step 4
 597 (as described in Algorithm 1) and thus influenced by an initial random noise vector ε , still induces
 598 substantial changes in the neural network’s decision boundary. Figure 4 visualizes these effects.
 599 The black lines in each subfigures depict the original decision boundary obtained with the trained
 600 parameters w . Figure 4 (a-c) show the perturbed decision boundaries (blue and orange lines) when
 601 distinct $\pm\delta_1$ with $\rho = 0.5$ is added to w . Each of these δ_1 vectors was computed using a different
 602 random initialization noise vector, denoted as $\varepsilon_1, \varepsilon_2$, and ε_3 , respectively. For a direct comparison of
 603 the perturbation’s nature, Figure 4(d) illustrates the decision boundary perturbed by directly adding
 604 the random noise vector ε to w . This vector ε is sampled from same distribution as initial vectors
 605 (e.g. ε_1) and, is scaled to $\|\varepsilon\|_2 = \rho$ same with δ_1 . As observed in Figure 4 (d), direct perturbation with
 606 such an arbitrary random noise vector does not meaning fully alter the decision boundary, even when
 607 its norm is equivalent to that of the δ_1 . This is sharply opposed with the significant changes induced
 608 by δ_1 perturbations shown in Figures 4 (a-c), underscoring that the direction derived by Algorithm
 609 1, even in a single step, is substantially more influential than arbitrary noise of the same magnitude.
 610 This result intuitively suggest that the perturbation δ_1 with single gradient step still meaningful and
 611 aligning with principle eigen vectors of FIM.
 612

613 To investigate the alignment between the single-step perturbation vector δ_1 and principle eigenspace 5
 614 of FIM, we explicitly calculate the FIM and its top three eigenvector v_1, v_2 , and v_3 , corresponding
 615 to largesst eigenvalues $\lambda_1 > \lambda_2 > \lambda_3$. The perturbation δ_1 , results from one normalized gradient
 616 ascent step applied to the KL divergence objective, starting from an initial random noise ε . In terms
 617 of power iteration algorithm, the δ_1 after first iteration without normalization, is sum of eigenvector
 618 of FIM weighted by $\lambda_i \alpha_i$.

619 Formally, let the initial random noise ε be expressed in the eigenbasis of $F(w)$ as $\varepsilon = \sum_i^m \alpha_i v_i$. 6
 620 $\varepsilon \sim \mathcal{N}(0, \sigma^2 I_m)$, then the coefficient α_i are i.i.d. as $\mathcal{N}(0, \sigma^2)$ since $\{v_i\}$ form an orthonormal basis.

$$\begin{aligned}
 F(\theta)\varepsilon &= \sum_i^m \lambda_i v_i v_i^\top \sum_i^m \alpha_i v_i \\
 &= \sum_i^m \lambda_i \alpha_i v_i
 \end{aligned}$$

621 So cosine similarity between δ_1 and v_i is $\lambda_i \alpha_i$. And $\frac{\|\sum_i^3 \delta_1^T v_i\|}{\|\delta_1\|}$, which indicates how much the δ_1 is
 622 in principle eigen space, $\{u|u = av_1 + bv_2 + cv_3, \quad abc \in [0, 1]\}$ of FIM, is $\frac{\|\sum_i^3 \alpha_i \lambda_i v_i\|}{\|\delta\|}$.




Figure 5: A synthetic classification example. δ_1 are align with top three eigen Vector of FIM sampling from 10000 gaussian noises ε

623 Figure 5 presents empirical results from this analysis. Figure 5 (a) shows histograms of the absolute
 624 cosine similarities between δ_1 (generated from 10,000 different ε samples) and each of the top
 625 three eigenvectors v_1, v_2 , and v_3 . We observe that δ_1 tends to have a higher cosine similarity with
 626 v_1 (corresponding to the largest eigenvalue λ_1) compared to v_2 , and v_3 . Furthermore, Figure 5
 627 (b) displays the distribution of the squared norm of the projection of δ_1 onto the top-3 eigenspace.
 628 The values are predominantly close to 1, indicating that δ vectors derived from different initial
 629 noise samples are largely confined to this principal subspace. These results empirically support the
 630 theoretical expectation that the single-step perturbation δ_1 is predominantly aligned with the principal
 631 eigenspace of the FIM.

632 D IAM algorithms 5

Algorithm 2 Inconsistency-Aware Minimization (SAM-like variant: IAM-S) 6

1: **Input:** Initial model parameters θ^0 ; Learning rate η ; neighborhood size ρ ; training set Z_n ; Batch
 size b ; Number of steps K for Algorithm 1.
 2: **while** not converged **do**
 3: Sample batch $\{(x_i, y_i)\}_{i=1}^b$.
 4: Compute δ_K from Algorithm 1 using current θ, ρ , and data $\{x_i\}_{i=1}^b$.
 5: Compute gradient $g = \nabla_{\theta} L(\theta)|_{\theta+\delta_K}$
 6: Update parameters: $\theta \leftarrow \theta - \eta g$.
 7: **end while**
 8: **Return** optimized parameters θ .

633 E Experimental Details 8

634 Practical Considerations in estimating $S_{\rho}(\theta)$ 9

- 635 • **Computational Efficiency:** Calculating the FIM explicitly and performing eigenvalue
 636 decomposition is computationally expensive ($O(m^2)$ or worse, where m is the number of
 637 parameters). Algorithm 1 avoids this by requiring only K gradient computations (forward
 638 and backward passes) per estimation, making its computational cost approximately $O(mK)$,
 639 which is significantly more feasible for large networks. 10

- **Number of Steps (K):** Empirical studies on neural network Hessians and FIMs suggest that the eigenspectrum is often dominated by a huge largest eigenvalues. Thus, the Power Iteration method can converge quickly to the dominant eigenvector. In practice, using a small number of steps, often just $K = 3$, is found to be sufficient to get a reasonable estimate of the maximizing direction. This makes the computation highly efficient.
- **Averaging for reduce Variance from initialization:** The estimate of $S_\rho(w)$ obtained from Algorithm 1 depends on the random initialization δ_0 with just $K = 1$. To obtain a more stable estimate, we compute the metric multiple times (e.g., 10 times) with different random initializations for δ_0 and report the average value: $\mathbb{E}_{\delta_0}[\text{Estimate from Alg 1}]$.

Infrastructure Experiments are implemented in PyTorch2.7 and executed on NVIDIA A40 and L4 GPUs.

E.1 Image classification

Each reported metric is the mean $\pm$ standard error computed over minimum test error from three independent runs.

Dataset. We evaluate on the **CIFAR-10** (50,000 training, 10,000 test images) and **CIFAR-100** (50,000 training, 10,000 test images). All images are resized to 32×32 and preprocessed with

- *RandomCrop*(32, padding= 4),
- *RandomHorizontalFlip*($p = 0.5$), and
- *Normalization* using the official mean and standard deviation.

No additional augmentation such as Cutout or Mixup is applied.

Optimization. Models are trained for **200 epochs** with mini-batch size **128**. We use SGD with momentum 0.9, weight decay 5×10^{-4} as an optimizer, and a multistep learning rate schedule that decays the initial rate 0.1 by 0.2 at epochs 60, 120, and 160.

Hyperparameters. The inconsistency weight β and neighborhood radius ρ are selected from $\beta \in \{0.1, 1.0, 5.0, 10.0, 20.0\}$ and $\rho \in \{0.01, 0.05, 0.1, 0.5, 1.0\}$ via grid search on the validation split using 10% of the training dataset. The best pairs are (1.0, 0.1) for CIFAR-10 and (10.0, 0.1) for CIFAR-100. For IAM-S, 0.1 and 0.5 were selected ρ value for CIFAR-10, 100 respectively.

Loss function. Cross-entropy with label smoothing ($\alpha = 0.1$) is used for all methods.

E.2 Semi-supervised learning

In semi-supervised learning experiment, we shared most of the settings with image classification. Each reported metric computed over minimum test error from three independent runs.

Optimization. Models are trained for **100 epochs** without learning rate scheduling.

Hyperparameters. We used $\beta = 1.0$ and $\rho = 0.1$ for both CIFAR-10 and CIFAR-100. SAM is also trained with $\rho = 0.1$.

E.3 Self-supervised learning

Each reported metric is the mean **test accuracy** obtained from three independent runs.

Dataset. We use the **CIFAR-10** benchmark. All images are resized to 32×32 and augmented with the SimCLR[6] pipeline:

- *RandomResizedCrop*(32, scale=(0.4, 1.0)),
- *RandomHorizontalFlip*($p = 0.5$),

- 680 • *ColorJitter*(0.4, 0.4, 0.2, 0.1) with probability 0.8, 1
681 • *RandomGrayscale*($p=0.2$), and
682 • *Normalization* using the official mean and standard deviation.

683 **Encoder&Projection Head.** We adopt a **ResNet-18** backbone with the first convolution modified 2
684 to 3×3 layer with stride = 1 and the max-pool removed. The projector is a two-layer MLP (hidden
685 size 512, output size 128) with ReLU activation.

686 **Optimization.** Models are trained for **200 epochs** with mini-batch size **1024**. We use SGD 3
687 (momentum 0.9, weight decay 1×10^{-4}) and a cosine-annealing learning-rate schedule starting at 1.0
688 after a 10-epoch warm-up.

689 **Contrastive Loss.** The NT-Xent loss is computed with temperature $\tau=0.5$. 4

690 **IAM Hyperparameters.** We set the inconsistency weight $\beta=1.0$, neighborhood radius $\rho=0.1$, and 5
691 noise-scale 3.0 (Gaussian initialization). The local inconsistency is computed between projection
692 head outputs with temperature $\tau=0.5$.

693 **Stability Heuristics.** Is identical to image classification setting. 6

694 **Linear Evaluation.** After every 5 epochs (and at the final epoch), a frozen encoder is evaluated via 7
695 a linear probe trained for 20 epochs with AdamW optimizer on the full training set (batch size 1024).
696 The reported metric is the probe’s test accuracy.

NeurIPS Paper Checklist

1. Claims 1

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? 2

Answer: [Yes] 3

Justification: The need to consider unlabeled and single model generalization measure is discussed in Section 1, 2. 4

Guidelines: 5

- The answer NA means that the abstract and introduction do not include the claims made in the paper. 6
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations 7

Question: Does the paper discuss the limitations of the work performed by the authors? 8

Answer: [Yes] 9

Justification: See supplemental materials. 10

Guidelines: 11

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper. 12
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs 13

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof? 14

Answer: [Yes] 15

Justification: See Appendix B, D. we provides the assumptions and proofs. 1

Guidelines: 2

- The answer NA means that the paper does not include theoretical results. 3
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility 4

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)? 5

Answer: [Yes] 6

Justification: See supplemental materials. we provide the codes for experiments. 7

Guidelines: 8

- The answer NA means that the paper does not include experiments. 9
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code 10

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material? 11

Answer: [Yes] 1

Justification: See supplemental materials. 2

Guidelines: 3

- The answer NA means that paper does not include experiments requiring code. 4
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details 5

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results? 6

Answer: [Yes] 7

Justification: See Appendix E 8

Guidelines: 9

- The answer NA means that the paper does not include experiments. 10
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance 11

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments? 12

Answer: [Yes] 13

Justification: See Section 5 14

Guidelines: 15

- The answer NA means that the paper does not include experiments. 16
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.

- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources 2

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments? 3

Answer: [Yes] 4

Justification: See Appendix E 5

Guidelines: 6

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics 8

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>? 9

Answer: [Yes] 10

Justification: Yes, we reviewed the Code of Ethics. 11

Guidelines: 12

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts 14

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed? 15

Answer: [NA] 16

Justification: there is no societal impact of the work performed 17

Guidelines: 18

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to

generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.

- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards 2

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)? 3

Answer: [NA] 4

Justification: The paper presents a theoretical analysis and does not release any data or models that carry a risk of misuse. 5

Guidelines: 6

- The answer NA means that the paper poses no such risks. 7
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets 8

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected? 9

Answer: [Yes] 10

Justification: See supplement materials. 11

Guidelines: 12

- The answer NA means that the paper does not use existing assets. 13
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets? 1

Answer: [Yes] 2

Justification: See supplemental materials. 3

Guidelines: 4

- The answer NA means that the paper does not release new assets. 5
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects 6

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)? 7

Answer: [NA] 8

Justification: the paper does not involve crowdsourcing nor research with human subjects. 9

Guidelines: 10

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 11
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects 12

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained? 13

Answer: [NA] 14

Justification: The paper does not involve crowdsourcing nor research with human subjects. 15

Guidelines: 16

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 17
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage 18

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigor, or originality of the research, declaration is not required. 19

1010 Answer: [NA] 1
1011 Justification: LLMs was not used to core method development. 2
1012 Guidelines: 3
1013 • The answer NA means that the core method development in this research does not 4
1014 involve LLMs as any important, original, or non-standard components.
1015 • Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) 5
1016 for what should or should not be described.