
Axiomatic Characterization of the Hamming and Jaccard Distances

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Measures of dissimilarity between a pair of objects can play a pivotal role in many
2 machine learning objectives such as clustering, outlier detection, or data visualiza-
3 tion. We focus on data in the form of binary vectors and analyze several methods
4 of measuring dissimilarity between them. We introduce several properties, *axioms*,
5 that a measure of dissimilarity can satisfy and characterize the *Hamming* and *Jac-*
6 *card* distances as the only measures satisfying particular subsets of our axioms.
7 Based on our analysis, we identify shortcomings of both distances, and propose
8 novel approaches that are better suited for certain applications. We complement
9 our theoretical findings by extensive empirical study. Our primary motivation is
10 the analysis of election data, in which the votes have the form of binary approval
11 of alternatives, but the applicability of our results reaches far beyond that.

12 1 Introduction

13 The Hamming and Jaccard distances find various applications throughout and beyond machine
14 learning. For example, they can play a key role in clustering [16, 15], keywords similarity analy-
15 sis [22], and recommendation systems [5]. Jaccard distance also serves as a loss function for object
16 detection in computer vision [12, 25, 30], while Hamming is often used in multi-label classification
17 problems [28, 27]. However, in many applications, it is unclear why this particular metric should
18 be used out of many possible measures of dissimilarity between binary vectors. This question
19 is of dire importance as the choice of the metric can lead to drastic differences in the produced
20 results [9, 25, 15].

21 To provide a deeper understanding of the nature of each measure and to offer guidance in selecting
22 an appropriate one for a given application, we adopt the *axiomatic method*. This approach involves
23 introducing simple and intuitive properties, called *axioms*, which characterize particular measures,
24 showcasing the distinctive behavior of each measure and highlighting the similarities and differences
25 between them. Axioms have been already employed in the theoretical analysis of machine learning
26 tools such as clustering [1, 3, 18], or classification [11]. Furthermore, they are cornerstones of the
27 related fields of computational social choice [7] and game theory [26].

28 Specifically, we introduce five *invariance* axioms (namely, *Anonymity*, *Independent Symmetry*, *Add*
29 *Zero*, *Zero-One Symmetry*, and *Scaling*) each requiring that a particular simple operation applied to
30 a pair of binary vectors does not affect the dissimilarity between them. Additionally, we consider
31 *Triangle Inequality*—a standard distance axiom; *Convergence* which asserts that the distance between
32 the vectors decreases as we introduce agreement between them; and *Normalization* that fixes the
33 dissimilarity value in a certain corner case.

34 We prove that the Hamming distance is the unique dissimilarity measure that satisfies Anonymity,
35 Independent Symmetry, Zero-One Symmetry, Scaling, Triangle Inequality, Convergence, and Nor-
36 malization axioms. Moreover, if in this characterization we exchange Zero-One Symmetry for Add

Table 1: Which dissimilarity measure satisfies which axiom.

	Anonym.	Scaling	Indep. Sym.	0-1 Sym.	Conv.	Tri. Ineq.	Norm.	Add 0
Hamming	✓	✓	✓	✓	✓	✓	✓	✗
Jaccard	✓	✓	✓	✗	✓	✓	✓	✓

Zero axiom, then such set of axioms uniquely characterize the Jaccard distance. This axiomatic approach thus offers a comprehensive comparison of both measures highlighting the key difference between them. Satisfaction of our axioms by the rules we consider is summarized in Table 1.

Furthermore, based on our analysis, we identify a common feature of the Hamming and Jaccard distances that may be undesirable in many scenarios. Indeed, both distances satisfy Independent Symmetry, which in practice implies that the vectors of different saturation will always be far from each other. We propose the modifications of both measures that escape this problem and show that they can be more suited for certain applications.

We complement our theoretical findings with extensive experiments on approval election data, in which a set of voters express binary preferences over a set of alternatives. Such data is one of the primary objects of study in the computational social choice (for more details see the work of Boehmer et al. [6]) and shows a high level of asymmetry: A voters' approval for an alternative is a much stronger signal than the lack thereof. We compare the behavior of the studied dissimilarity measures in several canonical classes of such elections and finish with the case study on real-world election instances from participatory budgeting. Our analysis confirms that the newly proposed measures can indeed be more efficient in certain settings. All missing proofs are available in Appendix A.

2 Preliminaries

For every $n \in \mathbb{N}$, we denote $[n] = \{1, \dots, n\}$. For two sets A, B , we denote their *symmetric difference* by $A \triangle B = (A \cup B) \setminus (A \cap B)$.

2.1 Binary Vectors

We consider a space of *binary vectors* of arbitrary (but finite) length. Each such vector, $x \in \{0, 1\}^n$, can equivalently be viewed as a subset of coordinates on which there are ones in the vector, which we denote by $A_x = \{i \in [n] : x_i = 1\}$. For a permutation $\pi : [n] \rightarrow [n]$, by $\pi(x)$ we denote the vector x with changed order of coordinates, i.e., $\pi(x)_i = x_{\pi(i)}$ for every $i \in [n]$. Drawing from game theory notation, for vector $x \in \{0, 1\}^n$, coordinate $i \in [n]$, and $b \in \{0, 1\}$ we write $y = (x_{-i}, b)$ to denote a vector such that $y_i = b$ and $y_j = x_j$ for every $j \in [n] \setminus \{i\}$. For vector $x \in \{0, 1\}^n$ by $\bar{x}$ we denote its bitwise negation, i.e., $\bar{x} = (1 - x_1, \dots, 1 - x_n)$. By $\circ$ we denote *concatenation* of vectors, i.e., for $x \in \{0, 1\}^n$ and $y \in \{0, 1\}^m$, if $z = x \circ y$, then $z_i = x_i$ for $i \in [n]$ and $z_i = y_i$ for every $i \in [m + n] \setminus [n]$. Finally, by x^k we denote the result of concatenating k copies of x one after another. For convenience, we will allow for x^0 to be the empty vector, which concatenated with any other vector y , results in y .

2.2 Dissimilarity Measures

A *dissimilarity measure* is a function that takes a pair of binary vectors of the same size as arguments and outputs some real nonnegative values, i.e., $f : (\{0, 1\}^n \times \{0, 1\}^n)_{n \in \mathbb{N}} \rightarrow \mathbb{R}_{\geq}$.

Examples include the (normalized) *Hamming distance*, also known as the ℓ_1 distance, which outputs the fraction of positions at which the corresponding elements in two vectors differ, i.e.,

$$H(x, y) = |A_x \triangle A_y| / n.$$

Jaccard distance is defined similarly, but instead of normalizing by the vector length, it divides the symmetric difference by the number of positions in which at least one vector has a positive entry, i.e.,

$$J(x, y) = |A_x \triangle A_y| / |A_x \cup A_y|.$$

Other examples include the *Euclidean distance* (or ℓ_2), which is just a square root of the Hamming distance, or the *discrete distance* (or ℓ_∞) that returns 0 if the vectors are identical and 1 otherwise.

73 3 Axiomatic Characterization of the Hamming Distance

74 We begin by introducing seven axioms that uniquely characterize the Hamming distance, arguably
75 the most popular notion of dissimilarity between binary vectors.

76 The first four axioms are so-called *invariance axioms*, which describe certain operations on a pair of
77 vectors and assert that this operation should not affect the dissimilarity value. This first such axiom is
78 called *Anonymity* and intuitively states that the order of coordinates does not matter. It is natural in
79 settings where there is no spatial or temporal relation between the bits of information that the vectors
80 represent. Rather, they come from independent sources, as, for example, when the vectors represent
81 support for a given candidate by various voters (indeed, the version of this axiom often appears in the
82 social choice literature [21]).

83 **Definition 3.1** (Anonymity). A dissimilarity measure, f , satisfies *Anonymity* if for every vectors
84 $x, y \in \{0, 1\}^n$ and permutation $\pi : [n] \rightarrow [n]$ it holds that

$$f(\pi(x), \pi(y)) = f(x, y).$$

85 For example, Anonymity would imply that $f((1, 0, 1, 1), (1, 1, 0, 0)) = f((0, 1, 1, 1), (1, 0, 1, 0))$.

86 The second axiom, *Scaling*, captures the intuition that the scale of the input should not be relevant
87 for the dissimilarity. Hence, we can copy the entries of the vectors several times and this will not
88 affect the value. This property was noted as crucial for some objectives such as measuring accuracy
89 in object detection tasks in computer vision [25].

90 **Definition 3.2** (Scaling). A dissimilarity measure, f , satisfies *Scaling* if for every vectors $x, y \in$
91 $\{0, 1\}^n$ and $k \in \mathbb{N}$ it holds that

$$f(x^k, y^k) = f(x, y).$$

92 For example, Scaling would imply that $f((1, 0), (1, 1)) = f((1, 0, 1, 0), (1, 1, 1, 1))$.

93 Our next axiom is a strengthening of a standard distance axiom of Symmetry. *Independent Symmetry*
94 states that on every coordinate we can switch the values between the two vectors, and this operation
95 should preserve the dissimilarity value.

96 **Definition 3.3** (Independent Symmetry). A dissimilarity measure, f , satisfies *Independent Symmetry*
97 if for every vectors $x, y \in \{0, 1\}^n$ and index $i \in [n]$ it holds that

$$f((x_{-i}, y_i), (y_{-i}, x_i)) = f(x, y).$$

98 For example, Independent Symmetry would imply that $f((1, 0, 1), (1, 1, 0)) = f((1, 1, 1), (1, 0, 0))$.

99 Our final invariance axiom, *Zero-One Symmetry*, states that the roles of 0 and 1 in the vectors are
100 symmetric, hence if we exchange all 0s for 1s in both vectors and vice-versa, we should get the same
101 dissimilarity. Indeed, in some applications the roles of zeros and ones are arbitrary.

102 **Definition 3.4** (Zero-One Symmetry). A dissimilarity measure, f , satisfies *Zero-one Symmetry* if for
103 every vectors $x, y \in \{0, 1\}^n$ it holds that

$$f(\bar{x}, \bar{y}) = f(x, y).$$

104 For example, Zero-One Symmetry would imply that $f((1, 0, 1), (1, 1, 0)) = f((0, 1, 0), (0, 0, 1))$.

105 Our next axiom is *Convergence*. It captures the intuition that as we add new information to the input
106 on which both vectors agree, then the dissimilarity between them should move towards zero, and that
107 this convergence to zero should be quick enough. Formally, we say that if we double the size of the
108 vectors and in the new entries there are only 1s in both of them, then the distance should decrease by
109 at least half.

110 **Definition 3.5** (Convergence). A dissimilarity measure, f , satisfies *Convergence* if for every vectors
111 $x, y \in \{0, 1\}^n$ it holds that

$$f(x \circ (1)^n, y \circ (1)^n) \leq \frac{1}{2} f(x, y).$$

112 For example, Convergence would imply that $f((1, 0, 1, 1), (1, 1, 1, 1)) \leq f((1, 0), (1, 1))/2$.

113 Next, we include a standard distance metric axiom of *Triangle Inequality*.

114 **Definition 3.6** (Triangle Inequality). A dissimilarity measure, f , satisfies *Triangle Inequality* if for
 115 every vectors $x, y, z \in \{0, 1\}^n$ it holds that

$$f(x, y) + f(y, z) \geq f(x, z).$$

116 Finally, *Normalization* specifies the dissimilarity in a basic case of length-1 vectors.

117 **Definition 3.7** (Normalization). A dissimilarity measure, f , satisfies *Normalization* if

$$f((0), (1)) = 1.$$

118 The above seven axioms uniquely characterize the Hamming distance.

119 **Theorem 3.8.** A dissimilarity measure, f , satisfies *Anonymity*, *Scaling*, *Independent Symmetry*,
 120 *Zero-One Symmetry*, *Convergence*, *Triangle Inequality*, and *Normalization* if and only if f is the
 121 *Hamming distance*.

122 *Proof sketch.* The proofs that Hamming satisfies each of the axioms are relatively straightforward.
 123 Thus, in this sketch we focus on showing that if a dissimilarity measure f satisfies all of the axioms,
 124 then it is the Hamming distance (the full proof can be found in Appendix A). The proof proceeds by
 125 considering growing subsets of axioms and characterizing the classes of dissimilarity measures that
 126 satisfy these axioms (in Appendix A, we present each such characterization as a separate lemma).

127 We begin by observing that dissimilarity measure satisfying *Anonymity* must be in fact a function
 128 of four arguments, which are the sizes of the intersection, both differences, and the intersection
 129 of complements of sets A_x and A_y for each pair of vectors x, y . Then, if we add *Independent*
 130 *Symmetry*, this function must depend on only three arguments, as instead of looking at $|A_x \setminus A_y|$
 131 and $|A_y \setminus A_x|$, we can simply look at the size of their symmetric difference. Adding *Zero-One*
 132 *symmetry* prevents distinguishing the intersection from the intersection of complements, so we end
 133 up with two arguments: $|A_x \triangle A_y|$ and $n - |A_x \triangle A_y|$, or, equivalently, $|A_x \triangle A_y|$ and n . Then,
 134 *Scaling* implies that $f(x, y)$ is actually a function of a single argument, i.e., there exists g such that
 135 $f(x, y) = g(|A_x \triangle A_y|/n) = g(H(x, y))$.

136 We use the remaining axioms to establish properties of g . *Triangle Inequality* means that g is
 137 subadditive. Next, we prove that with the addition of *Convergence*, g has to be a linear homogeneous
 138 function. In other words, there is $a \in \mathbb{R}$ such that $f(x, y) = a \cdot H(x, y)$. Then, adding *Normalization*
 139 concludes the proof. $\square$

140 We note that the set of axioms used for our characterization of the Hamming distance is minimal. In
 141 other words, the axioms are *independent*, as no subset of our axioms implies the remaining ones.

142 **Theorem 3.9.** For every axiom in the set *Anonymity*, *Scaling*, *Independent Symmetry*, *Zero-One*
 143 *Symmetry*, *Convergence*, *Triangle Inequality*, and *Normalization* there is a dissimilarity measure that
 144 satisfies all other axioms in this set except for this one.

145 4 Axiomatic Characterization of the Jaccard Distance

146 As mentioned in the previous section, *Zero-One Symmetry*, an axiom characterizing Hamming,
 147 requires that 0s and 1s have symmetric roles in all vectors. This is the case in some scenarios,
 148 however in many applications this is clearly not the case. For example, if the vectors represent the
 149 support for a candidate from different voters, then the support seems much more meaningful than the
 150 lack of it (as the lack of support may be a result of actual disapproval, but at the same time it can
 151 come from neutrality, or lack of knowledge about a candidate). Then, the fact that two candidates are
 152 supported by the same voters can be much more important for their perceived similarity than the fact
 153 that they agree on more entries.

154 Consider the following four vectors:

$$\begin{aligned} x &= (1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0), & a &= (1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0), \\ y &= (1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1), & b &= (0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1). \end{aligned}$$

155 Under the Hamming distance, the vectors a and b are more similar than the other two. While, if we
 156 treat ones as approvals and zeros as disapprovals, the sets of approved items by a and b are disjoint.

157 In many types of data there is a significant disproportion between the number of 1s and 0s, which
 158 suggests that they play very different roles. In such cases, the following *Add Zero* axiom might be
 159 desired, which basically says that coordinates in which both vectors have 0s might be disregarded.

160 **Definition 4.1** (Add Zero). A dissimilarity measure, f , satisfies *Add Zero* if for every vectors
 161 $x, y \in \{0, 1\}^n$ it holds that

$$f(x \circ (0), y \circ (0)) = f(x, y).$$

162 For example, Add Zero would imply that $f((1, 0, 1), (1, 1, 0)) = f((1, 0, 1, 0), (1, 1, 0, 0))$.

163 It turns out that if in the characterization of the Hamming distance we exchange Zero-One Symmetry
 164 for Add Zero axiom, then we obtain a unique characterization of the Jaccard distance.

165 **Theorem 4.2.** A dissimilarity measure, f , satisfies Anonymity, Independent Symmetry, Add Zero,
 166 Scaling, Convergence, Triangle Inequality, and Normalization if and only if f is the Jaccard distance.

167 *Proof Sketch.* The structure of the proof of the characterization of Jaccard resembles the one from
 168 the proof of the characterization of Hamming. Again, verifying that Jaccard satisfies all axioms is
 169 relatively straightforward, thus we focus on showing that a dissimilarity measure f satisfying the
 170 axioms is equal to the Jaccard distance.

171 From the proof of Theorem 3.8 we know that f satisfying Anonymity and Independent symmetry for
 172 each pair of vectors x, y is a function of three arguments, $|A_x \cap A_y|$, $|A_x \triangle A_y|$, and $n - |A_x \cup A_y|$.
 173 Add Zero allows us to disregard the last argument, thus we end up with a function of two arguments
 174 $|A_x \cap A_y|$ and $|A_x \triangle A_y|$, or, equivalently, $|A_x \triangle A_y|$ and $|A_x \cup A_y|$. Then, Scaling implies that $f(x, y)$
 175 is actually a function of a single argument, i.e., there exists g such that $f(x, y) = g(|A_x \triangle A_y| / |A_x \cup A_y|)$.
 176 Finally, using Triangle Inequality, Convergence, and Normalization, we prove
 177 that g has to be an identity. $\square$

178 Again, our axioms characterizing Jaccard are independent.

179 **Theorem 4.3.** For every axiom in the set Triangle Inequality, Anonymity, Independent Symmetry,
 180 Add Zero, Scaling, and Convergence, there is a dissimilarity measure that satisfies all other axioms in
 181 this set except for this one.

182 Finally, in Appendix B, we offer an additional axiomatic characterization of the discrete distance.

183 5 Rejecting Independent Symmetry

184 In this section, we argue that Independent Symmetry, which is used in characterization of both
 185 Hamming and Jaccard, might be actually not desirable in many situations. To see this, consider the
 186 following four vectors:

$$\begin{aligned} x &= (0, 0, 1, 1, 0, 0, 0, 0), & a &= (1, 1, 1, 1, 0, 0, 0, 0), \\ y &= (1, 1, 1, 1, 1, 1, 0, 0), & b &= (0, 0, 1, 1, 1, 1, 0, 0). \end{aligned}$$

187 Every dissimilarity measure f that satisfies Independent Symmetry must give $f(x, y) = f(a, b)$ by
 188 the definition of the axiom. However, observe that $A_x \subseteq A_y$, which is not true for a and b . If the
 189 vectors correspond to candidates and their coordinates to voters that can support them, this means
 190 that every supporter of candidate x also voted for y . This is a very strong signal that projects x and y
 191 are related. On the other hand, only 50% of supporters of a also voted for b and vice versa. Therefore,
 192 we might want to consider x and y as more similar than a and b . Alternatively, it is possible that in
 193 some settings actually a and b can be considered more similar, as they are of similar size. Either way,
 194 Independent Symmetry prevents distinguishing these two cases.

195 What could be a sensible dissimilarity metric that does not satisfy Independent Symmetry? Let us
 196 observe that the Hamming distance can be written as

$$H(x, y) = 2 \cdot \frac{|A_x \setminus A_y| + |A_y \setminus A_x|}{2} \cdot \frac{1}{n}.$$

197 I.e., it is proportional to the arithmetic mean of $|A_x \setminus A_y|$ and $|A_y \setminus A_x|$. Similarly, we can write

$$J(x, y) = 2 \cdot \frac{|A_x \setminus A_y| + |A_y \setminus A_x|}{2} \cdot \frac{1}{|A_x \cup A_y|}.$$

198 The arithmetic mean is invariant to subtracting a value from one argument and adding it to the other,
 199 which is why both distances satisfy Independent Symmetry. However, other means, like quadratic or
 200 geometric mean, lack this property. For $p \in \mathbb{R}_{\geq 0}$, we define a *generalized mean* of two numbers as

$$M_p(k, \ell) = \begin{cases} \sqrt[p]{(k^p + \ell^p)/2}, & \text{for } p \in (0, \infty), \\ \sqrt{k \cdot \ell}, & \text{for } p = 0. \end{cases}$$

201 Then, let us define p -Hamming and p -Jaccard as follows

$$H_p(x, y) = \frac{2M_p(|A_x \setminus A_y|, |A_y \setminus A_x|)}{n}, \quad J_p(x, y) = \frac{2M_p(|A_x \setminus A_y|, |A_y \setminus A_x|)}{2M_p(|A_x \setminus A_y|, |A_y \setminus A_x|) + |A_x \cap A_y|}.$$

202 Clearly $H_1(x, y) = H(x, y)$ and $J_1(x, y) = J(x, y)$. Observe that whenever $|A_x \setminus A_y| = |A_y \setminus A_x|$,
 203 we have that $H_p(x, y) = H_q(x, y)$, for every $p, q \in [0, \infty]$. However, if $|A_x \setminus A_y| \neq |A_y \setminus A_x|$,
 204 then $H_p(x, y) > H_q(x, y)$ for all $p > q$ (and similar properties hold for p -Jaccard). Therefore,
 205 assuming constant $|A_x \triangle A_y|$, whenever there is an imbalance in the sizes of the sets A_x and A_y ,
 206 $H_p(x, y)$ will give higher distances than the Hamming distance, when $p > 1$, or lower distance when
 207 $p < 1$. In particular, in the extreme, $H_0(x, y)$ and $J_0(x, y)$, which we call *geometric Hamming* and
 208 *geometric Jaccard*, returns dissimilarity 0, whenever $A_x \subseteq A_y$ or $A_y \subseteq A_x$ (but observe that this is
 209 incompatible with Triangle Inequality).

210 *Example 5.1.* Assume that we have three projects a, b , and c , where $|A_a| = |A_b| = 20$, $|A_c| = 6$,
 211 and where $|A_a \cap A_b| = 10$, $|A_a \cap A_c| = 0$, $|A_b \cap A_c| = 5$. And the question is whether b is more
 212 similar to a or c . On the one hand, the size of the intersection between A_a and A_b is twice as large as
 213 the intersection between A_b and A_c . On the other hand, almost all voters who approve c also approve
 214 b . Under the J_1 distance b is more similar to a , while under the J_0 the b is more similar to c .

215 6 Experiments

216 In this section, we illustrate the intuitive difference between the six following dissimilarity measures:
 217 H_0 (geometric Hamming), H_1 (standard Hamming), H_2 (quadratic Hamming), J_0 (geometric
 218 Jaccard), J_1 (standard Jaccard), and J_2 (quadratic Jaccard). Since one of the important potential
 219 applications of dissimilarity measures is the field of computational social choice, we use statistical
 220 models and real-life data from that field. In particular, we focus on approval elections, which consist
 221 of a set of candidates and a set of voters, where each voter approves a subset of the candidates [20].
 222 This structure can naturally be represented as a binary matrix, where rows correspond to voters and
 223 columns to candidates. Occasionally, we refer to such a matrix as a profile.

224 We begin by evaluating the measures on synthetic data and later we move to the real-life one. To
 225 generate synthetic data we use the Euclidean statistical cultures [6] (i.e., popular statistical models
 226 used in social choice). To sample d -Euclidean profile, we proceed as follows: first, sample ideal
 227 points for each voter and each candidate from d -dimensional Euclidean metric space; second, for
 228 each candidate, sample a radius (i.e., strength) from $\mathbb{R}_+$; third, a voter approves a given candidate if
 229 the Euclidean distance between their ideal points is less than or equal to the candidate's radius.

230 6.1 Ordering Accuracy

231 In the following experiment, we analyze how different metrics behave on 1D-Euclidean profiles.
 232 We consider two variants of the experiment, each variant involving 100 sampled profiles, with 100
 233 candidates and 100 voters sampled uniformly at random from $[0, 1]$ interval. In the first variant, the
 234 radius of each candidate is the same $r_1 = r_2 = 0.1$, while in the second variant, we sample radii
 235 uniformly at random from the $[0.015, 0.15]$ interval.

236 Let p_a denote the position of a candidate a on the interval. Given three candidates a, b and c such that
 237 $p_a < p_b < p_c$, we expect a well-behaved distance measure to satisfy the following two inequalities:
 238 $d(a, b) \leq d(a, c)$ and $d(b, c) \leq d(a, c)$. However, this is not necessarily the case. Therefore, for each
 239 of the studied distance measures, we compute the proportion of such inequalities that are satisfied.
 240 We refer to this proportion as the *ordering accuracy*.

241 The results are presented in Table 2. The Jaccard distance variants significantly outperform the
 242 Hamming distance variants in both cases. Notably, when all the candidates have equal radii, the
 243 Jaccard distance always satisfies the ordering inequalities.

244 **Proposition 6.1.** *Under Jaccard distance, if r is constant, then the ordering accuracy is equal to 1.*

Table 2: Average ordering accuracy.

Metric	$r = 0.1$	$r \in [0.015, 0.15]$
H_0	0.690 ± 0.067	0.719 ± 0.063
H_1	0.692 ± 0.066	0.724 ± 0.061
H_2	0.690 ± 0.066	0.719 ± 0.061
J_0	1.000 ± 0.000	0.997 ± 0.001
J_1	1.000 ± 0.000	0.997 ± 0.001
J_2	1.000 ± 0.000	0.997 ± 0.001

Table 3: Quality of clustering.

Metric	1D	2D
H_0	0.0013 ± 0.0014	0.0028 ± 0.0091
H_1	0.1178 ± 0.0875	0.0901 ± 0.0372
H_2	0.1888 ± 0.0903	0.0998 ± 0.0301
J_0	0.0015 ± 0.0019	0.0003 ± 0.0010
J_1	0.0028 ± 0.0033	0.0009 ± 0.0025
J_2	0.0042 ± 0.0070	0.0009 ± 0.0020

6.2 Clustering

Given a dissimilarity measure f , and two clusters $C = \{c_1, \dots, c_k\}$ and $D = \{d_1, \dots, d_\ell\}$, we define the distance between them as $\text{dist}(C, D) = 1/(|C| \cdot |D|) \sum_{c_i \in C} \sum_{d_j \in D} f(c_i, d_j)$.

Using this definition, we perform clustering on candidates in a Euclidean profile via the hierarchical method. It works as follows. Initially, each candidate forms its own cluster. Then, iteratively, at each step, we merge the two most similar clusters, i.e., such that the average distance between their members is minimal. We continue the procedure until exactly five clusters remain.

We begin with 1D-Euclidean space. We consider profiles with 1000 voters, and 100 candidates, sampled uniformly at random from $[0, 1]$ interval with radii sampled uniformly at random from $[0.015, 0.15]$ interval. The results are shown in Figure 1. Each candidate is depicted as a dot and a vertical line, which is proportional to its radius. The geometric variants (i.e., H_0 and J_0) successfully divide the space into five separate clusters. In contrast, for H_1 and H_2 , the green cluster overlaps significantly with other clusters, which is an undesirable outcome in this context. Finally, J_1 and J_2 are outperforming their Hamming counterparts, but still yield less interpretable clusters than the geometric variants.

Next, we consider a 2D-Euclidean space, with a similar setup: 1000 voters, and 100 candidates sampled uniformly at random from $[0, 1]^2$ square, but with candidate radii drawn from $[0.1, 0.33]$ interval. The results are shown in Figure 2. The grey circles denote the voters' positions. The colorful circles denote the candidates' positions, with the size being proportional to the radius. Each color marks a different cluster.

Overall, the results in the 2D case are similar to those in the 1D case. As expected, the Hamming distance performs poorly. It typically produces one dominant cluster that spans much of the space, grouping together a large number of disparate candidates. In contrast, the Jaccard variants (with the exception of the quadratic one) yield significantly better clustering quality. However, the standard Jaccard tends to produce one or two extremely small clusters, typically composed of candidates with very limited voter support.

We can verify the intuitive observations from the plots, by quantifying the number of incorrectly embedded points. For the 1D profiles, let A_3 denote the set of all triplets of candidates such that the first two belong to the same cluster and the third one does not. A triplet is considered invalid if the third candidate lies between the other two. Similarly, for the 2D profiles, let A_4 denote the set of all quadruplets of candidates, such that the first three of them belong to the same cluster and the fourth one does not. A quadruplet is considered invalid if the fourth candidate lies inside the triangle formed by the other three. In Table 3 we present averaged results over 100 profiles. As previously shown in the plots, for 1D profiles, geometric variants outperform the other methods. In the case of 2D profiles, the best performance is achieved by the geometric Jaccard, which results in three times fewer invalid quadruplets than the standard Jaccard.

6.3 Real-Life Data

Finally, we examine real-life data, which comes from Participatory Budgeting (PB) elections. PB elections are democratic processes in which some municipalities distribute a portion of their budget [24]. They usually involve citizens expressing approval preferences over potential projects, a selection of which is later implemented in practice.




Figure 1: Clustering of 1D-Euclidean profiles.




Figure 2: Clustering of 2D-Euclidean profiles.

286 We use participatory budgeting instances from Pabulib [13], focusing on district elections held in
 287 2023 and 2024 in three major Polish cities: Warsaw, Kraków, and Łódź. (Detailed parameters of
 288 these instances are available in Table 5 in Appendix C). In these cities, projects are labeled with
 289 categories. In Kraków and Łódź, each project belongs to exactly one category, while in Warsaw,
 290 projects can have multiple categories. Although these categories are assigned manually and may not
 291 always be accurate, we nonetheless expect that projects within the same category tend to be more
 292 similar to each other on average.

293 For each category, we compute two values: the intra distance d_{intra} , defined as the average pairwise
 294 distance between all projects within the same category, and the inter distance d_{inter} , defined as the
 295 average pairwise distance between projects in the category and those outside of it. For each instance,
 296 we compute the proportion of categories for which it holds that $d_{\text{inter}}/d_{\text{intra}} > 1.01$ (i.e., the inter
 297 distances are nontrivially larger than the intra ones). We then average these proportions across all
 298 districts (i.e., instances) for a given year and city. The results are presented in Table 4. Interestingly,
 299 geometric Jaccard consistently outperforms all other distance measures.

300 Besides categories, several of the Pabulib instances include geographic coordinates of some of the
 301 projects. We focus on district elections held between 2019 and 2024 in Warsaw. We conjectured that
 302 there might be a correlation between the Euclidean distance between two projects (based on location)
 303 and their distance under our measures.

304 Averaging across all 108 district elections held in Warsaw during that period (18 per year), the average
 305 Pearson correlation coefficient between the standard Jaccard distance and the Euclidean one is 0.219,

Table 4: Averaged share of categories with higher inter- than intra-category distance.

Metric	Warsaw		Kraków		Łódź	
	2023	2024	2023	2024	2023	2024
H_0 (geometric Hamming)	63.8%	47.2%	45.5%	40.9%	66.4%	69.6%
H_1 (standard Hamming)	68.6%	56.9%	49.7%	37.4%	65.7%	62.8%
H_2 (quadratic Hamming)	68.5%	62.5%	52.1%	41.9%	63.8%	63.8%
J_0 (geometric Jaccard)	82.6%	88.2%	79.4%	85.9%	82.8%	91.6%
J_1 (standard Jaccard)	67.3%	65.3%	56.9%	58.1%	75.4%	73.3%
J_2 (quadratic Jaccard)	70.2%	79.2%	66.3%	67.8%	78.4%	76.3%

and for geometric Jaccard it is slightly higher at 0.233. Interestingly, in 90% of the instances, the correlation was stronger for geometric Jaccard (detailed results, including comparison with other distance measures, are provided in Table 6 in Appendix C).

Beyond Warsaw, we also analyze a unique election, the green participatory budgeting held in 2023 in Wieliczka (a small Polish city). This case is particularly interesting as it was the first city to use in practice the recently developed method of equal shares [23]. Here, the average Pearson correlation between Jaccard and Euclidean distances is 0.369, while for geometric Jaccard it is 0.441. Wieliczka municipality consists of the central city and several small villages. It appears that voters in the villages tend to vote for their local projects, which may explain the higher correlation between geographic proximity and our distance measures in this case.

7 Related Work

As mentioned introduction, axiomatic approach has been already successfully employed to enhance understanding of various machine learning concepts. For example, a large body of work has concentrated on axioms for clustering. In particular, Kleinberg [18] proved that there is no clustering function that satisfies certain three desired properties. However, this impossibility could be avoided if we focus on functions that produce a specific, exogenously given number of clusters [1] or slightly modify the axioms [10, 19]. Furthermore, several works propose axiomatic characterizations of particular clustering rules [8, 29] or families of rules [2, 4].

Several distances between objects have also been studied axiomatically. A good example is an axiomatization of Kendall Tau distance by Kemeny [17]. The closest to our work, is an axiomatic characterization of Jaccard distance by Gerasimou [14]. However, this characterization considers only three axioms. Two of them are standard distance metric axioms, namely Identity and Triangle Inequality, and one axiom which specifies the marginal contribution of an additional bit of information on which two vectors do not agree. This last axiom is very specific to Jaccard, and there are no other widely known dissimilarity measures that satisfy it. As a result, this characterization does not really offer a view of the Jaccard distance as a consequence of several, not directly related assumptions. In contrast, we propose a larger family of simple axioms, each one of which is satisfied by many natural and popular dissimilarity measures. The Hamming distance was not considered in this paper.

8 Conclusions

We have developed an axiomatic characterization of the Hamming and Jaccard distances, which enables us to clearly capture the key properties of each measure. In particular, the two characterizations differ by a single axiom: Zero-One Symmetry for Hamming vs Add Zero for Jaccard, which highlights the most important difference between them.

Moreover, our axiomatic analysis allowed us to identify a potential shared shortcoming of the Hamming and Jaccard distances, as both satisfy Independent Symmetry, which may be undesirable in certain applications. To address this limitation, we introduced a family of novel measures based on generalized means. Empirical evaluation on both synthetic and real-life data showed that one of the new measures, i.e., geometric Jaccard, consistently outperforms other measures on various tasks, and is better at capturing some nuances. Future work may analyze the new measures in more detail.

References

- [1] M. Ackerman and S. Ben-David. Measures of clustering quality: A working set of axioms for clustering. *Advances in Neural Information Processing Systems*, 21, 2008.
- [2] M. Ackerman and S. Ben-David. Discerning linkage-based algorithms among hierarchical clustering methods. In *IJCAI*, pages 1140–1145, 2011.
- [3] M. Ackerman, S. Ben-David, S. Brânzei, and D. Loker. Weighted clustering: Towards solving the user’s dilemma. *Pattern Recognition*, 120:108152, 2021.
- [4] M. Ackerman, S. Ben-David, and D. Loker. Characterization of linkage-based clustering. In *COLT*, volume 2010, pages 270–281, 2010.
- [5] S. Bag, S. K. Kumar, and M. K. Tiwari. An efficient recommendation generation using relevant jaccard similarity. *Information Sciences*, 483:53–64, 2019.
- [6] N. Boehmer, P. Faliszewski, L. Janeczko, A. Kaczmarczyk, G. Lisowski, G. Pierczynski, S. Rey, D. Stolicki, S. Szufa, and T. Was. Guide to numerical experiments on elections in computational social choice. In *Proceedings of the 33rd International Joint Conference on Artificial Intelligence (IJCAI)*, pages 7962–7970, 2024.
- [7] F. Brandt, V. Conitzer, U. Endriss, J. Lang, and A. Procaccia, editors. *Handbook of Computational Social Choice*. Cambridge University Press, 2016.
- [8] G. E. Carlsson, F. Mémoli, et al. Characterization, stability and convergence of hierarchical clustering methods. *J. Mach. Learn. Res.*, 11(Apr):1425–1470, 2010.
- [9] S.-H. Cha. Comprehensive survey on distance/similarity measures between probability density functions. *International Journal of Mathematical models and Methods in Applied Sciences*, 1(4):300–307, 2007.
- [10] V. Cohen-Addad, V. Kanade, and F. Mallmann-Trenn. Clustering redemption—beyond the impossibility of kleinberg’s axioms. *Advances in Neural Information Processing Systems*, 31, 2018.
- [11] C. Dwork, M. Hardt, T. Pitassi, O. Reingold, and R. Zemel. Fairness through awareness. In *Proceedings of the 3rd innovations in theoretical computer science conference*, pages 214–226, 2012.
- [12] D. Erhan, C. Szegedy, A. Toshev, and D. Anguelov. Scalable object detection using deep neural networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 2147–2154, 2014.
- [13] P. Faliszewski, J. Flis, D. Peters, G. Pierczyński, P. Skowron, D. Stolicki, S. Szufa, and N. Talmon. Participatory budgeting: Data, tools and analysis. In *Proceedings of the 32nd International Joint Conference on Artificial Intelligence (IJCAI)*, pages 2667–2674, 8 2023.
- [14] G. Gerasimou. Characterization of the jaccard dissimilarity metric and a generalization. *Discrete Applied Mathematics*, 355:57–61, 2024.
- [15] Q. Huang, P. Luo, and A. K. H. Tung. A new sparse data clustering method based on frequent items. *Proceedings of the ACM on Management of Data*, 1(1):1–28, 2023.
- [16] L. Kaufman and P. J. Rousseeuw. *Finding groups in data: an introduction to cluster analysis*. John Wiley & Sons, 2009.
- [17] J. G. Kemeny. Mathematics without numbers. *Daedalus*, 88(4):577–591, 1959.
- [18] J. Kleinberg. An impossibility theorem for clustering. *Advances in neural information processing systems*, 15, 2002.
- [19] M. A. Kłopotek and R. A. Kłopotek. In-the-limit clustering axioms. In *International Conference on Artificial Intelligence and Soft Computing*, pages 199–209. Springer, 2020.

- 390 [20] M. Lackner and P. Skowron. *Multi-winner voting with approval preferences*. Springer Nature,
391 2023.
- 392 [21] H. Moulin. Implementing efficient, anonymous and neutral social choice functions. *Journal of*
393 *Mathematical Economics*, 7(3):249–269, 1980.
- 394 [22] S. Niwattanakul, J. Singthongchai, E. Naenudorn, and S. Wanapu. Using of jaccard coefficient
395 for keywords similarity. In *Proceedings of the International MultiConference of Engineers and*
396 *Computer Scientists*, volume 1, pages 380–384, 2013.
- 397 [23] D. Peters, G. Pierczyński, and P. Skowron. Proportional participatory budgeting with additive
398 utilities. *Advances in Neural Information Processing Systems*, 34:12726–12737, 2021.
- 399 [24] S. Rey, F. Schmidt, and J. Maly. The (computational) social choice take on indivisible participa-
400 tory budgeting. *arXiv preprint arXiv:2303.00621*, 2023.
- 401 [25] H. Rezatofighi, N. Tsoi, J. Gwak, A. Sadeghian, I. Reid, and S. Savarese. Generalized inter-
402 section over union: A metric and a loss for bounding box regression. In *Proceedings of the*
403 *IEEE/CVF conference on computer vision and pattern recognition*, pages 658–666, 2019.
- 404 [26] L. S. Shapley. A value for n-person games. *Contributions to the Theory of Games*, 2(28):307,
405 1953.
- 406 [27] A. N. Tarekegn, M. Giacobini, and K. Michalak. A review of methods for imbalanced multi-label
407 classification. *Pattern Recognition*, 118:107965, 2021.
- 408 [28] G. Tsoumakas and I. Katakis. Multi-label classification: An overview. *Data Warehousing and*
409 *Mining: Concepts, Methodologies, Tools, and Applications*, pages 64–74, 2008.
- 410 [29] R. B. Zadeh and S. Ben-David. A uniqueness theorem for clustering. In *Proceedings of the*
411 *Twenty-Fifth Conference on Uncertainty in Artificial Intelligence*, pages 639–646, 2009.
- 412 [30] Y.-F. Zhang, W. Ren, Z. Zhang, Z. Jia, L. Wang, and T. Tan. Focal and efficient iou loss for
413 accurate bounding box regression. *Neurocomputing*, 506:146–157, 2022.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [NA].
- [NA] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for NA).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will be asked to also include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While "[Yes]" is generally preferable to "[No]", it is perfectly acceptable to answer "[No]" provided a proper justification is given (e.g., "error bars are not reported because it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). In general, answering "[No]" or "[NA]" is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The main claims in Abstract & Introduction are about the axiomatic characterization of Hamming and Jaccard distance measures. And, indeed, in the paper, we first provide a set of axioms and then we provide the axiomatic characterization of Hamming and Jaccard distance measures, using these axioms.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: We have shown in the paper the limitations of the discussed distance measures. For instance, the Hamming distance is not the best choice for sparse vectors. On the other hand, geometric Jaccard is failing the triangle inequality.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: All the lemmas, propositions, and theorems stated in the paper are proved either in the main body or in the appendix.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The details needed for reproducing the main results of the paper are provided.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The real-life data used in the experiments is publicly available, while the synthetic data is described with enough detail for reproducibility. The code will be publicly available upon acceptance.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.

- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [NA]

Justification: No training has been conducted in this paper.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Yes, for experiments where it is relevant we provide the standard deviation.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: None of the experiments required extensive computational resources. All the experiments were computed on a standard MacBook Air with M1 chip.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: Yes, we fulfill the anonymity requirement as well as other parts of the NeurIPS Code of Ethics.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification: This work analyzes the properties of fundamental mathematical objects (dis-similarity measures), and as such it does not have direct societal impact.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: There is no high risk of misuse of the proposed data.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The only external data comes from Pabulib, and it is properly credited in the references.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA]

Justification: The paper does not introduce any new assets.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.

- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA]

Justification: LLMs were not involved in the core method development.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.