
Towards Reliable Code-as-Policies: A Neuro-Symbolic Framework for Embodied Task Planning¹

Sanghyun Ahn¹, Wonje Choi¹, Junyong Lee¹, Jinwoo Park², Honguk Woo^{1,2*}²

¹Department of Computer Science and Engineering, Sungkyunkwan University

²Department of Artificial Intelligence, Sungkyunkwan University
{shyuni5, wjchoi1995, lgy7488, pjw971022, hwoo}@skku.edu

Abstract³

Recent advances in large language models (LLMs) have enabled the automatic generation of executable code for task planning and control in embodied agents such as robots, demonstrating the potential of LLM-based embodied intelligence. However, these LLM-based code-as-policies approaches often suffer from limited environmental grounding, particularly in dynamic or partially observable settings, leading to suboptimal task success rates due to incorrect or incomplete code generation. In this work, we propose a neuro-symbolic embodied task planning framework that incorporates explicit symbolic verification and interactive validation processes during code generation. In the validation phase, the framework generates exploratory code that actively interacts with the environment to acquire missing observations while preserving task-relevant states. This integrated process enhances the grounding of generated code, resulting in improved task reliability and success rates in complex environments. We evaluate our framework on RL Bench and in real-world settings across dynamic, partially observable scenarios. Experimental results demonstrate that our framework improves task success rates by 46.2% over Code as Policies baselines and attains over 86.8% executability of task-relevant actions, thereby enhancing the reliability of task planning in dynamic environments.

1 Introduction⁵

Recent advances in embodied control have leveraged large language models (LLMs) to enable flexible, instruction following, effectively bridging natural language understanding with executable actions in physical environments. For instance, SayCan [1] combines LLM-based task interpretation with a reinforcement learning (RL) affordance model to construct a hybrid policy that grounds high-level language instructions, such as “bring me the sponge.” into sequences of low-level, predefined robotic skills. Building on this foundation, subsequent approaches have explored more expressive and compositional modes of action specification through code generation, introducing the paradigm of code-as-policies [2, 3, 4], where LLMs directly generate executable code to control embodied agents. This shift enables task planning that is more modular, interpretable, and adaptable to diverse environments, highlighting the potential of LLMs as general-purpose planners for robotic control.

While LLM-based code-as-policies approaches have demonstrated promising capabilities in fully observable and well-structured settings, their reliability deteriorates in dynamic or partially observable environments, where perceptual input is often sparse, delayed, or ambiguous. These limitations lead to incorrect or incomplete code generation, ultimately resulting in suboptimal task performance. For example, attempting to grasp a fragile object without access to accurate depth or height estimation may lead to dropping or damaging the object, preventing task completion. These challenges underscore

*Corresponding author




Figure 1: Concept of our NESYRO framework illustrated with an example of a room-cleaning task where drawer states are initially unknown. While (1) naive code generation fails without detecting missing observations, (2) NESYRO recursively probes the environment to recover drawer states, enabling the generation of grounded code that successfully completes the task.

the critical need for embodied agents to explicitly reason about uncertainty through exploratory yet safe interactions, and to verify the correctness of generated code prior to execution.

To address these challenges, we propose NESYRO, a neuro-symbolic robot task planning framework that incorporates explicit symbolic verification and interactive validation processes during code generation. Drawing inspiration from the long-standing software engineering principle of verification and validation (V&V) [5, 6, 7], our framework distinguishes between two key processes: verification ensures that the generated code is logically consistent and satisfies symbolic preconditions, while validation assesses whether the code is suitable for the current environment and task objectives. Specifically, symbolic verification statically checks code correctness using domain-specific symbolic tools, whereas interactive validation enables the agent to actively explore its environment to resolve ambiguities and acquire missing observations before task-specific execution.

Our NESYRO framework operates through a recursive composition of two phases: (i) Neuro-symbolic Code Verification, and (ii) Neuro-symbolic Code Validation. Following symbolic verification for code correctness, the interactive validation phase grounds each skill by identifying preconditions and invoking exploratory actions that establish those preconditions as effects, thereby transforming the environment state to enable the intended skill. This process resembles a form of backtracking search, where the agent navigates the environment to construct a valid execution path, progressively verifying and validating the code based on current observations and feedback from symbolic tools.

Figure 1 shows the concept of our framework through an object relocation task, comparing a naive code generation method and our approach. The naive method attempts to move an object without adequately accounting for uncertain factors, which results in the execution of actions that prevent task success. In contrast, our framework identifies the need for exploratory actions, ensures they are conducted safely, and successfully completes the task without causing damage.

We evaluate NESYRO on four task categories, including object relocation, object interaction, auxiliary manipulation, and long-horizon tasks, using both the RL Bench [8] simulation and real-world settings. Experimental results demonstrate that NESYRO improves task success rate by 46.2% over the state-of-the-art baseline, Code as Policies [2], while achieving over 86.8% executability of task-relevant actions in real-world settings. These underscore the enhanced reliability of our framework for robust task planning in dynamic, partially observable environments.

The contributions of this work are summarized as follows:

- We present the NESYRO framework to enable the automatic generation of executable code for task planning in dynamic, partially observable environments.
- We propose a novel recursive mechanism that combines symbolic verification and interactive validation to actively infer and satisfy task-relevant preconditions through exploratory code execution.

- Extensive evaluations on RLBench and real-world tasks demonstrate that NESYRO significantly improves both task success rates and the executability of task-relevant skills.

2 Related work

LLM-based embodied control. In the field of embodied control, there is an emerging trend of utilizing LLMs for reasoning and planning tasks [1, 9, 10, 11, 12, 13, 14, 15, 16, 3]. Building on the high-level reasoning capabilities of LLMs, recent approaches have explored generating executable code as a direct control policy, a paradigm often referred to as code-as-policies [2, 4, 17, 18, 19]. Rather than mapping instructions to predefined skills or discrete action primitives, these methods prompt LLMs to generate Python-like scripts that can be directly executed by embodied agents such as robots. This demonstrates that LLMs are capable of synthesizing low-level control logic, enabling greater flexibility and generalization across a diverse range of tasks. Yet, in dynamic or partially observable settings, the generated code often lacks proper grounding, resulting in incomplete or non-executable outputs. To mitigate this, NESYRO enhances the environmental grounding and reliability of generated code by integrating explicit feedback into the code generation process.

Code verification and validation. Verification and validation are foundational techniques in software engineering for ensuring the correctness and robustness of programs. Verification typically involves static analysis methods such as formal verification, theorem proving, and model checking [20, 21, 22, 23, 24], aiming to prove that a program satisfies its specification before execution. Validation assesses runtime behavior through unit testing, integration testing, system-level evaluation, and runtime verification [7, 25, 6, 26, 27], ensuring that the code performs as intended under real-world conditions. Recent works have explored combining these principles with LLMs to improve code reliability via various forms of static analysis and runtime feedback [28, 29, 5]. Still, existing approaches are limited to static or simulated settings and lack grounding in real-world environments, which is an essential requirement for embodied agents. Our framework addresses this by enabling agents to identify missing task-relevant observations in dynamic and partially observable environments.

Neuro-symbolic system. Recent neuro-symbolic systems combine the generalization capabilities of LLMs with the robustness and interpretability of symbolic reasoning tools. This hybrid approach has been actively investigated in areas such as symbolic problem solving, planning, and program synthesis [30, 31, 32, 33, 34, 35, 36]. Neuro-symbolic approaches in embodied agents commonly employ LLMs for perception and natural language instruction understanding, while utilizing symbolic tools to perform high-level task planning [37, 38, 39, 40, 41]. However, existing neuro-symbolic agents rely on fixed modular structures or pre-defined procedures, limiting their adaptability to missing observations and environmental uncertainty. NESYRO integrates symbolic reasoning with interactive validation and exploratory interactions, enabling reliable task planning in dynamic environments.

3 NESYRO Framework

3.1 Problem Formulation

We tackle the automatic generation of executable code for task planning and control in embodied agents operating in dynamic, partially observable environments. The environment is modeled as a Partially Observable Markov Decision Process (POMDP) $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{G}, T, R, \Omega, \mathcal{O})$ [42, 43, 14], where $s \in \mathcal{S}$ is a state, $a \in \mathcal{A}$ an action, and $g \in \mathcal{G}$ is a high-level goal (e.g., “pick up the red mug”). $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$ is the transition function describing dynamics. The reward function $R : \mathcal{S} \times \mathcal{A} \times \mathcal{G} \rightarrow \{0, 1\}$ returns a binary success signal, which is common in robotics where only task completion is observable. Due to partial observability, observations $o \in \Omega$ are received via $\mathcal{O} : \mathcal{S} \times \mathcal{A} \rightarrow \Omega$, where observations are represented in symbolic form, composed of structured predicate-based expressions (e.g., `is_locked(drawer)`, `on(object, surface)`). Under the code-as-policies paradigm, the LLM generates policy code π from observation history and goal as input and internally encodes the actions necessary to complete the task. When executed via tool `exe`, it yields policy `exe( $\pi$ )`. The set of policy codes is $\Pi = \{\pi \mid \text{exe}(\pi) : \Omega^* \rightarrow \mathcal{A}\}$, where Ω^* is the set of all finite observation histories $o_{\leq t} = (o_0, \dots, o_t)$ with each $o_i \in \Omega$. Our goal is to find the policy code π that maximizes the expected return:




Figure 2: The NESYRO framework with *Neuro-Symbolic Code Verification* and *Neuro-Symbolic Code Validation* phases. It recursively verifies and validates the policy code, while incrementally acquiring observations.

$$\pi^* = \operatorname{argmax}_{\pi \in \Pi} \mathbb{E}_{g \sim \mathcal{G}, \tau \sim \mathcal{P}(\operatorname{exe}(\pi), g)} \left[\sum_{t=0}^{\infty} R(s_t, \operatorname{exe}(\pi)(o_{\leq t}), g) \right]. \quad (1)$$

Here, $\tau = (s_0, o_0, a_0, s_1, o_1, a_1, \dots)$ denotes the trajectory generated by executing $\operatorname{exe}(\pi)$ in the environment, and $\mathcal{P}(\operatorname{exe}(\pi), g)$ is the resulting trajectory distribution induced by $\operatorname{exe}(\pi)$ under g , T , and $\mathcal{O}$. In our implementation, each action a_t in τ corresponds to a skill function composed of multiple low-level control APIs encoded within π . Since $\mathcal{M}$ is partially observable, π^* must balance exploration (to reduce uncertainty) and exploitation (to achieve goals), ensuring reliable task planning in dynamic environments.

3.2 Overall Framework

To achieve the objective described in Eq. (1), we introduce NESYRO, designed to achieve the generation of executable and grounded code through dynamic reconfiguration of reasoning components. As illustrated in Figure 2, NESYRO operates in two key phases: Phase i), *Neuro-Symbolic Code Verification*, which ensures the logical correctness of the policy code with respect to the generated task specification; and Phase ii), *Neuro-Symbolic Code Validation*, which ensures environmental feasibility by evaluating and refining skills based on their grounding.

In the verification phase i), given a language instruction g and current observation $o_{\leq t}$, the LLM generates a task specification $\mathcal{T}_{\text{spec}}$ along with the initial policy code π_{main} . The symbolic tool then verifies whether π_{main} satisfies $\mathcal{T}_{\text{spec}}$. If verification fails, the symbolic tool provides feedback to the LLM, which iteratively refines π_{main} until a verified version is obtained. In the validation phase ii), the sequence of skills defined in π_{main} is validated sequentially using a neuro-symbolic confidence score, NeSyConf, which integrates symbolic feasibility and commonsense plausibility. If a skill's confidence score falls below a threshold ϵ , NESYRO synthesizes a safe probe policy code π_{probe} to recover missing observations. π_{probe} is recursively processed through the composition of the verification and validation phases until all skills are grounded.

This recursive structure induces a policy tree rooted at π_{main} , where each π_{probe} serves as a subroutine that enables successful validation of its parent. The recursive process continues until all required observations have been acquired and every skill in π_{main} is validated. The final output is a grounded version of π_{main} , aligned with both $\mathcal{T}_{\text{spec}}$ and the current environment.

3.3 Neuro-symbolic Code Verification

LLM-based code generation. Given a language instruction g and observation $o_{\leq t}$, a verification LLM, denoted as Φ_{veri} , is prompted to reason in a chain-of-thought (CoT) manner [44, 45], synthesizing key objectives and constraints into a task specification $\mathcal{T}_{\text{spec}}$. It then uses this specification to generate the policy code π_{main}^i , for instance in Python, defining a sequence of skills along with their

parameters and required libraries. 1

$$\Phi_{\text{veri}} : (o_{\leq t}, g, l_{\text{cot}}, \mathcal{D}, \pi_{\text{main}}^{i-1}, \mathcal{F}_{\text{veri}}^{i-1}, n) \mapsto (\mathcal{T}_{\text{spec}}, \pi_{\text{main}}^i) \quad 2 \quad (2)$$

Here, $o_{\leq t}$ is the current observation, initially from o_0 and incrementally updated via probe. l_{cot} is the CoT prompt guiding the Φ_{veri} to generate the specification as an intermediate step. $\mathcal{D}$ denotes domain knowledge, consisting of available skills represented as parameterized function calls, each defined by its applicability conditions and resulting effects, as well as object types and attributes that map these skills to the environment. $\mathcal{F}_{\text{veri}}^{i-1}$ is verification feedback from the previous iteration, used by the LLM to generate the revised π_{main}^i . Importantly, π_{main}^{i-1} and $\mathcal{F}_{\text{veri}}^{i-1}$ are provided only when the previous verification attempt has failed. The index n indicates the skill function call order in π_{main} from which the code refinement begins, while calls prior to n remain unchanged. When $n=0$, it corresponds to the initial code generation. The resulting $\mathcal{T}_{\text{spec}}$ captures the high-level intent, constraints, and relevant subgoals derived from the g and $o_{\leq t}$. π_{main}^i is then passed to the symbolic verification tool.

Symbolic-based code verification. Next, a symbolic verification tool Ψ_{veri} (i.e., SMT solver) checks whether π_{main}^i satisfies $\mathcal{T}_{\text{spec}}$, identifying any violations of constraints defined in the specification. 4

$$\Psi_{\text{veri}} : (\mathcal{T}_{\text{spec}}, \pi_{\text{main}}^i) \mapsto \begin{cases} \text{verified } \pi_{\text{main}}, & \text{if verification succeeds} \\ \mathcal{F}_{\text{veri}}^i, & \text{if verification fails} \end{cases} \quad 5 \quad (3)$$

If verification fails, Ψ_{veri} provides detailed $\mathcal{F}_{\text{veri}}^i$ that identifies the specific parts of π_{main}^i violating $\mathcal{T}_{\text{spec}}$, such as incorrect parameter bindings or structural mismatches. This feedback is then passed to the next Φ_{veri} iteration to generate a revised π_{main}^i . Once π_{main}^i passes verification, resulting in a verified version of π_{main} , we proceed to the *Neuro-Symbolic Code Validation* phase. 6

3.4 Neuro-symbolic Code Validation 7

The verified policy code π_{main} is parsed into a sequence of skill function calls, $\pi_{\text{main}} = (f_0, f_1, \dots, f_N)$, where N denotes the maximum skill step. Unlike the *Neuro-symbolic code verification* phase, which reasons over the entire π_{main} holistically, the validation process evaluates and refines each skill sequentially to assess its feasibility in the current environment. The index n , as in Eq. (2), denotes the skill step under validation and represents the first unvalidated step in π_{main} .

Neuro-symbolic confidence score. To assess skill feasibility, we introduce Neuro-Symbolic Confidence score (NeSyConf), which combines Common Sense Confidence (CSC) from a validation LLM denoted Φ_{vali} and Logic Confidence (LC) from a symbolic validation tool Ψ_{vali} in parallel. 9

$$\Phi_{\text{vali}} : (\mathcal{D}, \mathcal{E}_{\text{demo}}, o_{\leq t}, g, f_n) \mapsto \text{CSC}_{f_n} \quad 10 \quad (4)$$

The CSC_{f_n} estimates the likelihood that a given skill $f_n \in \pi_{\text{main}}$ will succeed under the current observation $o_{\leq t}$ and instruction g , based on both domain knowledge $\mathcal{D}$ and retrieved demonstrations $\mathcal{E}_{\text{demo}}$. To compute this, we insert the code snippet corresponding to f_n into the LLM prompt along with $\mathcal{D}$, $o_{\leq t}$, g , and $\mathcal{E}_{\text{demo}}$. Φ_{vali} then assigns token-wise probabilities to the f_n , and we compute a perplexity-based score to estimate the skill’s plausibility. The cumulative log probabilities are normalized to produce a consistent confidence CSC_{f_n} . To reduce hallucinations and improve estimation accuracy, we retrieve skill-level demonstrations $\mathcal{E}_{\text{demo}}$ whose contexts closely resemble the current situation and include them in the prompt as guidance. 11

$$\Psi_{\text{vali}} : (\mathcal{D}, o_{\leq t}, g, f_n) \mapsto \text{LC}_{f_n} \quad 12 \quad (5)$$

The logic-based confidence LC_{f_n} is computed by the symbolic validation tool (i.e., PDDL planner), which assesses whether the f_n is symbolically feasible under the $o_{\leq t}$, g , and $\mathcal{D}$. We use a Ψ_{vali} to check whether f_n ’s preconditions hold in $o_{\leq t}$. If the planner successfully generates a plan including f_n , we set $\text{LC}_{f_n} = 1$; otherwise, we set $\text{LC}_{f_n} = 0$, indicating symbolic infeasibility under the $o_{\leq t}$. 13

$$\text{NeSyConf}_{f_n} = \text{CSC}_{f_n} \times \text{LC}_{f_n} \text{ with: } \begin{cases} \text{proceed to } \text{NeSyConf}_{f_{n+1}}, & \text{if } \text{NeSyConf}_{f_n} \geq \epsilon \\ \text{generate } \pi_{\text{probe}} \text{ using } \mathcal{F}_{\text{csc}}, \mathcal{F}_{\text{lc}} & \text{otherwise} \end{cases} \quad 14 \quad (6)$$

The NeSyConf_{f_n} represents the final confidence score for f_n , computed by multiplying CSC_{f_n} and LC_{f_n} . This score estimates whether f_n is correctly grounded in the environment and likely to succeed upon execution. If $\text{NeSyConf}_{f_n} < \epsilon$, our framework initiates a safe probe policy code π_{probe} using feedback from each component, namely $\mathcal{F}_{\text{csc}}$ and $\mathcal{F}_{\text{lc}}$. 15

Safe probe. If a f_n receives a low confidence score, our framework responds by generating a safe probe policy code π_{probe} to recover missing observations. Constructed using $\mathcal{F}_{\text{csc}}$ from CSC and $\mathcal{F}_{\text{lc}}$ from LC, π_{probe} undergoes the same verification and validation process as π_{main} . Because π_{probe} is validated before execution, the framework ensures that only safe and grounded code is deployed. This recursive structure generates a policy tree rooted at π_{main} , where each π_{probe} functions as a subroutine that enables successful validation of its parent skill. Once π_{probe} is executed, it collects new observations and updates the current observation to $o_{\leq t+1}$. This updated observation is then used in the subsequent *Policy code refinement* process to update π_{main} .

Policy code refinement. Following safe probe, the $o_{\leq t+1}$ is used to refine π_{main} at the skill level. Specifically, instead of regenerating the entire policy, our framework targets the current f_n and prompts the LLM to regenerate only its code segment using the $o_{\leq t+1}$. This regeneration is conducted through our *Neuro-symbolic Code Verification* process. The updated code for f_n is then evaluated using Eq. (6), where its confidence score NeSyConf_{f_n} is reassessed. This process of refinement and safe probe is repeated until $\text{NeSyConf}_{f_n} \geq \epsilon$. Once all skills $f_n \in \pi_{\text{main}}$ have been successfully validated, the grounded π_{main} is executed.

Further implementation details, and algorithmic pseudocode are provided in Appendix A.

4 Experiment

4.1 Experiment Setting

Environments. We conducted experiments in both RLBench [8] and real-world settings using a 7-DoF Franka Emika Research 3 robotic arm, enabling reproducible evaluations via randomized initial states and instructions to analyze safe probe strategies in dynamic, partially observable scenarios. In contrast, real-world experiments evaluated robustness and generalizability under real-world noise and variability. In dynamic, partially observable scenarios, we defined four observability levels based on initial observation availability: *High Incompleteness* condition removes more than half of the essential observations, constraining the task-solving process. *Low Incompleteness* condition retains most observations, though the observation remains partially incomplete. *Stochastic Incompleteness* condition provides a randomly selected subset of observations, with the incompleteness level varying across episodes. Finally, *Complete* condition offers full relevant observations, rendering probe unnecessary. We denote these four levels as *High*, *Low*, *Stochastic*, and *Complete*, respectively, each evaluated over ten randomized trials with varied initial conditions and instructions.

Table 1: Task types and their corresponding probe types

Task Type	Probe Type
<i>Object Relocation</i> (e.g., moving tomatoes on a plate)	Robot Pose Adjust (e.g., verifying which item is a tomato)
<i>Object Interaction</i> (e.g., opening a drawer)	Object State Check (e.g., checking whether a drawer is locked)
<i>Auxiliary Manipulation</i> (e.g., opening a drawer in a dark room)	Object State Change (e.g., turning on the light to locate the drawer)
<i>Long-Horizon</i> (e.g., placing a tomato inside a drawer)	Uses two or more of the above probe types depending on task structure and uncertainty

Task and probe types. In dynamic, partially observable environments, tasks often require acquiring missing observations before execution. To support structured analysis of the probe, we define task types based on manipulation goals and missing observation roles. Specifically, we distinguish whether the uncertainty concerns object identity, object state, or auxiliary conditions. These distinctions define three task types, each reflecting a distinct observation-seeking pattern. In addition, *long-horizon* tasks involve multi-step goals. Correspondingly, we define three probe types to characterize observation acquisition: (1) Robot Pose Adjust, adjusting viewpoint to resolve ambiguity; (2) Object State Check, identifying hidden object states relevant to the task; and (3) Object State Change, performing auxiliary skills to enable observation of otherwise inaccessible states. These categories are functionally defined

Table 2: Task performance under varying levels of observability incompleteness in RL Bench 1

Methods	High		Low		Stochastic		Complete	
	SR	GC	SR	GC	SR	GC	SR	GC
Task Type: Object Relocation								
CaP	25.0 $\pm$ 7.1	41.5 $\pm$ 8.8	30.0 $\pm$ 0.0	43.8 $\pm$ 1.8	10.0 $\pm$ 0.0	36.3 $\pm$ 1.8	90.0 $\pm$ 0.0	92.5 $\pm$ 3.5
CaP w/ Lemur	25.0 $\pm$ 7.1	43.8 $\pm$ 5.3	30.0 $\pm$ 0.0	43.8 $\pm$ 1.8	10.0 $\pm$ 0.0	36.3 $\pm$ 1.8	90.0 $\pm$ 0.0	96.3 $\pm$ 1.8
CaP w/ CodeSift	55.0 $\pm$ 7.1	72.5 $\pm$ 3.5	50.0 $\pm$ 0.0	57.5 $\pm$ 3.5	40.0 $\pm$ 0.0	52.5 $\pm$ 3.5	95.0 $\pm$ 7.1	95.0 $\pm$ 7.1
LLM-Planner	30.0 $\pm$ 0.0	35.0 $\pm$ 7.1	50.0 $\pm$ 0.0	58.8 $\pm$ 5.3	30.0 $\pm$ 0.0	43.8 $\pm$ 5.3	80.0 $\pm$ 0.0	88.8 $\pm$ 5.3
AutoGen	30.0 $\pm$ 0.0	35.0 $\pm$ 7.1	55.0 $\pm$ 7.1	60.0 $\pm$ 7.1	40.0 $\pm$ 14.1	47.5 $\pm$ 10.6	85.0 $\pm$ 7.1	87.5 $\pm$ 10.6
NEsYRO	70.0 $\pm$ 14.1	72.5 $\pm$ 10.6	75.0 $\pm$ 7.1	87.5 $\pm$ 3.5	65.0 $\pm$ 7.1	75.0 $\pm$ 0.0	95.0 $\pm$ 7.1	97.5 $\pm$ 3.5
Task Type: Object Interaction								
CaP	20.0 $\pm$ 14.1	35.0 $\pm$ 7.1	25.0 $\pm$ 7.1	40.0 $\pm$ 3.5	35.0 $\pm$ 7.1	51.3 $\pm$ 5.3	75.0 $\pm$ 7.1	77.5 $\pm$ 7.1
CaP w/ Lemur	35.0 $\pm$ 7.1	47.5 $\pm$ 7.1	35.0 $\pm$ 7.1	47.5 $\pm$ 3.5	30.0 $\pm$ 14.1	46.3 $\pm$ 12.1	85.0 $\pm$ 7.1	86.3 $\pm$ 8.8
CaP w/ CodeSift	40.0 $\pm$ 0.0	65.0 $\pm$ 7.1	50.0 $\pm$ 14.1	55.0 $\pm$ 7.1	40.0 $\pm$ 0.0	60.0 $\pm$ 14.1	90.0 $\pm$ 14.1	90.0 $\pm$ 14.1
LLM-Planner	5.0 $\pm$ 7.1	15.0 $\pm$ 0.0	40.0 $\pm$ 14.1	53.8 $\pm$ 5.3	35.0 $\pm$ 7.1	42.5 $\pm$ 14.1	55.0 $\pm$ 7.1	63.8 $\pm$ 8.8
AutoGen	40.0 $\pm$ 0.0	48.8 $\pm$ 1.8	50.0 $\pm$ 0.0	58.8 $\pm$ 5.3	50.0 $\pm$ 0.0	57.5 $\pm$ 0.0	75.0 $\pm$ 7.1	76.3 $\pm$ 8.8
NEsYRO	70.0 $\pm$ 0.0	76.3 $\pm$ 1.8	80.0 $\pm$ 0.0	83.8 $\pm$ 1.8	70.0 $\pm$ 14.1	73.8 $\pm$ 8.8	90.0 $\pm$ 0.0	92.5 $\pm$ 0.0
Task Type: Auxiliary Manipulation								
CaP	25.0 $\pm$ 7.1	25.0 $\pm$ 7.1	50.0 $\pm$ 0.0	51.3 $\pm$ 1.8	40.0 $\pm$ 0.0	45.8 $\pm$ 8.3	85.0 $\pm$ 7.1	90.0 $\pm$ 4.7
CaP w/ Lemur	30.0 $\pm$ 14.1	30.0 $\pm$ 14.1	50.0 $\pm$ 14.1	58.3 $\pm$ 7.1	30.0 $\pm$ 14.1	34.2 $\pm$ 15.3	85.0 $\pm$ 7.1	90.8 $\pm$ 3.5
CaP w/ CodeSift	5.0 $\pm$ 7.1	5.0 $\pm$ 7.1	55.0 $\pm$ 7.1	57.5 $\pm$ 3.5	35.0 $\pm$ 7.1	35.0 $\pm$ 7.1	90.0 $\pm$ 0.0	93.3 $\pm$ 0.0
LLM-Planner	15.0 $\pm$ 7.1	15.0 $\pm$ 7.1	30.0 $\pm$ 0.0	37.5 $\pm$ 3.5	10.0 $\pm$ 14.1	15.0 $\pm$ 7.1	75.0 $\pm$ 7.1	80.0 $\pm$ 0.0
AutoGen	15.0 $\pm$ 7.1	15.0 $\pm$ 7.1	35.0 $\pm$ 7.1	40.0 $\pm$ 0.0	20.0 $\pm$ 0.0	22.5 $\pm$ 3.5	80.0 $\pm$ 0.0	80.0 $\pm$ 0.0
NEsYRO	60.0 $\pm$ 0.0	80.8 $\pm$ 1.2	70.0 $\pm$ 14.1	74.2 $\pm$ 13.0	70.0 $\pm$ 14.1	85.8 $\pm$ 5.9	95.0 $\pm$ 7.1	96.7 $\pm$ 4.7
Task Type: Long-Horizon								
CaP	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	20.0 $\pm$ 0.0	40.4 $\pm$ 6.6	0.0 $\pm$ 0.0	0.7 $\pm$ 1.0	40.0 $\pm$ 14.1	53.8 $\pm$ 3.4
CaP w/ Lemur	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	30.0 $\pm$ 0.0	47.1 $\pm$ 0.0	0.0 $\pm$ 0.0	1.6 $\pm$ 0.2	55.0 $\pm$ 7.1	67.1 $\pm$ 5.1
CaP w/ CodeSift	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	30.0 $\pm$ 14.1	45.8 $\pm$ 7.9	5.0 $\pm$ 7.1	5.0 $\pm$ 7.1	65.0 $\pm$ 7.1	71.4 $\pm$ 6.7
LLM-Planner	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	10.0 $\pm$ 0.0	11.4 $\pm$ 0.0	5.0 $\pm$ 0.0	12.9 $\pm$ 8.1	35.0 $\pm$ 7.1	44.4 $\pm$ 9.9
AutoGen	0.0 $\pm$ 0.0	5.5 $\pm$ 3.0	30.0 $\pm$ 14.1	39.2 $\pm$ 10.3	20.0 $\pm$ 0.0	28.5 $\pm$ 7.2	50.0 $\pm$ 0.0	55.1 $\pm$ 0.8
NEsYRO	45.0 $\pm$ 7.1	65.2 $\pm$ 6.7	45.0 $\pm$ 7.1	58.1 $\pm$ 6.1	35.0 $\pm$ 7.1	41.9 $\pm$ 8.1	65.0 $\pm$ 7.1	73.7 $\pm$ 8.3

and implemented for distinct behavioral purposes, thereby enabling generalization to varied settings. We describe the task types and their associated probe types in Table 1.

Evaluation metrics. To assess the objectives in Section 3.1, we adopt metrics from prior work [46, 11, 47, 48]. Success Rate (SR) assigns 100% for full task completion and 0% otherwise. Goal Condition (GC) measures the percentage of sub-goals achieved. In real-world experiments, we report Irreversible Actions (IA), counting irreversible actions during task execution.

Baselines. We compare our approach against several state-of-the-art baselines that use LLMs to generate robot control code, covering different paradigms of V&V, reasoning, and replanning. **Code as Policies (CaP)** [2] generates reusable control code from task instructions via LLM. **CaP w/ Lemur** [49] integrates SMT verification into the code generation pipeline. **CaP w/ CodeSift** [5] improves the reliability of LLM-generated code through multi-stage syntactic and semantic validation, without relying on reference code or actual execution. **LLM-Planner** [11] introduces an execution-aware replanning framework, replans after failure using new observations from the environment. **AutoGen** [50] extends LLM-Planner by enabling multi-agent collaborative reasoning.

NEsYRO implementation. We employ GPT-4o-mini [51] for code generation and feedback generation. Additionally, Llama-3.2-3B [52] is used to compute the CSC. The decoding temperature is fixed at 0.0 for all generation steps. For the verification phase (i), the Z3 SMT solver [20] is employed as the symbolic verification tool, while for the validation phase (ii), the Fast Downward planner [53] is used as the symbolic validation tool.

Detailed descriptions of the experimental settings are provided in the Appendix B and C.

4.2 Main Result

Task performance on RL Bench. Table 2 reports the performance of robot control code across all task types in dynamic, partially observable scenarios. For each task, we consider four levels

Table 3: Task performance across task types in the real-world under partial observability (*High* and *Low Incompleteness* averaged). NESYRO-Complete reports results under *Complete*.

Real-World Task Type	CaP			CaP w/ CodeSift			NeSyRo			NeSyRo-Complete		
	SR (↑)	GC (↑)	IA (↓)	SR (↑)	GC (↑)	IA (↓)	SR (↑)	GC (↑)	IA (↓)	SR (↑)	GC (↑)	IA (↓)
Object Relocation	7.5 $\pm$ 3.5	11.3 $\pm$ 1.8	19	12.5 $\pm$ 3.5	19.4 $\pm$ 4.4	4	82.5 $\pm$ 3.5	83.8 $\pm$ 3.5	2	85.0 $\pm$ 7.1	90.0 $\pm$ 3.5	2
Object Interaction	30.0 $\pm$ 7.1	37.5 $\pm$ 7.1	12	20.0 $\pm$ 7.1	24.4 $\pm$ 9.7	4	75.0 $\pm$ 14.1	77.5 $\pm$ 17.7	0	90.0 $\pm$ 14.1	90.0 $\pm$ 14.1	0
Auxiliary Manipulation	0.0 $\pm$ 0.0	0.0 $\pm$ 0.0	4	2.5 $\pm$ 3.5	8.3 $\pm$ 4.7	5	20.0 $\pm$ 0.0	20.0 $\pm$ 0.0	2	20.0 $\pm$ 14.1	20.0 $\pm$ 14.1	2
Long-Horizon	5.0 $\pm$ 0.0	14.2 $\pm$ 3.5	18	7.5 $\pm$ 10.6	13.1 $\pm$ 7.4	16	52.5 $\pm$ 3.5	54.2 $\pm$ 3.5	3	60.0 $\pm$ 0.0	65.8 $\pm$ 8.3	2
Total	10.6 $\pm$ 0.9	15.7 $\pm$ 0.4	53	10.6 $\pm$ 4.4	16.3 $\pm$ 4.4	29	57.5 $\pm$ 3.5	58.9 $\pm$ 4.4	7	68.8 $\pm$ 5.3	71.5 $\pm$ 6.5	6




Figure 3: Representative failure scenarios under partial observability across real-world tasks

of observability incompleteness, requiring each method to recover missing observations. NESYRO consistently outperforms baselines (AutoGen and CaP w/ CodeSift) by 26.3% in SR and 24.3% in GC across all levels of observability incompleteness.

CaP w/ Lemur outperforms the base CaP model, indicating that pre-execution verification improves robustness even without explicit safe probe. LLM-based replanning methods, such as AutoGen and LLM-Planner, perform well compared to V&V-based approaches, such as CaP w/ CodeSift, under the *Low Incompleteness* condition, but their performance degrades significantly as uncertainty increases and critical observations are missing. In contrast, NESYRO sustains strong results across every observability regime and task type. It accurately detects missing observations and performs safe probes, avoiding irreversible actions. Under the *Complete* condition, all methods achieve high performance since no additional probe is required. As observability grows more incomplete, baseline performance drops sharply. Meanwhile, NESYRO detects the missing observations, explores safely, and maintains performance close to that of the Complete setting. This tendency becomes even more pronounced in long-horizon tasks, where extensive probe is required and bridging observation gaps becomes especially difficult.

Task performance on Real-World. Table 3 reports the real-world evaluation results under partial observability. The NESYRO-Complete configuration serves as an upper bound, representing the ideal performance achievable when all relevant observations are fully available. Across all task categories, NESYRO consistently achieves the highest SR and GC, outperforming existing baselines such as CaP and CaP w/ CodeSift. On average, NESYRO improves SR by an average of 47.0% and GC by 42.6% compared to CaP w/ CodeSift, while simultaneously reducing IA from 29 to 7.

In particular, for *Object Relocation* and *Object Interaction*, baseline methods achieve less than 30% SR, primarily due to their failure to account for missing observations. In contrast, NESYRO exceeds 75% SR by recovering these observations. Performance on *Auxiliary Manipulation* remains low across all models, including NESYRO-Complete, primarily due to failure in pressing the light switch with sufficient accuracy. Nevertheless, NESYRO achieves the same SR and GC as its Complete variant, indicating upper-bound performance despite the physical difficulty. This highlights the challenge as one of execution, rather than perceptual or reasoning limitations. On *Long-Horizon* tasks, baseline methods remain below 10% SR due to observation gaps and planning complexity. Nonetheless, NESYRO reaches 52.5% SR, approaching the 60.0% score of the *Complete* setting. This shows that our framework bridges long-range dependencies through safe probing. These results confirm that NESYRO enables robust real-world execution through safe recovery of missing observations, even under severe observability incompleteness.

Figure 3 illustrates failure modes associated with high IA scores. Among many failure cases, we present three representative scenarios where CaP fails due to irreversible actions, while NESYRO completes the tasks by recovering missing observations through safe probe. This selection illustrates that NESYRO not only improves success rates but also ensures safe execution grounded in safe probe.

4.3 Ablation 1

Table 4: Code generation performance across LLMs on long-horizon tasks under partial observability in RLbench

RLBench	CaP		NeSyRo		NeSyRo-Complete	
	SR	GC	SR	GC	SR	GC
GPT-4o-mini	10.0 $\pm$ 0.0	20.2 $\pm$ 3.3	45.0 $\pm$ 0.0	61.7 $\pm$ 0.3	65.0 $\pm$ 7.1	73.7 $\pm$ 8.3
o4-mini	12.5 $\pm$ 3.5	22.0 $\pm$ 1.4	42.5 $\pm$ 3.5	54.9 $\pm$ 2.2	50.0 $\pm$ 14.1	51.7 $\pm$ 11.8
GPT-4.1	32.5 $\pm$ 3.5	59.6 $\pm$ 4.8	50.0 $\pm$ 7.1	70.9 $\pm$ 1.8	75.0 $\pm$ 7.1	78.1 $\pm$ 2.7
o3	45.0 $\pm$ 0.0	64.0 $\pm$ 1.9	75.0 $\pm$ 7.1	87.6 $\pm$ 3.0	85.0 $\pm$ 7.1	93.8 $\pm$ 0.7




Figure 4: Real-world compile error rate over all task types

Effect of neuro-symbolic code verification Table 4 reports performance across different LLMs used for code generation, evaluated on long-horizon tasks and averaged over *High* and *Low Incompleteness* conditions. Stronger LLMs improve performance. On average, NeSyRo improves SR by 28.1% and GC by 27.3% over CaP across all LLMs. The SR difference between NeSyRo and the upper-bound NeSyRo-Complete remains at 15.6% on average, indicating that our framework consistently approaches the optimal performance achievable under complete observations. This shows that NeSyRo performs consistently across LLMs. Figure 4 reports real-world task failures due to compile errors across task types. NeSyRo consistently exhibits the lowest compile error rate, highlighting the robustness of its verification and validation pipeline. In contrast, CaP w/ CodeSift incurs more compile errors than the base CaP under partial observability, primarily due to hallucinated evaluations by the LLM in the absence of grounded feedback. NeSyRo addresses this issue through environment-aware validation, enabling reliable execution even with incomplete observations.

Table 5: Comparison of performance on long-horizon RLbench tasks when without LC or CSC from NeSyRo

RLBench	NeSyRo w/o LC		NeSyRo w/o CSC		NeSyRo	
	SR	GC	SR	GC	SR	GC
Object Relocation	50.0 $\pm$ 8.2	60.0 $\pm$ 4.1	35.0 $\pm$ 5.8	51.3 $\pm$ 4.8	67.5 $\pm$ 10.6	73.8 $\pm$ 5.3
Object Interaction	45.0 $\pm$ 7.1	56.3 $\pm$ 5.3	55.0 $\pm$ 7.1	62.5 $\pm$ 3.5	70.0 $\pm$ 7.1	75.0 $\pm$ 3.5
Auxiliary Manipulation	57.5 $\pm$ 3.5	65.4 $\pm$ 4.1	35.0 $\pm$ 7.1	39.2 $\pm$ 3.5	65.0 $\pm$ 7.1	77.5 $\pm$ 7.1
Long-Horizon	25.0 $\pm$ 0.0	36.0 $\pm$ 6.3	32.5 $\pm$ 3.5	52.9 $\pm$ 8.4	45.0 $\pm$ 0.0	61.7 $\pm$ 0.3
Total	44.3 $\pm$ 2.0	54.2 $\pm$ 0.1	37.1 $\pm$ 4.0	49.9 $\pm$ 4.4	61.9 $\pm$ 6.2	72.0 $\pm$ 3.9

Table 6: Effect of LLM parameter scale on CSC in RLbench

RLBench	Long-Horizon Tasks	
	SR	GC
Llama-3.2-1B	32.5 $\pm$ 3.5	55.4 $\pm$ 2.7
Llama-3.2-3B	45.0 $\pm$ 0.0	61.7 $\pm$ 0.3
Llama-3.1-8B	45.0 $\pm$ 7.1	57.7 $\pm$ 0.3
Qwen3-30B-A3B	45.0 $\pm$ 0.0	64.9 $\pm$ 5.1

Effect of neuro-symbolic code validation Tables 5 and 6 report results on long-horizon tasks, averaged over *High* and *Low Incompleteness* conditions. Table 5 compares performance when removing either LC (*w/o LC*) or CSC (*w/o CSC*) from our neuro-symbolic code validation phase. The results indicate that both LC and CSC contribute comparably to overall performance across task types, with an average SR drop of 21.2% and GC drop of 20.0% when either component is removed. This validates the importance of the parallel neuro-symbolic structure in reliably guiding execution under uncertainty. To analyze CSC robustness, Table 6 examines the effect of the LLM parameter scale used in computing the CSC. While smaller models (e.g., Llama-3.2-1B) show degraded performance, models with 3B parameters or more achieve the same SR and differ in GC by less than 7.2%. These results suggest that CSC computation is robust to LLM scaling beyond a moderate threshold, allowing flexible deployment depending on available compute resources.

4.4 Analysis 12

Real-world analysis. Figure 5 showcases how NeSyRo addresses real-world uncertainty in a partially observable setting involving a dark room and unknown drawer states. Initially, the LLM generates an ungrounded policy code without access to key observations, such as whether the drawers are locked, what objects are inside, or where to safely place the dice. As each skill is validated, NeSyRo computes NeSyConf. If the confidence falls below a threshold, the system initiates a targeted safe probe, such as checking individual drawers, to acquire missing observations. These safe probes trigger policy code refinement by updating parameters (e.g., selecting a different drawer) or regenerating the policy for the skill. This iterative validation process adapts the skill sequence and produces a fully grounded policy code that completes the task without causing irreversible failures. A full execution sequence and extended analysis of Figure 5 are provided in Appendix D.




Figure 5: Real-world example of safe probe and policy code refinement in partially observable setting

5 Conclusion

In this work, we presented NESYRO, a neuro-symbolic framework that integrates *Neuro-symbolic Code Verification* and *Neuro-Symbolic Code Validation* to generate reliable robot control code under dynamic, partially observable settings. The framework operates through a recursive process alternating between symbolic verification and interactive validation, ensuring that each skill is both logically consistent and environmentally grounded. By incorporating neuro-symbolic confidence estimation that combines commonsense and logic-based reasoning, NESYRO enables exploratory yet safe interactions and adaptive code refinement under uncertainty. Extensive evaluation in simulation and real-world environments demonstrates the strong performance of NESYRO across diverse tasks.

Limitation and future work. While the NeSyConf formulation is designed to allow LC and CSC to complement each other, similar to how SayCan [1] integrates LLM scores with affordance functions, the current implementation of NESYRO employs a binary LC and predefined domain knowledge, which limits its generality in real-world applications. Future work will address this limitation by incorporating probabilistic and temporal reasoning, such as probabilistic PDDL [54]. We also plan to relax these assumptions and explore the framework’s applicability to more diverse and dynamic domains by extending validation to skills that are not explicitly defined in the domain knowledge.

Ethical concern. LLM-based robot control may lead to unsafe behavior when interacting with hazardous tools (e.g., knives, scissors). To mitigate such risks, we incorporate explicit safety checks and enforce transparent safeguard mechanisms that verify tool affordances and action preconditions before execution, ensuring safe and interpretable operation.

Acknowledgement¹

This work was supported by Institute of Information & communications Technology Planning & Evaluation (IITP) grant funded by the Korea government (MSIT), (RS-2022-II220043 (2022-0-00043), Adaptive Personality for Intelligent Agents, RS-2022-II221045 (2022-0-01045), Self-directed multi-modal Intelligence for solving unknown, open domain problems, RS-2025-02218768, Accelerated Insight Reasoning via Continual Learning, RS-2025-25442569, AI Star Fellowship Support Program (Sungkyunkwan Univ.), and RS-2019-II190421, Artificial Intelligence Graduate School Program (Sungkyunkwan University)), the National Research Foundation of Korea (NRF) grant funded by the Korea government (MSIT) (No. RS-2023-00213118), IITP-ITRC (Information Technology Research Center) grant funded by the Korea government (MIST) (IITP-2025-RS-2024-00437633, 10%), IITP-ICT Creative Consilience Program grant funded by the Korea government (MSIT) (IITP-2025-RS-2020-II201821, 10%), and by Samsung Electronics.

References³

- [1] Anthony Brohan et al. “Do as I can, not as I say: Grounding language in robotic affordances”.⁴ In: *Proceedings of the 6th Conference on Robot Learning (CoRL)*. PMLR, 2023, pp. 287–318.
- [2] Jacky Liang et al. “Code as policies: language model programs for embodied control”. In: *Proceedings of the 40th International Conference on Robotics and Automation (ICRA)*. IEEE, 2023, pp. 9493–9500.
- [3] Wenlong Huang et al. “VoxPoser: Composable 3D Value Maps for Robotic Manipulation with Language Models”. In: *Proceedings of the 7th Conference on Robot Learning (CoRL)*. Vol. 229. 2023, pp. 540–562.
- [4] Yao Mu et al. “Robocodex: multimodal code generation for robotic behavior synthesis”. In: *Proceedings of the 41st International Conference on Machine Learning (ICML)*. PMLR, 2024, pp. 36434–36454.
- [5] Pooja Aggarwal et al. “CodeSift: An LLM-Based Reference-Less Framework for Automatic Code Validation”. In: *arXiv preprint arXiv:2408.15630* (2024).
- [6] Stefan Mitsch and André Platzer. “ModelPlex: verified runtime validation of verified cyber-physical system models”. In: *Formal Methods in System Design* 49 (2016), pp. 33–74.
- [7] Martin Leucker and Christian Schallhart. “A brief account of runtime verification”. In: *The Journal of Logic and Algebraic Programming* 78.5 (2009), pp. 293–303.
- [8] Stephen James et al. “RLBench: The Robot Learning Benchmark & Learning Environment”. In: *IEEE Robotics and Automation Letters* 5.2 (2020), pp. 3019–3026.
- [9] Wenlong Huang et al. “Inner Monologue: Embodied Reasoning through Planning with Language Models”. In: *Proceedings of the 6th Conference on Robot Learning (CoRL)*. Vol. 205. 2022, pp. 1769–1782.
- [10] Wenlong Huang et al. “Language models as zero-shot planners: Extracting actionable knowledge for embodied agents”. In: *Proceedings of the 39th International Conference on Machine Learning*. PMLR. 2022, pp. 9118–9147.
- [11] Chan Hee Song et al. “LLM-planner: Few-shot grounded planning for embodied agents with large language models”. In: *Proceedings of the 19th International Conference on Computer Vision (ICCV)*. 2023, pp. 2998–3009.
- [12] Zhenyu Wu et al. “Embodied task planning with large language models”. In: *arXiv preprint arXiv:2307.01848* (2023).
- [13] Zihao Wang et al. “Describe, explain, plan and select: Interactive planning with LLMs enables open-world multi-task agents”. In: *Proceedings of the 37th Advances in Neural Information Processing Systems (NeurIPS)*. 2023.
- [14] Ishika Singh et al. “ProgPrompt: Generating situated robot task plans using large language models”. In: *Proceedings of the 40th International Conference on Robotics and Automation (ICRA)*. IEEE, 2023, pp. 11523–11530.
- [15] Anthony Brohan et al. “RT-2: Vision-Language-Action Models Transfer Web Knowledge to Robotic Control”. In: *Proceedings of the 7th Conference on Robot Learning (CoRL)*. PMLR, 2023, pp. 2165–2183.
- [16] Danny Driess et al. “PaLM-E: an embodied multimodal language model”. In: *Proceedings of the 40th International Conference on Machine Learning (ICML)*. 2023, pp. 8469–8488.

- [17] Kaylee Burns et al. “GenCHiP: generating robot policy code for high-precision and contact-rich manipulation tasks”. In: *Proceedings of the 37th IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 2024, pp. 9596–9603.
- [18] Yin Li et al. “MCCoder: streamlining motion control with LLM-assisted code generation and rigorous verification”. In: *arXiv preprint arXiv:2410.15154* (2024).
- [19] Xiao Wang et al. “AutoMisty: a multi-agent LLM framework for automated code generation in the Misty social robot”. In: *arXiv preprint arXiv:2503.06791* (2025).
- [20] Leonardo De Moura and Nikolaj Bjørner. “Z3: An efficient smt solver”. In: *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer, 2008, pp. 337–340.
- [21] Xiang Yin and John Knight. “Formal verification of large software systems”. In: *Proceedings of the Second NASA Formal Methods Symposium (NFM)*. 2010.
- [22] Daniel Kroening, Peter Schrammel, and Michael Tautschnig. “CBMC – C Bounded Model Checker”. In: *Tools and Algorithms for the Construction and Analysis of Systems*. Vol. 8413. Springer, 2014, pp. 389–391.
- [23] Leonardo de Moura et al. “The Lean theorem prover (system description)”. In: *Automated Deduction – CADE-25*. Vol. 9195. Springer, 2015, pp. 378–388.
- [24] Jasmin Christian Blanchette and Lawrence C. Paulson. “Three Years of Experience with Sledgehammer, a Practical Link between Automatic and Interactive Theorem Provers”. In: *Proceedings of the 8th International Workshop on the Implementation of Logics (IWIL)*. 2010, pp. 1–11.
- [25] Sina Shamshiri et al. “Do automatically generated unit tests find real faults? An empirical study of effectiveness and challenges”. In: *Proceedings of the 30th IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 2015, pp. 201–211.
- [26] Yun Tang et al. “Systematic testing of autonomous driving systems using map topology-based scenario classification”. In: *Proceedings of the 36th IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 2021, pp. 1342–1346.
- [27] James A. Whittaker. “What is software testing? And why is it so hard?”. In: *IEEE Software* 17.1 (2000), pp. 70–79.
- [28] Jingkan Yang et al. “Octopus: embodied vision-language programmer from environmental feedback”. In: *Proceedings of the European Conference on Computer Vision (ECCV)*. Springer, 2024, pp. 20–38.
- [29] Greta Dolcetti et al. “Helping LLMs improve code generation using feedback from testing and static analysis”. In: *arXiv preprint arXiv:2412.14841* (2024).
- [30] Jundong Xu et al. “Faithful logical reasoning via symbolic chain-of-thought”. In: *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*. 2024.
- [31] Shraddha Barke et al. “HYSYNTH: context-free LLM approximation for guiding program synthesis”. In: *Proceedings of the 38th Annual Conference on Neural Information Processing Systems (NeurIPS)*. 2024.
- [32] Sumit Kumar Jha et al. “Counterexample Guided Inductive Synthesis Using Large Language Models and Satisfiability Solving”. In: *Proceedings of the IEEE Military Communications Conference (MILCOM)*. 2023, pp. 944–949.
- [33] Theo X. Olausson et al. “LINC: A Neurosymbolic Approach for Logical Reasoning by Combining Language Models with First-Order Logic Provers”. In: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, 2023, pp. 5153–5176.
- [34] Liangming Pan et al. “LOGIC-LM: Empowering Large Language Models with Symbolic Solvers for Faithful Logical Reasoning”. In: *Findings of the Association for Computational Linguistics (EMNLP)*. Association for Computational Linguistics, 2023, pp. 3806–3824.
- [35] Debargha Ganguly et al. “Proof of thought: neurosymbolic program synthesis allows robust and interpretable reasoning”. In: *Workshop on System 2 Reasoning at Scale, NeurIPS*. 2024.
- [36] Xin Quan et al. “Verification and refinement of natural language explanations through LLM-symbolic theorem proving”. In: *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, 2024, pp. 2933–2958.

- [37] Xinrui Lin et al. “CLMASP: Coupling Large Language Models with Answer Set Programming for Robotic Task Planning”. In: *arXiv preprint arXiv:2406.03367* (2024).
- [38] Alessio Capitanelli and Fulvio Mastrogiovanni. “A framework for neurosymbolic robot action planning using large language models”. In: *Frontiers in Neurorobotics* 18 (2024), p. 1342786.
- [39] Vinh Luong et al. “DANA: domain-aware neurosymbolic agents for consistency and accuracy”. In: *arXiv preprint arXiv:2410.02823* (2024).
- [40] Kaizhi Zheng et al. “JARVIS: a neuro-symbolic commonsense reasoning framework for conversational embodied agents”. In: *arXiv preprint arXiv:2208.13266* (2022).
- [41] Yujie Lu et al. “Neuro-symbolic procedural planning with commonsense prompting”. In: *Proceedings of the 11th International Conference on Learning Representations (ICLR)*. 2023.
- [42] Richard S. Sutton and Andrew G. Barto. *Reinforcement learning: an introduction*. MIT Press, 2018.
- [43] Lingfeng Sun et al. “Interactive planning using large language models for partially observable robotic tasks”. In: *Proceedings of the 41st International Conference on Robotics and Automation (ICRA)*. IEEE, 2024, pp. 14054–14061.
- [44] Jason Wei et al. “Chain-of-thought prompting elicits reasoning in large language models”. In: *Advances in neural information processing systems* 35 (2022), pp. 24824–24837.
- [45] Wonje Choi et al. “NeSyC: A Neuro-symbolic Continual Learner For Complex Embodied Tasks in Open Domains”. In: *Proceedings of the 13th International Conference on Learning Representations (ICLR)*. 2025.
- [46] Mohit Shridhar et al. “ALFRED: A Benchmark for Interpreting Grounded Instructions for Everyday Tasks”. In: *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2020.
- [47] So Yeon Min et al. “FILM: Following Instructions in Language with Modular Methods”. In: *Proceedings of the 10th International Conference on Learning Representations (ICLR)*. 2022.
- [48] Aishwarya Padmakumar et al. “TEACH: Task-Driven Embodied Agents That Chat”. In: *Proceedings of the 36th AAAI Conference on Artificial Intelligence*. Vol. 36. 2. 2022, pp. 2017–2025.
- [49] Haoze Wu, Clark Barrett, and Nina Narodytska. “LEMUR: Integrating Large Language Models in Automated Program Verification”. In: *Proceedings of the 12th International Conference on Learning Representations (ICLR)*. 2024.
- [50] Qingyun Wu et al. “AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation”. In: *arXiv preprint arXiv:2308.08155* (2023).
- [51] OpenAI. “GPT-4o system card”. In: *arXiv preprint arXiv:2410.21276* (2024).
- [52] AI@Meta Llama Team. “The Llama 3 Herd of Models”. In: *arXiv preprint arXiv:2407.21783* (2024).
- [53] Malte Helmert. “The Fast Downward Planning System”. In: *Journal of Artificial Intelligence Research* 26 (2006), pp. 191–246.
- [54] Håkan LS Younes and Michael L Littman. “PPDDL1. 0: An extension to PDDL for expressing planning domains with probabilistic effects”. In: *Techn. Rep. CMU-CS-04-162* 2 (2004), p. 99.
- [55] David Coleman et al. “Reducing the barrier to entry of complex robotic software: a moveit! case study”. In: *arXiv preprint arXiv:1404.3785* (2014).
- [56] Lirui Wang et al. “GenSim: Generating Robotic Simulation Tasks via Large Language Models”. In: *Proceedings of the 12th International Conference on Learning Representations (ICLR)*. 2024.

NeurIPS Paper Checklist¹

1. Claims²

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?³

Answer: [Yes]⁴

Justification: The main claims in the "Abstract" and "Introduction" clearly describe the proposed neuro-symbolic framework for generating grounded policy code and improving reliability in partially observable environments. These claims are supported by the results presented in the "Main Result" section.⁵

Guidelines:⁶

- The answer NA means that the abstract and introduction do not include the claims made in the paper.⁷
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations⁸

Question: Does the paper discuss the limitations of the work performed by the authors?⁹

Answer: [Yes]¹⁰

Justification: This paper discusses a key limitation of the framework in the "Conclusion", noting that the logic-based confidence score is binary and does not capture nuanced feasibility in real-world environments.¹¹

Guidelines:¹²

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.¹³
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs¹⁴

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof? 1

Answer: [NA] 2

Justification: [NA] 3

Guidelines: 4

- The answer NA means that the paper does not include theoretical results. 5
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility 6

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)? 7

Answer: [Yes] 8

Justification: This paper provides implementation details in the “Appendix”, and the source code is included in the supplementary material. 9

Guidelines: 10

- The answer NA means that the paper does not include experiments. 11
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code 1

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material? 2

Answer: [Yes] 3

Justification: We plan to release our project as open-source. 4

Guidelines: 5

- The answer NA means that paper does not include experiments requiring code. 6
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details 7

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results? 8

Answer: [Yes] 9

Justification: The experimental setup is described in “Experiment”, with additional details provided in the “Appendix”. 10

Guidelines: 11

- The answer NA means that the paper does not include experiments. 12
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance 13

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments? 14

Answer: [Yes] 15

Justification: All reported results include standard deviation error bars across trials, as shown in tables in “Main Result” and “Ablation”. 16

Guidelines: 17

- The answer NA means that the paper does not include experiments. 18
- The authors should answer “Yes” if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions). 1
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources 2

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments? 3

Answer: [Yes] 4

Justification: We specify the compute resources used for experiments in the “Appendix”. 5

Guidelines: 6

- The answer NA means that the paper does not include experiments. 7
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

9. Code of ethics 8

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>? 9

Answer: [Yes] 10

Justification: We have reviewed the NeurIPS Code of Ethics and confirm that our research complies with its guidelines. 11

Guidelines: 12

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics. 13
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts 14

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed? 15

Answer: [Yes] 16

Justification: This paper discusses potential risks associated with LLM-based robot control, such as unsafe behavior with hazardous tools, and mentions safeguards like safety checks and transparent execution control in the “Conclusion”. 17

Guidelines: 18

- The answer NA means that there is no societal impact of the work performed. 1
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards 2

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)? 3

Answer: [NA] 4

Justification: [NA] 5

Guidelines: 6

- The answer NA means that the paper poses no such risks. 7
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets 8

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected? 9

Answer: [Yes] 10

Justification: All third-party assets used in the paper, including RLBench, LLMs, and planners, are properly cited with references to their original papers. We ensure that their usage complies with publicly available licenses and terms of use. 11

Guidelines: 12

- The answer NA means that the paper does not use existing assets. 13
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided. 1
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets 2

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets? 3

Answer: [NA] 4

Justification: [NA] 5

Guidelines: 6

- The answer NA means that the paper does not release new assets. 7
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects 8

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)? 9

Answer: [NA] 10

Justification: [NA] 11

Guidelines: 12

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 13
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects 14

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained? 15

Answer: [NA] 16

Justification: [NA] 17

Guidelines: 18

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects. 19

- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper. 1
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage 2

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required. 3

Answer: [Yes] 4

Justification: LMs are used as core components of the proposed method, with detailed usage described in the “NeSyRo Framework”. 5

Guidelines: 6

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components. 7
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.

A Details of the NESYRO Framework¹

NESYRO enables exploration through a recursive composition of Neuro-symbolic Code Verification² and Neuro-symbolic Code Validation. A policy code produced by the LLM is first subjected to explicit symbolic verification, which statically checks logical consistency against the task specification. The verified policy code then enters interactive validation, where each skill is evaluated in neuro-symbolic manner; if an unmet precondition is detected, the system synthesizes exploratory safe probe policy code to gather the missing observations. Every probe is fed back through the same verification and validation (V&V) cycle, producing a policy tree whose nodes are recursively grounded until all skills achieve a satisfactory Neuro-symbolic Confidence Score. This recursive V&V framework guarantees that the policy code is both executable and environmentally grounded, even under dynamic, partially observable conditions.

A.1 NESYRO Algorithm³

We provide the full pseudocode for our neuro-symbolic task execution pipeline in Algorithm 1 and Algorithm 2. Below, we briefly describe the functional roles of each procedure and its interaction within the recursive planning framework.

Algorithm 1 Task Execution Pipeline⁵

Agent:⁶

env — environment interface $\mathcal{D}$ — domain knowledge
 $\mathcal{E}_{\text{demo}}$ — demonstration set ϵ — confidence threshold
 a — primitive action $o_{\leq t}$ — observation history up to t
 $\mathcal{A}$ — action set

Returns:⁷

τ — executed trajectory π_{main} — policy code

procedure RUNTASK(env, $\mathcal{D}$, $\mathcal{E}_{\text{demo}}$, ϵ)

$(o_{\leq 0}, g) \leftarrow \text{env.reset}()$
 $\tau \leftarrow ()$
 $(\pi_{\text{main}}, o_{\leq t}, \text{env}, \tau) \leftarrow \text{NESYRO}(g, o_{\leq 0}, \mathcal{D}, \mathcal{E}_{\text{demo}}, \epsilon, 0, \text{env}, \tau)$
 $(\tau_{\text{exe}}, o_{\leq t+1}, \text{env}) \leftarrow \text{EXE}(\pi_{\text{main}}, o_{\leq t}, \text{env})$
 $\tau \leftarrow \tau \cup \tau_{\text{exe}}$
return $(\tau, o_{\leq t+1}, \pi_{\text{main}})$

⁸

▷ Initialize empty trajectory

end procedure

procedure EXE(π , $o_{\leq t}$, env)

$\tau_{\text{exe}} \leftarrow []$
for each f **in** π **do**
 $\mathcal{A} \leftarrow \text{EXPANDSKILL}(f, o_{\leq t})$ ⁹
 for each a **in** $\mathcal{A}$ **do**
 $o_{\text{next}} \leftarrow \text{env.step}(a)$ ¹⁰
 $\tau_{\text{exe}} \leftarrow \tau_{\text{exe}} \cup (a, o_{\text{next}})$
 $o_{\leq t} \leftarrow o_{\leq t} \cup o_{\text{next}}$

end for

end for

return $(\tau_{\text{exe}}, o_{\leq t}, \text{env})$

end procedure

RunTask and execution loop. The RUNTASK procedure (Algorithm 1) initializes the environment¹¹ and launches the neuro-symbolic reasoning process. It first resets the environment to obtain the initial observation $o_{\leq 0}$ and instruction g , and initializes an empty trajectory τ . The main grounding routine NESYRO is then called to synthesize a grounded policy code π_{main} based on the instruction and current context. Once obtained, the policy is executed via EXE, which expands symbolic skills into primitive actions and steps through the environment. The complete trajectory τ and updated observation history $o_{\leq t}$ are returned. The EXE procedure handles the execution of symbolic skills. For

each skill f in π , it calls EXPANDSKILL to retrieve the sequence of corresponding low-level actions. These are executed sequentially in the environment, and the resulting observations are appended to both the trajectory and the observation history.

Algorithm 2 Recursive Neuro-symbolic Verification & Validation

```

procedure NESYRO( $g, o_{\leq t}, \mathcal{D}, \mathcal{E}_{\text{demo}}, \epsilon, k, \text{env}, \tau$ )
  ( $\mathcal{T}_{\text{spec}}, \pi_{\text{main}}$ )  $\leftarrow$  NEURO_SYMBOLIC_VERIFICATION( $g, o_{\leq t}, \mathcal{D}, k$ )
  ( $\pi_{\text{main}}, o_{\leq t}, \text{env}, \tau$ )  $\leftarrow$ 
    NEURO_SYMBOLIC_VALIDATION( $g, o_{\leq t}, \mathcal{D}, \mathcal{E}_{\text{demo}}, \pi_{\text{main}}, \epsilon, k, \text{env}, \tau$ )
  return ( $\pi_{\text{main}}, o_{\leq t}, \text{env}, \tau$ )  $\triangleright$  grounding policy code  $\pi_{\text{main}}$ 
end procedure

procedure NEURO_SYMBOLIC_VERIFICATION( $g, o_{\leq t}, \mathcal{D}, k$ )
  ( $\mathcal{T}_{\text{spec}}, \pi_{\text{main}}$ )  $\leftarrow$   $\Phi_{\text{veri}}(o_{\leq t}, g, l_{\text{cot}}, \mathcal{D}, k)$ 
  while  $\Psi_{\text{veri}}(\mathcal{T}_{\text{spec}}, \pi_{\text{main}}) = \text{fail}$  do
     $\mathcal{F}_{\text{veri}} \leftarrow \Psi_{\text{veri}}(\mathcal{T}_{\text{spec}}, \pi_{\text{main}})$ 
    ( $\mathcal{T}_{\text{spec}}, \pi_{\text{main}}$ )  $\leftarrow$   $\Phi_{\text{veri}}(o_{\leq t}, g, l_{\text{cot}}, \mathcal{D}, \pi_{\text{main}}, \mathcal{F}_{\text{veri}}, k)$ 
  end while
  return ( $\mathcal{T}_{\text{spec}}, \pi_{\text{main}}$ )
end procedure

procedure NEURO_SYMBOLIC_VALIDATION( $g, o_{\leq t}, \mathcal{D}, \mathcal{E}_{\text{demo}}, \pi_{\text{main}}, \epsilon, k, \text{env}, \tau$ )
   $n \leftarrow k$ 
  while  $n < |\pi_{\text{main}}|$  do
     $f_n \leftarrow \pi_{\text{main}}[n]$ 
     $\text{CSC} \leftarrow \Phi_{\text{vali}}(\mathcal{D}, \mathcal{E}_{\text{demo}}, o_{\leq t}, g, f_n)$ 
     $\text{LC} \leftarrow \Psi_{\text{vali}}(\mathcal{D}, o_{\leq t}, g, f_n)$ 
     $\text{NeSyConf} \leftarrow \text{CSC} \times \text{LC}$ 
    if  $\text{NeSyConf} < \epsilon$  then
       $g_{\text{probe}} \leftarrow \text{MAKEPROBEGOAL}(f_n, \mathcal{F}_{\text{csc}}, \mathcal{F}_{\text{lc}})$ 
      ( $\pi_{\text{probe}}, o_{\leq t}, \text{env}, \tau$ )  $\leftarrow$  NESYRO( $g_{\text{probe}}, o_{\leq t}, \mathcal{D}, \mathcal{E}_{\text{demo}}, \epsilon, 0, \text{env}, \tau$ )  $\triangleright$  recursive
      ( $\tau_{\text{exe}}, o_{\leq t+1}, \text{env}$ )  $\leftarrow$  EXE( $\pi_{\text{probe}}, o_{\leq t}, \text{env}$ )
       $\tau \leftarrow \tau \cup \tau_{\text{exe}}$ 
      ( $\mathcal{T}_{\text{spec}}, \pi_{\text{main}}$ )  $\leftarrow$  NEURO_SYMBOLIC_VERIFICATION( $g, o_{\leq t+1}, \mathcal{D}, n$ )
    else
       $n \leftarrow n + 1$ 
    end if
  end while
  return ( $\pi_{\text{main}}, o_{\leq t+\alpha}, \text{env}, \tau$ )  $\triangleright$   $\alpha$ : number of recursive  $\pi_{\text{probe}}$  executed during validation.
end procedure

```

Recursive neuro-symbolic reasoning. Algorithm 2 outlines the recursive grounding logic of NESYRO. The NESYRO procedure first invokes NEURO_SYMBOLIC_VERIFICATION to obtain a symbolic task specification $\mathcal{T}_{\text{spec}}$ and initial policy code π_{main} . Logical correctness is ensured through iterative verification using Φ_{veri} and Ψ_{veri} , which checks whether π_{main} satisfies $\mathcal{T}_{\text{spec}}$. After verification, the policy is passed to NEURO_SYMBOLIC_VALIDATION for skill-wise confidence assessment. For each skill f_n , the framework computes neuro-symbolic confidence score (NeSyConf). If NeSyConf falls below threshold ϵ , a probing goal g_{probe} is generated and recursively passed into NESYRO.

To construct this probing goal g_{probe} , the MAKEPROBEGOAL function synthesizes a new instruction that addresses the failure feedback. Specifically, it leverages NeSyConf feedback ($\mathcal{F}_{\text{csc}}, \mathcal{F}_{\text{lc}}$) to identify missing observations. This recursive routine allows a skill that fails to exceed the confidence threshold ϵ to be refined and validated using updated observations gathered from safe probe executions. Once all skills pass validation, the final grounded policy and accumulated trajectory are returned.

B Environment Settings¹

B.1 RLBench²

We use RLBench [8] as the simulation environment for our experiments. RLBench offers a wide range of tabletop manipulation tasks and provides realistic simulations of both robot control and visual observations. All experiments are performed using a 7-DoF Franka Emika Panda robotic arm, which is supported natively by RLBench. The environment is particularly suitable for evaluating planning and interaction under partial observability, as it supports randomized object configurations and sensor data, including RGB, depth, and segmentation masks. Its compatibility with Python also allows straightforward integration with our code generation and execution framework.




Figure 6: Example scenes illustrating the "open drawer" task in the RLBench. The top row shows the overhead view, and the bottom row shows the front view.

Object configuration. Each episode initializes a workspace containing seven unique objects placed on a table. The objects include two tomatoes, a piece of trash, a bin, a three-level drawer, a desk lamp, and a switch for the desk lamp. The position of each object is randomized in every episode, introducing perceptual variability and scene diversity across tasks. Figure 6 shows an example scene from the RLBench environment used in our experiments.

Task composition. We categorize the tasks into four types to enable structured evaluation, as summarized in Table 7. Each task type contains multiple language instructions, with associated probe targets indicating the source of uncertainty that must be resolved during execution.

Table 7: Task types, example of instructions, and associated probe targets in RLBench.

Task Type	Example Instructions	Probe Target
Object Relocation	Move two tomatoes onto plate	Tomato identity
	Put the trash into bin	Trash identity
Object Interaction	Open a drawer	Drawer locked/unlocked state
	Open two drawers	Drawer locked/unlocked state
Auxiliary Manipulation	Move two tomatoes onto plate in dark room	Missing visual observation (requires light activation)
	Open drawer in dark room	Missing visual observation (requires light activation)
Long-Horizon Tasks	Move a die into the drawer	Die identity, Drawer locked/unlocked state
	Move dice into the drawer	Dice identity, Drawer locked/unlocked state

B.2 Real-world¹⁰

Environment setup. We conducted our real-world experiments using a 7-DoF Franka Emika Research 3 robotic arm mounted on a tabletop workspace. An Intel RealSense D435 RGB-D camera was positioned above the table to provide top-down RGB and depth information. This input was processed by an object detection module to identify the categories and bounding boxes of task-relevant objects. Depth measurements were used to compute 3D coordinates, which were then transformed into the robot’s coordinate frame. This setup enabled accurate object localization and real-time observation grounding, providing the necessary perception for reliable execution.




Figure 7: Example scenes from the real-world environment used in our experiments.

Object configuration. The real-world environment contains ten unique objects arranged on a tabletop workspace. These include two dice, two pieces of trash, a bin, a three-level drawer, a small cardboard box, a paper cup filled with liquid, a roll of tape, and a light switch. The initial positions of all objects are randomized for each trial, introducing diverse spatial configurations and observation conditions across task instances. This variability supports evaluation under partial observability and enables direct comparison with the RLBench-based simulation setup. Figure 7 shows a representative setup of the real-world environment used in our experiments.

Task composition. We maintain the RLBench task categorization in the real-world setup to ensure consistency and enable direct comparison. Each of the four task types corresponds to a distinct source of uncertainty and is associated with multiple language instructions and probe targets, as summarized in Table 8.

Table 8: Task types, example of instructions, and associated probe targets in real-world.

Task Type	Example Instructions	Probe Target
Object Relocation	Place a die into drawer Move dice into drawer	Die identity Dice identity
Object Interaction	Open a drawer Open two drawers	Drawer locked/unlocked state Drawer locked/unlocked state
Auxiliary Manipulation	Place a die into drawer in dark room Open drawer in dark room	Missing visual observation (requires light activation) Missing visual observation (requires light activation)
Long-Horizon Tasks	Place a die into drawer Move a die into drawer	Die identity, Drawer locked/unlocked state Die identity, Drawer locked/unlocked, empty/occupied state

Low-Level Control. For motion planning and control in the real-world environment, we employed MoveIt [55], an open-source motion planning framework widely used for robotic manipulation. Once the target object positions were obtained from the perception pipeline, we invoked parameterized skill primitives such as `pick`, `place`, and `open`, which are designed to operate over arbitrary object poses. Each skill was instantiated using the transformed 3D coordinates of the corresponding object and passed to the planner as goal constraints. Trajectory optimization was handled by MoveIt’s built-in planners, which computed collision-free joint-space paths that respect the robot’s kinematic limits and workspace constraints. The resulting trajectories were executed using the robot’s internal controller through ROS. Although continuous force control was not used for the gripper, we implemented discrete grasping strategies based on object geometry and semantic role (e.g., trash, dice). This ensured consistent and safe execution across a variety of physical configurations.

Unlike the RLBench setting, where code execution is simulated through parameterized low-level APIs, our real-world system closes the loop by grounding skill calls with real sensor observations and executing planned trajectories on physical hardware. This setup allows us to evaluate the reliability of the proposed planning framework under real-world uncertainties.

C Experiment Details¹

C.1 Compute Resources²

Most experiments were conducted on a local machine with an Intel(R) Core(TM) i7-9700KF CPU³ and an NVIDIA GeForce RTX 4080 GPU (16GB VRAM). Each task instance used a single GPU, and RLBench simulation was executed with up to 32GB of system memory. Symbolic verification and PDDL planning were run on the CPU. For experiments using the larger language models listed in Table 6 in main paper, such as Llama-3.1-8B and Qwen3-30B-A3B, we used a cloud-based CUDA cluster with GPUs equipped with approximately 82GB of VRAM. All OpenAI models, including GPT-4o and GPT-4.1, were accessed via the OpenAI API.

C.2 NeSyRo Implementation⁴

LLM usage overview. Our framework utilizes LLM as core reasoning engines in three tightly⁵ integrated components of the code generation and validation pipeline:

- **Code Generation (Φ_{veri}):** Given the instruction g and the current observation history $o_{\leq t}$,⁶ the verification LLM Φ_{veri} performs chain-of-thought reasoning to produce an intermediate symbolic task specification $\mathcal{T}_{\text{spec}}$ and corresponding Python policy code π_{main} . If the code fails symbolic verification via Ψ_{veri} , structured feedback $\mathcal{F}_{\text{veri}}$ is returned and used by the LLM to iteratively revise the code. Only the unvalidated portion of π_{main} is regenerated at each step, preserving previously verified components.
- **CSC Computation (Φ_{vali}):** For each skill f_n in the primary policy π_{main} , the validation⁷ LLM Φ_{vali} computes CSC_{f_n} that estimates the likelihood of successful execution under the current observation $o_{\leq t}$ and instruction g . The LLM is prompted with the code for f_n , domain knowledge $\mathcal{D}$, retrieved single-skill demonstrations $\mathcal{E}_{\text{demo}}$, and task context. Token-level probabilities are aggregated and transformed into a negative log-likelihood score. This value is normalized to produce a scalar confidence score CSC_{f_n} used for validation. Before normalization, CSC_{f_n} ranges over $[0, \infty)$; after normalization, it is scaled to the interval $[0, 1]$.
- **CSC Feedback Generation ($\mathcal{F}_{\text{csc}}$):** If $\text{NeSyConf}_{f_n} < \epsilon$, the skill f_n is considered to⁸ require a safe probe. In such cases, CSC feedback $\mathcal{F}_{\text{csc}}$ is constructed based on f_n , the failure context, current observation $o_{\leq t}$, instruction g , and single-skill demonstrations $\mathcal{E}_{\text{demo}}$. This feedback is then used to prompt the LLM to generate a safe probe policy code π_{probe} . The resulting policy code is recursively verified and validated through the NeSyRo pipeline before execution.

Example of prompt. Below are the representative prompts used in each stage of our framework:⁹ generating executable robot code, computing CSC for each skill, and generating feedback when the NeSyConf falls below a threshold.

Code Generator Prompt¹⁰

Role: You are an AI assistant tasked with generating executable Python code that controls a¹¹ robot in a simulated environment.

Task: Complete the executable code using the provided inputs and by implementing any¹² missing skills. The goal is to ensure the robot can achieve the specified objective by executing a sequence of actions (plan).

Input Details:¹³

- **Domain PDDL:** Describes the available actions and predicates in the environment.¹⁴ It includes information about action parameters, preconditions, and effects. This provides the symbolic action space for the planner.
Provided domain PDDL: `{{domain_pddl}}`
- **Observation (Initial State Description):** Represents the initial state of the environment in PDDL format, including locations of objects, robot position, and other relevant state descriptions.
Provided observation: `{{observation}}`

- **Goal (Natural Language Description):** The goal specifies what the robot must accomplish in plain language.
Provided goal: `{{goal}}`
- **Specification (Code Generation Guidelines):** Provides strict rules and constraints that the generated code must follow.
Provided specification: `{{spec}}`
- **Skill Code (Python Implementations of Actions):** A set of predefined Python functions that implement low-level skills (e.g., move, pick, place).
Provided skill code: `{{skill_code}}`
- **Executable Code Skeleton:** A partially completed Python file containing environment setup and control flow scaffolding.
Provided skeleton: `{{skeleton_code}}`
- **Available Skill Names:** A list of all valid skill function names that may be called in the generated code.
Provided skills: `{{available_skills}}`
- **Object List:** A list of object names that exist in the environment.
Provided object list: `{{object_list_position}}`
- **Feedback:** Corrections or issues identified from the previous code generation.
Feedback: `{{feedback}}`
Previous code: `{{prev_code}}`
- **Exploration Knowledge:** Optional knowledge to infer or explore missing observations.
Exploration knowledge: `{{exploration_knowledge}}`
- **Frozen Code:** Code that must remain unchanged. New code must follow this segment.
[Frozen Code Start] `{{frozen_code_part}}` [Frozen Code End]

Implementation Requirements: 2

- Use only the predefined skills (e.g., move, pick, place) from `skill_code`; do not define new functions.
- Complete the provided skeleton by inserting plan logic that achieves the specified goal.
- Preserve all existing imports and [Frozen Code] segments.
- Output should be plain text only — do not use code blocks.
- Handle errors gracefully during skill execution (e.g., invalid arguments or missing objects).

- Use the following external modules as provided: 4

- `env`: Environment setup and shutdown.
- `skill_code`: Contains all callable action implementations.
- `video`: Tools for simulation recording.
- `object_positions`: For retrieving object location information.

CSC Computation Prompt 1

Role: You are an AI assistant responsible for estimating the likelihood that a specific robotic skill will succeed in the current environment. Your evaluation will be used to compute a token-level log-probability-based common sense confidence, rather than to generate output. 2

Input Details:

- **Instruction:** 3
{instruction}
- **Demonstrations:**
Demo 1: {demo_1}
Demo 2: {demo_2}
...
- **Observation:**
{observation}
- **Object List:**
{object_list}
- **Skill Code:**
{skill_code}

CSC Feedback Prompt 4

Role: You are an AI assistant responsible for analyzing why a robotic skill code is likely to fail and for generating feedback to guide its refinement. 5

Task: Based on the given context and the Neuro-Symbolic Confidence Score (NeSyConf) 6 being below threshold, provide structured feedback to help revise or improve the given skill code.

Input Details: 7

- **Demonstrations:** 8
Demo 1: {demo_1}
Demo 2: {demo_2}
...
- **Observation:**
{observation}
- **Skill Code:**
{skill_code}
- **Confidence Score**
Current Confidence Score: {NeSyConf}, Threshold: {threshold}
- **Instruction:**
{instruction}
- **Object List:**
{object_list}

Format: 9

- 1. Problem Identification 10
- 2. Justification (Why is it a problem?)
- 3. Proposed Solutions (High-level ideas)
- 4. Additional Notes (Optional)

Instruction to Model: Focus on issues such as force calibration, missing objects, logical 11 flaws, or unsafe execution. For example: “An object declared in the code is not in the actual object list.” Your feedback will guide the regeneration of this skill step. 12

Hyperparameter setting. The only hyperparameter in our framework is the confidence threshold ϵ used during neuro-symbolic validation. For each skill, we perform five safe exploration probes

under varied initial conditions to estimate its execution confidence. To determine a suitable value of ϵ 1 for a given environment, we exclude outlier trials in which the probe failed due to non-informative reasons, which could otherwise deflate confidence estimates. This ensures that ϵ reflects a realistic and actionable lower bound of confidence for successfully grounded skills.




Figure 8: Confidence scores over five safe probes for each of 7 skills in a long-horizon task. 3

Figure 8 illustrates this process using one of the long-horizon tasks in the RLBench environment. 4 Each of the 7 skills was probed five times, and confidence scores were recorded before and after each probe. The figure shows that while confidence increases with repeated probing, it typically saturates after a few trials, indicating convergence. Based on this observation, we compute the final confidence distribution by averaging only those probe outcomes that resulted in a successful skill grounding. We then set ϵ to the lower quartile of this filtered distribution, ensuring a conservative yet robust threshold that filters out unreliable executions while accepting skills with moderately confident grounding.

Single-skill demonstrations format. To support symbolic validation and LLM-based reasoning, 5 we constructed a general-purpose demonstration library consisting of synthetic examples. These demonstrations were generated entirely via a large language model (GPT-4o) [56], given domain-level PDDL definitions and representative symbolic contexts, without requiring environment-specific execution or human annotation. Each example encodes typical task-relevant transitions across common household activities (e.g., opening a drawer, placing an object, turning on a light), and captures both successful and failure cases under varied symbolic states. In total, we synthesized approximately 500 such demonstrations spanning over 15 diverse skill types. These examples are reused across all tasks to provide reusable prior knowledge for CSC computation and to guide safe probing decisions when symbolic grounding confidence is low.

Demonstration Structure 6

Each demonstration consists of a sequence of transitions. Each transition is represented as a 7 dictionary with the following fields:

- `initial_observation`: symbolic or structured observation before executing the 8 action.
- `action`: the skill function executed, such as a call to `pick(object)` or `open(drawer)`.
- `post_observation`: symbolic or structured observation after executing the action.
- `success`: boolean value indicating whether the action was successful.

C.3 Baselines Implementation 9

Code as Policies (CaP) [2] serves as the foundation of our framework and is implemented by invoking 10 the Code Generator Prompt with the frozen code length set to zero. Although a symbolic specification is also produced, this baseline does not include any verification process.

CaP w/ Lemur [49] extends CaP by performing verification over the generated specification. This process is conducted in the exact same manner as the Neuro-symbolic verification phase. 1

CaP w/ CodeSift [5] extends CaP by incorporating LLM-based verification and validation. In the verification stage, CodeSift performs static syntax checks using language-specific tools (pylint for Python, shellcheck for Bash) and prompts the LLM to summarize the code’s functionality. This summary is then used in the validation stage to assess semantic alignment with the original task instruction. The validation consists of multiple sub-steps: semantic similarity scoring, listing all functional mismatches, and determining whether the implementation is exact. If the code fails validation, the framework automatically generates refinement feedback and prompts the LLM to revise the code accordingly. The entire process is orchestrated via a modular pipeline that yields detailed diagnostic outputs and a refined version of the code when necessary. 2

LLM-Planner [11] follows the same initial procedure as CaP by generating code from the instruction using a Code Generator Prompt. During execution in the environment, if an action fails, the planner captures the current observation and provides it as additional context to the LLM. The previously executed portion of the code is marked as frozen, and a new code segment is generated to continue the task from the failure point. As in CaP, a symbolic specification is produced, but no verification or validation is performed throughout the process. 3

AutoGen [50] adopts the same iterative replanning strategy as LLM-Planner, where code is regenerated during execution upon failure by freezing the executed portion and providing the current observation as context. The key difference is that it uses a dedicated reasoning model, specifically o4-mini, to enhance task understanding and decision making. This improved reasoning enables more accurate replanning. As with LLM-Planner, no explicit verification or validation is performed. 4

D Real-world Experiment Details ¹

D.1 Figure 1 in Main Paper Details ²

In this section, we provide a detailed explanation of Figure 1 in main paper. The complete execution sequence depicted in Figure 1 including all safe probes is illustrated in Figure 9. To demonstrate the reliability of NESYRO in a real-world setting, we tasked an embodied agent with the instruction “Clean up the desk” and compared NESYRO against a naive code generation approach. As depicted in Figure 1, the naive approach failed in partially observable environments, leading it to execute an irreversible action by not recognizing that the middle drawer might be locked. In contrast, NESYRO addresses this uncertainty using a safe probe pipeline to acquire the missing observations. It initially plans a safe probe to determine whether the drawers are empty. However, through its recursive validation phase, it subsequently identifies the need to observe the locked status of the drawers. Consequently, Safe Probe 1, which checks the locked status of the drawers, is executed first, as shown in Figure 9. Upon its completion, the agent adds observations confirming that the middle drawer is locked and that the top and bottom drawers are unlocked. Subsequently, Safe Probe 2, which checks whether the drawers are empty, is executed and adds observations confirming that both the top and bottom drawers are empty. With these observations acquired, the policy code is now grounded. NESYRO proceeds to successfully execute the “Clean up the desk” instruction.

D.2 Figure 5 in Main Paper Details ⁴

This section provides a detailed explanation of Figure 5 in main paper. The complete execution sequence depicted in Figure 5 of main paper including all safe probes is illustrated in Figure 10. To further evaluate the robustness of NESYRO in a real-world setting, we implemented the instruction “Place one dice into a drawer in a dark room” which represents partially observable environments. This requires auxiliary manipulation such as turning on the light to restore visibility before executing the main task. The initial policy code π_{main} was ungrounded, missing observations regarding drawer visibility and lock status. To resolve this, NESYRO activates its safe probe pipeline and first generates Safe Probe 1 to turn on the light, enabling the agent to perceive object locations. However, even after Safe Probe 1, the NeSyConf for the skill `pick(middle_H, ...)` remains below a threshold. In response, NESYRO generates Safe Probe 2 to check the lock status of the drawers. This probe confirms that the top and bottom drawers are unlocked, while the middle drawer is locked. Based on these observations, the code is refined by replacing the initial skills that placed one dice into the middle drawer with new skills that place it into the top drawer. As a result, the skill `pick(top_H, ...)` becomes ready to execute. Subsequently, during the validation of `pick(dice, ...)`, the agent identifies the need to check whether the drawers are empty. Consequently, NESYRO generates Safe Probe 3 to check whether the drawers are empty. This probe detects trash inside the top drawer, leading to the insertion of additional code that removes it. These added skills also undergo the same validation phase in sequence. Once all skills are marked as ready to execute, the policy code is considered grounded. NESYRO proceeds to successfully execute the given instruction.

Instruction: Clean up the desk¹



Figure 9: Real-world execution sequence of the instruction “Clean up the desk”³

Instruction: Place one dice into a drawer in dark room



Figure 10: Real-world execution sequence of the instruction “Place one dice into a drawer in dark room”